



Presented by,
MySQL AB® & O'Reilly Media, Inc.

The Top 20 Design Tips For MySQL Enterprise Data Architects

Ronald Bradford
Principal
42SQL

April 2008



1. Know Your Technology Tools

- ❖ Generics are inefficient
- ❖ Product expertise in a different RDBMS is not enough
- ❖ You have chosen MySQL
 - ❖ Maximize it's strengths
 - ❖ Minimize it's weaknesses

Overview

- ❖ Table Structure
- ❖ SQL
- ❖ Indexes
- ❖ Enterprise Approaches

1. Know Your Technology Tools

- ❖ Maximize MySQL strengths
 - ❖ Scale out / HA Options
 - ❖ Different Storage Engines
 - ❖ Query Cache
- ❖ Minimize MySQL weaknesses
 - ❖ No Online Alter
 - ❖ Backup Strategies
 - ❖ Instrumentation

2. Know Your Disk Footprint

Disk = Memory = Performance

- ❖ Every single byte counts
- ❖ Average 25% - 30% saving on engagements
- ❖ Better 60% (200GB System)
- ❖ Best 78% (8GB per master with 12 masters)

Less disk accesses and more data in memory

3. Choose Your Numeric Data Type

- ❖ MySQL has 9 numeric data types
 - ❖ Oracle for example has only 1

3. Choose Your Numeric Data Type

- ❖ Integer: TINYINT, SMALLINT, MEDIUMINT, INT, BIGINT
- ❖ Floating Point: FLOAT, DOUBLE
- ❖ Fixed Point: DECIMAL
- ❖ Other: BIT, (ENUM maybe)

3. Choose Your Numeric Data Type

- ❖ Favorite signs of poor design
 - ❖ INT(1)
 - ❖ BIGINT AUTO_INCREMENT
 - ❖ no UNSIGNED used
 - ❖ DECIMAL(31,0)

3. Choose Your Numeric Data Type

- ❖ INT(1) - 1 does not mean 1 digit
 - ❖ (1) represents client output display format only
 - ❖ INT is 4 Bytes, TINYINT is 1 Byte
 - ❖ TINYINT UNSIGNED can store from 0 – 255
 - ❖ BIT is even better when values are 0 - 1

3. Choose Your Numeric Data Type

- ❖ BIGINT is not needed for AUTO_INCREMENT
- ❖ INT UNSIGNED stores 4.3 billion values
 - ❖ You should be partitioning when at billions of rows
- ❖ BIGINT is applicable for some columns
 - ❖ e.g. summation of values

3. Choose Your Numeric Data Type



Best
Practice

❖ Best Practice

- ❖ All integer columns UNSIGNED unless there is a reason otherwise
- ❖ Adds a level of data integrity for negative values

4. Other Data Type Efficiencies

❖ **TIMESTAMP v DATETIME**

- ❖ Suitable for EPOCH only values
- ❖ **TIMESTAMP** is 4 bytes
- ❖ **DATETIME** is 8 bytes
- ❖ FYI: **DATE** is 3 bytes, **TIME** is 3 bytes = 6 Bytes???

4. Other Data Type Efficiencies

- ❖ CHAR(n)
- ❖ Use VARCHAR(n) for variable values
- ❖ e.g. CHAR(128) when storing ~10 bytes

5. Application Data Type Efficiencies

❖ Using Codes or ENUM

- ❖ A description is a presentation layer function
- ❖ e.g. 'M', 'F' instead of 'Male', 'Female'
- ❖ e.g. 'A', 'I' instead of 'Active', 'Inactive'

❖ BINARY(16/20) v CHAR(32/40)

- ❖ MD5() or HASH() Hex value with twice the length

❖ INT UNSIGNED for IPv4 address

- ❖ VARCHAR(15) results in average 12 bytes v 4 bytes

6. NOT NULL

- ❖ Saves up to a byte per column per row of data
- ❖ Double benefit for indexed columns
- ❖ Don't use frameworks or tools
- ❖ NOT NULL DEFAULT " is bad design

**Always use NOT NULL unless
there is a reason why not**

Best
Practice

7. Know about character sets

- ❖ Default in MySQL 5 is UTF8
- ❖ Can be defined at database, schema, table or column level
- ❖ Only define columns that need UTF8
 - ❖ e.g. Codes, MD5 Value, web address
- ❖ MySQL internal buffers are fixed width
 - ❖ e.g. VARCHAR(255) utf8 is 765 bytes to store just 1 byte

8. When VARCHAR Is Bad

❖ VARCHAR(255)

- ❖ Poor Design - No understanding of underlying data
- ❖ Old Design - (4.x limitation, now 3-4 years old)
- ❖ Disk usage may be efficient
- ❖ MySQL internal memory usage is not

8. When VARCHAR is bad

```
CREATE TABLE `XXX` (  
  `orderHandle` varchar(255) NOT NULL default '',  
  `personName` varchar(255) default NULL,  
  `addressLines` varchar(255) default NULL,  
  `city` varchar(255) default NULL,  
  `state` varchar(255) default NULL,  
  `postalCode` varchar(255) default NULL,  
  `countryCode` varchar(255) default NULL,  
  `phone` varchar(255) default NULL,  
  `email` varchar(255) default NULL,  
  `shipMethod` varchar(255) default NULL,  
  `shipTo` varchar(255) default NULL,  
  `receiveByDate` date default NULL,  
  `currency` varchar(3) default NULL,  
  `price` varchar(255) default NULL,  
  `flags` int(11) default '0',  
  `lastUpdateTime` timestamp NOT NULL default CURRENT_TIMESTAMP on update CURRENT_TIMESTAMP,  
  `creationTime` timestamp NOT NULL default '0000-00-00 00:00:00',  
  PRIMARY KEY (`orderHandle`)  
) ENGINE=MyISAM DEFAULT CHARSET=utf8
```

9. Be Wary of TEXT/BLOB

- ❖ Using SELECT *
 - ❖ MySQL Internal Temporary table will force Temp Disk Table
- ❖ Internal storage (e.g. Innodb)
 - ❖ Stores first 768 bytes, then a separate 16k data page per row per TEXT/BLOB field

10. Know Every SQL Statement

- ❖ Developers don't write proper SQL statements
- ❖ SQL statements will directly affect your performance
- ❖ For Example
 - ❖ Repeating SQL statements for no benefit
 - ❖ 1000 very quick small unnecessary queries is worse than 1 slow query

10. Know Every SQL Statement

- ❖ Data collection options
- ❖ Incomplete Options
 - ❖ Slow Query Log
 - ❖ SHOW PROCESSLIST
 - ❖ Application level logging
- ❖ Impractical Options
 - ❖ General Log

10. Know Every SQL Statement

- ❖ Data collection options
 - ❖ MySQL Proxy
 - ❖ See histogram.lua
 - ❖ Firewall forwarding rules

11. Monitor Every SQL Statement

- ❖ Review Query Execution Plan (QEP)
 - ❖ EXPLAIN
- ❖ Time queries
- ❖ Row Count / Affected rows
- ❖ Result Set Size



Best
Practice

Review over time, things change

12. The Impact Of Indexes

❖ Good

- ❖ Dramatic performance improvements
- ❖ Improves memory usage
- ❖ Data Integrity

❖ Bad

- ❖ Slows performance for writes
- ❖ Wastes disk space for unused, duplicate or ineffective indexes
- ❖ In-effective usage of memory

13. Index Types For Design

- ❖ Concatenated Indexes
 - ❖ (col1, col2)
- ❖ Partial Indexes
 - ❖ (name(20))
- ❖ Covering Indexes
- ❖ Full Text Indexes
- ❖ No function based indexes

14. Minimizing internal MySQL processing

- ❖ Correctly design tables, indexes and SQL to eliminate
 - ❖ Using temporary table
 - ❖ Using filesort

15. Transactions

- ❖ Always design for transactions
- ❖ Always use transactions
- ❖ Use a transactional storage engine

16. Data Integrity is Key

- ❖ MySQL historically has been very lax
- ❖ Warnings (e.g. Truncations) are rarely every caught
- ❖ SQL_MODE=STRICT_ALL_TABLES
- ❖ Within Schema
 - ❖ NOT NULL
 - ❖ ENUM
 - ❖ UNSIGNED

17. Leverage The Query Cache

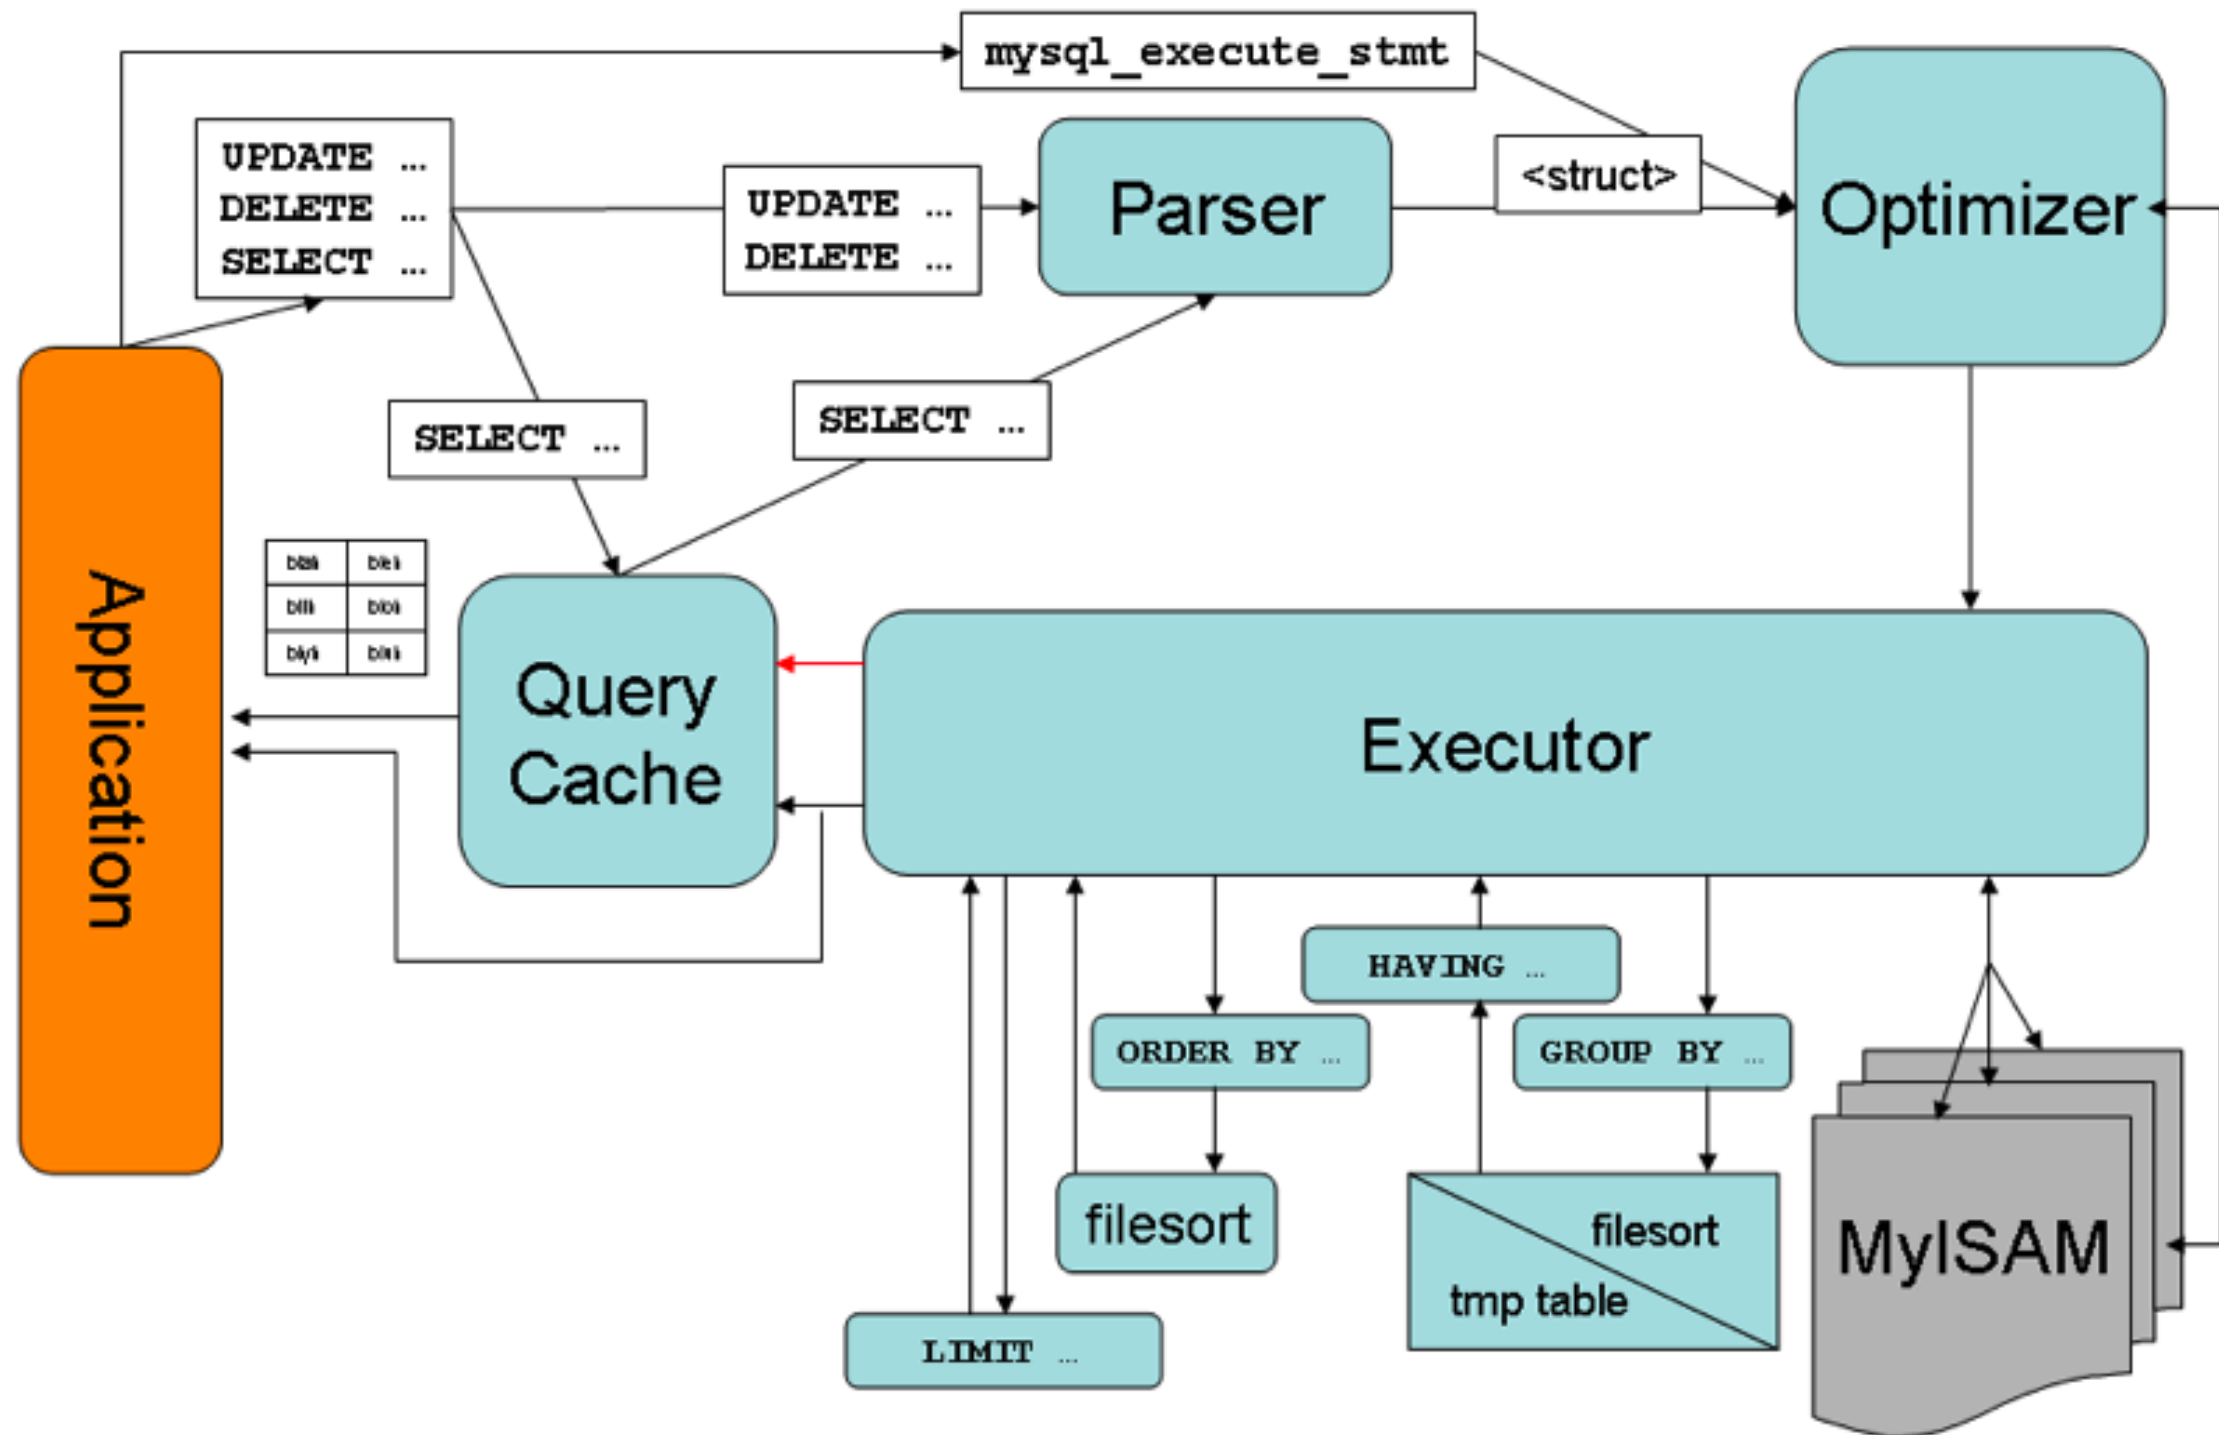
- ❖ Query Cache can be a great benefit
- ❖ Deterministic v Non Deterministic SQL



Best
Practice

**MySQL Query Cache is not the only type
of caching you should consider**

17. Leverage The Query Cache



17. Leverage The Query Cache

- ❖ SHOW PROFILE Path
- ❖ Simple SELECT
 - ❖ No Query Cache 17 steps
 - ❖ With Query Cache 5 steps
 - ❖ Does not perform parse
 - ❖ Does not perform optimize

The Top 20 Design Tips

Status	Duration	Source_function	Source_file	Source_line
(initialization)	0.000014	send_result_to_client	sql_cache.cc	1143
checking query cache for query	0.000042	open_tables	sql_base.cc	2652
Opening tables	0.000015	mysql_lock_tables	lock.cc	153
System lock	0.000009	mysql_lock_tables	lock.cc	163
Table lock	0.000034	mysql_select	sql_select.cc	2273
init	0.000041	optimize	sql_select.cc	765
optimizing	0.000008	optimize	sql_select.cc	924
statistics	0.000016	optimize	sql_select.cc	934
preparing	0.000012	exec	sql_select.cc	1594
executing	0.000008	exec	sql_select.cc	2114
Sending data	0.000163	mysql_select	sql_select.cc	2318
end	0.000021	mysql_execute_command	sql_parse.cc	5141
query end	0.000007	query_cache_end_of_result	sql_cache.cc	735
storing result in query cache	0.000007	mysql_parse	sql_parse.cc	6142
freeing items	0.000018	dispatch_command	sql_parse.cc	2146
closing tables	0.000009	log_slow_statement	sql_parse.cc	2204
logging slow query	0.000006	dispatch_command	sql_parse.cc	2169

17 rows in set (0.00 sec)

Status	Duration	Source_function	Source_file	Source_line
(initialization)	0.000012	send_result_to_client	sql_cache.cc	1143
checking query cache for query	0.000001	send_result_to_client	sql_cache.cc	1224
checking privileges on cached	0.000007	send_result_to_client	sql_cache.cc	1317
sending cached result to clien	0.000025	log_slow_statement	sql_parse.cc	2204
logging slow query	0.000007	dispatch_command	sql_parse.cc	2169

5 rows in set (0.00 sec)

18. Create Objects Appropriately

- ❖ Using 1 table instead of 'n' for same column structure
 - ❖ e.g. Code table for each type of code
- ❖ Splitting tables for optimal storage
 - ❖ e.g. Placing optional TEXT/BLOB columns in second table
- ❖ Use permanent tables instead of TEMPORARY tables

19. Naming Standards

- ❖ Name all Primary Key's Uniquely
 - ❖ e.g. customer_id, order_id not id
- ❖ Use Data Dictionary SQL to verify data types
 - ❖ Data Types & Lengths
- ❖ Be Descriptive
 - ❖ e.g. invoice_date not just date
- ❖ Avoid Reserved Words
 - ❖ e.g. date, time, timestamp

20. Testing, Testing, Testing

- ❖ You must have a testing environment
- ❖ Testing on a Production server is not an option



Best
Practice

**The goal of a testing environment
is not to test your software,
it is to break your software.**

Executive Summary

- ❖ Learn and know MySQL specifics
- ❖ Disk = Memory = Performance
- ❖ If you don't know your SQL you don't know your application. Log, Review & Monitor all SQL
- ❖ Know all benefits of different indexes
- ❖ You must test to failure in a dedicated test environment

Professional Help is Available

❖ PrimeBase Technologies

- ❖ Technology Experts
- ❖ Solution Experts

❖ 2 decades Expertise & Experience in Enterprise RDBMS Data Architecture

❖ 9 years in MySQL

www.ronaldbradford.com/contact