

MySQL Internals Manual

MySQL Internals Manual

Abstract

This is the MySQL Internals Manual.

Document generated on: 2012-08-01 (revision: 31)

For legal information, see the [Legal Notice](#).

Table of Contents

Preface and Legal Notice	xiii
1. A Guided Tour Of The MySQL Source Code	1
Getting the Source Tree	1
The Major Directories	2
Major Directories: <code>BUILD</code>	3
Major Directories: <code>client</code>	4
Major Directories: <code>myisam</code>	4
Major Directories: <code>mysys</code>	5
Major Directories: <code>sql</code>	5
Major Directories: <code>vio</code>	6
The Flow	7
The Open-Source Directories	7
The Internal and External Storage Engine Directories	8
The OS-Specific Directories	8
Odds and Ends	8
A Chunk of Code in <code>/sql/sql_update.cc</code>	9
The Skeleton of the Server Code	10
Recap	19
2. Coding Guidelines	21
General Development Guidelines	21
C/C++ Coding Guidelines of MySQL Server	21
How we maintain the server coding guidelines	21
Indentation and Spacing	22
Naming Conventions	25
Commenting Code	25
Header Files	27
Additional suggestions	27
Suggested mode in <code>emacs</code>	29
Basic <code>vim</code> setup	30
Another <code>vim</code> setup	30
An example setup for <code>ctags</code>	30
C++ Coding Guidelines for the NDB Storage Engine	30
Braces	31
Assignment	31
Use of <code>ndbrequire</code>	31
DEBUG Tags	31
3. Reusable Classes and Templates	33
Containers	33
Array	33
<code>I_P_List</code>	33
<code>I_List</code>	33
Memory Management	33
<code>MEM_ROOT</code>	33
How to Extend This Page	33
4. How to Build MySQL Server with CMake	35
Introduction	35
Prerequisites	35
How to install CMake	35
Very quick how-to-build	36
Build using the same options as in MySQL official release	37
Long description of how-to build	37
Create the build directory	37
Configuration step	37
Listing configuration parameters	38
Changing configuration parameters	38

Building Debug Configurations	39
./configure emulation	41
For Developers: how to write platform checks	42
FAQ / Miscellany	45
5. Autotools to CMake Transition Guide	47
Introduction	47
Command Invocation Syntax	47
Installation Layout Options	47
Storage Engine Options	48
Library Options	49
Miscellaneous Options	49
Debugging the Configuration Process	50
Interface to Third-Party Tools	50
6. Plugins	51
CMake Variables to Control Plugin Building (MySQL 5.5 and Later)	51
CMake Macro to Define the Plugin	51
Autotools configure Support (MySQL 5.1)	52
Autotools Plugin Macros	52
Specifying <code>mysqld</code> Variables Within a Plugin	55
Additional Accessors for Plugins	55
MySQL Services for Plugins	55
Plugin Services Components	55
Writing Services for Use from Within Plugins	56
7. Transaction Handling in the Server	57
Historical Note	57
Current Situation	57
Data Layout	58
Transaction Life Cycle	58
Roles and Responsibilities	59
Additional Notes on DDL and the Normal Transaction	60
8. The Optimizer	61
Code and Concepts	61
Definitions	61
The Optimizer Code	61
Primary Optimizations	62
Optimizing Constant Relations	62
Optimizing Joins	64
Transpositions	70
<code>ORDER BY</code> Clauses	71
<code>GROUP BY</code> and Related Conditions	72
Other Optimizations	73
9. Tracing the Optimizer	81
Typical Usage	81
System Variables Controlling the Trace	81
The <code>INFORMATION_SCHEMA.OPTIMIZER_TRACE</code> Table	82
Traceable Queries	82
Automatic Trace Purging	82
Tuning Trace Purging	82
Tracing Memory Usage	83
Privilege Checking	83
Interaction with the <code>--debug</code> Option	83
The <code>optimizer_trace</code> System Variable	83
The <code>end_markers_in_json</code> System Variable	83
Selecting Optimizer Features to Trace	84
General Trace Structure	84
Example	84
Nicely Displaying a Trace	89
Preventing Use of Optimizer Trace	90

Testing	90
Implementation	90
10. Memory Allocation	91
Memory Allocation in the MySQL Server (<code>sql</code> Directory)	91
Memory Allocation in a Library or Storage Engine	91
11. Important Algorithms and Structures	93
The <code>Item</code> Class	93
How MySQL Does Sorting (<code>filesort</code>)	94
Bulk Insert	95
How MySQL Does Caching	95
How MySQL Uses the Join Buffer Cache	96
How MySQL Handles <code>FLUSH TABLES</code>	97
Full-Text Search	97
<code>FLOAT</code> and <code>DOUBLE</code> Data Type Representation	101
Date and Time Data Type Representation	105
Threads	106
Character Sets and Collations	107
Error Flags and Functions	111
Functions in the <code>mysys</code> Library	112
Bitmaps	115
12. File Formats	117
MySQL <code>.frm</code> File Format	117
13. How MySQL Performs Different Selects	123
Steps of Select Execution	123
<code>select_result</code> Class	123
<code>SIMPLE</code> or <code>PRIMARY SELECT</code>	124
Structure Of Complex Select	124
Non-Subquery <code>UNION</code> Execution	125
Derived Table Execution	126
Subqueries	126
Single Select Engine	127
Union Engine	127
Special Engines	127
Explain Execution	128
14. How MySQL Transforms Subqueries	129
<code>Item_in_subselect::select_transformer</code>	129
Scalar <code>IN</code> Subquery	129
Row <code>IN</code> Subquery	131
<code>Item_allany_subselect</code>	132
<code>Item_singlerow_subselect</code>	132
15. MySQL Client/Server Protocol	135
Organization	136
Elements	136
The Packet Header	137
Packet Types	137
Handshake Initialization Packet	137
Client Authentication Packet	139
Password Functions	140
Passwords Before MySQL 4.0	140
Passwords in MySQL 4.1 and Later	140
Command Packet (Overview)	140
Command Packet (Detailed Description)	141
<code>COM_QUIT</code>	141
<code>COM_INIT_DB</code>	141
<code>COM_QUERY</code>	141
<code>COM_FIELD_LIST</code>	141
<code>COM_CREATE_DB</code>	141
<code>COM_DROP_DB</code>	142

COM_REFRESH	142
COM_SHUTDOWN	142
COM_STATISTICS	142
COM_PROCESS_INFO	142
COM_PROCESS_KILL	142
COM_DEBUG	143
COM_PING	143
COM_CHANGE_USER	143
COM_BINLOG_DUMP	143
COM_TABLE_DUMP	143
COM_REGISTER_SLAVE	143
COM_PREPARE	144
COM_EXECUTE	144
COM_LONG_DATA	144
COM_CLOSE_STMT	144
COM_RESET_STMT	144
COM_SET_OPTION	144
COM_FETCH_STMT	144
Result Packet Types	145
OK Packet	145
Error Packet	146
Result Set Header Packet	147
Field Packet	147
Row Data Packet	149
Row Data Packet: Binary (Tentative Description)	150
EOF Packet	150
OK for Prepared Statement Initialization Packet	151
Parameter Packet (Tentative Description)	152
Long Data Packet (Tentative Description)	152
Execute Packet (Tentative Description)	152
Compression	153
Encryption	154
16. Stored Programs	155
Overview	155
Persistent Representation	156
Stored Procedure and Stored Function Storage	156
Table Trigger Storage	156
Event Storage	157
Derived Attribute Storage	158
Internal Representation	158
Instructions	158
Parsing Context	159
Stored Program Parser	161
Parser Structure	161
Single-Pass Code Generation	163
Flow Analysis Optimizations	164
Dead Code Removal	165
Jump Shortcut Resolution	166
Stored Program Caches	167
Stored Procedure Cache	167
Stored Function Cache	168
Table Trigger Cache	168
Events and Caching	168
Stored Program Execution	169
Runtime Context	169
Executing One Instruction	171
Flow Control	171
Exception Handling	172

Call Nesting	173
17. Prepared Statement and Stored Routine Re-Execution	175
Statement Re-Execution Requirements	175
Preparation of a Prepared Statement	176
Execution of a Prepared Statement	177
Execution of a Stored Procedure Statement	178
18. Writing a Procedure	181
Extend Class <code>Procedure</code>	181
Constructor	181
<code>change_columns()</code>	181
<code>send_row()</code>	182
<code>add()</code>	182
<code>end_group()</code>	182
<code>end_of_records()</code>	183
Initialization Callback	183
Calling Sequence	183
Required Server Code Patches	183
Examples	184
<code>PROCEDURE ROWNUM()</code> – Adding a <code>RowNum</code> Pseudo Column to a Result Set	184
<code>PROCEDURE CALLTRACE()</code> – Simple Trace of Procedure Member Function Calls	185
19. Replication	187
Chapter Organization	187
Source Code Files	188
Principles	190
Binlog Formats	190
Differences Between Master and Slave	191
Crashes	194
Binlog Files and Indexes	194
Rules	194
Determining the Logging Format	194
Unsafe Statements	195
Logging Transactions	195
Logging Updates to <code>auto_increment</code> Columns	198
Logging Access to Variables and Functions	198
Other Unsafe Statements	198
<code>binlog_row_image</code>	198
Replication Locks	199
20. The Binary Log	201
Binary Log Overview	201
High-Level Binary Log Structure and Contents	202
Source Files Related to the Binary Log	203
Source File Archaeological Notes	203
Generating Browsable Binary Log Information	204
Event Classes and Types	204
Event Class Archaeological Notes	208
Event Meanings	209
Event Structure	211
Event Content-Writing Conventions	213
Event Header Fields	214
Event Data Fields (Event-Specific Information)	218
Binary Log Versions	219
Determining the Binary Log Version	221
Ensuring Compatibility of Future Binary Log Versions	222
Event Data for Specific Event Types	223
LOAD DATA INFILE Events	235
Row-Based Binary Logging	238
Additional Resources	239
21. MySQL Backup	241

Glossary	241
Architectural Design & APIs	241
Stub for WL#3169 material (Server kernel)	241
Stub for WL#3473 material (Backup engine API)	241
Stub for WL#3569 material (Backup synchronization algorithm)	242
Stub for WL#3574 material (Online backup: Service API for Metadata Objects II)	242
Backup Engines and Drivers	242
Stub for WL#3570 material (Default algorithm backup & restore (blocking))	242
Stub for WL#3776 material (Consistent snapshot backup for MVCC)	242
Stub for WL#866 material (Online backup: MyISAM native driver (redo-log algorithm, modify source file, log byte changes))	242
Stub for WL#4064 material (Online Backup: Adapt native MyISAM driver to new kernel architecture)	242
Stub for WL#4205 material (Online Backup: modify kernel to use object services API)	242
Stub for WL#3576 material (synchronization with binary log)	242
Stub for WL#4060 material (kernel updates)	242
Stub for WL#4062 material (DDL blocker, A.K.A metadata freeze)	243
Stub for WL#3956 material (commit blocker)	243
Error Handling	243
Stub for WL#3904 (Error Handling)	243
Stub for WL#4190 (Error Handling II)	243
Testing	243
Stub for WL#3324 (Test breakpoints)	243
Stub for WL#4120 (Breakpoints)	243
Data Formats	243
Stub for WL#3568 (Online backup: Stream format and server file handling)	243
WL#4063 (Online Backup: Finalize backup stream format)	243
22. MyISAM Storage Engine	245
MyISAM Record Structure	245
Introduction	245
Physical Attributes of Columns	247
Where to Look For More Information	251
The .MYI File	251
MyISAM Files	257
MyISAM Dynamic Data File Layout	257
Layout of the Record Storage Frame (Record Part, Record Block)	257
Record Contents	259
Packed Record Layout	259
In-memory Record Layout	259
MyISAM Compressed Data File Layout	259
Huffman Compression	259
The myisampack Program	259
Record and Blob Length Encoding	260
Code Tree Representation	260
Usage of the Index File	260
myisampack Tricks	260
Detailed Specification of the Decoding	260
MyISAM Key Cache	262
MyISAM Concurrent Insert	262
23. InnoDB Storage Engine	265
InnoDB Record Structure	265
High-Altitude Picture	265
Where to Look for More Information	268
InnoDB Page Structure	268
High-Altitude View	269
Example	272
Where to Look For More Information	273
24. Writing a Custom Storage Engine	275

Additional Resources	276
Overview	276
Creating Storage Engine Source Files	276
Adding Engine Specific Variables and Parameters	277
Creating the <code>handlerton</code>	277
Handling Handler Instantiation	280
Defining Filename Extensions	280
Creating Tables	281
Opening a Table	282
Implementing Basic Table Scanning	282
Implementing the <code>store_lock()</code> Method	283
Implementing the <code>external_lock()</code> Method	284
Implementing the <code>rnd_init()</code> Method	284
Implementing the <code>info(uinf flag)</code> Method	284
Implementing the <code>extra()</code> Method	285
Implementing the <code>rnd_next()</code> Method	285
Closing a Table	287
Adding Support for <code>INSERT</code> to a Storage Engine	287
Adding Support for <code>UPDATE</code> to a Storage Engine	288
Adding Support for <code>DELETE</code> to a Storage Engine	288
Supporting Non-Sequential Reads	289
Implementing the <code>position()</code> Method	289
Implementing the <code>rnd_pos()</code> Method	289
Supporting Indexing	289
Indexing Overview	290
Getting Index Information During <code>CREATE TABLE</code> Operations	290
Creating Index Keys	291
Parsing Key Information	291
Providing Index Information to the Optimizer	291
Preparing for Index Use with <code>index_init()</code>	293
Cleaning up with <code>index_end()</code>	293
Implementing the <code>index_read()</code> Method	293
Implementing the <code>index_read_idx()</code> Method	294
Implementing the <code>index_read_last()</code> Method	294
Implementing the <code>index_next()</code> Method	294
Implementing the <code>index_prev()</code> Method	294
Implementing the <code>index_first()</code> Method	294
Implementing the <code>index_last()</code> Method	295
Supporting Transactions	295
Transaction Overview	295
Starting a Transaction	295
Implementing ROLLBACK	297
Implementing COMMIT	297
Adding Support for Savepoints	298
API Reference	299
<code>bas_ext</code>	299
<code>close</code>	300
<code>create</code>	300
<code>delete_row</code>	301
<code>delete_table</code>	302
<code>external_lock</code>	303
<code>extra</code>	303
<code>index_end</code>	304
<code>index_first</code>	305
<code>index_init</code>	305
<code>index_last</code>	306
<code>index_next</code>	306
<code>index_prev</code>	307

index_read	308
index_read_idx	308
index_read_last	309
info	310
open	311
position	312
records_in_range	313
rnd_init	314
rnd_next	314
rnd_pos	315
start_stmt	316
store_lock	316
update_row	318
write_row	319
FAQ	320
25. Test Synchronization	321
Sleep	321
Wait Condition	322
Dbug Sleep	322
Error Injection	323
User-Level Locks	324
Debug Sync Point	325
Backup Breakpoint	326
Debug Sync Facility	326
Formal Syntax for DEBUG_SYNC Values	328
Debug Sync Activation/Deactivation	328
Debug Sync Implementation	329
A Typical Synchronization Pattern	329
Co-work With the DBUG Facility	330
Debug Sync Further Reading	330
26. Injecting Test Faults	333
Test Fault Macros	333
Test Fault Macro Usage	333
27. How to Create Good Test Cases	335
Formal Stuff	337
Coding Style	337
SQL Statement Example	337
Please Avoid too Dense Code	337
Header of Scripts	338
Header Example	338
Comments Within Boxes	338
Comments Everywhere	338
Complicated test Architecture or Tricky Code	339
A Subtest Case is Able to Reveal a Bug	339
Subtests	339
Make Test Protocols More Understandable	339
Better Protocol Example	340
Some Rules Outside of the Formal Stuff	340
Error Masking	340
Use of the Option <code>--disable_abort_on_error</code>	340
Perfect Cleanup at the End of a Test	341
Use of OS-Specific Commands	341
“Negative” Tests	341
Tests with Several Variants	342
Variation of the Storage Engine	342
Variation of the Protocol Option	343
Miscellaneous Tips	344
Stability Checks for Random Timing Problems	344

Stability Checks for Hostname-Related Problems	345
Stability Checks for Unexpected Problems	345
Row Order Within Result Sets	345
Test Case Behavior Should Depend on the Result of a <code>SHOW</code>	345
Does a Test Execution Hang?	345
If Your Test is Complicated, Support Debugging	345
You Need to Know Some Syntax, an Option, Etc., but Hate Reading Manuals	346
Notes on MTR2	346
Examples of “Suspicious” Tests and Scripts	346
Probably Mismatch of Focus of Test and Code Sequence	347
Too Greedy Test	347
Risky Handling Around Additional Sessions	348
28. Porting to Other Systems	351
Debugging a MySQL Server	352
Compiling MySQL for Debugging	352
Creating Trace Files	353
Debugging <code>mysqld</code> under <code>gdb</code>	354
Using a Stack Trace	355
Using Server Logs to Find Causes of Errors in <code>mysqld</code>	356
Making a Test Case If You Experience Table Corruption	356
Debugging a MySQL Client	357
The DBUG Package	357
Comments about RTS Threads	359
Differences Between Thread Packages	360
29. Error Messages	361
Adding New Error Messages to MySQL	361
Adding an Error Message to Multiple MySQL Versions	362
Adding Error Messages to MySQL 5.0.3 or Higher	363
Adding Error Messages to Old (< MySQL 5.0) Versions	365
Adding Storage Engine Error Messages	366
A. MySQL Source Code Distribution	367
Directory Listing	367
The <code>BUILD</code> Directory	369
The <code>client</code> Directory	369
The <code>config</code> Directory	369
The <code>cmd-line-utils</code> Directory	370
The <code>debug</code> Directory	371
The <code>Docs</code> Directory	371
The <code>extra</code> Directory	371
The <code>heap</code> Directory	372
The <code>include</code> Directory	373
The <code>innobase</code> Directory	373
The <code>libmysql</code> Directory	373
The <code>libmysql_r</code> Directory	374
The <code>libmysqld</code> Directory	374
The <code>man</code> Directory	374
The <code>myisam</code> Directory	374
The <code>myisammrg</code> Directory	376
The <code>mysql-test</code> Directory	377
The <code>mysys</code> Directory	377
The <code>ndb</code> Directory	381
The <code>netware</code> Directory	382
The <code>NEW-RPMS</code> Directory	382
The <code>os2</code> Directory	382
The <code>pstack</code> Directory	382
The <code>regex</code> Directory	383
The <code>SCCS</code> Directory	383
The <code>scripts</code> Directory	383

The <code>server-tools</code> Directory	384
The <code>sql</code> Directory	384
The <code>sql-bench</code> Directory	388
The <code>sql-common</code> Directory	388
The <code>SSL</code> Directory	388
The <code>strings</code> Directory	388
The <code>support-files</code> Directory	390
The <code>tests</code> Directory	390
The <code>tools</code> Directory	391
The <code>VC++Files</code> Directory	391
The <code>vio</code> Directory	392
The <code>zlib</code> Directory	392
B. <code>InnoDB</code> Source Code Distribution	395
Index	405

Preface and Legal Notice

This is a manual about MySQL internals. MySQL development personnel change it on an occasional basis. We make no guarantee that it is fully up to date. We do hope it illustrates how MySQL programmers work, and how MySQL Server works as a result.

Legal Notices

Copyright © 1997, 2012, Oracle and/or its affiliates. All rights reserved.

This software and related documentation are provided under a license agreement containing restrictions on use and disclosure and are protected by intellectual property laws. Except as expressly permitted in your license agreement or allowed by law, you may not use, copy, reproduce, translate, broadcast, modify, license, transmit, distribute, exhibit, perform, publish, or display any part, in any form, or by any means. Reverse engineering, disassembly, or decompilation of this software, unless required by law for interoperability, is prohibited.

The information contained herein is subject to change without notice and is not warranted to be error-free. If you find any errors, please report them to us in writing.

If this software or related documentation is delivered to the U.S. Government or anyone licensing it on behalf of the U.S. Government, the following notice is applicable:

U.S. GOVERNMENT RIGHTS Programs, software, databases, and related documentation and technical data delivered to U.S. Government customers are "commercial computer software" or "commercial technical data" pursuant to the applicable Federal Acquisition Regulation and agency-specific supplemental regulations. As such, the use, duplication, disclosure, modification, and adaptation shall be subject to the restrictions and license terms set forth in the applicable Government contract, and, to the extent applicable by the terms of the Government contract, the additional rights set forth in FAR 52.227-19, Commercial Computer Software License (December 2007). Oracle USA, Inc., 500 Oracle Parkway, Redwood City, CA 94065.

This software is developed for general use in a variety of information management applications. It is not developed or intended for use in any inherently dangerous applications, including applications which may create a risk of personal injury. If you use this software in dangerous applications, then you shall be responsible to take all appropriate fail-safe, backup, redundancy, and other measures to ensure the safe use of this software. Oracle Corporation and its affiliates disclaim any liability for any damages caused by use of this software in dangerous applications.

Oracle is a registered trademark of Oracle Corporation and/or its affiliates. MySQL is a trademark of Oracle Corporation and/or its affiliates, and shall not be used without Oracle's express written authorization. Other names may be trademarks of their respective owners.

This software and documentation may provide access to or information on content, products, and services from third parties. Oracle Corporation and its affiliates are not responsible for and expressly disclaim all warranties of any kind with respect to third-party content, products, and services. Oracle Corporation and its affiliates will not be responsible for any loss, costs, or damages incurred due to your access to or use of third-party content, products, or services.

This documentation is in prerelease status and is intended for demonstration and preliminary use only. It may not be specific to the hardware on which you are using the software. Oracle Corporation and its affiliates are not responsible for and expressly disclaim all warranties of any kind with respect to this documentation and will not be responsible for any loss, costs, or damages incurred due to the use of this documentation.

The information contained in this document is for informational sharing purposes only and should be considered in your capacity as a customer advisory board member or pursuant to your beta trial agreement only. It is not a commitment to deliver any material, code, or functionality, and should not be

relied upon in making purchasing decisions. The development, release, and timing of any features or functionality described in this document remains at the sole discretion of Oracle.

This document in any form, software or printed matter, contains proprietary information that is the exclusive property of Oracle. Your access to and use of this material is subject to the terms and conditions of your Oracle Software License and Service Agreement, which has been executed and with which you agree to comply. This document and information contained herein may not be disclosed, copied, reproduced, or distributed to anyone outside Oracle without prior written consent of Oracle or as specifically provided below. This document is not part of your license agreement nor can it be incorporated into any contractual agreement with Oracle or its subsidiaries or affiliates.

This documentation is NOT distributed under a GPL license. Use of this documentation is subject to the following terms:

You may create a printed copy of this documentation solely for your own personal use. Conversion to other formats is allowed as long as the actual content is not altered or edited in any way. You shall not publish or distribute this documentation in any form or on any media, except if you distribute the documentation in a manner similar to how Oracle disseminates it (that is, electronically for download on a Web site with the software) or on a CD-ROM or similar medium, provided however that the documentation is disseminated together with the software on the same medium. Any other use, such as any dissemination of printed copies or use of this documentation, in whole or in part, in another publication, requires the prior written consent from an authorized representative of Oracle. Oracle and/or its affiliates reserve any and all rights to this documentation not expressly granted above.

Chapter 1. A Guided Tour Of The MySQL Source Code

Table of Contents

Getting the Source Tree	1
The Major Directories	2
Major Directories: <code>BUILD</code>	3
Major Directories: <code>client</code>	4
Major Directories: <code>myisam</code>	4
Major Directories: <code>mysys</code>	5
Major Directories: <code>sql</code>	5
Major Directories: <code>vio</code>	6
The Flow	7
The Open-Source Directories	7
The Internal and External Storage Engine Directories	8
The OS-Specific Directories	8
Odds and Ends	8
A Chunk of Code in <code>/sql/sql_update.cc</code>	9
The Skeleton of the Server Code	10
Recap	19

What we're about to do in this section is pick up the latest copy of the MySQL source code off the Internet. Then we'll get a list of the directories and comment on why they're there. Next we'll open up some of the files that are vital to MySQL's working, and comment on specific lines in the source code. We'll close off with a few pictures of file formats.

Getting the Source Tree

We want to download the latest, the very latest, version of the MySQL server. So we won't click Downloads on the MySQL Developer Zone page that's usually a few weeks old. Instead we'll use Bazaar, which is a revision control package, vaguely like CVS or Perforce. This is what MySQL's developers use every day, so what we download with Bazaar is usually less than a day old. If you've ever submitted a bug report and gotten the response "thanks, we fixed the bug in the source code repository" that means you can get the fixed version with Bazaar.

The general syntax to create a local copy (branch) of a Bazaar source code repository is:

```
shell> bzz branch <remote repository> <local name>
```

To obtain a local copy of the MySQL Server 5.1 source tree, this would require the following commands:

```
shell> bzz init-repo $HOME/mysql-server
shell> cd $HOME/mysql-server
shell> bzz branch lp:mysql-server/5.1 mysql-5.1
```

(The `$HOME` directory is usually your personal area that you're allowed to write to. If that's not the case, replace `$HOME` with your personal choice whenever it appears.)

There is a lot of code, so the first time you do this the download will take over an hour. That's if you're lucky.

If you're glitch-prone, you'll need to read the manual: [Installing from the Development Source Tree](#).

On later occasions, you'll be doing what's called a *bzz pull* instead of a *bzz branch*, and it will go faster (as it will only fetch the recent changes that have been committed to the remote repository since the last time you pulled).

Directories, Alphabetical Order

After `bzr branch` finished you'll have some 40 new sets of files on your computer, as you'll be able to see with `ls` or `dir`.

```
BUILD
Docs
NEW-RPMS
SSL
VC++Files
bdb
client
cmd-line-utils
config
dbug
extra
heap
include
innobase
libmysql
libmysql_r
libmysqld
man
myisam
myisammrg
mysql-test
mysys
ndb
netware
os2
pstack
regex
scripts
server-tools
sql
sql-bench
sql-common
strings
support-files
tests
tools
vio
zlib
```

These will all be installed as directories below the directory that was created by the `bzr branch` command. At first all these directory names might intimidate you, and that's natural. After all, MySQL is a big project. But we're here to show you that there's order in this apparent chaos.

The Major Directories

1. `BUILD`
2. `client`
3. `Docs`
4. `myisam`
5. `mysys`
6. `sql`
7. `vio`

The orderly approach is to look first at the most important directories, then we'll look at the whole list in our second pass. So, first, let's look at what you'll find in just seven of the directories: `BUILD`, `client`, `Docs`, `myisam`, `mysys`, `sql`, and `vio`.

Major Directories: BUILD

The first major directory we'll look at is BUILD. It actually has very little in it, but it's useful, because one of the first things you might want to do with the source code is: compile and link it.

The example command line that we could use is

```
shell> ./BUILD/compile-pentium-debug --prefix=$HOME/mysql-bin
```

It invokes a batch file in the BUILD directory. When it's done, you'll have an executable MySQL server and client.

Or, um, well, maybe you won't. Sometimes people have trouble with this step because there's something missing in their operating system version, or whatever. Don't worry, it really does work, and there are people around who might help you if you have trouble with this step. Search for "build" in the archives of lists.mysql.com.

We, when we're done building, tend to install it with the following sequence:

```
shell> make
shell> make install
shell> $HOME/mysql-bin/bin/mysql_install_db\
--basedir=$HOME/mysql-bin\
--datadir=$HOME/mysql-bin/var
```

This puts the new MySQL installation files on

```
shell> $HOME/mysql-bin/libexec      -- for the server
shell> $HOME/mysql-bin/bin          -- for the mysql client
shell> $HOME/mysql-bin/var          -- for the databases
```

GNU Debugger

Once you've got something that runs, you can put a debugger on it. We recommend use of the GNU debugger

```
http://www.gnu.org/software/gdb/documentation/
```

And many developers use the graphical debugger tool DDD - Data Display Debugger

```
http://www.gnu.org/software/ddd/manual/
```

These are free and common, they're probably on your Linux system already.

There are debuggers for Windows and other operating systems, of course don't feel left out just because we're mentioning a Linux tool name! But it happens that we do a lot of things with Linux ourselves, so we happen to know what to say. To debug the mysqld server, say:

```
shell> ddd --gdb --args \
    $HOME/mysql-bin/libexec/mysqld \
    --basedir=$HOME/mysql-bin \
    --datadir=$HOME/mysql-bin/var\
    --skip-networking
```

From this point on, it may be tempting to follow along through the rest of the "guided tour" by setting breakpoints, displaying the contents of variables, and watching what happens after starting a client from another shell. That would be more fun. But it would require a detour, to discuss how to use the debugger. So we'll plow forward, the dull way, noting what's in the directories and using a text editor to note what's in the individual files.

Running a Test with the Debugger

To run a test named `some.test` with the debugger in embedded mode you could do this:

1. Run `libmysqld/examples/test_run --gdb some.test`. This creates a `libmysqld/examples/test-gdbinit` file which contains the required parameters for **mysqltest**.
2. Make a copy of the `test-gdbinit` file (call it, for example, `some-gdbinit`). The `test-gdbinit` file will be removed after `test-run --gdb` has finished.
3. Load `libmysqld/examples/mysqltest_embedded` into your favorite debugger, for example: `gdb mysqltest_embedded`.
4. In the debugger, for example in `gdb`, do: `--sou some-gdbinit`

Now `some.test` is running, and you can see if it's passing or not.

If you just want to debug some queries with the embedded server (not the test), it's easier to just run `libmysqld/examples/mysql`. It's the embedded server-based clone of the usual **mysql** tool, and works fine under **gdb** or whatever your favorite debugger is.

Major Directories: `client`

The next major directory is `mysql-5.0/client`.

```
size  name          comment
----  -
100034 mysql.cc          "The MySQL command tool"
 36913 mysqladmin.c  maintenance of MYSQL databases
 22829 mysqlshow.c   show databases, tables, or columns
+ 12 more .c and .cc programs
```

It has the source code of many of your familiar favorites, like `mysql`, which everybody has used to connect to the MySQL server at one time or another. There are other utilities too in fact, you'll find the source of most client-side programs here. There are also programs for checking the password, and for testing that basic functions such as threading or access via SSL are possible.

You'll notice, by the way, that we're concentrating on the files that have extension of `".c"` or `".cc"`. By now it's obvious that C is our principal language although there are some utilities written in Perl as well.

Major Directories: `myisam`

The next major directory is labelled `myisam`. We will begin by mentioning that `myisam` is one of what we call the MySQL storage engine directories.

```
The MySQL storage engine directories:
heap          -- also known as 'memory'
innodb        -- maintained by Innobase Oy
myisam        -- see next section!
ndb           -- ndb cluster
```

For example the `heap` directory contains the source files for the heap storage engine and the `ndb` directory contains the source files for the `ndb` storage engine.

But the files in those directories are mostly analogues of what's in the `myisam` directory, and the `myisam` directory is sort of a 'template'.

On the `myisam` directory, you'll find the programs that do file I/O. Notice that the file names begin with the letters `mi`, by the way. That stands for `MyISAM`, and most of the important files in this directory start with `mi`.

File handling programs on `mysql-5.0/myisam`:

```
size  name          comment
----  -
40301 mi_open.c      for opening
 3593 mi_close.c     for closing
 1951 mi_rename.c    for renaming
```

```
+ more mi_*.c programs
```

Row handling programs on mysql-5.0/myisam:

```
size  name          comment
----  -
29064 mi_delete.c    for deleting
 2562 mi_delete_all.c for deleting all
 6797 mi_update.c    for updating
32613 mi_write.c     for inserting
+ more mi_*.c programs
```

Drilling down a bit, you'll also find programs in the myisam directory that handle deleting, updating, and inserting of rows. The only one that's a little hard to find is the program for inserting rows, which we've called `mi_write.c` instead of `mi_insert.c`.

Key handling programs on mysql-5.0/myisam:

```
size  name          comment
----  -
 4668 mi_rkey.c      for random key searches
 3646 mi_rnext.c     for next-key searches
15440 mi_key.c       for managing keys
+ more mi_*.c programs
```

The final notable group of files in the myisam directory is the group that handles keys in indexes.

To sum up: (1) The myisam directory is where you'll find programs for handling files, rows, and keys. You won't find programs for handling columns we'll get to them a bit later. (2) The myisam directory is just one of the handler directories. The programs in the other storage engine directories fulfill about the same functions.

Major Directories: [mysys](#)

The next major directory is labelled `mysys`, which stands for MySQL System Library. This is the toolbox directory, for example it has low level routines for file access. The `.c` files in `mysys` have procedures and functions that are handy for calling by main programs, for example by the programs in the myisam directory. There are 115 `.c` files in `mysys`, so we only can note a sampling.

Sampling of programs on mysql-5.0/mysys

```
size  name          comment
----  -
17684 charset.c     character sets
 6165 mf_qsort.c     quicksort
 5609 mf__tempfile.c temporary files
+ 112 more *.c programs
```

Example one: with `charset.c` routines, you can change the character set.

Example two: `mf_qsort.c` contains our quicksort package.

Example three: `mf_tempfile.c` has what's needed for maintaining MySQL's temporary files.

You can see from these examples that `mysys` is a hodgepodge. That's why we went to the trouble of producing extra documentation in this document to help you analyze `mysys`'s contents.

Major Directories: [sql](#)

The next major directory is `mysql-5.0/sql`. If you remember your manual, you know that you must pronounce this: ess queue ell.

The "parser" programs on mysql-5.0/sql:

```
size  name          comment
----  -
```

```
51326 sql_lex.cc      lexer
230026 sql_yacc.yy    parser
+ many more *.cc programs
```

This is where we keep the parser. In other words, programs like `sql_lex.cc` and `sql_yacc.yy` are responsible for figuring out what's in an SQL command, and deciding what to do about it.

The "handler" programs on `mysql-5.0/sql`:

```
size  name          comment
----  ----
79798 ha_berkeley.cc  bdb
56687 ha_federated.cc federated (sql/med)
61033 ha_heap.cc     heap (memory)
214046 ha_innodb.cc  innodb
47361 ha_myisam.cc   myisam
14727 ha_myisammrg.cc merge
215091 ha_ndbcluster.cc ndb
```

This is also where we keep the handler programs. Now, you'll recall that the storage engine itself, for example `myisam`, is a separate directory. But here in the `sql` directory, we have programs which are responsible for determining which handler to call, formatting appropriate arguments, and checking results. In other words, the programs that begin with the letters `ha` are the handler interface programs, and there's one for each storage engine.

The "statement" routines in `mysql-5.0/sql`:

```
size  name          comment
----  ----
24212 sql_delete.cc  'delete ...' statement
1217  sql_do.cc        'do ...'
22362 sql_help.cc    'help ...'
75331 sql_insert.cc  'insert ...'
430486 sql_select.cc  'select ...'
130861 sql_show.cc   'show ...'
42346 sql_update.cc  'update ...'
+ many more sql_*.cc programs
```

Also in the `sql` directory, you'll find individual programs for handling each of the syntactical components of an SQL statement. These programs tend to have names beginning with `sql_`. So for the `SELECT` statement, check out `sql_select.cc`.

Thus, there are "statement" routines like `sql_delete.c`, `sql_load.c`, and `sql_help.c`, which take care of the `DELETE`, `LOAD DATA`, and `HELP` statements. The file names are hints about the SQL statements involved.

The "statement function" routines in `mysql-5.0/sql`:

```
size  name          comment
----  ----
19906 sql_string.cc  strings
6152  sql_olap.cc      olap (rollup)
14241 sql_udf.cc     user-defined functions
17669 sql_union.cc   unions
```

Then there are the routines for components of statements, such as strings, or online analytical processing which at this moment just means `ROLLUP`, or user-defined functions, or the `UNION` operator.

Major Directories: [vio](#)

The final major directory that we'll highlight is labelled `vio`, for "virtual I/O".

The `vio` routines are wrappers for the various network I/O calls that happen with different protocols. The idea is that in the main modules one won't have to write separate bits of code for each protocol. Thus `vio`'s purpose is somewhat like the purpose of Microsoft's `winsock` library.

That wraps up our quick look at the seven major directories. Just one summary chart remains to do.

The Flow

This is a diagram of the flow.

```
User enters "INSERT" statement      /* client */
|
|
Message goes over TCP/IP line        /* vio, various */
|
|
Server parses statement              /* sql */
|
|
Server calls low-level functions     /* myisam */
|
|
Handler stores in file               /* mysys */
```

The diagram is very simplified — it's a caricature that distorts important things, but remember that we've only discussed seven major directories so far: Docs, BUILD, and the five that you see here.

The flow works like this:

First, the client routines get an SQL statement from a user, allowing edit, performing initial checks, and so on.

Then, via the vio routines, the somewhat-massaged statement goes off to the server.

Next, the sql routines handle the parsing and call what's necessary for each individual part of the statement. Along the way, the sql routines will be calling the low level mysys routines frequently.

Finally, one of the ha (handler) programs in the sql directory will dispatch to an appropriate handler for storage. In this case we've assumed, as before, that the handler is myisam — so a myisam-directory program is involved. Specifically, that program is `mi_write.c`, as we mentioned earlier.

Simple, eh?

The Open-Source Directories

We're now getting into the directories which aren't "major." Starting with:

```
debug
pstack
regex
strings
zlib
```

Now it's time to reveal a startling fact, which is we didn't write all of the source code in all of the source code directories all by ourselves. This list is, in a sense, a tribute to the idea of open source.

There's `debug`, which is Fred Fish's debug library.

There's `pstack`, which displays the process stack.

There's `regex`, which is what we use for our regular expressions function.

There's `strings`, the meaning of which is obvious.

There's `zlib`, which is for Lempel-Ziv compression.

All of the programs in these directories were supplied by others, as open source code. We didn't just take them, of course. MySQL has checked and changed what's in these directories. But we acknowledge with thanks that they're the products of other projects, and other people's labor, and

we only regret that we won't have time to note all the contributed or publicly available components of MySQL, in this manual.

The Internal and External Storage Engine Directories

Continuing with our extract from the directory list ...

```
bdb                /* external */
heap
innobase           /* external */
myisam
myisammrg
ndb
```

Let's go through the idea of storage engines once more, this time with a list of all the storage engines, both the ones that we produce, and the ones that others produce. We've already mentioned the internal ones so now we'll remark on the directories of the two common external storage engines BDB and innobase.

The BDB, or Berkeley Database, handler, is strictly the product of Sleepycat software. Sleepycat has a web page at sleepycat.com, which contains, among other things, documentation for their product. So you can download Sleepycat's own documentation of the source code in the BDB directory.

As for the innobase handler, which many of you probably use, you'll be happy to know that the comments in the files are reasonably clear (the InnoDB Oy people are pretty strict about comments). There are two chapters about it in this document.

The OS-Specific Directories

```
netware
NEW-RPMS
os2
VC++Files
```

A few words are in order about the directories that contain files which relate to a particular environment that MySQL can run in.

The netware directory contains a set of files for interfacing with netware, and anyone who has an involvement with NetWare knows that we're allied with them, and so this is one of the directories that represents the joint enterprise.

The NEW-RPMS directory (empty at time of writing) is for Linux, and the os2 directory is for OS/2.

Finally, the VC++Files directory is for Windows. We've found that the majority of Windows programmers who download and build MySQL from source use Microsoft Visual C. In the VC++Files directory you will find a nearly complete replication of what's in all the other directories that we've discussed, except that the .c files are modified to account for the quirks of Microsoft tools.

Without endorsing by particular names, we should note that other compilers from other manufacturers also work.

Odds and Ends

Finally, for the sake of completeness, we'll put up a list of the rest of the directories those that we haven't had occasion to mention till now.

```
Source Code Administration Directories:
SCCS

Common .h Files:
include

GNU Readline library and related:
```

```
cmd-line-utils
```

```
Stand-alone Utility & Test Programs:  
extra  
mysql-test  
repl-tests  
support-files  
tests  
tools
```

You don't have to worry about the administration directories since they're not part of what you build.

You probably won't have to worry about the stand-alone programs either, since you just use them, you don't need to remake them.

There's an include directory that you **SHOULD** have a look at, because the common header files for programs from several directories are in here.

Finally, there are stand-alone utility and test programs. Strictly speaking they're not part of the "source code". But it's probably reassuring to know that there's a test suite, for instance. Part of the quality-assurance process is to run the scripts in the test suite before releasing.

And those are the last. We've now traipsed through every significant directory created during your download of the MySQL source package.

A Chunk of Code in `/sql/sql_update.cc`

Now, having finished with our bird's eye view of the source code from the air, let's take the perspective of the worms on the ground (which is another name for MySQL's developer staff -- turn on laugh track here).

```
int mysql_update(THD *thd, ...)  
{  
    ...  
    if ((lock_tables(thd, table_list)))  
        DEBUG_RETURN(1); ...  
    ...  
    init_read_record(&info, thd, table, select, 0, 1); ...  
    while (!(error=info.read_record(&info)) && !thd->killed)  
    {  
        ...  
        if (!(error=table->file->update_row((byte*) table->record[1]),  
                                              (byte*) table->record[0])))  
            updated++;  
        ...  
        if (table->triggers)  
            table->triggers->process_triggers(thd, TRG_EVENT_UPDATE, TRG_ACTION_AFTER);  
        ...  
    }  
    ...  
    if (updated && (error <= 0 || !transactional_table))  
    {  
        mysql_bin_log.write(&qinfo) && transactional_table;  
        ...  
    }  
}
```

Here's a snippet of code from a .c file in the sql directory, specifically from sql_update.cc, which as we mentioned earlier -- is invoked when there's an UPDATE statement to process.

The entire routine has many error checks with handlers for improbable conditions, and showing multiple screens would be tedious, so we've truncated the code a lot. Where you see an ellipsis (three dots in a row), that means "and so on".

So, what do we learn from this snippet of code? In the first place, we see that it's fairly conventional C code. A brace causes an indentation, instructions tend to be compact with few unnecessary spaces, and comments are sparse.

Abbreviations are common, for example `thd` stands for thread, you just have to get used to them. Typically a structure will be defined in a separate `.h` file.

Routine names are sometimes long enough that they explain themselves. For example, you can probably guess that this routine is opening and locking, allocating memory in a cache, initializing a process for reading records, reading records in a loop until the thread is killed or there are no more to read, storing a modified record for the table, and — after the loop is through — possibly writing to the log. Incidentally, a transactional table is usually a BDB or an InnoDB table.

Obviously we've picked out what's easy to follow, and we're not pretending it's all smooth sailing. But this is actual code and you can check it out yourself.

The Skeleton of the Server Code

And now we're going to walk through something harder, namely the server.

WARNING WARNING WARNING: code changes constantly, so names and parameters may have changed by the time you read this.

Important files we'll be walking through:

```
/sql/mysqld.cc  
/sql/sql_parse.cc  
/sql/sql_prepare.cc  
/sql/sql_insert.cc  
/sql/ha_myisam.cc  
/myisam/mi_write.c
```

This is not as simple as what we've just done. In fact we'll need multiple pages to walk through this one, and that's despite our use of truncation and condensation again. But the server is important, and if you can grasp what we're doing with it, you'll have grasped the essence of what the MySQL source code is all about.

We'll mostly be looking at programs in the `sql` directory, which is where `mysqld` and most of the programs for the SQL engine code are stored.

Our objective is to follow the server from the time it starts up, through a single INSERT statement that it receives from a client, to the point where it finally performs the low level write in the MyISAM file.

Walking Through The Server Code: `/sql/mysqld.cc`

```
int main(int argc, char **argv)  
{  
    _cust_check_startup();  
    (void) thr_setconcurrency(concurrency);  
    init_ssl();  
    server_init(); // 'bind' + 'listen'  
    init_server_components();  
    start_signal_handler();  
    acl_init((THD *)0, opt_noacl);  
    init_slave();  
    create_shutdown_thread();  
    create_maintenance_thread();  
    handle_connections_sockets(0); // !  
    DEBUG_PRINT("quit", ("Exiting main thread"));  
    exit(0);  
}
```

Here is where it all starts, in the main function of `mysqld.cc`.

Notice that we show a directory name and program name just above this snippet. We will do the same for all the snippets in this series.

By glancing at this snippet for a few seconds, you will probably see that the main function is doing some initial checks on startup, is initializing some components, is calling a function named

handle_connections_sockets, and then is exiting. It's possible that acl stands for "access control" and it's interesting that DBUG_PRINT is something from Fred Fish's debug library, which we've mentioned before. But we must not digress.

In fact there are 150 code lines in the main function, and we're only showing 13 code lines. That will give you an idea of how much we are shaving and pruning. We threw away the error checks, the side paths, the optional code, and the variables. But we did not change what was left. You will be able to find these lines if you take an editor to the mysqld.cc program, and the same applies for all the other routines in the snippets in this series.

The one thing you won't see in the actual source code is the little marker "// !". This marker will always be on the line of the function that will be the subject of the next snippet. In this case, it means that the next snippet will show the handle_connection_sockets function. To prove that, let's go to the next snippet.

Walking Through The Server Code: /sql/mysqld.cc

```
handle_connections_sockets (arg __attribute__((unused))
{
    if (ip_sock != INVALID_SOCKET)
    {
        FD_SET(ip_sock,&clientFDs);
        DBUG_PRINT("general",("Waiting for connections."));
        while (!abort_loop)
        {
            new_sock = accept(sock, my_reinterpret_cast(struct sockaddr*)
                (&cAddr), &length);
            thd= new THD;
            if (sock == unix_sock)
                thd->host=(char*) localhost;
            create_new_thread(thd);                // !
        }
    }
}
```

Inside handle_connections_sockets you'll see the hallmarks of a classic client/server architecture. In a classic client/server, the server has a main thread which is always listening for incoming requests from new clients. Once it receives such a request, it assigns resources which will be exclusive to that client. In particular, the main thread will spawn a new thread just to handle the connection. Then the main server will loop and listen for new connections — but we will leave it and follow the new thread.

As well as the sockets code that we chose to display here, there are several variants of this thread loop, because clients can choose to connect in other ways, for example with named pipes or with shared memory. But the important item to note from this section is that the server is spawning new threads.

Walking Through The Server Code: /sql/mysqld.cc

```
create_new_thread(THD *thd)
{
    pthread_mutex_lock(&LOCK_thread_count);
    pthread_create(&thd->real_id,&connection_attrib,
        handle_one_connection,                // !
        (void*) thd));
    pthread_mutex_unlock(&LOCK_thread_count);
}
```

Here is a close look at the routine that spawns the new thread. The noticeable detail is that, as you can see, it uses a mutex or mutual exclusion object. MySQL has a great variety of mutexes that it uses to keep actions of all the threads from conflicting with each other.

Walking Through The Server Code: /sql/sql_parse.cc

```
handle_one_connection(THD *thd)
{
    init_sql_alloc(&thd->mem_root, MEM_ROOT_BLOCK_SIZE, MEM_ROOT_PREALLOC);
    while (!net->error && net->vio != 0 && !thd->killed)
    {

```

```

        if (do_command(thd))          // !
            break;
    }
    close_connection(net);
    end_thread(thd,1);
    packet=(char*) net->read_pos;

```

With this snippet, we've wandered out of `mysqld.cc`. Now, we're in the `sql_parse` file, still in the `sql` directory. This is where the session's big loop is.

The loop repeatedly gets and does commands. When it ends, the connection closes. At that point, the thread will end and the resources for it will be deallocated.

But we're more interested in what happens inside the loop, when we call the `do_command` function.

Graphic:

```

client          <==== MESSAGE ====> server          <=====PACKETS =====>

Example:
INSERT INTO Table1 VALUES (1);

```

To put it graphically, at this point there is a long-lasting connection between the client and one server thread. Message packets will go back and forth between them through this connection. For today's tour, let's assume that the client passes the `INSERT` statement shown on the Graphic, for the server to process.

Walking Through The Server Code: `/sql/sql_parse.cc`

```

bool do_command(THD *thd)
{
    net_new_transaction(net);
    packet_length=my_net_read(net);
    packet=(char*) net->read_pos;
    command = (enum enum_server_command) (uchar) packet[0];
    dispatch_command(command,thd, packet+1, (uint) packet_length);
    // !
}

```

You've probably noticed by now that whenever we call a lower-level routine, we pass an argument named `thd`, which is an abbreviation for the word thread (we think). This is the essential context which we must never lose.

The `my_net_read` function is in another file called `net_serv.cc`. The function gets a packet from the client, uncompresses it, and strips the header.

Once that's done, we've got a multi-byte variable named `packet` which contains what the client has sent. The first byte is important because it contains a code identifying the type of message.

We'll pass that and the rest of the packet on to the `dispatch_command` function.

Walking Through The Server Code: `/sql/sql_parse.cc`

```

bool dispatch_command(enum enum_server_command command, THD *thd,
    char* packet, uint packet_length)
{
    switch (command) {
        case COM_INIT_DB:          ...
        case COM_REGISTER_SLAVE:   ...
        case COM_TABLE_DUMP:       ...
        case COM_CHANGE_USER:      ...
        case COM_EXECUTE:
            mysql_stmt_execute(thd,packet);
        case COM_LONG_DATA:        ...
        case COM_PREPARE:
            mysql_stmt_prepare(thd, packet, packet_length);    // !
        /* and so on for 18 other cases */
        default:

```

```
send_error(thd, ER_UNKNOWN_COM_ERROR);
break;
}
```

And here's just part of a very large switch statement in `sql_parse.cc`. The snippet doesn't have room to show the rest, but you'll see when you look at the `dispatch_command` function that there are more case statements after the ones that you see here.

There will be — we're going into list mode now and just reciting the rest of the items in the switch statement — code for prepare, close statement, query, quit, create database, drop database, dump binary log, refresh, statistics, get process info, kill process, sleep, connect, and several minor commands. This is the big junction.

We have cut out the code for all of the cases except for two, `COM_EXECUTE` and `COM_PREPARE`.

Walking Through The Server Code: `/sql/sql_prepare.cc`

We are not going to follow what happens with `COM_PREPARE`. Instead, we are going to follow the code after `COM_EXECUTE`. But we'll have to digress from our main line for a moment and explain what the prepare does.

```
"Prepare:
Parse the query
Allocate a new statement, keep it in 'thd->prepared statements' pool
Return to client the total number of parameters and result-set
metadata information (if any)"
```

The prepare is the step that must happen before execute happens. It consists of checking for syntax errors, looking up any tables and columns referenced in the statement, and setting up tables for the execute to use. Once a prepare is done, an execute can be done multiple times without having to go through the syntax checking and table lookups again.

Since we're not going to walk through the `COM_PREPARE` code, we decided not to show its code at this point. Instead, we have cut and pasted some code comments that describe prepare. All we're illustrating here is that there are comments in the code, so you will have aid when you look harder at the prepare code.

Walking Through The Server Code: `/sql/sql_parse.cc`

```
bool dispatch_command(enum enum_server_command command, THD *thd,
    char* packet, uint packet_length)
{
    switch (command) {
        case COM_INIT_DB:      ...
        case COM_REGISTER_SLAVE: ...
        case COM_TABLE_DUMP:   ...
        case COM_CHANGE_USER:  ...
        case COM_EXECUTE:
            mysql_stmt_execute(thd, packet);           // !
        case COM_LONG_DATA:    ...
        case COM_PREPARE:
            mysql_stmt_prepare(thd, packet, packet_length);
        /* and so on for 18 other cases */
    default:
        send_error(thd, ER_UNKNOWN_COM_ERROR);
        break;
    }
}
```

Let's return to the grand central junction again in `sql_parse.cc` for a moment. The thing to note on this snippet is that the line which we're really going to follow is what happens for `COM_EXECUTE`.

Walking Through The Server Code: `/sql/sql_prepare.cc`

```
void mysql_stmt_execute(THD *thd, char *packet)
{
    if (!(stmt=find_prepared_statement(thd, stmt_id, "execute")))
    {
```

```

        send_error(thd);
        DEBUG_VOID_RETURN;
    }
    init_stmt_execute(stmt);
    mysql_execute_command(thd);           // !
}

```

In this case, the line that we're following is the line that executes a statement.

Notice how we keep carrying the THD thread and the packet along with us, and notice that we expect to find a prepared statement waiting for us, since this is the execute phase. Notice as well that we continue to sprinkle error-related functions that begin with the letters DEBUG, for use by the debug library. Finally, notice that the identifier "stmt" is the same name that ODBC uses for the equivalent object. We try to use standard names when they fit.

Walking Through The Server Code: /sql/sql_parse.cc

```

void mysql_execute_command(THD *thd)
{
    switch (lex->sql_command) {
        case SQLCOM_SELECT: ...
        case SQLCOM_SHOW_ERRORS: ...
        case SQLCOM_CREATE_TABLE: ...
        case SQLCOM_UPDATE: ...
        case SQLCOM_INSERT: ...           // !
        case SQLCOM_DELETE: ...
        case SQLCOM_DROP_TABLE: ...
    }
}

```

In the mysql_execute_command function, we encounter another junction. One of the items in the switch statement is named SQLCOM_INSERT.

Walking Through The Server Code: /sql/sql_parse.cc

```

case SQLCOM_INSERT:
{
    my_bool update=(lex->value_list.elements ? UPDATE_ACL : 0);
    ulong privilege= (lex->duplicates == DUP_REPLACE ?
                      INSERT_ACL | DELETE_ACL : INSERT_ACL | update);
    if (check_access(thd,privilege,tables->db,&tables->grant.privilege))
        goto error;
    if (grant_option && check_grant(thd,privilege,tables))
        goto error;
    if (select_lex->item_list.elements != lex->value_list.elements)
    {
        send_error(thd,ER_WRONG_VALUE_COUNT);
        DEBUG_VOID_RETURN;
    }
    res = mysql_insert(thd,tables,lex->field_list,lex->many_values,
                      select_lex->item_list, lex->value_list,
                      (update ? DUP_UPDATE : lex->duplicates));
    // !
    if (thd->net.report_error)
        res= -1;
    break;
}

```

For this snippet, we've blown up the code around the SQLCOM_INSERT case in the mysql_execute_command function. The first thing to do is check whether the user has the appropriate privileges for doing an INSERT into the table, and this is the place where the server checks for that, by calling the check_access and check_grant functions. It would be tempting to follow those functions, but those are side paths. Instead, we'll follow the path where the work is going on.

Walking Through The Server Code: Navigation Aid

Some program names in the /sql directory:

Program Name	SQL statement type
-----	-----

sql_delete.cc	DELETE	
sql_do.cc	DO	
sql_handler.cc	HANDLER	
sql_help.cc	HELP	
sql_insert.cc	INSERT	// !
sql_load.cc	LOAD	
sql_rename.cc	RENAME	
sql_select.cc	SELECT	
sql_show.cc	SHOW	
sql_update.cc	UPDATE	

Question: Where will `mysql_insert()` be?

The line that we're following will take us next to a routine named `mysql_insert`. Sometimes it's difficult to guess what program a routine will be in, because MySQL has no consistent naming convention. However, here is one aid to navigation that works for some statement types. In the `sql` directory, the names of some programs correspond to statement types. This happens to be the case for `INSERT`, for instance. So the `mysql_insert` program will be in the program `sql_insert.cc`. But there's no reliable rule.

(Let's add here a few sentences about the tags 'ctags' program. When an editor supports ctags (and the list is long, but `vi` and `emacs` of course are there), the function definition is one key press away - no guessing involved. In the above case, a `vim` user could press `^]` on `mysql_insert` name and `vim` would open `sql_insert.cc` and position the cursor on the first line of the `mysql_insert()` function. The tags help can be indispensable in everyday work.)

Walking Through The Server Code: `/sql/sql_insert.cc`

```
int mysql_insert(THD *thd, TABLE_LIST *table_list, List<Item> &fields,
                List<List_item> &values_list, enum_duplicates duplic)
{
    table = open_table(thd, table_list, lock_type);
    if (check_insert_fields(thd, table, fields, *values, 1) ||
        setup_tables(table_list) ||
        setup_fields(thd, table_list, *values, 0, 0, 0))
        goto abort;
    fill_record(table->field, *values);
    error = write_record(table, &info); // !
    query_cache_invalidate3(thd, table_list, 1);
    if (transactional_table)
        error = ha_autocommit_or_rollback(thd, error);
    query_cache_invalidate3(thd, table_list, 1);
    mysql_unlock_tables(thd, thd->lock);
}
```

For the `mysql_insert` routine, we're just going to read what's in the snippet. What we're trying to do here is highlight the fact that the function names and variable names are nearly English.

Okay, we start by opening a table. Then, if a check of the fields in the `INSERT` fails, or if an attempt to set up the tables fails, or if an attempt to set up the fields fails, we'll abort.

Next, we'll fill the record buffer with values. Then we'll write the record. Then we'll invalidate the query cache. Remember, by the way, that MySQL stores frequently-used select statements and result sets in memory as an optimization, but once the insert succeeds the stored sets are invalid. Finally, we'll unlock the tables.

Walking Through The Server Code: `/sql/sql_insert.cc`

```
int write_record(TABLE *table, COPY_INFO *info)
{
    table->file->write_row(table->record[0]; // !
}
```

You can see from our marker that we're going to follow the line that contains the words 'write row'. But this is not an ordinary function call, so people who are just reading the code without the aid of a debugger can easily miss what the next point is in the line of execution here. The fact is, 'write_row' can take us to one of several different places.

Walking Through The Server Code: /sql/handler.h

```

/* The handler for a table type.
   Will be included in the TABLE structure */

handler(TABLE *table_arg) :
table(table_arg), active_index(MAX_REF_PARTS),
  ref(0), ref_length(sizeof(my_off_t)),
block_size(0), records(0), deleted(0),
  data_file_length(0), max_data_file_length(0),
index_file_length(0),
  delete_length(0), auto_increment_value(0), raid_type(0),
  key_used_on_scan(MAX_KEY),
  create_time(0), check_time(0), update_time(0), mean_rec_length(0),
  ft_handler(0)
{
}
...
virtual int write_row(byte * buf)=0;

```

To see what the `write_row` statement is doing, we'll have to look at one of the include files. In `handler.h` on the `sql` directory, we find that `write_row` is associated with a handler for a table. This definition is telling us that the address in `write_row` will vary it gets filled in at run time. In fact, there are several possible addresses.

There is one address for each handler. In our case, since we're using the default values, the value at this point will be the address of `write_row` in the MyISAM handler program.

Walking Through The Server Code: /sql/ha_myisam.cc

```

int ha_myisam::write_row(byte * buf)
{
  statistic_increment(ha_write_count, &LOCK_status);
  /* If we have a timestamp column, update it to the current time */
  if (table->time_stamp)
    update_timestamp(buf+table->time_stamp-1);
  /*
   If we have an auto_increment column and we are writing a changed row
   or a new row, then update the auto_increment value in the record.
  */
  if (table->next_number_field && buf == table->record[0])
    update_auto_increment();
  return mi_write(file, buf);    // !
}

```

And that brings us to `write_row` in the `ha_myisam.cc` program. Remember we told you that these programs beginning with the letters `ha` are interfaces to handlers, and this one is the interface to the `myisam` handler. We have at last reached the point where we're ready to call something in the handler package.

Walking Through The Server Code: /myisam/mi_write.c

```

int mi_write(MI_INFO *info, byte *record)
{
  _mi_readinfo(info, F_WRLCK, 1);
  _mi_mark_file_changed(info);
  /* Calculate and check all unique constraints */
  for (i=0 ; i < share->state.header.uniques ; i++)
  {
    mi_check_unique(info, share->uniqueinfo+i, record,
      mi_unique_hash(share->uniqueinfo+i, record),
      HA_OFFSET_ERROR);
  }
  ... to be continued in next snippet

```

Notice that at this point there is no more referencing of tables, the comments are about files and index keys. We have reached the bottom level at last. Notice as well that we are now in a C program, not a C++ program.

In this first half of the `mi_write` function, we see a call which is clearly commented. This is where checking happens for uniqueness (not the `UNIQUE` constraint, but an internal matter).

Walking Through The Server Code: `/myisam/mi_write.c`

```
... continued from previous snippet

/* Write all keys to indextree */
for (i=0 ; i < share->base.keys ; i++)
{
    share->keyinfo[i].ck_insert(info,i,buff,
        _mi_make_key(info,i,buff,record,filepos)
    )
}
(*share->write_record)(info,record);
if (share->base.auto_key)
    update_auto_increment(info,record);
}
```

In this second half of the `mi_write` function, we see another clear comment, to the effect that this is where the new keys are made for any indexed columns. Then we see the culmination of all that the last 20 snippets have been preparing, the moment we've all been waiting for, the writing of the record.

And, since the object of the `INSERT` statement is ultimately to cause a write to a record in a file, that's that. The server has done the job.

Walking Through The Server Code: Stack Trace

```
main in /sql/mysqld.cc
handle_connections_sockets in /sql/mysqld.cc
create_new_thread in /sql/mysqld.cc
handle_one_connection in /sql/sql_parse.cc
do_command in /sql/sql_parse.cc
dispatch_command in /sql/sql_parse.cc
mysql_stmt_execute in /sql/sql_prepare.cc
mysql_execute_command in /sql/sql_parse.cc
mysql_insert in /sql/mysql_insert.cc
write_record in /sql/mysql_insert.cc
ha_myisam::write_row in /sql/ha_myisam.cc
mi_write in /myisam/mi_write.c
```

And now here's a look at what's above us on the stack, or at least an idea of how we got here. We started with the main program in `mysqld.cc`. We proceeded through the creation of a thread for the client, the several junction processes that determined where we're heading, the parsing and initial execution of an SQL statement, the decision to invoke the MyISAM handler, and the writing of the row. We ended in a low level place, where we're calling the routines that write to the file. That's about as low as we should go today.

The server program would, of course, continue by returning several times in a row, sending a packet to the client saying "Okay", and ending up back in the loop inside the `handle_one_connection` function.

We, instead, will pause for a moment in awe at the amount of code we've just flitted past. And that will end our walk through the server code.

Graphic: A Chunk of MyISAM File

```
CREATE TABLE Table1 (
    column1 CHAR(1),
    column2 CHAR(1),
    column3 CHAR(1));

INSERT INTO Table1 VALUES ('a', 'b', 'c');

INSERT INTO Table1 VALUES ('d', NULL, 'e');

F1 61 62 63 00 F5 64 00 66 00 ... .abc..d e.
```

Continuing with our worm's-eye view, let's glance at the structure of a record in a MyISAM file.

The SQL statements on this graphic show a table definition and some insert statements that we used to populate the table.

The final line on the graphic is a hexadecimal dump display of the two records that we ended up with, as taken from the MyISAM file for Table1.

The thing to notice here is that the records are stored compactly. There is one byte at the start of each record F1 for the first record and F5 for the second record which contains a bit list.

When a bit is on, that means its corresponding field is NULL. That's why the second row, which has a NULL in the second column, or field, has a different header byte from the first row.

Complications are possible, but a simple record really does look this simple.

Graphic: A Chunk of InnoDB File

```
19 17 15 13 0C 06 Field Start Offsets /* First Row */
00 00 78 0D 02 BF Extra Bytes
00 00 00 00 04 21 System Column #1
00 00 00 00 09 2A System Column #2
80 00 00 00 2D 00 84 System Column #3
50 50 Field1 'PP'
50 50 Field2 'PP'
50 50 Field3 'PP'
```

If, on the other hand, you look at an InnoDB file, you'll find that it's got more complexities in the storage. The details are elsewhere in this document. But here's an introductory look.

The header here begins with offsets unlike MyISAM, which has no offsets. So you'd have to go through column 1 before getting to column 2.

Then there is a fixed header the extra bytes.

Then comes the record proper. The first fields of a typical record contain information that the user won't see, such as a row ID, a transaction ID, and a rollback pointer. This part would look different if the user had defined a primary key during the CREATE TABLE statement.

And finally there are the column contents the string of Ps at the end of the snippet here. You can see that InnoDB does more administrating.

There's been a recent change for InnoDB; what you see above is from a database made before version 5.0.

Graphic: A Packet

```
Header
Number Of Rows
ID
Status
Length
Message Content
```

Our final worm's-eye look at a physical structure will be a look at packets.

By packet, we mean: what's the format of a message that the client sends over the tcp/ip line to the server and what does the server send back?

Here we're not displaying a dump. If you want to see hexadecimal dumps of the contents of packets, this document is full of them. We're just going to note that a typical message will have a header, an identifier, and a length, followed by the message contents.

Admittedly this isn't following a standard like ISO's RDA or IBM's DRDA, but it's documented so if you want to go out and write your own type 4 JDBC driver, you've got what you need here. (Subject to license restrictions, of course.) But a word of advice on that last point: it's already been done. Mark Matthews wrote it originally, it's all in "MySQL Connector/J".

Recap

Okay, let's back up and restate. In this walkthrough, we've told you four main things.

One: How to get the MySQL source.

Two: What's in each directory in the source.

Three: The main sequence, as one walks through the server code.

Four: What physical structures look like.

We worked hard to make a description of the MySQL source that is simple, without distorting. If you were able to follow all that we've said, then that's wonderful, congratulations. If you ended up thinking that MySQL is really simple, well that's not what we wanted to convey, but we think you'll be disabused of that notion when you have a look at the code yourself.

Chapter 2. Coding Guidelines

Table of Contents

General Development Guidelines	21
C/C++ Coding Guidelines of MySQL Server	21
How we maintain the server coding guidelines	21
Indentation and Spacing	22
Naming Conventions	25
Commenting Code	25
Header Files	27
Additional suggestions	27
Suggested mode in <code>emacs</code>	29
Basic <code>vim</code> setup	30
Another <code>vim</code> setup	30
An example setup for <code>ctags</code>	30
C++ Coding Guidelines for the NDB Storage Engine	30
Braces	31
Assignment	31
Use of <code>ndbrequire</code>	31
DEBUG Tags	31

This section shows the guidelines that MySQL's developers follow when writing new code. Consistent style is important for us, because everyone must know what to expect. For example, after we become accustomed to seeing that everything inside an `if` is indented two spaces, we can glance at a listing and understand what's nested within what. Writing non-conforming code can be bad. For example, if we want to find where assignments are made to variable `mutex_count`, we might search for `mutex_count` with an editor and miss assignments that look like `mutex_count =` with a space before the equal sign (which is non-conforming). Knowing our rules, you'll find it easier to read our code, and when you decide to contribute (which we hope you'll consider!) we'll find it easier to read and review your code.

General Development Guidelines

- We use Bazaar for source management.
- You should use the TRUNK source tree (currently called "mysql-trunk") for all new developments. The public development branch can be downloaded with

```
shell> bzip2 -d mysql-trunk.tar.gz
```

- Before making big design decisions, please begin by posting a summary of what you want to do, why you want to do it, and how you plan to do it. This way we can easily provide you with feedback and also discuss it thoroughly. Perhaps another developer can assist you.

C/C++ Coding Guidelines of MySQL Server

This section covers guidelines for C/C++ code for the MySQL server. The guidelines do not necessarily apply for other projects such as MySQL Connector/J or Connector/ODBC.

How we maintain the server coding guidelines

We are committed to have a single coding style for core MySQL server. Storage engines, however, may have an own coding style: Falcon and NDB styles are documented later in this manual.

The server coding style is governed by a group of representatives from each technical team: Optimiser, Runtime, Replication, Backup Engines and the "general" team.

Currently these representatives are:

- Davi Arnaut - Runtime
- Mats Kindahl - Replication
- Chuck Bell - Backup
- Tor Didriksen - Optimizer
- Sergey Vojtovich - General

The group accepts and considers change proposals. Each proposal must include an implementation strategy, and is first published on Internals mailing list for a public discussion. When the discussion is over, the group of representatives holds a vote, and the change is accepted if it's approved by a simple majority of the ballots. The submitter of the change request then carries out its implementation.

Now to the coding style itself.

Indentation and Spacing

- For indentation use space; do not use the tab (`\t`) character. See the editor configuration tips at the end of this section for instructions on configuring a vim or emacs editor to use spaces instead of tabs.
- Avoid trailing whitespace, in code and comments.

Correct:

```
if (a)
```

Incorrect:

```
if (a)<SP><SP><TAB><SP>
```

Remove trailing spaces if you are already changing a line, otherwise leave existing code intact.

- Use line feed (`\n`) for line breaks. Do not use carriage return + line feed (`\r\n`); that can cause problems for other users and for builds. This rule is particularly important if you use a Windows editor.
- To begin indenting, add two spaces. To end indenting, subtract two spaces. For example:

```
{
  code, code, code
  {
    code, code, code
  }
}
```

- An exception to the rule above: namespaces (named or unnamed) do not introduce a new level of indentation. Example:

```
namespace foo
{
  class Bar
  {
    Bar();
  };
} // namespace foo
```

- The maximum line width is 80 characters. If you are writing a longer line, try to break it at a logical point and continue on the next line with the same indenting. Use of backslash is okay; however, multi-line literals might cause less confusion if they are defined before the function start.
- You may use empty lines (two line breaks in a row) wherever it seems helpful for readability. But never use two or more empty lines in a row. The only exception is after a function definition (see below).

- To separate two functions, use three line breaks (two empty lines). To separate a list of variable declarations from executable statements, use two line breaks (one empty line). For example:

```
int function_1()
{
    int i;
    int j;

    function0();
}

int function2()
{
    return;
}
```

- Matching '{}' (left and right braces) should be in the same column, that is, the closing '}' should be directly below the opening '{'. Do not put any non-space characters on the same line as a brace, not even a comment. Indent within braces. Exception: if there is nothing between two braces, i.e. '{}', they should appear together. For example:

```
if (code, code, code)
{
    code, code, code;
}
for (code, code, code)
{ }
```

- Indent `switch` like this:

```
switch (condition)
{
case XXX:
    statements;
case YYY:
    {
        statements;
    }
}
```

- You may align variable declarations like this:

```
Type      value;
int        var2;
ulonglong var3;
```

- When assigning to a variable, put zero spaces after the target variable name, then the assignment operator ('=' '+' etc.), then space(s). For single assignments, there should be only one space after the equal sign. For multiple assignments, add additional spaces so that the source values line up. For example:

```
a/= b;
return_value= my_function(arg1);
...
int x=          27;
int new_var=    18;
```

Align assignments from one structure to another, like this:

```
foo->member=      bar->member;
foo->name=         bar->name;
foo->name_length= bar->name_length;
```

- Put separate statements on separate lines. This applies for both variable declarations and executable statements. For example, this is wrong:

```
int x= 11; int y= 12;
```

```
z= x; y+= x;
```

This is right:

```
int x= 11;
int y= 12;

z= x;
y+= x;
```

- Put spaces both before and after binary comparison operators ('>', '==', '>=', etc.), binary arithmetic operators ('+' etc.), and binary Boolean operators ('||' etc.). Do not put spaces around unary operators like '!' or '++'. Do not put spaces around [de-]referencing operators like '->' or '[]'. Do not put space after '*' when '*' introduces a pointer. Do not put spaces after '('. Put one space after ')' if it ends a condition, but not if it ends a list of function arguments. For example:

```
int *var;

if ((x == y + 2) && !param->is_signed)
    function_call();
```

- When a function has multiple arguments separated by commas (','), put one space after each comma. For example:

```
ln= mysql_bin_log.generate_name(opt_bin_logname, "-bin", 1, buf);
```

- Put one space after a keyword which introduces a condition, such as `if` or `for` or `while`.
- After `if` or `else` or `while`, when there is only one instruction after the condition, braces are not necessary and the instruction goes on the next line, indented.

```
if (sig != MYSQL_KILL_SIGNAL && sig != 0)
    unireg_abort(1);
else
    unireg_end();
while (*val && my_isspace(mysql_d_charset, *val))
    *val++;
```

- In function declarations and invocations: there is no space between function name and '('; there is no space or line break between '(' and the first argument; if the arguments do not fit on one line then align them. Examples:

```
Return_value_type *Class_name::method_name(const char *arg1,
                                             size_t arg2, Type *arg3)
return_value= function_name(argument1, argument2, long_argument3,
                             argument4,
                             function_name2(long_argument5,
                                             long_argument6));

return_value=
    long_long_function_name(long_long_argument1, long_long_argument2,
                             long_long_long_argument3,
                             long_long_argument4,
                             long_function_name2(long_long_argument5,
                                             long_long_argument6));

Long_long_return_value_type *
Long_long_class_name::
long_long_method_name(const char *long_long_arg1, size_t long_long_arg2,
                      Long_long_type *arg3)
```

(You may but don't have to split `Class_name::method_name` into two lines.) When arguments do not fit on one line, consider renaming them.

- Format constructors in the following way:

```
Item::Item(int a_arg, int b_arg, int c_arg)
    :a(a_arg), b(b_arg), c(c_arg)
{}
```

But keep lines short to make them more readable:

```
Item::Item(int longer_arg, int more_longer_arg)
    :longer(longer_arg),
      more_longer(more_longer_arg)
{ }
```

If a constructor can fit into one line:

```
Item::Item(int a_arg) :a(a_arg) { }
```

Naming Conventions

- For identifiers formed from multiple words, separate each component with underscore rather than capitalization. Thus, use `my_var` instead of `myVar` or `MyVar`.
- Avoid capitalization except for class names; class names should begin with a capital letter.

```
class Item;
class Query_arena;
class Log_event;
```

- Avoid function names, structure elements, or variables that begin or end with `'_'`.
- Use long function and variable names in English. This will make your code easier to read for all developers.
- We used to have the rule: "Structure types are `typedef`'ed to an all-upper-case identifier." This has been deprecated for C++ code. Do not add typedefs for structs/classes in C++
- All `#define` declarations should be in upper case.

```
#define MY_CONSTANT 15
```

- Enumeration names should begin with `enum_`.
- Function declarations (forward declarations) have parameter names in addition to parameter types.

Commenting Code

- Comment your code when you do something that someone else may think is not trivial.
- Comments for pure virtual functions, documentation for API usage should be placed in front of (member, or non-member) function declarations. Description of implementation details, algorithms, anything that does not impact usage, should be put in front of implementation. Please try to not duplicate information. Make a reference to the declaration from the implementation if necessary. If the implementation and usage are too interleaved, put a reference from the interface to the implementation, and keep the entire comment in a single place.
- Class comments should be put in front of class declaration.
- When writing multi-line comments please put the `/*` and `*/` on their own lines, put the `*/` below the `/*`, put a line break and a two-space indent after the `/*`, do not use additional asterisks on the left of the comment.

```
/*
This is how a multi-line comment in the middle of code
should look. Note it not Doxygen-style if it's not at the
beginning of a code enclosure (function or class).
*/

/* ***** This comment is bad. It's indented incorrectly, it has
 *          additional asterisks. Don't write this way.
 * ***** */
```

- When writing single-line comments, the `/*` and `*/` are on the same line. For example:

```
/* We must check if stack_size = Solaris 2.9 can return 0 here. */
```

- Single-line comments like this are OK in C++

```
// We must check if stack_size = Solaris 2.9 can return 0 here.
```

- For a short comment at the end of a line, you may use either `/* ... */` or a `//` double slash. In C files or in header files used by C files, avoid `//` comments.
- Align short side `//` or `/* ... */` comments by 48 column (start the comment in column 49).

```
{ qc*= 2; /* double the estimation */ }
```

- When commenting members of a structure or a class, align comments by 48th column. If a comment doesn't fit into one line, move it to a separate line. Do not create multiline comments aligned by 48th column.

```
struct st_mysql_stmt
{
...
MYSQL_ROWS      *data_cursor;          /**< current row in cached result */
/* copy of mysql->affected_rows after statement execution */
my_ulonglong    affected_rows;
my_ulonglong    insert_id;              /**< copy of mysql->insert_id */
/*
mysql_stmt_fetch() calls this function to fetch one row (it's different
for buffered, unbuffered and cursor fetch).
*/
int              (*read_row_func)(struct st_mysql_stmt *stmt,
...
};
```

- All comments should be in English.
- Each standalone comment must start with a Capital letter.
- There is a `.` at the end of each statement in a comment paragraph (for the last one as well).

```
/*
This is a standalone comment. The comment is aligned to fit 79
characters per line. There is a dot at the end of each sentence.
Including the last one.
*/
```

- Every structure, class, method or function should have a description unless it is very short and its purpose is obvious.
- Use the below example as a template for function or method comments.
 - Please refer to the [Doxygen Manual](#) for additional information.
 - Note the IN and OUT parameters. IN is implicit, but can (but usually shouldn't) be specified with tag `@param[in]`. For OUT and INOUT parameters you should use tags `@param[out]` and `@param[in,out]` respectively.
 - Parameter specifications in `@param` section start with lowercase and are *not* terminated with a full stop/period.
 - Section headers are aligned at 2 spaces. This must be a sentence with a full stop/period at the end. If the sentence must express a subject that contains a full stop such that Doxygen would be fooled into stopping early, then use the `@brief` and `@details` to explicitly mark them.
 - Align `@retval` specifications at 4 spaces if they follow a `@return` description. Else, align at two spaces.

- Separate sections with an empty line.
- All function comments should be no longer than 79 characters per line.
- Put two line breaks (one empty line) between a function comment and its description.

```
/**
 * Initialize SHA1Context.
 *
 * Set initial values in preparation for computing a new SHA1 message digest.
 *
 * @param[in,out] context the context to reset
 *
 * @return Operation status
 *         @retval SHA_SUCCESS      OK
 *         @retval != SHA_SUCCESS  sha error Code
 */
int sha1_reset(SHA1_CONTEXT *context)
{
    ...
}
```

Header Files

- Use headers guards. Put the header guard in the first line of the header, before the copyright. Use all-uppercase name for the header guard. Derive the header guard name from the file name, and append `_INCLUDED` to create a macro name. Example: `sql_base.h` -> `SQL_BASE_INCLUDED`.
- Include directives shall be first in the file. In class implementation, include the header file with class declaration before all other header files, to make sure that the header is self-sufficient.
- Every header file should be self-sufficient in the sense that for a header file `my_header.h`, the following should compile without errors:

```
#include "my_header.h"
```

An exception is made for generated files, for example, those generated by Yacc and Lex, since it is not possible to re-write the generators to produce "correct" files.

Additional suggestions

- Try to write code in a lot of black boxes that can be reused or at least use a clean, easy to change interface.
- Reuse code; There are already many algorithms in MySQL that can be reused for list handling, queues, dynamic and hashed arrays, sorting, etc.
- Use the `my_*` functions like `my_read()`/`my_write()`/`my_malloc()` that you can find in the `mysys` library, instead of the direct system calls; This will make your code easier to debug and more portable.
- Use `libstring` functions (in the `strings` directory, declared in `include/m_string.h`) instead of standard `libc` string functions whenever possible.
- Try to always write optimized code, so that you don't have to go back and rewrite it a couple of months later. It's better to spend 3 times as much time designing and writing an optimal function than having to do it all over again later on.
- Avoid CPU wasteful code, even when its use is trivial, to avoid developing sloppy coding habits.
- If you can do something in fewer lines, please do so (as long as the code will not be slower or much harder to read).

- Do not check the same pointer for `NULL` more than once.
- Never use a macro when an (inline) function would work as well.
- Do not make a function inline if you don't have a very good reason for it. In many cases, the extra code that is generated is more likely to slow down the resulting code than give a speed increase because the bigger code will cause more data fetches and instruction misses in the processor cache.

It is okay to use inline functions are which satisfy most of the following requirements:

- The function is very short (just a few lines).
- The function is used in a speed critical place and is executed over and over again.
- The function is handling the normal case, not some extra functionality that most users will not use.
- The function is rarely called. (This restriction must be followed unless the function translates to fewer than 16 assembler instructions.)
- The compiler can do additional optimizations with inlining and the resulting function will be only a fraction of size of the original one.
- Think assembly - make it easier for the compiler to optimize your code.
- Avoid using `malloc()`, which is very slow. For memory allocations that only need to live for the lifetime of one thread, use `sql_alloc()` instead.
- All functions that can report an error (usually an allocation error), should return `0/FALSE/false` on success, `1/TRUE/true` on failure. Other return values should go in an output argument. If you have a predicate function which returns `bool`, and cannot fail, document that fact clearly (in the header file). Recommendation when writing new code: use return type `int`, to distinguish from predicate functions. Returning `true` on error, allows us to write:

```
if (a() || b() || c())
    error("something went wrong");
```

However, short-circuit evaluation like that above is not the best method for evaluating options.

- Beware of truncation when returning `TRUE/true` to indicate an error:

```
my_bool foo(int val)      { return val; } /* Bad. */
int      foo(longlong val){ return val; } /* Bad. */
```

```
my_bool foo(int val)      { return test(val); } /* Good. */
int      foo(longlong val){ return test(val); } /* Good. */
```

- Using `goto` is okay if not abused.
- If you have an 'if' statement that ends with a 'goto' or 'return' you should NOT have an else statement:

```
if (a == b)
    return 5;
else return 6;
```

->

```
if (a == b)
    return 5;
return 6;
```

- Avoid default variable initializations. Use `LINT_INIT()` if the compiler complains after making sure that there is really no way the variable can be used uninitialized.
- In C code, use `TRUE` and `FALSE` rather than `1/0`

- In C++ code, it is OK to use `true` and `false` (do not use `1/0`). You can use C++ `bool/true/false` when calling C functions (values will be safely promoted to `my_bool`).
- `bool` exists only in C++. In C, you have to use `my_bool` (which is `char`); it has different cast rules than `bool`:

```
int c= 256*2;
bool a= c;          /* a gets 'true' */
my_bool b= c;       /* b gets zero, i.e. 'false': BAD */
my_bool b= test(c); /* b gets 'true': GOOD */
```

- Do not instantiate a class if you do not have to.
- Use pointers rather than array indexing when operating on strings.
- Never pass parameters with the `&variable_name` construct in C++. Always use a pointer instead!

The reason is that the above makes it much harder for the one reading the caller function code to know what is happening and what kind of code the compiler is generating for the call.

- Do not use the `%p` marker of `printf()` (`fprintf()`, `vprintf()`, etc) because it leads to different outputs (for example on some Linux and Mac OS X the output starts with `0x` while it does not on some Solaris). In MySQL 5.5 and later, use `my_vsnprint`, `DEBUG_PRINT` with `%p` for pointer formatting consistent across different platforms. In earlier versions, use printf-family functions with `0x%lx`, but beware it truncates pointers on 64-bit Windows. Being sure that there is always `0x` enables us to quickly identify pointer values in the `DEBUG` trace.
- Relying on loop counter variables being local to the loop body if declared in the `for` statement is not portable. Some compilers still don't implement this ANSI C++ specification. The symptom of such use is an error like this:

```
c-1101 CC: ERROR File = listener.cc, Line = 187
      "i" has already been declared in the current scope.

      for (int i= 0; i < num_sockets; i++)
```

Suggested mode in `emacs`

```
(require 'font-lock)
(require 'cc-mode)
(setq global-font-lock-mode t) ;;colors in all buffers that support it
(setq font-lock-maximum-decoration t) ;;maximum color
(c-add-style "MY"
  '("K&R"
    (c-basic-offset . 2)
    (c-comment-only-line-offset . 0)
    (c-offsets-alist . ((statement-block-intro . +)
                        (knr-argdecl-intro . 0)
                        (substatement-open . 0)
                        (label . -)
                        (statement-cont . +)
                        (arglist-intro . c-lineup-arglist-intro-after-paren)
                        (arglist-close . c-lineup-arglist)
                        (innamespace . 0)
                        (inline-open . 0)
                        (statement-case-open . +)
                        ))
  ))

(defun mysql-c-mode-hook ()
  (interactive)
  (require 'cc-mode)
  (c-set-style "MY")
  (setq indent-tabs-mode nil)
  (setq comment-column 48))

(add-hook 'c-mode-common-hook 'mysql-c-mode-hook)
```

Basic vim setup

```
set tabstop=8
set shiftwidth=2
set backspace=2
set softtabstop
set smartindent
set cindent
set cinoptions=g0:0t0c2C1(0f0l1
set expandtab
```

Another vim setup

```
set tabstop=8
set shiftwidth=2
set bs=2
set et
set sts=2
set tw=78
set formatoptions=cqroal
set cinoptions=g0:0t0c2C1(0f0l1
set cindent

function InsertShiftTabWrapper()
  let num_spaces = 48 - virtcol('.')
  let line = ' '
  while (num_spaces > 0)
    let line = line . ' '
    let num_spaces = num_spaces - 1
  endwhile
  return line
endfunction
" jump to 48th column by Shift-Tab - to place a comment there
inoremap <S-tab> <c-r>=InsertShiftTabWrapper()<cr>
" highlight trailing spaces as errors
let c_space_errors=1
```

An example setup for ctags

Put this configuration into your ~/.ctags file:

```
--c++-kinds=+p
--fields=+iaS
--extra=+q
--langdef=errmsg
--regex-errmsg=/^(ER_[A-Z0-9_]+)/\1/
--langmap=errmsg:(errmsg*.txt),c:+.ic,yacc:+.yy
```

C++ Coding Guidelines for the NDB Storage Engine

The mysql handler part of NDB ([ha_ndbcluster.cc](#), [ha_ndbcluster_binlog.cc](#), etc.) uses the same coding style as the rest of the mysql code.

The non-mysql part of NDB code has a long history, and use a multitude of coding styles. When modifying and extending existing source files or modules, the coding style already used in that code should be followed in terms of indentations, naming conventions, etc. For completely new code, the mysql conventions (with exceptions below) should probably be followed.

Do *not* do any change to NDB code purely for the sake of changing from one formatting style to another. It just causes merge annoyances and makes patches harder to read, and we do not expect the style to ever become 100% consistent across all of the source code. It is however ok to fix inconsistent style in lines that are changed for other reasons.

One convention that should be followed for all new or modified code, in both mysql and non-mysql parts of the code, is that class member variables should be named with lowercase words separated by underscores '_ ', and pre-fixed with 'm_ '. Like this:

```
const char *m_my_class_member;
```

Braces

if, while, etc **must** always have braces.

eg. Good

```
if (a == b)
{
    dosomething();
}
```

Braces should be on separate line like above.

e.g BAD

```
if (a == b) {
    dosomething();
}
```

Inline methods inside class(struct) is ok to write like below, (i.e opening brace is on same line as function declaration)

```
struct A
{
    A() {
    }
}
```

Assignment

```
a = 3; // ok
a= 3; // not ok
```

Use of ndbrequire

In the NDB kernel code, the `ndbrequire()` facility has historically been widely used. However, most of this is now considered misuse, and use of `ndbrequire` should generally be avoided. Over time, we want to remove most or all `ndbrequires`.

There are three different classes of `ndbrequire()` usage, with corresponding replacement as follows:

- Verification of code logic, hitting this is a real bug, and the error message should be accordingly. For this one option is `ndbassert()` (only enabled in debug builds), or we might need to add `ndbchecklogic()` or similar.
- Hitting a configurable limit, which cannot be handled gracefully. For this one should use `ndbrequireErr()`. The error message should suggest config change to correct the problem, or refer to a section in the manual to read more.
- Hitting hardcoded limits; we should really try to avoid this, but if it is unavoidable, or if it is a limit we think we will never hit, use `ndbrequireErr()` and add appropriate error message.

DEBUG Tags

The full documentation of the DEBUG library is in files `debug/user.*` in the MySQL source tree. Here are some of the DEBUG tags we now use:

- `enter`

Arguments to the function.

- `exit`

Results from the function.

- `info`

Something that may be interesting.

- `warning`

When something doesn't go the usual route or may be wrong.

- `error`

When something went wrong.

- `loop`

Write in a loop, that is probably only useful when debugging the loop. These should normally be deleted when you are satisfied with the code and it has been in real use for a while.

Some tags specific to `mysqld`, because we want to watch these carefully:

- `trans`

Starting/stopping transactions.

- `quit`

`info` when `mysqld` is preparing to die.

- `query`

Print query.

Chapter 3. Reusable Classes and Templates

Table of Contents

Containers	33
Array	33
I_P_List	33
I_List	33
Memory Management	33
MEM_ROOT	33
How to Extend This Page	33

Containers

This section is yet to be written.

Array

This section is yet to be written.

I_P_List

This section is yet to be written.

I_List

This section is yet to be written.

Memory Management

This section is yet to be written.

MEM_ROOT

This section is yet to be written.

How to Extend This Page

For each new data structure, please make sure you add:

- the important trade-offs, such as speed vs. space, etc.
- best practices for use

Chapter 4. How to Build MySQL Server with CMake

Table of Contents

Introduction	35
Prerequisites	35
How to install CMake	35
Very quick how-to-build	36
Build using the same options as in MySQL official release	37
Long description of how-to build	37
Create the build directory	37
Configuration step	37
Listing configuration parameters	38
Changing configuration parameters	38
Building Debug Configurations	39
./configure emulation	41
For Developers: how to write platform checks	42
FAQ / Miscellany	45

Introduction

This page describes how to build MySQL distributions with CMake. Other resources that you might find useful:

- [Chapter 5, Autotools to CMake Transition Guide](#) : If you have previously built MySQL using the GNU autotools, this guide shows how to map common autotools/configure options to CMake.
- [CMake documentation](#) at cmake.org
- [CMake Useful Variables](#)

Prerequisites

- Usual MySQL build prerequisites:
 - Unix: compiler and make utility, curses dev package on Linux
 - Windows: Visual Studio (Express version okay)
 - Mac OS X: Xcode tools
 - Everywhere: bison, unless you're building from source package where bison output is already packed. On OpenSolaris/Solaris Express you would need to install m4 in addition to bison. NOTE: On Windows install bison into path without spaces, not into default location. Reason: [this bug](#) in bison 2.4.1
- CMake version 2.6.3 or later installed on your system.

How to install CMake

- Debian/Ubuntu Linux:

```
sudo apt-get install cmake cmake-gui
```

- Fedora Linux

```
sudo yum install cmake cmake-gui
```

- openSUSE Linux

```
sudo zypper install cmake cmake-gui
```

- Gentoo Linux

```
sudo emerge cmake
```

- OpenSolaris:

```
pfexec pkgadd install SUNWcmake
```

- Windows:

Download and install the latest distribution from <http://www.cmake.org/cmake/resources/software.html>. Download the installer .exe file and run it.

- Mac OS X:

Download and install the latest distribution from <http://www.cmake.org/cmake/resources/software.html>. Download the .dmg image and open it. Alternatively, if you have Darwinports installed, you can install the CMake port. The +gui variant causes cmake-gui to be installed as well. Omit +gui if you do not want to install cmake-gui.

```
port install cmake +gui
```

- Other Unixes:

Precompiled packages for other Unix flavors (HPUX, AIX) are available from <http://www.cmake.org/cmake/resources/software.html>

Alternatively, you can build from source. A source package is also available from the CMake download page.

Very quick how-to-build

This section describes how to build MySQL in release mode (with debug info)

- Unix (Makefiles)

To control which compiler is chosen, set the CC and CXX environment variables to point to C and C++ compilers. (This is optional, but if you have different versions of the compilers, it gives better control.)

```
mkdir bld
cd
bld cmake ..
make
```

- Windows (Visual Studio, from command line)

```
mkdir bld
cd
bld cmake ..
devenv mysql.sln /build relwithdebinfo
```

- Build Debug configuration

- Unix (Makefiles)

```
mkdir bld_debug
cd
bld_debug cmake .. -DCMAKE_BUILD_TYPE=Debug
```

- Visual Studio (using command line)

```
devenv mysql.sln /build debug
```

Build using the same options as in MySQL official release

Official MySQL releases add some compiler options. Also, some storage engines are linked statically into mysqld (for example, ARCHIVE). These build options for official releases are stored in `cmake/build_configurations/mysql_release.cmake`. To use the options, use the `-DBUILD_CONFIG=mysql_release` cmake parameter.

```
mkdir bld
cd
bld cmake -DBUILD_CONFIG=mysql_release ..
```

- Unix (Makefiles)

```
make
```

- Visual Studio (using command line)

```
devenv mysql.sln /build relwithdebinfo
```

Note that on Linux, the official release requires `libaio` to be installed on the build machine. For example:

- RedHat/Fedora

```
sudo yum install libaio-devel
```

- Debian/Ubuntu

```
sudo apt-get install libaio-dev
```

Long description of how-to build

Ensure that your compiler and cmake are in your PATH setting. The following description assumes that the current working directory is the top-level source directory.

Create the build directory

One of the nice CMake features is "out-of-source" build support, which means not building in the source directory, but in a dedicated build directory. This keeps the source directory clean and allows for more than a single build tree from the same source tree (e.g., debug and release, 32-bit and 64-bit, etc.). We'll create a subdirectory "bld" in the source directory for this purpose.

```
mkdir bld
cd
bld
```

Configuration step

On Unix machine, configure the build with

```
cmake ..
```

On Windows machine, to build with VS2008 and x64

```
cmake .. -G "Visual Studio 9 2008 Win64"
```

On OS X, if you want to use the Xcode IDE

```
cmake .. -G Xcode
```

You can add configuration parameters (see next section for description), e.g

```
cmake .. -DWITH_EMBEDDED_SERVER=1
```

Now, CMake runs system checks and generates Makefiles. CMake allows the configuration process to be iterative, so you can add more parameters after initial config has run. For example, running the following command after the initial preceding configuration step would add ARCHIVE to the list of statically compiled storage engines:

```
cmake .. -DWITH_ARCHIVE_STORAGE_ENGINE=1
```

System checks do not rerun after the initial configuration completes.

Listing configuration parameters

After the initial configuration step completes you can use

- short form

```
cmake . -L
```

- short form plus description

```
cmake . -LH
```

- long form (lists lots of parameters, including internal and advanced ones)

```
cmake . -LA
```

Better even, if you have cmake-gui installed, you can do

```
cmake-gui .
```

and see or change parameters here. On Unix, some people like to use ccmake (Curses based GUI for cmake):

```
ccmake .
```

Changing configuration parameters

The procedure above will build with default configuration. This configuration is likely not the most perfect for your needs: For example, the embedded library is not produced. Assume that you want to change the configuration parameters and compile embedded.

- You can provide parameters on the command line, like

```
cmake . -DWITH_EMBEDDED_SERVER=1
```

This can be done during the initial configuration or any time later. Note, that parameters are "sticky", that is they are remembered in the CMake cache (CMakeCache.txt file in the build directory)

- Configuration using cmake-gui (Windows, OS X, or Linux with cmake-gui installed)

From the build directory, issue

```
cmake-gui .
```

- Check the WITH_EMBEDDED_SERVER checkbox
- Click the "Configure" button
- Click the "Generate" button
- Close cmake-gui

Building Debug Configurations

Using Makefiles, debug build is done with `-DCMAKE_BUILD_TYPE=Debug` (shortcut for it is `-DWITH_DEBUG=1`). this would include `DEBUG` instrumentation, plus wrapper around pthread mutexes known as `SAFE_MUTEX` on Unixes.

If Visual Studio or XCode generators are used (you called `cmake` with `-G "Visual Studio ..."` or `-G Xcode`), then switching to release or debug configuration is done within IDE, or at the build time using command line switches, e.g

```
devenv MySQL.sln /build debug
```

Building With CMake

- Unix

```
make
```

Note: by default, `cmake` build is less verbose than `automake` build. Use

```
make VERBOSE=1
```

if you want to see how compiler is invoked.

- Windows (using "Visual Studio 9 2008" generator)

```
devenv MySQL.sln /build RelWithDebInfo
```

(alternatively, open `MySQL.sln` and build using the IDE)

- Mac OS X build with Xcode

```
xcodebuild -configuration RelWithDebInfo
```

(alternatively, open `MySQL.xcodeproj` and build using the IDE)

- Command line build with CMake 2.8

After creating project with `cmake` as above, issue

```
cmake --build .
```

this works with any CMake generator.

For Visual Studio and Xcode, you might want to add extra configuration parameters, to avoid building all configurations.

```
cmake --build . --config RelWithDebInfo
```

Build types

CMake has `CMAKE_BUILD_TYPE` variable for predefined build types. A build type affects optimization and whether the result of the build is debuggable.

The ones used by MySQL are `RelWithDebInfo` or `Debug`.

- `RelWithDebInfo` (optimizations are on, debug info is generated) is used in MySQL by default.
- `Debug` (optimizations are off, debug info is generated) is used if `WITH_DEBUG` variable is set.
- `CMAKE_BUILD_TYPE` is not set when custom compile flags are used (see next section)

How to control compiler flags

To specify your own compiler flags, in case compiler flags do not affect optimization you can

- Set environment variables: CFLAGS, CXXFLAGS
- Use CMake options: `cmake . -DCMAKE_C_FLAGS=your_c_flags -DCMAKE_CXX_FLAGS=your_cxx_flags`

When providing your own compiler flags, you might want to specify CMAKE_BUILD_TYPE as well.

For example, to create a 32-bit release build on a 64-bit Linux machine , you do:

```
cmake -DCMAKE_C_FLAGS=-m32 --DCMAKE_CXX_FLAGS=-m32 \
-DCMAKE_BUILD_TYPE=RelWithDebInfo.
```

If flags you set that affect optimization (`-Onumber`), then to specify a different optimization for default build (Relwithdebinfo), you need to set CMAKE_C_FLAGS_RELWITHDEBINFO and/or CMAKE_CXX_RELWITHDEBINFO. For example, to compile on Linux with -O3 and with debug symbols, do :

```
cmake "-DCMAKE_C_FLAGS_RELWITHDEBINFO=-O3 \
-g" "-DCMAKE_CXX_FLAGS_RELWITHDEBINFO=-O3 -g"
```

Predefined sets of options and compiler flags

It might be handy to specify a predefined set of options and do some compiler flag adjustments by passing just a single parameter to cmake. For MySQL ,this can be done using `cmake -DBUILD_CONFIG=some_config`. When set, cmake will execute script in `cmake/build_configurations/some_config.cmake`. Assuming we want to include embedded and exclude archive storage engine from build, this script could look like

```
SET(WITH_EMBEDDED_SERVER 1 CACHE BOOL "")
SET(WITHOUT_ARCHIVE_STORAGE_ENGINE 1 CACHE BOOL "")
```

Currently, there is just a single predefined configuration `mysql_release`, it reflects configuration parameters and compiler flags used by MySQL releases.

Creating binary packages

Packaging in form of tar.gz archives or .zip on Windows

1)If you're using "generic" Unix build with makefiles

```
make package
```

2)On Windows, using "Visual Studio" generator

```
devenv mysql.sln /build relwithdebinfo /project package
```

On Windows, current versions of CMake (2.8 and later) do not need any external tools to generate ZIP, with CMake 2.6 however 7Zip or Winzip must be installed and 7z.exe resp winzip.exe need to be in the PATH.

Another way to build packages is calling `cpack` executable directly like

```
cpack -G TGZ --config CPackConfig.cmake
```

(`-G TGZ` is for tar.gz generator, there is also `-GZIP`)

make install, make test

`install` target also provided for Makefile based generators. Installation directory can be controlled using configure-time parameter `CMAKE_INSTALL_PREFIX` (default is `/usr`). It is also possible to install to non-configured directory, using

```
make install DESTDIR="/some/absolute/path"
```

- "make test" runs unit tests (uses CTest for it)
- "make test-force" runs mysql-test-run.pl tests with --test-force parameter

Fine-tuning installation paths

IF you come from autotools background, you will be familiar with --bindir, --libdir, --sbindir etc parameters passed to configure script that allow for fine tuning the installation layout. A similar functionality is available with CMake build too.

- CMAKE_INSTALL_PREFIX: specifies the "root" directory of the installation, same as autotools --prefix
- INSTALL_BINDIR, INSTALL_SBINDIR, INSTALL_LIBDIR: correspond to autotools

--bindir, --sbindir, --libdir parameters. A subtle difference is that INSTALL_XXDIR should be paths relative to CMAKE_INSTALL_PREFIX, e.g INSTALL_BINDIR should be "bin" rather than "/usr/bin".

- there is INSTALL_LAYOUT parameter that allows to choose one of several predefined installation layouts
 - STANDALONE with layout is the same as in tar.gz/zip packages
 - RPM with layout similar to RPM packages - for example mysqld is in sbin subdirectory.
 - SVR4 - Solaris package layout
 - DEB (experimental)- Layout as in DEB package

Default layout is STANDALONE.

Here is an example on how to modify STANDLONE layout slightly and install libraries into "lib64" subdirectory instead of default "lib"

```
cmake . -DCMAKE_INSTALL_LAYOUT=STANDALONE -DCMAKE_INSTALL_LIBDIR=lib64
```

Packager-friendly build options (Unix)

MySQL source distribution contains sources for zlib (compression library), YaSSL (ssl library), readline and libedit. MySQL can be compiled using either libraries available on the system or, to minimize external dependencies, with bundled sources. For Unix/Linux packagers, using system libraries is a more natural option and CMake build has support for it, using options below

- -DWITH_ZLIB=system (link with system libz.so)
- -DWITH_SSL=system (link with system libssl.so, libcrypto.so)
- -DWITH_READLINE=system (link with system libreadline.so)
- On Linux, --Wl,--as-needed link option can also be used to remove unused dependencies. While CMake build tries to avoid unneeded dependencies, --as-needed brings better results, for example it removes unused dependency on libgcc_s.so
- --Wl,--no-undefined can *not* be used at the moment if plugins are built, because plugins have direct dependency (use symbols) exported by MySQL server.

./configure emulation

The legacy way to build MySQL on Unix was to run

```
BUILD/autorun.sh; ./configure lots of parameters; make
```

This will still work, however `./configure` created by `./BUILD/autorun.sh` is just a wrapper that translates old-style autotools parameters to new style cmake parameters. Beware that the script is neither perfect nor supported. It is meant to be a temporary solution for those who need time to rewrite `./configure` based scripts to native CMake.

Instead of running `BUILD/autorun.sh`, one can directly invoke `./cmake/configure.pl`

For Developers: how to write platform checks

If you modify MySQL source and want to add a new platform check, please read http://www.vtk.org/Wiki/CMake_HowToDoPlatformChecks first. In MySQL, most of the platform tests are implemented in `configure.cmake` and the template header file is `config.h.cmake`

Bigger chunks of functionality, for example, non-trivial macros, are implemented in files `src-root/cmake` subdirectory.

For people with autotools background, it is important to remember CMake does not provide autoheader functionality. That is, when you add a check

```
CHECK_FUNCTION_EXISTS(foo HAVE_FOO)
```

to `config.cmake`, then you will also need to add

```
#cmakedefine HAVE_FOO 1
```

to `config.h.cmake`

Useful bits:

- Check for existence of C/C++ compiler flags with `CHECK_{C,CXX}_COMPILER_FLAG`.

Here is an example of checking for (theoretical) `-foo` flag support in C compiler, and adding it to C flags, if the flag is supported.

```
INCLUDE(CheckCCompilerFlag)
CHECK_C_COMPILER_FLAG("-foo" HAVE_C_COMPILER_FOO)
IF(HAVE_C_COMPILER_FOO)
  SET(CMAKE_C_FLAGS "${CMAKE_C_FLAGS} -foo")
ENDIF()
```

Debug-only options

Sometimes, it is handy to add an option that is active only in Debug builds. When doing this one should keep in mind, that tests like `IF(WITH_DEBUG)` or `IF(CMAKE_BUILD_TYPE MATCHES "Debug")` do not work as expected. First, while `WITH_DEBUG` is an alias for `CMAKE_BUILD_TYPE=Debug`, the converse is not true.

Second, checking for `CMAKE_BUILD_TYPE` will not work everywhere, more precisely, it will **not** work with multi-configuration CMake generators, i.e. neither on Windows with Visual Studio and nor on OSX with Xcode.

So, when adding debug-only option consider extending `CMAKE_{C,CXX}_FLAGS_DEBUG` like for example:

```
# Works always
SET(CMAKE_C_FLAGS_DEBUG "${CMAKE_C_FLAGS_DEBUG} -DUNIV_DEBUG")
```

and do **NOT** do it like:

```
IF(WITH_DEBUG)
  # Does NOT work with CMAKE_BUILD_TYPE=Debug, Visual Studio or Xcode
  ADD_DEFINITIONS(-DUNIV_DEBUG)
ENDIF()
```


Adding platform checks/compiler flags for a specific OS

If you add a platform check for specific OS or want to modify compiler flags, rather than introducing `IF(CMAKE_SYSTEM_NAME MATCHES...)` in `configure.cmake`, add them to the appropriate section in `cmake/os/my_platform.cmake`. For example, Solaris specific adjustments are made in `cmake/os/SunOS.cmake`. This file will be included when you compile on Solaris.

Troubleshooting platform checks or configure errors

If you suspect that a platform check returned wrong result, examine `build-root/CMakeFiles/CMakeError.log` and `build-root/CMakeFiles/CMakeOutput.log`. These files contain compiler command line, and exact error messages.

Troubleshooting CMake code

While there are advanced flags for `cmake` like `-debug-trycompile` and `--trace`, a simple and efficient way to debug is to add `MESSAGE("interesting variable=${some_invariable}")` to the interesting places in `CMakeLists.txt`.

Tips for developers

- How to find out which compiler/linker flags are used

When using Makefile generator it is easy to examine which compiler flags are used to build. For example, compiler flags for `mysqld` are in `build-root/sql/CMakeFiles/mysqld.dir/flags.make` and the linker command line is in `build-root/sql/CMakeFiles/mysqld.dir/ulink.txt`.

- What is `CMakeCache.txt`?

`CMake` caches results of platform checks in `CMakeCache.txt`. It is a nice feature because tests do not rerun when reconfiguring (e.g. when a new test was added). The downside of caching is that when a platform test was wrong and was later corrected, the cached result is still used. If you encounter this situation, which should be a rare occasion, you need either to remove the offending entry from `CMakeCache.txt` (if test was for `HAVE_FOO`, remove lines containing `HAVE_FOO` from `CMakeCache.txt`) or just remove the cache file.

MySQL-Specific CMake Macros

MYSQL_ADD_EXECUTABLE

Almost the same as `ADD_EXECUTABLE`. Supports optional `DESTINATION` parameter which tells where to install the exe (if not specified, it goes to `${INSTALL_BINDIR}` directory). For executables not intended to be installed, use `ADD_EXECUTABLE` instead. On Windows, signs the executable if `SIGNCODE` option is set to `TRUE`.

Example usage

```
MYSQL_ADD_EXECUTABLE(mysqld ${MYSQLD_SOURCE} \
DESTINATION ${INSTALL_SBINDIR})
```

MYSQL_ADD_PLUGIN - build mysql plugin

```
MYSQL_ADD_PLUGIN(plugin_name source1...sourceN
[STORAGE_ENGINE]
[MANDATORY|DEFAULT]
[STATIC_ONLY|MODULE_ONLY]
[MODULE_OUTPUT_NAME module_name]
[STATIC_OUTPUT_NAME static_name]
[RECOMPILE_FOR_EMBEDDED]
[LINK_LIBRARIES lib1...libN]
[DEPENDENCIES target1...targetN])
```

Parameters:

- STORAGE_ENGINE

Define for storage engine. Causes shared library be built with ha_ prefix.

- MANDATORY

Define for mandatory plugins (like myisam). Causes plugin to be always built

- DEFAULT

Default plugin. Built unless WITHOUT_*plugin_name* option is defined. Note: innobase storage engine has this option starting with MySQL 5.5.5

- STATIC_ONLY

Can be only built as static library

- MODULE_ONLY

Can be only built as shared module

- MODULE_OUTPUT_NAME module_name

Defines plugin library name when it is built as shared module.

- STATIC_OUTPUT_NAME

Defines library name when it is built as static library.

- RECOMPILE_FOR_EMBEDDED

Needs to be recompiled with -DEMBEDDED_SERVER preprocessor flag for use with embedded server. Only few plugins need this - typically mandatory storage engines that depend on internal structures and on EMBEDDED_SERVER flag.

- LINK_LIBRARIES

Libraries to link with plugin

- DEPENDENCIES

Plugin dependencies

Example 1 - Simple plugin that is only built as shared module

```
MYSQL_ADD_PLUGIN(daemon_example daemon_example.cc MODULE_ONLY)
```

Example 2 - Innobase plugin. Storage engine, redefines output name of shared library to be ha_innodb rather than ha_innobase, depends on zlib library.

```
MYSQL_ADD_PLUGIN(innobase ${INNOBASE_SOURCES} STORAGE_ENGINE MODULE_OUTPUT_NAME ha_innodb LINK_LIBRARIES $
```

Interface to Third-Party Tools

Third-party tools that need to determine the MySQL version from the MySQL source can read the VERSION file in the top-level source directory. The file lists the pieces of the version separately. For example, if the version is 5.5.8, the file looks like this:

```
MYSQL_VERSION_MAJOR=5
MYSQL_VERSION_MINOR=5
MYSQL_VERSION_PATCH=8
MYSQL_VERSION_EXTRA=
```

If the source is not for a General Availability (GA) release, the MYSQL_VERSION_EXTRA value will be nonempty. For example, the value for a Release Candidate release would look like this:

```
MYSQL_VERSION_EXTRA=rc
```

To construct a five-digit number from the version components, use this formula:

```
MYSQL_VERSION_MAJOR*10000 + MYSQL_VERSION_MINOR*100 + MYSQL_VERSION_PATCH
```

FAQ / Miscellany

Running mysql-test-run.pl in out-of-source build

When using out-of-source build, use mysql-test-run.pl in the `builddir`/mysql-test. It is a wrapper script that calls mysql-test-run.pl in the source directory and tells it where to look for the binaries, via environment MTR_BINDIR variable. Attempts to run mysql-test-run.pl from the source directory will fail.

Running mysql-test-run.pl with Visual Studio or Xcode projects

If you build with Xcode, and you build more than a single configuration (e.g Debug and RelWithDebInfo), set environment variable MTR_VS_CONFIG=[cmake_configuration_name](#) to run tests for a specific configuration. The name of the variable, specially "VS" part in it just reflects the fact it was implemented for Visual Studio originally. When many configurations are build, MTR will prefer Release or RelWithDebInfo. To run debug configuration:

- On Mac OS X

```
cd builddir/mysql-test
MTR_VS_CONFIG=Debug perl mysql-test-run.pl parameters
```

- On Windows

```
cd builddir\mysql-test
set MTR_VS_CONFIG=Debug
perl mysql-test-run.pl parameters
```

make distclean

Unlike autotools, CMake [does not provide "distclean" target natively](#), nor there should be a need to use it, if you build out-of-source. But if you built in-source, use "bzip clean-tree" with --unknown and/or --ignored arguments. If you want to add new files to the tree, be sure to "bzip add" prior to "bzip clean-tree".

Compiling for different hardware architectures

- GCC (on Linux) or Sun Studio

Use compile option -m32 (force 32-bit build), -m64 (force 64-bit build)

- Windows, Visual Studio generator

Use cmake -G "Visual Studio 9 2008 Win64" [path_to_source_dir](#) to compile 64-bit (x64)

- Mac OS X

Use CMAKE_OSX_ARCHITECTURES CMake variable. You can set more than a single architecture to create universal binary, e.g

```
cmake "-DCMAKE_OSX_ARCHITECTURES=i386;ppc" path\_to\_source
```

will build universal binary with 32-bit intel / 32-bit powerpc.

```
cmake "-DCMAKE_OSX_ARCHITECTURES=x86_64" path\_to\_source
```

will create x86_64 binary.

Chapter 5. Autotools to CMake Transition Guide

Table of Contents

Introduction	47
Command Invocation Syntax	47
Installation Layout Options	47
Storage Engine Options	48
Library Options	49
Miscellaneous Options	49
Debugging the Configuration Process	50
Interface to Third-Party Tools	50

Introduction

This page shows some common autotools configuration options and how they map to CMake equivalents. It supplements the [CMake section](#) that describes how to build MySQL with CMake.

Command Invocation Syntax

The following table shows some common configure invocation syntax and the equivalent CMake commands. The "." should be replaced with the path to the top-level directory of the source tree if that directory is not your current working directory.

configure Command	CMake Command
./configure	cmake .
./configure --help	cmake . -LH or cmake .

To clean out old object files and cached information before reconfiguring and rebuilding:

Autotools:

```
make clean rm config.cache
```

CMake (Unix):

```
make clean rm CMakeCache.txt
```

CMake (Windows):

```
devenv MySQL.sln /clean del CMakeCache.txt
```

Installation Layout Options

These options control where to install various MySQL components.

In the following table, the CMAKE_INSTALL_PREFIX value is the installation base directory. Many other CMake layout options are interpreted relative to the prefix and their values are relative pathnames. Their values should not include the prefix.

Parameter	configure Option	CMake Option	CMake Notes
Installation base directory	--prefix=/usr	- DCMAKE_INSTALL_PREFIX=/usr	
mysqld directory	--libexecdir=/usr/sbin	- DINSTALL_SBINDIR=sbin	interpreted relative to prefix

Data directory	--localstatedir=/var/lib/mysql	-DMYSQL_DATADIR=/var/lib/mysql	
Config directory (for my.cnf)	--sysconfdir=/etc/mysql	-DSYSCONFDIR=/etc/mysql	
Plugin directory	--with-pluginlibdir=/usr/lib64/mysql/plugin	-DINSTALL_PLUGINLIBDIR=lib64/mysql/plugin	interpreted relative to prefix
Man page directory	--mandir=/usr/share/man	-DINSTALL_MANDIR=share/man	interpreted relative to prefix
Shared-data directory	--sharedstatedir=/usr/share/mysql	-DINSTALL_SHAREDIR=share	this is where a local/mysql.m4 should be installed
Library installation directory	--libdir=/usr/lib64/mysql	-DINSTALL_LIBDIR=lib64/mysql	interpreted relative to prefix
Header installation directory	--includedir=/usr/include/mysql	-DINSTALL_INCLUDEDIR=include/mysql	interpreted relative to prefix
Info doc directory	--infodir=/usr/share/info	-DINSTALL_INFODIR=share/info	interpreted relative to prefix

Storage Engine Options

Storage engines are plugins, so the options that control plugin building specify which storage engines to build.

The `--with-plugins` configure option accepts two constructs that have no direct equivalent in CMake:

- `--with-plugins` accepts a comma-separated list of engine names
- `--with-plugins` accepts a "group name" value that is shorthand for a set of engines

With CMake, engines are controlled with individual options.

Suppose that the configure option is:

```
--with-plugins=csv,mysam,mysammrg,heap,innobase,archive,blackhole
```

This builds the named engines as static plugins that are compiled into the server and need not be installed explicitly.

To convert this for CMake, omit these engine names because they are mandatory (always compiled in):

```
csv mysam mysammrg heap
```

Then use these options to enable the InnoDB, ARCHIVE, and BLACKHOLE engines:

```
-DWITH_INNOBASE_STORAGE_ENGINE=1 -DWITH_ARCHIVE_STORAGE_ENGINE=1 -DWITH_BLACKHOLE_STORAGE_ENGINE=1
```

You can also use ON rather than 1 as the option value.

If you used `--without-plugin=engine` in configure to exclude a storage engine from the build, use `-DWITHOUT_ENGINE_STORAGE_ENGINE` in CMake.

Examples:

```
-DWITHOUT_EXAMPLE_STORAGE_ENGINE=1 -DWITHOUT_FEDERATED_STORAGE_ENGINE=1 -DWITHOUT_PARTITION_STORAGE_ENG
```

If neither `-DWITH ENGINE STORAGE_ENGINE` nor `-DWITHOUT ENGINE STORAGE_ENGINE` are specified for a given storage engine, the engine is built as a shared module, or excluded if it cannot be built as a shared module. A shared module must be installed using the `INSTALL PLUGIN` statement or the `--plugin-load` option before it can be used.

For additional information about CMake options for plugins, see [Chapter 6, Plugins](#).

Library Options

These options shown in the following table indicate which libraries to use.

Parameter	configure Option	CMake Option	CMake Notes
readline library	<code>--with-readline</code>	<code>-DWITH_READLINE=1</code>	
SSL library	<code>--with-ssl=/usr</code>	<code>-DWITH_SSL=system</code>	
zlib library	<code>--with-zlib-dir=/usr</code>	<code>-DWITH_ZLIB=system</code>	
libwrap library	<code>--without-libwrap</code>	<code>-DWITH_LIBWRAP=0</code>	

Miscellaneous Options

Most of the previous MySQL build options are supported. The normal mapping between old and new is uppercase, remove leading dashes, replace dash with underscore.

Examples:

```
--with-debug => WITH_DEBUG=1 --with-embedded-server => WITH_EMBEDDED_SERVER
```

Parameter	configure Option	CMake Option	CMake Notes
TCP/IP port number	<code>--with-tcp-port=3306</code>	<code>-DMYSQL_TCP_PORT=3306</code>	
UNIX socket file	<code>--with-unix-socket-path=/tmp/mysql.sock</code>	<code>-DMYSQL_UNIX_ADDR=/tmp/mysql.sock</code>	
Enable LOCAL for LOAD DATA	<code>--enable-local-infile</code>	<code>-DENABLED_LOCAL_INFILE=1</code>	
Extra charsets	<code>--with-extra-charsets=all</code>	<code>-DEXTRA_CHARSETS=all</code>	default is "all"
Default charset	<code>--with-charset=utf8</code>	<code>-DDEFAULT_CHARSET=utf8</code>	
Default collation	<code>--with-collation=utf8_general_ci</code>	<code>-DDEFAULT_COLLATION=utf8_general_ci</code>	
Build the server	<code>--with-server</code>	none	
Build the embedded server	<code>--with-embedded-server</code>	<code>-DWITH_EMBEDDED_SERVER=1</code>	
libmysqld privilege control	<code>--with-embedded-privilege-control</code>	none	always enabled?
Install the documentation	<code>--without-docs</code>	none	
Big tables	<code>--with-big-tables, --without-big-tables</code>	none	tables are big by default
mysqld user	<code>--with-mysqld-user=mysql</code>	<code>-DMYSQL_USER=mysql</code>	mysql is the default

Debugging	--without-debug	-DWITH_DEBUG=0	default is debugging disabled
GIS support	--with-geometry	none	always enabled?
Community features	--enable-community-features	none	always enabled
Profiling	--disable-profiling	-DENABLE_PROFILING=0	enabled by default
pstack	--without-pstack	none	pstack is removed
Assembler string functions	--enable-assembler	none	
Build type	--build=x86_64-pc-linux-gnu	no equivalent	unneeded?
Cross-compile host	--host=x86_64-pc-linux-gnu	no equivalent	unneeded?
Client flag	--with-client-ldflags=-lstdc++	none	unneeded
Client flag	--enable-thread-safe-client	none	unneeded, clients are always thread safe
Comment	--with-comment='string'	-DWITH_COMMENT='string'	
Shared/static binaries	--enable-shared --enable-static	none	there is only DISABLE_SHARED
Memory use	--with-low-memory	none	unneeded

Debugging the Configuration Process

Configuration with autotools produces config.log and config.status files.

Configuration with CMake produces files under the CMakeFiles directory : CMakeFiles/CMakeError.log and CMakeFiles/CMakeOutput.log

Interface to Third-Party Tools

Previously, third-party tools that need to determine the MySQL version from the MySQL source read the configure.in file in the top-level source directory. For example, the AC_INIT line for 5.5.7-rc looked like this:

```
AC_INIT([MySQL Server], [5.5.7-rc], [], [mysql])
```

Such tools now can read the VERSION file. For example, if the version is 5.5.8, the file looks like this:

```
MYSQL_VERSION_MAJOR=5 MYSQL_VERSION_MINOR=5 MYSQL_VERSION_PATCH=8 MYSQL_VERSION_EXTRA=
```

If the source is not for a General Availability (GA) release, the MYSQL_VERSION_EXTRA value will be nonempty. For example, the value for a Release Candidate release would look like this:

```
MYSQL_VERSION_EXTRA=rc
```

To construct a five-digit number from the version components, use this formula:

```
MYSQL_VERSION_MAJOR*10000 + MYSQL_VERSION_MINOR*100 + MYSQL_VERSION_PATCH
```

Chapter 6. Plugins

Table of Contents

CMake Variables to Control Plugin Building (MySQL 5.5 and Later)	51
CMake Macro to Define the Plugin	51
Autotools configure Support (MySQL 5.1)	52
Autotools Plugin Macros	52
Specifying <code>mysqld</code> Variables Within a Plugin	55
Additional Accessors for Plugins	55
MySQL Services for Plugins	55
Plugin Services Components	55
Writing Services for Use from Within Plugins	56

Beginning with MySQL 5.1, the server supports a plugin architecture for loading plugins. For example, several storage engines have been converted to plugins, and they can be selected or disabled at configuration time. You can build a plugin as static (compiled into the server) or dynamic (built as a dynamic library that must be installed using the `INSTALL PLUGIN` statement or the `--plugin-load` option before it can be used). Some plugins might not support static or dynamic build.

This section describes the command-line options that are used to control which plugins get built, and the CMake/autotools macros that enable plugin configuration support to be described. The support for building plugins is different in MySQL 5.5 and MySQL 5.1 (the build tools are CMake in 5.5, autotools in 5.1).

CMake Variables to Control Plugin Building (MySQL 5.5 and Later)

- To build a plugin that is statically compiled into the server (assuming that the plugin supports static build), add `-DWITH_<PLUGIN>=1` to the CMake command line.
- To exclude a plugin from the build, use `-DWITHOUT_<PLUGIN>=1`
- If neither `WITH_<PLUGIN>` nor `WITHOUT_<PLUGIN>` are defined, the plugin will be built as a shared module, if plugin supports it. If the plugin does not support build as a shared module, it won't be built.
- `WITH_PLUGIN_<PLUGIN>` and `WITH_<PLUGIN>_STORAGE_ENGINE` are also supported.

CMake Macro to Define the Plugin

To define a plugin, you need to add the `MYSQL_ADD_PLUGIN()` macro into `CMakeList.txt`. Unlike in 5.1, there is no special `plug.in` script. If you need to perform system checks, use standard CMake techniques like `CHECK_FUNCTION_EXISTS` or `CHECK_INCLUDE_FILE` etc.

Note: There is NO autoheader-functionality. For example, `CHECK_FUNCTION_EXISTS(epoll_wait HAVE_EPOLL_WAIT)` will not automatically add `"#define HAVE_EPOLL_WAIT 1"` in `config.h`. Different plugins might chose different strategies to add plugin-specific defines

- Strategy 1 - use `ADD_DEFINITIONS`

This is similar to what InnoDB does in 5.5. An example

```
CHECK_FUNCTION_EXISTS(epoll_wait HAVE_EPOLL_WAIT)
IF(HAVE_EPOLL_WAIT)
    ADD_DEFINITIONS(-DHAVE_EPOLL_WAIT=1)
ENDIF()
```

- Strategy 2 - use own header template

It is more work than in Strategy 1 but result is a cleaner solution:

1) You need to have plugin specific <plugin>_config.h.in with content similar to

```
#cmakedefine HAVE_EPOLL_WAIT
#cmakedefine HAVE_EPOLL_CTL
```

2) In CMakeLists.txt, add system checks

```
CHECK_FUNCTION_EXISTS(epoll_wait HAVE_EPOLL_WAIT)
CHECK_FUNCTION_EXISTS(epoll_ctl HAVE_EPOLL_CTL)
```

After all system checks, add

```
CONFIGURE_FILE(plugin_config.h.in plugin_config.h)
```

3) use #include "plugin_config.h" in your source files

Autotools configure Support (MySQL 5.1)

Several **configure** options apply to plugin selection and building.

configure --help shows the following information pertaining to plugins:

- The plugin-related options
- The names of all available plugins
- For each plugin, a description of its purpose, which build types it supports (static or dynamic), and which plugin groups it is a part of.

The following **configure** options are used to select or disable plugins:

```
--with-plugins=PLUGIN[,PLUGIN]...
--with-plugins=GROUP
--with-plugin=PLUGIN
--without-plugin=PLUGIN
```

PLUGIN is an individual plugin name such as *csv* or *archive*.

As shorthand, *GROUP* is a configuration group name such as *none* (select no plugins), *all* (select all plugins), or *max* (select all plugins used in a **mysqld-max** server).

--with-plugins can take a list of one or more plugin names separated by commas, or a plugin group name. The named plugins are configured to be built as static plugins.

--with-plugin-*PLUGIN* configures the given plugin to be built as a static plugin.

--without-plugin-*PLUGIN* disables the given plugin from being built.

If a plugin is named both with a **--with** and **--without** option, the result is undefined.

For any plugin that is not explicitly or implicitly (as a member of a selected group) selected or disabled, it is selected to be built dynamically if it supports dynamic build, and is disabled if it does not support dynamic build. If no plugin options are given, **default** group is selected.

Autotools Plugin Macros

The following macros enable plugin support in the autotools configuration files.

- Declaring a plugin:

```
MYSQL_PLUGIN(name, long-name, description [,configlist])
```

Each plugin is required to have `MYSQL_PLUGIN()` declared first. `configlist` is an optional argument that is a comma-separated list of configurations of which the module is a member. Example:

```
MYSQL_PLUGIN(ftexample, [Simple Parser], [Simple full-text parser plugin])
```

- Declaring a storage engine plugin:

```
MYSQL_STORAGE_ENGINE(name, legacy-opt, long-name, description  
[, configlist])
```

This is a simple utility macro that calls `MYSQL_PLUGIN`. It performs the bare basics required to declare a storage engine plugin and provides support for handling the legacy **configure** command-line options. If `legacy-opt` is not specified, it will default to `--with-name-storage-engine`. Set the `legacy-opt` value to `no` if you do not want to handle any legacy option. This macro is roughly equivalent to:

```
MYSQL_PLUGIN(name, 'long-name, description')  
MYSQL_PLUGIN_DEFINE(name, WITH_NAME_STORAGE_ENGINE)
```

Example:

```
MYSQL_STORAGE_ENGINE(berkeley, berkeley-db, [BerkeleyDB Storage Engine],  
[Transactional Tables using BerkeleyDB], [max,max-no-ndb])
```

- Declaring a C preprocessor variable:

```
MYSQL_PLUGIN_DEFINE(name, define-name)
```

When a plugin will be included in a static build, this will set a preprocessor variable to 1. These preprocessor variables are defined in `config.h`. Example:

```
MYSQL_PLUGIN_DEFINE(innodb, WITH_INNOBASE_STORAGE_ENGINE)
```

- Declaring a source directory for a plugin:

```
MYSQL_PLUGIN_DIRECTORY(name, dir-name)
```

Includes the specified directory into the build. If a file named `configure` is detected in the directory, it will be executed as part of the **configure** build otherwise it is assumed that there is a `Makefile` to be built in that directory. Currently, there is only support for plugin directories to be specified in the `storage/` and `plugin/` subdirectories. Example:

```
MYSQL_PLUGIN_DIRECTORY(archive, [storage/archive])
```

- Declaring a static library name for a plugin:

```
MYSQL_PLUGIN_STATIC(name, dir-name)
```

Sets the **configure** substitution `@plugin_name_static_target@` to the supplied library name if the plugin is a static build. It also adds the library to the list of libraries to be linked into **mysqld**. It may either be just the name of the library (where, if there is a directory specified, the directory will be prepended for the link) or another **make** variable or substitution (in which case, it will be passed through as is). Example:

```
MYSQL_PLUGIN_STATIC(archive, [libarchive.a]) MYSQL_PLUGIN_STATIC(berkeley, [[\$(bdb_libs_with_path)])
```

- Declaring a dynamic library name for a plugin:

```
MYSQL_PLUGIN_DYNAMIC(name, dso-name)
```

Sets the **configure** substitution `@plugin_name_shared_target@` to the supplied dynamic shared object library name if the module is a dynamic build. Example:

```
MYSQL_PLUGIN_DYNAMIC(archive, [ha_archive.la])
```

- Declaring a plugin as a mandatory module:

```
MYSQL_PLUGIN_MANDATORY(name)
```

Mandatory plugins cannot be disabled. Example:

```
MYSQL_PLUGIN_MANDATORY(myisam)
```

- Declaring a plugin as disabled:

```
MYSQL_PLUGIN_DISABLED(name)
```

A disabled plugin will not be included in any build. If the plugin has been marked as **MANDATORY**, it will result in an **autoconf** error.

- Declaring additional plugin **configure** actions:

```
MYSQL_PLUGIN_ACTIONS(name, configure-actions)
```

This is useful if there are additional **configure** actions required for a plugin. The *configure-actions* argument may either be the name of an **autoconf** macro or more **autoconf** script. Example:

```
MYSQL_PLUGIN_ACTIONS(ndbcluster, [MYSQL_SETUP_NDBCLUSTER])
```

- Declaring plugin dependencies:

```
MYSQL_PLUGIN_DEPENDS(name, dependencies)
```

Declares all plugins, in a comma-separated list, that are required for the named plugin to be built. If the named plugin is selected, it will in turn enable all its dependencies. All plugins listed as a dependency must already have been declared with **MYSQL_PLUGIN()**. Example:

```
MYSQL_PLUGIN_DEPENDS(ndbcluster, [partition])
```

- Performing the magic:

```
MYSQL_CONFIGURE_PLUGINS(default-names)
```

Actually performs the task of generating the shell scripts for **configure** based upon the declarations made previously. It emits the shell code necessary to check the options and sets the variables accordingly. Example:

```
MYSQL_CONFIGURE_PLUGINS([none])
```

Plugin-related **configure** errors:

- When any plugin macro is called before **MYSQL_PLUGIN()** is declared for that plugin, **configure** aborts with an error.
- When any of the plugins specified in the dependency list don't exist, **configure** aborts with an error.
- When a mandatory plugin is specified in **--without-plugin-PLUGIN**, **configure** aborts with an error.
- When a disabled plugin is specified in **--with-modules=...** or **--with-plugin=PLUGIN**, **configure** reports an error.
- When an optional plugin that may only be built dynamically is specified in **--with-plugins=...** or **--with-plugin-PLUGIN**, **configure** emits a warning and continues to configure the plugin for dynamic build.
- When an optional plugin that may only be built statically is specified neither in **--with-plugins=...** nor **--without-plugin-PLUGIN**, **configure** emits a warning but should proceed anyway.

Avoiding **configure.in** changes:

- If a plugin source (which is located in a subdirectory of the `storage/` or `plugin/` directory) contains a `plug.in` file (for example, `storage/example/plugin.in`), this file will be included as a part of `configure.in`. This way, `configure.in` does not need to be modified to add a new plugin to the build.
- A `plug.in` file may contain everything, particularly all `MYSQL_PLUGIN_XXX` macros as just described. The `plug.in` file does not need to specify `MYSQL_PLUGIN_DIRECTORY`; it is set automatically to the directory of the `plug.in` file.

Specifying `mysqld` Variables Within a Plugin

A plugin can implement status and system variables. Information about this is in the [plugin API](#) section of the MySQL Reference Manual.

Additional Accessors for Plugins

Starting with MySQL 5.1.21 the following additional accessors are made available to all plug-ins:

- Full definition of `MYSQL_LEX_STRING` (identical to `LEX_STRING` from `m_string.h`)
- Full definition of `MYSQL_XID` (binary compatible with `XID` from `handler.h`)
- `mysql_tmpfile()`, creates a temporary file in `mysqld`'s `tmpdir`
- `thd_killed()`, to check killed state of connection
- `thd_alloc()` and similar allocation functions
- `thd_get_xid()`, to get `XID` of connection's transaction
- `mysql_query_cache_invalidate4()`, to invalidate a table's query cache entries

MySQL Services for Plugins

As of MySQL 5.5, plugins have access to server “services.” The services interface exposes server functionality that plugins can call. It complements the plugin API and has these characteristics:

- Services enable plugins to access code inside the server using ordinary function calls.
- Services are portable and work on multiple platforms.
- The interface includes a versioning mechanism so that plugin versions can be checked at load time against service versions supported by the server. Versioning protects against incompatibilities between the version of a service that the server provides and the version of the service expected or required by a plugin.

Plugin Services Components

On the plugin side of the services interface, the relevant information is provided in a set of header files. A plugin accesses this information by including the `plugin.h` file (which plugins must include anyway):

```
#include <mysql/plugin.h>
```

`plugin.h` includes the `services.h` file, which acts as an “umbrella” file that includes the service-specific headers with names of the form `service_XXX.h`. Within a MySQL source distribution, the header files are located in the `include/mysql` directory and have an inclusion hierarchy like this:

- `plugin.h` includes `services.h`.
- `services.h` is the “umbrella” header that includes all available service-specific header files.

- Service-specific headers have names like `service_my_snprintf.h` or `service_thd_alloc.h`.

The server side of the services interface uses the header files just listed, but also involves other files. Within a MySQL source distribution, these files are located in the `libservices` and `sql` directories:

The `libservices` directory contains the source files from which the `libmysqlservices` library is built. Files in this directory include:

- `HOWTO` : Instructions for writing plugin services.
- `xxx_services.h` : Service-specific interface files.

Current service files are `my_snprintf_service.c` and `thd_alloc_service.h`.

During the MySQL build and install process, the `libmysqlservices` library is compiled and installed in a directory where plugins can access it. All plugins should link in this library using the `-lmysqlservices` flag when they are built.

The `sql` directory contains `sql_plugin.cc`, which implements plugin functionality. This file includes `sql_plugin_services.h`, where each available service is registered.

Writing Services for Use from Within Plugins

Services are implemented by modifying server code in the `sql` directory to register the service in the plugin code, and by providing interface files in the `libservices` directory and under the `include` directory.

For complete instructions on writing a service, see the `HOWTO` file in the `libservices` directory within a MySQL source distribution. As you read the `HOWTO` file, you might find it useful to examine the files that implement existing services.

When you write a service, be sure to provide complete instructions on how to use it. A service named `xxx` will have a file named `service_xxx.h` in the `include/mysql` directory. For the benefit of plugin developers who use your service, this file should include comments that fully document the service interface:

- Its purpose
- Any applicable guidelines, including limitations or restrictions
- For each function, a description of what it does, its calling sequence, and return value

The goal for documentation in this file is that plugin developers should be able to look at the file and completely understand how to use the service.

Chapter 7. Transaction Handling in the Server

Table of Contents

Historical Note	57
Current Situation	57
Data Layout	58
Transaction Life Cycle	58
Roles and Responsibilities	59
Additional Notes on DDL and the Normal Transaction	60

In each client connection, MySQL maintains two transactional states:

- A statement transaction
- A standard transaction, also called a normal transaction

Historical Note

"Statement transaction" is a non-standard term that comes from the days when MySQL supported the BerkeleyDB storage engine.

First, observe that in BerkeleyDB the "auto-commit" mode causes automatic commit of operations that are atomic from the storage engine's perspective, such as a write of a record, but are too fine-grained to be atomic from the application's (MySQL's) perspective. One SQL statement could involve many BerkeleyDB auto-committed operations. So BerkeleyDB auto-commit was of little use to MySQL.

Second, observe that BerkeleyDB provided the concept of "nested transactions" instead of SQL standard savepoints. In a nutshell: transactions could be arbitrarily nested, but when the parent transaction was committed or aborted, all its child (nested) transactions were committed or aborted as well. Commit of a nested transaction, in turn, made its changes visible, but not durable: it destroyed the nested transaction, so all the nested transaction's changes would become visible to the parent and to other currently active nested transactions of the same parent.

So MySQL employed the mechanism of nested transactions to provide the "all or nothing" guarantee for SQL statements that the standard requires. MySQL would create a nested transaction at the start of each SQL statement, and destroy (commit or abort) the nested transaction at statement end. MySQL people internally called such a nested transaction a "statement transaction". And that's what gave birth to the term "statement transaction".

Current Situation

Nowadays a statement transaction is started for each statement that accesses transactional tables or uses the binary log. If the statement succeeds, the statement transaction is committed. If the statement fails, the transaction is rolled back. Commits of statement transactions are not durable -- each statement transaction is nested in the normal transaction, and if the normal transaction is rolled back, the effects of all enclosed statement transactions are undone as well. Technically, a statement transaction can be viewed as a transaction which starts with a savepoint which MySQL maintains automatically, in order to make the effects of one statement atomic.

The normal transaction is started by the user and is usually completed by a user request as well. The normal transaction encloses all statement transactions that are issued between its beginning and its end. In autocommit mode, the normal transaction is equivalent to the statement transaction.

Since MySQL supports pluggable storage engine architecture (PSEA), more than one transactional engine may be active at a time. So from the server point of view, transactions are always distributed.

In particular, MySQL maintains transactional state independently for each engine. To commit a transaction, MySQL employs a two-phase commit protocol.

Not all statements are executed in the context of a transaction. Administrative and status-information statements don't modify engine data, so they don't start a statement transaction and they don't affect a normal transaction. Examples of such statements are `SHOW STATUS` and `RESET SLAVE`.

Similarly, DDL statements are not transactional, and therefore a transaction is (almost) never started for a DDL statement. But there's a difference between a DDL statement and an administrative statement: the DDL statement always commits the current transaction (if any) before proceeding; the administrative statement doesn't.

Finally, SQL statements that work with non-transactional engines also have no effect on the transaction state of the connection. Even though they cause writes to the binary log, (and the binary log is by and large transactional), they write in "write-through" mode directly to the binlog file, then they do an OS cache sync -- in other words, they bypass the binlog undo log (translog). They do not commit the current normal transaction. A failure of a statement that uses non-transactional tables would cause a rollback of the statement transaction, but that's irrelevant if no non-transactional tables are used, because no statement transaction was started.

Data Layout

The server stores its transaction-related data in `thd->transaction`. This structure has two members of type `THD_TRANS`. These members correspond to the statement and normal transactions respectively:

- `thd->transaction.stmt` contains a list of engines that are participating in the given statement
- `thd->transaction.all` contains a list of engines that have participated in any of the statement transactions started within the context of the normal transaction. Each element of the list contains a pointer to the storage engine, engine-specific transactional data, and engine-specific transaction flags.

In autocommit mode, `thd->transaction.all` is empty. In that case, data of `thd->transaction.stmt` is used to commit/roll back the normal transaction.

The list of registered engines has a few important properties:

- No engine is registered in the list twice.
- Engines are present in the list in reverse temporal order -- new participants are always added to the beginning of the list.

Transaction Life Cycle

When a new connection is established, `thd->transaction` members are initialized to an empty state. If a statement uses any tables, all affected engines are registered in the statement engine list. In non-autocommit mode, the same engines are registered in the normal transaction list. At the end of the statement, the server issues a commit or a rollback for all engines in the statement list. At this point the transaction flags of an engine, if any, are propagated from the statement list to the list of the normal transaction. When commit/rollback is finished, the statement list is cleared. It will be filled in again by the next statement, and emptied again at the next statement's end.

The normal transaction is committed in a similar way (by going over all engines in `thd->transaction.all` list) but at different times:

- When the user issues an SQL `COMMIT` statement
- Implicitly, when the server begins handling a DDL statement or `SET AUTOCOMMIT={0|1}` statement

The normal transaction can be rolled back as well:

- When the user issues an SQL ROLLBACK statement
- When one of the storage engines requests a rollback by setting `thd->transaction_rollback_request`

For example, the latter condition may occur when the transaction in the engine was chosen as a victim of the internal deadlock resolution algorithm and rolled back internally. In such situations there is little the server can do and the only option is to roll back transactions in all other participating engines, and send an error to the user.

From the use cases above, it follows that a normal transaction is never committed when there is an outstanding statement transaction. In most cases there is no conflict, because commits of a normal transaction are issued by a stand-alone administrative or DDL statement, and therefore no outstanding statement transaction of the previous statement can exist. Besides, all statements that operate via a normal transaction are prohibited in stored functions and triggers, therefore no conflicting situation can occur in a sub-statement either. The remaining rare cases, when the server explicitly must commit a statement transaction prior to committing a normal transaction, are error-handling cases (see for example `SQLCOM_LOCK_TABLES`).

When committing a statement or a normal transaction, the server either uses the two-phase commit protocol, or issues a commit in each engine independently. The server uses the two-phase commit protocol only if:

- All participating engines support two-phase commit (by providing a `handler::prepare PSEA API` call), and
- Transactions in at least two engines modify data (that is, are not read-only)

Note that the two-phase commit is used for statement transactions, even though statement transactions are not durable anyway. This ensures logical consistency of data in a multiple-engine transaction. For example, imagine that some day MySQL supports unique constraint checks deferred until the end of the statement. In such a case, a commit in one of the engines could yield `ER_DUP_KEY`, and MySQL should be able to gracefully abort the statement transactions of other participants.

After the normal transaction has been committed, the `thd->transaction.all` list is cleared.

When a connection is closed, the current normal transaction, if any, is rolled back.

Roles and Responsibilities

The server has only one way to know that an engine participates in the statement and a transaction has been started in an engine: the engine says so. So, in order to be a part of a transaction, an engine must "register" itself. This is done by invoking the `trans_register_ha()` server call. Normally the engine registers itself whenever `handler::external_lock()` is called. Although `trans_register_ha()` can be invoked many times, it does nothing if the engine is already registered. If autocommit is not set, the engine must register itself twice -- both in the statement list and in the normal transaction list. A parameter of `trans_register_ha()` specifies which list to register.

Note: Although the registration interface in itself is fairly clear, the current usage practice often leads to undesired effects. For example, since a call to `trans_register_ha()` in most engines is embedded into an implementation of `handler::external_lock()`, some DDL statements start a transaction (at least from the server point of view) even though they are not expected to. For example `CREATE TABLE` does not start a transaction, since `handler::external_lock()` is never called during `CREATE TABLE`. But `CREATE TABLE ... SELECT` does, since `handler::external_lock()` is called for the table that is being selected from. This has no practical effects currently, but we must keep it in mind nevertheless.

Once an engine is registered, the server will do the rest of the work.

During statement execution, whenever any data-modifying PSEA API methods are used (for example, `handler::write_row()` or `handler::update_row()`), the read-write flag is raised in the statement transaction

for the relevant engine. Currently All PSEA calls are "traced", and the only way to change data is to issue a PSEA call. Important: Unless this invariant is preserved, the server will not know that a transaction in a given engine is read-write and will not involve the two-phase commit protocol!

The end of a statement causes invocation of the `ha_autocommit_or_rollback()` server call, which in turn invokes `handleron::prepare()` for every involved engine. After `handleron::prepare()`, there's a call to `handleron::commit_one_phase()`. If a one-phase commit will suffice, `handleron::prepare()` is not invoked and the server only calls `handleron::commit_one_phase()`. At statement commit, the statement-related read-write engine flag is propagated to the corresponding flag in the normal transaction. When the commit is complete, the list of registered engines is cleared.

Rollback is handled in a similar way.

Additional Notes on DDL and the Normal Transaction

DDL statements and operations with non-transactional engines do not "register" in `thd->transaction` lists, and thus do not modify the transaction state. Besides, each DDL statement in MySQL begins with an implicit normal transaction commit (a call to `end_active_trans()`), and thus leaves nothing to modify. However, as noted above for `CREATE TABLE .. SELECT`, some DDL statements can start a *new* transaction.

Behavior of the server in this case is currently badly defined. DDL statements use a form of "semantic" logging to maintain atomicity: If `CREATE TABLE t .. SELECT` fails, table `t` is deleted. In addition, some DDL statements issue interim transaction commits: for example, `ALTER TABLE` issues a commit after data is copied from the original table to the internal temporary table. Other statements, for example, `CREATE TABLE ... SELECT`, do not always commit after themselves. And finally there is a group of DDL statements such as `RENAME/DROP TABLE`, which don't start new transactions and don't commit.

This diversity makes it hard to say what will happen if by chance a stored function is invoked during a DDL statement -- it's not clear whether any modifications it makes will be committed or not. Fortunately, SQL grammar allows only a few DDL statements to invoke stored functions. Perhaps, for consistency, MySQL should always commit a normal transaction after a DDL statement, just as it commits a statement transaction at the end of a statement.

Chapter 8. The Optimizer

Table of Contents

Code and Concepts	61
Definitions	61
The Optimizer Code	61
Primary Optimizations	62
Optimizing Constant Relations	62
Optimizing Joins	64
Transpositions	70
ORDER BY Clauses	71
GROUP BY and Related Conditions	72
Other Optimizations	73

This chapter describes the operation of the MySQL Query optimizer, which is used to determine the most efficient means for executing queries.

Code and Concepts

This section discusses key optimizer concepts, terminology, and how these are reflected in the MySQL server source code.

Definitions

This description uses a narrow definition: The *optimizer* is the set of routines which decide what execution path the DBMS should take for queries.

MySQL changes these routines frequently, so you should compare what is said here with what's in the current source code. To make that easy, this description includes notes referring to the relevant file and routine, such as “See: `/sql/select_cc, optimize_cond()`”.

A *transformation* occurs when one query is changed into another query which delivers the same result. For example, a query could be changed from

```
SELECT ... WHERE 5 = a
```

to

```
SELECT ...WHERE a = 5
```

Most transformations are less obvious. Some transformations result in faster execution.

The Optimizer Code

This diagram shows the structure of the function `handle_select()` in `/sql/sql_select.cc` (the server code that handles a query):

```
handle_select()
mysql_select()
  JOIN::prepare()
  setup_fields()
  JOIN::optimize()           /* optimizer is from here ... */
  optimize_cond()
  opt_sum_query()
  make_join_statistics()
  get_quick_record_count()
  choose_plan()
  /* Find the best way to access tables */
```

```

/* as specified by the user.          */
optimize_straight_join()
    best_access_path()
/* Find a (sub-)optimal plan among all or subset */
/* of all possible query plans where the user    */
/* controls the exhaustiveness of the search.    */
greedy_search()
    best_extension_by_limited_search()
    best_access_path()
/* Perform an exhaustive search for an optimal plan */
find_best()
    make_join_select()          /* ... to here */
JOIN::exec()

```

The indentation in the diagram shows what calls what. Thus you can see that `handle_select()` calls `mysql_select()` which calls `JOIN::prepare()` which calls `setup_fields()`, and so on. The first part of `mysql_select()` is `JOIN::prepare()` which is for context analysis, metadata setup, and some subquery transformations. The optimizer is `JOIN::optimize()` and all its subordinate routines. When the optimizer finishes, `JOIN::exec()` takes over and does the job that `JOIN::optimize()` decides upon.

Although the word “JOIN” appears, these optimizer routines are applicable to all query types.

The `optimize_cond()` and `opt_sum_query()` routines perform transformations. The `make_join_statistics()` routine puts together all the information it can find about indexes that might be useful for accessing the query's tables.

Primary Optimizations

This section discusses the most important optimizations performed by the server.

Optimizing Constant Relations

Constant Propagation

A transformation takes place for expressions like this:

```
WHERE column1 = column2 AND column2 = 'x'
```

For such expressions, since it is known that, *if A=B and B=C then A=C* (the Transitivity Law), the transformed condition becomes:

```
WHERE column1='x' AND column2='x'
```

This transformation occurs for `column1 <operator> column2` conditions if and only if `<operator>` is one of these operators:

```
=, <, >, <=, >=, <>, <=>, LIKE
```

That is, transitive transformations don't apply for `BETWEEN`. Probably they should not apply for `LIKE` either, but that's a story for another day.

Constant propagation happens in a loop, so the output from one *propagation step* can be input for the next step.

See: `/sql/sql_select.cc, change_cond_ref_to_const()`. Or See: `/sql/sql_select.cc, propagate_cond_constants()`.

Eliminating “Dead” Code

A transformation takes place for conditions that are always true, for example:

```
WHERE 0=0 AND column1='y'
```

In this case, the first condition is removed, leaving

```
WHERE column1='y'
```

See: /sql/sql_select.cc, `remove_eq_conds()`.

A transformation also takes place for conditions that are always false. For example, consider this `WHERE` clause:

```
WHERE (0 = 1 AND s1 = 5) OR s1 = 7
```

Since the parenthesized part is always false, it is removed, reducing this expression to

```
WHERE s1 = 7
```

In some cases, where the `WHERE` clause represents an impossible condition, the optimizer might eliminate it completely. Consider the following:

```
WHERE (0 = 1 AND s1 = 5)
```

Because it is never possible for this condition to be true, the `EXPLAIN` statement will show the words `Impossible WHERE`. Informally, we at MySQL say that the `WHERE` has been optimized away.

If a column cannot be `NULL`, the optimizer removes any non-relevant `IS NULL` conditions. Thus,

```
WHERE not_null_column IS NULL
```

is an always-false situation, and

```
WHERE not_null_column IS NOT NULL
```

is an always-true situation so such columns are also eliminated from the conditional expression. This can be tricky. For example, in an `OUTER JOIN`, a column which is defined as `NOT NULL` might still contain a `NULL`. The optimizer leaves `IS NULL` conditions alone in such exceptional situations.

The optimizer will not detect all `Impossible WHERE` situations — there are too many possibilities in this regard. For example:

```
CREATE TABLE Table1 (column1 CHAR(1));
...
SELECT * FROM Table1 WHERE column1 = 'Canada';
```

The optimizer will not eliminate the condition in the query, even though the `CREATE TABLE` definition makes it an impossible condition.

Folding of Constants

A transformation takes place for this expression:

```
WHERE column1 = 1 + 2
```

which becomes:

```
WHERE column1 = 3
```

Before you say, “but I never would write `1 + 2` in the first place”, remember what was said earlier about constant propagation. It is quite easy for the optimizer to put such expressions together. This process simplifies the result.

Constants and Constant Tables

A MySQL *constant* is something more than a mere literal in the query. It can also be the contents of a *constant table*, which is defined as follows:

1. A table with zero rows, or with only one row
2. A table expression that is restricted with a `WHERE` condition, containing expressions of the form `column = constant`, for all the columns of the table's primary key, or for all the columns of any of the table's unique keys (provided that the unique columns are also defined as `NOT NULL`).

For example, if the table definition for `Table0` contains

```
... PRIMARY KEY (column1,column2)
```

then this expression

```
FROM Table0 ... WHERE column1=5 AND column2=7 ...
```

returns a constant table. More simply, if the table definition for `Table1` contains

```
... unique_not_null_column INT NOT NULL UNIQUE
```

then this expression

```
FROM Table1 ... WHERE unique_not_null_column=5
```

returns a constant table.

These rules mean that a constant table has at most one row value. MySQL will evaluate a constant table in advance, to find out what that value is. Then MySQL will plug that value into the query. Here's an example:

```
SELECT Table1.unique_not_null_column, Table2.any_column
FROM Table1, Table2
WHERE Table1.unique_not_null_column = Table2.any_column
AND Table1.unique_not_null_column = 5;
```

When evaluating this query, MySQL first finds that table `Table1` after restriction with `Table1.unique_not_null_column` is a constant table according to the second definition above. So it retrieves that value.

If the retrieval fails (there is no row in the table with `unique_not_null_column = EXPLAIN` for the statement:

```
Impossible WHERE noticed after reading const tables
```

Alternatively, if the retrieval succeeds (there is exactly one row in the table with `unique_not_null_column = MySQL` transforms the query to this:

```
SELECT 5, Table2.any_column
FROM Table1, Table2
WHERE 5 = Table2.any_column
AND 5 = 5;
```

Actually this is a grand-combination example. The optimizer does some of the transformation because of constant propagation, which we described earlier. By the way, we described constant propagation first because it happens *before* MySQL figures out what the constant tables are. The sequence of optimizer steps sometimes makes a difference.

Although many queries have no constant-table references, it should be kept in mind that whenever the word *constant* is mentioned hereafter, it refers either to a literal or to the contents of a constant table.

See: [/sql/sql_select.cc, make_join_statistics\(\)](#).

Optimizing Joins

This section discusses the various methods used to optimize joins.

Determining the Join Type

When evaluating a conditional expression, MySQL decides what *join type* the expression has. (Again: despite the word join, this applies for all conditional expressions, not just join expressions. A term like access type would be clearer.) These are the documented join types, in order from best to worst:

- `system` : a system table which is a constant table
- `const` : a constant table
- `eq_ref` : a unique or primary index with an equality relation
- `ref` : an index with an equality relation, where the index value cannot be `NULL`
- `ref_or_null` : an index with an equality relation, where it is possible for the index value to be `NULL`
- `range` : an index with a relation such as `BETWEEN`, `IN`, `>=`, `LIKE`, and so on.
- `index` : a sequential scan on an index
- `ALL` : a sequential scan of the entire table

See: `/sql/sql_select.h, enum join_type{}`. Notice that there are a few other (undocumented) join types too, for subqueries.

The optimizer can use the join type to pick a *driver expression*. For example, consider this query:

```
SELECT *
FROM Table1
WHERE indexed_column = 5 AND unindexed_column = 6
```

Since `indexed_column` has a better join type, it is more likely to be the driver. You'll see various exceptions as this description proceeds, but this is a simple first rule.

What is significant about a driver? Consider that there are two execution paths for the query:

The Bad Execution Plan: Read every row in the table. (This is called a *sequential scan* of `Table1` or simply *table scan*.) For each row, examine the values in `indexed_column` and in `unindexed_column`, to see if they meet the conditions.

The Good Execution Plan: Via the index, look up the rows which have `indexed_column = 5`. (This is called an *indexed search*.) For each row, examine the value in `unindexed_column` to see if it meets the condition.

An indexed search generally involves fewer accesses than a sequential scan, and far fewer accesses if the table is large but the index is unique. That is why it is better to access with the good execution plan, and why it is often good to choose `indexed_column` as the driver.

Joins and Access Methods

Bad join choices can cause more damage than bad choices in single-table searches, so MySQL developers have spent proportionally more time making sure that the tables in a query are joined in an optimal order and that optimal access methods (often called *access paths*) are chosen to retrieve table data. A combination of a fixed order in which tables are joined and the corresponding table access methods for each table is called *query execution plan (QEP)*. The goal of the query optimizer is to find an optimal QEP among all such possible plans. There are several general ideas behind join optimization.

Each plan (or part of plan) is assigned a *cost*. The cost of a plan reflects roughly the resources needed to compute a query according to the plan, where the main factor is the number of rows that will be accessed while computing a query. Once we have a way to assign costs to different QEPs we have a way to compare them. Thus, the goal of the optimizer is to find a QEP with minimal cost among all possible plans.

In MySQL, the search for an optimal QEP is performed in a bottom-up manner. The optimizer first considers all plans for one table, then all plans for two tables, and so on, until it builds a complete optimal QEP. Query plans that consist of only some of the tables (and predicates) in a query are called *partial plans*. The optimizer relies on the fact that the more tables that are added to a partial plan, the greater its cost. This allows the optimizer to expand with more tables only the partial plans with lower cost than the current best complete plan.

The key routine that performs the search for an optimal QEP is `sql/sql_select.cc, find_best()`. It performs an exhaustive search of all possible plans and thus guarantees it will find an optimal one.

Below we represent `find_best()` in an extremely free translation to pseudocode. It is recursive, so some input variables are labeled so far to indicate that they come from a previous iteration.

```
remaining_tables = {t1, ..., tn}; /* all tables referenced in a query */

procedure find_best(
    partial_plan in,          /* in, partial plan of tables-joined-so-far */
    partial_plan_cost,       /* in, cost of partial_plan */
    remaining_tables,        /* in, set of tables not referenced in partial_plan */
    best_plan_so_far,        /* in/out, best plan found so far */
    best_plan_so_far_cost) /* in/out, cost of best_plan_so_far */
{
    for each table T from remaining_tables
    {
        /* Calculate the cost of using table T. Factors that the
           optimizer takes into account may include:
           Many rows in table (bad)
           Many key parts in common with tables so far (very good)
           Restriction mentioned in the WHERE clause (good)
           Long key (good)
           Unique or primary key (good)
           Full-text key (bad)
           Other factors that may at some time be worth considering:
           Many columns in key
           Short average/maximum key length
           Small table file
           Few levels in index
           All ORDER BY / GROUP columns come from this table */
        cost = complex-series-of-calculations;
        /* Add the cost to the cost so far. */
        partial_plan_cost += cost;

        if (partial_plan_cost >= best_plan_so_far_cost)
            /* partial_plan_cost already too great, stop search */
            continue;

        partial_plan = expand partial_plan by best_access_method;
        remaining_tables = remaining_tables - table T;
        if (remaining_tables is not an empty set)
        {
            find_best(partial_plan, partial_plan_cost,
                      remaining_tables,
                      best_plan_so_far, best_plan_so_far_cost);
        }
        else
        {
            best_plan_so_far_cost = partial_plan_cost;
            best_plan_so_far = partial_plan;
        }
    }
}
```

Here the optimizer applies a *depth-first search algorithm*. It performs estimates for every table in the `FROM` clause. It will stop a search early if the estimate becomes worse than the best estimate so far. The order of scanning will depend on the order that the tables appear in the `FROM` clause.

See: `/sql/table.h, struct st_table`.

`ANALYZE TABLE` may affect some of the factors that the optimizer considers.

See also: `/sql/sql_sselect.cc, make_join_statistics()`.

The straightforward use of `find_best()` and `greedy_search()` will not apply for `LEFT JOIN` or `RIGHT JOIN`. For example, starting with MySQL 4.0.14, the optimizer may change a left join to a straight join and swap the table order in some cases. See also [LEFT JOIN and RIGHT JOIN Optimization](#).

The **range** Join Type

Some conditions can work with indexes, but over a (possibly wide) range of keys. These are known as *range* conditions, and are most often encountered with expressions involving these operators: `>`, `>=`, `<`, `<=`, `IN`, `LIKE`, `BETWEEN`

To the optimizer, this expression:

```
column1 IN (1,2,3)
```

is the same as this one:

```
column1 = 1 OR column1 = 2 OR column1 = 3
```

and MySQL treats them the same — there is no need to change `IN` to `OR` for a query, or vice versa.

The optimizer will use an index (range search) for

```
column1 LIKE 'x%'
```

but not for

```
column1 LIKE '%x'
```

That is, there is no range search if the first character in the pattern is a wildcard.

To the optimizer,

```
column1 BETWEEN 5 AND 7
```

is the same as this expression

```
column1 >= 5 AND column1 <= 7
```

and again, MySQL treats both expressions the same.

The optimizer may change a **Range** to an **ALL** join type if a condition would examine too many index keys. Such a change is particularly likely for `<` and `>` conditions and multiple-level secondary indexes. See: (for **MyISAM** indexes) `/myisam/mi_range.c, mi_records_in_range()`.

The **index** Join Type

Consider this query:

```
SELECT column1 FROM Table1;
```

If `column1` is indexed, then the optimizer may choose to retrieve the values from the index rather than from the table. An index which is used this way is called a *covering index* in most texts. MySQL simply uses the word “index” in **EXPLAIN** descriptions.

For this query:

```
SELECT column1, column2 FROM Table1;
```

the optimizer will use `join type = index` only if the index has this definition:

```
CREATE INDEX ... ON Table1 (column1, column2);
```

In other words, all columns in the select list must be in the index. (The order of the columns in the index does not matter.) Thus it might make sense to define a multiple-column index strictly for use as a covering index, regardless of search considerations.

The Index Merge Join Type

Overview

[Index Merge](#) is used when table condition can be converted to form:

```
cond_1 OR cond_2 ... OR cond_N
```

The conditions for conversion are that each [cond_i](#) can be used for a range scan, and no pair ([cond_i](#), [cond_j](#)) uses the same index. (If [cond_i](#) and [cond_j](#) use the same index, then [cond_i OR cond_j](#) can be combined into a single range scan and no merging is necessary.)

For example, [Index Merge](#) can be used for the following queries:

```
SELECT * FROM t WHERE key1=c1 OR key2<c2 OR key3 IN (c3,c4);
SELECT * FROM t WHERE (key1=c1 OR key2<c2) AND nonkey=c3;
```

[Index Merge](#) is implemented as a “container” for range key scans constructed from [cond_i](#) conditions. When doing [Index Merge](#), MySQL retrieves rows for each of the keyscans and then runs them through a duplicate elimination procedure. Currently the [Unique](#) class is used for duplicate elimination.

Index Merge Optimizer

A single [SEL_TREE](#) object cannot be constructed for conditions that have different members of keys in the [OR](#) clause, like in condition:

```
key1 < c1 OR key2 < c2
```

Beginning with MySQL 5.0, these conditions are handled with the [Index Merge](#) method, and its range optimizer structure, class [SEL_IMERGE](#). [SEL_IMERGE](#) represents a disjunction of several [SEL_TREE](#) objects, which can be expressed as:

```
sel_imerge_cond = (t_1 OR t_1 OR ... OR t_n)
```

where each of [t_i](#) stands for a [SEL_TREE](#) object, and no pair ([t_i](#), [t_j](#)) of distinct [SEL_TREE](#) objects can be combined into single [SEL_TREE](#) object.

The current implementation builds [SEL_IMERGE](#) only if no single [SEL_TREE](#) object can be built for the part of the query condition it has analyzed, and discards [SEL_TREE](#) immediately if it discovers that a single [SEL_TREE](#) object can be constructed. This is actually a limitation, and can cause worse row retrieval strategy to be used. E.g. for query:

```
SELECT * FROM t WHERE (goodkey1=c1 OR goodkey1=c2) AND badkey=c3
```

scan on [badkey](#) will be chosen even if [Index Merge](#) on ([goodkey1](#), [goodkey](#)) would be faster.

The [Index Merge](#) optimizer collects a list of possible ways to access rows with [Index Merge](#). This list of [SEL_IMERGE](#) structures represents the following condition:

```
(t_11 OR t_12 OR ... OR t_1k) AND
(t_21 OR t_22 OR ... OR t_2l) AND
...
(t_M1 OR t_M2 OR ... OR t_mp)
```

where [t_ij](#) is one [SEL_TREE](#) and one line is for one [SEL_IMERGE](#) object.

The [SEL_IMERGE](#) object with *minimal* cost is used for row retrieval.

In `sql/opt_range.cc`, see `imerge_list_and_list()`, `imerge_list_or_list()`, and `SEL_IMERGE` class member functions for more details of `Index Merge` construction.

See the `get_index_merge_params` function in the same file for `Index Merge` cost calculation algorithm.

The `range` Optimizer

For `range` queries, the MySQL optimizer builds a `SEL_TREE` object which represents a condition in this form:

```
range_cond = (cond_key_1 AND cond_key_2 AND ... AND cond_key_N)
```

Each of `cond_key_i` is a condition that refers to components of one key. MySQL creates a `cond_key_i` condition for each of the usable keys. Then the cheapest condition `cond_key_i` is used for doing range scan.

A single `cond_key_i` condition is represented by a pointer-linked network of `SEL_ARG` objects. Each `SEL_ARG` object refers to particular part of the key and represents the following condition:

```
sel_arg_cond= (inf_val < key_part_n AND key_part_n < sup_val) (1)
               AND next_key_part_sel_arg_cond                (2)
               OR left_sel_arg_cond                          (3)
               OR right_sel_arg_cond                         (4)
```

1. is for an interval, possibly without upper or lower bound, either including or not including boundary values.
2. is for a `SEL_ARG` object with condition on next key component.
3. is for a `SEL_ARG` object with an interval on the same field as this `SEL_ARG` object. Intervals between the current and left objects are disjoint and `left_sel_arg_cond.sup_val <= inf_val`.
4. is for a `SEL_ARG` object with an interval on the same field as this `SEL_ARG` object. Intervals between the current and right objects are disjoint and `left_sel_arg_cond.min_val >= max_val`.

MySQL is able to convert arbitrary-depth nested AND-OR conditions to the above conjunctive form.

Row Retrieval Algorithm

`Index Merge` works in two steps:

Preparation step:

```
activate 'index only';
foreach key_i in (key_scans \ clustered_pk_scan)
{
    while (retrieve next (key, rowid) pair from key_i)
    {
        if (no clustered PK scan ||
            row doesn't match clustered PK scan condition)
            put rowid into Unique;
    }
}
deactivate 'index only';
```

Row retrieval step:

```
for each rowid in Unique
{
    retrieve row and pass it to output;
}
if (clustered_pk_scan)
{
```

```
while (retrieve next row for clustered_pk_scan)
    pass row to output;
}
```

See: [sql/opt_range.cc](#), [QUICK_INDEX_MERGE_SELECT](#) class members for [Index Merge](#) row retrieval code.

Transpositions

MySQL supports *transpositions* (reversing the order of operands around a relational operator) for simple expressions only. In other words:

```
WHERE - 5 = column1
```

becomes:

```
WHERE column1 = -5
```

However, MySQL does not support transpositions where arithmetic exists. Thus:

```
WHERE 5 = -column1
```

is *not* treated the same as:

```
WHERE column1 = -5
```

Transpositions to expressions of the form `column = constant` are ideal for index lookups. If an expression of this form refers to an indexed column, then MySQL always uses the index, regardless of the table size. (**Exception:** If the table has only zero rows or only one row, it is a constant table and receives special treatment. See [Constants and Constant Tables](#).)

AND Relations

An [AND](#)ed search has the form `condition1 AND condition2`, as in this example:

```
WHERE column1 = 'x' AND column2 = 'y'
```

Here, the optimizer's decision process can be described as follows:

1. If (neither condition is indexed) use sequential scan.
2. Otherwise, if (one condition has better join type) then pick a driver based on join type (see [Determining the Join Type](#)).
3. Otherwise, since (both conditions are indexed and have equal join type) pick a driver based on the first index that was created.

The optimizer can also choose to perform an [index_merge](#) index intersection, as described in [Index Merge Optimizer](#).

Here's an example:

```
CREATE TABLE Table1 (s1 INT, s2 INT);
CREATE INDEX Index1 ON Table1 (s2);
CREATE INDEX Index2 ON Table1 (s1);
...
SELECT * FROM Table1 WHERE s1 = 5 AND s2 = 5;
```

When choosing a strategy to solve this query, the optimizer picks `s2 = 5` as the driver because the index for `s2` was created first. Regard this as an accidental effect rather than a rule, it could change at any moment.

OR Relations

An [OR](#)ed search has the form `"condition1" OR "condition2"`, as in this example:

```
WHERE column1 = 'x' OR column2 = 'y'
```

Here the optimizer's decision is to use a sequential scan.

There is also an option to use index merge under such circumstances. See [Overview](#) and [Index Merge Optimization](#), for more information.

The above warning does not apply if the same column is used in both conditions. For example:

```
WHERE column1 = 'x' OR column1 = 'y'
```

In such a case, the search is indexed because the expression is a range search. This subject will be revisited during the discussion of the IN predicate.

UNION Queries

All [SELECT](#) statements within a [UNION](#) are optimized separately. Therefore, for this query:

```
SELECT * FROM Table1 WHERE column1 = 'x'  
UNION ALL  
SELECT * FROM Table1 WHERE column2 = 'y'
```

if both [column1](#) and [column2](#) are indexed, then each [SELECT](#) is done using an indexed search, and the result sets are merged. Notice that this query might produce the same results as the query used in the [OR](#) example, which uses a sequential scan.

NOT (<>) Relations

It is a logical rule that

```
column1 <> 5
```

is the same as

```
column1 < 5 OR column1 > 5
```

However, MySQL does not transform in this circumstance. If you think that a range search would be better, then you should do your own transforming in such cases.

It is also a logical rule that

```
WHERE NOT (column1 != 5)
```

is the same as

```
WHERE column1 = 5
```

However, MySQL does not transform in this circumstance either.

We expect to add optimizations for both the previous cases.

ORDER BY Clauses

In general, the optimizer will skip the sort procedure for the [ORDER BY](#) clause if it sees that the rows will be in order anyway. But let's examine some exceptional situations.

For the query:

```
SELECT column1 FROM Table1 ORDER BY 'x';
```

the optimizer will throw out the [ORDER BY](#) clause. This is another example of dead code elimination.

For the query:

```
SELECT column1 FROM Table1 ORDER BY column1;
```

the optimizer will use an index on `column1`, if it exists.

For the query:

```
SELECT column1 FROM Table1 ORDER BY column1+1;
```

the optimizer will use an index on `column1`, if it exists. But don't let that fool you! The index is only for finding the values. (It's cheaper to do a sequential scan of the index than a sequential scan of the table, that's why `index` is a better join type than `ALL` see [optimizer.html#optimizer-index-join-type Section?? 3.2.2.4, The `index` Join Type].) There will still be a full sort of the results.

For the query:

```
SELECT * FROM Table1 WHERE column1 > 'x' AND column2 > 'x' ORDER BY column2;
```

if both `column1` and `column2` are indexed, the optimizer will choose an index on ... `column1`. The fact that ordering takes place by `column2` values does not affect the choice of driver in this case.

See: `/sql/sql_select.cc`, `test_if_order_by_key()`, and `/sql/sql_select.cc`, `test_if_skip_sort_order()`.

`ORDER BY Optimization`, provides a description of the internal sort procedure which we will not repeat here, but urge you to read, because it describes how the buffering and the quicksort mechanisms operate.

See: `/sql/sql_select.cc`, `create_sort_index()`.

GROUP BY and Related Conditions

These are the main optimizations that take place for `GROUP BY` and related items (`HAVING`, `COUNT()`, `MAX()`, `MIN()`, `SUM()`, `AVG()`, `DISTINCT()`).

- `GROUP BY` will use an index, if one exists.
- `GROUP BY` will use sorting, if there is no index. The optimizer may choose to use a hash table.
- For the case `GROUP BY x ORDER BY x`, the optimizer will realize that the `ORDER BY` is unnecessary, because the `GROUP BY` comes out in order by `x`.
- The optimizer contains code for shifting certain `HAVING` conditions to the `WHERE` clause; however, this code is not operative at time of writing. See: `/sql/sql_select.cc`, `JOIN::optimize()`, after `#ifdef HAVE_REF_TO_FIELDS`.
- If the table handler has a quick row-count available, then the query

```
SELECT COUNT(*) FROM Table1;
```

gets the count without going through all the rows. This is true for `MyISAM` tables, but not for `InnoDB` tables. Note that the query

```
SELECT COUNT(column1) FROM Table1;
```

is not subject to the same optimization, unless `column1` is defined as `NOT NULL`.

- New optimizations exist for `MAX()` and `MIN()`. For example, consider the query

```
SELECT MAX(column1)
FROM Table1
WHERE column1 < 'a';
```

If `column1` is indexed, then it's easy to find the highest value by looking for 'a' in the index and going back to the key before that.

- The optimizer transforms queries of the form

```
SELECT DISTINCT column1 FROM Table1;
```

to

```
SELECT column1 FROM Table1 GROUP BY column1;
```

if and only if both of these conditions are true:

1. The `GROUP BY` can be done with an index. (This implies that there is only one table in the `FROM` clause, and no `WHERE` clause.)
2. There is no `LIMIT` clause.

Because `DISTINCT` is not always transformed to `GROUP BY`, do not expect that queries with `DISTINCT` will always cause ordered result sets. (You can, however, rely on that rule with `GROUP BY`, unless the query includes `ORDER BY NULL`.)

See: `/sql/sql_select.cc, opt_sum_query()`, and `/sql/sql_select.cc, remove_duplicates()`.

Other Optimizations

In this section, we discuss other, more specialized optimizations performed in the MySQL server.

NULLs Filtering for `ref` and `eq_ref` Access

This section discusses the `NULLs` filtering optimization used for `ref` and `eq_ref` joins.

Early `NULLs` Filtering

Suppose we have a join order such as this one:

```
..., tblX, ..., tblY, ...
```

Suppose further that table `tblY` is accessed via `ref` or `eq_ref` access on

```
tblY.key_column = tblX.column
```

or, in the case of `ref` access using multiple key parts, via

```
... AND tblY.key_partN = tblX.column AND ...
```

where `tblX.column` can be `NULL`. Here the early `NULLs` filtering for `ref` (or `eq_ref`) access is applied. We make the following inference:

```
(tblY.key_partN = tblX.column) => (tblX.column IS NOT NULL)
```

The original equality can be checked only after we've read the current rows of both tables `tblX` and `tblY`. The `IS NOT NULL` predicate can be checked after we've read the current row of table `tblX`. If there are any tables in the join order between `tblX` and `tblY`, the added `IS NOT NULL` check will allow us to skip accessing those tables.

This feature is implemented in these places in the server code:

- The `ref` analyzer (contained in such functions as `update_ref_and_keys()`) detects and marks equalities like that shown above by setting `KEY_FIELD::null_rejecting=TRUE`.
- After the join order has been chosen, `add_not_null_conds()` adds appropriate `IS NOT NULL` predicates to the conditions of the appropriate tables.

It is possible to add `IS NOT NULL` predicates for all equalities that could be used for `ref` access (and not for those that are actually used). However, this is currently not done.

Late **NULLs** Filtering

Suppose we have a query plan with table `tblX` being accessed via the `ref` access method:

```
tblX.key_part1 = expr1 AND tblX.key_part2 = expr2 AND ...
```

Before performing an index lookup, we determine whether any of the `expri` values is **NULL**. If it is, we don't perform the lookup, but rather immediately return that the matching tuple is not found.

This optimization reuses the `null_rejecting` attribute produced by the early **NULLs** filtering code (see [Early **NULLs** Filtering](#)). The check itself is located in the function `join_read_always_key()`.

Partitioning-Related Optimizations

This section discusses optimizations relating to MySQL Partitioning. See [Partitioning](#) for general information about the partitioning implementation in MySQL 5.1 and later.

Partition Pruning

The operation of *partition pruning* is defined as follows:

Given a query over partitioned table, match the table DDL against any **WHERE** or **ON** clauses, and find the minimal set of partitions that must be accessed to resolve the query.

The set of partitions thus obtained (hereafter referred to as *used*) can be smaller than the set of all table partitions. Partitions that did not get into this set (that is, those that were pruned away) will not be accessed at all: this is how query execution is made faster.

Non-Transactional Table Engines. With non-transactional tables such as `MyISAM`, locks are placed on entire partitioned table. It is theoretically possible to use partition pruning to improve concurrency by placing locks only on partitions that are actually used, but this is currently not implemented.

Partition pruning doesn't depend on what table engine is used. Therefore its implementation is a part of the MySQL Query Optimizer. The next few sections provide a detailed description of partition pruning.

Partition Pruning Overview

Partition pruning is performed using the following steps:

1. Analyze the **WHERE** clause and construct an *interval graph* describing the results of this analysis.
2. Walk the graph, and find sets of partitions (or subpartitions, if necessary) to be used for each interval in the graph.
3. Construct a set of partitions used for the entire query.

The description represented by the interval graph is structured in a bottom-up fashion. In the discussion that follows, we first define the term *partitioning interval*, then describe how partitioning interval are combined to make an interval graph, and then describe the graph walking process.

Partitioning Intervals

Single-Point Intervals

Let's start from simplest cases. Suppose that we have a partitioned table with *N* columns, using partitioning type `p_type` and the partitioning function `p_func`, represented like this:

```
CREATE TABLE t (columns)  
PARTITION BY p_type(p_func(col1, col2,... colN)...);
```

Suppose also that we have a **WHERE** clause of the form

```
WHERE t.col1=const1 AND t.col2=const2 AND ... t.colN=constN
```


We can calculate `p_func(const1, const2 ... constN)` and discover which partition can contain records matching the `WHERE` clause. Note that this process works for all partitioning types and all partitioning functions.

Note: This process works only if the `WHERE` clause is of the exact form given above — that is, each column in the table must be tested for equality with some arbitrary constant (not necessarily the same constant for each column). For example, if `col1=const1` were missing from the example `WHERE` clause, then we would not be able to calculate the partitioning function value and so would be unable to restrict the set of partitions to those actually used.

Interval Walking

Let a partitioned table `t` be defined with a set of column definitions `columns`, a partitioning type `p_type` using a partitioning function `p_func` taking an integer column `int_col`, as shown here:

```
CREATE TABLE t (columns)
PARTITION BY
p_type(p_func(int_col))
...
```

Now suppose that we have a query whose `WHERE` clause is of the form

```
WHERE const1 <= int_col <= const2
```

We can reduce this case to a number of cases of single-point intervals by converting the `WHERE` clause into the following relation:

```
int_field=const1 OR int_field=const1 + 1 OR int_field=const1 + 2 OR ... OR int_field=const2
```

In the source code this conversion is referred to as *interval walking*. Walking over short intervals is not very expensive, since we can reduce the number of partitions to scan to a small number. However, walking over long intervals may not be very efficient there will be lots of numbers to examine, and we are very likely to find that all partitions need to be scanned.

The threshold for interval walking is determined by

```
#define MAX_RANGE_TO_WALK=10
```

Note

The logic of the previous example also applies to a relation such as this one:

```
const1
```

Interval Mapping

Let a partitioned table `t` be defined as follows:

```
CREATE TABLE t (columns)
PARTITION BY RANGE|LIST(unary_ascending_function(column))
```

Suppose we have a query on table `t` whose `WHERE` clause is of one of the forms shown here:

- `const1 <= t.column <= const2`
- `t.column <= const2`
- `const1 <= t.column`

Since the partitioning function is ascending, the following relationship holds:

```
const1 <= t.col <= const2
=> p_func(const1) <=
```

```
p_func(t.column) <= p_func(const2)
```

Using A and B to denote the leftmost and rightmost parts of this relation, we can rewrite it like this:

```
A <= p_func(t.column) <= B
```

Note: In this instance, the interval is closed and has two bounds. However, similar inferences can be performed for other kinds of intervals.

For **RANGE** partitioning, each partition occupies one interval on the partition function value axis, and the intervals are disjoint, as shown here:

	p0	p1	p2	
table partitions	-----x-----x-----x----->			
search interval	---x=====x----->			
	A	B		

A partition needs to be accessed if and only if its interval has a non-empty intersection with the search interval [A , B].

For **LIST** partitioning, each partition covers a set of points on the partition function value axis. Points produced by various partitions may be interleaved, as shown here:

	p0	p1	p2	p1	p1	p0	
table partitions	--+--+--+--+--+--+--+--+--+>						
search interval	---x=====x----->	A	B				

A partition needs to be accessed if it has at least one point in the interval [A, B]. The set of partitions used can be determined by running from A to B and collecting partitions that have their points within this range.

Subpartitioning Intervals

In the previous sections we've described ways to infer the set of used partitions from "elementary" WHERE clauses. Everything said there about partitions also applies to subpartitions (with exception that subpartitioning by RANGE or LIST is currently not possible).

Since each partition is subpartitioned in the same way, we'll find which subpartitions should be accessed within each partition.

From WHERE Clauses to Intervals

Previous sections deal with inferring the set of partitions used from WHERE clauses that represent partitioning or subpartitioning intervals. Now we look at how MySQL extracts intervals from arbitrary WHERE clauses.

The extraction process uses the *Range Analyzer* a part of the MySQL optimizer that produces plans for the range access method. This is because the tasks are similar. In both cases we have a WHERE clause as input: the range access method needs index ranges (that is, intervals) to scan; partition pruning module needs partitioning intervals so that it can determine which partitions should be used.

For range access, the Range Analyzer is invoked with the WHERE clause and descriptions of table indexes. Each index is described by an ordered list of the columns which it covers:

```
(keypart1, keypart2, ..., keypartN)
```

For partition pruning, Range Analyzer is invoked with the WHERE clause and a list of table columns used by the partitioning and subpartitioning functions:

```
(part_col1, part_col2, ... part_colN,  
subpart_col1, subpart_col2, ... subpart_colM)
```

The result of the Range Analyzer's work is known as a `SEL_ARG` graph. *This is a complex (and not yet fully documented) structure, which we will not attempt to describe here. What's important for the current discussion is that we can walk over it and collect partitioning and subpartitioning intervals.*

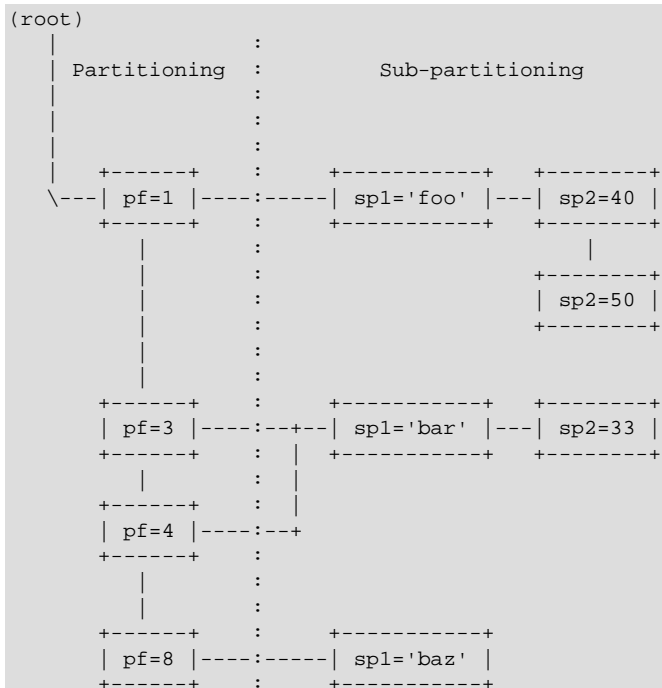
The following example illustrates the structure and the walking process. Suppose a table `t` is partitioned as follows:

```
CREATE TABLE t (... , pf INT, sp1 CHAR(5), sp2 INT, ... )
  PARTITION BY LIST (pf)
  SUBPARTITION BY HASH(sp1, sp2) (
    PARTITION p0 VALUES IN (1),
    PARTITION p1 VALUES IN (2),
    PARTITION p2 VALUES IN (3),
    PARTITION p3 VALUES IN (4),
    PARTITION p4 VALUES IN (5),
  );
```

Now suppose that a query on table `t` has a highly complex `WHERE` clause, such as this one:

```
pf=1 AND (sp1='foo' AND sp2 IN (40,50))
OR
(pf=3 OR pf=4) AND sp1='bar' AND sp2=33
OR
((pf=3 OR pf=4) AND sp1=5)
OR
p=8
```

The `SEL_ARG` graph for this is shown here:



In the previous diagram, vertical edges (`|`) represent `OR` and the horizontal ones (`-`) represent `AND` (the line with both horizontal and vertical segments also represents `AND`).

The partition-pruning code walks the graph top to bottom and from left to right, making these inferences:

1. Start with an empty set of used partitions at the topmost and leftmost interval.
 - a. Perform interval analysis for `pf=1`; find a corresponding set of partitions P0; move right.
 - b. Move right again, to `sp2=40`.
 - c. Analyze the interval `sp1='foo' AND sp2=40` interval; find that it covers rows in some subpartition SP1. Make first inference: within each partition making up set P0, mark subpartition SP1 as used.
 - d. Move down to `sp2=50`.
 - e. Analyze the interval `sp1='foo' AND sp2=50`, finding that it covers rows in some subpartition SP2. Make another inference: within each partition of set P0, mark subpartition SP2 as used.
 - f. Move back to `pf=1`, and then down to `pf=3`.
 - g. Perform interval analysis for `pf=3`; find a corresponding set of partitions P1; move right.
 - h. Move right again, to `sp2=33`.
 - i. Analyze the interval `sp1='foo' AND sp2=33`, find that it covers rows in a subpartition SP3. Make another inference: within each partition from set P1, mark subpartition SP3 as used.
 - j. Move back to `pf=3`, then down to `pf=4`.
 - k. Perform interval analysis for `pf=4`; find a corresponding set of partitions P2; move right.
 - l. Perform moves and inferences analogous to what we did to the right of `pf=3`. There is some potential inefficiency due to the fact that that we will analyze the interval for `sp1='foo' AND sp2=33` again, but this should not have much impact on overall performance.
 - m. Move back to `pf=3`, then down to `pf=8`.
 - n. Perform interval analysis for `pf=8`; find a corresponding set of partitions P3, move right.
 - o. Now we've arrived at `sp1='baz'`, and find that we can't move any further to the right and can't construct a subpartitioning interval. We remember this, and move back to `pf=8`.
 - p. In the previous step we could not limit the set of subpartitions, so we make this inference: for every partition in set P3, assume that all subpartitions are active, and mark them as such.
2. Try to move down from `pf=8`; find that there is nothing there; this completes the graph analysis.

Note: In certain cases the result of the `RANGE` optimizer will be several `SEL_ARG` graphs that are to be combined using `OR` or `AND` operators. This happens for `WHERE` clauses which either are very complicated or do not allow for the construction of a single list of intervals. In such cases, the partition pruning code takes appropriate action, an example being this query:

```
SELECT * FROM t1 WHERE partition_id=10 OR subpartition_id=20
```

No single list of intervals can be constructed in this instance, but the partition pruning code correctly infers that the set of partitions used is a union of:

1. All subpartitions within the partition containing rows with `partition_id=10`; and `subpartition_id=20` within each partition.

Partition Pruning in the Source Code

Here is a short walkthrough of what is where in the code:

- `sql/opt_range.cc` :

This file contains the implementation of what is described in [From WHERE Clauses to Intervals](#). The entry point is the function `prune_partitions()`. There are also detailed code-level comments about partition pruning; search for `PartitionPruningModule` to find the starting point.

- `sql/partition_info.h`:

```
class partition_info {
...
/*
  Bitmap of used (i.e. not pruned away) partitions. This is where result
  of partition pruning is stored.
*/
MY_BITMAP used_partitions;

/*
  "virtual function" pointers to functions that perform interval analysis
  on this partitioned table (used by the code in opt_range.cc)
*/
get_partitions_in_range_iter get_part_iter_for_interval;
get_partitions_in_range_iter get_subpart_iter_for_interval;
};
```

- `sql/sql_partition.cc`:

This file contains the functions implementing all types of partitioning interval analysis.

Partition Selection

If a partitioned table is accessed in a series of index lookups (that is, using the `ref`, `eq_ref`, or `ref_or_null` access methods), MySQL checks to see whether it needs to make index lookups in all partitions or that it can limit access to a particular partition. This is performed for each index lookup.

Consider this example:

```
CREATE TABLE t1 (a INT, b INT);

INSERT INTO t1 VALUES (1,1),(2,2),(3,3);

CREATE TABLE t2 (
  keypart1 INT,
  keypart2 INT,
  KEY(keypart1, keypart2)
)
PARTITION BY HASH(keypart2);

INSERT INTO t2 VALUES (1,1),(2,2),(3,3);
```

The query

```
SELECT * FROM t1, t2
  WHERE t2.keypart1=t1.a
  AND   t2.keypart2=t1.b;
```

is executed using this algorithm:

```
(for each record in t1:)
{
  t2->index_read({current-value-of(t1.a), current-value-of(t1.b)});
  while( t2->index_next_same() )
    pass row combination to query output;
}
```

In the `index_read()` call, the partitioned table handler will discover that the value of all partitioning columns (in this case, the single column `b`) is fixed, and find a single partition to access. If this partition was pruned away, then no partitions will be accessed at all.

Chapter 9. Tracing the Optimizer

Table of Contents

Typical Usage	81
System Variables Controlling the Trace	81
The <code>INFORMATION_SCHEMA.OPTIMIZER_TRACE</code> Table	82
Traceable Queries	82
Automatic Trace Purging	82
Tuning Trace Purging	82
Tracing Memory Usage	83
Privilege Checking	83
Interaction with the <code>--debug</code> Option	83
The <code>optimizer_trace</code> System Variable	83
The <code>end_markers_in_json</code> System Variable	83
Selecting Optimizer Features to Trace	84
General Trace Structure	84
Example	84
Nicely Displaying a Trace	89
Preventing Use of Optimizer Trace	90
Testing	90
Implementation	90

Since milestone version 5.6.3, a new tracing capability has been added to the MySQL optimizer. The interface is provided by a set of `optimizer_trace_xxx` system variables and the `INFORMATION_SCHEMA.OPTIMIZER_TRACE` table, but is subject to change.

Typical Usage

```
# Turn tracing on (it's off by default):
SET optimizer_trace="enabled=on";
SELECT ...; # your query here
SELECT * FROM INFORMATION_SCHEMA.OPTIMIZER_TRACE;
# possibly more queries...
# When done with tracing, disable it:
SET optimizer_trace="enabled=off";
```

A session can trace only statements which it executes; it cannot see a trace of another session.

System Variables Controlling the Trace

A brief overview from "mysqld --verbose --help":

```
--optimizer-trace=name
    Controls tracing of the Optimizer:
--optimizer-trace-features=name
    Enables/disables tracing of selected features of the
    Optimizer:
    optimizer_trace_features=option=val[,option=val...],
    where option is one of {greedy_search, range_optimizer,
    dynamic_range, repeated_subselect} and val is one of {on,
    off, default}
--optimizer-trace-limit=#
    Maximum number of shown optimizer traces
--optimizer-trace-max-mem-size=#
    Maximum allowed cumulated size of stored optimizer traces
--optimizer-trace-offset=#
    Offset of first optimizer trace to show; see manual
--end-markers-in-json=#
    In JSON output ("EXPLAIN FORMAT=JSON" and optimizer
    trace), if set to 1, repeats the structure's key (if it
    has one) near the closing bracket
```

Those will be described in more detail below.

The `INFORMATION_SCHEMA.OPTIMIZER_TRACE` Table

The `OPTIMIZER_TRACE` table contains information about traced statements. The table has these columns:

- `QUERY`: the statement text
- `TRACE`: the trace, in JSON format (see json.org: basically it has scalars (number, string, bool) and structures (either arrays or associative arrays))
- `MISSING_BYTES_BEYOND_MAX_MEM_SIZE` (explained further below)
- `MISSING_PRIVILEGES` (explained further below).

Traceable Queries

They are: `SELECT`; `INSERT` or `REPLACE` (with `VALUES` or `SELECT`); `UPDATE/DELETE` and their multi-table variants; all the previous ones prefixed by `EXPLAIN`; `SET` (unless it manipulates the `optimizer_trace` system variable); `DO`; `DECLARE/CASE/IF/RETURN` (stored routines language elements); `CALL`. If one of those statements is prepared and executed in separate steps, preparation and execution are separately traced.

Automatic Trace Purging

By default, each new trace overwrites the previous trace. Thus, if a statement contains substatements (example: invokes stored procedures, stored functions, triggers), the top statement and substatements each generate one trace, but at the execution's end only the last substatement's trace is visible. A user who wants to see the trace of another substatement, can enable/disable tracing around the desired substatement, but this requires editing the routine's code, which may not be possible. Another solution is to tune trace purging.

Tuning Trace Purging

This is done with

```
SET optimizer_trace_offset=<OFFSET>, optimizer_trace_limit=<LIMIT>
```

where `OFFSET` is a signed integer, and `LIMIT` is a positive integer. The default for `optimizer_trace_offset` is -1; the default for `optimizer_trace_limit` is 1. The `SET` statement has the following effects:

- All remembered traces are cleared
- A later `SELECT` on the `OPTIMIZER_TRACE` table returns the first `LIMIT` traces of the `OFFSET` oldest remembered traces (if `OFFSET >= 0`), or the first `LIMIT`

traces of the `(-OFFSET)` newest remembered traces (if `OFFSET < 0`).

For example, a combination of `OFFSET=-1` and `LIMIT=1` will make the last trace be shown (as is default), `OFFSET=-2` and `LIMIT=1` will make the next-to-last be shown, `OFFSET=-5` and `LIMIT=5` will make the last five traces be shown. Such negative `OFFSET` can be useful when one knows that the interesting substatements are the few last ones of a stored routine, like this:

```
SET optimizer_trace_offset=-5, optimizer_trace_limit=5;
CALL stored_routine(); # more than 5 substatements in this routine
SELECT * FROM information_schema.OPTIMIZER_TRACE; # see only last 5 traces
```

On the opposite, a positive `OFFSET` can be useful when one knows that the interesting substatements are the few first ones of a stored routine.

The more accurately those two variables are adjusted, the less memory is used. For example, OFFSET=0 and LIMIT=5 will use memory to remember 5 traces, so if only the three first are needed, OFFSET=0 and LIMIT=3 is better (tracing stops after LIMIT traces, so the 4th and 5th trace are not created and take up no memory). A stored routine may have a loop which executes many substatements and thus generates many traces, which would use a lot of memory; proper OFFSET and LIMIT can restrict tracing to one iteration of the loop for example. This also gains speed, as tracing a substatement impacts performance.

If OFFSET>=0, only LIMIT traces are kept in memory. If OFFSET<0, that is not true: instead, (-OFFSET) traces are kept in memory; indeed even if LIMIT is smaller than (-OFFSET), so excludes the last statement, the last statement must still be traced because it will be inside LIMIT after executing one more statement (remember that OFFSET<0 is counted from the end: the "window" slides as more statements execute).

Such memory and speed gains are the reason why optimizer_trace_offset and optimizer_trace_limit, which are restrictions at the trace producer level, are offered. They are better than using

```
SELECT * FROM OPTIMIZER_TRACE LIMIT <LIMIT> OFFSET <OFFSET>;
```

which is a restriction on the trace consumer level and saves almost nothing.

Tracing Memory Usage

Each remembered trace is a String. It is extended (with realloc()) as optimization progresses and appends data to it. The optimizer_trace_max_mem_size variable sets a limit on the total amount of memory used by all currently remembered traces: If this limit is reached, the current trace isn't extended (so it will be incomplete), and the MISSING_BYTES_BEYOND_MAX_MEM_SIZE column will show the number of bytes missing from this trace. The optimizer_trace_max_mem_size variable has a small default (16kB) to protect the innocent, it's often necessary to grow it for real-life traces.

Privilege Checking

In complex scenarios where the query uses SQL SECURITY DEFINER views or stored routines, it may be that a user is denied from seeing the trace of its query because it lacks some extra privileges on those objects. In that case, the trace will be shown as empty and the MISSING_PRIVILEGE column will show "1".

Interaction with the --debug Option

Anything written to the trace is automatically written to the --debug file.

The `optimizer_trace` System Variable

The optimizer_trace system variable has these on/off switches:

- enabled: allows to enable/disable tracing
- one_line: if on, the trace will have no whitespace; it's unreadable for humans but readable for JSON parsers (they ignore whitespace); the only advantage is a saving on space.

The `end_markers_in_json` System Variable

If a JSON structure is large, it's difficult to pair its closing bracket with its opening bracket; to help the reader (but this is not JSON-compliant), setting @@end_markers_in_json=on repeats the structure's key (if it has one) near the closing bracket. It also affects 'EXPLAIN FORMAT=JSON' in the same way. Note that before MySQL 5.6.6, this variable didn't exist and was rather a switch in the @@optimizer_trace variable (it was set with 'set optimizer_trace="end_marker=on";').

Selecting Optimizer Features to Trace

Some features in the optimizer can be invoked many times during statement optimization and execution, and thus can make the trace grow beyond reason. They are:

- Greedy search: with a N-table join, this could explore factorial(N) plans
- Range optimizer
- Dynamic range optimization (known as "range checked for each record" in EXPLAIN output: each outer row causes a re-run of the range optimizer)
- Subqueries: a subquery in WHERE clause may be executed once per row.

Those features can be excluded from tracing using the `optimizer_trace_features` system variable, which has these on/off switches:

- `greedy_search`: if off, greedy search is not traced
- `range_optimizer`: if off, range optimizer is not traced
- `dynamic_range`: if off, only the first call to the range optimizer on this JOIN_TAB::SQL_SELECT is traced
- `repeated_subselect`: if off, only the first execution of this Item_subselect is traced.

General Trace Structure

A trace follows closely the actual execution path: there is a join-preparation object, a join-optimization object, a join-execution object, for each JOIN. Query transformations (IN->EXISTS, outer join to inner join...), simplifications (elimination of clauses), equality propagation are shown in subobjects. Calls to the range optimizer, cost evaluations, reasons why an access path is chosen over another one, or why a sorting method is chosen over another one, are shown too. It is far from showing everything happening in the optimizer, but we plan to show more information in the future.

Example

Here we take an example from the test suite.

```
#
# Tracing of ORDER BY & GROUP BY simplification.
#
SET OPTIMIZER_TRACE="enabled=on",END_MARKERS_IN_JSON=on; # be readable
SET OPTIMIZER_TRACE_MAX_MEM_SIZE=1000000; # avoid small default
CREATE TABLE t1 (
  pk INT, col_int_key INT,
  col_varchar_key VARCHAR(1), col_varchar_nokey VARCHAR(1)
);
INSERT INTO t1 VALUES
(10,7,'v','v'),(11,0,'s','s'),(12,9,'l','l'),(13,3,'y','y'),(14,4,'c','c'),
(15,2,'i','i'),(16,5,'h','h'),(17,3,'q','q'),(18,1,'a','a'),(19,3,'v','v'),
(20,6,'u','u'),(21,7,'s','s'),(22,5,'y','y'),(23,1,'z','z'),(24,204,'h','h'),
(25,224,'p','p'),(26,9,'e','e'),(27,5,'i','i'),(28,0,'y','y'),(29,3,'w','w');
CREATE TABLE t2 (
  pk INT, col_int_key INT,
  col_varchar_key VARCHAR(1), col_varchar_nokey VARCHAR(1),
  PRIMARY KEY (pk)
);
INSERT INTO t2 VALUES
(1,4,'b','b'),(2,8,'y','y'),(3,0,'p','p'),(4,0,'f','f'),(5,0,'p','p'),
(6,7,'d','d'),(7,7,'f','f'),(8,5,'j','j'),(9,3,'e','e'),(10,188,'u','u'),
(11,4,'v','v'),(12,9,'u','u'),(13,6,'i','i'),(14,1,'x','x'),(15,5,'l','l'),
(16,6,'q','q'),(17,2,'n','n'),(18,4,'r','r'),(19,231,'c','c'),(20,4,'h','h'),
(21,3,'k','k'),(22,3,'t','t'),(23,7,'t','t'),(24,6,'k','k'),(25,7,'g','g'),
```

```
(26,9,'z','z'),(27,4,'n','n'),(28,4,'j','j'),(29,2,'l','l'),(30,1,'d','d'),
(31,2,'t','t'),(32,194,'y','y'),(33,2,'i','i'),(34,3,'j','j'),(35,8,'r','r'),
(36,4,'b','b'),(37,9,'o','o'),(38,4,'k','k'),(39,5,'a','a'),(40,5,'f','f'),
(41,9,'t','t'),(42,3,'c','c'),(43,8,'c','c'),(44,0,'r','r'),(45,98,'k','k'),
(46,3,'l','l'),(47,1,'o','o'),(48,0,'t','t'),(49,189,'v','v'),(50,8,'x','x'),
(51,3,'j','j'),(52,3,'x','x'),(53,9,'k','k'),(54,6,'o','o'),(55,8,'z','z'),
(56,3,'n','n'),(57,9,'c','c'),(58,5,'d','d'),(59,9,'s','s'),(60,2,'j','j'),
(61,2,'w','w'),(62,5,'f','f'),(63,8,'p','p'),(64,6,'o','o'),(65,9,'f','f'),
(66,0,'x','x'),(67,3,'q','q'),(68,6,'g','g'),(69,5,'x','x'),(70,8,'p','p'),
(71,2,'q','q'),(72,120,'q','q'),(73,25,'v','v'),(74,1,'g','g'),(75,3,'l','l'),
(76,1,'w','w'),(77,3,'h','h'),(78,153,'c','c'),(79,5,'o','o'),(80,9,'o','o'),
(81,1,'v','v'),(82,8,'y','y'),(83,7,'d','d'),(84,6,'p','p'),(85,2,'z','z'),
(86,4,'t','t'),(87,7,'b','b'),(88,3,'y','y'),(89,8,'k','k'),(90,4,'c','c'),
(91,6,'z','z'),(92,1,'t','t'),(93,7,'o','o'),(94,1,'u','u'),(95,0,'t','t'),
(96,2,'k','k'),(97,7,'u','u'),(98,2,'b','b'),(99,1,'m','m'),(100,5,'o','o');
SELECT SUM(alias2.col_varchar_nokey) , alias2.pk AS field2 FROM t1 AS alias1
STRAIGHT_JOIN t2 AS alias2 ON alias2.pk = alias1.col_int_key WHERE alias1.pk
GROUP BY field2 ORDER BY alias1.col_int_key,alias2.pk ;
SUM(alias2.col_varchar_nokey) field2
0 1
0 2
0 3
0 4
0 5
0 6
0 7
0 9
```

Now we look at the trace:

```
SELECT * FROM INFORMATION_SCHEMA.OPTIMIZER_TRACE;
QUERY TRACE MISSING_BYTES_BEYOND_MAX_MEM_SIZE INSUFFICIENT_PRIVILEGES
SELECT SUM(alias2.col_varchar_nokey) , alias2.pk AS field2 FROM t1 AS alias1
STRAIGHT_JOIN t2 AS alias2 ON alias2.pk = alias1.col_int_key WHERE alias1.pk
GROUP BY field2 ORDER BY alias1.col_int_key,alias2.pk {
```

This was the first column: it repeats the query (this is a useful mark when several traces are remembered thanks to `optimizer_trace_offset` and `optimizer_trace_limit`). Now the trace. The statement's execution is naturally made of "steps":

```
"steps": [
  {
    "join_preparation": {
```

This is a join's preparation

```
"select#": 1,
```

for the first SELECT of the statement (which has only one, here). Here are steps of the join's preparation:

```
"steps": [
  {
    "expanded_query": "/* select#1 */ select \
      sum(`test`.`alias2`.`col_varchar_nokey`) AS \
      `SUM(alias2.col_varchar_nokey)`,`test`.`alias2`.`pk` AS `field2` \
    from (`test`.`t1` `alias1` straight_join `test`.`t2` `alias2` \
    on((`test`.`alias2`.`pk` = `test`.`alias1`.`col_int_key`))) \
    where `test`.`alias1`.`pk` \
    group by `test`.`alias2`.`pk` \
    order by `test`.`alias1`.`col_int_key`,`test`.`alias2`.`pk`"
  }
]
```

Above is the query as it is in the join's preparation: fields have been resolved to their database and table, and each SELECT is annotated with its number (useful with subqueries).

```
] /* steps */
  } /* join_preparation */
},
{
```

Off to optimization:

```
"join_optimization": {
  "select#": 1,
  "steps": [
    {
      "condition_processing": {
        "condition": "WHERE",
        "original_condition": "(`test`.`alias1`.`pk` and \
(`test`.`alias2`.`pk` = `test`.`alias1`.`col_int_key`))",
        "steps": [
          {
            "transformation": "equality_propagation",
            "resulting_condition": "(`test`.`alias1`.`pk` and \
multiple equal(`test`.`alias2`.`pk`, \
`test`.`alias1`.`col_int_key`))"
          },
          {
            "transformation": "constant_propagation",
            "resulting_condition": "(`test`.`alias1`.`pk` and \
multiple equal(`test`.`alias2`.`pk`, \
`test`.`alias1`.`col_int_key`))"
          },
          {
            "transformation": "trivial_condition_removal",
            "resulting_condition": "(`test`.`alias1`.`pk` and \
multiple equal(`test`.`alias2`.`pk`, \
`test`.`alias1`.`col_int_key`))"
          }
        ]
      }
    ]
  }
}
```

Not much happened in equality propagation above.

```
    }
  ] /* steps */
} /* condition_processing */
},
{
  "ref_optimizer_key_uses": [
    {
      "database": "test",
      "table": "alias2",
      "field": "pk",
      "equals": "`test`.`alias1`.`col_int_key`",
      "null_rejecting": true
    }
  ]
}
```

A possible ref access has been identified, and it is NULL-rejecting: any NULL value in ``test`.`alias1`.`col_int_key`` cannot have a match (it could have a match if the operator were `<=>`).

```
    }
  ] /* ref_optimizer_key_uses */
},
{
  "records_estimation": [
    {
      "database": "test",
      "table": "alias1",
      "const_keys_added": {
        "keys": [
          ] /* keys */,
        "cause": "group_by"
      } /* const_keys_added */,
      "range_analysis": {
        "table_scan": {
          "records": 20,
          "cost": 8.1977
        } /* table_scan */
      } /* range_analysis */
    }
  ]
}
```

Now for every table in the query we estimate the cost of, and number of records returned by, a table scan, a range access,

```
"records_estimation": [
  {
    "database": "test",
    "table": "alias1",
    "const_keys_added": {
      "keys": [
        ] /* keys */,
      "cause": "group_by"
    } /* const_keys_added */,
    "range_analysis": {
      "table_scan": {
        "records": 20,
        "cost": 8.1977
      } /* table_scan */
    } /* range_analysis */
  },
  {
    "database": "test",
    "table": "alias2",
    "const_keys_added": {
      "keys": [
        ] /* keys */,
      "cause": "group_by"
    } /* const_keys_added */,
    "range_analysis": {
      "table_scan": {
        "records": 20,
        "cost": 8.1977
      } /* table_scan */
    } /* range_analysis */
  }
]
```

```
{
  "database": "test",
  "table": "alias2",
  "const_keys_added": {
    "keys": [
      "PRIMARY"
    ] /* keys */,
    "cause": "group_by"
  } /* const_keys_added */,
  "range_analysis": {
    "table_scan": {
      "records": 100,
      "cost": 24.588
    } /* table_scan */,
    "potential_range_indices": [
      {
        "index": "PRIMARY",
        "usable": true,
        "key_parts": [
          "pk"
        ] /* key_parts */
      }
    ] /* potential_range_indices */,
    "setup_range_conditions": [
    ] /* setup_range_conditions */,
    "group_index_range": {
      "chosen": false,
```

Not possible to use GROUP_MIN_MAX because it can handle only one table, and we have two in the join:

```
"cause": "not_single_table"
} /* group_index_range */
```

No range access is possible.

```
} /* range_analysis */
}
] /* records_estimation */
},
{
```

Finding an optimal order for tables (greedy search); actually as this is STRAIGHT_JOIN only the requested order is explored, and access methods are selected:

```
"considered_execution_plans": [
  {
    "database": "test",
    "table": "alias1",
    "best_access_path": {
      "considered_access_paths": [
        {
          "access_type": "scan",
          "records": 20,
          "cost": 2.0977,
          "chosen": true
```

Table scan be chosen!

```
        }
      ] /* considered_access_paths */
    } /* best_access_path */,
    "cost_for_plan": 6.0977,
    "records_for_plan": 20,
```

We estimate that reading the first table, and applying any conditions to it, will yield 20 rows.

```
"rest_of_plan": [
  {
    "database": "test",
    "table": "alias2",
```

```

    "best_access_path": {
      "considered_access_paths": [
        {
          "access_type": "ref",
          "index": "PRIMARY",
          "records": 1,
          "cost": 20.2,
          "chosen": true

```

We choose ref access on the primary key of alias2.

```

      },
      {
        "access_type": "scan",
        "using_join_cache": true,
        "records": 75,
        "cost": 7.4917,
        "chosen": false

```

But not table scan, because its amount of records (75) is far greater than that of ref access (1).

```

      }
    ] /* considered_access_paths */
  } /* best_access_path */
  "cost_for_plan": 30.098,
  "records_for_plan": 20,
  "chosen": true
}
] /* rest_of_plan */
}
] /* considered_execution_plans */
},
{

```

Now that the order of tables is fixed, we can split the WHERE condition in chunks which can be tested early ("push down of conditions down the join tree"):

```

"attaching_conditions_to_tables": {
  "original_condition": "((`test`.`alias2`.`pk` = \
`test`.`alias1`.`col_int_key`) and `test`.`alias1`.`pk`)",
  "attached_conditions_computation": [
    ] /* attached_conditions_computation */
  "attached_conditions_summary": [
    {
      "database": "test",
      "table": "alias1",
      "attached": "(`test`.`alias1`.`pk` and \
(`test`.`alias1`.`col_int_key` is not null))"

```

This condition above can be tested on rows of alias1 without even reading rows of alias2.

```

    },
    {
      "database": "test",
      "table": "alias2",
      "attached": null
    }
  ] /* attached_conditions_summary */
} /* attaching_conditions_to_tables */
},
{

```

Now we try to simplify ORDER BY:

```

"clause_processing": {
  "clause": "ORDER BY",
  "original_clause": "`test`.`alias1`.`col_int_key`,`test`.`alias2`.`pk`",
  "items": [
    {
      "item": "`test`.`alias1`.`col_int_key`"
    }
  ],

```

```
{
  "item": "`test`.`alias2`.`pk`",
  "eq_ref_to_preceding_items": true
}
```

Because the WHERE clause contains `alias2`.`pk`=`alias1`.`col_int_key`, ordering by both columns is a waste: can just order by the first column, the second will always be equal to it.

```
}
] /* items */,
"resulting_clause_is_simple": true,
"resulting_clause": "`test`.`alias1`.`col_int_key`"
```

So we get a shorter ORDER BY clause - and this is not visible in EXPLAIN or EXPLAIN EXTENDED!! This simplification can be worth it: this shorter clause, being single-column and single-table, could be implemented by an index scan...

```
    } /* clause_processing */
  },
  {
    "clause_processing": {
      "clause": "GROUP BY",
      "original_clause": "`test`.`alias2`.`pk`",
      "items": [
        {
          "item": "`test`.`alias2`.`pk`"
        }
      ] /* items */,
      "resulting_clause_is_simple": false,
      "resulting_clause": "`test`.`alias2`.`pk`"
    } /* clause_processing */
  },
  {
    "refine_plan": [
      {
        "database": "test",
        "table": "alias1",
        "scan_type": "table"
      },
      {
        "database": "test",
        "table": "alias2"
      }
    ] /* refine_plan */
  }
] /* steps */
} /* join_optimization */
},
{
```

Now we execute the join, and nothing interesting happens here:

```
  "join_execution": {
    "select#": 1,
    "steps": [
      ] /* steps */
    } /* join_execution */
  }
] /* steps */
} 0 0
```

This was just an example. All traces have the same basic structure, but if a statement uses subqueries, there are several join preparations/optimizations/executions, subquery-specific transformations not shown here...

Nicely Displaying a Trace

Looking at a trace in the "mysql" command-line client can be cumbersome (though the command "pager less" makes it better). An alternative can be to send the trace to a file:

```
SELECT TRACE INTO DUMPFILE <filename> FROM INFORMATION_SCHEMA.OPTIMIZER_TRACE;
```

and then passing this file to some JSON viewer, for example the JsonView Firefox add-on (shows objects in colours, allows object collapsing/expanding). Make sure to use INTO DUMPFILE and not INTO OUTFILE because the latter escapes newlines. Also, have the @@end_marker_in_json variable off: when it's on, the trace is more human-readable but not JSON-compliant.

Preventing Use of Optimizer Trace

If for some reason, as DBA of a MySQL Server, you wish to prevent all users from seeing traces of their queries, start the server with

```
--maximum-optimizer-trace-max-mem-size=0 --optimizer-trace-max-mem-size=0
```

This will set the maximum size to 0, and prevent users from changing this limit, thus all traces will be truncated to 0 bytes.

Testing

This feature is tested in mysql-test/suite/opt_trace and unittest/gunit/opt_trace-t.

Implementation

See files sql/opt_trace* , starting from sql/opt_trace.h. A trace is started by creating an instance of Opt_trace_start; information is added to this trace by creating instances of Opt_trace_object and Opt_trace_array, and by using the add() functions of those classes...

Chapter 10. Memory Allocation

Table of Contents

Memory Allocation in the MySQL Server (sql Directory)	91
Memory Allocation in a Library or Storage Engine	91

Memory Allocation in the MySQL Server ([sql](#) Directory)

The basic logic to use:

All things that are used only for the duration of a query are allocated in THD::mem_root through `sql_alloc()` or `thd->alloc()` except:

- Things that may grow, like string buffers of type `String`. See `sql/sql_string.cc`.
- Large blocks of memory used in one state of the query that can be released early. These are things like sort buffers, range trees, etc.
- Things in libraries that are outside of MySQL's control (like hash tables).

Things that are needed a longer time should be allocated with `my_malloc()` or through another `MEMROOT`.

Some objects have their own `MEMROOT`:

- `TABLE`
- `TABLE_SHARE`
- `Query_arena`
- `st_transactions`

Memory Allocation in a Library or Storage Engine

For the simple case, use the functions in `mysys/my_malloc.c`:

- `void *my_malloc(size_t size, myf my_flags);`
- `void *my_memdup(const void *from, size_t length, myf my_flags);`
- `char *my_strdup(const char *from, myf my_flags);`
- `char *my_strndup(const char *from, size_t length, myf my_flags);`
- `void my_free(void *ptr, myf my_flags);`

Alternatively, if you want to allocate many pieces at once, use `my_multi_malloc()` from `mysys/mulalloc.c`

- `void *my_multi_malloc(myf myFlags, ...)`

For the complex case where you want to allocate a lot and free things at once, use the `MEM_ROOT` object defined in `mysys/my_alloc.c`

Functions to use:

- `void init_alloc_root(MEM_ROOT *mem_root, size_t block_size, size_t pre_alloc_size);`

- `void *alloc_root(MEM_ROOT *mem_root, size_t length);`
- `void *multi_alloc_root(MEM_ROOT *root, ...);`
- `void free_root(MEM_ROOT *root, myf MyFlags);`

Chapter 11. Important Algorithms and Structures

Table of Contents

The <code>Item</code> Class	93
How MySQL Does Sorting (<code>filesort</code>)	94
Bulk Insert	95
How MySQL Does Caching	95
How MySQL Uses the Join Buffer Cache	96
How MySQL Handles <code>FLUSH TABLES</code>	97
Full-Text Search	97
<code>FLOAT</code> and <code>DOUBLE</code> Data Type Representation	101
Date and Time Data Type Representation	105
Threads	106
Character Sets and Collations	107
Error Flags and Functions	111
Functions in the <code>mysys</code> Library	112
Bitmaps	115

MySQL uses many different algorithms and structures. This chapter tries to describe some of them.

The `Item` Class

To us, the word *Item* means more than just “thingamabob”; it is a technical term with a precise definition in the context of our source code. `Item` is a class. Each instance of the `Item` class has:

- an analogue in the SQL language
- a value
- a data type descriptor

All of the following SQL “thingamabobs” are modeled in the `Item` class:

- literals
- column references
- session or global variables
- procedure variables
- parameters
- SQL functions (not a surprise since SQL functions have data types and return values).

In the *function* category we include operators such as `+` and `||`, because operators are merely functions that return values. We also include operators such as `=` and `LIKE`, which are operators that return boolean values. Consider the following statement:

```
SELECT UPPER(column1) FROM t WHERE column2 = @x;
```

For this statement, MySQL will need to store a list of items for the select list ('column1' column reference and `UPPER` function), and a list of items for the `WHERE` clause ('column2' column reference and '@x' variable and '=' operator).

Terminology: an `Item` instance in a MySQL program roughly corresponds to a “site”, which according to the standard SQL definition is “a place that holds an instance of a value of a specified data type”. Another word that you’ll see often in MySQL code is “field”, which means column reference, and the `Item_field` subclass is generally for column values that occur for the intersection of a row and column in a table.

MySQL's Item class is defined in `.../sql/item.h`, and its subclasses are defined in `.../sql/item*.h` (that is, in `item.h`, `item_cmpfunc.h`, `item_func.h`, `item_geofunc.h`, `item_row.h`, `item_strfunc.h`, `item_subselect.h`, `item_sum.h`, `item_timefunc.h`). Page-width limitations prevent us from displaying the whole tree, but these are the main Item subclasses, and the subclasses of the subclasses:

```
Item_ident (Item_field, Item_ref)
Item_null
Item_num (Item_int, Item_real)
Item_param
Item_string (Item_static_string_func, Item_datetime, Item_empty_string)
Item_hex_string (Item_bin_string)
Item_result_field (all "item_func.h" "item_subselect.h" "item_sub.h" classes)
Item_copy_string
Item_cache (Item_cache_int, Item_cache_real, Item_cache_str, Item_cache_row)
Item_type_holder
Item_row
```

There's no formal classification of subclasses, but the main distinctions are by use (field, parameter, function) and by data type (num, string).

So, how does MySQL use items? You'll find that nearly every `.cc` program in the `/sql` directory makes some use of the Item class and its subclasses, so this list of programs is only partial and very general:

```
sql_parse.cc:    Makes new items in add_field_to_list()
item_sum.cc:     Uses item_func subclasses for COUNT, AVG, SUM
item_buff.cc:    Where buffers for item values can be stored
item_cmpfunc.cc: Comparison functions with item_func subclasses
item_create.cch: For creating items that the lex might use
item_subselect.cc: Subqueries are another type of function
mysqld.cc:      When main() ends, it uses clean_up() for items
opt_range.cc:   Uses field, compare-condition, and value subclasses
procedure.cc:   Notice Procedure * has a pointer to an item list
protocol.cc:    Uses send_fields() to pass item values back to users
sys_var.cc:     System variables have Item associations too
sql_base.cc:    Thread-specific Item searchers like find_field_in_table()
sql_class.cc:   Look at cleanup_after_query()
sql_delete.cc:  This (like sql_insert.cc etc.) has field references
sql_error.cc:   Has one of many examples of SHOW's use of items
sql_lex.cc:     Notice "add...to_list" functions
sql_select.cc:  The largest program that uses items, apparently
udf_example.cc: The comments in this program are extensive
```

Whenever there's a need for an SQL operation that assigns, compares, aggregates, accepts, sends, or validates a site, you'll find a MySQL use of Item and its subclasses.

How MySQL Does Sorting ([filesort](#))

[NOTE] This description is also present in `refman`

MySQL has two [filesort](#) algorithms for sorting and retrieving results. The original method uses only the [ORDER BY](#) columns. The modified method uses not just the [ORDER BY](#) columns, but all the columns used in the query.

The optimizer selects which [filesort](#) algorithm to use. Prior to MySQL 4.1, it uses the original algorithm. As of MySQL 4.1, it normally uses the modified algorithm except when [BLOB](#) or [TEXT](#) columns are involved, in which case it uses the original algorithm.

The original [filesort](#) algorithm works as follows:

1. Read all rows according to key or by table scanning. Rows that do not match the [WHERE](#) clause are skipped.
2. For each row, store a pair of values in a buffer (the sort key and the row pointer). The size of the buffer is the value of the [sort_buffer_size](#) system variable.
3. When the buffer gets full, run a qsort (quicksort) on it and store the result in a temporary file. Save a pointer to the sorted block. (If all pairs fit into the sort buffer, no temporary file is created.)

4. Repeat the preceding steps until all rows have been read.
5. Do a multi-merge of up to `MERGEBUFF` (7) regions to one block in another temporary file. Repeat until all blocks from the first file are in the second file.
6. Repeat the following until there are fewer than `MERGEBUFF2` (15) blocks left.
7. On the last multi-merge, only the pointer to the row (the last part of the sort key) is written to a result file.
8. Read the rows in sorted order by using the row pointers in the result file. To optimize this, we read in a big block of row pointers, sort them, and use them to read the rows in sorted order into a row buffer. The size of the buffer is the value of the `read_rnd_buffer_size` system variable. The code for this step is in the `sql/records.cc` source file.

One problem with this approach is that it reads rows twice: One time when evaluating the `WHERE` clause, and again after sorting the pair values. And even if the rows were accessed successively the first time (for example, if a table scan is done), the second time they are accessed randomly. (The sort keys are ordered, but the row positions are not.)

The modified `filesort` algorithm incorporates an optimization such that it records not only the sort key value and row position, but also the columns required for the query. This avoids reading the rows twice. The modified `filesort` algorithm works like this:

1. Read the rows that match the `WHERE` clause.
2. For each row, record a tuple of values consisting of the sort key value and row position, and also the columns required for the query.
3. Sort the tuples by sort key value
4. Retrieve the rows in sorted order, but read the required columns directly from the sorted tuples rather than by accessing the table a second time.

Using the modified `filesort` algorithm, the tuples are longer than the pairs used in the original method, and fewer of them fit in the sort buffer (the size of which is given by `sort_buffer_size`). As a result, it is possible for the extra I/O to make the modified approach slower, not faster. To avoid a slowdown, the optimization is used only if the total size of the extra columns in the sort tuple does not exceed the value of the `max_length_for_sort_data` system variable. (A symptom of setting the value of this variable too high is that you should see high disk activity and low CPU activity.)

Bulk Insert

The logic behind bulk insert optimization is simple.

Instead of writing each key value to B-tree (that is, to the key cache, although the bulk insert code doesn't know about the key cache), we store keys in a balanced binary (red-black) tree, in memory. When this tree reaches its memory limit, we write all keys to disk (to key cache, that is). But since the key stream coming from the binary tree is already sorted, inserting goes much faster, all the necessary pages are already in cache, disk access is minimized, and so forth.

How MySQL Does Caching

MySQL has the following caches. (Note that some of the filenames contain an incorrect spelling of the word "cache.")

- **Key Cache**

A shared cache for all B-tree index blocks in the different NISAM files. Uses hashing and reverse linked lists for quick caching of the most recently used blocks and quick flushing of changed entries for a specific table. (`mysys/mf_keycash.c`)

- **Record Cache**

This is used for quick scanning of all records in a table. ([mysys/mf_iocash.c](#) and [isam/_cash.c](#))

- **Table Cache**

This holds the most recently used tables. ([sql/sql_base.cc](#))

- **Hostname Cache**

For quick lookup (with reverse name resolving). This is a must when you have a slow DNS. ([sql/hostname.cc](#))

- **Privilege Cache**

To allow quick change between databases, the last used privileges are cached for each user/database combination. ([sql/sql_acl.cc](#))

- **Heap Table Cache**

Many uses of [GROUP BY](#) or [DISTINCT](#) cache all found rows in a [HEAP](#) table. (This is a very quick in-memory table with hash index.)

- **Join Buffer Cache**

For every “full join” in a [SELECT](#) statement the rows found are cached in a join cache. (A “full join” here means there were no keys that could be used to find rows for the next table in the list.) In the worst case, one [SELECT](#) query can use many join caches.

How MySQL Uses the Join Buffer Cache

Basic information about the join buffer cache:

- The size of each join buffer is determined by the value of the [join_buffer_size](#) system variable.
- This buffer is used only when the join is of type [ALL](#) or [index](#) (in other words, when no possible keys can be used).
- A join buffer is never allocated for the first non-const table, even if it would be of type [ALL](#) or [index](#).
- The buffer is allocated when we need to do a full join between two tables, and freed after the query is done.
- Accepted row combinations of tables before the [ALL/index](#) are stored in the cache and are used to compare against each read row in the [ALL](#) table.
- We only store the used columns in the join buffer, not the whole rows.

Assume you have the following join:

Table name	Type
t1	range
t2	ref
t3	ALL

The join is then done as follows:

- While rows in t1 matching range
- Read through all rows in t2 according to reference key
- Store used fields from t1, t2 in cache
- If cache is full
- Read through all rows in t3
- Compare t3 row against all t1, t2 combinations in cache
- If row satisfies join condition, send it to client
- Empty cache

```
- Read through all rows in t3
- Compare t3 row against all stored t1, t2 combinations in cache
  - If row satisfies join condition, send it to client
```

The preceding description means that the number of times table `t3` is scanned is determined as follows:

```
S = size-of-stored-row(t1,t2)
C = accepted-row-combinations(t1,t2)
scans = (S * C)/join_buffer_size + 1
```

Some conclusions:

- The larger the value of `join_buffer_size`, the fewer the scans of `t3`. If `join_buffer_size` is already large enough to hold all previous row combinations, there is no speed to be gained by making it larger.
- If there are several tables of join type `ALL` or `index`, then we allocate one buffer of size `join_buffer_size` for each of them and use the same algorithm described above to handle it. (In other words, we store the same row combination several times into different buffers.)

How MySQL Handles `FLUSH TABLES`

- `FLUSH TABLES` is handled in `sql/sql_base.cc::close_cached_tables()`.
- The idea of `FLUSH TABLES` is to force all tables to be closed. This is mainly to ensure that if someone adds a new table outside of MySQL (for example, by copying files into a database directory with `cp`), all threads will start using the new table. This will also ensure that all table changes are flushed to disk (but of course not as optimally as simply calling a `sync` for all tables)!
- When you do a `FLUSH TABLES`, the variable `refresh_version` is incremented. Every time a thread releases a table, it checks if the refresh version of the table (updated at open) is the same as the current `refresh_version`. If not, it will close it and broadcast a signal on `COND_refresh` (to await any thread that is waiting for all instances of a table to be closed).
- The current `refresh_version` is also compared to the open `refresh_version` after a thread gets a lock on a table. If the refresh version is different, the thread will free all locks, reopen the table and try to get the locks again. This is just to quickly get all tables to use the newest version. This is handled by `sql/lock.cc::mysql_lock_tables()` and `sql/sql_base.cc::wait_for_tables()`.
- When all tables have been closed, `FLUSH TABLES` returns an okay to the client.
- If the thread that is doing `FLUSH TABLES` has a lock on some tables, it will first close the locked tables, then wait until all other threads have also closed them, and then reopen them and get the locks. After this it will give other threads a chance to open the same tables.

Full-Text Search

MySQL uses Ranking with Vector Spaces for ordinary full-text queries.

Rank, also known as relevance rank, also known as relevance measure, is a number that tells us how good a match is.

Vector Space, which MySQL sometimes calls "natural language", is a well-known system based on a metaphor of lines that stretch in different dimensions (one dimension per term) for varying distances (one distance unit per occurrence of term). The value of thinking of it this way is: once you realize that term occurrences are lines in a multi-dimensional space, you can apply basic trigonometry to calculate "distances", and those distances are equatable with similarity measurements. A comprehensible discussion of vector space technology is here: <http://www.miislita.com/term-vector/term-vector-1.html>. And a text which partly inspired our original developer is here: <ftp://ftp.cs.cornell.edu/pub/smart/smart.11.0.tar.Z> ("SMART").

But let's try to describe the classic formula:

```
w = tf * idf
```

This means "weight equals term frequency times inverse of document frequency", or "increase weight for number of times term appears in one document, decrease weight for number of documents the term appears in". (For historical reasons we're using the word "weight" instead of "distance", and we're using the information-retrieval word "document" throughout; when you see it, think of "the indexed part of the row".)

For example: if "rain" appears three times in row #5, weight goes up; but if "rain" also appears in 1000 other documents, weight goes down.

MySQL uses a variant of the classic formula, and adds on some calculations for "the normalization factor". In the end, MySQL's formula looks something like:

```
w = (log(dtf)+1)/sumdtf * U/(1+0.0115*U) * log((N-nf)/nf)
```

Where:

dtf	is the number of times the term appears in the document
sumdtf	is the sum of (log(dtf)+1)'s for all terms in the same document
U	is the number of Unique terms in the document
N	is the total number of documents
nf	is the number of documents that contain the term

The formula has three parts: base part, normalization factor, global multiplier.

The base part is the left of the formula, "(log(dtf)+1)/sumdtf".

The normalization factor is the middle part of the formula. The idea of normalization is: if a document is shorter than average length then weight goes up, if it's average length then weight stays the same, if it's longer than average length then weight goes down. We're using a pivoted unique normalization factor. For the theory and justification, see the paper "Pivoted Document Length Normalization" by Amit Singhal and Chris Buckley and Mandar Mitra ACM SIGIR'96, 21-29, 1996: <http://ir.iit.edu/~dagr/cs529/files/handouts/singhal96pivoted.pdf>. The word "unique" here means that our measure of document length is based on the unique terms in the document. We chose 0.0115 as the pivot value, it's PIVOT_VAL in the MySQL source code header file myisam/ftdefs.h.

If we multiply the base part times the normalization factor, we have the term weight. The term weight is what MySQL stores in the index.

The global multiplier is the final part of the formula. In the classic Vector Space formula, the final part would be the inverse document frequency, or simply

```
log(N/nf)
```

We have replaced it with

```
log((N-nf)/nf)
```

This variant is more often used in "probabilistic" formulas. Such formulas try to make a better guess of the probability that a term will be relevant. To go back to the old system, look in myisam/ftdefs.h for "#define GWS_IN_USE GWS_PROB" (i.e. global weights by probability) and change it to "#define GWS_IN_USE GWS_IDF" (i.e. global weights by inverse document frequency).

Then, when retrieving, the rank is the product of the weight and the frequency of the word in the query:

```
R = w * qf;
```

Where:

w	is the weight (as always)
qf	is the number of times the term appears in the query

In vector-space speak, the similarity is the product of the vectors.

And R is the floating-point number that you see if you say: `SELECT MATCH(...) AGAINST (...) FROM t`.

To sum it up, w, which stands for weight, goes up if the term occurs more often in a row, goes down if the term occurs in many rows, goes up / down depending whether the number of unique words in a row is fewer / more than average. Then R, which stands for either Rank or Relevance, is w times the frequency of the term in the AGAINST expression.

The Simplest Possible Example

First, make a fulltext index. Follow the instructions in the "MySQL Full-Text Functions" section of the MySQL Reference Manual. Succinctly, the statements are:

```
CREATE TABLE articles (
  id INT UNSIGNED AUTO_INCREMENT NOT NULL PRIMARY KEY,
  title VARCHAR(200),
  body TEXT,
  FULLTEXT (title,body) );
INSERT INTO articles (title,body) VALUES
('MySQL Tutorial','DBMS stands for DataBase ...'),
('How To Use MySQL Well','After you went through a ...'),
('Optimizing MySQL','In this tutorial we will show ...'),
('1001 MySQL Tricks','1. Never run mysqld as root. 2. ...'),
('MySQL vs. YourSQL','In the following database comparison ...'),
('MySQL Security','When configured properly, MySQL ...');
```

Now, let's look at the index.

There's a utility for looking at the fulltext index keys and their weights. The source code is `myisam/myisam_ftdump.c`, and the executable comes with the binary distribution. So, if `exedir` is where the executable is, and `datadir` is the directory name that you get with `"SHOW VARIABLES LIKE 'datadir %'"`, and `dbname` is the name of the database that contains the articles table, then this works:

```
>/exedir/myisam_ftdump /datadir/dbname/articles 1 -d
b8      0.9456265 1001
f8      0.9560229 comparison
140     0.8148246 configured
0       0.9456265 database
f8      0.9560229 database
0       0.9456265 dbms
0       0.9456265 mysql
38      0.9886308 mysql
78      0.9560229 mysql
b8      0.9456265 mysql
f8      0.9560229 mysql
140     1.3796179 mysql
b8      0.9456265 mysqld
78      0.9560229 optimizing
140     0.8148246 properly
b8      0.9456265 root
140     0.8148246 security
78      0.9560229 show
0       0.9456265 stands
b8      0.9456265 tricks
0       0.9456265 tutorial
78      0.9560229 tutorial
f8      0.9560229 yoursq
```

Let's see how one of these numbers relates to the formula.

The term 'tutorial' appears in document 0. The full document is "MySQL Tutorial / DBMS stands for DataBase ...". The word "tutorial" appears once in the document, so `dtf` = The word "for" is a stopword, so there are only 5 unique terms in the document ("mysql", "tutorial", "dbms", "stands", "database"), so `U` = Each of these terms appears once in the document, so `sumdtf` is the sum of `log(1)+1`, five times. So, taking the first two parts of the formula (the term weight), we have:

```
(log(dtf)+1)/sumdtf * U/(1+0.0115*U)
```

which is

```
(log(1)+1)/( (log(1)+1)*5) * 5/(1+0.0115*5)
```

which is

```
0.9456265
```

which is what myisam_ftdump says. So the term weight looks good.

Now, what about the global multiplier? Well, myisam_ftdump could calculate it, but you'll see it with the mysql client. The total number of rows in the articles table is 6, so $N = 6$. And "tutorial" occurs in two rows, in row 0 and in row 78, so $nf = 2$. So, taking the final (global multiplier) part of the formula, we have:

```
log((N-nf)/nf)
```

which is

```
log((6-2)/2)
```

which is

```
0.6931472
```

So what would we get for row 0 with a search for 'tutorial'? Well, first we want w , so: Multiply the term weight of tutorial (which is 0.9456265) times the global multiplier (which is 0.6931472). Then we want R , so: Multiply w times the number of times that the word 'tutorial' appears in the search (which is 1). In other words, $R = 0.9456265 * 0.6931472 * 1$. Here's the proof:

```
mysql> select round(0.9456265 * 0.6931472 * 1, 7) as R;
+-----+
| R      |
+-----+
| 0.6554583 |
+-----+
1 row in set (0.00 sec)

mysql> select round(match(title,body) against ('tutorial'), 7) as R
-> from articles limit 1;
+-----+
| R      |
+-----+
| 0.6554583 |
+-----+
1 row in set (0.00 sec)
```

You'll need memory

The MySQL experience is that many users appreciate the full-text precision or recall, that is, the rows that MySQL returns are relevant and the rows that MySQL misses are rare, in the judgment of some real people. That means that the weighting formula is probably justifiable for most occasions. Since it's the product of lengthy academic research, that's understandable.

On the other hand, there are occasional complaints about speed. Here, the tricky part is that the formula depends on global factors -- specifically N (the number of documents) and nf (the number of documents that contain the term). Every time that insert/update/delete occurs for any row in the table, these global weight factors change for all rows in the table.

If MySQL was a search engine and there was no need to update in real time, this tricky part wouldn't matter. With occasional batch runs that redo the whole index, the global factors can be stored in the index. Search speed declines as the number of rows increases, but search engines work.

However, MySQL is a DBMS. So when updates happen, users expect the results to be visible immediately. It would take too long to replace the weights for all keys in the fulltext index, for every single update/insert/delete. So MySQL only stores the local factors in the index. The global factors are more dynamic. So MySQL stores an in-memory binary tree of the keys. Using this tree, MySQL can calculate the count of matching rows with reasonable speed. But speed declines logarithmically as the number of terms increases.

Weighting in boolean mode

The basic idea is as follows: In an expression of the form `A or B or (C and D and E)`, either `A` or `B` alone is enough to match the whole expression, whereas `C`, `D`, and `E` should **all** match. So it's reasonable to assign weight 1 to each of `A`, `B`, and `(C and D and E)`. Furthermore, `C`, `D`, and `E` each should get a weight of 1/3.

Things become more complicated when considering boolean operators, as used in MySQL full-text boolean searching. Obviously, `+A +B` should be treated as `A and B`, and `A B` - as `A or B`. The problem is that `+A B` can **not** be rewritten in and/or terms (that's the reason why this extended set of operators was chosen). Still, approximations can be used. `+A B C` can be approximated as `A or (A and (B or C))` or as `A or (A and B) or (A and C) or (A and B and C)`. Applying the above logic (and omitting mathematical transformations and normalization) one gets that for `+A_1 +A_2 ... +A_N B_1 B_2 ... B_M` the weights should be: $A_i = N$, $B_j = 1$ if $N=0$, and, otherwise, in the first rewriting approach $B_j = B_j = (1 + (M-1) * 2^M) / (M * (2^{(M+1)} - 1))$.

The second expression gives a somewhat steeper increase in total weight as number of matched `B_j` values increases, because it assigns higher weights to individual `B_j` values. Also, the first expression is much simpler, so it is the first one that is implemented in MySQL.

FLOAT and DOUBLE Data Type Representation

The MySQL Reference Manual has a discussion of floating-point numbers in Section 11.2 Numeric Types, including details about the storage. Let us now take up the story from where the MySQL Reference Manual leaves off.

The following discussion concentrates on the case where no display width and decimals are given. This means that `FLOAT` is stored as whatever the C type `float` is and `REAL` or `DOUBLE [PRECISION]` is stored as whatever the C type `double` is. The field length is selected by the MySQL code.

This document was created when <http://bugs.mysql.com/4457> (Different results in SQL-Statements for the same record) was fixed at the end of August 2004. Until then there was some confusion in the double-to-string conversion at different places in the code.

The bugfix for <http://bugs.mysql.com/4937> (`INSERT + SELECT + UNION ALL + DATE to VARCHAR(8)` conversion problem) produced a conversion function which was a promising approach to the conversion problems. Unfortunately it was only used for direct field conversions and not for function results etc. It did not take small numbers (absolute value less than 1) and negative numbers into account. It did not take the limited precision of `float` and `double` data types into account. The bugfix was developed in two steps: The first attempt looked like this (in principle):

```
length= sprintf(buff, "%.g", field_length, nr);
if (length > field_length)
    length= sprintf(buff, "%.g", field_length-5, nr);
```

If the `libc` conversion produces too many characters, the precision is reduced by the space required for the scientific notation (1.234e+05). Thus the `printf()` conversion is forced to switch to the scientific notation, since the value would not fit otherwise. Or, if it was scientific already, the precision is reduced and also uses less space. I left out some important stuff around limit checking just to show the idea. This simple algorithm should work quite well in most cases, but has been discarded for the sake of performance. The double call to the slow `printf()` conversion `%g` didn't seem reasonable, though it would only be used for extreme values and small fields. During my explorations of the code I didn't find places where `float` or `double` were to be converted into small fields. Remember that I talk only of conversions where field length and precision are not given. In this case a sufficient field length is selected at several places, except for a bug where it was selected wrongly. If a field length is given, a different conversion is used anyway. But since the code is quite complex, I don't claim to grasp it in full, and therefore may be in error. So let us look further:

The second attempt to fix the bug looked like this:

```
bool use_scientific_notation=TRUE;
if (field_length < 32 && nr > 1)
```

```
{
    double e[]={1, 1e1, 1e2, 1e4, 1e8, 1e16 }, p=1;
    for (int i=sizeof(e), j=1<i-- ; j; i--, j>=1 )
    {
        if (field_length & j)
            p*=e[i];
    }
    use_scientific_notation=(p < nr);
}
length= sprintf(buff, "%.*g", use_scientific_notation ?
                field_length-5 : field_length, nr);
```

Here we evaluate if the string representation of a given number fits into `field_length` characters. If not, we reduce the precision to make it fit. Again, I left out important details. For example, the evaluation is done only once per field for the sake of performance. The downside here is the unconditional reduction of precision for `field_length > 31` (which doesn't really matter), for negative numbers and for small numbers (absolute value less than 1).

Both algorithms do not take the limited precision of `float` and `double` values into account. This could lead to conversions with ridiculous bogus precision output. For example a value of 0.7 converted with `%.30g` will give a lot of digits, which pretend to tell about deviations from the value 0.7 and are completely absurd: 0.699999988079071044921875. To understand more about the `%g` conversion, I quote from a comment introduced in the source at the beginning of bugfixing #4937 (this comment was removed because it mainly describes, how the `printf()` conversion works, but I think it's valuable enough to include it here):

```
/*
    Let's try to pretty print a floating point number. Here we use
    '%-*.g' conversion string:
    '-' stands for right-padding with spaces, if such padding will take
    place
    '*' is a placeholder for the first argument, field_length, and
    signifies minimal width of result string. If result is less than
    field length it will be space-padded. Note, however, that we'll not
    pass spaces to Field_string::store(const char *, ...), due to
    strcend in the next line.
    '.' is a placeholder for DBL_DIG and defines maximum number of
    significant digits in the result string. DBL_DIG is a hardware
    specific C define for maximum number of decimal digits of a floating
    point number, such that rounding to hardware floating point
    representation and back to decimal will not lead to loss of
    precision. That is: if DBL_DIG is 15, number 123456789111315 can be
    represented as double without precision loss. As one can judge from
    this description, choosing DBL_DIG here is questionable, especially
    because it introduces a system dependency.
    'g' means that conversion will use [-]ddd.ddd (conventional) style,
    and fall back to [-]d.ddde[+|i]ddd (scientific) style if there is not
    enough space for all digits.
    Maximum length of result string (not counting spaces) is (I guess)
    DBL_DIG + 8, where 8 is 1 for sign, 1 for decimal point, 1 for
    exponent sign, 1 for exponent, and 4 for exponent value.
    XXX: why do we use space-padding and trim spaces in the next line?
*/
sprintf(to,"%-*.g",(int) field_length,DBL_DIG,nr);
to=strcend(to,' ');
```

There is one small misapprehension in the comment. `%g` does not switch to scientific notation when there is 'not enough space for all digits'. As the commentator says, the field length gives the minimal output length. `printf()` happily outputs more characters if required to produce a result with 'precision' digits. In fact it switches to scientific when the value can no longer be represented by 'precision' digits in conventional notation. The man page says "Style e is used if the exponent from its conversion is less than -4 or greater than or equal to the precision." In explanation, a precision of 3 digits can print a value of 345 in conventional notation, but 3456 needs scientific notation, as it would require 4 digits (a precision of 4) in conventional notation. Thus, it is printed as 3.46e+03 (rounded).

Since we don't want spaces in the output, we should not give a field length, but always use `"%.*g"`. However, the precision matters, as seen above. It is worth its own paragraph.

Since MySQL uses the machine-dependent binary representation of `float` and `double` to store values in the database, we have to care about these. Today, most systems use the IEEE standard 754 for binary floating-point arithmetic. It describes a representation for single precision numbers as 1 bit for sign, 8 bits for biased exponent and 23 bits for fraction and for double precision numbers as 1-bit sign, 11-bit biased exponent and 52-bit fraction. However, we can not rely on the fact that every system uses this representation. Luckily, the ISO C standard requires the standard C library to have a header `float.h` that describes some details of the floating point representation on a machine. The comment above describes the value `DBL_DIG`. There is an equivalent value `FLT_DIG` for the C data type `float`.

So, whenever we print a floating-point value, we must not specify a precision above `DBL_DIG` or `FLT_DIG` respectively. Otherwise we produce a bogus precision, which is wrong. For the honor of the writer of the first attempt above, I must say that his complete algorithm took `DBL_DIG` into account, if however only for the second call to `sprintf()`. But `FLT_DIG` has never been accounted for. At the conversion section of the code, it was not even known whether the value came from a `float` or `double` field.

My attempt to solve the problems tries to take all this into account. I tried to concentrate all `float/double`-to-string conversions in one function, and to bring the knowledge about `float` versus `double` to this function wherever it is called. This solution managed to keep the test suite happy while solving the new problem of [\[http://bugs.mysql.com/4457 Bug#4457\]](http://bugs.mysql.com/4457). Luckily the first problem was not big, as the test cases have been very carefully selected, so that they succeed as long as the machine uses IEEE 754.

Nevertheless, the function is still not perfect. It is not possible to guess how many significant digits a number has. Given that, it is not simple to tell how long the resulting string would be. This applies to numbers with an absolute value smaller than 1. There are probably ways to figure this out, but I doubt that we would win in terms of performance over the simple solution of the first attempt, and besides we might cause new bugs. The compromise taken here is to accept that the resulting string may exceed the destination field length by five characters in the worst case.

```
if (nr < 0.0)
{
    abs_nr= -nr;
    extra_space= 1;
}
else
{
    abs_nr= nr;
    extra_space= 0;
}
precision= is_float ? FLT_DIG : DBL_DIG;
if (precision > field_length)
    precision= field_length;

if (! initialized)
{
    /* Better switch to scientific too early than too late. */
    double mult;
    mult= 1e0;
    for (length= 0; length < DBL_DIG; length++)
        mult/= 1e1;
    mult= 1e1 - mult;

    double val;
    val= 1.0;
    for (int idx= 0; idx < DBL_DIG+1; idx++)
    {
        DEBUG_PRINT("info", ("double_to_string_conv: big[%d] %.*g",
                               idx, DBL_DIG+3, val));
        big_number[idx]= val;
        val*= mult;
    }
    small_number[0]= 1e0;
    small_number[1]= 1e0;
    small_number[2]= 1e0;
    small_number[3]= 1e-1;
```

```

    small_number[4]= 1e-2;
    small_number[5]= 1e-3;
    small_number[6]= 1e-4;
    /* %g switches to scientific when exponent < -4. */
    for (int idx= 7; idx < DBL_DIG+1; idx++)
        small_number[idx]= 1e-4;
    initialized= TRUE;
}

use_scientific_notation= (abs_nr != 0.0) &&
    ((abs_nr >  big_number[precision]) ||
     (abs_nr < small_number[precision]));

if (use_scientific_notation)
{
    if ((nr >= 0.0) && ((nr >= 1e+100) || (nr <= 1e-100))) ||
        ((nr < 0.0) && ((nr <= -1e+100) || (nr >= -1e-100)))
        extra_space+= 6; /* .e+100 or .e-100 */
    else
        extra_space+= 5; /* .e+99 or .e-99 */
}

if (field_length < extra_space)
    precision= 0;
else if (precision > (field_length - extra_space))
    precision= field_length - extra_space;

length= sprintf(buff, "%.*g", precision, nr);

```

This solution takes performance into account by initializing the limiting numbers arrays only once into static space. It copes with negative numbers and tries to decide even over small numbers. The latter has only small implications, as the prefix 0.000 is exactly the same size as the postfix e-100. But knowing if scientific notation will be selected by `sprintf()` allows for saving one digit when the exponent is larger than -100.

The calculations for the big number array are less precise than in the second attempt, but faster. The precision is sufficient for the guess whether `sprintf()` uses scientific notation. There may be number to field length combinations which exploit the gap, but these won't emerge anyway as I found no situation where this function is called with small field lengths. Remember again that it is not called with user-supplied field lengths.

However in the current stable releases (including gamma) we have some places where the field length is too small by one character. Thus, the precision is sometimes one digit smaller than `DBL_DIG` would allow for. Consequently, we cannot use the simple algorithm in the stable releases. There is a chance of doing it in a development release, though.

Addendum:

There turned out to be a new solution to the "big number array" problem. We have a statically initialized array `log_10`, which holds the necessary values. But I did not check whether these values are safe. Even if computed by the compiler, they could carry values slightly above the decimal powers, which would be bad. In this case we needed to initialize by 9.99999999e+xxx, where the number of nines is equal to `DBL_DIG`. This must be protected by `#if DBL_DIG == YY`, so that a new `DBL_DIG` on a new platform is detected. And the array is of limited length. We must at least protect it by a `DEBUG_ASSERT(sizeof(log_10)/sizeof(log_10[0]) > DBL_DIG)`.

But all of this is probably completely unnecessary, since we are only speaking of cases where no user-supplied field length is given. So MySQL selects the field length on its own. So it is totally possible, indeed highly desirable, that MySQL selects a field length, which allows for a maximum of precision for all possible values. And these are `DBL_DIG+7` or `FLT_DIG+6` respectively as far as IEEE 754 is used. In this case we can have values of about +/-1e-307 to +/-1e+308 for `double` and +/-1e-37 to +/-1e+38 for `float`. That is, for example -1.<DBL_DIG-1 digits>e+100. For cases where a precision above IEEE 754 is possible, we may need +8 instead. We can detect this with `#if DBL_MAX_10_EXP >=`. So using a field length of `DBL_DIG+8` in all cases should be sufficient for a simple `sprintf(buff, "%. *g", DBL_DIG, nr)` or `sprintf(buff, "%. *g", FLT_DIG, nr)`, respectively. To be safe,

we should not use the machine dependent constants everywhere, but instead concentrate them into definitions like these:

```
#if (DBL_MAX_10_EXP > 9999) || (DBL_MIN_10_EXP < -9999)
# error "Need new definition for UNSPECIFIED_DOUBLE_FIELD_LENGTH"
#elif (DBL_MAX_10_EXP > 999) || (DBL_MIN_10_EXP < -999)
# define UNSPECIFIED_DOUBLE_FIELD_LENGTH (DBL_DIG+8)
#else
# define UNSPECIFIED_DOUBLE_FIELD_LENGTH (DBL_DIG+7)
#endif

#if (FLT_MAX_10_EXP > 999) || (FLT_MIN_10_EXP < -999)
#error "Need new definition for UNSPECIFIED_FLOAT_FIELD_LENGTH"
#elif (FLT_MAX_10_EXP > 99) || (FLT_MIN_10_EXP < -99)
# define UNSPECIFIED_FLOAT_FIELD_LENGTH (FLT_DIG+7)
#else
# define UNSPECIFIED_FLOAT_FIELD_LENGTH (FLT_DIG+6)
#endif
```

These definitions should be used wherever an item or field of type `float` or `double` without an explicit field length specification is encountered. We have to propagate these lengths through all derived items and fields and we have to select the maximum of all field lengths wherever in two or more of them are used in an expression or a function.

We need to treat the precision (`DBL_DIG/FLT_DIG`) similarly, but have to select the minimum in expressions or functions.

Date and Time Data Type Representation

The following table shows the storage requirements for date and time data types.

Type	Storage before MySQL 5.6.4	Storage as of MySQL 5.6.4
YEAR	1 byte, little endian	Unchanged
DATE	3 bytes, little endian	Unchanged
TIME	3 bytes, little endian	3 bytes + fractional-seconds storage, big endian
TIMESTAMP	4 bytes, little endian	4 bytes + fractional-seconds storage, big endian
DATETIME	8 bytes, little endian	5 bytes + fractional-seconds storage, big endian

Before MySQL 5.6.4, date and time data types have these encodings:

- YEAR: A one-byte integer
- DATE: A three-byte integer packed as $YYYY \times 16 \times 32 + MM \times 32 + DD$
- TIME: A three-byte integer packed as $DD \times 24 \times 3600 + HH \times 3600 + MM \times 60 + SS$
- TIMESTAMP: A four-byte integer representing seconds UTC since the epoch ('1970-01-01 00:00:00' UTC)
- DATETIME: Eight bytes: A four-byte integer for date packed as $YYYY \times 10000 + MM \times 100 + DD$ and a four-byte integer for time packed as $HH \times 10000 + MM \times 100 + SS$

As of MySQL 5.6.4 the TIME, TIMESTAMP, and DATETIME types can have a fractional seconds part. Storage for these types is big endian (for memcmp() compatibility purposes), with the nonfractional part followed by the fractional part. (Storage and encoding for the YEAR and DATE types remains unchanged.)

- TIME encoding for nonfractional part:

```

1 bit sign      (1= non-negative, 0= negative)
1 bit unused    (reserved for future extensions)
10 bits hour    (0-838)
6 bits minute   (0-59)
6 bits second   (0-59)
-----
24 bits = 3 bytes

```

- **TIMESTAMP** encoding for nonfractional part: Same as before 5.6.4, except big endian rather than little endian
- **DATETIME** encoding for nonfractional part:

```

1 bit sign      (1= non-negative, 0= negative)
17 bits year*13+month (year 0-9999, month 0-12)
5 bits day      (0-31)
5 bits hour     (0-23)
6 bits minute   (0-59)
6 bits second   (0-59)
-----
40 bits = 5 bytes

```

The sign bit is always 1. A value of 0 (negative) is reserved.

- Fractional-part encoding depends on the fractional seconds precision (FSP).

FSP	Storage
0	0 bytes
1,2	1 byte
3,4	2 bytes
4,5	3 bytes

Threads

Threads in `mysqld` can run at four different priorities, defined in `mysql_priv.h`:

```

#define INTERRUPT_PRIOR 10
#define CONNECT_PRIOR   9
#define WAIT_PRIOR      8
#define QUERY_PRIOR     6

```

Some threads try to set their priority; others don't. These calls are passed along to `pthread_setschedparam()` if the native threading library implements it.

The different threads are:

- The main thread. Runs at `CONNECT_PRIOR` priority. Calls `thr_setconcurrency()` if it is available at compile time; this call is generally assumed to exist only on Solaris, its value should reflect the number of physical CPUs.
- The "bootstrap" thread. See `handle_bootstrap()` in `sql_parse.cc`. The `mysql_install_db` script starts a server with an option telling it to start this thread and read commands in from a file. Used to initialize the grant tables. Runs once and then exits.
- The "maintenance" thread. See `sql_manager.cc`. Like the old "sync" daemon in unix, this thread occasionally flushes MyISAM tables to disk. InnoDB has a separate maintenance thread, but BDB also uses this one to occasionally call `berkeley_cleanup_log_files()`. Begins at startup and persists until shutdown.
- The "handle TCP/IP sockets" thread. See `handle_connections_sockets()` in `mysqld.cc`. Loop with a `select()` function call, to handle incoming connections.
- The "handle named pipes" thread. Only on Windows.

- The "handle shared memory connections" thread. Only on Windows.
- Signal handler ("interrupt") thread. See `signal_hand()` in `mysqld.cc`. Runs at `INTERRUPT_PRIOR` priority. Sets up to receive signals, and then handles them as they come in. Begins at server startup and persists until shutdown.
- The "shutdown" thread. See `kill_server()` in `mysqld.cc`. Created by the signal handling thread. Closes all connections with `close_connections()`, then ends.
- Active and cached per-connection threads. See `handle_one_connection()` in `sql_parse.cc`. These can run at `QUERY_PRIOR` priority or `WAIT_PRIOR` priority depending on what they are doing.
- The "delayed" thread. See `handle_delayed_insert()` in `sql_insert.cc`. Used for MyISAM's delayed inserts.
- The two slave threads, in `slave.cc`. One thread connects to the master and handles network IO. The other reads queries from the relay log and executes them.

InnoDB, all thread management is handled through `os/os0thread.c`. InnoDB's threads are:

- The I/O handler threads, See `io_handler_thread()`.
- Two "watchmen" threads: `srv_lock_timeout_and_monitor_thread()`, and `srv_error_monitor_thread()`.
- The master thread "which does purge and other utility operations", See `srv_master_thread()`.

InnoDB's internal `os_thread_set_priority()` function implements three priorities (Background, normal, and high) but only on windows. The function is a no-op on unix.

Character Sets and Collations

Character sets are used by MySQL when storing information, both to ensure that the information is stored (and returned) in the correct format, but also for the purposes of collation and sorting. Each character set supports one or more collations, and so these are collectively known as [Collation Sets](#), rather than character sets.

Character sets are recorded against individual tables and returned as part of the field data. For example, the `MYSQL_FIELD` data type definition includes the field `charsetnr`:

```
typedef struct st_mysql_field {
    char *name;                /* Name of column */
    char *org_name;             /* Original column name, if an alias */
    char *table;                /* Table of column if column was a field */
    char *org_table;            /* Org table name, if table was an alias */
    char *db;                   /* Database for table */
    char *catalog;              /* Catalog for table */
    char *def;                   /* Default value (set by mysql_list_fields) */
    unsigned long length;        /* Width of column (create length) */
    unsigned long max_length;    /* Max width for selected set */
    unsigned int name_length;
    unsigned int org_name_length;
    unsigned int table_length;
    unsigned int org_table_length;
    unsigned int db_length;
    unsigned int catalog_length;
    unsigned int def_length;
    unsigned int flags;         /* Div flags */
    unsigned int decimals;      /* Number of decimals in field */
    unsigned int charsetnr;     /* Character set */
    enum enum_field_types type; /* Type of field. See mysql_com.h for types */
} MYSQL_FIELD;
```

Character set and collation information are specific to a server version and installation, and are generated automatically from the [sql/share/charsets/Index.xml](#) file in the source distribution.

You can obtain a list of the available character sets configured within a server by running [SHOW COLLATION](#), or by running a query on the [INFORMATION_SCHEMA.COLLATION](#) table. A sample of the information from that table has been provided here for reference.

Collation Id	Charset	Collation	Default	Sortlen
64	armscii8	armscii8_bin	??	1
32	armscii8	armscii8_general_ci	Yes	1
65	ascii	ascii_bin	??	1
11	ascii	ascii_general_ci	Yes	1
84	big5	big5_bin	??	1
1	big5	big5_chinese_ci	Yes	1
63	binary	binary	Yes	1
66	cp1250	cp1250_bin	??	1
44	cp1250	cp1250_croatian_ci	??	1
34	cp1250	cp1250_czech_cs	??	2
26	cp1250	cp1250_general_ci	Yes	1
50	cp1251	cp1251_bin	??	1
14	cp1251	cp1251_bulgarian_ci	??	1
52	cp1251	cp1251_general_cs	??	1
23	cp1251	cp1251_ukrainian_ci	??	1
51	cp1251	cp1251_general_ci	Yes	1
67	cp1256	cp1256_bin	??	1
57	cp1256	cp1256_general_ci	Yes	1
58	cp1257	cp1257_bin	??	1
29	cp1257	cp1257_lithuanian_ci	??	1
59	cp1257	cp1257_general_ci	Yes	1
80	cp850	cp850_bin	??	1
4	cp850	cp850_general_ci	Yes	1
81	cp852	cp852_bin	??	1
40	cp852	cp852_general_ci	Yes	1
68	cp866	cp866_bin	??	1
36	cp866	cp866_general_ci	Yes	1
96	cp932	cp932_bin	??	1
95	cp932	cp932_japanese_ci	Yes	1
69	dec8	dec8_bin	??	1
3	dec8	dec8_swedish_ci	Yes	1
98	eucjpms	eucjpms_bin	??	1
97	eucjpms	eucjpms_japanese_ci	Yes	1
85	euckr	euckr_bin	??	1
19	euckr	euckr_korean_ci	Yes	1
86	gb2312	gb2312_bin	??	1
24	gb2312	gb2312_chinese_ci	Yes	1
87	gbk	gbk_bin	??	1

Character Sets and Collations

28	gbk	gbk_chinese_ci	Yes	1
93	geostd8	geostd8_bin	??	1
92	geostd8	geostd8_general_ci	Yes	1
70	greek	greek_bin	??	1
25	greek	greek_general_ci	Yes	1
71	hebrew	hebrew_bin	??	1
16	hebrew	hebrew_general_ci	Yes	1
72	hp8	hp8_bin	??	1
6	hp8	hp8_english_ci	Yes	1
73	keybcs2	keybcs2_bin	??	1
37	keybcs2	keybcs2_general_ci	Yes	1
74	koi8r	koi8r_bin	??	1
7	koi8r	koi8r_general_ci	Yes	1
75	koi8u	koi8u_bin	??	1
22	koi8u	koi8u_general_ci	Yes	1
47	latin1	latin1_bin	??	1
15	latin1	latin1_danish_ci	??	1
48	latin1	latin1_general_ci	??	1
49	latin1	latin1_general_cs	??	1
5	latin1	latin1_german1_ci	??	1
31	latin1	latin1_german2_ci	??	2
94	latin1	latin1_spanish_ci	??	1
8	latin1	latin1_swedish_ci	Yes	1
77	latin2	latin2_bin	??	1
27	latin2	latin2_croatian_ci	??	1
2	latin2	latin2_czech_cs	??	4
21	latin2	latin2_hungarian_ci	??	1
9	latin2	latin2_general_ci	Yes	1
78	latin5	latin5_bin	??	1
30	latin5	latin5_turkish_ci	Yes	1
79	latin7	latin7_bin	??	1
20	latin7	latin7_estonian_cs	??	1
42	latin7	latin7_general_cs	??	1
41	latin7	latin7_general_ci	Yes	1
43	macce	macce_bin	??	1
38	macce	macce_general_ci	Yes	1
53	macroman	macroman_bin	??	1
39	macroman	macroman_general_ci	Yes	1
88	sjis	sjis_bin	??	1
13	sjis	sjis_japanese_ci	Yes	1
82	swe7	swe7_bin	??	1
10	swe7	swe7_swedish_ci	Yes	1

Character Sets and Collations

89	tis620	tis620_bin	??	1
18	tis620	tis620_thai_ci	Yes	4
90	ucs2	ucs2_bin	??	1
138	ucs2	ucs2_czech_ci	??	8
139	ucs2	ucs2_danish_ci	??	8
145	ucs2	ucs2_esperanto_ci	??	8
134	ucs2	ucs2_estonian_ci	??	8
146	ucs2	ucs2_hungarian_ci	??	8
129	ucs2	ucs2_icelandic_ci	??	8
130	ucs2	ucs2_latvian_ci	??	8
140	ucs2	ucs2_lithuanian_ci	??	8
144	ucs2	ucs2_persian_ci	??	8
133	ucs2	ucs2_polish_ci	??	8
131	ucs2	ucs2_romanian_ci	??	8
143	ucs2	ucs2_roman_ci	??	8
141	ucs2	ucs2_slovak_ci	??	8
132	ucs2	ucs2_slovenian_ci	??	8
142	ucs2	ucs2_spanish2_ci	??	8
135	ucs2	ucs2_spanish_ci	??	8
136	ucs2	ucs2_swedish_ci	??	8
137	ucs2	ucs2_turkish_ci	??	8
128	ucs2	ucs2_unicode_ci	??	8
35	ucs2	ucs2_general_ci	Yes	1
91	ujis	ujis_bin	??	1
12	ujis	ujis_japanese_ci	Yes	1
83	utf8	utf8_bin	??	1
202	utf8	utf8_czech_ci	??	8
203	utf8	utf8_danish_ci	??	8
209	utf8	utf8_esperanto_ci	??	8
198	utf8	utf8_estonian_ci	??	8
210	utf8	utf8_hungarian_ci	??	8
193	utf8	utf8_icelandic_ci	??	8
194	utf8	utf8_latvian_ci	??	8
204	utf8	utf8_lithuanian_ci	??	8
208	utf8	utf8_persian_ci	??	8
197	utf8	utf8_polish_ci	??	8
195	utf8	utf8_romanian_ci	??	8
207	utf8	utf8_roman_ci	??	8
205	utf8	utf8_slovak_ci	??	8
196	utf8	utf8_slovenian_ci	??	8
206	utf8	utf8_spanish2_ci	??	8
199	utf8	utf8_spanish_ci	??	8

200	utf8	utf8_swedish_ci	??	8
201	utf8	utf8_turkish_ci	??	8
192	utf8	utf8_unicode_ci	??	8
33	utf8	utf8_general_ci	Yes	1

Note that it is the collation ID, not the character set ID, that is used to identify the unique combination of character set and collation. Thus, when requesting character set information using one of the character set functions in `mysys/charset.c`, such as `get_charset()`, different IDs may return the same base character set, but a different collation set.

The following functions provide an internal interface to the collation and character set information, enabling you to access the information by name or ID:

```
static uint get_collation_number_internal(const char *name)
uint get_collation_number(const char *name)
uint get_charset_number(const char *charset_name, uint cs_flags)
const char *get_charset_name(uint charset_number)
static CHARSET_INFO *get_internal_charset(uint cs_number, myf flags)
CHARSET_INFO *get_charset(uint cs_number, myf flags)
CHARSET_INFO *get_charset_by_name(const char *cs_name, myf flags)
CHARSET_INFO *get_charset_by_csname(const char *cs_name,
                                     uint cs_flags,
                                     myf flags)
```

The table below details the functions, the key argument that is supplied, and the return value.

Function	Supplied Argument	Return Value
<code>get_collation_number_internal()</code>	Collation name	Collation ID
<code>get_collation_number()</code>	Collation name	Collation ID
<code>get_charset_number()</code>	Character set name	Collation ID
<code>get_charset_name()</code>	Collation ID	Character set name
<code>get_internal_charset()</code>	Collation ID	Character datatype
<code>get_charset()</code>	Collation ID	Character datatype

An example of using the collation/character set functions is available in the `extras/charset2html.c`, which outputs an HTML version of the internal collation set table.

Error Flags and Functions

The following flags can be examined or set to alter the behavior during error handling:

- `thd->net.report_error`

`thd->net.report_error` is set in `my_message_sql()` if the error message was registered. (`my_message_sql()` is called by `my_error()`, `my_printf_error()`, `my_message()`).

- `thd->query_error`

Like `net.report_error`, but is always set to 1 in `my_message_sql()` if error was not caught by an error handler. Used by replication to see if a query generated any kind of errors.

- `thd->no_warnings_for_error`

Normally an error also generates a warning. The warning can be disabled by setting `thd->no_warnings_for_error`. (This allows one to catch all error messages generated by a statement)

- `thd->lex->current_select->no_error`

This is set to in case likes `INSERT IGNORE ... SELECT`. In this case we ignore all not fatal errors generated by the select.

- `thd->is_fatal_error`

Set this if we should abort the current statement (and any multi-line statements) because something went fatally wrong. (for example, a stored procedure continue handler should not be able to catch this). This is reset by `mysql_reset_thd_for_next_command()`.

- `thd->abort_on_warning`

Strict mode flag, which means that we should abort the statement if we get a warning. In the `field::store` function this changes the warning level from `WARN` to `ERROR`. In other cases, this flag is mostly tested with `thd->really_abort_on_warning()` to ensure we don't abort in the middle of an update with not transactional tables.

- `thd->count_cuted_fields`

If set, we generate warning for field conversions (normal case for `INSERT/UPDATE/DELETE`). This is mainly set to 0 when doing internal copying of data between fields and we don't want to generate any conversion errors at any level.

- `thd->killed`

Set in case of error in connection protocol or in case of 'kill'. In this case we should abort the query and kill the connection.

Error functions

- `thd->really_abort_on_warning()`

This function returns 1 if a warning should be converted to an error, like in strict mode when all tables are transactional. The conversion is handled in `sql_error.cc::push_warning()`.

- `thd->fatal_error()`

Should be called if we want to abort the current statement and any multi-line statement.

- `thd->clear_error()`

Resets `thd->net.report_error` and `thd->query_error`.

Functions in the `mysys` Library

Functions in `mysys`: (For flags see `my_sys.h`)

- `int my_copy_A((const char *from, const char *to, myf MyFlags));`

Copy file from `from` to `to`.

- `int my_rename_A((const char *from, const char *to, myf MyFlags));`

Rename file from `from` to `to`.

- `int my_delete_A((const char *name, myf MyFlags));`

Delete file `name`.

- `int my_redel_A((const char *from, const char *to, int MyFlags));`

Delete `from` before rename of `to` to `from`. Copies state from old file to new file. If `MY_COPY_TIME` is set, sets old time.

- `int my_getwd _A((string buf, uint size, myf MyFlags));, int my_setwd _A((const char *dir, myf MyFlags));`

Get and set working directory.

- `string my_tempnam _A((const char *dir, const char *pfx, myf MyFlags));`

Make a unique temporary file name by using `dir` and adding something after `pfx` to make the name unique. The file name is made by adding a unique six character string and `TMP_EXT` after `pfx`. Returns pointer to `malloc()`'ed area for filename. Should be freed by `free()`.

- `File my_open _A((const char *FileName, int Flags, myf MyFlags));, File my_create _A((const char *FileName, int CreateFlags, int AccsesFlags, myf MyFlags));, int my_close _A((File Filedes, myf MyFlags));, uint my_read _A((File Filedes, byte *Buffer, uint Count, myf MyFlags));, uint my_write _A((File Filedes, const byte *Buffer, uint Count, myf MyFlags));, ulong my_seek _A((File fd, ulong pos, int whence, myf MyFlags));, ulong my_tell _A((File fd, myf MyFlags));`

Use instead of `open`, `open-with-create-flag`, `close`, `read`, and `write` to get automatic error messages (flag `MYF_WME`) and only have to test for `!= MY_NABP`.

- `FILE *my_fopen _A((const char *FileName, int Flags, myf MyFlags));, FILE *my_fdopen _A((File Filedes, int Flags, myf MyFlags));, int my_fclose _A((FILE *fd, myf MyFlags));, uint my_fread _A((FILE *stream, byte *Buffer, uint Count, myf MyFlags));, uint my_fwrite _A((FILE *stream, const byte *Buffer, uint Count, myf MyFlags));, ulong my_fseek _A((FILE *stream, ulong pos, int whence, myf MyFlags));, ulong my_ftell _A((FILE *stream, myf MyFlags));`

Same read-interface for streams as for files.

- `gp_ptr _mymalloc _A((uint uSize, const char *sFile, uint uLine, myf MyFlag));, gp_ptr _myrealloc _A((string pPtr, uint uSize, const char *sFile, uint uLine, myf MyFlag));, void _myfree _A((gp_ptr pPtr, const char *sFile, uint uLine));, int _sanity _A((const char *sFile, unsigned int uLine));, gp_ptr _myget_copy_of_memory _A((const byte *from, uint length, const char *sFile, uint uLine, myf MyFlag));`

`malloc(size, myflag)` is mapped to these functions if not compiled with `-DSAFEMALLOC`.

- `void TERMINATE _A((void));`

Writes `malloc()` info on `stdout` if compiled with `-DSAFEMALLOC`.

- `int my_chsize _A((File fd, ulong newlength, myf MyFlags));`

Change size of file `fd` to `newlength`.

- `void my_error _D((int nr, myf MyFlags, ...));`

Writes message using error number (see `mysys/errors.h`) on `stdout`, or using curses, if `MYSYS_PROGRAM_USES_CURSES()` has been called.

- `void my_message _A((const char *str, myf MyFlags));`

Writes `str` on `stdout`, or using curses, if `MYSYS_PROGRAM_USES_CURSES()` has been called.

- `void my_init _A((void));`

Start each program (in `main()`) with this.

- `void my_end _A((int infoflag));`

Gives info about program. If `infoflag & MY_CHECK_ERROR`, prints if some files are left open. If `infoflag & MY_GIVE_INFO`, prints timing info and `malloc()` info about program.

- `int my_copystat _A((const char *from, const char *to, int MyFlags));`

Copy state from old file to new file. If `MY_COPY_TIME` is set, sets old time.

- `string my_filename _A((File fd));`

Returns filename of open file.

- `int dirname _A((string to, const char *name));`

Copy name of directory from filename.

- `int test_if_hard_path _A((const char *dir_name));`

Test if `dir_name` is a hard path (starts from root).

- `void convert_dirname _A((string name));`

Convert `dirname` according to system. On Windows, changes all characters to capitals and changes `'/'` to `'\'`.

- `string fn_ext _A((const char *name));`

Returns pointer to extension in filename.

- `string fn_format _A((string to, const char *name, const char *dsk, const char *form, int flag));`

Format a filename with replacement of library and extension and convert between different systems. The `to` and `name` parameters may be identical. Function doesn't change name if `name != to`. `flag` may be:

1	Force replace filenames library with 'dsk'
2	Force replace extension with 'form' */
4	Force unpack filename (replace ~ with home directory)
8	Pack filename as short as possible for output to user

All open requests should always use at least `open(fn_format(temp_buffer, name, " ", " ", 4), ...)` to unpack home and convert filename to system-form.

- `string fn_same _A((string toname, const char *name, int flag));`

Copies directory and extension from `name` to `toname` if needed. Copying can be forced by same flags used in `fn_format()`.

- `int wild_compare _A((const char *str, const char *wildstr));`

Compare if `str` matches `wildstr`. `wildstr` can contain `'*'` and `'?'` as wildcard characters. Returns 0 if `str` and `wildstr` match.

- `void get_date _A((string to, int timeflag));`

Get current date in a form ready for printing.

- `void soundex _A((string out_pntr, string in_pntr))`

Makes `in_pntr` to a 5 char long string. All words that sound alike have the same string.

- `int init_key_cache _A((ulong use_mem, ulong leave_this_much_mem));`

Use caching of keys in MISAM, PISAM, and ISAM. `KEY_CACHE_SIZE` is a good size. Remember to lock databases for optimal caching.

- `void end_key_cache _A((void));`

End key caching.

Bitmaps

Inside the `mysys` directory is a file named `my_bitmap.c`. It contains functions for manipulating bitmaps. Specifically there are functions for setup or teardown (`bitmap_init`, `bitmap_free`), for setting and clearing individual bits or whole sections of the bitmap (`bitmap_set_bit`, `bitmap_fast_test_and_set`, `bitmap_clear_all`, `bitmap_set_all`, `bitmap_set_prefix`, `bitmap_set_above`), and for performing comparisons and set operations on two bitmaps (`bitmap_cmp`, `bitmap_intersect`, `bitmap_subtract`, `bitmap_union`). Bitmaps are useful, so the functions are called from several places (`opt_range.cc`, `slave.cc`, `mysqld.c`, `sql_insert.cc`, `log_event.cc`, `sql_show.cc`) and we're expecting to make more use of them in the next version of MySQL, MySQL 5.1.

There are a few warnings and limitations that apply for the present bitmap implementation. First: the allocation is an integral number of bytes, and it is not possible to determine whether the last few bits are meaningful. Second: the whole bitmap might have to be protected by a mutex for manipulations; this is settable by passing appropriate flag values. Third: the bitmap is allocated with a 'uint' size, which means that ordinarily it can't have more than 2^{32} bytes. Fourth: when unioning two bitmaps, they must be of the same size.

Chapter 12. File Formats

Table of Contents

MySQL `.frm` File Format 117

MySQL uses a number of different file formats for the storage of information. This section covers the different file formats and how to read, write and understand the contents.

MySQL `.frm` File Format

Regardless of the storage engine you choose, every MySQL table you create is represented, on disk, by a `.frm` file, which describes the table's format (i.e. the table definition). The file bears the same name as the table, with a `.frm` extension. The `.frm` format is the same on all platforms but in the description of the `.frm` format that follows, the examples come from tables created under the Linux operating system.

First, let's create an example table, using the `mysql` client:

```
mysql> CREATE TABLE table1 (column1 CHAR(5)) ENGINE=MYISAM COMMENT '*';
Query OK, 0 rows affected (0.00 sec)
```

The `.frm` file associated with `table1` can be located in the directory that represents the database (or schema) to which the table belongs. The `datadir` variable contains the name of this directory:

```
mysql> SHOW VARIABLES LIKE 'datadir';
+-----+-----+
| Variable_name | Value                               |
+-----+-----+
| datadir       | /usr/local/mysql/var/             |
+-----+-----+
1 row in set (0.00 sec)
```

The `DATABASE()` function contains the name of the relevant database:

```
mysql> SELECT DATABASE();
+-----+
| DATABASE() |
+-----+
| ff         |
+-----+
1 row in set (0.00 sec)
```

Since MySQL stores `.frm` files in `datadir/database_name`, it's a simple matter to locate the corresponding `.frm` file for `table1`. For example, within a Linux shell:

```
shell> su root
shell> cd /usr/local/mysql/var/ff
shell> ls table1.*
```

You'll see a response like the following:

```
shell> su root
shell> cd /usr/local/mysql/var/ff
shell> ls table1.*
table1.frm  table1.MYD  table1.MYI
shell> ls -l table1.*
-rw-rw---- 1 root root 8566 2006-09-22 11:22 table1.frm
-rw-rw---- 1 root root    0 2006-09-22 11:22 table1.MYD
-rw-rw---- 1 root root 1024 2006-09-22 11:22 table1.MYI
```

The `.MYD` and `.MYI` files are not our concern here; they are described elsewhere in this MySQL internals manual. To understand the `.frm` format, let's look at `table1.frm` using a hexadecimal-dump utility; the contents are show below.

```
shell> hexdump -v -C table1.frm
```

```

00000000 fe 01 09 09 03 00 00 10 01 00 00 30 00 00 10 00 | .....0....|
00000010 06 00 00 00 00 00 00 00 00 00 00 02 08 00 08 00 | .....|
00000020 00 05 00 00 00 00 08 00 00 00 00 00 00 00 00 10 | .....|
00000030 00 00 00 c0 c3 00 00 10 00 00 00 00 00 00 00 00 | .....|
00000040 2f 2f 00 00 20 00 00 00 00 00 00 00 00 00 00 00 | //.. ....|
... | (many 0s) |
00001000 00 00 00 00 02 00 ff 00 00 00 00 00 00 00 00 00 | .....|
00001010 ff 20 20 20 20 20 00 00 06 00 4d 79 49 53 41 4d | . ....MyISAM|
... | (many 0s) |
00002000 6c 01 00 10 00 00 00 00 00 00 00 00 00 00 00 00 | l.....|
00002010 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 | .....|
00002020 00 00 00 00 00 00 00 00 00 00 00 00 00 01 2a | .....*|
... |
00002100 01 00 01 00 3b 00 05 00 00 00 06 00 0a 00 00 00 | ....i.....|
00002110 00 00 00 00 00 00 50 00 16 00 01 00 00 00 00 00 | .....P.....|
00002120 3b 00 02 01 02 14 29 20 20 20 20 20 20 20 20 20 | i.....)|
00002130 20 20 20 20 20 20 20 20 20 20 20 20 20 20 20 20 | |
00002140 20 20 20 20 20 20 20 20 20 20 20 20 20 20 20 00 | |
00002150 04 00 08 63 6f 6c 75 6d 6e 31 00 04 08 05 05 00 | ...column1.....|
00002160 02 00 00 00 80 00 00 00 fe 08 00 00 ff 63 6f 6c | .....col|
00002170 75 6d 6e 31 ff 00 |umn1..|
00002176

```

The details shown above might change, especially since there is a transition underway from an old ("binary") format to a new ("text based") `.frm` format. You can confirm that the details are correct by comparing this description with the statements in [sql/table.cc](#), [create_frm\(\)](#). The table below explains the meaning of each byte in the hexadecimal dump shown in the preceding example. The Offset column shows the byte position in the file; Length is the number of bytes; Value is what's in that byte position for that length (remember that storage is "low byte first" so 0010 means 1000!); and Explanation provides a brief explanation of the contents.

Offset	Length	Value	Explanation
header	??	??	??
0000	1	fe	Always
0001	1	01	Always
0002	1	09	FRM_VER (which is in <code>include/mysql_version.h</code>) +3 +test(create_info->varchar)
0003	1	09	See enum <code>legacy_db_type</code> in <code>sql/handler.h</code> . e.g. 09 is <code>DB_TYPE_MYISAM</code> , but 14 if MyISAM with partitioning.
0004	1	03	??
0005	1	00	Always
0006	2	0010	IO_SIZE
0008	2	0100	??
000a	4	00300000	Length, based on <code>key_length + rec_length + create_info->extra_size</code>
000e	2	1000	"tmp_key_length", based on <code>key_length</code>

0010	2	0600	rec_length
0012	4	00000000	create_info->max_rows
0016	4	00000000	create_info->min_rows
001b	1	02	Always (means "use long pack-fields")
001c	2	0800	key_info_length
001e	2	0800	create_info->table_options also known as db_create_options? one possible option is HA_LONG_BLOB_PTR
0020	1	00	always
0021	1	05	Always (means "version 5 frm file")
0022	4	00000000	create_info->avg_row_length
0026	1	08	create_info->default_table_charset
0027	1	00	Always
0028	1	00	create_info->row_type
0029	6	00..00	Always (formerly used for RAID support)
002f	4	10000000	key_length
0033	4	c0c30000	MYSQL_VERSION_ID from include/ mysql_version.h
0037	4	10000000	create_info->extra_size
003b	2	0000	Reserved for extra_rec_buf_length
003d	1	00	Reserved for default_part_db_type, but 09 if MyISAM with partitioning
003e	2	0000	create_info->key_block_size
key_info	??	??	??
1000	1	00	Always 00 when there are no keys i.e. indexes
...	??	??	??
101a	6	"MyISAM"	Name of engine. If partitioning, the partition clauses are here
comment	??	??	??

202e	1	01	Length of comment
202f	40	"*"	The string in the COMMENT clause
columns	??	??	??
2100	2	01	Always
2102	2	0100	<code>share->fields</code> (number of columns)
2104	2	3b00	<code>pos</code> ("length of all screens"). Goes up if <code>column-name</code> length increases. Doesn't go up if add comment.
2106	2	0500	Based on number of bytes in row.
210c	2	0500	<code>n_length</code> . Goes up if row length increases
210e	2	0000	<code>interval_count</code> . Number of different enum/set columns
2110	2	0000	<code>interval_parts</code> . Number of different strings in enum/set columns
2112	2	0000	<code>int_length</code>
211a	2	0100	<code>share->null_fields</code> . Number of nullable columns
211c	2	0000	<code>com_length</code>
2152	1	08	Length of column-name including '\0' termination
2153	3	"column1\0"	column-name
215b	1	04	??
215c	1	03	??
215d	1	05	Number of bytes in column
215e	1	05	Number of bytes in column
215f	4	00020000	??
2163	1	00	Flags for zerofill, unsigned, etc.
2164	1	80	Additional flags, and scale if decimal/numeric
2168	1	fe	Data type (fe= <code>char</code> , 02= <code>smallint</code> , 03= <code>int</code> , etc.) see <code>enum field_types</code> in <code>include/mysql_com.h</code>

2169	1	08	Character set or geometry type
(later)	??	??	Column names again, defaults, enum/ set strings, column comments ... at end of row. not shown.

The `.frm` file for a partitioned table contains partition information, in clear text, in addition to the usual table definition details. Let's create a partitioned table and do a hexadecimal dump of its `.frm`.

```
mysql> CREATE TABLE table2 (column1 INT) ENGINE=MYISAM COMMENT '*'
      PARTITION BY HASH(column1) PARTITIONS 2;
Query OK, 0 rows affected (0.00 sec)
```

The hexadecimal dump from `table2` is shown below:

```
00000000 fe 01 09 14 03 00 00 10 01 00 00 30 00 00 10 00 | .....0....|
00000010 05 00 00 00 00 00 00 00 00 00 00 02 08 00 08 00 | .....|
00000020 00 05 00 00 00 00 08 00 00 00 00 00 00 00 10 | .....|
00000030 00 00 00 c0 c3 00 00 3d 00 00 00 00 00 09 00 00 | .....=.....|
00000040 2f 2f 00 00 20 00 00 00 00 00 00 00 00 00 00 | //.. ....|
00000050 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 | .....|
...
00001000 00 00 00 00 02 00 ff 00 00 00 00 00 00 00 00 | .....|
00001010 ff 00 00 00 00 00 00 09 00 70 61 72 74 69 74 69 | .....partiti|
00001020 6f 6e 2a 00 00 00 20 50 41 52 54 49 54 49 4f 4e | on*... PARTITION|
00001030 20 42 59 20 48 41 53 48 20 28 63 6f 6c 75 6d 6e | BY HASH (column|
00001040 31 29 20 50 41 52 54 49 54 49 4f 4e 53 20 32 20 | 1) PARTITIONS 2 |
00001050 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 | .....|
...
00002000 76 01 00 10 00 00 00 00 00 00 00 00 00 00 00 | v.....|
00002010 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 | .....|
00002020 00 00 00 00 00 00 00 00 00 00 00 00 01 2a | .....*|
00002030 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 | .....|
...
00002100 01 00 01 00 3b 00 0b 00 00 00 05 00 0a 00 00 00 | ...;.....|
00002110 00 00 00 00 00 00 50 00 16 00 01 00 00 00 00 00 | .....P.....|
00002120 3b 00 02 01 02 14 29 20 20 20 20 20 20 20 20 | ;.....)|
00002130 20 20 20 20 20 20 20 20 20 20 20 20 20 20 20 | .....|
00002140 20 20 20 20 20 20 20 20 20 20 20 20 20 20 00 | .....|
00002150 04 00 08 63 6f 6c 75 6d 6e 31 00 04 08 0b 0b 00 | ...column1.....|
00002160 02 00 00 1b 80 00 00 00 03 08 00 00 ff 63 6f 6c | .....col|
00002170 75 6d 6e 31 ff 00 | umn1..|
00002176
```

In the example output, notice that position 00001010 and following contains the clear text of the `CREATE TABLE ... PARTITION` clause and not just the `MYISAM` engine information, as in `table1`, which shows the `.frm` of a non-partitioned table.

Finally, `CREATE VIEW` also causes creation of a `.frm` file, but a view `.frm` bears no resemblance to a base table `.frm`; it's purely textual. Here's an example of a `.frm` for a view made with:

```
mysql> CREATE VIEW v AS SELECT 5;
Query OK, 0 rows affected (0.00 sec)
```

Just looking at the text will tell you what it's about. For example let's do another hexadecimal dump; the contents are shown below.

```
linux:/usr/local/mysql/var/d # hexdump -v -C v.frm
00000000 54 59 50 45 3d 56 49 45 57 0a 71 75 65 72 79 3d | TYPE=VIEW.query=|
00000010 73 65 6c 65 63 74 20 35 20 41 53 20 60 35 60 0a | select 5 AS `5`.|
00000020 6d 64 35 3d 38 64 39 65 32 62 62 66 64 35 33 35 | md5=8d9e2bbfd535|
00000030 66 35 37 39 64 34 61 39 34 39 62 39 65 62 37 64 | f579d4a949b9eb7d|
00000040 32 33 34 39 0a 75 70 64 61 74 61 62 6c 65 3d 30 | 2349.updatable=0|
00000050 0a 61 6c 67 6f 72 69 74 68 6d 3d 30 0a 64 65 66 | .algorithm=0.def|
```

```

00000060 69 6e 65 72 5f 75 73 65 72 3d 72 6f 6f 74 0a 64 |iner_user=root.d|
00000070 65 66 69 6e 65 72 5f 68 6f 73 74 3d 6c 6f 63 61 |efiner_host=loca|
00000080 6c 68 6f 73 74 0a 73 75 69 64 3d 32 0a 77 69 74 |lhost.suid=2.wit|
00000090 68 5f 63 68 65 63 6b 5f 6f 70 74 69 6f 6e 3d 30 |h_check_option=0|
000000a0 0a 72 65 76 69 73 69 6f 6e 3d 31 0a 74 69 6d 65 |.revision=1.time|
000000b0 73 74 61 6d 70 3d 32 30 30 36 2d 30 39 2d 32 32 |stamp=2006-09-22|
000000c0 20 31 32 3a 31 34 3a 34 38 0a 63 72 65 61 74 65 | 12:14:48.create|
000000d0 2d 76 65 72 73 69 6f 6e 3d 31 0a 73 6f 75 72 63 |-version=1.sourc|
000000e0 65 3d 73 65 6c 65 63 74 20 35 0a |e=select 5.|

```

Chapter 13. How MySQL Performs Different Selects

Table of Contents

Steps of Select Execution	123
<code>select_result</code> Class	123
<code>SIMPLE</code> or <code>PRIMARY SELECT</code>	124
Structure Of Complex Select	124
Non-Subquery <code>UNION</code> Execution	125
Derived Table Execution	126
Subqueries	126
Single Select Engine	127
Union Engine	127
Special Engines	127
Explain Execution	128

Steps of Select Execution

Every select is performed in these base steps:

- `JOIN::prepare`
 - Initialization and linking `JOIN` structure to `st_select_lex`.
 - `fix_fields()` for all items (after `fix_fields()`, we know everything about item).
 - Moving `HAVING` to `WHERE` if possible.
 - Initialization procedure if there is one.
- `JOIN::optimize`
 - Single select optimization.
 - Creation of first temporary table if needed.
- `JOIN::exec`
 - Performing select (a second temporary table may be created).
- `JOIN::cleanup`
 - Removing all temporary tables, other cleanup.
- `JOIN::reinit`
 - Prepare all structures for execution of `SELECT` (with `JOIN::exec`).

`select_result`Class

This class has a very important role in `SELECT` performance with `select_result` class and classes inherited from it (usually called with a `select_` prefix). This class provides the interface for transmitting results.

The key methods in this class are the following:

- `send_fields` sends given item list headers (type, name, etc.).
- `send_data` sends given item list values as row of table of result.

- `send_error` is used mainly for error interception, making some operation and then `::send_error` will be called.

For example, there are the following `select_result` classes:

- `select_send` used for sending results though network layer.
- `select_export` used for exporting data to file.
- `multi_delete` used for multi-delete.
- `select_insert` used for `INSERT ... SELECT ...`
- `multi_update` used for multi-update.
- `select_singlerow_subselect` used for row and scalar subqueries..
- `select_exists_subselect` used for `EXISTS/IN/ALL/ANY/SOME` subqueries.
- `select_max_min_finder_subselect` used for min/max subqueries (`ALL/ANY` subquery optimization).

SIMPLE or PRIMARY SELECT

For performing single primary select, `SELECT` uses the `mysql_select` function, which does:

- allocate `JOIN`
- `JOIN::prepare`
- `JOIN::optimize`
- `JOIN::exec`
- `JOIN::cleanup`

In previous versions of MySQL, all `SELECT` operations were performed with the help of this function and `mysql_select()` was not divided into parts.

Structure Of Complex Select

There are two structures that describe selects:

- `st_select_lex` (`SELECT_LEX`) for representing `SELECT` itself
- `st_select_lex_unit` (`SELECT_LEX_UNIT`) for grouping several selects in a bunch

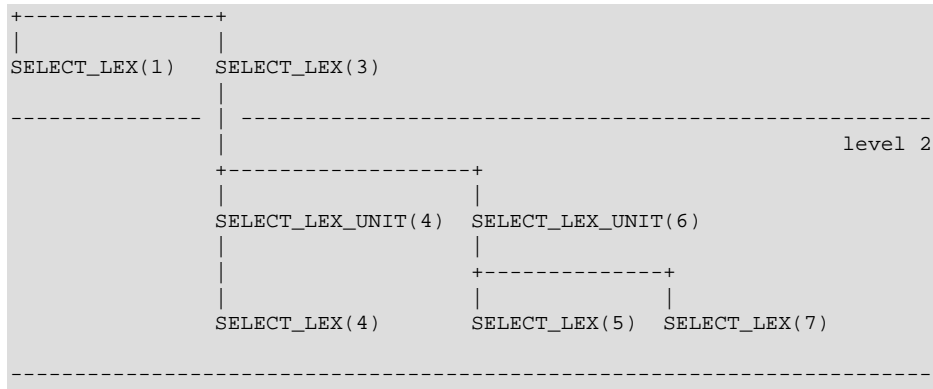
The latter item represents `UNION` operation (the absence of `UNION` is a union with only one `SELECT` and this structure is present in any case). In the future, this structure will be used for `EXCEPT` and `INTERSECT` as well.

For example:

```
(SELECT ...) UNION (SELECT ... (SELECT...)...(SELECT...UNION...SELECT))
  1           2       3           4           5           6           7
```

will be represented as:

```
-----
SELECT_LEX_UNIT(2)                                     level 1
|
```



Note: Single subquery 4 has its own `SELECT_LEX_UNIT`.

The uppermost `SELECT_LEX_UNIT` (#2 in example) is stored in LEX. The first and uppermost `SELECT_LEX` (#1 in example) is stored in LEX, too. These two structures always exist.

At the time of creating or performing any `JOIN::*` operation, `LEX::current_select` points to an appropriate `SELECT_LEX`.

Only during parsing of global `ORDER BY` and `LIMIT` clauses (for the whole `UNION`), `LEX::current_select` points to `SELECT_LEX_UNIT` of this unit, in order to store this parameter in this `SELECT_LEX_UNIT`. `SELECT_LEX` and `SELECT_LEX_UNIT` are inherited from `st_select_lex_node`.

Non-Subquery UNION Execution

Non-subquery unions are performed with the help of `mysql_union()`. For now, it is divided into the following steps:

- `st_select_lex_unit::prepare` (the same procedure can be called for single `SELECT` for derived table => we have support for it in this procedure, but we will not describe it here):
 - Create `select_union` (inherited from `select_result`) which will write select results in this temporary table, with empty temporary table entry. We will need this object to store in every `JOIN` structure link on it, but we have not (yet) temporary table structure.
 - Allocate `JOIN` structures and execute `JOIN::prepare()` for every `SELECT` to get full information about types of elements of `SELECT` list (results). Merging types of result fields and storing them in special items (`Item_type_holder`) will be done in this loop, too. Result of this operation (list of types of result fields) will be stored in `st_select_lex_unit::types`.
 - Create a temporary table for storing union results (if `UNION` without `ALL` option, 'distinct' parameter will be passed to the table creation procedure).
 - Assign a temporary table to the `select_union` object created in the first step.
- `st_select_lex_unit::exec`
 - Delete rows from the temporary table if this is not the first call.
 - if this is the first call, call `JOIN::optimize` else `JOIN::reinit` and then `JOIN::exec` for all `SELECTs` (`select_union` will write a result for the temporary table). If union is cacheable and this is not the first call, the method will do nothing.
 - Call `mysql_select` on temporary table with global `ORDER BY` and `LIMIT` parameters after collecting results from all `SELECTs`. A special `fake_select_lex` (`SELECT_LEX`) which is created for every `UNION` will be passed for this procedure (this `SELECT_LEX` also can be used to store global `ORDER BY` and `LIMIT` parameters if brackets used in a query).

Derived Table Execution

Derived tables is the internal name for subqueries in the `FROM` clause.

The processing of derived tables is now included in the table opening process (`open_and_lock_tables()` call). Routine of execution derived tables and substituting temporary table instead of it (`mysql_handle_derived()`) will be called just after opening and locking all real tables used in query (including tables used in derived table query).

If `lex->derived_tables` flag is present, all `SELECT_LEX` structures will be scanned (there is a list of all `SELECT_LEX` structures in reverse order named `lex->all_selects_list`, the first `SELECT` in the query will be last in this list).

There is a pointer for the derived table, `SELECT_LEX_UNIT` stored in the `TABLE_LIST` structure (`TABLE_LIST::derived`). For any table that has this pointer, `mysql_derived()` will be called.

`mysql_derived()`:

- Creates `union_result` for writing results in this table (with empty table entry, same as for `UNIONS`).
- call `unit->prepare()` to get list of types of result fields (it work correctly for single `SELECT`, and do not create temporary table for `UNION` processing in this case).
- Creates a temporary table for storing results.
- Assign this temporary table to `union_result` object.
- Calls `mysql_select` or `mysql_union` to execute the query.
- If it is not explain, then cleanup `JOIN` structures after execution (`EXPLAIN` needs data of optimization phase and cleanup them after whole query processing).
- Stores pointer to this temporary table in `TABLE_LIST` structure, then this table will be used by outer query.
- Links this temporary table in `thd->derived_tables` for removing after query execution. This table will be closed in `close_thread_tables` if its second parameter (`bool skip_derived`) is true.

Subqueries

In expressions, subqueries (that is, subselects) are represented by `Item` inherited from `Item_subselect`.

To hide difference in performing single `SELECT`s and `UNIONS`, `Item_subselect` uses two different engines, which provide uniform interface for access to underlying `SELECT` or `UNION` (`subselect_single_select_engine` and `subselect_union_engine`, both are inherited from `subselect_engine`).

The engine will be created at the time `Item_subselect` is constructed (`Item_subselect::init` method).

On `Item_subselect::fix_fields()`, `engine->prepare()` will be called.

Before calling any value-getting method (`val`, `val_int`, `val_str`, `bring_value` (in case of row result)) `engine->exec()` will be called, which executes the query or just does nothing if subquery is cacheable and has already been executed.

Inherited items have their own `select_result` classes. There are two types of them:

- `select_singlerow_subselect`, to store values of given rows in `Item_singlerow_subselect` cache on `send_data()` call, and report error if `Item_subselect` has 'assigned' attribute.

- `select_exists_subselect` just store 1 as value of `Item_exists_subselect` on `send_data()` call. Since `Item_in_subselect` and `Item_allany_subselect` are inherited from `Item_exists_subselect`, they use the same `select_result` class.

`Item_subselect` will never call the `cleanup()` procedure for `JOIN`. Every `JOIN::cleanup` will call `cleanup()` for inner `JOINS`. The uppermost `JOIN::cleanup` will be called by `mysql_select()` or `mysql_union()`.

Single Select Engine

`subselect_single_select_engine`:

- `constructor` allocate `JOIN` and store pointers on `SELECT_LEX` and `JOIN`.
- `prepare()` call `JOIN::prepare`.
- `fix_length_and_dec()` prepare cache and receive type and parameters of returning items (called only by `Item_singlerow_subselect`).
- `exec()` drop 'assigned' flag of `Item_subselect`. If this is the first time, call `JOIN::optimize` and `JOIN::exec()`, else do nothing or `JOIN::reinit()` `JOIN::exec()` depending on type of subquery.

Union Engine

`subselect_union_engine`:

- `constructor` just store pointer to `st_select_lex_union` (`SELECT_LEX_UNION`).
- `prepare()` call `st_select_lex_unit::prepare`.
- `fix_length_and_dec()` prepare cache and receive type and parameters (maximum of length) of returning items (called only by `Item_singlerow_subselect`).
- `exec()` call `st_select_lex_unit::exec()`. `st_select_lex_unit::exec()` can drop 'assigned' flag of `Item_subselect` if `st_select_lex_unit::item` is not 0.

Special Engines

There are special engines used for optimization purposes. These engines do not have a full range of features. They can only fetch data. The normal engine can be replaced with such special engines only during the optimization process.

Now we have two such engines:

`subselect_uniquesubquery_engine` used for:

```
left_expression IN (SELECT primary_key FROM table WHERE conditions)
```

This looks for the given value once in a primary index, checks the `WHERE` condition, and returns was it found or not?

- `subselect_indexsubquery_engine` used for:

```
left_expression IN (SELECT any_key FROM table WHERE conditions)
```

This first looks up the value of the left expression in an index (checking the `WHERE` condition), then if value was not found, it checks for `NULL` values so that it can return `NULL` correctly (only if a `NULL` result makes sense, for example if an `IN` subquery is the top item of the `WHERE` clause then `NULL` will not be sought)

The decision about replacement of the engine happens in `JOIN::optimize`, after calling `make_join_readinfo`, when we know what the best index choice is.

Explain Execution

For an `EXPLAIN` statement, for every `SELECT`, `mysql_select` will be called with option `SELECT_DESCRIBE`.

For main `UNION`, `mysql_explain_union` will be called.

For every `SELECT` in a given union, `mysql_explain_union` will call `mysql_explain_select`.

`mysql_explain_select` will call `mysql_select` with option `SELECT_DESCRIBE`.

`mysql_select` creates a `JOIN` for select if it does not already exist (it might already exist because if it called for subquery `JOIN` can be created in `JOIN::optimize` of outer query when it decided to calculate the value of the subquery). Then it calls `JOIN::prepare`, `JOIN::optimize`, `JOIN::exec` and `JOIN::cleanup` as usual.

`JOIN::exec` is called for `SELECT` with `SELECT_DESCRIBE` option call `select_describe`.

`select_describe` returns the user description of `SELECT` and calls `mysql_explain_union` for every inner `UNION`.

PROBLEM: how it will work with global query optimization?

Chapter 14. How MySQL Transforms Subqueries

Table of Contents

<code>Item_in_subselect::select_transformer</code>	129
Scalar <code>IN</code> Subquery	129
Row <code>IN</code> Subquery	131
<code>Item_allany_subselect</code>	132
<code>Item_singlerow_subselect</code>	132

The `Item_subselect` virtual method `select_transformer` is used to rewrite subqueries. It is called from `Item_subselect::init` (which is called just after the call to the `fix_fields()` method for all items in `JOIN::prepare`).

`Item_in_subselect::select_transformer`

`Item_in_subselect::select_transformer` is divided into two parts, one for the scalar left part and one for the row left part.

Scalar `IN` Subquery

To rewrite a scalar `IN` subquery, the `Item_in_subselect::single_value_transformer` method is used. The scalar `IN` subquery will be replaced with an `Item_in_optimizer` item.

An `Item_in_optimizer` item is a special boolean function. On a value request (one of `val`, `val_int`, or `val_str` methods) it evaluates the left expression of the `IN` by storing its value in a cache item (one of `Item_cache*` items), then it tests the cache to see whether it is `NULL`. If left expression (cache) is `NULL`, then `Item_in_optimizer` returns `NULL`, else it evaluates `Item_in_subselect`.

Example queries.

```
a) SELECT * from t1 where t1.a in (SELECT t2.a FROM t2);
b) SELECT * from t1 where t1.a in (SELECT t2.a FROM t2 GROUP BY t2.a);
```

- `Item_in_subselect` inherits the mechanism for getting a value from `Item_exists_subselect`.
- `Select_transformer` stores a reference to the left expression in its conditions:

```
(in WHERE and HAVING in case 'a' and in HAVING in case 'b')
```

- Item from item list of this select (`t2.a`) can be referenced with a special reference (`Item_ref_null_helper` or `Item_null_helper`). This reference informs `Item_in_optimizer` whether item (`t2.a`) is `NULL` by setting the `was_null` flag.
- The return value from `Item_in_subselect` will be evaluated as follows:
 - If `TRUE`, return `true`
 - If `NULL`, return `null` (that is, unknown)
 - If `FALSE`, and `was_null` is set, return `null`
 - Return `FALSE`

`<left_expression> IN (SELECT <item> ...)` will be represented as follows:

```
+-----+
```



```
Item_ref(<cached_left_expression>) = Item_ref_null_helper(item))
```

- Example 3:

```
<left_expression> IN (SELECT <item> UNION ...)
```

will become

```
(SELECT 1
  HAVING Item_ref(<cached_left_expression>)= <Item_null_helper(<Item>)>
 UNION ...)
```

(**HAVING** without **FROM** is a syntax error, but a **HAVING** condition is checked even for subquery without **FROM**)

- Example 4:

```
<left_expression> IN (select <item>)
```

will be completely replaced with `<left_expression> = <item>`

Now conditions (**WHERE** (a) or **HAVING** (b)) will be changed, depending on the select, in the following way:

If subquery contains a **HAVING** clause, **SUM()** function or **GROUP BY** (example 1), then the item list will be unchanged and an `Item_ref_null_helper` reference will be created on item list element. A condition will be added to the **HAVING**.

If the subquery does not contain **HAVING**, **SUM()** function, or **GROUP BY** (example 2), then:

- `item list` will be replaced with 1.
- `left_expression cache> = <item> or is null <item>` will be added to the **WHERE** clause and a special `is_not_null(item)` will be added to the **HAVING**, so null values will be registered. If returning **NULL** wouldn't make correct sense, then only `left_expression cache> = <item>` will be added to the **WHERE** clause. If this subquery does not contain a **FROM** clause or if the subquery contains **UNION** (example 3), then `left_expression cache> = Item_ref_null_helper(<item>)` will be added to the **HAVING** clause.

A single select without a **FROM** clause will be reduced to just `<left_expression> = <item>` without use of `Item_in_optimizer`.

Row IN Subquery

To rewrite a row **IN** subquery, the method used is `Item_in_subselect::row_value_transformer`. It works in almost the same way as the scalar analog, but works with `Item_cache_row` for caching left expression and uses references for elements of `Item_cache_row`. To refer to the item list, it uses `Item_ref_null_helper(ref_array+i)`.

A subquery with **HAVING**, **SUM()** function, or **GROUP BY** will transformed in the following way:

```
ROW(l1, l2, ... lN) IN (SELECT i1, i2, ... iN FROM t HAVING <having_expr>)
```

will become:

```
(SELECT i1, i2, ... iN FROM t
  HAVING <having_expr> and <cache_l0> = <Item_ref_null_helper(ref_array[0]> AND <cache_l1> =
  ...
  <cache_lN-1> = <Item_ref_null_helper(ref_array[N-1]>)
```

SELECT without **FROM** will be transformed in this way, too.

It will be the same for other subqueries, except for the **WHERE** clause.

Item_allany_subselect

`Item_allany_subselect` is inherited from `Item_in_subselect`. `ALL/ANY/SOME` use the same algorithm (and the same method of `Item_in_subselect`) as scalar `IN`, but use a different function instead of `=`.

`ANY/SOME` use the same function that was listed after the left expression.

`ALL` uses an inverted function, and all subqueries passed as arguments to `Item_func_not_all` (`Item_func_not_all` is a special `NOT` function used in optimization, see following).

But before above transformation ability of independent `ALL/ANY/SOME` optimization will be checked (query is independent, operation is one of `<`, `<=`, `>`, `>=`, returning correct `NULL` have no sense (top level of `WHERE` clause) and it is not row subquery).

For such queries, the following transformation can be done:

```
val > ALL (SELECT...) -> val >  MAX (SELECT...)
val < ALL (SELECT...) -> val <  MIN (SELECT...)
val > ANY (SELECT...) -> val >  MIN (SELECT...)
val < ANY (SELECT...) -> val <  MAX (SELECT...)
val >= ALL (SELECT...) -> val >= MAX (SELECT...)
val <= ALL (SELECT...) -> val <= MIN (SELECT...)
val >= ANY (SELECT...) -> val >= MIN (SELECT...)
val <= ANY (SELECT...) -> val <= MAX (SELECT...)
```

`ALL` subqueries already have `NOT` before them. This problem can be solved with help of special `NOT`, which can bring 'top' tag to its argument and correctly process `NULL` if it is 'top' item (return `TRUE` if argument is `NULL` if it is 'top' item). Let's call this operation `_NOT_`. Then we will have following table of transformation:

```
val > ALL (SELECT...) -> _NOT_ val >= MAX (SELECT...)
val < ALL (SELECT...) -> _NOT_ val <= MIN (SELECT...)
val > ANY (SELECT...) -> val <  MIN (SELECT...)
val < ANY (SELECT...) -> val >  MAX (SELECT...)
val >= ALL (SELECT...) -> _NOT_ val >  MAX (SELECT...)
val <= ALL (SELECT...) -> _NOT_ val <  MIN (SELECT...)
val >= ANY (SELECT...) -> val <= MIN (SELECT...)
val <= ANY (SELECT...) -> val >= MAX (SELECT...)
```

If subquery does not contain grouping and aggregate function, above subquery can be rewritten with `MAX()/MIN()` aggregate function, for example:

```
val > ANY (SELECT item ...) -> val < (SELECT MIN(item)...)

```

For queries with aggregate function and/or grouping, special `Item_maxmin_subselect` will be used. This subquery will return the maximum (minimum) value of the result set.

Item_singlerow_subselect

`Item_singlerow_subselect` will be rewritten only if it contains no `FROM` clause, and it is not part of `UNION`, and it is a scalar subquery. For now, there will be no conversion of subqueries with field or reference on top of item list (on the one hand, we can't change the name of such items, but on the other hand, we should assign to it the name of the whole subquery which will be reduced).

The following will not be reduced:

```
SELECT a;
SELECT 1 UNION SELECT 2;
SELECT 1 FROM t1;
```

The following select will be reduced:

```
SELECT 1;
```

```
SELECT a+2;
```

Such a subquery will be completely replaced by its expression from item list and its `SELECT_LEX` and `SELECT_LEX_UNIT` will be removed from `SELECT_LEX`'s tree.

But every `Item_field` and `Item_ref` of that expression will be marked for processing by a special `fix_fields()` procedure. The `fix_fields()` procedures for such `Items` will be performed in the same way as for items of an inner subquery. Also, if this expression is `Item_fields` or `Item_ref`, then the name of this new item will be the same as the name of this item (but not `(SELECT ...)`). This is done to prevent broken references on such items from more inner subqueries.

Chapter 15. MySQL Client/Server Protocol

Table of Contents

Organization	136
Elements	136
The Packet Header	137
Packet Types	137
Handshake Initialization Packet	137
Client Authentication Packet	139
Password Functions	140
Passwords Before MySQL 4.0	140
Passwords in MySQL 4.1 and Later	140
Command Packet (Overview)	140
Command Packet (Detailed Description)	141
COM_QUIT	141
COM_INIT_DB	141
COM_QUERY	141
COM_FIELD_LIST	141
COM_CREATE_DB	141
COM_DROP_DB	142
COM_REFRESH	142
COM_SHUTDOWN	142
COM_STATISTICS	142
COM_PROCESS_INFO	142
COM_PROCESS_KILL	142
COM_DEBUG	143
COM_PING	143
COM_CHANGE_USER	143
COM_BINLOG_DUMP	143
COM_TABLE_DUMP	143
COM_REGISTER_SLAVE	143
COM_PREPARE	144
COM_EXECUTE	144
COM_LONG_DATA	144
COM_CLOSE_STMT	144
COM_RESET_STMT	144
COM_SET_OPTION	144
COM_FETCH_STMT	144
Result Packet Types	145
OK Packet	145
Error Packet	146
Result Set Header Packet	147
Field Packet	147
Row Data Packet	149
Row Data Packet: Binary (Tentative Description)	150
EOF Packet	150
OK for Prepared Statement Initialization Packet	151
Parameter Packet (Tentative Description)	152
Long Data Packet (Tentative Description)	152
Execute Packet (Tentative Description)	152
Compression	153
Encryption	154

Organization

The topic is: the contents of logical packets in MySQL version 5.0 client/server communication.

The description is of logical packets. There will be only passing mention of non-logical considerations, such as physical packets, transport, buffering, and compression. If you are interested in those topics, you may wish to consult another document: "MySQL Client - Server Protocol Documentation" in the file `net_doc.txt` in the `internals` directory of the `mysqldoc` MySQL documentation repository.

The description is of the version-5.0 protocol at the time of writing. Most of the examples show version-4.1 tests, which is okay because the changes from version-4.1 to version-5.0 were small.

A typical description of a packet will include:

"Bytes and Names". This is intended as a quick summary of the lengths and identifiers for every field in the packet, in order of appearance. The "Bytes" column contains the length in bytes. The Names column contains names which are taken from the MySQL source code whenever possible. If the version-4.0 and version-4.1 formats differ significantly, we will show both formats.

Descriptions for each field. This contains text notes about the usage and possible contents.

(If necessary) notes about alternative terms. Naming in this document is not authoritative and you will often see different words used for the same things, in other documents.

(If necessary) references to program or header files in the MySQL source code. An example of such a reference is: `sql/protocol.cc net_store_length()` which means "in the `sql` subdirectory, in the `protocol.cc` file, the function named `net_store_length`".

An Example. All examples have three columns:

```
-- the field name
-- a hexadecimal dump
-- an ascii dump, if the field has character data
```

All spaces and carriage returns in the hexadecimal dump are there for formatting purposes only.

In the later sections, related to prepared statements, the notes should be considered unreliable and there are no examples.

Elements

Null-Terminated String: used for some variable-length character strings. The value `'\0'` (sometimes written `0x00`) denotes the end of the string.

Length Coded Binary: a variable-length number. To compute the value of a Length Coded Binary, one must examine the value of its first byte.

Value Of First Byte	# Of Bytes Following	Description
0-250	0	= value of first byte
251	0	column value = NULL only appropriate in a Row Data Packet
252	2	= value of following 16-bit word
253	4	= value of following 32-bit word
254	8	= value of following 64-bit word

Thus the length of a Length Coded Binary, including the first byte, will vary from 1 to 9 bytes. The relevant MySQL source program is `sql/protocol.cc net_store_length()`.

All numbers are stored with the least significant byte first. All numbers are unsigned.

Length Coded String: a variable-length string. Used instead of Null-Terminated String, especially for character strings which might contain `'\0'` or might be very long. The first part of a Length Coded String is a Length Coded Binary number (the length); the second part of a Length Coded String is the actual

data. An example of a short Length Coded String is these three hexadecimal bytes: 02 61 62, which means "length = 2, contents = 'ab'".

The Packet Header

Bytes	Name
-----	----
3	Packet Length
1	Packet Number

Packet Length: The length, in bytes, of the packet that follows the Packet Header. There may be some special values in the most significant byte. Since $2^{24} = 16\text{MB}$, the maximum packet length is 16MB.

Packet Number: A serial number which can be used to ensure that all packets are present and in order. The first packet of a client query will have Packet Number = 0. Thus, when a new SQL statement starts, the packet number is re-initialized.

The Packet Header will not be shown in the descriptions of packets that follow this section. Think of it as always there. But logically, it "precedes the packet" rather than "is included in the packet".

Indeed, if the packet length is equal or greater than $(2^{24} - 1)$ Bytes, this packet must be split into two or more packets.

Alternative terms: Packet Length is also called "packet size". Packet Number is also called "Packet no".

Relevant MySQL Source Code:
include/my_global.h int3store()
sql/net_serv.cc my_net_write(), net_flush(), net_write_command(), my_net_read()

Packet Types

This is what happens in a typical session:

The Handshake (when client connects):

Server Sends To Client: Handshake Initialisation Packet

Client Sends To Server: Client Authentication Packet

Server Sends To Client: OK Packet, or Error Packet

The Commands (for every action the client wants the server to do):

Client Sends To Server: Command Packet

Server Sends To Client: OK Packet, or Error Packet, or Result Set Packet

In the rest of this chapter, you will find a description for each packet type, in separate sections.

Alternative terms: The Handshake is also called "client login" or "login procedure" or "connecting".

Handshake Initialization Packet

From server to client during initial handshake. The follow is taken from sql/sql_acl.cc:

Bytes	Name
-----	----
1	protocol_version
n (Null-Terminated String)	server_version
4	thread_id
8	scramble_buff

```

1          (filler) always 0x00
2          server_capabilities
1          server_language
2          server_status
13         (filler) always 0x00 ...

protocol_version:  The server takes this from PROTOCOL_VERSION
                   in /include/mysql_version.h. Example value = 10.

server_version:    The server takes this from MYSQL_SERVER_VERSION
                   in /include/mysql_version.h. Example value = "4.1.1-alpha".

thread_number:     ID of the server thread for this connection.

scramble_buff:     The password mechanism uses this.
                   (See "Password functions" section elsewhere in this document.)

server_capabilities: CLIENT_XXX options. The possible flag values at time of
writing (taken from include/mysql_com.h):
CLIENT_LONG_PASSWORD 1 /* new more secure passwords */
CLIENT_FOUND_ROWS 2 /* Found instead of affected rows */
CLIENT_LONG_FLAG 4 /* Get all column flags */
CLIENT_CONNECT_WITH_DB 8 /* One can specify db on connect */
CLIENT_NO_SCHEMA 16 /* Don't allow database.table.column */
CLIENT_COMPRESS 32 /* Can use compression protocol */
CLIENT_ODBC 64 /* Odbc client */
CLIENT_LOCAL_FILES 128 /* Can use LOAD DATA LOCAL */
CLIENT_IGNORE_SPACE 256 /* Ignore spaces before '(' */
CLIENT_PROTOCOL_41 512 /* New 4.1 protocol */
CLIENT_INTERACTIVE 1024 /* This is an interactive client */
CLIENT_SSL 2048 /* Switch to SSL after handshake */
CLIENT_IGNORE_SIGPIPE 4096 /* IGNORE sigpipes */
CLIENT_TRANSACTIONS 8192 /* Client knows about transactions */
CLIENT_RESERVED 16384 /* Old flag for 4.1 protocol */
CLIENT_SECURE_CONNECTION 32768 /* New 4.1 authentication */
CLIENT_MULTI_STATEMENTS 65536 /* Enable/disable multi-stmt support */
CLIENT_MULTI_RESULTS 131072 /* Enable/disable multi-results */

server_language:   current server character set number

server_status:     SERVER_STATUS_xxx flags: e.g. SERVER_STATUS_AUTOCOMMIT

```

Alternative terms: Handshake Initialization Packet is also called "greeting packet". Protocol version is also called "Prot. version". server_version is also called "Server Version String". thread_number is also called "Thread Number". current server charset number is also called "charset_no". scramble_buff is also called "crypt seed". server_status is also called "SERVER_STATUS_xxx flags" or "Server status variables".

Example Handshake Initialization Packet	Hexadecimal	ASCII
protocol_version	0a	.
server_version	34 2e 31 2e 31 2d 71 6c	4.1.1-al
	70 68 61 2d 64 65 62 75	pha-debu
	67 00	g.
thread_number	01 00 00 00
scramble_buff	3a 23 3d 4b 43 4a 2e 43
(filler)	00	.
server_capabilities	2c 82	..
server_language	08	.
server_status	02 00	..
(filler)	00 00 00 00 00 00 00 00
	00 00 00 00 00	

In the example, the server is telling the client that its server_capabilities include CLIENT_MULTI_RESULTS, CLIENT_SSL, CLIENT_COMPRESS, CLIENT_CONNECT_WITH_DB, CLIENT_FOUND_ROWS.

The "server_language" (or "charset") corresponds to the character_set_server variable in the MySQL server. This number also contains the collation used. Technically this number determines the collation

and the character set is implicit for the collation. You can use the following SQL statement to get the cleartext information:

```
mysql> SELECT CHARACTER_SET_NAME, COLLATION_NAME -> FROM INFORMATION_SCHEMA.COLLATIONS -> WHERE ID=8; +
```

Client Authentication Packet

From client to server during initial handshake.

```

VERSION 4.0
Bytes      Name
-----
2          client_flags
3          max_packet_size
n (Null-Terminated String) user
8          scramble_buff
1          (filler) always 0x00

VERSION 4.1
Bytes      Name
-----
4          client_flags
4          max_packet_size
1          charset_number
23         (filler) always 0x00...
n (Null-Terminated String) user
8          scramble_buff
1          (filler) always 0x00
n (Null-Terminated String) databasename

client_flags:      CLIENT_xxx options. The list of possible flag
                    values is in the description of the Handshake
                    Initialisation Packet, for server_capabilities.
                    For some of the bits, the server passed "what
                    it's capable of". The client leaves some of the
                    bits on, adds others, and passes back to the server.
                    One important flag is: whether compression is desired.

max_packet_size:   the maximum number of bytes in a packet for the client

charset_number:    in the same domain as the server_language field that
                    the server passes in the Handshake Initialisation packet.

user:              identification

scramble_buff:     the password, after encrypting using the scramble_buff
                    contents passed by the server (see "Password functions"
                    section elsewhere in this document)

databasename:      name of schema to use initially

```

The `scramble_buff` and `databasename` fields are optional. The length-coding byte for the `scramble_buff` will always be given, even if it's zero.

Alternative terms: "Client authentication packet" is sometimes called "client auth response" or "client auth packet" or "login packet". "Scramble_buff" is sometimes called "crypt password".

Relevant MySQL Source Code:

```

- On the client side: libmysql/libmysql.c::mysql_real_connect().
- On the server side: sql/sql_parse.cc::check_connections()

```

Example Client Authentication Packet

	Hexadecimal	ASCII
	-----	-----
client_flags	85 a6 03 00
max_packet_size	00 00 00 01
charset_number	08	.
(filler)	00 00 00 00 00 00 00 00
	00 00 00 00 00 00 00 00
	00 00 00 00 00 00 00

user	70 67 75 6c 75 74 7a 61 6e 00	pgulutzan.
------	----------------------------------	------------

Password Functions

The Server Initialization Packet and the Client Authentication Packet both have an 8-byte field, `scramble_buff`. The value in this field is used for password authentication.

Relevant MySQL Source Code: `libmysql/password.c`, see also comments at start of file. It works thus:

Passwords Before MySQL 4.0

- The server sends a random string to the client, in `scramble_buff`.
- The client encrypts the `scramble_buff` value using the hash of a password that the user has entered. This happens in `sql/password.c:scramble()` function.
- The client sends the encrypted `scramble_buff` value to the server.
- The server encrypts the original random string using a value in the mysql database, `mysql.user.Password`.
- The server compares its encrypted random string to what the client sent in `scramble_buff`.
- If they are the same, the password is okay.

In this protocol, snooping on the wire doesn't reveal the password. But note the problem - if the client doesn't know the password, but knows a hash of it (as stored in `mysql.user.Password`) it can connect to the server. In other words, the hash of a password **is** the real password; if one can get the value of `mysql.user.Password` - he can connect to the server.

Passwords in MySQL 4.1 and Later

Remember that `mysql.user.Password` stores `SHA1(SHA1(password))`

- The server sends a random string (scramble) to the client
- the client calculates:
 - `stage1_hash = SHA1(password)`, using the password that the user has entered.
 - `token = SHA1(scramble + SHA1(stage1_hash)) XOR stage1_hash`
- the client sends the token to the server
- the server calculates
 - `stage1_hash' = token XOR SHA1(scramble + mysql.user.Password)`
- the server compares `SHA1(stage1_hash')` and `mysql.user.Password`
- If they are the same, the password is okay.

(Note `SHA1(A+B)` is the SHA1 of the concatenation of A with B.)

This protocol fixes the flaw of the old one, neither snooping on the wire nor `mysql.user.Password` are sufficient for a successful connection. But when one has both `mysql.user.Password` and the intercepted data on the wire, he has enough information to connect.

Command Packet (Overview)

From client to server whenever the client wants the server to do something.

Bytes	Name
-----	----
1	command
n	arg

The command byte is stored in the `thd` structure for the MySQL worker threads and is shown in the Command column for `SHOW PROCESSLIST`. An inactive thread gets 0x00 (Sleep). The dedicated thread to execute `INSERT DELAYED` gets 0x10.

The replication requests (0x12 .. 0x15) cannot be send from regular clients, only from another server or from the `mysqlbinlog` program.

Relevant MySQL source code:

```
sql-common/client.c cli_advanced_command(), mysql_send_query().
libmysql/libmysql.c mysql_real_query(), simple_command(), net_field_length().
```

Example Command Packet

	Hexadecimal	ASCII
	-----	-----
command	02	.
arg	74 65 73 74	test

In the example, the value 02 in the command field stands for `COM_INIT_DB`. This is the packet that the client puts together for "use test;".

Command Packet (Detailed Description)

COM_QUIT

Closes the current connection. No arguments.

COM_INIT_DB

Functional equivalent to the SQL statement `USE <database>`. Exported by many clients, i.e. in PHP as `mysqli::select_db()`

Bytes	Name
-----	----
n	database name (up to end of packet, no termination character)

COM_QUERY

The most common request type. Used to execute nonprepared SQL statements.

Bytes	Name
-----	----
n	SQL statement (up to end of packet, no termination character)

COM_FIELD_LIST

Functional equivalent to `SHOW [FULL] FIELDS FROM ...`

Bytes	Name
-----	----
n	table name (null terminated)
n	column name or wildcard (optional)

COM_CREATE_DB

results from a call of the C-API function `mysql_create_db()`. This function is marked deprecated; the recommended way to create a database is to use the SQL statement `CREATE DATABASE`.

Bytes	Name
-----	----

n	database name (up to end of packet, no termination character)
---	--

COM_DROP_DB

results from a call of the C-API function *mysql_drop_db()*. This function is marked deprecated; the recommended way to drop a database is to use the SQL statement DROP DATABASE.

Bytes	Name
-----	----
n	database name (up to end of packet, no termination character)

COM_REFRESH

Parameter is one byte, evaluated as bitmap. Some (but not all) options are available via the FLUSH statement or via *mysqladmin flush-foo*. This is used by the C-API call *mysql_refresh()*.

Bytes	Name
-----	----
1	bitmap of refresh options 0x01...REFRESH_GRANT 0x02...REFRESH_LOG 0x04...REFRESH_TABLES 0x08...REFRESH_HOSTS 0x10...REFRESH_STATUS 0x20...REFRESH_THREADS 0x40...REFRESH_SLAVE 0x80...REFRESH_MASTER defined in mysql_com.h

COM_SHUTDOWN

Asks the MySQL server to shutdown. Parameter is one byte, optional. This packet can be sent with *mysqladmin shutdown*.

Bytes	Name
-----	----
1	shutdown option: 0x00...SHUTDOWN_DEFAULT 0x01...SHUTDOWN_WAIT_CONNECTIONS 0x02...SHUTDOWN_WAIT_TRANSACTIONS 0x08...SHUTDOWN_WAIT_UPDATES 0x10...SHUTDOWN_WAIT_ALL_BUFFERS 0x11...SHUTDOWN_WAIT_CRITICAL_BUFFERS 0xFE...KILL_QUERY 0xFF... KILL_CONNECTION defined in mysql_com.h

COM_STATISTICS

Asks the MySQL server to compile a text message with some server statistics (uptime, queries per second, etc.). This packet can be sent with *mysqladmin status*. No arguments.

COM_PROCESS_INFO

Functional equivalent to the SQL statement SHOW PROCESSLIST. This packet can be sent by *mysqladmin processlist*. No arguments.

COM_PROCESS_KILL

Functional equivalent to the SQL statement KILL <id>.

Bytes	Name
-----	----
4	Process ID (little endian)

COM_DEBUG

Asks the MySQL server to dump some debug information. The amount of data depends on compile time options (debug=no|yes|full). This packet can be sent with *mysqladmin debug*. No arguments.

COM_PING

This packet can be used to test the connection and to reset the connection inactivity counter in the MySQL server (wait_timeout). This packet can be sent with *mysqladmin ping*. Also exported by almost any client API. No arguments.

COM_CHANGE_USER

This packet is effectively a re-login without closing/opening the connection. Important side effect: this packet destroys the session context (temporary tables, session variables, etc.) in the MySQL server.

Some connection pool implementations use this to clean up the session context.

Bytes	Name
-----	----
n	user name (Null-terminated string)
n	password
	3.23 scramble - Null-terminated string (9 bytes)
	4.1 scramble - Length (1 byte) coded string (21 byte)
n	database name (Null-terminated string)
2	character set number (since 5.1.23?)

COM_BINLOG_DUMP

This request is the last request sent from slave to master when a replication connection is established. The master answers with a stream of response packets, each containing one binlog event. If the master goes down, it sends an EOF packet.

Bytes	Name
-----	----
4	binlog position to start at (little endian)
2	binlog flags (currently not used; always 0)
4	server_id of the slave (little endian)
n	binlog file name (optional)

If the binlog file name is not given, it defaults to the first binlog available on the master.

COM_TABLE_DUMP

This request is sent from slave to master for a LOAD TABLE ... FROM MASTER statement. This feature is marked deprecated. Do not use!

Bytes	Name
-----	----
n	schema name (length coded string)
n	table name (length coded string)

COM_REGISTER_SLAVE

If the *report_host* variable is set on the slave, it sends this packet when it establishes the replication connection.

Bytes	Name
-----	----
4	server_id on the slave (little endian)
n	report_host (length coded string)
n	report_user (length coded string)
n	report_password (length coded string)
2	report_port
4	rpl_recovery_rank

4	server_id on the master (always 0)
---	------------------------------------

The *rpl_recovery_rank* is a MySQL server variable that can be set, but is not yet used. In the future this will be used for replication failover.

COM_PREPARE

Prepare a SQL statement. This request is answered with a special OK packet (documented elsewhere), sending the statement handle. All the other request packets for prepared statements use this statement handle.

Bytes	Name
-----	----
n	query string with '?' place holders (up to end of packet, no termination character)

COM_EXECUTE

This is documented elsewhere.

COM_LONG_DATA

result of a call to *mysql_stmt_send_long_data()* to send a BLOB in pieces.

Bytes	Name
-----	----
4	Statement ID (little endian)
2	Parameter number (little endian)
n	payload (up to end of packet, no termination character)

COM_CLOSE_STMT

Destroy a prepared statement. The statement handle becomes invalid.

Bytes	Name
-----	----
4	Statement ID (little endian)

COM_RESET_STMT

Reset (empty) the parameter buffers for a prepared statement. Mostly used in connection with COM_LONG_DATA.

Bytes	Name
-----	----
4	Statement ID (little endian)

COM_SET_OPTION

The parameter is a 16-bit integer. There is an ENUM type *enum_mysql_set_option* defined in *mysql_com.h*:

- MYSQL_OPTION_MULTI_STATEMENTS_ON
- MYSQL_OPTION_MULTI_STATEMENTS_OFF

Bytes	Name
-----	----
2	option to be set (little endian)

COM_FETCH_STMT

Fetch result rows from a prepared statement. Can fetch a variable amount of rows.

Bytes	Name
-----	----
4	Statement ID (little endian)
4	number of rows to fetch (little endian)

Result Packet Types

A "result packet" is a packet that goes from the server to the client in response to a Client Authentication Packet or Command Packet. To distinguish between the types of result packets, a client must look at the first byte in the packet. We will call this byte "field_count" in the description of each individual package, although it goes by several names.

Type Of Result Packet	Hexadecimal Value Of First Byte (field_count)
-----	-----
OK Packet	00
Error Packet	ff
Result Set Packet	1-250 (first byte of Length-Coded Binary)
Field Packet	1-250 (")
Row Data Packet	1-250 (")
EOF Packet	fe

OK Packet

From server to client in response to command, if no error and no result set.

VERSION 4.0	
Bytes	Name
-----	----
1 (Length Coded Binary)	field_count, always = 0
1-9 (Length Coded Binary)	affected_rows
1-9 (Length Coded Binary)	insert_id
2	server_status
n (until end of packet)	message
VERSION 4.1	
Bytes	Name
-----	----
1 (Length Coded Binary)	field_count, always = 0
1-9 (Length Coded Binary)	affected_rows
1-9 (Length Coded Binary)	insert_id
2	server_status
2	warning_count
n (until end of packet)	message
field_count:	always = 0
affected_rows:	= number of rows affected by INSERT/UPDATE/DELETE
insert_id:	If the statement generated any AUTO_INCREMENT number, the number is returned here. Otherwise this field contains 0. Note: when using for example a multiple row INSERT the insert_id will be from the first row inserted, not from last.
server_status:	= The client can use this to check if the command was inside a transaction.
warning_count:	number of warnings
message:	For example, after a multi-line INSERT, message might be "Records: 3 Duplicates: 0 Warnings: 0"

The message field is optional.

Alternative terms: OK Packet is also known as "okay packet" or "ok packet" or "OK-Packet". field_count is also known as "number of rows" or "marker for ok packet". message is also known as "Messagetext". OK Packets (and result set packets) are also called "Result packets".

Relevant files in MySQL source:
 (client) sql/client.c mysql_read_query_result()
 (server) sql/protocol.cc send_ok()

Example OK Packet		
	Hexadecimal	ASCII
	-----	-----
field_count	00	.
affected_rows	01	.
insert_id	00	.
server_status	02 00	..
warning_count	00 00	..

In the example, the optional message field is missing (the client can determine this by examining the packet length). This is a packet that the server returns after a successful INSERT of a single row that contains no auto_increment columns.

Error Packet

From server to client in response to command, if error.

VERSION 4.0	
Bytes	Name
----	----
1	field_count, always = 0xff
2	errno (little endian)
n	message
VERSION 4.1	
Bytes	Name
----	----
1	field_count, always = 0xff
2	errno
1	(sqlstate marker), always '#'
5	sqlstate (5 characters)
n	message
field_count:	Always 0xff (255 decimal).
errno:	The possible values are listed in the manual, and in the MySQL source code file /include/mysqld_error.h.
sqlstate marker:	This is always '#'. It is necessary for distinguishing version-4.1 messages.
sqlstate:	The server translates errno values to sqlstate values with a function named mysql_errno_to_sqlstate(). The possible values are listed in the manual, and in the MySQL source code file /include/sql_state.h.
message:	The error message is a string which ends at the end of the packet, that is, its length can be determined from the packet header. The MySQL client (in the my_net_read() function) always adds '\0' to a packet, so the message may appear to be a Null-Terminated String. Expect the message to be between 0 and 512 bytes long.

Alternative terms: field_count is also known as "Status code" or "Error Packet marker". errno is also known as "Error Number" or "Error Code".

Relevant files in MySQL source: (client) client.c net_safe_read() (server) sql/protocol.cc send_error()

Example of Error Packet		
	Hexadecimal	ASCII
	-----	-----
field_count	ff	.
errno	1b 04	..
(sqlstate marker)	23	#
sqlstate	34 32 53 30 32	42S02
message	55 63 6b 6e 6f 77 6e 20	Unknown


```
74 61 62 6c 6c 65 20 27    table '
71 27                      q'
```

Note that some error messages past MySQL 4.1 are still returned without SQLState. For example, error 1043 'Bad handshake'.

Result Set Header Packet

From server to client after command, if no error and result set -- that is, if the command was a query which returned a result set.

The Result Set Header Packet is the first of several, possibly many, packets that the server sends for result sets. The order of packets for a result set is:

(Result Set Header Packet)	the number of columns
(Field Packets)	column descriptors
(EOF Packet)	marker: end of Field Packets
(Row Data Packets)	row contents
(EOF Packet)	marker: end of Data Packets

Bytes	Name
-----	----
1-9 (Length-Coded-Binary)	field_count
1-9 (Length-Coded-Binary)	extra
field_count: See the section "Types Of Result Packets" to see how one can distinguish the first byte of field_count from the first byte of an OK Packet, or other packet types.	
extra: For example, SHOW COLUMNS uses this to send the number of rows in the table.	

The "extra" field is optional and never appears for ordinary result sets.

Alternative terms: a Result Set Packet is also called "a result packet for a command returning rows" or "a field description packet".

Relevant MySQL source code:

```
libmysql/libmysql.c (client):
mysql_store_result() Read a result set from the server to memory
mysql_use_result()   Read a result set row by row from the server.
See also my_net_write() which describes local data loading.
```

Example of Result Set Header Packet	Hexadecimal	ASCII
	-----	-----
field_count	03	.

In the example, we see what the packet would contain after "SELECT * FROM t7" if table t7 has 3 columns.

Field Packet

From Server To Client, part of Result Set Packets. One for each column in the result set. Thus, if the value of field_columns in the Result Set Header Packet is 3, then the Field Packet occurs 3 times.

Bytes	Name
-----	----
n (Length Coded String)	table
n (Length Coded String)	name
4 (Length Coded Binary)	length
2 (Length Coded Binary)	type
2 (Length Coded Binary)	flags
1	decimals
n (Length Coded Binary)	default

```

VERSION 4.1
Bytes      Name
-----
n (Length Coded String)  catalog
n (Length Coded String)  db
n (Length Coded String)  table
n (Length Coded String)  org_table
n (Length Coded String)  name
n (Length Coded String)  org_name
1                      (filler)
2                      charsetnr
4                      length
1                      type
2                      flags
1                      decimals
2                      (filler), always 0x00
n (Length Coded Binary)  default

```

In practice, since identifiers are almost always 250 bytes or shorter, the Length Coded Strings look like: (1 byte for length of data) (data)

```

catalog:      Catalog. For 4.1, 5.0 and 5.1 the value is "def".
db:           Database identifier, also known as schema name.
table:        Table identifier, after AS clause (if any).
org_table:    Original table identifier, before AS clause (if any).
name:         Column identifier, after AS clause (if any).
org_name:     Column identifier, before AS clause (if any).
charsetnr:    Character set number.
length:       Length of column, according to the definition.
               Also known as "display length". The value given
               here may be larger than the actual length, for
               example an instance of a VARCHAR(2) column may
               have only 1 character in it.
type:         The code for the column's data type. Also known as
               "enum_field_type". The possible values at time of
               writing (taken from include/mysql_com.h), in hexadecimal:
0x00  FIELD_TYPE_DECIMAL
0x01  FIELD_TYPE_TINY
0x02  FIELD_TYPE_SHORT
0x03  FIELD_TYPE_LONG
0x04  FIELD_TYPE_FLOAT
0x05  FIELD_TYPE_DOUBLE
0x06  FIELD_TYPE_NULL
0x07  FIELD_TYPE_TIMESTAMP
0x08  FIELD_TYPE_LONGLONG
0x09  FIELD_TYPE_INT24
0x0a  FIELD_TYPE_DATE
0x0b  FIELD_TYPE_TIME
0x0c  FIELD_TYPE_DATETIME
0x0d  FIELD_TYPE_YEAR
0x0e  FIELD_TYPE_NEWDATE
0x0f  FIELD_TYPE_VARCHAR (new in MySQL 5.0)
0x10  FIELD_TYPE_BIT (new in MySQL 5.0)
0xf6  FIELD_TYPE_NEWDECIMAL (new in MySQL 5.0)
0xf7  FIELD_TYPE_ENUM
0xf8  FIELD_TYPE_SET
0xf9  FIELD_TYPE_TINY_BLOB
0xfa  FIELD_TYPE_MEDIUM_BLOB
0xfb  FIELD_TYPE_LONG_BLOB
0xfc  FIELD_TYPE_BLOB
0xfd  FIELD_TYPE_VAR_STRING
0xfe  FIELD_TYPE_STRING
0xff  FIELD_TYPE_GEOMETRY

flags:        The possible flag values at time of
               writing (taken from include/mysql_com.h), in hexadecimal:
0001  NOT_NULL_FLAG
0002  PRI_KEY_FLAG
0004  UNIQUE_KEY_FLAG
0008  MULTIPLE_KEY_FLAG
0010  BLOB_FLAG
0020  UNSIGNED_FLAG

```

```

0040 ZEROFILL_FLAG
0080 BINARY_FLAG
0100 ENUM_FLAG
0200 AUTO_INCREMENT_FLAG
0400 TIMESTAMP_FLAG
0800 SET_FLAG

decimals:      The number of positions after the decimal
                point if the type is DECIMAL or NUMERIC.
                Also known as "scale".
default:       For table definitions. Doesn't occur for
                normal result sets. See mysql_list_fields().

```

Alternative Terms: Field Packets are also called "Header Info Packets" or "field descriptor packets" (that's a better term but it's rarely used). In non-MySQL contexts Field Packets are more commonly known as "Result Set Metadata".

Relevant MySQL source code:
(client) client/client.c unpack_fields().
(server) sql/sql_base.cc send_fields().

Example of Field Packet	Hexadecimal	ASCII
	-----	-----
catalog	03 73 74 64	.std
db	03 64 62 31	.dbl
table	02 54 37	.T7
org_table	02 74 37	.t7
name	02 53 31	.S1
org_name	02 73 31	.s1
(filler)	0c	.
charsetnr	08 00	..
length	01 00 00 00
type	fe	.
flags	00 00	..
decimals	00	.
(filler)	00 00	..

In the example, we see what the server returns for "SELECT s1 AS S1 FROM t7 AS T7" where column s1 is defined as CHAR(1).

Row Data Packet

From server to client. One packet for each row in the result set.

Bytes	Name
-----	----
n (Length Coded String) (column value)	
...	
(column value):	The data in the column, as a character string. If a column is defined as non-character, the server converts the value into a character before sending it. Since the value is a Length Coded String, a NULL can be represented with a single byte containing 251(see the description of Length Coded Strings in section "Elements" above).

The (column value) fields occur multiple times. All (column value) fields are in one packet. There is no space between each (column value).

Alternative Terms: Row Data Packets are also called "Row Packets" or "Data Packets".

Relevant MySQL source code:
(client) client/client.c read_rows

Example of Row Data Packet	Hexadecimal	ASCII
	-----	-----
(first column)	01 58	.X

```
(second column)    02 35 35                .55
```

In the example, we see what the packet contains after a SELECT from a table defined as "(s1 CHAR, s2 INTEGER)" and containing one row where s1='X' and s2=55.

Row Data Packet: Binary (Tentative Description)

From server to client, or from client to server (if the client has a prepared statement, the "result set" packet format is used for transferring parameter descriptors and parameter data).

Recall that in the description of Row Data we said that: "If a column is defined as non-character, the server converts the value into a character before sending it." That doesn't have to be true. If it isn't true, it's a Row Data Packet: Binary.

Bytes	Name
-----	----
1	0 (packet header)
(col_count+7+2)/8	Null Bit Map with first two bits = 01
n	column values

Null Bit Map: The first 2 bits are reserved. Since there is always one bit on and one bit off, this can't be confused with the first byte of an Error Packet (255), the first byte of a Last Data Packet (254), or the first byte of an OK Packet (0).
NOTE: MySQL 5.x these 2 bits are always 0.

(column value): The column order and organization are the same as for conventional Row Data Packets. The difference is that each column value is sent just as it is stored. It's now up to the client to convert numbers to strings if that's desirable. For a description of column storage, see "Physical Attributes Of Columns" elsewhere in this document.

Only non-NULL parameters are passed.

Because no conversion takes place, fixed-length data items are as described in the "Physical Attributes of Columns" section: one byte for TINYINT, two bytes for FLOAT, four bytes for FLOAT, etc. Strings will appear as packed-string-length plus string value. DATETIME, DATE and TIME will be as follows:

Type	Size	Comment
----	----	-----
date	1 + 0-11	Length + 2 byte year, 1 byte MMDDHHMMSS, 4 byte billionth of a second
datetime	1 + 0-11	Length + 2 byte year, 1 byte MMDDHHMMSS, 4 byte billionth of a second
time	1 + 0-11	Length + sign (0 = pos, 1= neg), 4 byte days, 1 byte HHMMDD, 4 byte billionth of a second

If the sub-second part is 0, it isn't sent.
If the time-part is 00:00:00 too, it isn't sent either.
If all fields are 0, nothing is sent, but the length byte.

Alternative Terms: Row Data Packet: Binary is also called "Binary result set packet".

Except for the different way of signalling NULLs, the server/client parameter interaction here proceeds the same way that the server sends result set data to the client. Since the data is not sent as a string, the length and meaning depend on the data type. The client must make appropriate conversions given its knowledge of the data type.

EOF Packet

From Server To Client, at the end of a series of Field Packets, and at the end of a series of Data Packets. With prepared statements, EOF Packet can also end parameter information, which we'll describe later.

```
VERSION 4.0
```

Bytes	Name
----	----
1	field_count, always = 0xfe
VERSION 4.1	
Bytes	Name
----	----
1	field_count, always = 0xfe
2	warning_count
2	Status Flags
field_count:	The value is always 0xfe (decimal 254). However ... recall (from the section "Elements", above) that the value 254 can begin a Length-Encoded-Binary value which contains an 8-byte integer. So, to ensure that a packet is really an EOF Packet: (a) check that first byte in packet = 0xfe, (b) check that size of packet < 9.
warning_count:	Number of warnings. Sent after all data has been sent to the client.
server_status:	Contains flags like SERVER_MORE_RESULTS_EXISTS

Alternative terms: EOF Packet is also known as "Last Data Packet" or "End Packet".

Relevant MySQL source code:
(server) protocol.cc send_eof()

Example of EOF Packet

	Hexadecimal	ASCII
	-----	-----
field_count	fe	.
warning_count	00 00	..
server_status	00 00	..

OK for Prepared Statement Initialization Packet

From server to client, in response to prepared statement initialization packet.

It is made up of:

- a PREPARE_OK packet
- if "number of parameters" > 0
 - (field packets) as in a Result Set Header Packet
 - (EOF packet)
- if "number of columns" > 0
 - (field packets) as in a Result Set Header Packet
 - (EOF packet)

The PREPARE_OK packet is:

Bytes	Name
----	----
1	0 - marker for OK packet
4	statement_handler_id
2	number of columns in result set
2	number of parameters in query
1	filler (always 0)
2	warning count

Alternative terms: statement_handler_id is called "statement handle" or "hstmt" everywhere but at MySQL. Prepared statement initialization packet is also called "prepared statement init packet".

Parameter Packet (Tentative Description)

From server to client, for prepared statements which contain parameters.

The Parameter Packets follow a Prepared Statement Initialization Packet which has a positive value in the parameters field.

Bytes	Name
-----	----
2	type
2	flags
1	decimals
4	length
type:	Same as for type field in a Field Packet.
flags:	Same as for flags field in a Field Packet.
decimals:	Same as for decimals field in a Field Packet.
length:	Same as for length field in a Field Packet.

Notice the similarity to a Field Packet.

The parameter data will be sent in a packet with the same format as Row Data Packet: Binary.

Long Data Packet (Tentative Description)

From client to server, for long parameter values.

Bytes	Name
-----	----
4	statement_handler_id
2	parameter_number
2	type
n	data
statement_handler_id:	ID of statement handler
parameter_number:	Parameter number.
type:	Parameter data type. Not used at time of writing.
data:	Value of parameter, as binary string. The length of data is implicit from the packet length.

This is used by `mysql_send_long_data()` to set any parameter to a string value. One can call `mysql_send_long_data()` multiple times for the same parameter; The server will concatenate the results to one big string.

The server will not send an ok or error packet in response to this. If there is an error (for example the string is too big), one will see the error when calling "execute".

Relevant MySQL Source Code:
(server) `mysql_send_long_data`

Execute Packet (Tentative Description)

From client to server, to execute a prepared statement.

Bytes	Name
-----	----
1	code
4	statement_id
1	flags
4	iteration_count
if param_count > 0:	

```
(param_count+7)/8    null_bit_map
1                    new_parameter_bound_flag
    if new_params_bound == 1:
n*2                  type of parameters
n                    values for the parameters

code:                always COM_EXECUTE

statement_id:        statement identifier

flags:               reserved for future use. In MySQL 4.0, always 0.
                    In MySQL 5.0:
                        0: CURSOR_TYPE_NO_CURSOR
                        1: CURSOR_TYPE_READ_ONLY
                        2: CURSOR_TYPE_FOR_UPDATE
                        4: CURSOR_TYPE_SCROLLABLE

iteration_count:      reserved for future use. Currently always 1.

null_bit_map:        A bitmap indicating parameters that are NULL.
                    Bits are counted from LSB, using as many bytes
                    as necessary ((param_count+7)/8)
                    i.e. if the first parameter (parameter 0) is NULL, then
                    the least significant bit in the first byte will be 1.

new_parameter_bound_flag:  Contains 1 if this is the first time
                        that "execute" has been called, or if
                        the parameters have been rebound.

type:                Occurs once for each parameter;
                    The highest significant bit of this 16-bit value
                    encodes the unsigned property. The other 15 bits
                    are reserved for the type (only 8 currently used).
                    This block is sent when parameters have been rebound
                    or when a prepared statement is executed for the
                    first time.

values:              for all non-NULL values, each parameters appends its value
                    as described in Row Data Packet: Binary (column values)
```

The Execute Packet is also known as "COM_EXECUTE Packet".

In response to an Execute Packet, the server should send back one of: an OK Packet, an Error Packet, or a series of Result Set Packets in which all the Row Data Packets are binary.

Relevant MySQL Source Code: `libmysql/libmysql.c cli_read_prepare_result()`

Compression

This chapter does not discuss compression, but you should be aware of its existence.

Compression is of one or more logical packets. The `packet_number` field that is in each packet header is an aid for keeping track.

The opposite of "compressed" is "raw".

Compression is used if both client and server support zlib compression, and the client requests compression.

A compressed packet header is: packet length (3 bytes), packet number (1 byte), and Uncompressed Packet Length (3 bytes). The Uncompressed Packet Length is the number of bytes in the original, uncompressed packet. If this is zero then the data is not compressed.

When the compressed protocol is in use (that is, when the client has requested it by setting the flag bit in the Client Authentication Packet and the server has accepted it), either the client or the server may compress packets. However, compression will not occur if the compressed length is greater than the original length. Thus, some packets will be compressed while other packets are not compressed.

Encryption

If the server advertises SSL (*aka* TLS) capability by setting the CLIENT_SSL flag in the initial greeting packet, then the client can request that the connection be encrypted.

For an encrypted connection, the client first sends an *abbreviated* client authentication packet containing only the flags word, with the CLIENT_SSL bit set. Without waiting for a server response, the client then begins the [TLS handshake](#) (starting with the ClientHello message). Once the SSL/TLS session has been established, the client sends the full client authentication packet over the newly established channel (including the flags field again, as well as the remaining authentication fields). The server responds to this client authentication packet as usual.

Chapter 16. Stored Programs

Table of Contents

Overview	155
Persistent Representation	156
Stored Procedure and Stored Function Storage	156
Table Trigger Storage	156
Event Storage	157
Derived Attribute Storage	158
Internal Representation	158
Instructions	158
Parsing Context	159
Stored Program Parser	161
Parser Structure	161
Single-Pass Code Generation	163
Flow Analysis Optimizations	164
Dead Code Removal	165
Jump Shortcut Resolution	166
Stored Program Caches	167
Stored Procedure Cache	167
Stored Function Cache	168
Table Trigger Cache	168
Events and Caching	168
Stored Program Execution	169
Runtime Context	169
Executing One Instruction	171
Flow Control	171
Exception Handling	172
Call Nesting	173

Overview

Stored Programs in general refers to:

- [PROCEDURE](#)
- [FUNCTION](#)
- [TABLE TRIGGER](#)
- [EVENT](#) (Starting with 5.1)

When developing, there are a couple of tools available in the server itself that are helpful. These tools are only available in builds compiled with debugging support:

- [SHOW PROCEDURE CODE](#)
- [SHOW FUNCTION CODE](#)

The equivalent for triggers or events is not available at this point.

The internal implementation of Stored Programs in the server depends on several components:

- The storage layer, used to store in the database itself a program (hence the name *stored* program)
- The internal memory representation of a Stored Program, used within the server implementation

- The SQL parser, used to convert a Stored Program from its persistent representation to its internal form
- A flow analyser, used to optimize the code representing a stored program
- Various caches, used to improve performance by avoiding the need to load and parse a stored program at every invocation
- The Stored Program runtime execution itself, which interprets the code of the program and executes its statements

Persistent Representation

Storage of Stored Programs is implemented using either tables in the database (in the `mysql` schema), or physical files.

Stored Procedure and Stored Function Storage

The table `mysql.proc` contains one record per Stored Procedure or Stored Function. Note that this table design is a mix of relational and non relational (blob) content:

- Attributes that are part of the interface of a stored procedure or function (like its name, return type, etc), or that are global to the object (implementation language, deterministic properties, security properties, sql mode, etc) are stored with a dedicated column in the `mysql.proc` table.
- The body of a stored procedure or function, which consists of the original code expressed in SQL, including user comments if any, is stored as-is preserving the original indentation in blob column 'body'.

This design choice allows the various attributes to be represented in a format that is easy to work with (relational model), while allowing a lot of flexibility for the content of the body.

A minor exception to this is the storage of the parameters of a stored procedure or function (which are part of its interface) inside the blob column `param_list` (instead of using a child `proc_param` table).

Table `mysql.procs_priv` describes privileges granted for a given Stored Procedure or Stored Function in `mysql.proc`.

The code used to encapsulate database access is:

- `db_create_routine()`
- `db_load_routine()`
- `db_drop_routine()`
- `mysql_routine_grant()`
- `grant_load()`
- `grant_reload()`

Table Trigger Storage

Trigger definitions are stored in plain text files in the directory that contains the schema objects.

The file `<schema>/<trigger>.TRN` is the `TRIGGERNAME` file. It represents the fact that the object named "trigger" is a table trigger, and points to the table the trigger is attached to. Every trigger has a dedicated `*.TRN` file. This design decision is used to facilitate operating system filesystem services to enforce the SQL standard requirement that all triggers in a given schema must be unique.

The file `<schema>/<table>.TRG` is the `TRIGGERS` file. It represents all the table triggers attached to a given table, so this file can contain triggers for multiple events (`BEFORE/AFTER`, `INSERT/UPDATE/DELETE`). Currently it is impossible to have more than one trigger per table for a given trigger action time and type, hence this file may contain at most six trigger definitions.

The code used to encapsulate file access is:

- `Table_triggers_list::create_trigger()`
- `Table_triggers_list::drop_trigger()`
- `Table_triggers_list::check_n_load()`
- `Table_triggers_list::drop_all_triggers()`
- `Table_triggers_list::change_table_name()`
- See the C++ class `Table_triggers_list` in general.

Using files for triggers is due to historical reasons, and follows the same design as `*.FRM` files for table metadata. This approach has several drawbacks:

- Each file has yet another text file format, which is necessary to print and parse back correctly. Custom code has to be implemented, which is consuming in terms of resources, and introduces technical risk or code duplication.
- Tables are replicated, values in columns are checked for data validity, integrity constraints can be defined ... where none of the above is available with a file based implementation.
- With tables, the default locking, transaction and isolation mechanism used by the server in general can be leveraged, but the same is not available with files.
- Cluster support for any new metadata operation that operates on files will require a custom solution. E.g. to propagate `CREATE TABLE` statement across MySQL Cluster `mysqld` nodes we use a so-called “.FRM shipping” technique. There is no similar solution implemented for triggers at this point, and thus a trigger created in one `mysqld` node does not automatically become visible on other nodes. Potentially, if data is stored in tables, cluster support may be added as simply as by issuing `ALTER TABLE mysql.triggers ENGINE=NDB;`

Note

Various drawbacks of filesystem based solution are provided in this chapter only for a sake of example. Other advantages and disadvantages of two approaches may be found in relevant worklog entries and design documents.

Warning

The current implementation of the storage layer for table triggers is considered *private* to the server, and might change without warnings in future releases.

Event Storage

Events storage is very similar to Stored Procedure and Stored Function storage, and shares the same design. Since more attributes are needed to represent an event, a different table is used: `mysql.event`.

The code used to encapsulate the database access is:

- `Event_db_repository::create_event()`
- `Event_db_repository::update_event()`
- `Event_db_repository::drop_event()`

- See the C++ class `Event_db_repository` in general.

Derived Attribute Storage

Some critical attributes, such as `SQL_MODE`, are explicitly part of the storage format.

Other attributes, that also impact significantly the behavior in general of Stored Programs, can be *implicitly derived* from other properties of the storage layer. In particular:

- The `USE <database>` in effect for a stored program is the `schema` the stored object belongs to.
- The statement `DECLARE v CHAR(10)` does not intrinsically convey any notion of character set or collation. The character set and collation of this local variable, in a stored program, derives from the character set and collation of the `schema` the stored object belongs to.

Internal Representation

A Stored Program is represented in memory by two major parts:

- The code of the stored program, including SQL statements and control flow logic (`IF`, `WHILE`, ...),
- A symbol table that describes all the local variables, cursors, labels, conditions ... declared in the code.

Individual instructions of various kind are implemented by all the C++ classes that inherit from class `sp_instr`. The symbol table ('symbol table' is a term used in conjunction with compilers or interpreters, in `MySQL` the term 'Parsing Context' is used instead) is implemented by the C++ class `sp_pcontext`. A Stored Program as a whole is represented by the C++ class `sp_head`, which contains the instructions (array `m_instr`) and the root parsing context (member `m_pcont`).

Caution

Class `sp_head` contains concepts from different areas: it represents both what a stored program *is*, which is the topic of this section, and how a stored program logic *is used* during runtime interpretation, which is the subject of other sections.

Instructions

Data Definition Language and Data Manipulation Language SQL statements are represented as-is, by a single instruction. For flow control statements and exception handlers, several instructions are used to implement in the low level `sp_instr` language the semantic of the SQL construct.

Let's see an example with a stored procedure:

```
delimiter $$

CREATE PROCEDURE proc_1(x int)
BEGIN
  IF x < 0 THEN
    INSERT INTO t1 VALUES ("negative");
  ELSEIF x = 0 THEN
    INSERT INTO t1 VALUES ("zero");
  ELSE
    INSERT INTO t1 VALUES ("positive");
  END IF;
END$$
```

The resulting code, displayed by `SHOW PROCEDURE CODE`, is:

```
SHOW PROCEDURE CODE proc_1;
Pos      Instruction
0        jump_if_not 3(7) (x@0 < 0)
1        stmt 5 "INSERT INTO t1 VALUES ("negative")"
2        jump 7
```

```

3      jump_if_not 6(7) (x@0 = 0)
4      stmt 5 "INSERT INTO t1 VALUES ('zero')"
5      jump 7
6      stmt 5 "INSERT INTO t1 VALUES ('positive')"

```

Instructions are numbered sequentially. Position 0 is the start of the code. The position 7 that is one past the last instruction in this example represents the end of the code.

Note that the instruction `jump_if_not 3(7)` at position 0 can actually jump to *three* locations:

- When the evaluation of the condition " $x < 0$ " is true, the next instruction will be position 1 (the "then" branch),
- When the evaluation of the condition " $x < 0$ " is false, the next instruction will be position 3 (the "else" branch),
- When the evaluation of the condition " $x < 0$ " results in an error, and when a continue handler exists for the error, the next instruction will be position 7, known as the "continuation" destination.

Now, let's see how exception handlers are represented. The following code contains just a very basic handler, protecting a `BEGIN/END` block in the SQL logic:

```

CREATE PROCEDURE proc_2(x int)
BEGIN
    SELECT "Start";

    INSERT INTO t1 VALUES (1);

    BEGIN
        DECLARE CONTINUE HANDLER FOR SQLEXCEPTION
        BEGIN
            SELECT "Oops";
        END;

        INSERT INTO t1 VALUES (2);
        INSERT INTO t1 VALUES (2);
    END;

    INSERT INTO t1 VALUES (3);
    SELECT "Finish";
END$$

```

The internal instructions for this stored procedure are:

```

SHOW PROCEDURE CODE proc_2;
Pos      Instruction
0        stmt 0 "SELECT 'Start'"
1        stmt 5 "INSERT INTO t1 VALUES (1)"
2        hpush_jump 5 1 CONTINUE
3        stmt 0 "SELECT 'Oops'"
4        hreturn 1
5        stmt 5 "INSERT INTO t1 VALUES (2)"
6        stmt 5 "INSERT INTO t1 VALUES (2)"
7        hpop 1
8        stmt 5 "INSERT INTO t1 VALUES (3)"
9        stmt 0 "SELECT 'Finish'"

```

Note the flow of control in the code: there is not a single if. The couple of `hpush_jump` / `hpop` represent the installation and the removal of the exception handler. The body of the exception handler starts at position 3, whereas the code protected by the handler starts at position 5. `hpush_jump 5 1` means: add a handler for "1" condition (`sqlexception`), where "1" stands for the index of declared conditions in the parsing context, and execute the code starting at position "5".

Parsing Context

A parsing context is a tree of nodes, where each node contains symbols (variables, cursors, labels, ...) declared locally in the same name visibility scope.

For example, with the following SQL code:

```
CREATE PROCEDURE proc_3(x int, y int)
BEGIN
  -- This is the root parsing context
  DECLARE v1 INT;
  DECLARE v2 INT;
  DECLARE v3 INT;

  IF (x > 0) THEN
    BEGIN
      -- This is the child context A
      DECLARE v1 INT;
      DECLARE v4 INT DEFAULT 100;

      set v4:= 1;
      set v1:= x;
    END;
  ELSE
    BEGIN
      -- This is the child context B
      DECLARE v2 INT;
      DECLARE v4 INT DEFAULT 200;

      set v4:= 2;
      set v2:= y;
      set v3:= 3;
    END;
  END IF;

  set v1 := 4;
END$$
```

The parsing contexts match exactly the nesting of **BEGIN/END** blocks:

- The root parsing context contains parameters **x**, **y**, and local variables **v1**, **v2**, **v3**,
- The **BEGIN/END** block in the **THEN** part defines a child parsing context (let's call it 'A'), that contains local variables **v1** and **v4**,
- Likewise, the **ELSE** block defines a parsing context (let's call it 'B') which is a child of the root, and contains local variables **v2** and **v4**.

The total number of symbols is 9: 5 for the root + 2 for A + 2 for B. All the symbols are numbered internally (starting at offset 0), by walking the parsing context tree in a depth first manner, resulting in the following:

- Root:x --> 0, Root:y --> 1, Root:v1 --> 2, Root:v2 --> 3, Root:v3 --> 4,
- A:v1 --> 5, A:v4 --> 6,
- B:v2 --> 7, B:v4 --> 8,

There is no tool to dump the parsing context tree explicitly. However, the internal numbering of symbols is apparent when printing the code:

```
SHOW PROCEDURE CODE proc_3;
Pos      Instruction
0        set v1@2 NULL
1        set v2@3 NULL
2        set v3@4 NULL
3        jump_if_not 9(14) (x@0 > 0)
4        set v1@5 NULL
5        set v4@6 100
6        set v4@6 1
7        set v1@5 x@0
8        jump 14
9        set v2@7 NULL
```

```

10      set v4@8 200
11      set v4@8 2
12      set v2@7 y@1
13      set v3@4 3
14      set v1@2 4

```

The points of interest are that:

- There are *two* variables named `v1`, where the variable `v1` from block A (represented as `v1@5`) eclipses the variable `v1` from the root block (represented as `v1@2`).
- There are *two* variables named `v4`, which are independent. The variable `v4` from block A is represented as `v4@6`, whereas the variable `v4` from block B is represented as `v4@8`.

The parsing context C++ class, `sp_pcontext`, contains much more information related to each symbol, notably *data types* of variables (unfortunately not printable with `SHOW PROCEDURE CODE`).

Stored Program Parser

There is no “Stored Program Parser” as such, there is only one parser in the SQL layer in the server. This parser is capable of understanding every SQL statement, including statements related to Stored Programs. The parser is implemented as an ascendant parser, using bison. The source code is located in the file `sql/sql_yacc.yy`.

The parts of the parser dedicated more specially to Stored Programs are starting at the following rules:

- `CREATE PROCEDURE` : see rule `sp_tail`,
- `CREATE FUNCTION` : see rule `sp_tail`,
- `CREATE TRIGGER` : see rule `trigger_tail`,
- `CREATE EVENT` : see rule `event_tail`.

In every case, the parser reads the SQL text stream that represents the code as input, and creates an internal representation of the Stored Program as output, with one C++ object of type `sp_head`. A limiting consequence of this approach is that a stored program does not support nesting: it is impossible to embed one `CREATE PROCEDURE` into another, since the parser currently may only support one `sp_head` object at a time.

Parser Structure

Conceptually, there are many different layers involved during parsing:

- Lexical analysis (making words or tokens from a character stream),
- Syntactic analysis (making "sentences" or an abstract syntax tree from the tokens),
- Semantic analysis (making sure these sentences do make sense),
- Code generation (for compilers) or evaluation (for interpreters).

From the implementation point of view, many different concepts from different layers actually collide in the same code base, so that the actual code organization is as follows:

- The lexical analysis is implemented in `sql/sql_lex.cc`, as when parsing regular statements.
- Syntactic analysis, semantic analysis, and code generation -- *all of them* -- are done at *once*, during parsing of the code. From that perspective, the parser behaves as a single pass compiler. In other words, both the code and the symbol table for the final result are generated on the fly, interleaved with syntactic analysis.

This is both very efficient from a performance point of view, but difficult to understand, from a maintenance point of view.

Let's illustrate for example how the following SQL statement is parsed:

```
DECLARE my_cursor CURSOR FOR SELECT coll FROM t1;
```

The corresponding part of the grammar in the parser for DECLARE CURSOR statements is the following (with annotated line numbers):

```
[ 1] sp_decl:
[ 2]   DECLARE_SYM ident CURSOR_SYM FOR_SYM sp_cursor_stmt
[ 3]   {
[ 4]       LEX *lex= Lex;
[ 5]       sp_head *sp= lex->sphead;
[ 6]       sp_pcontext *ctx= lex->spcont;
[ 7]       uint offp;
[ 8]       sp_instr_cpush *i;
[ 9]
[10]       if (ctx->find_cursor(&$2, &offp, TRUE))
[11]       {
[12]           my_error(ER_SP_DUP_CURS, MYF(0), $2.str);
[13]           delete $5;
[14]           MYSQL_YYABORT;
[15]       }
[16]       i= new sp_instr_cpush(sp->instructions(), ctx, $5,
[17]                             ctx->current_cursor_count());
[18]       sp->add_instr(i);
[19]       ctx->push_cursor(&$2);
[20]       $$vars= $$conds= $$hdlrs= 0;
[21]       $$curs= 1;
[22]   }
[23] ;
```

The lines [1], [2] and [23] are bison code that express the structure of the grammar. These lines belong to the syntactic parsing realm.

The lines [3] and [22] are bison delimiters for the associated action to execute, when parsing of the syntax succeeds. Everything between lines [3] and [22] is C++ code, executed when the parser finds a syntactically correct `DECLARE CURSOR` statement.

The lines [4] to [8] could be considered syntactic parsing: what the code does is find what is the current Stored Program being parsed, find the associated part of the syntax tree under construction (`sp_head`), and find the associated current context in the symbol table (`sp_pcontext`).

Note that there is some black magic here: since we are still currently parsing the *content* of a Stored Program (the `DECLARE CURSOR` statement), the final “syntax” tree for the Stored Program (`sp_head`) is not supposed to exist yet. The reason the `sp_head` object is already available is that the actions in the `CREATE PROCEDURE`, `CREATE FUNCTION`, `CREATE TRIGGER`, or `CREATE EVENT` are implemented as a *descendant* parser (it created an empty `sp_head` object first, filling the content later). Mixing code that way (descendant actions with ascendant parsing) is extremely sensitive to changes.

The line [10] is a semantic check. The statement might be syntactically correct (it parsed), but to be semantically correct, the *name* or the cursor must be unique in the symbol table.

Line [12] is reporting a semantic error back to the client (duplicate cursor). The code at line [14] forces the syntactic parser (bison) to abort.

By line [16], we have verified that the code is syntactically valid, and semantically valid: it's now time for code generation, implemented by creating a new `sp_instr_cpush` to represent the cursor in the compiled code. Note that variable allocation is done on the fly, by looking up the current cursor count in the symbol table (`sp_pcontext::current_cursor_count()`).

Line [18] adds the generated code to the object representing the stored program (code generation).

Line [19] maintains the symbol table (semantic parsing) by adding the new cursor in the current local context.

Lines [20] and [21] return to bison a fragment of a fake syntax tree, indicating that one cursor was found.

By looking at the complete implementation of this action in bison, one should note that the target code was generated, the symbol table for the Stored Program was looked up and updated, while at no point in time a syntax node was even created. Note that the `sp_instr_cpush` object should really be considered generated code: the fact that there is a one-to-one correspondence with the syntax is incidental.

Single-Pass Code Generation

All the code generated by the parser is emitted in a *single pass*. For example, consider the following SQL logic:

```
CREATE FUNCTION func_4(i int)
RETURNS CHAR(10)
BEGIN
    DECLARE str CHAR(10);

    CASE i
        WHEN 1 THEN SET str="1";
        WHEN 2 THEN SET str="2";
        WHEN 3 THEN SET str="3";
        ELSE SET str="unknown";
    END CASE;

    RETURN str;
END$$
```

The compiled program for this Stored Function is:

```
SHOW FUNCTION CODE func_4;
Pos      Instruction
0        set str@1 NULL
1        set_case_expr (12) 0 i@0
2        jump_if_not 5(12) (case_expr@0 = 1)
3        set str@1 _latin1'1'
4        jump 12
5        jump_if_not 8(12) (case_expr@0 = 2)
6        set str@1 _latin1'2'
7        jump 12
8        jump_if_not 11(12) (case_expr@0 = 3)
9        set str@1 _latin1'3'
10       jump 12
11       set str@1 _latin1'unknown'
12       freturn 254 str@1
```

Note the instruction at position 4: `jump 12`. How can the compiler generate this instruction in a single pass, when the destination (12) is not known yet? This instruction is a *forward jump*. What happens during code generation is that, by the time the compiler has generated the code for positions [0] to [11], the generated code looks like this:

```
Pos      Instruction
0        set str@1 NULL
1        set_case_expr ( ?? ) 0 i@0
2        jump_if_not 5( ?? ) (case_expr@0 = 1)
3        set str@1 _latin1'1'
4        jump ??
5        jump_if_not 8( ?? ) (case_expr@0 = 2)
6        set str@1 _latin1'2'
7        jump ??
8        jump_if_not 11( ?? ) (case_expr@0 = 3)
9        set str@1 _latin1'3'
10       jump ??
```

```
11      set str@1 _latin1'unknown'
...
```

The final destination of the label for the `END CASE` is not known yet, and the list of all the instructions (1, 2, 4, 5, 7, 8 and 10) that need to point to this unknown destination (represented as `??`) is maintained in a temporary structure used during code generation only. This structure is called the context back patch list.

When the destination label is finally resolved to a destination (12), all the instructions pointing to that label, which have been already generated (but with a bogus destination) are *back patched* to point to the correct location. See the comments marked `BACKPATCH` in the code for more details.

As a side note, this generated code also shows that some temporary variables can be generated implicitly, such as the operand of the `CASE` expression, labeled `case_expr@0`.

Caution

Numbering of case expressions in the symbol table uses a different namespace than variables, so that `case_expr@0` and `i@0` are two different variables, even when both internally numbered with offset zero.

Flow Analysis Optimizations

After code is generated, the low level `sp_instr` instructions are optimized. The optimization focuses on two areas:

- Dead code removal,
- Jump shortcut resolution.

These two optimizations are performed together, as they both are a problem involving *flow analysis* in the *graph* that represents the generated code.

The code that implements these optimizations is `sp_head::optimize()`.

Caution

Do not confuse `sp_head::optimize()` with the component named the *optimizer*, as they are very different. The former is specific to Stored Programs, and focuses on improving the flow of statements, whereas the latter is general to queries, and focuses on finding the best execution plan when executing a single statement. For the optimizer, see <http://dev.mysql.com/doc/en/optimization.html>.

The (Stored Program) optimizer is invoked from only one place, in the following code:

```
db_load_routine(..., sp_head **sphp, ...)
{
    ...
    (*sphp)->optimize();
    ...
}
```

Tip

By disabling the call to `sp_head::optimize()` and recompiling the code, `SHOW PROCEDURE CODE` will display the code *before* flow optimization.

Caution

When investigating issues related to this area, you may want to use a `DEBUG_EXECUTE_IF` to avoid recompiling the server with or without flow

optimization every time. Be careful to shutdown and restart the server with or without the call to `sp_head::optimize()` for each test, or you will find that *caching* of a Stored Program code does interfere.

Dead Code Removal

Dead code is also known as *unreachable code*: code that cannot possibly be executed, because no path in the logic leads to it.

For example, consider the following SQL code:

```
CREATE PROCEDURE proc_5()
BEGIN
  DECLARE i INT DEFAULT 0;

  again:
  WHILE TRUE DO
    BEGIN
      set i:= i+1;

      SELECT "This code is alive";

      IF (i = 100) THEN
        LEAVE again;
      END IF;

      ITERATE again;

      SELECT "This code is dead";
    END;
  END WHILE;
END$$
```

Before flow optimization, the compiled code is:

```
SHOW PROCEDURE CODE proc_5;
Pos      Instruction
0        set i@0 0
1        jump_if_not 10(10) 1
2        set i@0 (i@0 + 1)
3        stmt 0 "SELECT "This code is alive""
4        jump_if_not 7(7) (i@0 = 100)
5        jump 10
6        jump 7
7        jump 1
8        stmt 0 "SELECT "This code is dead""
9        jump 1
```

Note the instruction at position 8: the previous instruction is an *unconditional* jump, so the flow of control can never reach 8 by coming from 7. Because there exists no jump in the entire code that leads to 8 either, the instruction at 8 is unreachable. By looking further in the flow, because 8 is unreachable and there are no jumps to position 9, the instruction at position 9 is also unreachable.

The instruction at position 6 is also unreachable, for a similar reason: the `THEN` part of the if contains a jump, due to the statement `LEAVE again;`, so that the code never executes the jump generated by the compiler to go from the end of the `THEN` block to the statement following the `IF`.

After detecting all the unreachable instructions, and simplifying the code, the result after flow optimization is:

```
SHOW PROCEDURE CODE proc_5;
Pos      Instruction
0        set i@0 0
1        jump_if_not 10(10) 1
2        set i@0 (i@0 + 1)
3        stmt 0 "SELECT "This code is alive""
4        jump_if_not 1(1) (i@0 = 100)
5        jump 10
```

The flow optimizer is good at detecting *most* of the dead code, but has limitations. For example, coding in SQL `IF FALSE THEN ... END IF;` leads to code that can never be executed, but since the flow optimizer does neither propagate constants nor consider impossible conditional jumps, this code will not be removed.

The goal of the flow optimizer is mostly to perform *simple* local optimizations with a low cost. It's not a fully featured code optimizer, and does not guard against poor SQL.

Jump Shortcut Resolution

The term *jump shortcut* refers to the following optimization: when instruction A is a jump (conditional or not) that goes to position B, and when B is an unconditional jump to position C, the code can be changed so that A can jump to C directly, taking a *shortcut* to avoid the unnecessary B. Consider the following SQL code:

```
CREATE PROCEDURE proc_6(x int, y int, z int)
BEGIN
  SELECT "Start";

  IF (x > 0)
  THEN
    BEGIN
      SELECT "x looks ok";
      IF (y > 0)
      THEN
        BEGIN
          SELECT "so does y";
          IF (z > 0)
          THEN
            SELECT "even z is fine";
          ELSE
            SELECT "bad z";
          END IF;
        END;
      ELSE
        SELECT "bad y";
      END IF;
    END;
  ELSE
    SELECT "bad x";
  END IF;

  SELECT "Finish";
END$$
```

Before flow optimization, the compiled code is:

```
SHOW PROCEDURE CODE proc_6;
Pos      Instruction
0        stmt 0 "SELECT "Start""
1        jump_if_not 12(13) (x@0 > 0)
2        stmt 0 "SELECT "x looks ok""
3        jump_if_not 10(11) (y@1 > 0)
4        stmt 0 "SELECT "so does y""
5        jump_if_not 8(9) (z@2 > 0)
6        stmt 0 "SELECT "even z is fine""
7        jump 9
8        stmt 0 "SELECT "bad z""
9        jump 11
10       stmt 0 "SELECT "bad y""
11       jump 13
12       stmt 0 "SELECT "bad x""
13       stmt 0 "SELECT "Finish""
```

Note the `jump 9` at position 7: since the instruction at position 9 is `jump 11`, the code at position 7 can be simplified to `jump 11`. The optimization is also recursive: since the instruction 11 is `jump 13`, the final jump destination for the instruction at position 7 is `jump 13`. Conditional jumps are optimized also, so that the instruction 5: `jump_if_not 8(9)` can be optimized to `jump_if_not 8(13)`.

After flow optimization, the compiled code is:

```
SHOW PROCEDURE CODE proc_6;
Pos      Instruction
0        stmt 0 "SELECT "Start""
1        jump_if_not 12(13) (x@0 > 0)
2        stmt 0 "SELECT "x looks ok""
3        jump_if_not 10(13) (y@1 > 0)
4        stmt 0 "SELECT "so does y""
5        jump_if_not 8(13) (z@2 > 0)
6        stmt 0 "SELECT "even z is fine""
7        jump 13
8        stmt 0 "SELECT "bad z""
9        jump 13
10       stmt 0 "SELECT "bad y""
11       jump 13
12       stmt 0 "SELECT "bad x""
13       stmt 0 "SELECT "Finish""
```

Note the differences with every jump instruction.

Caution

For clarity, this example has been designed to *not* involve dead code. Note that in general, an instruction that was reachable *before* taking a shortcut might become unreachable *after* the shortcut, so that the optimizations for jump shortcuts and dead code are tightly intertwined.

Stored Program Caches

The *goal* of the Stored Program cache is to keep a parsed `sp_head` in memory, for future reuse. Reuse means:

- To be able to execute concurrently the same Stored Program in different `THD` threads,
- To be able to execute the same Stored Program multiple times (for recursive calls) in the same `THD` thread.

To achieve this, the implementation of `sp_head` must be both thread-safe and stateless. Unfortunately, it is neither:

- The class `sp_head` is composed of `sp_instr` instructions to represent the code, and these instructions in turn depend on `Item` objects, used to represent the internal structure of a statement. The various C++ `Item` classes are *not* currently thread-safe, since the evaluation of an `Item` at runtime involves methods like `Item::fix_fields()`, which modify the internal state of items, making them impossible to safely evaluate concurrently.
- The class `sp_head` itself contains attributes that describe the *SQL logic* of a Stored Program (which are safe to share), mixed with attributes that relate to the *evaluation* of this logic in a given instance to a Stored Program call (mostly the `MEM_ROOT` memory pool used during execution), which by definition cannot be shared.

The consequence of these restrictions is less than optimal code. What is *currently* implemented in the server is detailed in the following subsections, to help maintenance.

Warning

Needless to say, the current implementation of Stored Program caching is by no mean final, and could be re factored in future releases.

Stored Procedure Cache

The `PROCEDURE` cache is maintained on a *per thread* basis, in `THD::sp_proc_cache`.

The function used to lookup the cache is `sp_find_routine`. It relies on the C++ class `sp_cache` for the low level implementation.

There is a global mechanism to invalidate all the caches of all the `THD` threads at once, implemented with the variable `Cversion` in file `sp_cache.cc`, which is incremented by function `sp_cache_invalidate()`. This global invalidation is used when the server executes `DROP PROCEDURE` or `UPDATE PROCEDURE` statements.

Each entry in the cache is keyed by name, and consists of a linked list of stored procedure instances which are all duplicates of the same object. The reason for the list is recursion, when the runtime needs to evaluate several calls to the same procedure at once.

The runtime behavior of this caching mechanism has some limitations, and in particular:

- *each* `THD` has its own cache, so *each* separate client connection to the server uses its own cache. Multiple client connections calling the same Stored Procedure will cause the parser to be invoked multiple times, and memory to be consumed multiple times.
- If a given client constantly opens and closes a new connection to the server, and invokes Stored Procedures, the cache will be always empty, causing excessive parsing of used stored procedures on every invocation.
- If a given client constantly keeps an existing connection to the server for a long time, and invokes Stored Procedures, the cache size will grow, consuming and retaining memory. In other words, memory limits or expulsion of cold members of the stored procedure cache is not implemented.
- Calling `sp_cache_invalidate()` does *not* reclaim the cache memory. This memory will be reclaimed only *if* a Stored Procedure is looked up in the cache again, causing the cache to flush.

Stored Function Cache

The `FUNCTION` cache is implemented exactly like the `PROCEDURE` cache, in the thread member in `THD::sp_func_cache`.

Note that because `THD::sp_proc_cache` and `THD::sp_func_cache` are both invalidated based on the *same* `Cversion` counter, executing `DROP PROCEDURE` happens to invalidate the `FUNCTION` cache as well, while `DROP FUNCTION` also invalidates the `PROCEDURE` cache. In practice, this has no consequences since DDL statements like this are not executed typically while an application is running, only when it is deployed.

Table Trigger Cache

For table triggers, all the triggers that relate to a given table are grouped in the C++ class `Table_triggers_list`, which in particular contains the member `sp_head *bodies[TRG_EVENT_MAX][TRG_ACTION_MAX]`.

Note that at most one trigger per event (`BEFORE`, `AFTER`) and per action (`INSERT`, `UPDATE`, `DELETE`) can be defined currently.

The `Table_triggers_list` itself is a part of the structure `struct st_table`, which is better known as a `TABLE` handle.

As a result, each table trigger body is duplicated in each table handle, which is necessary to properly evaluate them. `TABLE` handles are globally cached and reused across threads, so the table triggers are effectively reused across different clients connections manipulating the same physical table.

Events and Caching

For events, the `sp_head` object that represents the body of an `EVENT` is part of the C++ class `Event_parse_data`.

There is no caching of `sp_head` for multiple scheduling of an event. The method `Event_job_data::execute()` invokes the parser every time an event is executed.

Stored Program Execution

Executing a Stored Program consists of *interpreting* the low level `sp_instr` code. The runtime interpreter itself is implemented in the method `sp_head::execute()`. Wrappers for different kinds of Stored Programs are implemented in the following methods:

- `PROCEDURE` : see `sp_head::execute_procedure()`,
- `FUNCTION` : see `sp_head::execute_function()`,
- `TRIGGER` : see `sp_head::execute_trigger()`,
- `EVENT` : see `Event_job_data::execute()`.

Runtime Context

An interpreter needs to be able to represent the *state* of the SQL program being executed: this is the role of the C++ class `sp_rcontext`, or runtime context.

Local Variables

Values of local variables in an SQL Stored Program are stored within the `sp_rcontext`. When the code enters a new scope, the `sp_instr` contains explicit statements to initialize the local variable `DEFAULT` value, if any. Since initialization of values is done in the *code*, and since no logic needs to be executed when an SQL variable goes out of scope, space allocation to represent the *data* does not need to follow the nesting of `BEGIN/END` blocks during runtime.

Another important point regarding the representation of local SQL variables is that, conceptually, a local variable can be considered to be an SQL table with a single column (of the variable type), with a single row (to represent the value).

As a result, *all* the local variables of a Stored Program are represented by a row in a table internally. For example, consider the following SQL code:

```
CREATE PROCEDURE proc_7(x int)
BEGIN
  DECLARE v1 INT;
  DECLARE v2 VARCHAR(10);
  DECLARE v3 TEXT;

  IF (x > 0) THEN
    BEGIN
      DECLARE v4 BLOB;
      DECLARE v5 VARCHAR(20);
    END;
  ELSE
    BEGIN
      DECLARE v6 DECIMAL(10, 2);
      DECLARE v7 BIGINT;
    END;
  END IF;
END$$
```

Internally, a temporary table is created, with the following structure:

```
CREATE TEMPORARY TABLE `proc_7_vars` (
  `v1` int(11) DEFAULT NULL,
  `v2` varchar(10) DEFAULT NULL,
  `v3` text,
  `v4` blob,
  `v5` varchar(20) DEFAULT NULL,
```

```
`v6` decimal(10,2) DEFAULT NULL,
`v7` bigint(20) DEFAULT NULL
) ENGINE=MyISAM DEFAULT CHARSET=latin1
```

The real name of the table and the columns are purely internal, and the table is not accessible to regular statements for DDL or DML operations: `proc_7_vars` and `v1 ... v7` are just a notation used in this example. The `TABLE` handle that implements all the local variable storage is the member `sp_rcontext::m_var_table`

Inside a statement, local variables in a Stored Program are represented by the dedicated C++ class `Item_splocal`. `Item_splocal` really is a proxy exposing the interface needed to support `Item`, which delegates to the underlying `sp_rcontext` for reading or writing local variable values. The coupling between `Item_splocal` and `sp_rcontext` is based on `Item_splocal::m_var_idx`, which is the variable *index* in the symbol table computed by the *parser*, and maintained in `sp_pcontext`.

Cursors

Unlike local variables, some action is needed in the interpreter when a `CURSOR` goes out of scope: the cursor must be closed, to prevent leaks of the underlying `TABLE` resources.

As a result, cursor allocation (and really, *deallocation* so they can be properly *closed*) needs to follow tightly the BEGIN-END block structure of the code, so a *stack* is used, implemented by `sp_rcontext::m_cstack` and `sp_rcontext::m_ccount`.

Case Expressions

For `CASE` expressions, temporary variables are generated automatically. Like `CURSOR`, there are some constraints that prevent treating these special local variables like regular local variables.

The difficulty with `CASE` is that the real type of the expression is only known when the case statement is executed, so that allocating space in a statically computed `TABLE` is not practical. For example, `CASE (SELECT col1 FROM t1 WHERE ...)` is a case expression that involves a single row subselect. During parsing, the table might not even exist, so evaluating the type of `col1` is impossible. Creation of the table can be delayed until execution, with statements like `CREATE TEMPORARY TABLE`.

Instead, an *array* of `Item *` is used, implemented by `sp_rcontext::m_case_expr_holders`. The size of the array is static (it's the total number of cases), but the content of each element is dynamic (to account for the type of the case expression).

Caution

Note the wording used here: “static” means something that can be evaluated when compiling the code, in the parser, whereas “dynamic” means something that can be evaluated only when interpreting the code, during runtime. Of course, from a C++ coding point of view, everything is *dynamic*.

Inside a `CASE` statement, temporary local variables in a Stored Program are represented by the dedicated C++ class `Item_case_expr`. The class `Item_case_expr` is also a proxy, similar in nature to `Item_splocal`, and delegates to `sp_rcontext` for accessing the underlying case expression value. The coupling between `Item_case_expr` and `sp_rcontext` is based on `Item_case_expr::m_case_expr_id`, which is the case expression *index* in the symbol table (see `sp_pcontext`).

Exception Handlers

Warning

As of `MySQL 5.1`, which is the code base described by this documentation, the statements `SIGNAL`, `RESIGNAL` and `GET DIAGNOSTICS` are not supported.

Implementing these features will have some impact on the material described here, so changes in later versions are anticipated.

When the code *enters* a block of logic guarded by an SQL exception handler, the *state* or the runtime context in the interpreter *changes*, to represent this fact. The state change is not apparent immediately, it will only become apparent if an exception is raised. The internal runtime state of the engine also changes when the code *leaves* a block that contains an exception handler.

How exception handlers work during runtime is the subject of another section (see [Exception Handling](#)). What is described here is the state maintained internally, to represent *which* `HANDLER` is currently “active”, and *what* `CONDITION` is protected against.

The SQL precedence rules for `HANDLER` dictates that the last installed (inner most) handler is always considered first, so the natural structure to represent what handler is active is a *stack*, implemented by `sp_rcontext::m_handler` and `sp_rcontext::m_hcount`.

In addition, some extra information is required for `CONTINUE` handlers: the “address” in the code, or instruction pointer in the `sp_instr` array, of where to resume execution when the handler returns. This data is maintained in `sp_rcontext::m_hstack` and `sp_rcontext::m_hsp`, which again is a *stack* because exception handlers can be *nested* (exceptions can be raised and trapped *during* the execution of the body of an exception handler, too).

Executing One Instruction

Executing an instruction consists of calling the virtual method `sp_instr::execute()`, which is implemented for each instruction.

For instructions that can be executed directly, and don't depend on the evaluation of a general SQL statement or expression, the execution is very simple. See for example `sp_instr_jump::execute()`, `sp_instr_hpush_jump::execute()` or `sp_instr_hpop::execute()`. In all cases, the implementation of the `execute()` method is purely internal to the runtime interpreter.

For instructions that need to evaluate a general expression, like `sp_instr_jump_if_not::execute()`, or general instructions that need to execute an SQL statement, suchh as `sp_instr_stmt::execute()`, things are more complex. The implementation needs to leverage the existing code that is already capable of evaluating an expression or executing a *query*, and is implemented by the function `mysql_execute_command()`.

The function `mysql_execute_command()`, for historical reasons (it was implemented *before* Stored Programs), is mostly designed to consume directly the result of the parser, which is passed as input in `THD::lex`.

To comply with this interface, the runtime for stored program has to provide a `THD::lex` structure before executing each instruction, to prepare an execution environment which *looks as if* the statement to execute was just parsed. Dealing with the existing interface for re-entrant execution of SQL statements is the role of the C++ class `sp_lex_keeper`. The wrapper method to used to execute instructions is `sp_lex_keeper::reset_lex_and_exec_core()`, which ultimately invokes the `sp_instr::exec_core()` instructions implementation.

Flow Control

Instructions are numbered sequentially, and the current position in the code is represented by an “instruction pointer”, which is just an integer. In the main execution loop in `sp_head::execute()`, this instruction pointer is represented by the local variable `ip`.

When executing each instruction, the method `sp_head::execute()` is also responsible to return the address of the next instruction to execute. Most of the time, this corresponds to the “next” instruction (implemented by `m_ip+1`), except for absolute jumps (see `sp_instr_jump::execute()`) or conditional jumps (see `sp_instr_jump_if_not::execute()`).

Exception Handling

When the code *enters* a block protected by a `HANDLER`, the execution leads to `sp_instr_hpush_jump::execute()`, which installs the exception handler in the runtime handler stack, by calling `sp_rcontext::push_handler()`.

In a similar way, when the code *leaves* a block protected by a `HANDLER`, `sp_instr_hpop::execute()` removes the handlers installed by the matching `sp_instr_hpush_jump`, by calling `sp_rcontext::pop_handlers()`.

During the execution of *any* statement, different `CONDITION` can be raised at runtime, which are reported by the implementation of each statement by calling `push_warning()`, `my_error()` or similar functions. All these entry points ultimately leads to the error handler hook callback function implemented by `error_handler_hook` in `mysys/my_error.c`. In case of the server itself, this hook points to the function `my_message_sql()`.

Under normal circumstances, `my_message_sql()` just reports a warning or an error to the client application, and for errors causes the query to abort.

When executing a stored program, `THD::spcont` points to the runtime context of the program currently executed. When a `HANDLER` is active, the runtime context contains in its handler stack the list of all the `CONDITIONS` currently trapped, giving a chance to the call to `sp_rcontext::handle_error()` to intercept error handling.

If the condition reported does not match any of the conditions for which an exception handler is active, `sp_rcontext::handle_error()` returns false, and `my_message_sql()` raises the error or warning as usual.

When the condition reported *does* match an active `HANDLER`, that handler is *called*, but the technical nature of this call is special: the call is *asynchronous*. Instead of invoking the exception handler *directly*, `sp_rcontext::handle_error()` marks which exception handler is *to be called*, by setting the member variable `sp_rcontext::m_hfound`, and then returns true, so that `my_message_sql()` returns without reporting anything: at this point, the error condition has been totally masked, except for the fact that `sp_rcontext::m_hfound` is set.

Once `my_message_sql()` returns, the implementation of a given statement continues, either by proceeding if only a warning was reported, or by aborting the current execution if an error was raised. The execution of code in the server will eventually return from the implementation of a statement, and return from the call to `sp_instr::execute()` for that statement, returning control to the loop located in `sp_head::execute()`. Note that during the execution of the code that follows a call to `my_message_sql()`, error conditions are propagated in the call stack through the function's return value. It is transparent to the implementation of statements in general whether an exception was caught by an error handler.

After an instruction is executed in `sp_head::execute()`, the main interpreter loop checks for any pending exception handler code to call, by invoking `sp_rcontext::found_handler()`. If an exception was caught, the handler to call has been found by `sp_rcontext::handle_error()`, and is invoked.

In case of `CONTINUE HANDLER`, the instruction to return to after the handler code is executed needs to be saved in the runtime context. Finding the continuation destination is accomplished by the call to `sp_instr::get_cont_dest()` for the current instruction, whereas preserving this destination is done with `sp_rcontext::push_hstack()`. The matching call to `sp_rcontext::pop_hstack()`, which is executed when the exception handler is done, is located in `sp_instr_hreturn::execute()`.

Caution

To integrate properly with exception handling in general, the code should avoid testing for `thd->net.report_error`, or worse inspecting the content of the

error stack (displayed by `SHOW ERRORS`), because doing this actually assumes not only that an error was raised, but also that it was not caught. Instead, the proper way to implement error handling in the server is to return error status values and check for them.

Call Nesting

In the following example, the Stored Procedure `proc_1` makes a nested call to `proc_2`.

```
CREATE TABLE my_debug(
  seq int NOT NULL AUTO_INCREMENT,
  msg varchar(80),
  PRIMARY KEY(seq)
);

delimiter $$
CREATE PROCEDURE proc_1()
BEGIN
  INSERT INTO my_debug(msg) VALUES ("entering p1");
  CALL proc_2();
  INSERT INTO my_debug(msg) VALUES ("leaving p1");
END$$

CREATE PROCEDURE proc_2()
BEGIN
  INSERT INTO my_debug(msg) VALUES ("inside p2");
END$$

delimiter ;
CALL proc_1();
```

Tip

Oh, yes, we don't have a debugger yet, so this is old school `printf`-like debugging into a table.

By setting a breakpoint in `mysql_insert` in the server, the current thread stack at the first insert will look like this:

```
#0 mysql_insert () at sql_insert.cc:351
#1 in mysql_execute_command () at sql_parse.cc:2643
#2 in sp_instr_stmt::exec_core () at sp_head.cc:2609
#3 in sp_lex_keeper::reset_lex_and_exec_core () at sp_head.cc:2455
#4 in sp_instr_stmt::execute () at sp_head.cc:2560
#5 in sp_head::execute () at sp_head.cc:1077
#6 in sp_head::execute_procedure () at sp_head.cc:1726
#7 in mysql_execute_command () at sql_parse.cc:3807
#8 in mysql_parse () at sql_parse.cc:5274
#9 in dispatch_command () at sql_parse.cc:896
#10 in do_command () at sql_parse.cc:662
#11 in handle_one_connection () at sql_connect.cc:1089
#12 in start_thread () from /lib/libpthread.so.0
#13 in clone () from /lib/libc.so.6
```

By the time the second `INSERT` is executed, the stack will look like this:

```
#0 mysql_insert () at sql_insert.cc:351
#1 in mysql_execute_command () at sql_parse.cc:2643
#2 in sp_instr_stmt::exec_core () at sp_head.cc:2609
#3 in sp_lex_keeper::reset_lex_and_exec_core () at sp_head.cc:2455
#4 in sp_instr_stmt::execute () at sp_head.cc:2560
#5 in sp_head::execute () at sp_head.cc:1077
#6 in sp_head::execute_procedure () at sp_head.cc:1726
#7 in mysql_execute_command () at sql_parse.cc:3807
#8 in sp_instr_stmt::exec_core () at sp_head.cc:2609
#9 in sp_lex_keeper::reset_lex_and_exec_core () at sp_head.cc:2455
#10 in sp_instr_stmt::execute () at sp_head.cc:2560
#11 in sp_head::execute () at sp_head.cc:1077
#12 in sp_head::execute_procedure () at sp_head.cc:1726
```

```
#13 in mysql_execute_command () at sql_parse.cc:3807
#14 in mysql_parse () at sql_parse.cc:5274
#15 in dispatch_command () at sql_parse.cc:896
#16 in do_command () at sql_parse.cc:662
#17 in handle_one_connection () at sql_connect.cc:1089
#18 in start_thread () from /lib/libpthread.so.0
#19 in clone () from /lib/libc.so.6
```

In this stack trace, `sp_head::execute_procedure()` at #12 corresponds to `CALL proc_1()`, whereas `sp_head::execute_procedure()` at #6 corresponds to `CALL proc_2()`. In other words, recursive calls in the *user* SQL code are implemented by performing matching recursive calls in the *system* C++ code (the server).

This is actually a severe limitation of the implementation, which causes problems for the following reasons:

- User logic can be arbitrarily nested, with a long chain of Stored Programs calling other Stored Programs. The total depth of calls can be greater than one would expect, especially considering that a `VIEW` can invoke a `FUNCTION`, and that a `TRIGGER` can also invoke other `PROCEDURE`, `FUNCTION`, or `TRIGGER` objects.
- The amount of memory that can be consumed in the *stack* for a thread is not infinite. In fact, it's quite limited because $\{\text{MAX_NUMBER_OF_THREADS}\} * \{\text{MAX_THREAD_STACK}\} = \{\text{TOTAL_STACK}\}$. Note the catch in the equation here: *MAX* thread stack, which is dependent on the nesting of stored program in the user SQL code, for the worst case. Since MySQL currently does not use a thread pool but is compiled sometimes with a *BIG* number of threads, this can be a problem affecting scalability.
- As a result, the Stored Program interpreter has to protect itself against stack overflow. This is implemented by `check_stack_overflow()`

What *should* be implemented instead, is representing the user *SQL stack* on the C++ *heap*, and have the interpreter *loop* instead of making *recursive* calls.

There are also other good reasons to use the heap. For example, for error reporting, the current implementation has no way to tell that `proc_2` was called from `proc_1`, since this data is not available to the code; it's hidden in the C++ stack.

Nesting calls also has some impact on SQL exception handlers. The member `THD::spcont` for the current thread is not pointing to a single `sp_rcontext`, but to a *stack* of runtime contexts. This is implemented internally as a linked list, with `sp_rcontext::m_prev_runtime_ctx`.

With the example used, when the code is executing `proc_1`, `THD::spcont` points to the runtime context for `proc_1`. When the code is inside `proc_2`, the current thread `THD::spcont` points to `sp_rcontext{proc_2}`, which member `m_prev_runtime_ctx` points to `sp_rcontext{proc_1}`. This chain allows a parent Stored Program to catch exceptions raised by children Stored Programs.

Chapter 17. Prepared Statement and Stored Routine Re-Execution

Table of Contents

Statement Re-Execution Requirements	175
Preparation of a Prepared Statement	176
Execution of a Prepared Statement	177
Execution of a Stored Procedure Statement	178

Let us start with a general description of the MySQL statement processing workflow in order to provide the reader with understanding of the problem of reexecution and vocabulary for the following sections.

Conventional statements, that is, SQL queries sent in `COM_QUERY` protocol packet, are the only statements present in MySQL server prior to version 4.1. Execution of such statements is performed in a batch mode, one query processed by the server at a time. The original implementation is streamlined for this mode and has a single global connection state `THD` shared among all operational steps.

When executing a query in conventional mode, the server sequentially parses its text, acquires table level locks, analyzes the parsed tree, builds an execution plan, executes the built plan and releases the locks.

Memory for parsing is allocated using block allocator `MEM_ROOT` in 4k chunks and freed once in the end of execution. Memory for execution is allocated in the memory root of the parsed tree, as well as in the system heap, and in some cases in local "memory roots" of execution modules.

The role of the parser is to create a set of objects to represent the query. E.g. for a `SELECT` statement, this set includes a list of Item's for `SELECT` list, a list of tables (`TABLE_LIST` object for each table) for `FROM` clause, and a tree of Item's for `WHERE` clause.

During context analysis phase, links are established from the parsed tree to the physical objects of the database, such as open tables and table columns. A physical table is represented by a heir of class handler that corresponds to the storage engine the table belongs to, and is saved in `TABLE_LIST::file`.

When context analysis is done, the query optimizer is run. It performs two major tasks:

- Query transformation — a transformation of the parsed tree to an equivalent one, which is simpler and more efficient to execute.
- Creation of an execution plan, including evaluation of an order of joins and initialization of methods to access the used tables. At this step parts of the execution plan are attached to the parsed tree.

Finally, the query is passed to the execution runtime — an interpreter that operates with and modifies both the parsed tree and the execution plan in order to execute the query.

It should be noted that the overall procedure is infamous for breaking borders between abstraction layers. For example, MySQL performs [sub]query transformation during context analysis; moreover, most parts of the code rely on the fact that `THD` is processing only one statement at a time.

Statement Re-Execution Requirements

Features of MySQL 4.1 and 5.0 put a new demand on the execution process: prepared statements and stored routines need to reuse the same parsed tree to execute a query many times.

So far no easy mechanism that would allow query reexecution using the conventional query processing code has been found. For instance, copying of the parsed tree before each reexecution is not simple to

implement as a parsed tree, which can contain instances of more than 300 different classes, has a lot of cross-references between its objects.

The present solution introduces a concept of change management for the changes of the parsed tree and is largely a unification of numerous fixes of bugs in reexecution. The solution has two aspects.

The first one is that modifications of the parsed tree are tracked and a way to restore the tree to a state that allows reexecution is introduced.

The second aspect is that a dedicated block allocator (memory root) is used to store the parsed tree, and the memory allocated in this memory root is freed only when the parsed tree is destroyed. Later this memory root will be denoted as the permanent memory root of a statement.

In order to properly restore the parsed tree to a usable state, all modifications of the tree are classified as destructive or non-destructive and an appropriate action is taken for every type of modification.

A non-destructive modification does not depend on actual values of prepared statement placeholders or contents of the tables used in a query. Such modification is [and should be, for future changes] made only once and the memory for it is allocated in the permanent memory root of the statement.

As a result, the modified parsed tree remains usable.

Examples of non-destructive and useful modifications of the parsed tree are:

- `WHERE /ON` clause flattening
- `NOT` elimination
- `LEFT JOIN` elimination, when it can be done based on the constants explicitly specified in the query

The rest of modifications are destructive, generally because they are based on actual contents of tables or placeholders.

Examples of destructive modifications are:

- Equality propagation
- Sorting of members of `IN` array for quick evaluation of `IN` expression.

Destructive modifications are (and should be for all future changes) allocated in a memory root dedicated to execution, are registered in `THD::change_list` and rolled back in the end of each execution. Later the memory root dedicated to execution of a statement will be denoted as the runtime memory root of the statement. Because allocations are done indirectly via `THD::mem_root`, `THD::mem_root` at any given moment of time can point either to the permanent or to the runtime memory root of the statement. Consequently, `THD::mem_root` and `THD::free_list` can be denoted as 'currently active arena' of THD.

Preparation of a Prepared Statement

As mentioned above, `THD` is currently a required argument and the runtime context for every function in the server. Therefore, in order to call the parser and allocate memory in the statement memory root we perform several save-restore steps with `THD::mem_root` and `THD::free_list` (the active arena of `THD`).

1. In order to parse a statement, we save the currently active arena of `THD` and assign its members from the permanent arena of the statement. This is achieved by calling `THD::set_and_backup_active_arena`. This way `alloc_query` and `yyparse` operate on the permanent arena.

2. We don't want the garbage which is created during statement validation to be left in the permanent arena of the statement. For that, after parse but before validation of the statement, we restore the THD arena saved in (1). In other words, we use the arena of `THD` that was active when `Prepared_statement::prepare` was invoked as the runtime arena of the statement when it is validated.
3. Statement validation is performed in function `check_prepared_statement()`. This function will subsequently call `st_select_lex_unit::prepare()` and `setup_fields()` for the main LEX unit, create `JOINS` for every unit, and call `JOIN::prepare` for every join (`JOINS` in MySQL represents a part of the execution plan). Our prepared statement engine does not save the execution plan in a prepared statement for reuse, and ideally we should not create it at prepare stage. However, currently there is no other way to validate a statement except to call `JOIN::prepare` for all its units.
4. During validation we may perform a transformation of the parsed tree. In a clean implementation this would belong to a separate step, but in our case the majority of the server runtime was not refactored to support reexecution of statements, and a permanent transformation of the parsed tree can happen at any moment during validation. Such transformations **absolutely must** use the permanent arena of the prepared statement. To make this arena accessible, we save a pointer to it in `thd->stmt_arena` before calling `check_prepared_statement`.

Later, whenever we need to perform a permanent transformation, we first call `THD::activate_stmt_arena_if_needed` to make the permanent arena active, transform the tree, and restore the runtime arena.

5. Some parts of the execution do not distinguish between preparation of a prepared statement and its execution and perform destructive optimizations of the parsed tree even during validation. These changes of the parsed tree are recorded in `THD::change_list` using method `THD::register_item_tree_change`.
6. After the validation is done, we rollback the changes registered in `THD::change_list` and free new items and other memory allocated by destructive transformations.

Execution of a Prepared Statement

In order to call `mysql_execute_command` (the function that executes a statement) for a prepared statement and not damage its parse tree, we backup and restore the active `Query_arena` of `THD`.

- We don't want the garbage created during execution to be left in the permanent arena of the statement. To ensure that, every statement is executed in the runtime arena of `THD`. In other words, the arena which was active when `mysql_stmt_execute` was called is used as the runtime arena of the statement during its execution.
- Before calling `mysql_stmt_execute`, we allocate `thd->query` with parameter markers ('?') replaced with their values: the new query is allocated in the runtime arena. We'll need this query for general, binary, error and slow logs.
- The execution plan created at prepare stage is not saved (see [Preparation of a Prepared Statement](#)), and at execute we simply create a new set of `JOINS` and then prepare and optimize it. During the first execution of the prepared statement the server may perform non-destructive transformations of statement's parsed tree: normally that would belong to a separate step executed at statement prepare, but once again, this haven't been done in 4.1 or 5.0. Such transformations **absolutely must** use the permanent arena of the prepared statement (saved in `thd->stmt_arena`). Whenever we need to perform a permanent transformation, we first call `THD::activate_stmt_arena_if_needed` to make the permanent arena active, transform the tree, and restore the runtime arena. To avoid double transformations in such cases, we track current state of the parsed tree in `Query_arena::state`.

This state may be one of the following:

- `INITIALIZED` — we're in statement `PREPARE`.

- `INITIALIZED_FOR_SP` — we're in first execution of a stored procedure statement.
- `PREPARED` — we're in first execution of a prepared statement.
- `EXECUTED` — we're in a subsequent execution of a prepared statement or a stored procedure statement.
- `CONVENTIONAL_EXECUTION` — we're executing a pre-4.1 query.

One can use helper methods of `Query_arena` to check this state

```
(is_conventional_execution(), is_stmt_prepare(), is_stmt_execute(),  
is_stmt_prepare_or_first_sp_execute()).
```

Additionally, `st_select_lex_unit::first_execution` contains a flag for the state of each subquery in a complex statement. A separate variable is needed because not all subqueries may get executed during the first execution of a statement.

- Some optimizations damage the parsed tree, e.g. replace leafs and subtrees of items with other items or leave item objects cluttered with runtime data. To allow re-execution of a prepared statement the following mechanisms are currently employed:
 1. A hierarchy of `Item::cleanup()` and `st_select_lex::cleanup()` methods to restore the parsed tree to the condition of right-after-parse. These cleanups are called in `Prepared_statement::cleanup_stmt()` after the statement has been executed.
 2. In order to roll back destructive transformations of the parsed tree, every replacement of one item with another is registered in `THD::change_list` by using `THD::change_item_tree()`. In the end of execution all such changes are rolled back in reverse order.

Example:

```
if (!(fld= new Item_field(from_field)))  
goto error;  
thd->change_item_tree(reference, fld);
```

If a transformation is a non-destructive, it should not be registered, but performed only once in the permanent memory root. Additionally, be careful to not supply a pointer to stack as the first argument of `change_item_tree()`; that will lead to stack corruption when a tree is restored.

3. `AND/OR` subtrees of `WHERE` and `ON` clauses are created anew for each execution. It was easier to implement in 4.1, and the approach with change record list used in (b) could not have been used for `AND/OR` transformations, because these transformations not only replace one item with another, but also can remove a complete subtree. Leafs of `AND/OR` subtrees are not copied by this mechanism because currently they are not damaged by the transformation. For details, see `Item::copy_andor_structure()`.
 4. **No** other mechanism exists in the server at the moment to allow re-execution. If the code that you're adding transforms the parsed tree, you must use one of the mechanisms described above, or propose and implement a better approach.
- When execution is done, we rollback the damage of the parsed tree.

Execution of a Stored Procedure Statement

Execution of a stored procedure statement is similar to execution of a prepared statement. The few existing exceptions are described below.

During execution of a stored procedure, `THD::stmt_arena` points to the permanent query arena of the stored procedure. This arena happens to be also the permanent query arena of every instruction of the procedure, as the parser creates all instructions in the same arena. More generally,

`THD::stmt_arena` is always set and always points to the permanent arena of a statement. If the statement is a conventional query, then the permanent arena simply points to the runtime arena of the query.

An own runtime memory root is set up for execution of every stored procedure statement and freed in the end of execution. This is a necessary measure to avoid memory leaks if a stored procedure statement is executed in a loop.

With regard to the transformations and restoration of the parsed tree, execution of a stored procedure statement follows the path of execution of a prepared statement, with the exception that there is no separate prepare step. `THD::is_first_sp_execute()` is used to determine whether it's the first execution, and in this case non-destructive permanent transformations of the parsed tree are made in the permanent memory root of the statement that is currently being executed.

During subsequent executions no non-destructive transformations are performed, while all destructive ones are rolled back in the end of execution using the same algorithm as in prepared statements.

Chapter 18. Writing a Procedure

Table of Contents

Extend Class <code>Procedure</code>	181
Constructor	181
<code>change_columns()</code>	181
<code>send_row()</code>	182
<code>add()</code>	182
<code>end_group()</code>	182
<code>end_of_records()</code>	183
Initialization Callback	183
Calling Sequence	183
Required Server Code Patches	183
Examples	184
<code>PROCEDURE ROWNUM()</code> – Adding a <code>RowNum</code> Pseudo Column to a Result Set	184
<code>PROCEDURE CALLTRACE()</code> – Simple Trace of Procedure Member Function Calls	185

Note

This section is not about SQL stored procedures but about MySQL procedures like `PROCEDURE ANALYSE`. For internal information on SQL stored procedures see the [Chapter 16, Stored Programs](#).

Extend Class `Procedure`

Each new procedure needs to extend the **Procedure** class. For a minimal dummy procedure that doesn't actually change the result set it would look like this:

```
class proc_dummy: public Procedure { }
```

In a real procedure you'd extend at least some of the member functions below:

Constructor

Prototype: **n/a**

The class constructors prototype signature is completely up to you. The only place where objects are instantiated is your own init callback.

To initialize your derived procedure object you have to pass on the `select_result` pointer the init callback was called with to the base class constructor together with a flag parameter which specifies what kind of procedure you are going to implement. So a minimal constructor would look like this:

```
your_proc::your_proc(select_result *res)
    :Procedure(res, PROC_NO_SORT)
{
}
```

Possible flag values are **PROC_NO_SORT** and **PROC_GROUP**. I have no real idea yet what the two flags are doing but found that for simple procedures `PROC_NO_SORT` seems to be the right flag to use.

See also [Initialization Callback](#).

`change_columns()`

Prototype: **virtual bool change_columns(List<Item> &field_list);**

Here you can change the structure of the result field list, e.g. you can add fields to the **field_list** or replace the queries result fields by something completely different altogether (like PROCEDURE ANALYSE() does).

An example that adds an INTEGER field at the end of the field list:

```
bool proc_rownum::change_columns(List<Item> &field_list)
{
    DEBUG_ENTER("proc_rownum::change_columns");

    // create a new column item
    row_num_column = new Item_proc_int("RowNum");

    // and attach it to the column list
    field_list.push_back(row_num_column);

    DEBUG_RETURN(0);
}
```

send_row()

Prototype: **virtual int send_row(List<Item> &fields);**

This member is called for every result row in the original result set. Whatever you do here is up to you, it is important to note though that to pass on the result row to the client you have to call **result->send_data()** yourself.

PROCEDURE ANALYSE() for example does not send any data here, it only produces result rows after aggregating information across all result rows so its **send_row()** member only aggregates but doesn't send anything.

A simple example which modifies the result value for a single field in the field list before sending it on to the client:

```
int proc_rownum::send_row(List<Item> &field_list __attribute__((unused)))
{
    DEBUG_ENTER("proc_rownum::send_row");

    // increment row count and set its new value in result row
    row_num_column->set(++row_num);

    // now send the modified results
    if (result->send_data(field_list))
        DEBUG_RETURN(-1);

    DEBUG_RETURN(0);
}
```

add()

Prototype: **virtual void add(void);**

This member function is called once for every source row for a GROUP BY query.

See also [end_group\(\)](#).

end_group()

Prototype: **virtual void end_group(void);**

This member function is called whenever the end of a group in a GROUP BY is detected, it is called after the call to [add\(\)](#) for the last source row in the group but before sending the actual aggregated result row for the group with [send_row\(\)](#).

end_of_records()

Prototype: **virtual bool end_of_records(void);**

This member function is called at the very end after all result rows have been processed with calls to `send_row()`. This is where you can send extra summary result rows as e.g. PROCEDURE ANALYSE() does.

Initialization Callback

The initialization callback is registered together with the procedure name in the **sql_procs** array in **procedure.cc**.

The initialization callback prototype signature looks like this:

```
Procedure *(*init)(THD *thd,ORDER *param,select_result *result,
List<Item> &field_list);
```

In the initialization callback you usually just create and return an instance of your derived Procedure class:

```
// Create and register the actual procedure object
Procedure *proc_rownum_init(THD *thd,
                           ORDER *param,
                           select_result *result,
                           List<Item> &field_list)
{
    DEBUG_ENTER("proc_rownum_init");

    proc_rownum *pc = new proc_rownum(result);

    DEBUG_RETURN(pc);
}
```

Calling Sequence

The **init callback** is always called first at the beginning of a query, followed by a call to the `change_columns()`. The `end_of_records()` member is always called last at the very end.

For simple non-grouping queries only `send_row()` is called once for each result row.

For grouping queries `add()` is called once for each source row. `end_group()` is called once at the end of each group followed by a call to `send_row()`.

You can use the CALLTRACE() procedure to check in which sequence the member functions are called on any query. This procedure will send one result row for each member function call, every row contains a single text field with one of the values **add**, **end_group**, **end_of_records** or **send_row**.

Required Server Code Patches

Currently procedures can only be compiled into the server statically. There is no dynamic procedure for loading them dynamically yet like we have for UDFs and plugins.

So you have to perform the following steps to register your procedure with the server:

- register the **.cc** and **.h** file in the **mysqld_SOURCES** list in **sql/Makefile.am**, e.g. right after the line that the **sql_analyse.*** files are on:

```
diff -ruN 5.0/sql/Makefile.am 5.0-myproc/sql/Makefile.am
--- 5.0/sql/Makefile.am 2007-08-09 12:11:16.000000000 +0200
+++ 5.0-myproc/sql/Makefile.am 2007-08-09 00:22:19.000000000 +0200
@@ -94,6 +94,7 @@
    sql_db.cc sql_table.cc sql_rename.cc sql_crypt.cc \
```

```

    sql_load.cc mf_iocache.cc field_conv.cc sql_show.cc \
    sql_udf.cc sql_analyse.cc sql_analyse.h sql_cache.cc \
+   procedure_rownum.cc procedure_rownum.h \
    slave.cc sql_repl.cc sql_union.cc sql_derived.cc \
    client.c sql_client.cc mini_client_errors.c pack.c\
    stacktrace.c repl_failsafe.h repl_failsafe.cc \

```

- register the **.cc** source file in the **sqlsources** list in **libmysqld/Makefile.am**

```

diff -rUN 5.0/libmysqld/Makefile.am 5.0-myproc/libmysqld/Makefile.am
--- 5.0/libmysqld/Makefile.am 2007-08-09 12:18:16.000000000 +0200
+++ 5.0-myproc/libmysqld/Makefile.am 2007-08-09 00:21:55.000000000 +0200
@@ -54,6 +54,7 @@
    opt_sum.cc procedure.cc records.cc sql_acl.cc \
    sql_load.cc discover.cc sql_locale.cc \
    sql_analyse.cc sql_base.cc sql_cache.cc sql_class.cc \
+   procedure_rownum.cc \
    sql_crypt.cc sql_db.cc sql_delete.cc sql_error.cc sql_insert.cc \
    sql_lex.cc sql_list.cc sql_manager.cc sql_map.cc sql_parse.cc \
    sql_prepare.cc sql_derived.cc sql_rename.cc \

```

- include your **.h** file in **sql/procedure.cc**

```

diff -rUN 5.0/sql/procedure.cc 5.0-myproc/sql/procedure.cc
--- 5.0/sql/procedure.cc 2007-08-09 12:11:16.000000000 +0200
+++ 5.0-myproc/sql/procedure.cc 2007-08-09 00:21:04.000000000 +0200
@@ -23,6 +23,7 @@
    #include "mysql_priv.h"
    #include "procedure.h"
    #include "sql_analyse.h"    // Includes procedure
+   #include "procedure_rownum.h" // Includes procedure
    #ifdef USE_PROC_RANGE
    #include "proc_range.h"
    #endif

```

- register your procedures init callback in the **sql_procs** array in **sql/procedure.cc**

```

diff -rUN 5.0/sql/procedure.cc 5.0-myproc/sql/procedure.cc
--- 5.0/sql/procedure.cc 2007-08-09 12:11:16.000000000 +0200
+++ 5.0-myproc/sql/procedure.cc 2007-08-09 00:21:04.000000000 +0200
@@ -37,6 +38,7 @@
    { "split_count", proc_count_range_init }, // Internal procedure at TCX
    { "matris_ranges", proc_matris_range_init }, // Internal procedure at TCX
    #endif
+   { "rownum", proc_rownum_init }, // Add RowNum column to result
    { "analyse", proc_analyse_init } // Analyse a result
};

```

- run **automake** and **autoconf** to regenerate the **Makefile.in**

files and the [configure](#) script

- run **configure** or refresh your previous configuration by running **config.status**
- run **make**
- if your build fails with **don't know how to make xxx.cc required by xxx.o** in the **libmysqld/** directory you might need to create a symlink from **libmysqld/xxx.cc** to **sql/xxx.cc** yourself. The build system is not too clever about creating these symlinks for files added after the first compile ...

Examples

PROCEDURE ROWNUM() – Adding a RowNum Pseudo Column to a Result Set

A [patch](#) against current MySQL 5.0 (should work with MySQL 5.1, too).

Example:

```
mysql> SELECT User, Host FROM mysql.user PROCEDURE ROWNUM();
```

User	Host	RowNum
root	127.0.0.1	0
root	linux	1
root	localhost	2

```
3 rows in set (0.00 sec)
```

How to apply:

- [download the patch](#)
- in the MySQL 5.0 source directory to **patch -p1 < proc-rownum-1.0.0.patch**
- run **automake** and **autoconf** to regenerate **configure**
- run **./configure** as usual or refresh your previous configuration by running **./config.status**
- run **make** and **make install** to compile and install the server
- restart the server
- try appending **PROCEDURE ROWNUM()** to a query

PROCEDURE CALLTRACE() – Simple Trace of Procedure Member Function Calls

A [patch](#) against current MySQL 5.0 (should work with MySQL 5.1, too).

This procedure will replace the actual query results with one row per procedure object member function call, so showing the sequence these functions are called on a result set.

Example:

```
mysql> select * from t2;
```

i	j
1	1
2	1
2	2
3	1
3	2
3	3

```
6 rows in set (0.00 sec)
```



```
mysql> select i from t2 group by i procedure calltrace();
```

Call
add
end_group
send_row
add
add
end_group
send_row
add
add
add
end_group
send_row
end_of_records

```
13 rows in set (0.00 sec)
```

How to apply: see the PROCEDURE ROWNUM() example above, the patch for this procedure needs to be applied in exactly the same way.

Chapter 19. Replication

Table of Contents

Chapter Organization	187
Source Code Files	188
Principles	190
Binlog Formats	190
Differences Between Master and Slave	191
Crashes	194
Binlog Files and Indexes	194
Rules	194
Determining the Logging Format	194
Unsafe Statements	195
Logging Transactions	195
Logging Updates to <code>auto_increment</code> Columns	198
Logging Access to Variables and Functions	198
Other Unsafe Statements	198
<code>binlog_row_image</code>	198
Replication Locks	199

Status of this section: up to date 2009-10-21

This chapter describes MySQL replication principles, rules, and code, as it is supposed to work in version 5.1.

The MySQL replication feature allows a server - the **master** - to send all changes to another server - the **slave** - and the slave tries to apply all changes to keep up-to-date with the master. Replication works as follows:

- Whenever the master's database is modified, the change is written to a file, the so-called **binary log**, or **binlog**. This is done by the client thread that executed the query that modified the database.
- The master has a thread, called the **dump thread**, that continuously reads the master's binlog and sends it to the slave.
- The slave has a thread, called the **IO thread**, that receives the binlog that the master's dump thread sent, and writes it to a file: the **relay log**.
- The slave has another thread, called the **SQL thread**, that continuously reads the relay log and applies the changes to the slave server.

Chapter Organization

We distinguish between two levels of the architecture: *principles*, and *rules*:

principle	High-level goal that declares how the program shall work, from an external point of view. Principles do not discuss how the code works. An example of a principle is "no row events shall be written to the binlog if <code>binlog_format=statement</code> ".
rule	Defines how to implement the principles. Rules can be formulated on a very high level, but they describe implementation (how the code works from the inside) rather than interface (how the program behaves when looking from the outside). An example of a rule is "at ROLLBACK when the transaction cache only contains updates to

transactional tables, clear the transaction cache and do not write to the binlog".

In the next section, we list the source code files used by replication and describe what each file contains. The section after describes the replication principles and the following section describes the replication rules. The last section contains very old, obsolete documentation. It will be removed after we have verified that anything useful has been transferred to the main sections.

Source Code Files

Status of this section: up to date 2009-12-16

Files in the **sql** directory:

File	Description
slave.h/.cc	Contains the slave IO and SQL threads . This is the high-level administrative logic for the slave threads - i.e., not the low-level functions for parsing the dump from the net, parsing the relay log, formatting the relay log, or executing events. Also contains the code to execute SHOW SLAVE STATUS .
log.h/.cc	The high-level binary logging mechanism for organizing events into a sequence so that it becomes a binary log. Routines for creating, writing, and deleting binary log files. Also the handler callbacks for the binlog.
log_event.h/.cc	The Log_event class and subclasses, for creating, writing, reading, printing, and applying events of every event type. Reading and writing here is at a low level, that is, serializing values into records.
old_log_event.h/.cc	Contains classes to read and execute old versions of the row log events.
rpl_rli.h/.cc	Implementation of the data structure Relay_log_info that holds the state of the SQL thread, and also some auxiliary methods used by the slave SQL thread. (The slave SQL thread is in slave.cc and it also uses other auxiliary functions.)
rpl_mi.h/.cc	Contains the data structure Master_info that holds some state of the IO thread (the IO thread is in slave.cc and it also uses other auxiliary functions).
sql_repl.cc	The dump thread , where the master sends its binary log to the slave. This is also where the code resides for RESET SLAVE, CHANGE MASTER, RESET MASTER, SHOW BINLOG EVENTS, SHOW MASTER STATUS, SHOW BINARY LOGS, PURGE BINARY LOGS , as well as some replication-related system variables.
sql_binlog.cc	Contains code to execute BINLOG statements (i.e., the base64-encoded things that mysqlbinlog prints when it sees row events).
rpl_record.h/.cc	Utilities for encoding and decoding table rows into and out of the row event format .

repl_failsafe.h/.cc	Utilities to initialize and register slaves on the master . Also unfinished and unused code dealing with "failsafe" (master election if the primary master fails).
replication.h	Observer class declarations, which together constitute the binary log interface .
rpl_constants.h	Enumeration of incidents (events that occur during replication). Also some constants that are local to the replication code.
rpl_filter.h/.cc	Implements the table and database filters used by the --{binlog,replicate}--{do,ignore}-db, --replicate[-wild]--{do,ignore}-table, and --replicate-rewrite-db flags.
rpl_handler.h/.cc	Coordination classes used by plugins to register to the binary log interface.
rpl_injector.h/.cc	The injector class that allows external insertions into the binary log. This is used for cluster replication binary logging.
rpl_reporting.h/.cc	Utilities for reporting replication conditions and reporting errors, warnings, and informational messages on the slave.
rpl_tblmap.h/.cc	Utilities to generate a mapping from numbers to tables . The mapping is used by the row logging system to identify tables.
rpl_utility.h/.cc	Auxiliary classes and functions used for Table_map_events, and also an auxiliary class for smart pointers.
sql_base.cc	Prior to Bug#39934 , the function decide_logging_format() that determines if statements should be written row-based or statement-based to the binlog. After Bug#39934 , there is nothing related to replication here.
sql_class.cc	The function binlog_query() , called from commands that need to log a query_log_event. After Bug#39934 , also decide_logging_format() .
sql_lex.h/.cc	List of all types of unsafe statements , and functions for marking statements unsafe.

Files in the **client** directory:

File	Description
mysqlbinlog.cc	The mysqlbinlog program. This file mainly contains a loop that calls auxiliary functions (members of Log_event defined in log_event.cc) that read and print events.

Files in the **plugin/semisync** directory:

File	Description
semisync.h/.cc	Auxiliary code, particularly for tracing, that is used by both the master semisync module and the slave semisync module.

semisync_master.h/.cc	The master semisync module.
semisync_master_plugin.cc	The callbacks invoked by the server to use the master semisync module, as well as code to register the master semisync module.
semisync_slave.h/.cc	The slave semisync module.
semisync_slave_plugin.cc	The callbacks invoked by the server to use the slave semisync module, as well as code to register the master semisync module.

Principles

In this section, we describe the architectural principles of replication. These are high-level goals that replication shall achieve. The principles have been used as guidelines to construct the Rules of replication (next section).

Binlog Formats

Status of this subsection: Complete but not reviewed 2009-10-21

The binlog is organized as a linear sequence of **events**. An SQL query that modifies the database will generate one or more events and append them to the binlog. There are also auxiliary event types that describe the structure of the binlog.

Queries can be logged in two ways:

- In statement format: the SQL query is written to the binlog in text.
- In row format: rows that changed are written to the binlog in a binary format. Each row may consist of a **Before Image (BI)** and/or an **After Image (AI)**. The BI identifies the row to modify and the AI describes the row after the change. There are three types of log_events:
 - Write_rows_log_event: adds a new row to a table. Has only AI.
 - Update_rows_log_event: modifies an existing row in a table. Has both BI and AI.
 - Delete_rows_log_event: removes an existing row from a table. Has only BI.

Which of the two formats to use is configured with the @@session.binlog_format variable, which takes the values STATEMENT, ROW, or MIXED. The following principles shall hold:

(P-binlog_format-statement) @@session.binlog_format=STATEMENT:

- We do not guarantee correct logging.
- The client may not generate row events.
- If the server cannot determine that a statement is correctly logged, a warning or error shall be issued:
 - If it is possible that the user (through application-specific logic) knows that the statement is correctly logged, then a warning shall be issued.
 - If it is inherently impossible for the user to determine that the statement will be correctly logged, an error shall be issued and the statement shall not execute.

(P-binlog_format-row) @@session.binlog_format=ROW:

- We guarantee correct logging. If a statement cannot be correctly logged, then an error shall be generated and the statement shall not execute.

- DML changes may only be logged in row format, not in statement format. If a DML change cannot be logged in row format, then an error shall be generated and the statement shall not execute.

(P-binlog_format-mixed) @@session.binlog_format=MIXED:

- We guarantee correct logging. If a statement cannot be correctly logged, then an error shall be generated and the statement shall not execute.
- If correct logging can be guaranteed by logging in statement format, then statement format shall be used. Otherwise, row format shall be used.

Clarification: If it cannot be determined in a practical manner that statement format leads to correct logging, then row format shall be used.

Differences Between Master and Slave

Status of this subsection: In progress 2009-10-21

What does the term **correct replication** really mean? To clarify the notion, we make the following preliminary definitions:

(D-identical-environments) Two environments are **identical** if all the following are identical:

- The hardware representation of floating point numbers and the hardware implementation of floating point arithmetic
- The case sensitivity of the file systems
- The versions of all used components of MySQL

Note: The following are examples of things not taken into account by this definition:

- The hardware's word size, as long as it is supported by MySQL
- The hardware's endianness

(D-identical-server-states) Two server states are **identical** if all the following are identical:

- The sets of databases (a.k.a. schemas)
- The table definitions (including table names, table options, column definitions) of all tables outside the mysql and information_schema databases
- The table contents, modulo row order (in mathematical language: the unordered multisets of rows are equal), of all tables outside the mysql and information_schema databases
- The definitions of all functions, procedures, triggers, views, prepared statements, and events

Note: The following are examples of things not taken into account in this definition:

- System variables and user variables
- The state of clients, including the replication slave. The state includes session variables and temporary tables.
- Binlog files, relay log files, binlog indexes, relay log indexes
- The internal state of the random number generator
- Which plugins are installed

TODO: Question:

- Should user privileges count as server state?

(D-rpl-correct) Replication is **correct** if both the following hold:

- Any change on the master eventually results in the same change on the slave.
- Any intermediate state of the slave is identical to some intermediate state of the master.

Note: it is not required that each intermediate state of the master is identical to some intermediate state of the slave.

We now state the architectural principles that define when replication shall be correct. The following is the main rule:

(P-rpl-correct) If a replication master and slave reside on identical environments (D-identical-environments) and the server states are identical (D-identical-server-states), and @@client.binlog_format!=STATEMENT, then replication shall be correct.

Exceptions: Situations Where We Do Not Guarantee Correct Replication

There are some exceptions to (P-rpl-correct) where we do not guarantee correct replication.

(P-exception-federated) If a table uses a federated table on the master, then TODO

(P-exception-table-definition) If a CREATE TABLE uses a DATA DIRECTORY or INDEX DIRECTORY clause, then the table may not be correctly replicated.

(P-exception-plugins) TODO: figure out principles for plugin replication

Additions: Special Situations Where We Do Support Correct Replication

In addition to what we guarantee in (P-rpl-correct), we also guarantee correct replication in the following scenarios:

(P-rpl-different-file-system-case-sensitivities) TODO: allowed differences in file system case sensitivity

(P-rpl-different-versions) Replication shall be correct even if master has version a.b.c and slave has version A.B.C, where $A.B.C \geq a.b.c$ and $A \leq a+1$.

(P-rpl-different-table-definitions) Table options may differ in the following fields: TODO: (comments, data/index directories, various hints)

(P-rpl-different-engines) TODO: allowed differences in storage engines

(P-rpl-different-column-definitions) Replication shall be correct even if the table definitions differ in one or more of the following ways:

(rpl-extra-slave-columns) binlog_format=ROW and the slave has extra columns *after* the columns of the master. Notes:

- Extra indexes still must follow the rules for indexes - see below.
- For purposes of defining "correct replication", the tables are considered equal if they are equal on the common columns.

(rpl-missing-slave-columns) binlog_format=ROW and the master has extra columns *after* the columns of the slave, as long as the following rules apply:

1. The slave cannot have both missing columns and extra columns (see above) at the same time.

2. If the master uses `--binlog_row_image=minimal` or `binlog_row_image=noblob`, then the BI must contain at least one column that exists on the slave. Moreover, the set of columns that are logged in BI must not match two different rows on the slave (but it may match two or more identical rows). This can be ensured, e.g., by one of the following strategies:
 - a. The master has a PK that only includes columns that the slave has.
 - b. The master has a PK that includes all columns that the slave has, and possibly other columns too.

Note: For purposes of defining "correct replication", the tables are considered equal if they are equal on the common columns.

(rpl-type-promotion) The data type of a column differs as allowed in [the manual](#).

(rpl-different-keys) Keys, indexes, and NOT NULL attributes may differ freely between master and slave, as long as the following rule applies:

If the slave has an enabled key, and the master does not have an enabled key of the same type over the exact same set of columns (e.g., because the key is missing/disabled or of a different type; or because the columns only exist on the slave), then the semantics of the slave's key must be ensured before the rows are inserted on the slave. Specifically:

1. If the slave has a uniqueness constraint (PK or UK), then uniqueness must be guaranteed before a row is inserted on the slave. This can be done, for example, through the following strategies:
 - a. Have a uniqueness constraint (UK or PK) on the master, over the same columns or over a subset of the columns.
 - b. If a column only exists on the slave (or if application-specific logic ensures that only NULL values are inserted on the master), then the AUTOINCREMENT attribute can be used on the slave.
 - c. Use application-specific logic on the master that ensures that rows inserted are unique in the key's columns.
 - d. Use BEFORE INSERT and/or BEFORE UPDATE triggers on the slave that ensure (through application-specific logic) that the rows are unique.
2. If the slave has a non-NULL constraint (PK or NOT NULL), then the absence of NULL values must be ensured before a row is inserted on the slave. This can be done analogously to how uniqueness constraints are satisfied above:
 - a. Have non-NULL constraints (PK or NOT NULL) on the master covering all the columns.
 - b. Use AUTOINCREMENT as above.
 - c. Use application-specific logic on the master that ensures no NULL values are inserted.
 - d. Use BEFORE INSERT and/or BEFORE UPDATE triggers on the slave that ensure (through application-specific logic) that no NULL values are inserted.

Note: There are no restrictions on extra keys on the master.

(rpl-column-names) If `binlog_format=ROW`, then column names may differ: columns are identified only by their position.

(P-rpl-different-rows) TODO: allowed extra rows on master or slave

(P-rpl-different-default-values) TODO:

Crashes

Status of this subsection: Not started 2009-10-21

Binlog Files and Indexes

Status of this subsection: Not started 2009-10-21

Rules

Determining the Logging Format

For each statement, we must determine the logging format: row or statement. This is done as follows.

1. At parse time, it is detected if the statement is unsafe to log in statement format (i.e., requires row format). If this is the case, the `THD::Lex::set_stmt_unsafe()` function is called. This must be done prior to the call to `THD::decide_logging_format()` (i.e., prior to `lock_tables`). As a special case, some types of unsafeness are detected inside `THD::decide_logging_format()`, before the logging format is decided. Note that statements shall be marked unsafe even if `binlog_format!=mixed`.
2. `THD::decide_logging_format()` determines the logging format, based on the value of `binlog_format` and the unsafeness of the current statement.
3. `THD::decide_logging_format()` also determines if the statement is impossible to log, in which case it generates an error and the statement is not executed. The statement may be impossible to log for the following reasons:
 - both row-incapable engines and statement-incapable engines are involved (`ER_BINLOG_ROW_ENGINE_AND_STMT_ENGINE`)
 - `BINLOG_FORMAT = ROW` and at least one table uses a storage engine limited to statement-logging (`ER_BINLOG_ROW_MODE_AND_STMT_ENGINE`)
 - statement is unsafe, `BINLOG_FORMAT = MIXED`, and storage engine is limited to statement-logging and (`ER_BINLOG_UNSAFE_AND_STMT_ENGINE`)
 - statement is a row injection (i.e., a row event executed by the slave SQL thread or a BINLOG statement) and at least one table uses a storage engine limited to statement-logging (`ER_BINLOG_ROW_INJECTION_AND_STMT_ENGINE`)
 - `BINLOG_FORMAT = STATEMENT` and at least one table uses a storage engine limited to row-logging (`ER_BINLOG_STMT_MODE_AND_ROW_ENGINE`)
 - statement is a row injection (i.e., a row event executed by the slave SQL thread or a BINLOG statement) and `BINLOG_FORMAT = STATEMENT` (`ER_BINLOG_ROW_INJECTION_AND_STMT_MODE`)
 - more than one engine is involved and at least one engine is self-logging (`ER_BINLOG_MULTIPLE_ENGINES_AND_SELF_LOGGING_ENGINE`)

See the comment above `decide_logging_format` for details.

4. `THD::decide_logging_format()` also determines if a warning shall be issued. A warning is issued for unsafe statements if `binlog_format=STATEMENT`. Warnings are not issued immediately; instead, `THD::binlog_stmt_unsafe_flags` is set and the warning is issued in `THD::binlog_query()`. This prevents warnings in the case that the statement generates an error later so that it is not logged.

Sub-statements. Let T be a statement that invokes an unsafe sub-statement S (S may be a stored function, stored procedure, trigger, view, or prepared statement). Each sub-statement is cached as an `sp_head` object. The `sp_head` object stores the Lex that was generated when the statement defining

the sub-statement was parsed (i.e., when CREATE FUNCTION/CREATE PROCEDURE/CREATE TRIGGER/CREATE VIEW/PREPARE was parsed). Hence, this cached Lex has the unsafe flag set. When T is parsed, it fetches S from the cache. At this point, it calls `sp_head::propagate_attributes()`, which marks the current Lex object as unsafe if the cached Lex object was unsafe.

Unsafe Statements

NOTE: the following list is incomplete; it does not take into account changes made in 2010 or later (roughly).

A statement may be flagged as **unsafe**. An unsafe statement will be logged in row format if `binlog_format=MIXED` and will generate a warning if `binlog_format=STATEMENT`.

The following types of sub-statements are currently marked unsafe:

- System functions that may return a different value on slave, including: `FOUND_ROWS`, `GET_LOCK`, `IS_FREE_LOCK`, `IS_USED_LOCK`, `LOAD_FILE`, `MASTER_POS_WAIT`, `RAND`, `RELEASE_LOCK`, `ROW_COUNT`, `SESSION_USER`, `SLEEP`, `SYSDATE`, `SYSTEM_USER`, `USER`, `UUID`, `UUID_SHORT`.

Note: the following non-deterministic functions are not marked unsafe:

- `CONNECTION_ID` (Query_log_events contain the connection number)
- `CURDATE`, `CURRENT_DATE`, `CURRENT_TIME`, `CURRENT_TIMESTAMP`, `CURTIME`, `LOCALTIME`, `LOCALTIMESTAMP`, `NOW`, `UNIX_TIMESTAMP`, `UTC_DATE`, `UTC_TIME`, `UTC_TIMESTAMP` (Query_log_event contain timezone and the time when the statement was executed)
- `LAST_INSERT_ID` (this is replicated in an `Intvar_log_event`)

Also note that most floating-point math functions will return a hardware-dependent result. We do not mark such function unsafe, because we only support replication between platforms that use identical floating point math.

- System variables, with some exceptions listed at <http://dev.mysql.com/doc/en/binary-log-mixed.html>
- UDFs: since we have no control over what the UDF does, it may be doing something unsafe.
- Update from a sub-statement of a table that has an autoincrement column. This is unsafe because the `Intvar_log_event` is limited to only hold autoincrement values for one table.
- `INSERT DELAYED`, since the rows inserted may interleave with concurrently executing statements.
- Updates using `LIMIT`, since the order in which rows are retrieved is not specified.
- Statements referencing system log tables, since the contents of those tables may differ between master and slave.
- Non-transactional reads or writes executing after transactional reads or writes in a transaction (see [Logging Transactions](#)).
- Reads or writes to self-logging tables, and all statements executing after reads or writes to self-logging tables in the same transaction.

The following has not yet been implemented:

- Statements using fulltext parser plugins (cf. [Bug#48183](#))

Logging Transactions

Status of this subsection: complete but not reviewed 2009-10-21

There are several types of statements that require attention because of their special behavior in transactions:

- Non-transactional updates that take place inside a transaction present problems for logging because (1) they are visible to other clients before the transaction is committed, and (2) they are not rolled back even if the transaction is rolled back. It is not always possible to log correctly in statement format when both transactional and non-transactional tables are used in the same transaction.
- Statements that do an implicit commit (i.e., most but not all DDL, and some utility commands) are logged specially due to unspecified requirements by NDB.
- Statements that update temporary tables need special treatment since they are not logged in row format.

Definitions

To reason about logging different table types, we make some preliminary definitions.

(D-T-table) A table that has a transactional engine is called a **T-table**.

(D-N-table) A table that has a non-transactional engine is called an **N-table**.

(D-N-write) A statement makes an **N-write** if it makes any type of change to the server state that will not be changed by a ROLLBACK.

Note: N-writes include updates to N-tables, but also CREATE and DROP statements.

(D-log-target) Events are either appended to the **Transaction Cache (TC)** or to the **Statement Cache (SC)** or written directly to the binlog.

Preliminary Rules

The following preliminary rules are actually consequences of the principle that statements shall be correctly logged when `binlog_format=MIXED` or `ROW`. They also apply when `binlog_format=STATEMENT`: this makes statement format work in many practical cases.

(PR-causality) If statement A is executed before statement B, and B is logged in statement format, and B reads tables that A may modifies, then B shall be logged after A.

(PR-durability) Events shall be written to the binary log at the moment they take effect. In particular, changes to N-tables shall be written to the binary log when they have been executed, and changes to T-tables shall be written to the binary log on commit. If `--sync-binlog` has been specified, then it suffices that events are be written to the binary log at the next synchronization point.

(PR-causality-precedence) If P-causality and P-durability cannot both be fulfilled, then P-causality is considered more important.

Rules for Non-committing Statements, Except `CREATE TEMPORARY TABLE...SELECT`

The preliminary rules above, together with the principles for logging format, have been used to construct the following rules.

CALL statements are unrolled (see `???``TODO`: add section about unrolling`???`), so that each statement executed by the stored procedure is logged separately. (If a stored procedure A invokes a stored procedure B, then B is unrolled recursively). In the following, we assume that unrolling has already been done, and the word "statement" refers to a non-CALL top-level statement or a non-CALL sub-statement.

Let S be a logged statement that does not have an implicit commit, except `CREATE TEMPORARY TABLE...SELECT` (*This includes all "pure DML": INSERT, UPDATE, DELETE, REPLACE, TRUNCATE, SELECT, DO, CALL, EXECUTE, LOAD DATA INFILE, and BINLOG. It also includes CREATE TEMPORARY TABLE without SELECT, and DROP TEMPORARY TABLE. CREATE TEMPORARY TABLE...SELECT is handled in the next subsection*).

- Before executing S, determine unsafeness:

(R-unsafe-nontransactional) If S either makes N-writes or reads from an N-table, and either S or a previous statement in the same transaction reads or writes to a T-table then S is marked unsafe.

(R-unsafe-self-logging) If either S or a previous statement in the same transaction reads or writes to a self-logging table, then S is marked unsafe.

- When logging S, determine where to log it by applying the following rules in order:

(R-log-statement-format) If S is to be logged in statement format (*i.e.*, if one of the following holds: (1) @@session.binlog_format=STATEMENT; (2) @@session.binlog_format=MIXED and S is safe; (3) S is of DDL type, *i.e.*, CREATE TEMPORARY TABLE):

1. If S produces an error and does not do any N-write, do not log.
2. Otherwise, if either S or any previous statement in the same transaction reads or writes in any T-tables, log to TC.
3. Otherwise, log to SC.

(R-log-row-format) If S is to be logged in row format (*i.e.*, if S is DML and one of the following holds: (1) @@session.binlog_format=ROW; (2) @@session.binlog_format=MIXED and S is unsafe)

1. Do not log row events that write to temporary tables.
2. Log row events that write to non-temporary N-tables to SC.
3. Log row events that write to non-temporary T-tables to TC, except rows that are rolled back due to an error. (*Note: if there is an error, rows written to a T-table are kept if there are subsequent rows written to an N-table.*)

(R-flush-SC) At the end of S, write BEGIN + SC + COMMIT to the binlog and clear the SC.

- At end of transaction:

(R-log-commit) At COMMIT or implicit commit, where all XA tables in the transaction succeed in the "prepare" phase:

1. If the TC is non-empty, write BEGIN + TC + COMMIT to the binlog.
2. If the TC is empty, do nothing.

(R-log-rollback) At ROLLBACK; or at COMMIT or implicit commit where some XA table fails in the "prepare" phase:

1. If the TC contains any N-write, write BEGIN + TC + ROLLBACK to the binlog.
2. If the TC does not contain any N-write, do nothing.

(R-log-rollback-to-savepoint) At ROLLBACK TO SAVEPOINT:

1. If the TC contains any N-write after the savepoint, write ROLLBACK TO SAVEPOINT to the TC.
2. Otherwise, clear the part of the TC that starts at the savepoint and extends to the end of the TC. ([Bug#47327](#) breaks this rule)

(R-clear-TC) Clear the TC at the end of the transaction.

Rules for CREATE [TEMPORARY] TABLE...SELECT

First, unsafeness is determined as above (R-unsafe-transaction). Then the logging format is decided. Then the following rules apply.

(R-log-create-select-statement-format) If logging in statement format (*i.e.*, one of the following holds: (1) `@ @session.binlog_format=STATEMENT`; (2) `@ @session.binlog_format=MIXED` and statement is safe):

1. If there is an error, do not write anything.
2. If there is no error and the `TEMPORARY` keyword is used, write the entire `CREATE...SELECT` statement to the TC.
3. If there is no error and the `TEMPORARY` keyword is not used, write the entire `CREATE...SELECT` directly to the binlog.

(R-log-create-select-row-format) If logging in row format (*i.e.*, one of the following holds: (1) `@ @session.binlog_format=ROW`; (2) `@ @session.binlog_format=MIXED` and statement is unsafe):

1. If the `TEMPORARY` keyword is used, do not write anything.
2. If the `TEMPORARY` keyword is not used, write `CREATE TABLE` (without select) + `BEGIN` + row events + `COMMIT` to the TC. If there is an error, clear the TC; otherwise flush the TC to the binlog at the end of the statement and then clear the TC. (Note: currently [Bug#47899](#) breaks this rule)

Note: this breaks **D-rpl-correct** rule, because the slave will have an intermediate state that never existed on the master (namely, a state where the new table exists and is empty).

Rules for Committing Statements, Except `CREATE [TEMPORARY] TABLE...SELECT`

(R-log-commit-statement) All other statements that have a pre-commit are written directly to the binlog. (Note: this is semantically equivalent to writing it to the SC and flushing the SC. However, due to requirements by NDB (which have not been clarified), we write directly to the binlog.)

Logging Updates to `auto_increment` Columns

Status of this subsection: not started 2009-10-21

Logging Access to Variables and Functions

Status of this subsection: not started 2009-10-21

User variables: User variables (`@variable`) are logged as

User-defined functions

Server variables

Built-in functions

Other Unsafe Statements

Status of this subsection: not started 2009-10-21

INSERT DELAYED

LIMIT

System tables

`binlog_row_image`

Status of this subsection: finished, not reviewed, not fully implemented 2009-10-21

The sets of columns recorded in the BI and AI are determined by the value of `binlog_row_image`. To specify the sets of columns, we define the PKE (for Primary Key Equivalent), as follows:

- If a PK exists, the PKE is equal to the PK.
- Otherwise, if there exists a UK where all columns have the NOT NULL attribute, then that is the PKE (if there are more than one such UKs, then one is chosen arbitrarily).
- Otherwise, the PKE is equal to the set of all columns.

The set of columns included in the BI and AI are defined as in the following tables:

- write event

binlog_row_image	Before image	After image
minimal	-	All columns where a value was specified, and the autoincrement column if there is one
noblob	-	All columns where a value was specified, and the autoincrement column if there is one, and all non-blob columns
full	-	All columns

- update event

binlog_row_image	Before image	After image
minimal	PKE	All columns where a value was specified
noblob	PKE + all non-blob columns	All columns where a value was specified, and all non-blob columns
full	All columns	All columns

- delete event

binlog_row_image	Before image	After image
minimal	PKE	-
noblob	PKE + all non-blob columns	-
full	All columns	-

Cf. WL#5092.

Replication Locks

Status of this subsection: not started 2009-10-21

Chapter 20. The Binary Log

Table of Contents

Binary Log Overview	201
High-Level Binary Log Structure and Contents	202
Source Files Related to the Binary Log	203
Source File Archaeological Notes	203
Generating Browsable Binary Log Information	204
Event Classes and Types	204
Event Class Archaeological Notes	208
Event Meanings	209
Event Structure	211
Event Content-Writing Conventions	213
Event Header Fields	214
Event Data Fields (Event-Specific Information)	218
Binary Log Versions	219
Determining the Binary Log Version	221
Ensuring Compatibility of Future Binary Log Versions	222
Event Data for Specific Event Types	223
LOAD DATA INFILE Events	235
Row-Based Binary Logging	238
Additional Resources	239

This chapter describes the MySQL binary log, which contains information about data modifications made to a MySQL server instance.

Most of the information here applies equally to the relay log used on replication slave servers because a relay log has the same format as a binary log.

Binary Log Overview

[Some information information in this section is derived from [The Binary Log](#), in the MySQL Reference Manual.]

The binary log is a set of log files that contain information about data modifications made to a MySQL server instance. The log is enabled by starting the server with the `--log-bin` option.

The binary log was introduced in MySQL 3.23.14. It contains all statements that update data. It also contains statements that potentially could have updated it (for example, a `DELETE` which matched no rows), unless row-based logging is used. Statements are stored in the form of "events" that describe the modifications. The binary log also contains information about how long each statement took that updated data.

The binary log also contains some other metadata, including:

- Information about the state of the server that is needed to reproduce statements correctly
- Error codes
- Metadata needed for the maintenance of the binary log itself (for example, rotate events)

The binary log is a trace of changes of the server's global state generated during its operation. The events that it contains describe changes of this state. More precisely, binary log events describe actions that can be used to reproduce the same changes of global state which have happened on server.

The binary log has two important purposes:

- For replication, the binary log is used on master replication servers as a record of the statements to be sent to slave servers. Many details of binary log format and handling are specific to this purpose. The master server sends the events contained in its binary log to its slaves, which execute those events to make the same data changes that were made on the master. A slave stores events received from the master in its relay log until they can be executed. The relay log has the same format as the binary log.
- Certain data recovery operations require use of the binary log. After a backup file has been restored, the events in the binary log that were recorded after the backup was made are re-executed. These events bring databases up to date from the point of the backup.

There are two types of binary logging:

- Statement-based logging: Events contain SQL statements that produce data changes (inserts, updates, deletes)
- Row-based logging: Events describe changes to individual rows

Mixed logging uses statement-based logging by default but switches to row-based logging automatically as necessary.

Row-based (and mixed) logging is available beginning with MySQL 5.1.

For more details about row-based logging, see [Row-Based Binary Logging](#).

The `mysqlbinlog` utility can be used to print binary or relay log contents in readable form.

High-Level Binary Log Structure and Contents

The binary log is a set of files that contain information about data modifications made to a MySQL server instance.

- The log consists of a set of binary log files, plus an index file.
- Each log file contains a 4-byte magic number, followed by a set of events that describe data modifications:
 - The magic number bytes are 0xfe 0x62 0x69 0x6e = 0xfe 'b' 'i' 'n' (this is the `BINLOG_MAGIC` constant in `log_event.h`).
 - Each event contains header bytes followed by data bytes:
 - The header bytes provide information about the type of event, when it was generated, by which server, and so forth.
 - The data bytes provide information specific to the type of event, such as a particular data modification.
 - The first event is a descriptor event that describes the format version of the file (the format used to write events in the file).
 - The remaining events are interpreted according to the version.
 - The final event is a log-rotation event that specifies the next binary log filename.
- The index file is a text file that lists the current binary log files.

The details about event structure have changed over time, which gives rise to different versions of the binary log format. Currently, there are three binary log format versions, described in [#Binary Log Versions](#).

The log files are sequentially numbered using a `.NNNNNN` suffix. The index file has a suffix of `.index`. All files share a common basename. The default binary log file-naming basename is `"HOSTNAME-bin"`. With the default basename, the binary log has files with names like this:

```
...HOSTNAME-bin.0000101
HOSTNAME-bin.0000102
HOSTNAME-bin.0000103
...
HOSTNAME-bin.index
```

Relay log file naming is similar to that of the binary log files. The default relay log file-naming basename is `"HOSTNAME-relay"`. With the default basename, the relay log has files with names like this:

```
...HOSTNAME-relay.0000101
HOSTNAME-relay.0000102
HOSTNAME-relay.0000103
...
HOSTNAME-relay.index
```

Source Files Related to the Binary Log

This section describes the files in a MySQL source tree that are most relevant to binary log processing.

sql directory:

- `log.h/log.cc`: The high-level binary logging mechanism for organizing events into a sequence so that it becomes a binary log. Routines for creating, writing, deleting binary log files.
- `log_event.h/log_event.cc`: The low-level binary logging mechanism for serializing of values into records. The `Log_event` class and subclasses, for creating, writing, reading, printing, and applying events of every event type. Reading and writing here is at a low level, that is, serializing values into records.
- `rpl_constants.h`: Contains codes for `INCIDENT_EVENT` incident types.
- `slave.cc`: Contains some logic for processing binary logs in a replication slave (the IO and SQL threads).
- `rpl_injector.h/rpl_injector.cc`: Contains the injector class that allows external insertions into the binary log. This is used for cluster replication binary logging.
- `rpl_record.h/rpl_record.cc`: Utilities for encoding and decoding table rows into and out of the format used by row events.
- `rpl_tblmap.h/rpl_tblmap.cc`: Contains a mapping from numbers to tables. The mapping is used by the row logging system to identify tables.
- `rpl_utility.h/rpl_utility.cc`: Contains auxiliary classes and functions used for `Table_map_events`, and also an auxiliary class for smart pointers
- `sql_binlog.cc`: Code to execute `BINLOG` statements (the base64-encoded values that `mysqlbinlog` prints when it sees row events).
- `sql_base.cc`: The function `decide_logging_format()` that determines whether statements should be written to the binary log using row-based or statement-based format.

client directory:

- `mysqlbinlog.cc`: The source for the `mysqlbinlog` utility that reads binary log files and displays them in text format. It shares some event-interpretation code with the server.

Source File Archaeological Notes

`log_event_old.h/log_event_old.cc` (present in MySQL 5.1.18 and up): Classes to read and execute old versions of the row log events:

- [Write_rows_log_event_old](#)
- [Update_rows_log_event_old](#)
- [Delete_rows_log_event_old](#)

Prior to MySQL 5.1.17, those classes were known as:

- [Write_rows_log_event](#)
- [Update_rows_log_event](#)
- [Delete_rows_log_event](#)

For information about the relationship of the [Xxx_rows_log_event](#) and [Xxx_row_log_event_old](#) classes, see [#Event Classes and Types](#).

Generating Browseable Binary Log Information

Source files in the `sql` directory of a MySQL source tree or source distribution contain comments that can be processed with doxygen to generate HTML files that describe classes, files, and so forth. Those files can be viewed in your Web browser.

To generate the HTML files and view information about the classes related to binary logging, do this:

1. In your MySQL source tree, change location to the `sql` directory:

```
shell> cd sql
```

2. Run doxygen to generate the HTML files. These commands create a subdirectory named `html` containing the HTML output:

```
shell> doxygen -g
```

```
shell> doxygen
```

3. To view the top-level index page, load the [html/index.html](#) file into your browser.
4. To view the classes for binary logging, load the [html/class_log__event.html](#) page. The [Log_event](#) class is the main event class, and the others inherit from [Log_event](#).

The pages also contain links that take you to other related pages. For example, to navigate from [index.html](#) to [class_log__event.html](#), click on the [Classes](#) tab. On the next page, scroll down to [Log_event](#) and click on it.

Event Classes and Types

Internally, the server uses C++ classes to represent binary log events. Prototypes are in [log_event.h](#). Code for methods of these classes is in [log_event.cc](#).

[Log_event](#) is the base class. Other more specific event subclasses are derived from it. Type codes are associated with subclasses because class instance contents are written to binary or relay logs or are sent over the network from master to slave. In those contexts, an event is just a sequence of bytes, not a class structure, so a type code is needed to allow recognition of the event type from the byte sequence.

An event byte sequence has a header part and a data part. The type code appears in the header part of each event.

The possible type codes for events are listed in the [Log_event_type](#) enumeration:

```
enum Log_event_type {
    UNKNOWN_EVENT= 0,
    START_EVENT_V3= 1,
    QUERY_EVENT= 2,
```

```

STOP_EVENT= 3,
ROTATE_EVENT= 4,
INTVAR_EVENT= 5,
LOAD_EVENT= 6,
SLAVE_EVENT= 7,
CREATE_FILE_EVENT= 8,
APPEND_BLOCK_EVENT= 9,
EXEC_LOAD_EVENT= 10,
DELETE_FILE_EVENT= 11,
NEW_LOAD_EVENT= 12,
RAND_EVENT= 13,
USER_VAR_EVENT= 14,
FORMAT_DESCRIPTION_EVENT= 15,
XID_EVENT= 16,
BEGIN_LOAD_QUERY_EVENT= 17,
EXECUTE_LOAD_QUERY_EVENT= 18,
TABLE_MAP_EVENT = 19,
PRE_GA_WRITE_ROWS_EVENT = 20,
PRE_GA_UPDATE_ROWS_EVENT = 21,
PRE_GA_DELETE_ROWS_EVENT = 22,
WRITE_ROWS_EVENT = 23,
UPDATE_ROWS_EVENT = 24,
DELETE_ROWS_EVENT = 25,
INCIDENT_EVENT= 26,
HEARTBEAT_LOG_EVENT= 27,
ENUM_END_EVENT
/* end marker */
};

```

The `INTVAR_EVENT` type has "subtypes," listed in the `Int_event_type` enumeration:

```

enum Int_event_type {
    INVALID_INT_EVENT = LAST_INSERT_ID_EVENT = INSERT_ID_EVENT = 2
};

```

The following table summarizes the relationship between event classes and type codes. Each class is derived from `Log_event` unless otherwise indicated. As can be seen, an event class is associated with a single type code in most cases, although there are some exceptions:

- Some classes are not associated with any type code because they are used only as a base class for which to derive subclasses or because they are never written to binary or relay logs or sent from master to slave. For example, `Log_event` has no type code because it is used only as a base class.
- A class may be associated with multiple type codes: `Load_log_event` may contain a type code of either `LOAD_EVENT` or `NEW_LOAD_EVENT`.

Value	Type Code	Class
		Log_event <ul style="list-style-type: none"> • Base class for most other classes
		Muted_query_log_event <ul style="list-style-type: none"> • Added in 5.0.23 • Removed in 6.0.4
		Rows_log_event <ul style="list-style-type: none"> • Added in 5.1.5 • Base class for <ul style="list-style-type: none"> Write_rows_log_event, Update_rows_log_event, Delete_rows_log_event
		Old_rows_log_event

		<ul style="list-style-type: none"> • Added in 5.1.22 • Base class for <code>Write_rows_log_event_old</code>, <code>Update_rows_log_event_old</code>, <code>Delete_rows_log_event_old</code>
0	UNKNOWN_EVENT	Unknown_log_event
1	START_EVENT_V3	Start_log_event_v3 <ul style="list-style-type: none"> • Renamed from <code>START_EVENT/Start_log_event</code> in 5.0.0 • Base class for <code>Format_description_log_event</code>
2	QUERY_EVENT	Query_log_event <ul style="list-style-type: none"> • Base class for <code>Execute_load_query_log_event</code>
3	STOP_EVENT	Stop_log_event
4	ROTATE_EVENT	Rotate_log_event
5	INTVAR_EVENT	Intvar_log_event
6	LOAD_EVENT	Load_log_event <ul style="list-style-type: none"> • Base class for <code>Create_file_log_event</code>
7	SLAVE_EVENT	Slave_log_event <ul style="list-style-type: none"> • Added in 4.0.0
8	CREATE_FILE_EVENT	Create_file_log_event <ul style="list-style-type: none"> • Added in 4.0.0 • Derived from <code>Load_log_event</code>
9	APPEND_BLOCK_EVENT	Append_block_log_event <ul style="list-style-type: none"> • Added in 4.0.0 • Base class for <code>Begin_load_query_log_event</code>
10	EXEC_LOAD_EVENT	Execute_load_log_event <ul style="list-style-type: none"> • Added in 4.0.0
11	DELETE_FILE_EVENT	Delete_file_log_event <ul style="list-style-type: none"> • Added in 4.0.0
12	NEW_LOAD_EVENT	Load_log_event <ul style="list-style-type: none"> • Added in 4.0.0
13	RAND_EVENT	Rand_log_event <ul style="list-style-type: none"> • Added in 4.0.5
14	USER_VAR_EVENT	User_var_log_event

		<ul style="list-style-type: none"> Added in 4.1.0
15	FORMAT_DESCRIPTION_EVENT	Format_description_log_event <ul style="list-style-type: none"> Added in 5.0.0 Derived from Start_log_event_v3
16	XID_EVENT	Xid_log_event <ul style="list-style-type: none"> Added in 5.0.3
17	BEGIN_LOAD_QUERY_EVENT	Begin_load_query_log_event <ul style="list-style-type: none"> Added in 5.0.3 Derived from Append_block_log_event
18	EXECUTE_LOAD_QUERY_EVENT	Execute_load_query_log_event <ul style="list-style-type: none"> Added in 5.0.3 Derived from Query_log_event
19	TABLE_MAP_EVENT	Table_map_log_event <ul style="list-style-type: none"> Added in 5.1.5
20	PRE_GA_WRITE_ROWS_EVENT	Write_rows_log_event_old <ul style="list-style-type: none"> Added in 5.1.5 as WRITE_ROWS_EVENT/Write_rows_log_event and derived from Rows_log_event Renamed in 5.1.18 to PRE_GA_WRITE_ROWS_EVENT/Write_rows_log_event and derived from Write_rows_log_event As of 5.1.22, derived from Old_rows_log_event
21	PRE_GA_UPDATE_ROWS_EVENT	Update_rows_log_event_old <ul style="list-style-type: none"> Added in 5.1.5 as UPDATE_ROWS_EVENT/Update_rows_log_event and derived from Rows_log_event Renamed in 5.1.18 to PRE_GA_UPDATE_ROWS_EVENT/Update_rows_log_event and derived from Update_rows_log_event As of 5.1.22, derived from Old_rows_log_event
22	PRE_GA_DELETE_ROWS_EVENT	Delete_rows_log_event_old <ul style="list-style-type: none"> Added in 5.1.5 as DELETE_ROWS_EVENT/Delete_rows_log_event

		and derived from Rows_log_event <ul style="list-style-type: none"> Renamed in 5.1.18 to PRE_GA_DELETE_ROWS_EVENT/Delete_r and derived from Delete_rows_log_event As of 5.1.22, derived from Old_rows_log_event
23	WRITE_ROWS_EVENT	Write_rows_log_event <ul style="list-style-type: none"> Derived from Rows_log_event Renumbered in 5.1.18 from 20 to 23
24	UPDATE_ROWS_EVENT	Update_rows_log_event <ul style="list-style-type: none"> Derived from Rows_log_event Renumbered in 5.1.18 from 21 to 24
25	DELETE_ROWS_EVENT	Delete_rows_log_event <ul style="list-style-type: none"> Derived from Rows_log_event Renumbered in 5.1.18 from 22 to 25
26	INCIDENT_EVENT	Incident_log_event <ul style="list-style-type: none"> Added in 5.1.18
27	HEARTBEAT_LOG_EVENT	Heartbeat_log_event <ul style="list-style-type: none"> Added in 6.0.5

Event Class Archaeological Notes

Despite the "V3" in the type code name, [START_EVENT_V3](#) currently is used as the type code not only for v3 start events, but also for v1 start events. The original symbol for type code 1 was [START_EVENT](#) in the format now known as v1. Later, when v3 was developed, type code 1 was reused and the symbol associated with it was renamed from [START_EVENT](#) to [START_EVENT_V3](#). The start events for both v1 and v3 therefore have a type code of 1, although the event structures differ and must be distinguished by examining their contents.

Up to MySQL 5.1.17, event type codes from 20 to 22 were associated with symbols and classes as follows:

Value	Type Code	Class
20	WRITE_ROWS_EVENT	Write_rows_log_event
21	UPDATE_ROWS_EVENT	Update_rows_log_event
22	DELETE_ROWS_EVENT	Delete_rows_log_event

In 5.1.18, the symbols and classes were renamed:

Value	Type Code	Class
20	PRE_GA_WRITE_ROWS_EVENT	Write_rows_log_event_old
21	PRE_GA_UPDATE_ROWS_EVENT	Update_rows_log_event_old
22	PRE_GA_DELETE_ROWS_EVENT	Delete_rows_log_event_old

Also in 5.1.18, the original symbols were reused with different values and new implementations of the classes that used the original names:

Value	Type Code	Class
23	WRITE_ROWS_EVENT	Write_rows_log_event
24	UPDATE_ROWS_EVENT	Update_rows_log_event
25	DELETE_ROWS_EVENT	Delete_rows_log_event

Events with type codes 20 to 22 are obsolete now and appear only in binary logs created by servers from MySQL 5.1.5 to 5.1.17.

Event Meanings

The following descriptions briefly summarize the meaning of each event type:

- [UNKNOWN_EVENT](#)

This event type should never occur. It is never written to a binary log. If an event is read from a binary log that cannot be recognized as something else, it is treated as [UNKNOWN_EVENT](#).

- [START_EVENT_V3](#)

A descriptor event that is written to the beginning of the each binary log file. (In MySQL 4.0 and 4.1, this event is written only to the first binary log file that the server creates after startup.) This event is used in MySQL 3.23 through 4.1 and superseded in MySQL 5.0 by [FORMAT_DESCRIPTION_EVENT](#).

- [QUERY_EVENT](#)

Written when an updating statement is done.

- [STOP_EVENT](#)

Written when `mysqld` stops.

- [ROTATE_EVENT](#)

Written when `mysqld` switches to a new binary log file. This occurs when someone issues a [FLUSH LOGS](#) statement or the current binary log file becomes too large. The maximum size is determined by `max_binlog_size`.

- [INTVAR_EVENT](#)

Written every time a statement uses an [AUTO_INCREMENT](#) column or the [LAST_INSERT_ID\(\)](#) function; precedes other events for the statement. This is written only before a [QUERY_EVENT](#) and is not used with row-based logging. An [INTVAR_EVENT](#) is written with a "subtype" in the event data part:

- [INSERT_ID_EVENT](#) indicates the value to use for an [AUTO_INCREMENT](#) column in the next statement.
- [LAST_INSERT_ID_EVENT](#) indicates the value to use for the [LAST_INSERT_ID\(\)](#) function in the next statement.
- [LOAD_EVENT](#)

Used for `LOAD DATA INFILE` statements in MySQL 3.23. See [LOAD DATA INFILE Events](#).

- `SLAVE_EVENT`

Not used.

- `CREATE_FILE_EVENT`

Used for `LOAD DATA INFILE` statements in MySQL 4.0 and 4.1. See [LOAD DATA INFILE Events](#).

- `APPEND_BLOCK_EVENT`

Used for `LOAD DATA INFILE` statements as of MySQL 4.0. See [LOAD DATA INFILE Events](#).

- `EXEC_LOAD_EVENT`

Used for `LOAD DATA INFILE` statements in 4.0 and 4.1. See [LOAD DATA INFILE Events](#).

- `DELETE_FILE_EVENT`

Used for `LOAD DATA INFILE` statements as of MySQL 4.0. See [LOAD DATA INFILE Events](#).

- `NEW_LOAD_EVENT`

Used for `LOAD DATA INFILE` statements in MySQL 4.0 and 4.1. See [LOAD DATA INFILE Events](#).

- `RAND_EVENT`

Written every time a statement uses the `RAND()` function; precedes other events for the statement. Indicates the seed values to use for generating a random number with `RAND()` in the next statement. This is written only before a `QUERY_EVENT` and is not used with row-based logging.

- `USER_VAR_EVENT`

Written every time a statement uses a user variable; precedes other events for the statement. Indicates the value to use for the user variable in the next statement. This is written only before a `QUERY_EVENT` and is not used with row-based logging.

- `FORMAT_DESCRIPTION_EVENT`

A descriptor event that is written to the beginning of the each binary log file. This event is used as of MySQL 5.0; it supersedes `START_EVENT_V3`.

- `XID_EVENT`

Generated for a commit of a transaction that modifies one or more tables of an XA-capable storage engine. Normal transactions are implemented by sending a `QUERY_EVENT` containing a `BEGIN` statement and a `QUERY_EVENT` containing a `COMMIT` statement (or a `ROLLBACK` statement if the transaction is rolled back).

- `BEGIN_LOAD_QUERY_EVENT`

Used for `LOAD DATA INFILE` statements as of MySQL 5.0. See [LOAD DATA INFILE Events](#).

- `EXECUTE_LOAD_QUERY_EVENT`

Used for `LOAD DATA INFILE` statements as of MySQL 5.0. See [LOAD DATA INFILE Events](#).

- `TABLE_MAP_EVENT`

Used for row-based binary logging. This event precedes each row operation event. It maps a table definition to a number, where the table definition consists of database and table names and column definitions. The purpose of this event is to enable replication when a table has different

definitions on the master and slave. Row operation events that belong to the same transaction may be grouped into sequences, in which case each such sequence of events begins with a sequence of `TABLE_MAP_EVENT` events: one per table used by events in the sequence.

- `PRE_GA_WRITE_ROWS_EVENT`

Obsolete version of `WRITE_ROWS_EVENT`.

- `PRE_GA_UPDATE_ROWS_EVENT`

Obsolete version of `UPDATE_ROWS_EVENT`.

- `PRE_GA_DELETE_ROWS_EVENT`

Obsolete version of `DELETE_ROWS_EVENT`.

- `WRITE_ROWS_EVENT`

Used for row-based binary logging. This event logs inserts of rows in a single table.

- `UPDATE_ROWS_EVENT`

Used for row-based binary logging. This event logs updates of rows in a single table.

- `DELETE_ROWS_EVENT`

Used for row-based binary logging. This event logs deletions of rows in a single table.

- `INCIDENT_EVENT`

Used to log an out of the ordinary event that occurred on the master. It notifies the slave that something happened on the master that might cause data to be in an inconsistent state.

- `HEARTBEAT_LOG_EVENT`

Sent by a master to a slave to let the slave know that the master is still alive. Not written to log files.

Event Structure

This section describes the general properties of events as byte sequences as they are written to binary or relay log files.

All events have a common general structure consisting of an event header followed by event data:

```

+=====+
| event header |
+=====+
| event data   |
+=====+
```

The details about what goes in the header and data parts have changed over time, which gives rise to different versions of the binary log format:

- v1: Used in MySQL 3.23
- v3: Used in MySQL 4.0.2 through 4.1
- v4: Used in MySQL 5.0 and up

A v2 format was used briefly (in early MySQL 4.0.x versions), but it is obsolete and no longer supported.

Some details of event structure are invariant across binary log versions; others depend on the version. Within any given version, different types of events vary in the structure of the data part.

The first event in a log file is special. It is a descriptor event that provides information such as the binary log version and the server version. The information in the descriptor event enables programs to determine which version of the binary log format applies to the file so that the remaining events in the file can be properly read and interpreted.

For details about the initial descriptor event and how to use it to determine the format of a binary log file, see [Binary Log Versions](#). For additional information about other types of events, see [Event Data for Specific Event Types](#).

The following event diagrams contain field descriptions written using these conventions:

- A field line has a name describing the contents of the field.
- The name is followed by two numbers in *offset : length* format, where *offset* is the 0-based offset (position) of the field within the event and *length* is the length of the field. Both values are given in bytes.

The overall structure for events in the different binary versions is shown here. The following sections describe the header and data parts in more detail.

v1 event structure:

+-----+			
event	timestamp	0 : 4	
header	+-----+		
	type_code	4 : 1	
	+-----+		
	server_id	5 : 4	
	+-----+		
	event_length	9 : 4	
+-----+			
event	fixed part	13 : y	
data	+-----+		
	variable part		
+-----+			

header length = 13 bytes

data length = (event_length - 13) bytes

y is specific to the event type.

v3 event structure:

+-----+			
event	timestamp	0 : 4	
header	+-----+		
	type_code	4 : 1	
	+-----+		
	server_id	5 : 4	
	+-----+		
	event_length	9 : 4	
	+-----+		
	next_position	13 : 4	
	+-----+		
	flags	17 : 2	
+-----+			
event	fixed part	19 : y	
data	+-----+		
	variable part		
+-----+			

header length = 19 bytes

data length = (event_length - 19) bytes

y is specific to the event type.

v4 event structure:

+=====+		
event	timestamp	0 : 4
header	+-----+	
	type_code	4 : 1
	+-----+	
	server_id	5 : 4
	+-----+	
	event_length	9 : 4
	+-----+	
	next_position	13 : 4
	+-----+	
	flags	17 : 2
	+-----+	
	extra_headers	19 : x-19
+=====+		
event	fixed part	x : y
data	+-----+	
	variable part	
+=====+		

header length = x bytes

data length = (event_length - x) bytes

fixed data length = y bytes variable data length = (event_length - (x + y)) bytes

x is given by the header_length field in the format description event (FDE). Currently, x is 19, so the extra_headers field is empty.

y is specific to the event type, and is given by the FDE. The fixed-part length is the same for all events of a given type, but may vary for different event types.

The fixed part of the event data is sometimes referred to as the "post-header" part. The variable part is sometimes referred to as the "payload" or "body."

For information about how to use the FDE to interpret v4 events, see [Binary Log Formats](#).

Event Content-Writing Conventions

Event contents are written using these conventions:

- Numbers are written in little-endian format (least significant byte first), unless otherwise indicated.
- Values that represent positions or lengths are given in bytes and should be considered unsigned.
- Some numbers are written as Packed Integers. The format is described later in this section.
- Strings are written in varying formats:
 - A string may be written to a fixed-length field and null-padded (with 0x00 bytes) on the right.
 - A variable-length string may be preceded by a length field that indicates the length of the string.
 - Some variable-length strings are null-terminated; others are not. The descriptions for individual string fields indicates which is the case.
 - For null-terminated strings that are preceded by a length field, the length does not include the terminating null byte, unless otherwise indicated.
 - If there is a variable-length string at the end of an event and no length field precedes it, its length may be determined as the event length minus the length of the other fields in the event.

Some events use Packed Integers, a special format for efficient representation of unsigned integers. A Packed Integer has the capacity of storing up to 8-byte integers, while small integers still can use 1,

3, or 4 bytes. The value of the first byte determines how to read the number, according to the following table.

First byte	Format
0-250	The first byte is the number (in the range 0-250). No additional bytes are used.
252	Two more bytes are used. The number is in the range 251-0xffff.
253	Three more bytes are used. The number is in the range 0xffff-0xffffffff.
254	Eight more bytes are used. The number is in the range 0xffffffff-0xffffffffffffffff.

Packed Integer format derives from the "Length Coded Binary" representation used in the MySQL client/server network protocol (see [Chapter 15, MySQL Client/Server Protocol](#)). That representation allows a first byte value of 251 to represent the SQL NULL value, but 251 is apparently unused for Packed Integers in the binary log.

Event Header Fields

Each event starts with a header of size `LOG_EVENT_HEADER_LEN`. The value of this constant is 13 in MySQL 3.23 (v1 format), and 19 in MySQL 4.0 and up (v3 format and up). The value is larger as of 4.0 because next position and flags fields were added to the header format then:

- v1: 13 bytes: timestamp + type code + server ID + event length
- v3: 19 bytes: v1 fields + next position + flags
- v4: 19 bytes or more: v3 fields + possibly other information

The header for any version is a superset of the header for all earlier versions:

- The first 13 bytes for v3 and v4 are the same as those for v1.
- The first 19 bytes for v4 are the same as those for v3.

Because the event header in a newer binary log format starts with the header of the old formats, headers in different formats are backward compatible.

v1 event header:

```

+-----+
| timestamp      0 : 4 |
+-----+
| type_code      4 : 1 |
+-----+
| server_id      5 : 4 |
+-----+
| event_length    9 : 4 |
+-----+

```

The 13 bytes of the v1 header are also present in the header of all subsequent binary log versions.

v3 event header:

```

+-----+
| timestamp      0 : 4 |
+-----+
| type_code      4 : 1 |
+-----+
| server_id      5 : 4 |
+-----+
| event_length    9 : 4 |
+-----+

```

```

+-----+
| next_position    13 : 4 |
+-----+
| flags           17 : 2 |
+-----+

```

Compared to v1, the header in v3 and up contains two additional fields, for a total of 19 bytes.

v4 event header:

```

+-----+
| timestamp        0 : 4 |
+-----+
| type_code        4 : 1 |
+-----+
| server_id        5 : 4 |
+-----+
| event_length     9 : 4 |
+-----+
| next_position    13 : 4 |
+-----+
| flags           17 : 2 |
+-----+
| extra_headers    19 : x-19 |
+-----+

```

The v4 format includes an `extra_headers` field; this is a mechanism for adding extra fields to the header without breaking the format. This extension mechanism is implemented via the format description event that appears as the first event in the file. (See [Binary Log Versions](#) for details.) Currently, $x = 19$, so the `extra_headers` field is empty; thus, the v4 header is the same as the v3 header.

Note: The `extra_headers` field does not appear in the `FORMAT_DESCRIPTION_EVENT` or `ROTATE_EVENT` header.

The offsets of several fields within the event header are available as constants in `log_event.h`:

- `EVENT_TYPE_OFFSET = 4`
- `SERVER_ID_OFFSET = 5`
- `EVENT_LEN_OFFSET = 9`
- `LOG_POS_OFFSET = 13`
- `FLAGS_OFFSET = 17`

The header fields contain the following information:

- **timestamp**

4 bytes. This is the time at which the statement began executing. It is represented as the number of seconds since 1970 (UTC), like the `TIMESTAMP` SQL data type.


- **type_code**

1 byte. The type of event. 1 means `START_EVENT_V3`, 2 means `QUERY_EVENT`, and so forth. These numbers are defined in the `enum Log_event_type` enumeration in `log_event.h`. (See [Event Classes and Types](#).)

- **server_id**

4 bytes. The ID of the `mysqld` server that originally created the event. It comes from the `server-id` option that is set in the server configuration file for the purpose of replication. The server ID enables endless loops to be avoided when circular replication is used (with option `--log-slave-updates` on). Suppose that M1, M2, and M3 have server ID values of 1, 2, and 3, and that they are replicating in

circular fashion: M1 is the master for M2, M2 is the master for M3, and M3 is that master for M1. The master/server relationships look like this:



A client sends an `INSERT` statement to M1. This is executed on M1 and written to its binary log with an event server ID of 1. The event is sent to M2, which executes it and writes it to its binary log; the event is still written with server ID 1 because that is the ID of the server that originally created the event. The event is sent to M3, which executes it and writes it to its binary log, still with server ID 1. Finally, the event is sent to M1, which sees server ID = 1 and understands this event originated from itself and therefore must be ignored.

- **event_length**

4 bytes. The total size of this event. This includes both the header and data parts. Most events are less than 1000 bytes, except when using `LOAD DATA INFILE` (where events contain the loaded file, so they can be big).

- **next_position** (not present in v1 format).

4 bytes. The position of the next event in the master's binary log. The format differs between binlogs and relay logs, and depending on the version of the server (and for relay logs, depending on the version of the master):

- binlog on a v3 server: Offset to the beginning of the event, counting from the beginning of the binlog file. In other words, equal to the value of `tell()` just before the event was written.

So the first event of the binlog has `next_position = 4`, and for events `n` and `n+1`, it holds that `next_position(n+1) = next_position(n) + event_length(n)`.

- relay log on a v3 server where the master uses v1: Probably 0, but I can't test this because I don't know how to run a 3.23 server.
- relay log on a v3 server where the master uses v3: Offset to the beginning of the event as it was in the master's binlog file, counting from the beginning of the master's binlog file.

The slave's relay log can be different from the master's binlog, so `next_position` can be different from the offset of the event in the relay log, counting from the beginning of the relay log file. However, if both event `n` and event `n+1` originate from the master, it holds that `next_position(n+1) = next_position(n) + event_length(n)`;

- binlog on a v4 server: Offset to the end of the event, counting from the beginning of the binlog file. In other words, equal to the value of `tell()` just after the event was written.

So the first event of the binlog has `next_position = 4 + event_length`, and for events number `n` and `n+1`, it holds that `next_position(n+1) = next_position(n) + event_length(n+1)`.

- relay log on a v4 server: Offset to the end of the event as it was in the master's binlog file, counting from the beginning of the master's binlog file.

The slave's relay log can be different from the master's binlog, so `next_position` can be different from the offset of the event in the relay log, counting from the beginning of the relay log file. However, if both event `n` and event `n+1` originate from the master, it holds that `next_position(n+1) = next_position(n) + event_length(n+1)`.

The `next_position` is used on the slave in two cases:

- for `SHOW SLAVE STATUS` to be able to show coordinates of the last executed event *in the master's coordinate system*.

- for `START SLAVE UNTIL MASTER_LOG_FILE=x, MASTER_LOG_POS=y`, so that the master's coordinates can be used.

In 5.0 and up, `next_position` is called "end_log_pos" in the output from `mysqlbinlog` and `SHOW BINLOG EVENTS`. In 4.1, `next_position` is called "log_pos" in the output from `mysqlbinlog` and "orig_log_pos" in the output from `SHOW BINLOG EVENTS`.

- **flags** (not present in v1 format)

2 bytes. The possible flag values are described at [Event Flags](#).

- **extra_headers** (not present in v1, v3 formats)

Variable-sized. The size of this field is determined by the format description event that occurs as the first event in the file. Currently, the size is 0, so, in effect, this field never actually occurs in any event. At such time as the size becomes non-zero, this field still will not appear in events of type `FORMAT_DESCRIPTION_EVENT` or `ROTATE_EVENT`.

Event Flags

Event headers for v3 format and up contain event flags in the two flag bytes at position `FLAGS_OFFSET` = 17. There are comments about these flags in `log_event.h`, in addition to the remarks in this section.

Current event flags:

- `LOG_EVENT_BINLOG_IN_USE_F` = 0x1 (New in 5.0.3)

Used to indicate whether a binary log file was closed properly. This flag makes sense only for `FORMAT_DESCRIPTION_EVENT`. It is set when the event is written to the log file. When the log file is closed later, the flag is cleared. (This is the only case when MySQL modifies an already written part of a binary log file).

- `LOG_EVENT_THREAD_SPECIFIC_F` = 0x4 (New in 4.1.0)

Used only by `mysqlbinlog` (not by the replication code at all) to be able to deal properly with temporary tables. `mysqlbinlog` displays events from the binary log in printable format, so that you can feed the output into `mysql` (the command-line interpreter), to achieve incremental backup recovery. But suppose that the binary log is as follows, where two simultaneous threads used temporary tables with the same name (which is allowed because temporary tables are visible only in the thread which created them):

```
<thread id 1>
CREATE TEMPORARY TABLE t (a INT);
<thread id 2>
CREATE TEMPORARY TABLE t (a INT);
```

In this case, simply feeding this into `mysql` will lead to a "table t already exists" error. This is why events that use temporary tables are marked with the flag, so that `mysqlbinlog` knows it has to set the `pseudo_thread_id` system variable before, like this:

```
SET PSEUDO_THREAD_ID=1;
CREATE TEMPORARY TABLE t (a INT);
SET PSEUDO_THREAD_ID=2;
CREATE TEMPORARY TABLE t (a INT);
```

This way there is no confusion for the server that receives these statements. Always printing `SET PSEUDO_THREAD_ID`, even when temporary tables are not used, would cause no bug, it would just slow down.

- `LOG_EVENT_SUPPRESS_USE_F` = 0x8 (New in 4.1.7)

Suppresses generation of a USE statement before the actual statement to be logged. This flag should be set for any event that does not need to have the default database set to function correctly,

such as CREATE DATABASE and DROP DATABASE. This flag should only be used in exceptional circumstances because it introduces a significant change in behavior regarding the replication logic together with the `--binlog-do-db` and `--replicate-do-db` options.

- `LOG_EVENT_UPDATE_TABLE_MAP_VERSION_F` = 0x10 (New in 5.1.4)

Causes the table map version internal to the binary log to be increased after the event has been written to the log.

Obsolete event flags:

- `LOG_EVENT_TIME_F` (obsolete as of 4.1.1). This flag was never set.
- `LOG_EVENT_FORCED_ROTATE_F` (obsolete as of 4.1.1). This flag was set in events of type `ROTATE_EVENT` on the master, but was not used for anything useful

They are now commented out in `log_event.h` and their values are available for reuse or have already been reused. (But see the associated comments in `log_event.h` for various cautions!)

Event Data Fields (Event-Specific Information)

The structure of an event's data part depends on the event type:

- In v1 and v3, the event type entirely determines the data format
- In v4, interpretation of the data part depends on the event type in conjunction with information from the format description event. This is because v4 allows for an `extra_headers` field, the size of which is defined in the format description event. In practice, the `extra_headers` field currently is empty.

The data part of an event consists of a fixed-size part and a variable-size part. Either or both parts may be empty, depending on the event type. (For example, a `STOP_EVENT` consists only of the header part; the fixed and variable data parts are both empty.)

The size of the event data part is the event size (contained in the header) minus the header size. The size of the fixed data part is a function of the event type. The size of the variable data part is the event size minus the size of the header, minus the size of the fixed data part.

The following principles hold across all events in a binary log file:

- The fixed part of the event data is the same size for all events of a given type.
- The variable part of the event data can differ in size among events of a given type.

For details about the fixed and variable parts of event data for different events, see [Event Data for Specific Event Types](#).

Event Data Field Notational Caveat

The fixed part of the event data goes under different names, depending on which source file, work log, bug report, etc. you are reading:

- Sometimes it is called the "fixed data" part, as in this discussion.
- Sometimes it is called the "post-headers" part.
- To make things notationally interesting, sometimes the fixed data part is referred to as the "event-specific headers" part of the event. That is, the word "header" is used in reference to a portion of the data part. One manifestation of this notational phenomenon appears in `log_event.h`, where you will find the symbol `LOG_EVENT_MINIMAL_HEADER_LEN` defined as 19 (the header length for v3 and v4), plus other symbols with names of the form `XXX_HEADER_LEN` for different event types.

The former symbol is the size of the event header (always 19). The latter symbols define the size of the fixed portion of the data part that is to be treated as the event-specific headers. For example, `ROTATE_HEADER_LEN` is 8 because a `ROTATE_EVENT` has an 8-byte field in the fixed data part that indicates the position in the next log file of the first event in that file.

The variable part of event data also goes under different names, such as the event "payload" or "body."

Binary Log Versions

There are several versions of the binary log file format:

- v1: Used in MySQL 3.23
- v3: Used in MySQL 4.0.2 though 4.1
- v4: Used in MySQL 5.0 and up

A v2 format was used briefly (in early MySQL 4.0.x versions), but it is obsolete and no longer supported.

Programs that process the binary log must be able to account for each of the supported binary log formats. This section describes how the server distinguishes each format to identify which one a binary log file uses. `mysqlbinlog` uses the same principles.

Important constants:

- `START_EVENT_V3` = 1
- `FORMAT_DESCRIPTION_EVENT` = 15
- `EVENT_TYPE_OFFSET` = 4
- `EVENT_LEN_OFFSET` = 9
- `ST_SERVER_VER_LEN` = 50

A binary log file begins with a 4-byte magic number followed by an initial descriptor event that identifies the format of the file.

- In v1 and v3, this event is called a "start event."
- In v4, it is called a "format description event."

Elsewhere you may see either term used generically to refer collectively to both types of event. This discussion uses "descriptor event" as the collective term.

The header and data parts of the descriptor event for each binary log format version are shown following. The diagrams use the same conventions as those described earlier in [Event Structure](#).

v1 start event (size = 69 bytes):

+=====+				
event	timestamp	0 : 4		
header	+-----+			
	type_code	4 : 1		= START_EVENT_V3 = 1
	+-----+			
	server_id	5 : 4		
	+-----+			
	event_length	9 : 4		= 69
+=====+				
event	binlog_version	13 : 2		= 1
data	+-----+			
	server_version	15 : 50		

```

+-----+
| create_timestamp 65 : 4 |
+-----+

```

v3 start event (size = 75 bytes):

```

+-----+
| event | timestamp      0 : 4 |
| header +-----+
|       | type_code       4 : 1 | = START_EVENT_V3 = 1
|       +-----+
|       | server_id      5 : 4 |
|       +-----+
|       | event_length  9 : 4 | = 75
|       +-----+
|       | next_position 13 : 4 |
|       +-----+
|       | flags       17 : 2 |
+-----+
| event | binlog_version 19 : 2 | = 3
| data  +-----+
|       | server_version 21 : 50 |
|       +-----+
|       | create_timestamp 71 : 4 |
+-----+

```

v4 format description event (size >= 91 bytes; the size is 76 + the number of event types):

```

+-----+
| event | timestamp      0 : 4 |
| header +-----+
|       | type_code       4 : 1 | = FORMAT_DESCRIPTION_EVENT = 15
|       +-----+
|       | server_id      5 : 4 |
|       +-----+
|       | event_length  9 : 4 | >= 91
|       +-----+
|       | next_position 13 : 4 |
|       +-----+
|       | flags       17 : 2 |
+-----+
| event | binlog_version 19 : 2 | = 4
| data  +-----+
|       | server_version 21 : 50 |
|       +-----+
|       | create_timestamp 71 : 4 |
|       +-----+
|       | header_length  75 : 1 |
|       +-----+
|       | post-header    76 : n | = array of n bytes, one byte per event
|       | lengths for all |   type that the server knows about
|       | event types    |
+-----+

```

In all binary log versions, the event data for the descriptor event begins with a common set of fields

- **binlog_version**

The binary log version number (1, 3, or 4).

- **server_version**

The server version as a string.

- **create_timestamp**

The creation timestamp, if non-zero, is the time in seconds when this event was created; it indicates the moment when the binary log was created. This field is actually of no value: If non-zero, it is redundant because it has the same value that is in the header timestamp.

Note: In practice, the creation timestamp field should be considered reserved for future use and programs should not rely on its value. This field may be commandeered in the future to serve another purpose.

The v4 format descriptor event data contains two additional fields that enable interpretation of other types of events:

- **header_length**

The length of the event header. This value includes the `extra_headers` field, so this header length - 19 yields the size of the `extra_headers` field.

Currently in v4, the header length (at offset 75) is 19, which means that in other events, no extra headers will follow the `flags` field. If in the future the header length is some value $x > 19$, then $x - 19$ extra header bytes will appear in other events in the `extra_headers` field following the `flags` field.

Note: The `FORMAT_DESCRIPTION_EVENT` itself contains no `extra_headers` field. Suppose that the FDE did have a `header_length` field after the `flags` field. That would introduce this problem:

- The value of x is given in the `header_length` field, which occurs in a position later than where the `extra_headers` field would be.
- Until you know the value of x , you cannot know the exact offset of the `header_length` field.

In other words, you would need to know x to find the `header_length` field, but you cannot know x until you read the `header_length` field. (A circular dependency.) This means that the event extensibility mechanism afforded by the FDE does not apply to the FDE itself, which therefore is not itself extensible.

- **post-header lengths**

The lengths for the fixed data part of each event. This is an array that provides post-header lengths for all events beginning with `START_EVENT_V3` (type code 1). The array includes no length for `UNKNOWN_EVENT` (type code 0).

Determining the Binary Log Version

Given any binary log file, the information in this section describes how to determine the format in which it is written.

Some important points about descriptor event formats:

- The v1 header fields are common to all formats. (v3 and v4 headers begin with the v1 header fields, and add `next_position` and `flags` fields.)
- The v3 and v4 headers contain the same fields. The data part for v3 and v4 differs, such that the v4 data part enables extensions to the format without having to modify the header.
- It would be possible to ascertain the binary log version simply by reading the two `binlog_version` bytes, were it not for the fact that these bytes occur at a different position in v1 compared to v3/v4 (position 13 versus 19). Therefore, it's necessary to determine whether the first event in a file represents a v1-format start event.

To determine the version of a binary log file, use the following procedure:

1) The file begins with a 4-byte magic number. Skip over that to get to the first event in the file (which in most cases is a start event or format description event).

2) From the first event, read two values:

- The 1-byte type code at position `EVENT_TYPE_OFFSET` (4) within the event.

- The 4-byte event length at position `EVENT_LEN_OFFSET` (9) within the event.

3) If the type code is not `START_EVENT_V3` or `FORMAT_DESCRIPTION_EVENT`, the file format is v3. (See Exceptional Condition 1 later in this section.)

4) If the type code is `START_EVENT_V3` (1), check the event length. If the length is less than 75, the file format is v1, and v3 otherwise. Why the value 75? Because that is the length of a v3 start event:

- header (19 bytes)
- binlog version (2 bytes)
- server version (`ST_SERVER_VER_LEN` = 50 bytes)
- timestamp (4 bytes)

Summing those lengths yields $19 + 2 + 50 + 4 = 75$

Therefore, if the event is shorter than 75 bytes, it must be from a v1 file because that will have a shorter first event than a v3 file.

5) If the type code is `FORMAT_DESCRIPTION_EVENT` (15), the file format is v4.

The preceding steps describe the general binary log format-recognition principles. However, there are some exceptional conditions that must be accounted for:

Exceptional Condition 1: In MySQL 4.0 and 4.1, the initial event in a binary log file might not be a start event. This occurs because the server writes the start event only to the first binary log file that it creates after startup. For subsequent files, the server writes an event of type `ROTATE_EVENT` to the end of the current log file, closes it, and then begins the next file without writing a start event to it. If a log file begins with an event that is not `START_EVENT_V3` or `FORMAT_DESCRIPTION_EVENT`, it can be assumed to be a v3 file because this behavior occurs only in MySQL 4.0 and 4.1, and all servers in those versions use v3 format.

Exceptional Condition 2: In MySQL 5.1 and 5.2, several early versions wrote binary log files using v4 format, but using different event numbers from those currently used in v4. Therefore, when the FDE is read and discovered to be v4, it is also necessary to read the server version, which is a string that occurs at position 21. If the version is one of those in the set of affected versions, event renumbering occurs such that events read from the file are mapped onto the current v4 event numbering.

Ensuring Compatibility of Future Binary Log Versions

To enable any future binary log formats to be correctly understood, the following conventions must hold:

- a) The binary log file must start with a descriptor event
- b) The descriptor event must start with a v3 header (19 bytes)
- c) The 2 bytes following the header (at position 19) must contain the binary log format version number

With respect to the current formats, only a) holds for v1. However, as indicated earlier, v1-format files can be recognized from the initial event in the file, by a type code of `START_EVENT_V3` and an event length less than 75.

The v4 format description event is designed so that it can handle future format updates. A new format with the same layout of event packets as in v4 but with additional fields in the header and post-headers can use this format description event to correctly describe itself. Actually, it is (theoretically) possible to have different "flavors" of v4 format that have different (larger) header lengths and even a different number of events.

The current code is written to handle this possibility. That is, any code that parses a binary log and discovers that it is v4 uses the header lengths as given by the format description event (thus potentially different lengths from the values hard-wired in the server code).

Note: Although headers of events in v4 format can be longer than 19 bytes, the format description event is an exception. Its header is always 19 bytes long to meet the preceding backward compatibility requirements. That is, the `FORMAT_DESCRIPTION_EVENT` does not include an `extra_headers` field.

Event Data for Specific Event Types

The following sections provide details about what appears in the fixed and variable parts of the event data for each event type.

`LOAD DATA INFILE` statements have been associated over time with several different events. The event contents are detailed in this section, and [LOAD DATA INFILE Events](#) provides a historical perspective on which events were used when.

`Start_log_event_v3/START_EVENT_V3`

This event occurs at the beginning of v1 or v3 binary log files. See [Binary Log Versions](#) for how it is used.

In MySQL 4.0 and 4.1, such events are written only to the first binary log file that `mysqld` creates after startup. Log files created subsequently (when someone issues a `FLUSH LOGS` statement or the current binary log file becomes too large) do not contain this event. In MySQL 5.0 and up, all binary log files start with a `FORMAT_DESCRIPTION_EVENT`.

Fixed data part:

- 2 bytes. The binary log format version. This is 1 in MySQL 3.23 and 3 in MySQL 4.0 and 4.1. (In MySQL 5.0 and up, `FORMAT_DESCRIPTION_EVENT` is used instead of `START_EVENT_V3`.)
- 50 bytes. The MySQL server's version (example: 4.0.14-debug-log), padded with 0x00 bytes on the right.
- 4 bytes. Timestamp in seconds when this event was created (this is the moment when the binary log was created). This value is redundant; the same value occurs in the `timestamp` header field.

Variable data part:

- Empty

`Query_log_event/QUERY_EVENT`

Fixed data part:

- 4 bytes. The ID of the thread that issued this statement. Needed for temporary tables. This is also useful for a DBA for knowing who did what on the master.
- 4 bytes. The time in seconds that the statement took to execute. Only useful for inspection by the DBA.
- 1 byte. The length of the name of the database which was the default database when the statement was executed. This name appears later, in the variable data part. It is necessary for statements such as `INSERT INTO t VALUES(1)` that don't specify the database and rely on the default database previously selected by `USE`.
- 2 bytes. The error code resulting from execution of the statement on the master. Error codes are defined in `include/mysqld_error.h`. 0 means no error. How come statements with a non-zero error code can exist in the binary log? This is mainly due to the use of non-transactional tables within transactions. For example, if an `INSERT ... SELECT` fails after inserting 1000 rows into a `MyISAM`

table (for example, with a duplicate-key violation), we have to write this statement to the binary log, because it truly modified the `MyISAM` table. For transactional tables, there should be no event with a non-zero error code (though it can happen, for example if the connection was interrupted (Control-C)). The slave checks the error code: After executing the statement itself, it compares the error code it got with the error code in the event, and if they are different it stops replicating (unless `--slave-skip-errors` was used to ignore the error).

- 2 bytes (not present in v1, v3). The length of the status variable block.

Variable part:

- Zero or more status variables (not present in v1, v3). Each status variable consists of one byte code identifying the variable stored, followed by the value of the variable. The format of the value is variable-specific, as described later.
- The default database name (null-terminated).
- The SQL statement. The slave knows the size of the other fields in the variable part (the sizes are given in the fixed data part), so by subtraction it can know the size of the statement.

Each entry in the status variable block has a code and a value, where the value format is as indicated in the following list. The list provides basic information about each variable. For additional details, see `log_event.h`.

- `Q_FLAGS2_CODE` = 0. Value is a 4-byte bit-field. This variable is written only as of MySQL 5.0.
- `Q_SQL_MODE_CODE` = 1. Value is an 8-byte SQL mode value.
- `Q_CATALOG_CODE` = 2. Value is the catalog name: a length byte followed by that many bytes, plus a terminating null byte. This variable is present only in MySQL 5.0.0 to 5.0.3. It was replaced with `Q_CATALOG_NZ_CODE` in MySQL 5.0.4 because the terminating null is unnecessary.
- `Q_AUTO_INCREMENT` = 3. Value is two 2-byte unsigned integers representing the `auto_increment_increment` and `auto_increment_offset` system variables. This variable is present only if `auto_increment` is greater than 1.
- `Q_CHARSET_CODE` = 4. Value is three 2-byte unsigned integers representing the `character_set_client`, `collation_connection`, and `collation_server` system variables.
- `Q_TIME_ZONE_CODE` = 5. Value is the time zone name: a length byte followed by that many bytes. This variable is present only if the time zone string is non-empty.
- `Q_CATALOG_NZ_CODE` = 6. Value is the catalog name: a length byte followed by that many bytes. Value is always `std`. This variable is present only if the catalog name is non-empty.
- `Q_LC_TIME_NAMES_CODE` = 7. Value is a 2-byte unsigned integer representing the `lc_time_names` number. This variable is present only if the value is not 0 (that is, not `en_US`).
- `Q_CHARSET_DATABASE_CODE` = 8. Value is a 2-byte unsigned integer representing the `collation_database` system variable.
- `Q_TABLE_MAP_FOR_UPDATE_CODE` = 9. Value is 8 bytes representing the table map to be updated by a multiple-table update statement. Each bit of this variable represents a table, and is set to 1 if the corresponding table is to be updated by the statement.

`Table_map_for_update` is used to evaluate the filter rules specified by `--replicate-do-table` / `--replicate-ignore-table`.

Stop_log_event/STOP_EVENT

A `Stop_log_event` is written under these circumstances:

- A master writes the event to the binary log when it shuts down
- A slave writes the event to the relay log when it shuts down or when a `RESET SLAVE` statement is executed

Fixed data part:

- Empty

Variable data part:

- Empty

`Rotate_log_event/ROTATE_EVENT`

When a binary log file exceeds a size limit, a `ROTATE_EVENT` is written at the end of the file that points to the next file in the sequence. This event is information for the slave to know the name of the next binary log it is going to receive.

`ROTATE_EVENT` is generated locally and written to the binary log on the master. It is written to the relay log on the slave when `FLUSH LOGS` occurs, and when receiving a `ROTATE_EVENT` from the master. In the latter case, there will be two rotate events in total originating on different servers.

There are conditions under which the terminating log-rotation event does not occur. For example, the server might crash.

Fixed data part:

- 8 bytes. The position of the first event in the next log file. Always contains the number 4 (meaning the next event starts at position 4 in the next binary log). This field is not present in v1; presumably the value is assumed to be 4.

Variable data part:

- The name of the next binary log. The filename is not null-terminated. Its length is the event size minus the size of the fixed parts.

`Intvar_log_event/INTVAR_EVENT`

Fixed data part:

- Empty

Variable data part:

- 1 byte. A value indicating the variable type: `LAST_INSERT_ID_EVENT` = 1 or `INSERT_ID_EVENT` = 2.
- 8 bytes. An unsigned integer indicating the value to be used for the `LAST_INSERT_ID()` invocation or `AUTO_INCREMENT` column.

`Load_log_event/LOAD_EVENT`

This event is used for `LOAD DATA INFILE` statements. See also [LOAD DATA INFILE Events](#).

`LOAD_EVENT` shares the `Load_log_event` class with `NEW_LOAD_EVENT`. The primary difference is that `LOAD_EVENT` allows only single-character field and line option values, whereas `NEW_LOAD_EVENT` allows multiple-character values. Also, `LOAD_EVENT` has no file ID or data block because with this event, the slave asks the master to transfer the data file at event execution time.

The format for this event is more complicated than for others, because it contains information about many `LOAD DATA INFILE` statement clauses.

Fixed data part:

- 4 bytes. The ID of the thread on the master that issued this `LOAD DATA INFILE` statement. Needed for temporary tables. This is also useful for a DBA for knowing who did what on the master.
- 4 bytes. The time in seconds which the `LOAD DATA INFILE` took for execution. Only useful for inspection by the DBA.
- 4 bytes. The number of lines to skip at the beginning of the file (corresponds to the `IGNORE N LINES` clause of `LOAD DATA INFILE`).
- 1 byte. The length of the name of the table to load.
- 1 byte. The length of the name of the database containing the table.
- 4 bytes. The number of columns to load (`(col_name, ...)` clause). Will be non-zero only if the columns to load were explicitly mentioned in the statement.

Variable data part:

- 1 byte. The field-terminating character (`FIELDS TERMINATED BY` option).
- 1 byte. The field-enclosing character (`FIELDS ENCLOSED BY` option).
- 1 byte. The line-terminating character (`LINES TERMINATED BY` option).
- 1 byte. The line-starting character (`LINES STARTING BY` option).
- 1 byte. The escaping character (`FIELDS ESCAPED BY` option).
- 1 byte. Flags that indicate whether certain keywords are present in the statement:
 - `DUMPFIELD_FLAG` = 0x1 (unused; this flag appears to be a botch because it would apply to `SELECT ... INTO OUTFILE`, not `LOAD DATA INFILE`)
 - `OPT_ENCLOSED_FLAG` = 0x2 (`FIELD OPTIONALLY ENCLOSED BY` option)
 - `REPLACE_FLAG` = 0x4 (`LOAD DATA INFILE REPLACE`)
 - `IGNORE_FLAG` = 0x8 (`LOAD DATA INFILE IGNORE`)
- 1 byte. Flags that indicate whether each of the field and line options are empty. The low-order five bits are 1 to indicate an empty option (has a length of 0) or 0 to indicate a non-empty option (has a length of 1).
 - `FIELD_TERM_EMPTY` = 0x1
 - `ENCLOSED_EMPTY` = 0x2
 - `LINE_TERM_EMPTY` = 0x4
 - `LINE_START_EMPTY` = 0x8
 - `ESCAPED_EMPTY` = 0x10
- 1 byte. The length of the name of the first column to load.
- ...
- 1 byte. The length of the name of the last column to load.
- Variable-sized. The name of the first column to load (null-terminated).
- ...

- Variable-sized. The name of the last column to load (null-terminated).
- Variable-sized. The name of the table to load (null-terminated).
- Variable-sized. The name of the database that contains the table (null-terminated).
- Variable-sized. The name of the file that was loaded (the original name on the master, not the name of the temporary file created on the slave). The length of the data filename is the event size minus the size of all other parts.

Note: Because this event allows only single-character field and line option values, presumably `LOAD DATA INFILE` statements will not replicate correctly if any such option contains multiple characters.

`Slave_log_event/SLAVE_EVENT`

This event is never written, so it cannot exist in a binary log file. It was meant for failsafe replication, which has never been implemented.

`Create_file_log_event/CREATE_FILE_EVENT`

This event is used for `LOAD DATA INFILE` statements. See also [LOAD DATA INFILE Events](#).

This event tells the slave to create a temporary file and fill it with a first data block. Later, zero or more `APPEND_BLOCK_EVENT` events append blocks to this temporary file. `EXEC_LOAD_EVENT` tells the slave to load the temporary file into the table, or `DELETE_FILE_EVENT` tells the slave not to do the load and to delete the temporary file. `DELETE_FILE_EVENT` occurs when the `LOAD DATA` failed on the master: On the master we start to write loaded blocks to the binary log before the end of the statement. If for some reason there is an error, we must tell the slave to abort the load.

The format for this event is more complicated than for others, because it contains information about many `LOAD DATA INFILE` statement clauses.

Fixed data part:

- 4 bytes. The ID of the thread on the master that issued this `LOAD DATA INFILE` statement. Needed for temporary tables. This is also useful for a DBA for knowing who did what on the master.

Variable data part:

- 4 bytes. The time in seconds which the `LOAD DATA INFILE` took for execution. Only useful for inspection by the DBA.
- 4 bytes. The number of lines to skip at the beginning of the file (corresponds to the `IGNORE N LINES` clause of `LOAD DATA INFILE`).
- 1 byte. The length of the name of the table to load.
- 1 byte. The length of the name of the database containing the table.
- 4 bytes. The number of columns to load (`((col_name,...)` clause). Will be non-zero only if the columns to load were explicitly mentioned in the statement.
- 4 bytes. An ID for the data file. This is necessary in case several `LOAD DATA INFILE` statements occur in parallel on the master. In that case, the binary log may contain intermixed events for the statements. The ID resolves which file the blocks in each `APPEND_BLOCK_EVENT` must be appended, and the file that must be loaded by the `EXEC_LOAD_EVENT` or deleted by the `DELETE_FILE_EVENT`.
- 1 byte. The length of the field-terminating string (`FIELDS TERMINATED BY` option).
- Variable-sized. The field-terminating string.
- 1 byte. The length of the field-enclosing string (`FIELDS ENCLOSED BY` option).

- Variable-sized. The field-enclosing string.
- 1 byte. The length of the line-terminating string (`LINES TERMINATED BY` option).
- Variable-sized. The line-terminating string.
- 1 byte. The length of the line-starting string (`LINES STARTING BY` option).
- Variable-sized. The line-starting string.
- 1 byte. The length of the escaping string (`FIELDS ESCAPED BY` option).
- Variable-sized. The escaping string.
- 1 byte. Flags that indicate whether certain keywords are present in the statement:
 - `DUMPFIELD_FLAG` = 0x1 (unused; this flag appears to be a botch because it would apply to `SELECT ... INTO OUTFILE`, not `LOAD DATA INFILE`)
 - `OPT_ENCLOSED_FLAG` = 0x2 (`FIELD OPTIONALLY ENCLOSED BY` option)
 - `REPLACE_FLAG` = 0x4 (`LOAD DATA INFILE REPLACE`)
 - `IGNORE_FLAG` = 0x8 (`LOAD DATA INFILE IGNORE`)
- 1 byte. The length of the name of the first column to load.
- ...
- 1 byte. The length of the name of the last column to load.
- Variable-sized. The name of the first column to load (null-terminated).
- ...
- Variable-sized. The name of the last column to load (null-terminated).
- Variable-sized. The name of the table to load (null-terminated).
- Variable-sized. The name of the database that contains the table (null-terminated).
- Variable-sized. The name of the file that was loaded (the original name on the master, not the name of the temporary file created on the slave) (null-terminated). The length of the data filename is not explicit in the event. However, it is null-terminated, so the length can be determined by reading to the null byte.
- Variable-sized. The block of raw data to load. If the file size exceeds a threshold, additional `APPEND_BLOCK_EVENT` instances will follow, each containing a data block. The size of the raw data is the event size minus the size of all other parts.

`Append_block_log_event/APPEND_BLOCK_EVENT`

This event is used for `LOAD DATA INFILE` statements. See also [LOAD DATA INFILE Events](#).

Contains data to be written to the data file for a `LOAD DATA INFILE` statement.

Fixed data part:

- 4 bytes. The ID of the file to append this block to.

Variable data part:

- The raw data to load. The raw data size is the event size minus the size of all the fixed event parts.

Execute_log_event/EXEC_LOAD_EVENT

This event is used for `LOAD DATA INFILE` statements. See also [LOAD DATA INFILE Events](#).

Indicates the end of a successful `LOAD DATA INFILE` statement and that the data file should be loaded.

Fixed data part:

- 4 bytes. The ID of the file to load.

Variable data part:

- Empty

Delete_file_log_event/DELETE_FILE_EVENT

This event is used for `LOAD DATA INFILE` statements. See also [LOAD DATA INFILE Events](#).

Indicates the end of an unsuccessful `LOAD DATA INFILE` statement and that the data file should not be loaded.

Fixed data part:

- 4 bytes. The ID of the file to be deleted.

Variable data part:

- Empty

Load_log_event/NEW_LOAD_EVENT

This event is used for `LOAD DATA INFILE` statements. See also [LOAD DATA INFILE Events](#).

The format for this event is more complicated than for others, because it contains information about many `LOAD DATA INFILE` statement clauses.

`LOAD_EVENT` shares the `Load_log_event` class with `NEW_LOAD_EVENT`. The primary difference is that `LOAD_EVENT` allows only single-character field and line option values, whereas `NEW_LOAD_EVENT` allows multiple-character values. Each of these is encoded as a length followed by a string rather than as a single character. Because of that, `NEW_LOAD_DATA` does not have the flags byte that indicates whether each option is empty.

Fixed data part:

- 4 bytes. The ID of the thread on the master that issued this `LOAD DATA INFILE` statement. Needed for temporary tables. This is also useful for a DBA for knowing who did what on the master.
- 4 bytes. The time in seconds which the `LOAD DATA INFILE` took for execution. Only useful for inspection by the DBA.
- 4 bytes. The number of lines to skip at the beginning of the file (corresponds to the `IGNORE N LINES` clause of `LOAD DATA INFILE`).
- 1 byte. The length of the name of the table to load.
- 1 byte. The length of the name of the database containing the table.
- 4 bytes. The number of columns to load (`(col_name, ...)` clause). Will be non-zero only if the columns to load were explicitly mentioned in the statement.

Variable data part:

- 1 byte. The length of the field-terminating string (`FIELDS TERMINATED BY` option).
- Variable-sized. The field-terminating string.
- 1 byte. The length of the field-enclosing string (`FIELDS ENCLOSED BY` option).
- Variable-sized. The field-enclosing string.
- 1 byte. The length of the line-terminating string (`LINES TERMINATED BY` option).
- Variable-sized. The line-terminating string.
- 1 byte. The length of the line-starting string (`LINES STARTING BY` option).
- Variable-sized. The line-starting string.
- 1 byte. The length of the escaping string (`FIELDS ESCAPED BY` option).
- Variable-sized. The escaping string.
- 1 byte. Flags that indicate whether certain keywords are present in the statement:
 - `DUMPFIELD_FLAG` = 0x1 (unused; this flag appears to be a botch because it would apply to `SELECT ... INTO OUTFILE`, not `LOAD DATA INFILE`)
 - `OPT_ENCLOSED_FLAG` = 0x2 (`FIELD OPTIONALLY ENCLOSED BY` option)
 - `REPLACE_FLAG` = 0x4 (`LOAD DATA INFILE REPLACE`)
 - `IGNORE_FLAG` = 0x8 (`LOAD DATA INFILE IGNORE`)
- 1 byte. The length of the name of the first column to load.
- ...
- 1 byte. The length of the name of the last column to load.
- Variable-sized. The name of the first column to load (null-terminated).
- ...
- Variable-sized. The name of the last column to load (null-terminated).
- Variable-sized. The name of the table to load (null-terminated).
- Variable-sized. The name of the database that contains the table (null-terminated).
- Variable-sized. The name of the file that was loaded (the original name on the master, not the name of the temporary file created on the slave). The length of the data filename is not explicit in the event. It is determined as the remaining characters to the end of the event.

There is no file ID or data block in the variable data part. The slave is supposed to request the file from the master in a separate connection.

`Rand_log_event/RAND_EVENT`

`RAND()` in MySQL generates a random number. A `RAND_EVENT` contains two seed values that set the `rand_seed1` and `rand_seed2` system variables that are used to compute the random number.

Fixed data part:

- Empty

Variable data part:

- 8 bytes. The value for the first seed.
- 8 bytes. The value for the second seed.

`User_var_log_event`/`USER_VAR_EVENT`

Fixed data part:

- Empty

Variable data part:

- 4 bytes. the size of the user variable name.
- The user variable name.
- 1 byte. Non-zero if the variable value is the SQL `NULL` value, 0 otherwise. If this byte is 0, the following parts exist in the event.
- 1 byte. The user variable type. The value corresponds to elements of `enum Item_result` defined in `include/mysql_com.h` (`STRING_RESULT=0`, `REAL_RESULT=1`, `INT_RESULT=2`, `ROW_RESULT=3`, `DECIMAL_RESULT=4`).
- 4 bytes. The number of the character set for the user variable (needed for a string variable). The character set number is really a collation number that indicates a character set/collation pair.
- 4 bytes. The size of the user variable value (corresponds to member `val_len` of class `Item_string`).
- Variable-sized. For a string variable, this is the string. For a float or integer variable, this is its value in 8 bytes. For a decimal this value is a packed value - 1 byte for the precision, 1 byte for the scale, and `$size - 2` bytes for the actual value. See the `decimal2bin` function in `strings/decimal.c` for the format of this packed value.

`Format_description_log_event`/`FORMAT_DESCRIPTION_EVENT`

This event occurs at the beginning of v4 binary log files. See [Binary Log Versions](#) for how it is used.

In MySQL 5.0 and up, all binary log files start with a `FORMAT_DESCRIPTION_EVENT`, but there will be a way to distinguish between a `FORMAT_DESCRIPTION_EVENT` created at `mysqld` startup and other `FORMAT_DESCRIPTION_EVENT` instances; such a distinction is needed because the first category of `FORMAT_DESCRIPTION_EVENT` (which means the master has started) should trigger some cleaning tasks on the slave. (Suppose the master died brutally and restarted: the slave must delete old replicated temporary tables.)

Fixed data part:

- 2 bytes. The binary log format version. This is 4 in MySQL 5.0 and up.
- 50 bytes. The MySQL server's version (example: 5.0.14-debug-log), padded with 0x00 bytes on the right.
- 4 bytes. Timestamp in seconds when this event was created (this is the moment when the binary log was created). This value is redundant; the same value occurs in the `timestamp` header field.
- 1 byte. The header length. This length - 19 gives the size of the `extra headers` field at the end of the header for other events.
- Variable-sized. An array that indicates the post-header lengths for all event types. There is one byte per event type that the server knows about.

Variable data part:

- Empty

[Xid_log_event/XID_EVENT](#)

An XID event is generated for a commit of a transaction that modifies one or more tables of an XA-capable storage engine. Strictly speaking, [Xid_log_event](#) is used if `thd->transaction.xid_state.xid.get_my_xid()` returns non-zero.

Here is an example of how to generate an XID event (it occurs whether or not [innodb_support_xa](#) is enabled):

```
CREATE TABLE t1 (a INT) ENGINE = INNODB;  
START TRANSACTION;  
INSERT INTO t1 VALUES (1);  
COMMIT;
```

Fixed data part:

- Empty

Variable data part:

- 8 bytes. The XID transaction number.

Note: Contrary to all other numeric fields, the XID transaction number is not always written in little-endian format. The bytes are copied unmodified from memory to disk, so the format is machine-dependent. Hence, when replicating from a little-endian to a big-endian machine (or vice versa), the numeric value of transaction numbers will differ. In particular, the output of [mysqlbinlog](#) differs. This should does not cause inconsistencies in replication because the only important property of transaction numbers is that different transactions have different numbers (relative order does not matter).

[Begin_load_query_log_event/BEGIN_LOAD_QUERY_EVENT](#)

This event is used for [LOAD DATA INFILE](#) statements. See also [LOAD DATA INFILE Events](#).

Fixed data part:

- 4 bytes. An ID for the data file. This is necessary in case several [LOAD DATA INFILE](#) statements occur in parallel on the master. In that case, the binary log may contain intermixed events for the statements. The ID resolves which file the blocks in each [APPEND_BLOCK_EVENT](#) must be appended to, and the file that must be loaded by the [EXEC_LOAD_QUERY_EVENT](#) or deleted by the [DELETE_FILE_EVENT](#).

Variable data part:

- Variable-sized. The first block of data to load. The size is the event size minus the size of all other fields in the event. If the file size exceeds a threshold, additional [APPEND_BLOCK_EVENT](#) instances will follow, each containing a data block.

[Execute_load_query_log_event/EXECUTE_LOAD_QUERY_EVENT](#)

This event is used for [LOAD DATA INFILE](#) statements. See also [LOAD DATA INFILE Events](#).

Indicates the end of a successful [LOAD DATA INFILE](#) statement and that the data file should be loaded. It is similar to [QUERY_EVENT](#), but before executing the statement, it substitutes the original filename in the statement with the name of the slave-side temporary file. The first 13 bytes of the fixed data part are the same as for [QUERY_EVENT](#), as is the initial status variable block in the variable data part. See the description of that event type for additional information.

Fixed data part:

- 4 bytes. The ID of the thread that issued this statement.

- 4 bytes. The time in seconds that the statement took to execute.
- 1 byte. The length of the name of the database which was the default database when the statement was executed.
- 2 bytes. The error code resulting from execution of the statement on the master.
- 2 bytes. The length of the status variable block.
- 4 bytes. The ID of the file to load.
- 4 bytes. The start position within the statement for filename substitution.
- 4 bytes. The end position within the statement for filename substitution.
- 1 byte. How to handle duplicates: `LOAD_DUP_ERROR` = 0, `LOAD_DUP_IGNORE` = 1, `LOAD_DUP_REPLACE` = 2

Variable data part:

- Zero or more status variables. Each status variable consists of one byte identifying the variable stored, followed by the value of the variable.
- Variable-sized. The database name (null-terminated).
- The `LOAD DATA INFILE` statement. The length is the event size minus the size of all other fields.

`Table_map_log_event/TABLE_MAP_EVENT`

Used for row-based binary logging beginning with MySQL 5.1.5.

Fixed data part:

- 6 bytes. The table ID.
- 2 bytes. Reserved for future use.

Variable data part:

- 1 byte. The length of the database name.
- Variable-sized. The database name (null-terminated).
- 1 byte. The length of the table name.
- Variable-sized. The table name (null-terminated).
- Packed integer. The number of columns in the table.
- Variable-sized. An array of column types, one byte per column.
- Packed integer. The length of the metadata block.
- Variable-sized. The metadata block; see `log_event.h` for contents and format.
- Variable-sized. Bit-field indicating whether each column can be `NULL`, one bit per column. For this field, the amount of storage required for N columns is $\text{INT}((N+7)/8)$ bytes.

`Write_rows_log_event_old/PRE_GA_WRITE_ROWS_EVENT`

Used for row-based binary logging from MySQL 5.1.5 to 5.1.17, using the old implementation of `Write_rows_log_event/WRITE_ROWS_EVENT`. The structure is similar to that for the newer event.

`Update_rows_log_event_old/PRE_GA_UPDATE_ROWS_EVENT`

Used for row-based binary logging from MySQL 5.1.5 to 5.1.17, using the old implementation of [Update_rows_log_event/UPDATE_ROWS_EVENT](#). The structure is similar to that for the newer event.

[Delete_rows_log_event_old/PRE_GA_DELETE_ROWS_EVENT](#)

Used for row-based binary logging from MySQL 5.1.5 to 5.1.17, using the old implementation of [Delete_rows_log_event/DELETE_ROWS_EVENT](#). The structure is similar to that for the newer event.

[Write_rows_log_event/WRITE_ROWS_EVENT](#)

Used for row-based binary logging beginning with MySQL 5.1.18.

[TODO: following needs verification; it's guesswork]

Fixed data part:

- 6 bytes. The table ID.
- 2 bytes. Reserved for future use.

Variable data part:

- Packed integer. The number of columns in the table.
- Variable-sized. Bit-field indicating whether each column is used, one bit per column. For this field, the amount of storage required for N columns is $\text{INT}((N+7)/8)$ bytes.
- Variable-sized (for [UPDATE_ROWS_LOG_EVENT](#) only). Bit-field indicating whether each column is used in the [UPDATE_ROWS_LOG_EVENT](#) after-image; one bit per column. For this field, the amount of storage required for N columns is $\text{INT}((N+7)/8)$ bytes.
- Variable-sized. A sequence of zero or more rows. The end is determined by the size of the event. Each row has the following format:
 - Variable-sized. Bit-field indicating whether each field in the row is [NULL](#). Only columns that are "used" according to the second field in the variable data part are listed here. If the second field in the variable data part has N one-bits, the amount of storage required for this field is $\text{INT}((N+7)/8)$ bytes.
 - Variable-sized. The row-image, containing values of all table fields. This only lists table fields that are used (according to the second field of the variable data part) and non-[NULL](#) (according to the previous field). In other words, the number of values listed here is equal to the number of zero bits in the previous field (not counting padding bits in the last byte).

The format of each value is described in the [log_event_print_value\(\)](#) function in [log_event.cc](#).

- (for [UPDATE_ROWS_EVENT](#) only) the previous two fields are repeated, representing a second table row.

For each row, the following is done:

- For [WRITE_ROWS_LOG_EVENT](#), the row described by the row-image is inserted.
- For [DELETE_ROWS_LOG_EVENT](#), a row matching the given row-image is deleted.
- For [UPDATE_ROWS_LOG_EVENT](#), a row matching the first row-image is removed, and the row described by the second row-image is inserted.

[Update_rows_log_event/UPDATE_ROWS_EVENT](#)

Used for row-based binary logging beginning with MySQL 5.1.18.

See the description for [WRITE_ROWS_EVENT](#).

[Delete_rows_log_event/DELETE_ROWS_EVENT](#)

Used for row-based binary logging beginning with MySQL 5.1.18.

See the description for [WRITE_ROWS_EVENT](#).

[Incident_log_event/INCIDENT_EVENT](#)

Fixed data part:

- 1 byte. The incident number.
- 1 byte. The message length.

Variable data part:

- The incident message, if present.

Incident number codes are listed in [rpl_constant.h](#). The only code currently used is [INCIDENT_LOST_EVENTS](#), which indicates that there may be lost events (a "gap") in the replication stream that requires databases to be resynchronized.

[Heartbeat_log_event/HEARTBEAT_LOG_EVENT](#)

A [Heartbeat_log_event](#) is sent by a master to a slave to let the slave know that the master is still alive. Events of this type do not appear in the binary or relay logs. They are generated on a master server by the thread that dumps events and sent straight to the slave without ever being written to the binary log. The slave examines the event contents and then discards it without writing it to the relay log.

Fixed data part:

- Empty

Variable data part:

- Empty

[Muted_query_log_event](#)

This is a subclass of [Query_log_event](#) that is not written to the log. It is used as a means of flushing a transaction without logging an event.

This event class was added in MySQL 5.0.23 and removed in 6.0.4. It was a solution to [Bug#16206](#) that became unnecessary with the fix for [Bug#29020](#).

LOAD DATA INFILE Events

[LOAD DATA INFILE](#) is not written to the binary log like other statements. It is written as one or more events in a packed format, not as a cleartext statement in the binary log. The events indicate what options are present in the statement and how to process the data file.

Historically, there seem to have been at least four event sequences for representing [LOAD DATA INFILE](#) operations.

- 1) In MySQL 3.23, there was only one event: [Load_log_event](#) (type code [LOAD_EVENT](#) = 6). [Load_log_event](#) only contains the filename, not the file itself. When the slave sees a [Load_log_event](#), it requests that the master send the file in a separate connection. This has the

drawback that the binary log is not self-contained: If the file has been removed on the master or the slave cannot access the master, the file transfer fails.

2) In MySQL 4.0.0, the file contents were included in the binary log. Several new event types were introduced: `Create_file_log_event` (type code `CREATE_FILE_EVENT` = 8), `Append_block_log_event` (type code `APPEND_BLOCK_EVENT` = 9), `Execute_load_log_event` (type code `EXEC_LOAD_EVENT` = 10), and `Delete_file_log_event` (type code `DELETE_FILE_EVENT` = 11). The event sequence is:

- `Create_file_log_event` : 1 instance
- `Append_block_log_event` : 0 or more instances
- `Execute_load_log_event` (success) or `Delete_file_log_event` (failure): 1 instance

The `Create_file_log_event` contains the options to `LOAD DATA INFILE`. This was a design flaw since the file cannot be loaded until the `Exec_load_log_event` is seen. To work around this, the slave, when executing the `Create_file_log_event`, wrote the `Create_file_log_event` to a temporary file. When the `Execute_load_log_event` was seen, this temporary file was read back so that the `LOAD DATA INFILE` statement could be constructed.

`Append_block_log_event` is used for files larger than a threshold. In this case, the file is split and the pieces are sent in separate events. The threshold is around $2^{17} = 131072$ bytes.

`CREATE_FILE_EVENT` tells the slave to create a temporary file and fill it with a first data block. Later, zero or more `APPEND_BLOCK_EVENT` events append blocks to this temporary file. `EXEC_LOAD_EVENT` tells the slave to load the temporary file into the table, or `DELETE_FILE_EVENT` tells the slave not to do the load and to delete the temporary file. `DELETE_FILE_EVENT` occurs when the `LOAD DATA` failed on the master: On the master we start to write loaded blocks to the binary log before the end of the statement. If for some reason there is an error, we must tell the slave to abort the load.

MySQL 4.0.0 also introduced the `NEW_LOAD_EVENT` = 12 type code.

If a slave reads a `NEW_LOAD_EVENT` from a binlog, it will use it as a `LOAD_EVENT` (but allowing longer separator names). Lost in the mysteries of time is the knowledge of whether there was ever a server version capable of writing `NEW_LOAD_EVENT`.

3) The original `Load_log_event` used one character for each of the delimiters (`FIELDS TERMINATED BY`, and so forth). At an unknown point in the version history, the format was modified to allow multiple-character strings as separators. This uses the same class, `Load_log_event`, but has the type code `NEW_LOAD_EVENT` = 12. This affects `Create_file_log_event`, since that inherits from `Load_log_event`. So the new feature in `Load_log_event` allows `Create_file_log_event` to use multiple-character delimiters.

[A guess for when this occurred would be MySQL 4.0.0: That is when the `sql_ex` structure that holds the single-character field/line options was renamed to `old_sql_ex` and a new `sql_ex` structure was created that allows multiple-character values.]

4) In 5.0.3, the event sequence was changed again with the addition of two new event types: `Begin_load_query_log_event` (type code `BEGIN_LOAD_QUERY_EVENT` = 17) and `Execute_load_query_log_event` (type code `EXECUTE_LOAD_QUERY_EVENT` = 18). The event sequence is:

- `Begin_load_query_log_event` : 1 instance
- `Append_block_log_event` : 0 or more instances
- `Execute_load_query_log_event` (success) or `Delete_file_log_event` (failure): 1 instance

With the new sequence, information about the options to `LOAD DATA INFILE` is moved from the first event to the last event. Consequently, `Begin_load_query_log_event` is almost the same as

[Append_file_log_event](#) (it contains only file data), whereas [Execute_load_query_log_event](#) contains the text of the `LOAD DATA INFILE` statement. The revised event sequence fixes the design flaw in the 4.0 format.

Also, the temp file that stores the parameters to `LOAD DATA INFILE` is not needed anymore. There is still a temp file containing all the data to be loaded.

Here is a concrete example (it applies to MySQL 4.0 and 4.1):

On the master we have a file named `/m/tmp/u.txt` that contains these lines:

```
>1,2,3
>4,5,6
>7,8,9
>10,11,12
```

And we issue this statement on the master while the default database is `test`:

```
load data infile '/m/tmp/u.txt' replace into table x fields
terminated by ',' optionally enclosed by '"' escaped by '\\'
lines starting by '>' terminated by '\n' ignore 2 lines (a,b,c);
```

Then in the master's binary log we have this [CREATE_FILE_EVENT](#) (hexadecimal dump):

```
00000180:          db4f 153f 0801 0000  ....O.?...
00000190: 006f 0000 0088 0100 0000 0004 0000 0000  .O.....
000001a0: 0000 0002 0000 0001 0403 0000 0003 0000  .....
000001b0: 0001 2c01 2201 0a01 3e01 5c06 0101 0161  .,."...>.\...a
000001c0: 0062 0063 0078 0074 6573 7400 2f6d 2f74  .b.c.x.test./m/t
000001d0: 6d70 2f75 2e74 7874 003e 312c 322c 330a  mp/u.txt.>1,2,3.
000001e0: 3e34 2c35 2c36 0a3e 372c 382c 390a 3e31  >4,5,6.>7,8,9.>1
000001f0: 302c 3131 2c31 32db 4f15 3f0a 0100 0000  0,11,12.O.?....
00000200: 1700 0000 f701 0000 0000 0300 0000  ....
```

- Line 180:
 - Timestamp (db4f 153f)
 - Event type (08)
 - Server ID (01 0000 00)
- Line 190:
 - Event size (6f 0000 00)
 - Position in the binary log (88 0100 00) (that's 392 in decimal base)
 - Flags (00 00)
 - Thread ID (04 0000 00)
 - Time it took (00 0000 00)
- Line 1a0:
 - Number of lines to skip at the beginning of the file (02 0000 00)
 - Length of the table name (01)
 - Length of the database name (04)
 - Number of columns to load (03 0000 00)
 - The file ID (03 0000 00)

- Line 1b0:
 - Length of the field terminating string (01)
 - Field terminating string (2c = ,)
 - Length of the field enclosing string (01)
 - Field enclosing string (22 = ")
 - Length of the line terminating string (01)
 - Line terminating string (0a = newline)
 - Length of the line starting string (01)
 - Line starting string (3e = >)
 - Length of the escaping string (01)
 - Escaping string (5c = \)
 - Flags (06) (that's `OPT_ENCLOSED_FLAG` | `REPLACE_FLAG`)
 - Length of the name of the first column to load (01)
 - Length of the name of the second column to load (01)
 - Length of the name of the third column to load (01)
 - Name of the first column to load (61 00 = "a")
- Line 1c0:
 - Name of the second column to load (62 00 = "b")
 - Name of the third column to load (63 00 = "c")
 - Name of the table to load (78 00 = "x"), name of the database that contains the table (74 6573 7400 = "test")
 - Name of the file on the master (2f6d 2f74 6d70 2f75 2e74 7874 00 = "/m/tmp/u.txt")
- Line 1d0 and following:
 - Raw data to load (3e 312c 322c 330a 3e34 2c35 2c36 0a3e 372c 382c 390a 3e31 302c 3131 2c31 32)

The next byte in the file is the beginning of the `EXEC_LOAD_EVENT` event.

Row-Based Binary Logging

Originally, the binary log was written using statement-based logging. Row-based logging was added in MySQL 5.1.5.

- Statement-based logging: Events contain SQL statements that produce data changes (inserts, updates, deletes)
- Row-based logging: Events describe changes to individual rows

Several event types are used specific to row-based logging:

- `TABLE_MAP_EVENT`

- [WRITE_ROWS_EVENT](#)
- [UPDATE_ROWS_EVENT](#)
- [DELETE_ROWS_EVENT](#)

It's common to refer to "row-based logging" (RBL) as "row-based replication" (RBR), but RBR is in fact a misnomer. Logging in this format can be done independent of whether the log is used for replication. That is, replication need not figure into the use of the binary log at all. Similar remarks apply to the terms "statement-based logging" (SBL) versus "statement-based replication" (SBR). Unfortunately, by now the RBR/SBR terminology is probably too well entrenched for there to be much hope of a return to the proper RBL/SBL terms, so I must content myself with including this nomenclatural screed here. :-)

Additional Resources

The following resources provide additional information on the structure, content, or use of the binary log.

- In this document, see [Chapter 19, Replication](#): Replication is based on the transfer of the binary log contents from master to slave.
- MySQL Reference Manual, [The Binary Log](#)
- MySQL Reference Manual, [Replication Relay and Status Files](#)
- MySQL Reference Manual, [mysqlbinlog — Utility for Processing Binary Log Files](#)

Chapter 21. MySQL Backup

Table of Contents

Glossary	241
Architectural Design & APIs	241
Stub for WL#3169 material (Server kernel)	241
Stub for WL#3473 material (Backup engine API)	241
Stub for WL#3569 material (Backup synchronization algorithm)	242
Stub for WL#3574 material (Online backup: Service API for Metadata Objects II)	242
Backup Engines and Drivers	242
Stub for WL#3570 material (Default algorithm backup & restore (blocking))	242
Stub for WL#3776 material (Consistent snapshot backup for MVCC)	242
Stub for WL#866 material (Online backup: MyISAM native driver (redo-log algorithm, modify source file, log byte changes))	242
Stub for WL#4064 material (Online Backup: Adapt native MyISAM driver to new kernel architecture)	242
Stub for WL#4205 material (Online Backup: modify kernel to use object services API)	242
Stub for WL#3576 material (synchronization with binary log)	242
Stub for WL#4060 material (kernel updates)	242
Stub for WL#4062 material (DDL blocker, A.K.A metadata freeze)	243
Stub for WL#3956 material (commit blocker)	243
Error Handling	243
Stub for WL#3904 (Error Handling)	243
Stub for WL#4190 (Error Handling II)	243
Testing	243
Stub for WL#3324 (Test breakpoints)	243
Stub for WL#4120 (Breakpoints)	243
Data Formats	243
Stub for WL#3568 (Online backup: Stream format and server file handling)	243
WL#4063 (Online Backup: Finalize backup stream format)	243

This chapter describes MySQL Backup.

Mostly it consists of stubs for material yet to be added.

Glossary

Definitions of important terms. Examples:

- Server instance
- Backup kernel
- Backup engine
- Backup/restore driver

Architectural Design & APIs

Stub for WL#3169 material (Server kernel)

This section is yet to be written.

Stub for WL#3473 material (Backup engine API)

This section is yet to be written.

Stub for WL#3569 material (Backup synchronization algorithm)

This section is yet to be written.

Stub for WL#3574 material (Online backup: Service API for Metadata Objects II)

This section is yet to be written.

Backup Engines and Drivers

Stub for WL#3570 material (Default algorithm backup & restore (blocking))

There is a default backup engine that provides default backup and restore drivers, to be used if a better backup engine cannot be found. This engine's drivers use blocking algorithms.

Stub for WL#3776 material (Consistent snapshot backup for MVCC)

The consistent snapshot engine provides a backup similar to that created by `mysqldump --single-transaction`.

Stub for WL#866 material (Online backup: MyISAM native driver (redo-log algorithm, modify source file, log byte changes))

This section is yet to be written.

Stub for WL#4064 material (Online Backup: Adapt native MyISAM driver to new kernel architecture)

This section is yet to be written.

Stub for WL#4205 material (Online Backup: modify kernel to use object services API)

This section is yet to be written.

Stub for WL#3576 material (synchronization with binary log)

A backup image must have contents that are consistent with the binary log coordinates taken from the time of the backup. Otherwise, point-in-time recovery using the backup image plus the binary log contents beginning from the coordinates will not work correctly. Online Backup synchronizes with binary logging to make sure that the backup image and binary log are consistent with each other.

(Discuss implementation of synchronization with binary logging)

Stub for WL#4060 material (kernel updates)

- Character support for identifiers (database/table metadata)
- Table name mangling
- Implementation of the catalog services API
- Implementation of the Image_info class
- Memory management

- Memory (de)allocator for backup stream library

Stub for WL#4062 material (DDL blocker, A.K.A metadata freeze)

During a backup or restore operation, it is not allowable to modify the structure of databases or tables being backed up. Consequently, several DDL statements are blocked during these operations.

(Discuss implementation of blocking)

Stub for WL#3956 material (commit blocker)

Rationale is similar to that for DDL blocking.

(Discuss implementation of blocking)

Error Handling

Stub for WL#3904 (Error Handling)

This section is yet to be written.

Stub for WL#4190 (Error Handling II)

This section is yet to be written.

Testing

The code allows for breakpoints to be inserted into the backup or restore process

Stub for WL#3324 (Test breakpoints)

This section is yet to be written.

Stub for WL#4120 (Breakpoints)

This section is yet to be written.

Data Formats

Stub for WL#3568 (Online backup: Stream format and server file handling)

This section is yet to be written.

WL#4063 (Online Backup: Finalize backup stream format)

There are two parts to this:

- Definitions of the data formats
- API for library that applications use to read/write these formats

Chapter 22. [MyISAM](#) Storage Engine

Table of Contents

MyISAM Record Structure	245
Introduction	245
Physical Attributes of Columns	247
Where to Look For More Information	251
The .MYI File	251
MyISAM Files	257
MyISAM Dynamic Data File Layout	257
Layout of the Record Storage Frame (Record Part, Record Block)	257
Record Contents	259
Packed Record Layout	259
In-memory Record Layout	259
MyISAM Compressed Data File Layout	259
Huffman Compression	259
The myisampack Program	259
Record and Blob Length Encoding	260
Code Tree Representation	260
Usage of the Index File	260
myisampack Tricks	260
Detailed Specification of the Decoding	260
MyISAM Key Cache	262
MyISAM Concurrent Insert	262

[MyISAM](#) Record Structure

Introduction

When you say:

```
CREATE TABLE Table1 ...
```

MySQL creates files named [Table1.MYD](#) ("MySQL Data"), [Table1.MYI](#) ("MySQL Index"), and [Table1.frm](#) ("Format"). These files will be in the directory:

```
/<datadir>/<database>/
```

For example, if you use Linux, you might find the files in the [/usr/local/var/test](#) directory (assuming your database name is [test](#)). If you use Windows, you might find the files in the [\mysql\data\test\](#) directory.

Let's look at the [.MYD](#) Data file ([MyISAM](#) SQL Data file) more closely. There are three possible formats — fixed, dynamic, and packed. First, let's discuss the fixed format.

- **Page Size**

Unlike most DBMSs, MySQL doesn't store on disk using pages. Therefore you will not see filler space between rows. (Reminder: This does not refer to [BDB](#) and [InnoDB](#) tables, which do use pages).

- **Record Header**

The minimal record header is a set of flags:

- "X bit" = 0 if row is deleted, = 1 if row is not deleted

- "Null Bits" = 1 if row contains any null fields, or = 0 otherwise.
- "Filler Bits" = 1

The length of the record header is thus:

```
(1 + number of NULL columns + 7) / 8 bytes
```

After the header, all columns are stored in the order that they were created, which is the same order that you would get from `SHOW COLUMNS`.

Here's an example. Suppose you say:

```
CREATE TABLE Table1 (column1 CHAR(1), column2 CHAR(1), column3 CHAR(1));
INSERT INTO Table1 VALUES ('a', 'b', 'c');
INSERT INTO Table1 VALUES ('d', NULL, 'e');
```

A `CHAR(1)` column takes precisely one byte (plus one bit of overhead that is assigned to every column — I'll describe the details of column storage later). So the file `Table1.MYD` looks like this:

Hexadecimal Display of `Table1.MYD` file

```
F1 61 62 63 00 F5 64 00 66 00          ... .abc..d e.
```

Here's how to read this hexadecimal-dump display:

- The hexadecimal numbers `F1 61 62 63 00 F5 64 20 65 00` are byte values and the column on the right is an attempt to show the same bytes in ASCII.
- The `F1` byte means that there are no null fields in the first row.
- The `F5` byte means that the second column of the second row is `NULL`.

(It's probably easier to understand the flag setting if you restate `F5` as `11110101` binary, and (a) notice that the third flag bit from the right is `on`, and (b) remember that the first flag bit is the X bit.)

There are complications — the record header is more complex if there are variable-length fields — but the simple display shown in the example is exactly what you'd see if you looked at the MySQL Data file with a debugger or a hexadecimal file dumper.

So much for the fixed format. Now, let's discuss the dynamic format.

The dynamic file format is necessary if rows can vary in size. That will be the case if there are `BLOB` columns, or "true" `VARCHAR` columns. (Remember that MySQL may treat `VARCHAR` columns as if they're `CHAR` columns, in which case the fixed format is used.) A dynamic row has more fields in the header. The important ones are "the actual length", "the unused length", and "the overflow pointer". The actual length is the total number of bytes in all the columns. The unused length is the total number of bytes between one physical record and the next one. The overflow pointer is the location of the rest of the record if there are multiple parts.

For example, here is a dynamic row:

```
03          start of header - Block type, see mi_dynrec.c, _mi_get_block_info()
04, 00      actual length
0c          unused length
01, fc      flags + overflow pointer
****       data in the row
*****     unused bytes
          <-- next row starts here)
```

In the example, the actual length and the unused length are short (one byte each) because the table definition says that the columns are short — if the columns were potentially large, then the actual

length and the unused length could be two bytes each, three bytes each, and so on. In this case, actual length plus unused length is 10 hexadecimal (sixteen decimal), which is a minimum.

In a dynamic row, there is no deleted bit. Instead, deleted rows are marked with a block of type 0.

As for the third format — packed — we will only say briefly that:

- Numeric values are stored in a form that depends on the range (start/end values) for the data type.
- All columns are packed using either Huffman or enum coding.

For details, see the source files `/myisam/mi_statrec.c` (for fixed format), `/myisam/mi_dynrec.c` (for dynamic format), and `/myisam/mi_packrec.c` (for packed format).

Note: Internally, MySQL uses a format much like the fixed format which it uses for disk storage. The main differences are:

1. `BLOB` values have a length and a memory pointer rather than being stored inline.
2. "True `VARCHAR`" (a column storage which will be fully implemented in version 5.0) will have a 16-bit length plus the data.
3. All integer or floating-point numbers are stored with the low byte first. Point (3) does not apply for `ISAM` storage or internals.

Physical Attributes of Columns

Next I'll describe the physical attributes of each column in a row. The format depends entirely on the data type and the size of the column, so, for every data type, I'll give a description and an example.

All the types are defined within the `include/mysql_com.h` file within the `enum_field_types` enumerated structure. Here's a sample of the key values and corresponding numbers:

```
MYSQL_TYPE_BIT 16 MYSQL_TYPE_BLOB 252 MYSQL_TYPE_DATE 10 MYSQL_TYPE_DATETIME 12 MYSQL_TYPE_DECIMAL 0 M
```

• The character data types

`CHAR`

- Storage: fixed-length string with space padding on the right.
- Example: a `CHAR(5)` column containing the value 'A' looks like: hexadecimal 41 20 20 20 20 -- (length = 5)

`VARCHAR`

- Storage: variable-length string with a preceding length.
- Example: a `VARCHAR(7)` column containing 'A' looks like: hexadecimal 01 41 -- (length = 1)
- In MySQL 4.1 the length is always 1 byte. In MySQL 5.0 the length may be either 1 byte (for up to 255) or 2 bytes (for 256 to 65535). Some further random notes about the new format: In old tables (from MySQL 4.1 and earlier), `VARCHAR` columns have type `MYSQL_TYPE_VAR_STRING`, which works exactly like a `CHAR` with the exception that if you do an `ALTER TABLE`, it's converted to a true `VARCHAR` (`MYSQL_TYPE_VARCHAR`). (This means that old tables will work as before for users.) ... Apart from the above case, there are no longer any automatic changes from `CHAR` to `VARCHAR` or from `VARCHAR` to `CHAR`. MySQL will remember the declared type and stick to it ... `VARCHAR` is implemented in `field.h` and `field.cc` through the new class `Field_varstring` ... `MyISAM` implements `VARCHAR` both for dynamic-length and fixed-length rows (as signaled with the `ROW_FORMAT` flag) ... `VARCHAR` now stores trailing spaces. (If they don't fit, that's an error in strict mode.) Trailing spaces are not significant in comparisons ... In `table->record`, the space is reserved for length (1 or 2 bytes) plus data ... The number of bytes used to store

the length is in the field `Field_varchar->length_bytes`. Note that internally this can be 2 even if `Field_varchar->field_length < 256` (for example, for a shortened key to a `varchar(256)`) ... There is a new macro, `HA_VARCHAR_PACKLENGTH(field_length)`, that can be used on `field->length` in `write_row / read_row` to check how many length bytes are used. (In this context we can't have a `field_length < 256` with a 2-byte pack length) ... When creating a key for the handler, `HA_KEYTYPE_VARTTEXT1` and `HA_KEYTYPE_BINARY1` are used for a key on a column that has a 1-byte length prefix and `HA_KEYTYPE_VARTTEXT2` and `HA_KEYTYPE_BINARY2` for a column that has a 2-byte length prefix. (In the future we will probably delete `HA_KEYTYPE_BINARY#`, as this can be instead be done by just using the `binary` character set with `HA_KEYTYPE_VARTTEXT#`.) ... When sending a key to the handler for `index_read()` or `records_in_range`, we always use a 2-byte length for the `VARCHAR` to make things simpler. (For version 5.1 we intend to change `CHARS` to also use a 2-byte length for these functions, as this will speed up and simplify the key handling code on the handler side.) ... The test case file `mysql-test/include/varchar.inc` should be included in the code that tests the handler. See `t/myisam.test` for how to use this. You should verify the result against the one in `mysql-test/t/myisam.result` to ensure that you get the correct results ... A client sees both the old and new `VARCHAR` type as `MYSQL_TYPE_VAR_STRING`. It will never (at least for 5.0) see `MYSQL_TYPE_VARCHAR`. This ensures that old clients will work as before ... If you run MySQL 5.0 with the `--new` option, MySQL will show old `VARCHAR` columns as `'CHAR'` in `SHOW CREATE TABLE`. (This is useful when testing whether a table is using the new `VARCHAR` type or not.)

- **The numeric data types**

Important: MySQL almost always stores multi-byte binary numbers with the low byte first. This is called "little-endian" numeric storage; it's normal on Intel x86 machines; MySQL uses it even for non-Intel machines so that databases will be portable.

`TINYINT`

- Storage: fixed-length binary, always one byte.
- Example: a `TINYINT` column containing 65 looks like: `hexadecimal 41 --` (length = 1, value = 65)

`SMALLINT`

- Storage: fixed-length binary, always two bytes.
- Example: a `SMALLINT` column containing 65 looks like: `hexadecimal 41 00 --` (length = 2, value = 65)

`MEDIUMINT`

- Storage: fixed-length binary, always three bytes.
- Example: a `MEDIUMINT` column containing 65 looks like: `hexadecimal 41 00 00 --` (length = 3, value = 65)

`INT`

- Storage: fixed-length binary, always four bytes.
- Example: an `INT` column containing 65 looks like: `hexadecimal 41 00 00 00 --` (length = 4, value = 65)

`BIGINT`

- Storage: fixed-length binary, always eight bytes.
- Example: a `BIGINT` column containing 65 looks like: `hexadecimal 41 00 00 00 00 00 00 00 --` (length = 8, value = 65)

FLOAT

- Storage: fixed-length binary, always four bytes.
- Example: a **FLOAT** column containing approximately 65 looks like: `hexadecimal 00 00 82 42` -- (length = 4, value = 65)

DOUBLE PRECISION

- Storage: fixed-length binary, always eight bytes.
- Example: a **DOUBLE PRECISION** column containing approximately 65 looks like: `hexadecimal 00 00 00 00 00 40 50 40` -- (length = 8, value = 65)

REAL

- Storage: same as **FLOAT**, or same as **DOUBLE PRECISION**, depending on the setting of the `--ansi` option.

DECIMAL

- MySQL 4.1 Storage: fixed-length string, with a leading byte for the sign, if any.
- Example: a **DECIMAL(2)** column containing 65 looks like: `hexadecimal 20 36 35` -- (length = 3, value = '65')
- Example: a **DECIMAL(2) UNSIGNED** column containing 65 looks like: `hexadecimal 36 35` -- (length = 2, value = '65')
- Example: a **DECIMAL(4,2) UNSIGNED** column containing 65 looks like: `hexadecimal 36 35 2E 30 30` -- (length = 5, value = '65.00')
- MySQL 5.0 Storage: high byte first, four-byte chunks. We call the four-byte chunks "decimal digits". Since $2^{32} =$ There is an implied decimal point. Details are in `/strings/decimal.c`.
- Example: a MySQL 5.0 **DECIMAL(21,9)** column containing `111222333444.555666777` looks like: `hexadecimal 80 6f 0d 40 8a 04 21 1e cd 59` -- (flag + '111', '222333444', '555666777').

NUMERIC

- Storage: same as **DECIMAL**.

BOOL

- Storage: same as **TINYINT**.
- **The temporal data types**

DATE

- Storage: 3 byte integer, low byte first. Packed as: 'day + month*32 + year*16*32'
- Example: a **DATE** column containing '1962-01-02' looks like: `hexadecimal 22 54 0F`

DATETIME

- Storage: eight bytes.
- Part 1 is a 32-bit integer containing $\text{year} \times 10000 + \text{month} \times 100 + \text{day}$.

- Part 2 is a 32-bit integer containing $\text{hour} \times 10000 + \text{minute} \times 100 + \text{second}$.
- Example: a `DATETIME` column for `'0001-01-01 01:01:01'` looks like: hexadecimal `B5 2E 11 5A 02 00 00 00`

TIME

- Storage: 3 bytes, low byte first. This is stored as seconds: $\text{days} \times 24 \times 3600 + \text{hours} \times 3600 + \text{minutes} \times 60 + \text{seconds}$
- Example: a `TIME` column containing `'1 02:03:04'` (1 day 2 hour 3 minutes and 4 seconds) looks like: hexadecimal `58 6E 01`

TIMESTAMP

- Storage: 4 bytes, low byte first. Stored as unix `time()`, which is seconds since the Epoch (00:00:00 UTC, January 1, 1970).
- Example: a `TIMESTAMP` column containing `'2003-01-01 01:01:01'` looks like: hexadecimal `4D AE 12 23`

YEAR

- Storage: same as unsigned `TINYINT` with a base value of 0 = 1901.
- **Others**

SET

- Storage: one byte for each eight members in the set.
- Maximum length: eight bytes (for maximum 64 members).
- This is a bit list. The least significant bit corresponds to the first listed member of the set.
- Example: a `SET('A', 'B', 'C')` column containing `'A'` looks like: `01 --` (length = 1, value = 'A')

ENUM

- Storage: one byte if less than 256 alternatives, else two bytes.
- This is an index. The value 1 corresponds to the first listed alternative. (Note: `ENUM` always reserves the value 0 for an erroneous value. This explains why `'A'` is 1 instead of 0.)
- Example: an `ENUM('A', 'B', 'C')` column containing `'A'` looks like: `01 --` (length = A')
- **The Large-Object data types**

Warning: Because `TINYBLOB`'s preceding length is one byte long (the size of a `TINYINT`) and `MEDIUMBLOB`'s preceding length is three bytes long (the size of a `MEDIUMINT`), it's easy to think there's some sort of correspondence between the `BLOB` and `INT` types. There isn't a `BLOB`'s preceding length is not four bytes long (the size of an `INT`). `TINYBLOB`

- Storage: variable-length string with a preceding one-byte length.
- Example: a `TINYBLOB` column containing `'A'` looks like: hexadecimal `01 41 --` (length = A')

TINYTEXT

- Storage: same as `TINYBLOB`.

BLOB

- Storage: variable-length string with a preceding two-byte length.
- Example: a `BLOB` column containing 'A' looks like: hexadecimal 01 00 41 -- (length = A')

`TEXT`

- Storage: same as `BLOB`.

`MEDIUMBLOB`

- Storage: variable-length string with a preceding length.
- Example: a `MEDIUMBLOB` column containing 'A' looks like: hexadecimal 01 00 00 41 -- (length = A')

`MEDIUMTEXT`

- Storage: same as `MEDIUMBLOB`.

`LONGBLOB`

- Storage: variable-length string with a preceding four-byte length.
- Example: a `LONGBLOB` column containing 'A' looks like: hexadecimal 01 00 00 00 41 -- (length = A')

`LONGTEXT`

- Storage: same as `LONGBLOB`.

Where to Look For More Information

References:

Most of the formatting work for `MyISAM` columns is visible in the program `/sql/field.cc` in the source code directory. And in the `MyISAM` directory, the files that do formatting work for different record formats are: `/myisam/mi_statrec.c`, `/myisam/mi_dynrec.c`, and `/myisam/mi_packrec.c`.

The `.MYI` File

A `.MYI` file for a `MyISAM` table contains the table's indexes.

The `.MYI` file has two parts: the header information and the key values. So the next sub-sections will be "The `.MYI` Header" and "The `.MYI` Key Values".

The `.MYI` Header

A `.MYI` file begins with a header, with information about options, about file sizes, and about the "keys". In MySQL terminology, a "key" is something that you create with `CREATE [UNIQUE] INDEX`.

Program files which read and write `.MYI` headers are in the `./myisam` directory: `mi_open.c` has the routines that write each section of the header, `mi_create.c` has a routine that calls the `mi_open.c` routines in order, and `myisamdef.h` has structure definitions corresponding to what we're about to describe.

These are the main header sections:

Section	Occurrences
-----	-----
state	Occurs 1 time

base	Occurs 1 time
keydef (including keysegs)	Occurs once for each key
recinfo	Occurs once for each field

Now we will look at each of these sections, showing each field.

We are going to use an example table throughout the description. To make the example table, we executed these statements:

```
CREATE TABLE T (S1 CHAR(1), S2 CHAR(2), S3 CHAR(3));
CREATE UNIQUE INDEX I1 ON T (S1);
CREATE INDEX I2 ON T (S2,S3);
INSERT INTO T VALUES ('1', 'aa', 'b');
INSERT INTO T VALUES ('2', 'aa', 'bb');
INSERT INTO T VALUES ('3', 'aa', 'bbb');
DELETE FROM T WHERE S1 = '2';
```

We took a hexadecimal dump of the resulting file, [T.MYI](#).

In all the individual descriptions below, the column labeled “Dump From Example File” has the exact bytes that are in [T.MYI](#). You can verify that by executing the same statements and looking at a hexadecimal dump yourself. With Linux this is possible using `od -h T.MYI`; with Windows you can use the command-line debugger.

Along with the typical value, we may include a comment. The comment usually explains why the value is what it is. Sometimes the comment is derived from the comments in the source code.

state

This section is written by [mi_open.c](#), [mi_state_info_write\(\)](#).

Name	Size	Dump From Example File	Comment
----	----	-----	-----
file_version	4	FE FE 07 01	from myisam_file_magic
options	2	00 02	HA_OPTION_COMPRESS_RECORD etc.
header_length	2	01 A2	this header example has 0x01A2 bytes
state_info_length	2	00 B0	= MI_STATE_INFO_SIZE defined in myisamdef.h
base_info_length	2	00 64	= MI_BASE_INFO_SIZE defined in myisamdef.h
base_pos	2	00 D4	= where the base section starts
key_parts	2	00 03	a key part is a column within a key
unique_key_parts	2	00 00	key-parts+unique-parts
keys	1	02	here are 2 keys -- I1 and I2
uniques	1	00	number of hash unique keys used internally in temporary tables (nothing to do with 'UNIQUE' definitions)
language	1	08	"language for indexes"
max_block_size	1	01	
fulltext_keys	1	00	# of fulltext keys. = 0 if version <= 4.0
not_used	1	00	to align to 8-byte boundary
state->open_count	2	00 01	
state->changed	1	39	set if table updated; reset if shutdown (so one can examine this to see if there was an update without proper shutdown)

state->sortkey	1	FF	"sorted by this key" (not used)
state->state.records	8	00 00 00 00 00 00 00 02	number of actual, un-deleted, records
state->state.del	8	00 00 00 00 00 00 00 01	# of deleted records
state->split	8	00 00 00 00 00 00 00 03	# of "chunks" (e.g. records or spaces left after record deletion)
state->dellink	8	00 00 00 00 00 00 00 07	"Link to next removed "block". Initially = HA_OFFSET_ERROR
state->state.key_file_length	8	00 00 00 00 00 00 0c 00	2048
state->state.data_file_length	8	00 00 00 00 00 00 00 15	= size of .MYD file
state->state.empty	8	00 00 00 00 00 00 00 00	
state->state.key_empty	8	00 00 00 00 00 00 00 00	
state->auto_increment	8	00 00 00 00 00 00 00 00	
state->checksum	8	00 00 00 00 00 00 00 00	
state->process	4	00 00 09 E6	from getpid(). process of last update
state->unique	4	00 00 00 0B	initially = 0
state->status	4	00 00 00 00	
state->update_count	4	00 00 00 04	updated for each write lock (there were 3 inserts + 1 delete, total 4 operations)
state->key_root	8	00 00 00 00 00 00 04 00	offset in file where 11 keys start, can be = HA_OFFSET_ERROR
		00 00 00 00 00 00 08 00	state->key_root occurs twice because there are two keys
state->key_del	8	FF FF FF FF FF FF FF FF	delete links for keys (occurs many times if many delete links)
state->sec_index_changed	4	00 00 00 00	sec_index = secondary index (presumably) not currently used
state->sec_index_used	4	00 00 00 00	"which extra indexes are in use" not currently used
state->version	4	3F 3F EB F7	"timestamp of create"
state->key_map	8	00 00 00 03	"what keys are in use"
state->create_time	8	00 00 00 00 3F 3F EB F7	"time when database created" (actually: time when file made)
state->recover_time	8	00 00 00 00 00 00 00 00	"time of last recover"
state->check_time	8	00 00 00 00 3F 3F EB F7	"time of last check"
state->rec_per_key_rows	8	00 00 00 00 00 00 00 00	
state->rec_per_key_parts	4	00 00 00 00	(key_parts = 3, so rec_per_key_parts occurs 3 times)
		00 00 00 00	
		00 00 00 00	

base

This section is written by `mi_open.c`, `mi_base_info_write()`. The corresponding structure in `myisamdef.h` is `MI_BASE_INFO`.

In our example `T.MYI` file, the first byte of the `base` section is at offset `0x00d4`. That's where it's supposed to be, according to the header field `base_pos` (above).

Name	Size	Dump	From Example File	Comment
----	----	-----	-----	-----
base->keystart	8	00 00 00 00 00 00 04 00		keys start at offset 1024 (0x0400)
base->max_data_file_length	8	00 00 00 00 00 00 00 00		
base->max_key_file_length	8	00 00 00 00 00 00 00 00		
base->records	8	00 00 00 00 00 00 00 00		
base->reloc	8	00 00 00 00 00 00 00 00		
base->mean_row_length	4	00 00 00 00		
base->reclength	4	00 00 00 07		length(s1)+length(s2)

```

base->pack_reclength      4    00 00 00 07
base->min_pack_length     4    00 00 00 07
base->max_pack_length     4    00 00 00 07
base->min_block_length    4    00 00 00 14
base->fields              4    00 00 00 04
                                4 fields: 3 defined,
                                plus 1 extra

base->pack_fields         4    00 00 00 00
base->rec_reclength       1    04
base->key_reclength       1    04
base->keys                1    02
                                was 0 at start
base->auto_key            1    00
base->pack_bits           2    00 00
base->blobs               2    00 00
base->max_key_block_length 2    04 00
                                length of block = 1024
                                bytes (0x0400)
                                including length of
                                pointer

base->max_key_length       2    00 10

base->extra_alloc_bytes   2    00 00
base->extra_alloc_procent 1    00
base->raid_type           1    00
base->raid_chunks         2    00 00
base->raid_chunksize      4    00 00 00 00
[extra] i.e. filler       6    00 00 00 00 00 00

```

keydef

This section is written by `mi_open.c`, `mi_keydef_write()`. The corresponding structure in `myisamdef.h` is `MI_KEYDEF`.

This is a multiple-occurrence structure. Since there are two indexes in our example (I1 and I2), we will see that `keydef` occurs two times below. There is a subordinate structure, `keyseg`, which also occurs multiple times (once within the `keydef` for I1 and two times within the `keydef` for I2).

Name	Size	Dump	From Example File	Comment

/* key definition for I1 */				
keydef->keysegs	1	01		there is 1 keyseg (for column S1).
keydef->key_alg	1	01		algorithm = Rtree or Btree
keydef->flag	2	00 49		HA_NOSAME + HA_SPACE_PACK_USED + HA_NULL_PART_KEY
keydef->block_length	2	04 00		i.e. 1024
key def->keylength	2	00 06		field-count+sizeof(S1) sizeof(ROWID)
keydef->minlength	2	00 06		
keydef->maxlength	2	00 06		
/* keyseg for S1 in I1 */				
keyseg->type	1	01		/* I1(S1) size(S1)=1, column = 1 */ = HA_KEYTYPE_TEXT
keyseg->language	1	08		
keyseg->null_bit	1	02		
keyseg->bit_start	1	00		
keyseg->bit_end	1	00		
[0] i.e. filler	1	00		
keyseg->flag	2	00 14		HA_NULL_PART + HA_PART_KEY
keyseg->length	2	00 01		length(S1) = 1
keyseg->start	4	00 00 00 01		offset in the row
keyseg->null_pos	4	00 00 00 00		
/* key definition for I2 */				
keydef->keysegs	1	02		keysegs=2, for columns S2 and S3
keydef->key_alg	1	01		algorithm = Rtree or

```

keydef->flag                2    00 48
keydef->block_length        2    04 00
key def->keylength          2    00 0B
keydef->minlength           2    00 0B
keydef->maxlength           2    00 0B
/* keyseg for S2 in I2 */
keyseg->type                 1    01
keyseg->language             1    08
keyseg->null_bit             1    04
keyseg->bit_start            1    00
keyseg->bit_end              1    00
[0] i.e. filler             1    00
keyseg->flag                 2    00 14
keyseg->length               2    00 02
keyseg->start                4    00 00 00 02
keyseg->null_pos             4    00 00 00 00
/* keyseg for S3 in I2 */
keyseg->type                 1    01
keyseg->language             1    08
keyseg->null_bit             1    08
keyseg->bit_start            1    00
keyseg->bit_end              1    00
[0] i.e. filler             1    00
keyseg->flag                 2    00 14
keyseg->length               2    00 03
keyseg->start                4    00 00 00 04
keyseg->null_pos             4    00 00 00 00

```

Btree
 HA_SPACE_PACK_USED +
 HA_NULL_PART_KEY
 i.e. 1024
 field-count+ sizeof(all fields)+
 sizeof(RID)

/* I2(S2) size(S2)=2,
 column = 2 */

HA_NULL_PART +
 HA_PART_KEY
 length(S2) = 2

/* I2(S3) size(S3)=3,
 column = 3 */

HA_NULL_PART +
 HA_PART_KEY
 length(S3) = 3

recinfo

The `recinfo` section is written by `mi_open.c`, `mi_recinfo_write()`. The corresponding structure in `mysamdef.h` is `MI_COLUMNDEF`.

This is another multiple-occurrence structure. It appears once for each field that appears in a key, including an extra field that appears at the start and has flags (for deletion and for null fields).

Name	Size	Dump From Example File	Comment
recinfo->type	2	00 00	extra
recinfo->length	2	00 01	
recinfo->null_bit	1	00	
recinfo->null_pos	2	00 00	
recinfo->type	2	00 00	I1 (S1)
recinfo->length	2	00 01	
recinfo->null_bit	1	02	
recinfo->null_pos	2	00 00	
recinfo->type	2	00 00	I2 (S2)
recinfo->length	2	00 02	
recinfo->null_bit	1	04	
recinfo->null_pos	2	00 00	
recinfo->type	2	00 00	I2 (S3)
recinfo->length	2	00 03	
recinfo->null_bit	1	08	
recinfo->null_pos	2	00 00	

We are now at offset 0xA2 within the file `T.MYI`. Notice that the value of the third field in the header, `header_length`, is 0xA2. Anything following this point, up till the first key value, is filler.

The .MYI Key Values

And now we look at the part which is not the information header: we look at the key values. The key values are in blocks (MySQL's term for pages). A block contains values from only one index. To continue our example: there is a block for the I1 key values, and a block for the I2 key values.

According to the header information (`state->key_root` above), the I1 block starts at offset 0x0400 in the file, and the I2 block starts at offset 0x0800 in the file.

At offset 0x0400 in the file, we have this:

Name ----	Size	Dump	From Example File	Comment
(block header)	2	00 0E		= size (inclusive) (first bit of word = 0 meaning this is a B-Tree leaf, see the mi_test_if_nod macro)
(first key value)	2	01 31		Value is "1" (0x31).
(first key pointer)	4	00 00 00 00		Pointer is to Record #0000.
(second key value)	2	01 33		Value is "3" (0x33).
(second key pointer)	4	00 00 00 02		Pointer is to Record #0002.
(junk)	1010		rest of the 1024-byte block is unused

At offset 0800x in the file, we have this:

Name ----	Size	Dump	From Example File	Comment
(block header)	2	00 18		= size (inclusive)
(first key value)	7	01 61 61 01 62 20 20		Value is "aa/b "
(first key pointer)	4	00 00 00 00		Pointer is to Record #0000.
(second key value)	7	01 61 61 01 62 62 62		Value is "aa/bbb"
(second key pointer)	4	00 00 00 02		Pointer is to Record #0002.
(junk)	1000		rest of the 1024-byte block is unused

From the above illustrations, these facts should be clear:

- Each key contains the entire contents of all the columns, including trailing spaces in `CHAR` columns. There is no front truncation. There is no back truncation. (There can be space truncation if `keyseg->flag HA_SPACE_PACK` flag is on.)
- For fixed-row tables: The pointer is a fixed-size (4-byte) number which contains an ordinal row number. The first row is Record #0000. This item is analogous to the ROWID, or RID (row identifier), which other DBMSs use. For dynamic-row tables: The pointer is an offset in the `.MYD` file.
- The normal block length is 0x0400 (1024) bytes.

These facts are not illustrated, but are also clear:

- If a key value is `NULL`, then the first byte is 0x00 (instead of 001 as in the preceding examples) and that's all. Even for a fixed `CHAR(3)` column, the size of the key value is only 1 byte.
- Initially the junk at the end of a block is filler bytes, value = A5. If MySQL shifts key values up after a `DELETE`, the end of the block is not overwritten.
- A normal block is at least 65% full, and typically 80% full. (This is somewhat denser than the typical B-tree algorithm would cause, it is thus because `myisamchk -rq` will make blocks nearly 100% full.)
- There is a pool of free blocks, which increases in size when deletions occur. If all blocks have the same normal block length (1024), then MySQL will always use the same pool.

- The maximum number of keys is 32 (`MI_MAX_KEY`). The maximum number of segments in a key is 16 (`MI_MAX_KEY_SEG`). The maximum key length is 500 (`MI_MAX_KEY_LENGTH`). The maximum block length is 16384 (`MI_MAX_KEY_BLOCK_LENGTH`). All these `MI_...` constants are expressed by `#defines` in the `myisamdef.h` file.

MyISAM Files

Some notes about `MyISAM` file handling:

- If a table is never updated, MySQL will never touch the table files, so it would never be marked as closed or corrupted.
- If a table is marked readonly by the OS, it will only be opened in readonly mode. Any updates to it will fail.
- When a normal table is opened for reading by a `SELECT`, MySQL will open it in read/write mode, but will not write anything to it.
- A table can be closed during one of the following events:
 - Out of space in table cache
 - Someone executed flush tables
 - MySQL was shut down
 - `flush_time` expired (which causes an automatic flush-tables to be executed)
- When MySQL opens a table, it checks if the table is clean. If it isn't and the server was started with the `--myisam-recover` option, check the table and try to recover it if it's crashed. (The safest automatic recover option is probably `--myisam-recover=BACKUP`.)

MyISAM Dynamic Data File Layout

Variable length records are contained in "frames". A record can be put in one or more frames, also called the record "parts" or "blocks". The sense of the frames is to allow reuse of the space of deleted records. Starting with an empty data file, records are put in a single frame each, unless a record is bigger than the maximum frame size (16MB - 4). When a record is deleted, its frame is marked deleted. When a record is inserted after this, it reuses the old frame. If the new record is smaller, the frame is split. The unused part is marked deleted. If a new record is bigger than the old frame, the frame is filled with the record as much as fits. The rest is inserted in other old frames, or, if non is available, in a new frame at end of file.

Layout of the Record Storage Frame (Record Part, Record Block)

```
MI_MIN_BLOCK_LENGTH  20          /* 20 bytes are required for the biggest frame type: Deleted block.
MI_MAX_BLOCK_LENGTH  16777212    /* 16MB - 4, max 3 bytes for length available, 4-byte aligned. */
MI_DYN_ALIGN_SIZE    4          /* Frames start a 4-byte boundaries. */
```

Part header[0] (decimal/hexadecimal, one byte):

```
0/00: Deleted block
      block_len 3 bytes [1-3]
      next_filepos 8 bytes [4-11]
      prev_filepos 8 bytes [12-19]
      => header length 20
```

```
1/01: Full small record, full block
      rec_len,data_len,block_len 2 bytes [1-2]
      => header length 3
```

```
2/02: Full big record, full block
      rec_len,data_len,block_len 3 bytes [1-3]
      => header length 4
```

```
3/03: Full small record, unused space
      rec_len,data_len 2 bytes [1-2]
      unused_len 1 byte [3]
      => header length 4
```

```
4/04: Full big record, unused space
      rec_len,data_len 3 bytes [1-3]
      unused_len 1 byte [4]
      => header length 5
```

```
5/05: Start small record
      rec_len 2 bytes [1-2]
      data_len,block_len 2 bytes [3-4]
      next_filepos 8 bytes [5-12]
      => header length 13
```

```
6/06: Start big record
      rec_len 3 bytes [1-3]
      data_len,block_len 3 bytes [4-6]
      next_filepos 8 bytes [7-14]
      => header length 15
```

```
7/07: End small record, full block
      data_len,block_len 2 bytes [1-2]
      => header length 3
```

```
8/08: End big record, full block
      data_len,block_len 3 bytes [1-3]
      => header length 4
```

```
9/09: End small record, unused space
      data_len 2 bytes [1-2]
      unused_len 1 byte [3]
      => header length 4
```

```
10/0A: End big record, unused space
      data_len 3 bytes [1-3]
      unused_len 1 byte [4]
      => header length 5
```

```
11/0B: Continue small record
      data_len,block_len 2 bytes [1-2]
      next_filepos 8 bytes [3-10]
      => header length 11
```

```
12/0C: Continue big record
      data_len,block_len 3 bytes [1-3]
      next_filepos 8 bytes [4-11]
      => header length 12
```

```
13/0D: Start giant record
      rec_len 4 bytes [1-4]
      data_len,block_len 3 bytes [5-7]
      next_filepos 8 bytes [8-15]
      => header length 16
```

block_len does not include the header length except of deleted blocks. In deleted blocks block_len includes the header length.

data_len is the number of bytes within this block that are part of the record.

rec_len is the total record length including the data lengths of all belonging blocks.

In deleted blocks next_filepos and prev_filepos make a doubly linked list over all deleted blocks. At list start, prev_filepos is HA_OFFSET_ERROR (all bits set). At list end, next_filepos is HA_OFFSET_ERROR (all bits set).

In non-deleted blocks next_filepos points to the next part of the record.

The read function for the block header of dynamic records is mi_dynrec.c:_mi_get_block_info().

Record Contents

This section is still to be written.

Packed Record Layout

```
pack bits (!= NULL bits!): One bit per packable column:
    FIELD_BLOB: Set if blob is empty. No blob data in record.
    FIELD_SKIP_ZERO: Set if all bytes zero, no data in record.
    FIELD_SKIP_ENDSPACE,
    FIELD_SKIP_PRESPACE: Set if some blanks stripped in record.
=> The "pack bits" are rounded up to the next byte boundary.
```

Record contents from the in-memory representation.
Each field is copied verbatim unless packed according
to the "pack bits" paragraph.

In-memory Record Layout

```
null bits:    One bit per column that can be NULL.
              The "null bits" are rounded up to the next byte boundary.
              the number of "null bytes" are also referred as data_offset.
```

```
fields:      One field per column. No separators, length fields, etc.
              Length (pack_length) and layout of the fields depend on the
              field type:
```

```
-- to be added --
```

Note: The "in-memory record layout" is used by the SQL layer. It is independent from storage engines. All storage engines have to accept and produce this in-memory format at the handler interface.

MyISAM Compressed Data File Layout

This chapter describes the layout for the data file of compressed [MyISAM](#) tables.

Huffman Compression

[MyISAM](#) compression is based on Huffman compression. In his article from 1952 Huffman proved that his algorithm uses the least possible number of bits to encode a sequence of messages. The number of bits assigned to each message depends on its probability to appear in the sequence.

Huffman did not specify exactly, what those "messages" are. One could take all possible values - say of a table column - as "messages". But if there are too many of them, the code tables could become bigger than the uncompressed table. One would need to specify every possible value once and the code tree with its indexes and offsets. Not to forget the effort to step through big binary trees for every value and - on the encoding side - the comparison of each value against the already collected distinct values.

The usual way to define "Huffman messages" is to take the possible 256 values, which a byte can express, as the "messages". That way the code trees are of limited size. On the other hand, the theoretical maximum compression is 1:8 (12.5%) in this case.

The [myisampack](#) Program

[myisampack](#) tries both ways to compress the column values. When starting to analyze the existing uncompressed data, it collects distinct column values up to a limit of 8KB. If there are more, it falls back to byte value compression for this column.

This means also that [myisampack](#) may use different algorithms for different columns. Besides a couple of other tricks, [myisampack](#) determines for every column if distinct column value compression or byte value compression is better. After that it tries to combine byte value compression trees of different columns into one or more code trees. This means that finally we may have less code trees

than columns. Therefore the column information in the file header contains the number of the code tree used for each column. Some columns might not need a code tree at all. This happens for columns which have the same value in all records.

Record and Blob Length Encoding

Since the compressed data file should be usable for read-only purposes by the MySQL database management system, every record starts on a byte boundary. For easier handling by the system, every record begins with a length information for the compressed record and a length information for the total size of all uncompressed blobs of this record. Both lengths are encoded in 1 to 5 bytes, depending on its value.

A length from 1 to 253 bytes is represented in one byte. A length of 254 to 65535 bytes (64KB-1) is represented by three bytes. The first contains the value 254 and the next two bytes contain the plain length. The low order byte goes first. A length of 65536 to 4294967295 bytes (4GB-1) is represented by five bytes. The first contains the value 255 and the next four bytes contain the plain length. The low order byte goes first.

The encoded compressed record length does not include these length bytes. It tells the number of bytes which follow behind the length bytes for this record.

Code Tree Representation

The code trees are binary trees. Every node has exactly two children. The children can be leaves or branches. A leaf contains one original, uncompressed value. A branch contains a pointer to another node. The Huffman codes represent the navigation through the tree. Every left branch gets a 0 bit, every right branch gets a 1 bit.

The in-memory representation of the trees are two unsigned integers per node. Each describes either a leaf value or an offset (in unsigned integers relative from this node) to another node. To distinguish values from offsets, the 15th bit (decimal value 32768) is set together with offsets. This is safe as the size of the trees is limited by either having a maximum of 256 elements for byte value compression or 4096 elements for distinct column value compression.

The representation of the trees in the compressed data file is almost the same. But instead of writing all bits of the unsigned integers, only as many bits are written as are required to represent the highest value or offset respectively. One more bit per value/offset is written in advance, to distinguish both. The number of bits required per value and per offset is computed in advance and part of the code tree description.

Usage of the Index File

While the header of the compressed data file contains a lot of information, there are still some things which need to be taken from the index file. These are the number of columns of the table and the length of each column. The latter is required for columns with suppressed leading spaces or suppressed trailing spaces or zeros.

`myisampack` Tricks

As already mentioned, `myisampack` uses some tricks to decrease the amount of data to be encoded. These cope with leading and trailing spaces or zeros or with all blank or `NULL` fields.

I do not describe these in detail here. They do not materialize in the compressed data files other than the later mentioned field and pack types. They are however important to know for decoding the records.

Detailed Specification of the Decoding

Below follows the detailed specification of the encoding:

Datafile fixed header (32 bytes):

```

4 byte  magic number
4 byte  total header length (fixed + column info + code trees)
4 byte  minimum packed record length
4 byte  maximum packed record length
4 byte  total number of elements in all code trees
4 byte  total number of bytes collected for distinct column values
2 byte  number of code trees
1 byte  maximum number of bytes required to represent record+blob lengths
1 byte  record pointer length, number of bytes for compressed data file length
4 byte  zeros

```

Column Information. For every column in the table:

```

5 bits  field type
        FIELD_NORMAL          0
        FIELD_SKIP_ENDSPACE   1
        FIELD_SKIP_PRESPACE   2
        FIELD_SKIP_ZERO       3
        FIELD_BLOB             4
        FIELD_CONSTANT         5
        FIELD_INTERVALL        6
        FIELD_ZERO             7
        FIELD_VARCHAR          8
        FIELD_CHECK            9

6 bits  pack type as a set of flags
        PACK_TYPE_SELECTED     1
        PACK_TYPE_SPACE_FIELDS 2
        PACK_TYPE_ZERO_FILL    4

5 bits  if pack type contains PACK_TYPE_ZERO_FILL
        minimum number of trailing zero bytes in this column
        else
        number of bits to encode the number of
        packed bytes in this column (length_bits)

x bits  number of the code tree used to encode this column
        x is the minimum number of bits required to represent the highest
        tree number.

```

Alignment:

```

x bits alignment to the next byte border

```

Code Trees. For every tree:

```

1 bit  compression type
        0 = byte value compression
            8 bits  minimum byte value coded by this tree
            9 bits  number of byte values encoded by this tree
            5 bits  number of bits used to encode the byte values
            5 bits  number of bits used to encode offsets to next tree elements
        1 = distinct column value compression
            15 bits number of distinct column values encoded by this tree
            16 bits length of the buffer with all column values
            5 bits  number of bits used to encode the index of the column value
            5 bits  number of bits used to encode offsets to next tree elements

For each code tree element:
    1 bit  IS_OFFSET
    x bits the announced number of bits for either a value or an offset
x bits alignment to the next byte border
If compression by distinct column values:
    The number of 8-bit values that make up the column value buffer

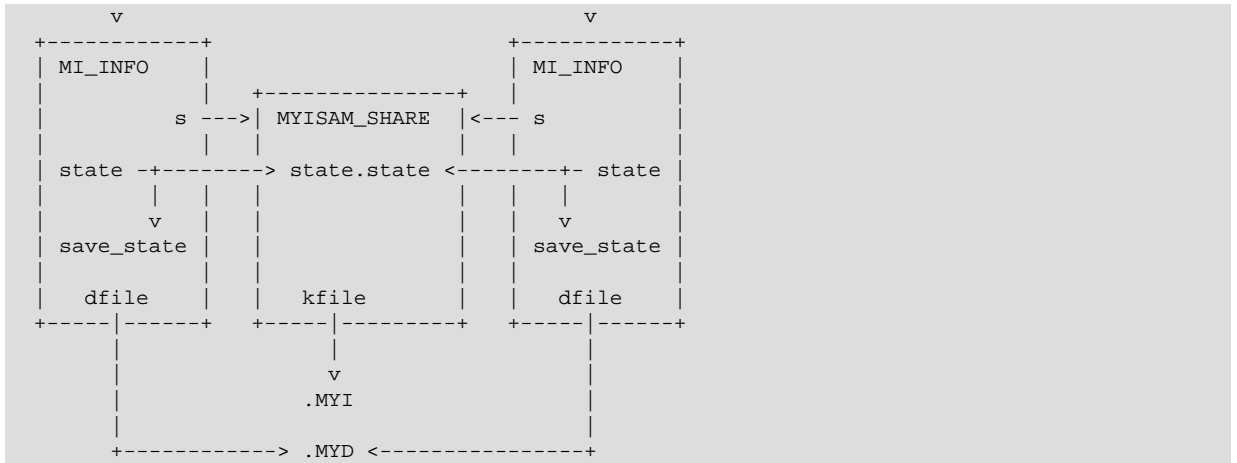
```

Compressed Records. For every record:

```

1-5 bytes length of the compressed record in bytes
    1. byte  0..253 length
           254  length encoded in the next two bytes little endian

```

MI_INFO::state may either point to MI_INFO::save_state or to MYISAM_SHARE::state.state. The latter is the normal case. Amongst others, state contains data_file_length.

To support concurrent inserts, every statement starts with copying MYISAM_SHARE::state.state to MI_INFO::save_state and lets MI_INFO::state point to the copy. This is done in `mi_get_status()`. This is called from the hook `THR_LOCK::(*get_status)()`. Some of the hooks are explained in `thr_lock.c`. `(*get_status)()` is called when a thread gets a lock.

The copy back is done in `mi_update_status()`, which is called from `mi_lock_database()` when unlocking from a write lock. This, in turn, is called from `ha_myisam::external_lock()` from `unlock_external()` from `mysql_unlock_tables()`. Until 5.1 this was done after `thr_multi_unlock()`. So it was possible that another thread (or even multiple of them) could `thr_lock` the table and work with it before the first thread updates the status. In 6.0 the order is reversed. The status should now be accurate when another thread acquires a `thr_lock`.

However, with concurrent inserts, the trick is that read locks are allowed to proceed while a concurrent insert holds a write lock. So it can copy outdated information when entering the lock. Since it works on its local copy of the state, it won't notice rows that are made available through `mi_update_status()` after it got the lock.

But there is another chance to miss the row(s). See also [Bug#36618](#) (myisam insert not immediately visible to select from another client). When the concurrent insert ends, it reports success to its client application before it unlocks the table. So there is non-deterministic time span between the seemingly successful ended insert and the final update of the MyISAM status.

Chapter 23. InnoDB Storage Engine

Table of Contents

InnoDB Record Structure	265
High-Altitude Picture	265
Where to Look for More Information	268
InnoDB Page Structure	268
High-Altitude View	269
Example	272
Where to Look For More Information	273

InnoDB Record Structure

This page contains:

- A high-altitude "summary" picture of the parts of a MySQL/InnoDB record structure.
- A description of each part.
- An example.

After reading this page, you will know how MySQL/InnoDB stores a physical record.

High-Altitude Picture

The chart below shows the three parts of a physical record.

Name	Size
Field Start Offsets	(F*1) or (F*2) bytes
Extra Bytes	6 bytes
Field Contents	depends on content

Legend: The letter 'F' stands for 'Number Of Fields'.

The meaning of the parts is as follows:

- The FIELD START OFFSETS is a list of numbers containing the information "where a field starts".
- The EXTRA BYTES is a fixed-size header.
- The FIELD CONTENTS contains the actual data.

An Important Note About The Word "Origin"

The "Origin" or "Zero Point" of a record is the first byte of the Field Contents --- not the first byte of the Field Start Offsets. If there is a pointer to a record, that pointer is pointing to the Origin. Therefore the first two parts of the record are addressed by subtracting from the pointer, and only the third part is addressed by adding to the pointer.

FIELD START OFFSETS

The Field Start Offsets is a list in which each entry is the position, relative to the Origin, of the start of the next field. The entries are in reverse order, that is, the first field's offset is at the end of the list.

An example: suppose there are three columns. The first column's length is 1, the second column's length is 2, and the third column's length is 4. In this case, the offset values are, respectively, 1, 3

(1+2), and 7 (1+2+4). Because values are reversed, a core dump of the Field Start Offsets would look like this: 07, 03, 01.

There are two complications for special cases:

- Complication #1: The size of each offset can be either one byte or two bytes. One-byte offsets are only usable if the total record size is less than 127. There is a flag in the "Extra Bytes" part which will tell you whether the size is one byte or two bytes.
- Complication #2: The most significant bits of an offset may contain flag values. The next two paragraphs explain what the contents are.

When The Size Of Each Offset Is One Byte

- 1 bit = `NULL`, = NULL
- 7 bits = the actual offset, a number between 0 and 127

When The Size Of Each Offset Is Two Bytes

- 1 bit = `NULL`, = NULL
- 1 bit = 0 if field is on same page as offset, = 1 if field and offset are on different pages
- 14 bits = the actual offset, a number between 0 and 16383

It is unlikely that the "field and offset are on different pages" unless the record contains a large `BLOB`.

EXTRA BYTES

The Extra Bytes are a fixed six-byte header.

Name	Size	Description
info_bits:	??	??
()	1 bit	unused or unknown
()	1 bit	unused or unknown
deleted_flag	1 bit	1 if record is deleted
min_rec_flag	1 bit	1 if record is predefined minimum record
n_owned	4 bits	number of records owned by this record
heap_no	13 bits	record's order number in heap of index page
n_fields	10 bits	number of fields in this record, 1 to 1023
1byte_offs_flag	1 bit	1 if each Field Start Offsets is 1 byte long (this item is also called the "short" flag)
next 16 bits	16 bits	pointer to next record in page
TOTAL	48 bits	??

Total size is 48 bits, which is six bytes.

If you're just trying to read the record, the key bit in the Extra Bytes is `1byte_offs_flag` — you need to know if `1byte_offs_flag` is 1 (i.e.: "short 1-byteoffsets") or 0 (i.e.: "2-byte offsets").

Given a pointer to the Origin, `InnoDB` finds the start of the record as follows:

- Let $X = n_fields$ (the number of fields is by definition equal to the number of entries in the Field Start Offsets Table).
- If `1byte_offs_flag` equals 0, then let $X = X * 2$ because there are two bytes for each entry instead of just one.

- Let $X = X + 6$, because the fixed size of Extra Bytes is 6.
- The start of the record is at (pointer value minus X).

FIELD CONTENTS

The Field Contents part of the record has all the data. Fields are stored in the order they were defined in.

There are no markers between fields, and there is no marker or filler at the end of a record.

Here's an example.

- I made a table with this definition:

```
CREATE TABLE T
  (FIELD1 VARCHAR(3), FIELD2 VARCHAR(3), FIELD3 VARCHAR(3))
  Type=InnoDB;
```

To understand what follows, you must know that table `T` has six columns — not three — because `InnoDB` automatically added three "system columns" at the start for its own housekeeping. It happens that these system columns are the row ID, the transaction ID, and the rollback pointer, but their values don't matter now. Regard them as three black boxes.

- I put some rows in the table. My last three `INSERT` statements were:

```
INSERT INTO T VALUES ('PP', 'PP', 'PP');
INSERT INTO T VALUES ('Q', 'Q', 'Q');
INSERT INTO T VALUES ('R', NULL, NULL);
```

- I ran Borland's TDUMP to get a hexadecimal dump of the contents of `\mysql\data\ibdata1`, which (in my case) is the MySQL/`InnoDB` data file (on Windows).

Here is an extract of the dump:

Address Values in Hexadecimal	Values in ASCII
0D4280: 00 00 2D 00 84 4F 4F 4F 4F 4F 4F 4F 4F 4F 19 17	..-..0000000000..
0D4290: 15 13 0C 06 00 00 78 0D 02 BF 00 00 00 00 04 21x.....!
0D42A0: 00 00 00 00 09 2A 80 00 00 00 2D 00 84 50 50 50*.....-..PPP
0D42B0: 50 50 50 16 15 14 13 0C 06 00 00 80 0D 02 E1 00	PPP.....
0D42C0: 00 00 00 04 22 00 00 00 00 09 2B 80 00 00 00 2D".....+.....-
0D42D0: 00 84 51 51 51 94 94 14 13 0C 06 00 00 88 0D 00	..QQQ.....
0D42E0: 74 00 00 00 00 04 23 00 00 00 00 09 2C 80 00 00	t.....#.....,
0D42F0: 00 2D 00 84 52 00 00 00 00 00 00 00 00 00 00 00	.-..R.....

A reformatted version of the dump, showing only the relevant bytes, looks like this (I've put a line break after each field and added labels):

Reformatted Hexadecimal Dump

```
19 17 15 13 0C 06 Field Start Offsets /* First Row */
```

```

00 00 78 0D 02 BF Extra Bytes
00 00 00 00 04 21 System Column #1
00 00 00 00 09 2A System Column #2
80 00 00 00 2D 00 84 System Column #3
50 50 Field1 'PP'
50 50 Field2 'PP'
50 50 Field3 'PP'

16 15 14 13 0C 06 Field Start Offsets /* Second Row */
00 00 80 0D 02 E1 Extra Bytes
00 00 00 00 04 22 System Column #1
00 00 00 00 09 2B 80 System Column #2
00 00 00 2D 00 84 System Column #3
51 Field1 'Q'
51 Field2 'Q'
51 Field3 'Q'

94 94 14 13 0C 06 Field Start Offsets /* Third Row */
00 00 88 0D 00 74 Extra Bytes
00 00 00 00 04 23 System Column #1
00 00 00 00 09 2C System Column #2
80 00 00 00 2D 00 84 System Column #3
52 Field1 'R'

```

You won't need explanation if you followed everything I've said, but I'll add helpful notes for the three trickiest details.

- Helpful Notes About "Field Start Offsets":

Notice that the sizes of the record's fields, in forward order, are: 6, 6, 7, 2, 2, 2. Since each offset is for the start of the "next" field, the hexadecimal offsets are 06, 0c (6+6), 13 (6+6+7), 15 (6+6+7+2), 17 (6+6+7+2+2), 19 (6+6+7+2+2+2). Reversing the order, the Field Start Offsets of the first record are: [19,17,15,13,0c,06](#).

- Helpful Notes About "Extra Bytes":

Look at the Extra Bytes of the first record: [00 00 78 0D 02 BF](#). The fourth byte is [0D hexadecimal](#), which is [1101 binary](#) ... the 110 is the last bits of `n_fields` ([110 binary](#) is 6 which is indeed the number of fields in the record) and the final 1 bit is `1byte_offs_flag`. The fifth and sixth bytes, which contain [02 BF](#), constitute the "next" field. Looking at the original hexadecimal dump, at address [0D42BF](#) (which is position [02BF](#) within the page), you'll see the beginning bytes of System Column #1 of the second row. In other words, the "next" field points to the "Origin" of the following row.

- Helpful Notes About NULLs:

For the third row, I inserted [NULLs](#) in FIELD2 and FIELD3. Therefore in the Field Start Offsets the top bit is [on](#) for these fields (the values are [94 hexadecimal](#), [94 hexadecimal](#), instead of [14 hexadecimal](#), [14 hexadecimal](#)). And the row is shorter because the [NULLs](#) take no space.

Where to Look for More Information

References:

The most relevant [InnoDB](#) source-code files are [rem0rec.c](#), [rem0rec.ic](#), and [rem0rec.h](#) in the [rem](#) ("Record Manager") directory.

InnoDB Page Structure

[InnoDB](#) stores all records inside a fixed-size unit which is commonly called a "page" (though [InnoDB](#) sometimes calls it a "block" instead). Currently all pages are the same size, 16KB.

A page contains records, but it also contains headers and trailers. I'll start this description with a high-altitude view of a page's parts, then I'll describe each part of a page. Finally, I'll show an example. This discussion deals only with the most common format, for the leaf page of a data file.

High-Altitude View

An [InnoDB](#) page has seven parts:

- Fil Header
- Page Header
- Infimum + Supremum Records
- User Records
- Free Space
- Page Directory
- Fil Trailer

As you can see, a page has two header/trailer pairs. The inner pair, "Page Header" and "Page Directory", are mostly the concern of the `\page` program group, while the outer pair, "Fil Header" and "Fil Trailer", are mostly the concern of the `\fil` program group. The "Fil" header also goes by the name of "File Page Header".

Sandwiched between the headers and trailers, are the records and the free (unused) space. A page always begins with two unchanging records called the Infimum and the Supremum. Then come the user records. Between the user records (which grow downwards) and the page directory (which grows upwards) there is space for new records.

Fil Header

The Fil Header has eight parts, as follows:

Name	Size	Remarks
FIL_PAGE_SPACE	4	4 ID of the space the page is in
FIL_PAGE_OFFSET	4	ordinal page number from start of space
FIL_PAGE_PREV	4	offset of previous page in key order
FIL_PAGE_NEXT	4	offset of next page in key order
FIL_PAGE_LSN	8	log serial number of page's latest log record
FIL_PAGE_TYPE	2	current defined types are: FIL_PAGE_INDEX , FIL_PAGE_UNDO_LOG , FIL_PAGE_INODE , FIL_PAGE_IBUF_FREE_LIST
FIL_PAGE_FILE_FLUSH_LSN	8	"the file has been flushed to disk at least up to this lsn" (log serial number), valid only on the first page of the file
FIL_PAGE_ARCH_LOG_NO	4	the latest archived log file number at the time that FIL_PAGE_FILE_FLUSH_LSN was written (in the log)

- [FIL_PAGE_SPACE](#) is a necessary identifier because different pages might belong to different (table) spaces within the same file. The word "space" is generic jargon for either "log" or "tablespace".
- [FIL_PAGE_PREV](#) and [FIL_PAGE_NEXT](#) are the page's "backward" and "forward" pointers. To show what they're about, I'll draw a two-level B-tree.

```

-----
- root -
-----
  |
  |-----
  |
  |

```

```

|           |
|-----|   |-----|
| - leaf - <--> - leaf - |
|-----|   |-----|

```

Everyone has seen a B-tree and knows that the entries in the root page point to the leaf pages. (I indicate those pointers with vertical '|' bars in the drawing.) But sometimes people miss the detail that leaf pages can also point to each other (I indicate those pointers with a horizontal two-way pointer '<-->' in the drawing). This feature allows [InnoDB](#) to navigate from leaf to leaf without having to back up to the root level. This is a sophistication which you won't find in the classic B-tree, which is why [InnoDB](#) should perhaps be called a B+-tree instead.

- The fields [FIL_PAGE_FILE_FLUSH_LSN](#), [FIL_PAGE_PREV](#), and [FIL_PAGE_NEXT](#) all have to do with logs, so I'll refer you to my article "How Logs Work With MySQL And InnoDB" on [devarticles.com](#).
- [FIL_PAGE_FILE_FLUSH_LSN](#) and [FIL_PAGE_ARCH_LOG_NO](#) are valid only for the first page of a data file.

Page Header

The Page Header has 14 parts, as follows:

Name	Size	Remarks
PAGE_N_DIR_SLOTS	2	number of directory slots in the Page Directory part; initial value = 2
PAGE_HEAP_TOP	2	record pointer to first record in heap
PAGE_N_HEAP	2	number of heap records; initial value = 2
PAGE_FREE	2	record pointer to first free record
PAGE_GARBAGE	2	"number of bytes in deleted records"
PAGE_LAST_INSERT	2	record pointer to the last inserted record
PAGE_DIRECTION	2	either PAGE_LEFT , PAGE_RIGHT , or PAGE_NO_DIRECTION
PAGE_N_DIRECTION	2	number of consecutive inserts in the same direction, e.g. "last 5 were all to the left"
PAGE_N_RECS	2	number of user records
PAGE_MAX_TRX_ID	8	the highest ID of a transaction which might have changed a record on the page (only set for secondary indexes)
PAGE_LEVEL	2	level within the index (0 for a leaf page)
PAGE_INDEX_ID	8	identifier of the index the page belongs to
PAGE_BTR_SEG_LEAF	10	"file segment header for the leaf pages in a B-tree" (this is irrelevant here)
PAGE_BTR_SEG_TOP	10	"file segment header for the non-leaf pages in a B-tree" (this is irrelevant here)

(Note: I'll clarify what a "heap" is when I discuss the User Records part of the page.)

Some of the Page Header parts require further explanation:

- [PAGE_FREE](#) :

Records which have been freed (due to deletion or migration) are in a one-way linked list. The [PAGE_FREE](#) pointer in the page header points to the first record in the list. The "next" pointer in the record header (specifically, in the record's Extra Bytes) points to the next record in the list.

- [PAGE_DIRECTION](#) and [PAGE_N_DIRECTION](#):

It's useful to know whether inserts are coming in a constantly ascending sequence. That can affect [InnoDB's](#) efficiency.

- [PAGE_HEAP_TOP](#) and [PAGE_FREE](#) and [PAGE_LAST_INSERT](#):

Warning: Like all record pointers, these point not to the beginning of the record but to its Origin (see the earlier discussion of Record Structure).

- [PAGE_BTR_SEG_LEAF](#) and [PAGE_BTR_SEG_TOP](#):

These variables contain information (space ID, page number, and byte offset) about index node file segments. [InnoDB](#) uses the information for allocating new pages. There are two different variables because [InnoDB](#) allocates separately for leaf pages and upper-level pages.

The Infimum and Supremum Records

"Infimum" and "supremum" are real English words but they are found only in arcane mathematical treatises, and in [InnoDB](#) comments. To [InnoDB](#), an infimum is lower than the lowest possible real value (negative infinity) and a supremum is greater than the greatest possible real value (positive infinity). [InnoDB](#) sets up an infimum record and a supremum record automatically at page-create time, and never deletes them. They make a useful barrier to navigation so that "get-prev" won't pass the beginning and "get-next" won't pass the end. Also, the infimum record can be a dummy target for temporary record locks.

The [InnoDB](#) code comments distinguish between "the infimum and supremum records" and the "user records" (all other kinds).

It's sometimes unclear whether [InnoDB](#) considers the infimum and supremum to be part of the header or not. Their size is fixed and their position is fixed, so I guess so.

User Records

In the User Records part of a page, you'll find all the records that the user inserted.

There are two ways to navigate through the user records, depending whether you want to think of their organization as an unordered or an ordered list.

An unordered list is often called a "heap". If you make a pile of stones by saying "whichever one I happen to pick up next will go on top" — rather than organizing them according to size and colour — then you end up with a heap. Similarly, [InnoDB](#) does not want to insert new rows according to the B-tree's key order (that would involve expensive shifting of large amounts of data), so it inserts new rows right after the end of the existing rows (at the top of the Free Space part) or wherever there's space left by a deleted row.

But by definition the records of a B-tree must be accessible in order by key value, so there is a record pointer in each record (the "next" field in the Extra Bytes) which points to the next record in key order. In other words, the records are a one-way linked list. So [InnoDB](#) can access rows in key order when searching.

Free Space

I think it's clear what the Free Space part of a page is, from the discussion of other parts.

Page Directory

The Page Directory part of a page has a variable number of record pointers. Sometimes the record pointers are called "slots" or "directory slots". Unlike other DBMSs, [InnoDB](#) does not have a slot for every record in the page. Instead it keeps a sparse directory. In a fullish page, there will be one slot for every six records.

The slots track the records' logical order (the order by key rather than the order by placement on the heap). Therefore, if the records are 'A' 'B' 'F' 'D' the slots will be (pointer to 'A') (pointer

to 'B') (pointer to 'D') (pointer to 'F'). Because the slots are in key order, and each slot has a fixed size, it's easy to do a binary search of the records on the page via the slots.

(Since the Page Directory does not have a slot for every record, binary search can only give a rough position and then InnoDB must follow the "next" record pointers. InnoDB's "sparse slots" policy also accounts for the `n_owned` field in the Extra Bytes part of a record: `n_owned` indicates how many more records must be gone through because they don't have their own slots.)

Fil Trailer

The Fil Trailer has one part, as follows:

Name	Size	Remarks
<code>FIL_PAGE_END_LSN</code>	8	low 4 bytes = checksum of page, last 4 bytes = same as <code>FIL_PAGE_LSN</code>

The final part of a page, the fil trailer (or File Page Trailer), exists because InnoDB's architect worried about integrity. It's impossible for a page to be only half-written, or corrupted by crashes, because the log-recovery mechanism restores to a consistent state. But if something goes really wrong, then it's nice to have a checksum, and to have a value at the very end of the page which must be the same as a value at the very beginning of the page.

Example

For this example, I used Borland's TDUMP again, as I did for the earlier chapter on Record Format. This is what a page looked like:

Address Values in Hexadecimal	Values in ASCII
0D4000: 00 00 00 00 00 00 00 00 35 FF FF FF FF FF FF FF FF5.....
0D4010: 00 00 00 00 00 00 00 00 E2 64 45 BF 00 00 00 00 00 00dE.....
0D4020: 00 00 00 00 00 00 00 00 05 02 F5 00 12 00 00 00 00
0D4030: 02 E1 00 02 00 0F 00 10 00 00 00 00 00 00 00 00 00
0D4040: 00 00 00 00 00 00 00 00 00 00 14 00 00 00 00 00 00
0D4050: 00 02 16 B2 00 00 00 00 00 00 00 00 02 15 F2 08 01
0D4060: 00 00 03 00 89 69 6E 66 69 6D 75 6D 00 09 05 00infimum....
0D4070: 08 03 00 00 73 75 70 72 65 6D 75 6D 00 22 1D 18supremum."..
0D4080: 13 0C 06 00 00 10 0D 00 B7 00 00 00 00 00 04 14 00
0D4090: 00 00 00 09 1D 80 00 00 00 00 2D 00 84 41 41 41 41-...AAAA
0D40A0: 41 41 41 41 41 41 41 41 41 41 41 41 1F 1B 17 13 0C	AAAAAAAAAAAAA.....
...	??
...	??

0D7FE0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 74t
0D7FF0: 02 47 01 AA 01 0A 00 65 3A E0 AA 71 00 00 E2 64	.G.....e:...q...d

Let's skip past the first 38 bytes, which are Fil Header. The bytes of the Page Header start at location `0d4026` hexadecimal:

Location	Name	Description
00 05	PAGE_N_DIR_SLOTS	There are 5 directory slots.
02 F5	PAGE_HEAP_TOP	At location <code>0402F5</code> , not shown, is the beginning of free space. Maybe a better name would have been <code>PAGE_HEAP_END</code> .
00 12	PAGE_N_HEAP	There are 18 (hexadecimal 12) records in the page.
00 00	PAGE_FREE	There are zero free (deleted) records.
00 00	PAGE_GARBAGE	There are zero bytes in deleted records.
02 E1	PAGE_LAST_INSERT	The last record was inserted at location <code>02E1</code> , not shown, within the page.
00 02	PAGE_DIRECTION	A glance at <code>page0page.h</code> will tell you that 2 is the #defined value for <code>PAGE_RIGHT</code> .
00 0F	PAGE_N_DIRECTION	The last 15 (hexadecimal 0F) inserts were all done "to the right" because I was inserting in ascending order.
00 10	PAGE_N_RECS	There are 16 (hexadecimal 10) user records. Notice that <code>PAGE_N_RECS</code> is smaller than the earlier field, <code>PAGE_N_HEAP</code> .
00 00 00 00 00 00 00 00	PAGE_MAX_TRX_ID	??
00 00	PAGE_LEVEL	Zero because this is a leaf page.
00 00 00 00 00 00 00 00 00 14	PAGE_INDEX_ID	This is index number 20.
00 00 00 00 00 00 00 00 00 02 16 B2	PAGE_BTR_SEG_LEAF	??
00 00 00 00 00 00 00 00 00 02 15 F2	PAGE_BTR_SEG_TOP	??

Immediately after the page header are the infimum and supremum records. Looking at the "Values In ASCII" column in the hexadecimal dump, you will see that the contents are in fact the words "infimum" and "supremum" respectively.

Skipping past the User Records and the Free Space, many bytes later, is the end of the 16KB page. The values shown there are the two trailers.

- The first trailer (`00 74, 02 47, 01 AA, 01 0A, 00 65`) is the page directory. It has 5 entries, because the header field `PAGE_N_DIR_SLOTS` says there are 5.
- The next trailer (`3A E0 AA 71, 00 00 E2 64`) is the fil trailer. Notice that the last four bytes, `00 00 E2 64`, appeared before in the fil header.

Where to Look For More Information

References:

The most relevant InnoDB source-code files are `page0page.c`, `page0page.ic`, and `page0page.h` in the `page` directory.

Chapter 24. Writing a Custom Storage Engine

Table of Contents

Additional Resources	276
Overview	276
Creating Storage Engine Source Files	276
Adding Engine Specific Variables and Parameters	277
Creating the <code>handlerton</code>	277
Handling Handler Instantiation	280
Defining Filename Extensions	280
Creating Tables	281
Opening a Table	282
Implementing Basic Table Scanning	282
Implementing the <code>store_lock()</code> Method	283
Implementing the <code>external_lock()</code> Method	284
Implementing the <code>rnd_init()</code> Method	284
Implementing the <code>info(uinf flag)</code> Method	284
Implementing the <code>extra()</code> Method	285
Implementing the <code>rnd_next()</code> Method	285
Closing a Table	287
Adding Support for <code>INSERT</code> to a Storage Engine	287
Adding Support for <code>UPDATE</code> to a Storage Engine	288
Adding Support for <code>DELETE</code> to a Storage Engine	288
Supporting Non-Sequential Reads	289
Implementing the <code>position()</code> Method	289
Implementing the <code>rnd_pos()</code> Method	289
Supporting Indexing	289
Indexing Overview	290
Getting Index Information During <code>CREATE TABLE</code> Operations	290
Creating Index Keys	291
Parsing Key Information	291
Providing Index Information to the Optimizer	291
Preparing for Index Use with <code>index_init()</code>	293
Cleaning up with <code>index_end()</code>	293
Implementing the <code>index_read()</code> Method	293
Implementing the <code>index_read_idx()</code> Method	294
Implementing the <code>index_read_last()</code> Method	294
Implementing the <code>index_next()</code> Method	294
Implementing the <code>index_prev()</code> Method	294
Implementing the <code>index_first()</code> Method	294
Implementing the <code>index_last()</code> Method	295
Supporting Transactions	295
Transaction Overview	295
Starting a Transaction	295
Implementing ROLLBACK	297
Implementing COMMIT	297
Adding Support for Savepoints	298
API Reference	299
<code>bas_ext</code>	299
<code>close</code>	300
<code>create</code>	300
<code>delete_row</code>	301
<code>delete_table</code>	302
<code>external_lock</code>	303
<code>extra</code>	303

index_end	304
index_first	305
index_init	305
index_last	306
index_next	306
index_prev	307
index_read	308
index_read_idx	308
index_read_last	309
info	310
open	311
position	312
records_in_range	313
rnd_init	314
rnd_next	314
rnd_pos	315
start_stmt	316
store_lock	316
update_row	318
write_row	319
FAQ	320

With MySQL 5.1, MySQL AB has introduced a pluggable storage engine architecture that makes it possible to create new storage engines and add them to a running MySQL server without recompiling the server itself.

This architecture makes it easier to develop new storage engines for MySQL and deploy them.

This chapter is intended as a guide to assist you in developing a storage engine for the new pluggable storage engine architecture.

Additional Resources

- A forum dedicated to custom storage engines is available at <http://forums.mysql.com/list.php?94>.

Overview

The MySQL server is built in a modular fashion:

The storage engines manage data storage and index management for MySQL. The MySQL server communicates with the storage engines through a defined API.

Each storage engine is a class with each instance of the class communicating with the MySQL server through a special [handler](#) interface.

Handlers are instantiated on the basis of one handler for each thread that needs to work with a specific table. For example: If three connections all start working with the same table, three handler instances will need to be created.

Once a handler instance is created, the MySQL server issues commands to the handler to perform data storage and retrieval tasks such as opening a table, manipulating rows, and managing indexes.

Custom storage engines can be built in a progressive manner: Developers can start with a read-only storage engine and later add support for [INSERT](#), [UPDATE](#), and [DELETE](#) operations, and even later add support for indexing, transactions, and other advanced operations.

Creating Storage Engine Source Files

The easiest way to implement a new storage engine is to begin by copying and modifying the [EXAMPLE](#) storage engine. The files [ha_example.cc](#) and [ha_example.h](#) can be found in the [storage/](#)

`example` directory of the MySQL 5.1 source tree. For instructions on how to obtain the 5.1 source tree, see [MySQL Installation Using a Source Distribution](#).

When copying the files, change the names from `ha_example.cc` and `ha_example.h` to something appropriate to your storage engine, such as `ha_foo.cc` and `ha_foo.h`.

After you have copied and renamed the files you must replace all instances of `EXAMPLE` and `example` with the name of your storage engine. If you are familiar with `sed`, these steps can be done automatically (in this example, the name of your storage engine would be `FOO`):

```
sed -e s/EXAMPLE/FOO/g -e s/example/foo/g ha_example.h > ha_foo.h
sed -e s/EXAMPLE/FOO/g -e s/example/foo/g ha_example.cc > ha_foo.cc
```

Adding Engine Specific Variables and Parameters

For more information on adding system variables, status variables, and options supported on the command line or configuration file, see [Specifying `mysqld` Variables Within a Plugin](#).

Creating the `handler_ton`

The `handler_ton` (short for handler singleton) defines the storage engine and contains method pointers to those methods that apply to the storage engine as a whole, as opposed to methods that work on a per-table basis. Some examples of such methods include transaction methods to handle commits and rollbacks.

Here's an example from the `EXAMPLE` storage engine:

```
handler_ton example_hton= {
    "EXAMPLE",
    SHOW_OPTION_YES,
    "Example storage engine",
    DB_TYPE_EXAMPLE_DB,
    NULL,      /* Initialize */
    0,         /* slot */
    0,         /* savepoint size. */
    NULL,      /* close_connection */
    NULL,      /* savepoint */
    NULL,      /* rollback to savepoint */
    NULL,      /* release savepoint */
    NULL,      /* commit */
    NULL,      /* rollback */
    NULL,      /* prepare */
    NULL,      /* recover */
    NULL,      /* commit_by_xid */
    NULL,      /* rollback_by_xid */
    NULL,      /* create_cursor_read_view */
    NULL,      /* set_cursor_read_view */
    NULL,      /* close_cursor_read_view */
    example_create_handler, /* Create a new handler */
    NULL,      /* Drop a database */
    NULL,      /* Panic call */
    NULL,      /* Release temporary latches */
    NULL,      /* Update Statistics */
    NULL,      /* Start Consistent Snapshot */
    NULL,      /* Flush logs */
    NULL,      /* Show status */
    NULL,      /* Replication Report Sent Binlog */
    HTON_CAN_RECREATE
};
```

This is the definition of the `handler_ton` from `handler.h`:

```
typedef struct
{
    const char *name;
    SHOW_COMP_OPTION state;
```

```
const char *comment;
enum db_type db_type;
bool (*init)();
uint slot;
uint savepoint_offset;
int (*close_connection)(THD *thd);
int (*savepoint_set)(THD *thd, void *sv);
int (*savepoint_rollback)(THD *thd, void *sv);
int (*savepoint_release)(THD *thd, void *sv);
int (*commit)(THD *thd, bool all);
int (*rollback)(THD *thd, bool all);
int (*prepare)(THD *thd, bool all);
int (*recover)(XID *xid_list, uint len);
int (*commit_by_xid)(XID *xid);
int (*rollback_by_xid)(XID *xid);
void *(*create_cursor_read_view)();
void (*set_cursor_read_view)(void *);
void (*close_cursor_read_view)(void *);
handler *(*create)(TABLE *table);
void (*drop_database)(char* path);
int (*panic)(enum ha_panic_function flag);
int (*release_temporary_latches)(THD *thd);
int (*update_statistics)();
int (*start_consistent_snapshot)(THD *thd);
bool (*flush_logs)();
bool (*show_status)(THD *thd, stat_print_fn *print, enum ha_stat_type stat);
int (*repl_report_sent_binlog)(THD *thd, char *log_file_name, my_off_t end_offset);
uint32 flags;
} handlerton;
```

There are a total of 30 handlerton elements, only a few of which are mandatory (specifically the first four elements and the `create()` method).

1. The name of the storage engine. This is the name that will be used when creating tables (`CREATE TABLE ... ENGINE = FOO;`).
2. The value to be displayed in the `status` field when a user issues the `SHOW STORAGE ENGINES` command.
3. The storage engine comment, a description of the storage engine displayed when using the `SHOW STORAGE ENGINES` command.
4. An integer that uniquely identifies the storage engine within the MySQL server. The constants used by the built-in storage engines are defined in the `handler.h` file. Custom engines should use `DB_TYPE_UNKNOWN`.
5. A method pointer to the storage engine initializer. This method is only called once when the server starts to allow the storage engine class to perform any housekeeping that is necessary before handlers are instantiated.
6. The slot. Each storage engine has its own memory area (actually a pointer) in the `thd`, for storing per-connection information. It is accessed as `thd->ha_data[foo_hnton.slot]`. The slot number is initialized by MySQL after `foo_init()` is called. For more information on the `thd`, see [#Implementing ROLLBACK](#).
7. The savepoint offset. To store per-savepoint data the storage engine is provided with an area of a requested size (0, if no savepoint memory is necessary).

The savepoint offset must be initialized statically to the size of the needed memory to store per-savepoint information. After `foo_init` it is changed to be an offset to the savepoint storage area and need not be used by the storage engine. For more information, see [#Specifying the Savepoint Offset](#).

1. Used by transactional storage engines, clean up any memory allocated in their slot.
2. A method pointer to the handler's `savepoint_set()` method. This is used to create a savepoint and store it in memory of the requested size.

For more information, see [#Implementing the `savepoint_set` Method](#).

1. A method pointer to the handler's `rollback_to_savepoint()` method. This is used to return to a savepoint during a transaction. It's only populated for storage engines that support savepoints.

For more information, see [#Implementing the `savepoint_rollback\(\)` Method](#).

1. A method pointer to the handler's `release_savepoint()` method. This is used to release the resources of a savepoint during a transaction. It's optionally populated for storage engines that support savepoints.

For more information, see [Implementing the `savepoint_release\(\)` Method](#).

1. A method pointer to the handler's `commit()` method. This is used to commit a transaction. It's only populated for storage engines that support transactions.

For more information, see [#Implementing COMMIT](#).

1. A method pointer to the handler's `rollback()` method. This is used to roll back a transaction. It's only populated for storage engines that support transactions.

For more information, see [#Implementing ROLLBACK](#).

1. Required for XA transactional storage engines. Prepare transaction for commit.
2. Required for XA transactional storage engines. Returns a list of transactions that are in the prepared state.
3. Required for XA transactional storage engines. Commit transaction identified by XID.
4. Required for XA transactional storage engines. Rollback transaction identified by XID.
5. Called when a cursor is created to allow the storage engine to create a consistent read view.
6. Called to switch to a specific consistent read view.
7. Called to close a specific read view.
8. *MANDATORY* - Construct and return a handler instance.

For more information, see [Handling Handler Instantiation](#).

1. Used if the storage engine needs to perform special steps when a schema is dropped (such as in a storage engine that uses tablespaces).
2. Cleanup method called during server shutdown and crashes.
3. *InnoDB* -specific method.
4. *InnoDB* -specific method called at start of `SHOW ENGINE InnoDB STATUS`.
5. Method called to begin a consistent read.
6. Called to indicate that logs should be flushed to reliable storage.
7. Provides human readable status information on the storage engine for `SHOW ENGINE foo STATUS`.
8. *InnoDB* -specific method used for replication.
9. Handlernton flags that indicate the capabilities of the storage engine. Possible values are defined in `sql/handler.h` and copied here:

```
#define HTON_NO_FLAGS 0
```

```
#define HTON_CLOSE_CURSORS_AT_COMMIT (1 << 0)
#define HTON_ALTER_NOT_SUPPORTED     (1 << 1)
#define HTON_CAN_RECREATE             (1 << 2)
#define HTON_FLUSH_AFTER_RENAME      (1 << 3)
#define HTON_NOT_USER_SELECTABLE     (1 << 4)
```

`HTON_ALTER_NOT_SUPPORTED` is used to indicate that the storage engine cannot accept `ALTER TABLE` statements. The `FEDERATED` storage engine is an example. `HTON_FLUSH_AFTER_RENAME` indicates that `FLUSH LOGS` must be called after a table rename. `HTON_NOT_USER_SELECTABLE` indicates that the storage engine should not be shown when a user calls `SHOW STORAGE ENGINES`. Used for system storage engines such as the dummy storage engine for binary logs.

Handling Handler Instantiation

The first method call your storage engine needs to support is the call for a new handler instance.

Before the `handler_ton` is defined in the storage engine source file, a method header for the instantiation method must be defined. Here is an example from the `CSV` engine:

```
static handler* tina_create_handler(TABLE *table);
```

As you can see, the method accepts a pointer to the table the handler is intended to manage, and returns a handler object.

After the method header is defined, the method is named with a method pointer in the `create() handler_ton` element, identifying the method as being responsible for generating new handler instances.

Here is an example of the `MyISAM` storage engine's instantiation method:

```
static handler *myisam_create_handler(TABLE *table)
{
    return new ha_myisam(table);
}
```

This call then works in conjunction with the storage engine's constructor. Here is an example from the `FEDERATED` storage engine:

```
ha_federated::ha_federated(TABLE *table_arg)
:handler(&federated_pton, table_arg),
mysql(0), stored_result(0), scan_flag(0),
ref_length(sizeof(MYSQL_ROW_OFFSET)), current_position(0)
{ }
```

And here's one more example from the `EXAMPLE` storage engine:

```
ha_example::ha_example(TABLE *table_arg)
:handler(&example_pton, table_arg)
{ }
```

The additional elements in the `FEDERATED` example are extra initializations for the handler. The minimum implementation required is the `handler()` initialization shown in the `EXAMPLE` version.

Defining Filename Extensions

Storage engines are required to provide the MySQL server with a list of extensions used by the storage engine with regard to a given table, its data and indexes.

Extensions are expected in the form of a null-terminated string array. The following is the array used by the `CSV` engine:

```
static const char *ha_tina_exts[] = {
    ".CSV",
```

```
NullS
};
```

This array is returned when the `[custom-engine.html#custom-engine-api-reference-bas_ext_bas_ext()]` method is called:

```
const char **ha_tina::bas_ext() const
{
    return ha_tina_exts;
}
```

By providing extension information you can also omit implementing `DROP TABLE` functionality as the MySQL server will implement it for you by closing the table and deleting all files with the extensions you specify.

Creating Tables

Once a handler is instanced, the first operation that will likely be required is the creation of a table.

Your storage engine must implement the `[custom-engine.html#custom-engine-api-reference-create_create()]` virtual method:

```
virtual int create(const char *name, TABLE *form, HA_CREATE_INFO *info)=0;
```

This method should create all necessary files but does not need to open the table. The MySQL server will call for the table to be opened later on.

The `*name` parameter is the name of the table. The `*form` parameter is a `TABLE` structure that defines the table and matches the contents of the `tablename.frm` file already created by the MySQL server. Storage engines must not modify the `tablename.frm` file.

The `*info` parameter is a structure containing information on the `CREATE TABLE` statement used to create the table. The structure is defined in `handler.h` and copied here for your convenience:

```
typedef struct st_ha_create_information
{
    CHARSET_INFO *table_charset, *default_table_charset;
    LEX_STRING connect_string;
    const char *comment,*password;
    const char *data_file_name, *index_file_name;
    const char *alias;
    unsigned long max_rows,min_rows;
    unsigned long auto_increment_value;
    unsigned long table_options;
    unsigned long avg_row_length;
    unsigned long raid_chunksize;
    unsigned long used_fields;
    SQL_LIST merge_list;
    enum db_type db_type;
    enum row_type row_type;
    uint null_bits; /* NULL bits at start of record */
    uint options; /* OR of HA_CREATE_ options */
    uint raid_type,raid_chunks;
    uint merge_insert_method;
    uint extra_size; /* length of extra data segment */
    bool table_existed; /* 1 in create if table existed */
    bool frm_only; /* 1 if no ha_create_table() */
    bool varchar; /* 1 if table has a VARCHAR */
} HA_CREATE_INFO;
```

A basic storage engine can ignore the contents of `*form` and `*info`, as all that is really required is the creation and possibly the initialization of the data files used by the storage engine (assuming the storage engine is file-based).

For example, here is the implementation from the `CSV` storage engine:

```
int ha_tina::create(const char *name, TABLE *table_arg,
    HA_CREATE_INFO *create_info)
{
    char name_buff[FN_REFLLEN];
    File create_file;
    DEBUG_ENTER("ha_tina::create");

    if ((create_file= my_create(fn_format(name_buff, name, "", ".CSV",
        MY_REPLACE_EXT|MY_UNPACK_FILENAME),0,
        O_RDWR | O_TRUNC,MYF(MY_WME))) < 0)
        DEBUG_RETURN(-1);

    my_close(create_file,MYF(0));

    DEBUG_RETURN(0);
}
```

In the preceding example, the [CSV](#) engine does not refer at all to the `*table_arg` or `*create_info` parameters, but simply creates the required data files, closes them, and returns.

The `my_create` and `my_close` methods are helper methods defined in `src/include/my_sys.h`.

Opening a Table

Before any read or write operations are performed on a table, the MySQL server will call the `[custom-engine.html#custom-engine-api-reference-open handler::open()]` method to open the table data and index files (if they exist).

```
int open(const char *name, int mode, int test_if_locked);
```

The first parameter is the name of the table to be opened. The second parameter determines what file to open or what operation to take. The values are defined in `handler.h` and are copied here for your convenience:

```
O_RDONLY - Open read only
O_RDWR - Open read/write
```

The final option dictates whether the handler should check for a lock on the table before opening it. The following options are available:

```
#define HA_OPEN_ABORT_IF_LOCKED 0 /* default */
#define HA_OPEN_WAIT_IF_LOCKED 1
#define HA_OPEN_IGNORE_IF_LOCKED 2
#define HA_OPEN_TMP_TABLE 4 /* Table is a temp table */
#define HA_OPEN_DELAY_KEY_WRITE 8 /* Don't update index */
#define HA_OPEN_ABORT_IF_CRASHED 16
#define HA_OPEN_FOR_REPAIR 32 /* open even if crashed */
```

Typically your storage engine will need to implement some form of shared access control to prevent file corruption in a multi-threaded environment. For an example of how to implement file locking, see the `get_share()` and `free_share()` methods of `sql/examples/ha_tina.cc`.

Implementing Basic Table Scanning

The most basic storage engines implement read-only table scanning. Such engines might be used to support SQL queries of logs and other data files that are populated outside of MySQL.

The implementation of the methods in this section provide the first steps toward the creation of more advanced storage engines.

The following shows the method calls made during a nine-row table scan of the [CSV](#) engine:

```
ha_tina::store_lock
ha_tina::external_lock
```



```

ha_tina::info
ha_tina::rnd_init
ha_tina::extra - ENUM HA_EXTRA_CACHE    Cache record in HA_rrnd()
ha_tina::rnd_next
ha_tina::rnd_next
ha_tina::rnd_next
ha_tina::rnd_next
ha_tina::rnd_next
ha_tina::rnd_next
ha_tina::rnd_next
ha_tina::rnd_next
ha_tina::rnd_next
ha_tina::extra - ENUM HA_EXTRA_NO_CACHE  End caching of records (def)
ha_tina::external_lock
ha_tina::extra - ENUM HA_EXTRA_RESET    Reset database to after open

```

Implementing the `store_lock()` Method

The [custom-engine.html#custom-engine-api-reference-store_lock `store_lock()`] method is called before any reading or writing is performed.

Before adding the lock into the table lock handler **mysqld** calls store lock with the requested locks. Store lock can modify the lock level, for example change blocking write lock to non-blocking, ignore the lock (if we don't want to use MySQL table locks at all) or add locks for many tables (like we do when we are using a MERGE handler).

When releasing locks, `store_lock()` is also called. In this case, one usually doesn't have to do anything.

If the argument of `store_lock` is `TL_IGNORE`, it means that MySQL requests the handler to store the same lock level as the last time.

The potential lock types are defined in `includes/thr_lock.h` and are copied here:

```

enum thr_lock_type
{
    TL_IGNORE=-1,
    TL_UNLOCK,           /* UNLOCK ANY LOCK */
    TL_READ,             /* Read lock */
    TL_READ_WITH_SHARED_LOCKS,
    TL_READ_HIGH_PRIORITY, /* High prior. than TL_WRITE. Allow concurrent insert */
    TL_READ_NO_INSERT,    /* READ, Don't allow concurrent insert */
    TL_WRITE_ALLOW_WRITE, /* Write lock, but allow other threads to read / write. */
    TL_WRITE_ALLOW_READ,  /* Write lock, but allow other threads to read / write. */
    TL_WRITE_CONCURRENT_INSERT, /* WRITE lock used by concurrent insert. */
    TL_WRITE_DELAYED,     /* Write used by INSERT DELAYED. Allows READ locks */
    TL_WRITE_LOW_PRIORITY, /* WRITE lock that has lower priority than TL_READ */
    TL_WRITE,             /* Normal WRITE lock */
    TL_WRITE_ONLY         /* Abort new lock request with an error */
};

```

Actual lock handling will vary depending on your locking implementation and you may choose to implement some or none of the requested lock types, substituting your own methods as appropriate. The following is the minimal implementation, for a storage engine that does not need to downgrade locks:

```

THR_LOCK_DATA **ha_tina::store_lock(THD *thd,
                                     THR_LOCK_DATA **to,
                                     enum thr_lock_type lock_type)
{
    if (lock_type != TL_IGNORE && lock_type == TL_UNLOCK)
        lock_type=lock_type;
    *to++ = &lock;
    return to;
}

```

See also `ha_myisammrg::store_lock()` for a more complex implementation.

Implementing the `external_lock()` Method

The [custom-engine.html#custom-engine-api-reference-external_lock `external_lock()`] method is called at the start of a statement or when a `LOCK TABLES` statement is issued.

Examples of using `external_lock()` can be found in the `sql/ha_innodb.cc` file, but most storage engines can simply return 0, as is the case with the `EXAMPLE` storage engine:

```
int ha_example::external_lock(THD *thd, int lock_type)
{
    DEBUG_ENTER("ha_example::external_lock");
    DEBUG_RETURN(0);
}
```

Implementing the `rnd_init()` Method

The method called before any table scan is the [custom-engine.html#custom-engine-api-reference-rnd_init `rnd_init()`] method. The `rnd_init()` method is used to prepare for a table scan, resetting counters and pointers to the start of the table.

The following example is from the `CSV` storage engine:

```
int ha_tina::rnd_init(bool scan)
{
    DEBUG_ENTER("ha_tina::rnd_init");

    current_position= next_position= 0;
    records= 0;
    chain_ptr= chain;

    DEBUG_RETURN(0);
}
```

If the `scan` parameter is true, the MySQL server will perform a sequential table scan, if false the MySQL server will perform random reads by position.

Implementing the `info(uinf flag)` Method

Prior to commencing a table scan, the [custom-engine.html#custom-engine-api-reference-info `info()`] method is called to provide extra table information to the optimizer.

The information required by the optimizer is not given through return values but instead by populating certain properties of the stats member of the handler class, which the optimizer reads after the `info()` call returns. The stats member is an instance of the `ha_statistics` class which is also defined in `handler.h`.

In addition to being used by the optimizer, many of the values set during a call to the `info()` method are also used for the `SHOW TABLE STATUS` statement. The flag argument is a bitfield that conveys for which context the info method was called.

The flags are defined in `include/my_base.h`. The ones that are used are:

- `HA_STATUS_NO_LOCK` - the handler may use outdated info if it can prevent locking the table shared
- `HA_STATUS_TIME` - only update of stats->update_time required
- `HA_STATUS_CONST` - update the immutable members of stats (max_data_file_length, max_index_file_length, create_time, sortkey, ref_length, block_size, data_file_name, index_file_name)
- `HA_STATUS_VARIABLE` - records, deleted, data_file_length, index_file_length, delete_length, check_time, mean_rec_length

- `HA_STATUS_ERRKEY` - status pertaining to last error key (errkey and dupp_ref)
- `HA_STATUS_AUTO` - update autoincrement value

The public properties are listed in full in `sql/handler.h`; several of the more common ones are copied here:

```

ulonglong data_file_length;      /* Length off data file */
ulonglong max_data_file_length; /* Length off data file */
ulonglong index_file_length;
ulonglong max_index_file_length;
ulonglong delete_length;        /* Free bytes */
ulonglong auto_increment_value;
ha_rows records;                /* Records in table */
ha_rows deleted;                /* Deleted records */
ulong raid_chunksize;
ulong mean_rec_length;          /* physical reclength */
time_t create_time;             /* When table was created */
time_t check_time;
time_t update_time;

```

For the purposes of a table scan, the most important property is `records`, which indicates the number of records in the table. The optimizer will perform differently when the storage engine indicates that there are zero or one rows in the table than it will when there are two or more. For this reason it is important to return a value of two or greater when you do not actually know how many rows are in the table before you perform the table scan (such as in a situation where the data may be externally populated).

The bare minimum implementation for the info method is probably something like what is used for the CSV (tina) engine:

```

int ha_tina::info(uint flag)
{
    DEBUG_ENTER("ha_tina::info");
    /* This is a lie, but you don't want the optimizer to see zero or 1 */
    if (!records_is_known && stats.records < 2)
        stats.records = 2;
    DEBUG_RETURN(0);
}

```

Implementing the `extra()` Method

Prior to some operations, the [custom-engine.html#custom-engine-api-reference-extra `extra()`] method is called to provide extra hints to the storage engine on how to perform certain operations.

Implementation of the hints in the `extra` call is not mandatory, and most storage engines return 0:

```

int ha_tina::extra(enum ha_extra_function operation)
{
    DEBUG_ENTER("ha_tina::extra");
    DEBUG_RETURN(0);
}

```

Implementing the `rnd_next()` Method

After the table is initialized, the MySQL server will call the handler's [custom-engine.html#custom-engine-api-reference-rnd_next `rnd_next()`] method once for every row to be scanned until the server's search condition is satisfied or an end of file is reached, in which case the handler returns `HA_ERR_END_OF_FILE`.

The `rnd_next()` method takes a single byte array parameter named `*buf`. The `*buf` parameter must be populated with the contents of the table row in the internal MySQL format.

The server uses three data formats: fixed-length rows, variable-length rows, and variable-length rows with BLOB pointers. In each format, the columns appear in the order in which they were defined by the

CREATE TABLE statement. (The table definition is stored in the `.frm` file, and the optimizer and the handler are both able to access table metadata from the same source, its `TABLE` structure).

Each format begins with a NULL bitmap of one bit per nullable column. A table with as many as eight nullable columns will have a one-byte bitmap; a table with nine to sixteen nullable columns will have a two-byte bitmap, and so forth. One exception is fixed-width tables, which have an additional starting bit so that a table with eight nullable columns would have a two-byte bitmap.

After the NULL bitmap come the columns, one by one. Each column is of the size indicated in [MySQL Data Types](#). In the server, column data types are defined in the `sql/field.cc` file. In the fixed length row format, the columns are simply laid out one by one. In a variable-length row, `VARCHAR` columns are coded as a one or two-byte length, followed by a string of characters. In a variable-length row with `BLOB` columns, each blob is represented by two parts: first an integer representing the actual size of the `BLOB`, and then a pointer to the `BLOB` in memory.

Examples of row conversion (or packing) can be found by starting at `rnd_next()` in any table handler. In `ha_tina.cc`, for example, the code in `find_current_row()` illustrates how the `TABLE` structure (pointed to by `table`) and a string object (named `buffer`) can be used to pack character data from a CSV file. Writing a row back to disk requires the opposite conversion, unpacking from the internal format.

The following example is from the `CSV` storage engine:

```
int ha_tina::rnd_next(byte *buf)
{
    DEBUG_ENTER("ha_tina::rnd_next");

    statistic_increment(table->in_use->status_var.ha_read_rnd_next_count, &LOCK_status);

    current_position= next_position;
    if (!share->mapped_file)
        DEBUG_RETURN(HA_ERR_END_OF_FILE);
    if (HA_ERR_END_OF_FILE == find_current_row(buf) )
        DEBUG_RETURN(HA_ERR_END_OF_FILE);

    records++;
    DEBUG_RETURN(0);
}
```

The conversion from the internal row format to CSV row format is performed in the `find_current_row()` method:

```
int ha_tina::find_current_row(byte *buf)
{
    byte *mapped_ptr= (byte *)share->mapped_file + current_position;
    byte *end_ptr;
    DEBUG_ENTER("ha_tina::find_current_row");

    /* EOF should be counted as new line */
    if ((end_ptr= find_eoln(share->mapped_file, current_position,
                          share->file_stat.st_size)) == 0)
        DEBUG_RETURN(HA_ERR_END_OF_FILE);

    for (Field **field=table->field ; *field ; field++)
    {
        buffer.length(0);
        mapped_ptr++; // Increment past the first quote
        for(;mapped_ptr != end_ptr; mapped_ptr++)
        {
            // Need to convert line feeds!
            if (*mapped_ptr == '"' &&
                (((mapped_ptr[1] == ',') && (mapped_ptr[2] == '"')) ||
                 (mapped_ptr == end_ptr - 1)))
            {
                mapped_ptr += Move past the , and the "
                break;
            }
            if (*mapped_ptr == '\\') && mapped_ptr != (end_ptr - 1))
            {
```

```

        mapped_ptr++;
        if (*mapped_ptr == 'r')
            buffer.append('\r');
        else if (*mapped_ptr == 'n' )
            buffer.append('\n');
        else if ((*mapped_ptr == '\\') || (*mapped_ptr == ''))
            buffer.append(*mapped_ptr);
        else /* This could only happen with an externally created file */
        {
            buffer.append('\\');
            buffer.append(*mapped_ptr);
        }
    }
    else
        buffer.append(*mapped_ptr);
}
(*field)->store(buffer.ptr(), buffer.length(), system_charset_info);
}
next_position= (end_ptr - share->mapped_file)+1;
/* Maybe use \N for null? */
memset(buf, 0, table->s->>null_bytes); /* We do not implement nulls! */

DEBUG_RETURN(0);
}

```

Closing a Table

When the MySQL server is finished with a table, it will call the [\[custom-engine.html#custom-engine-api-reference-close close\(\)\]](#) method to close file pointers and release any other resources.

Storage engines that use the shared access methods seen in the [CSV](#) engine and other example engines must remove themselves from the shared structure:

```

int ha_tina::close(void)
{
    DEBUG_ENTER("ha_tina::close");
    DEBUG_RETURN(free_share(share));
}

```

Storage engines using their own share management systems should use whatever methods are needed to remove the handler instance from the share for the table opened in their handler.

Adding Support for [INSERT](#) to a Storage Engine

Once you have read support in your storage engine, the next feature to implement is support for [INSERT](#) statements. With [INSERT](#) support in place, your storage engine can handle WORM (write once, read many) applications such as logging and archiving for later analysis.

All [INSERT](#) operations are handled through the [\[custom-engine.html#custom-engine-api-reference-write_row write_row\(\)\]](#) method:

```

int ha_foo::write_row(byte *buf)

```

The `*buf` parameter contains the row to be inserted in the internal MySQL format. A basic storage engine could simply advance to the end of the data file and append the contents of the buffer directly (this would also make reading rows easier as you could read the row and pass it directly into the buffer parameter of the `rnd_next()` method).

The process for writing a row is the opposite of the process for reading one: take the data from the MySQL internal row format and write it to the data file. The following example is from the [MyISAM](#) storage engine:

```

int ha_myisam::write_row(byte * buf)
{
    statistic_increment(table->in_use->status_var.ha_write_count, &LOCK_status);

```

```

/* If we have a timestamp column, update it to the current time */
if (table->timestamp_field_type & TIMESTAMP_AUTO_SET_ON_INSERT)
    table->timestamp_field->set_time();

/*
    If we have an auto_increment column and we are writing a changed row
    or a new row, then update the auto_increment value in the record.
*/
if (table->next_number_field && buf == table->record[0])
    update_auto_increment();
return mi_write(file,buf);
}

```

Three items of note in the preceding example include the updating of table statistics for writes, the setting of the timestamp prior to writing the row, and the updating of **AUTO_INCREMENT** values.

Adding Support for **UPDATE** to a Storage Engine

The MySQL server executes **UPDATE** statements by performing a (table/index/range/etc.) scan until it locates a row matching the **WHERE** clause of the **UPDATE** statement, then calling the [custom-engine.html#custom-engine-api-reference-update_row `update_row()` method:]

```
int ha_foo::update_row(const byte *old_data, byte *new_data)
```

The `*old_data` parameter contains the data that existed in the row prior to the update, while the `*new_data` parameter contains the new contents of the row (in the MySQL internal row format).

Performing an update will depend on row format and storage implementation. Some storage engines will replace data in-place, while other implementations delete the existing row and append the new row at the end of the data file.

Non-indexed storage engines can typically ignore the contents of the `*old_data` parameter and just deal with the `*new_data` buffer. Transactional engines may need to compare the buffers to determine what changes have been made for a later rollback.

If the table being updated contains timestamp columns, the updating of the timestamp will have to be managed in the `update_row()` call. The following example is from the **CSV** engine:

```

int ha_tina::update_row(const byte * old_data, byte * new_data)
{
    int size;
    DEBUG_ENTER("ha_tina::update_row");

    statistic_increment(table->in_use->status_var.ha_read_rnd_next_count,
                        &LOCK_status);

    if (table->timestamp_field_type & TIMESTAMP_AUTO_SET_ON_UPDATE)
        table->timestamp_field->set_time();

    size= encode_quote(new_data);

    if (chain_append())
        DEBUG_RETURN(-1);

    if (my_write(share->data_file, buffer.ptr(), size, MYF(MY_WME | MY_NABP)))
        DEBUG_RETURN(-1);
    DEBUG_RETURN(0);
}

```

Note the setting of the timestamp in the previous example.

Adding Support for **DELETE** to a Storage Engine

The MySQL server executes **DELETE** statements using the same approach as for **UPDATE** statements: It advances to the row to be deleted using the `rnd_next()` method and then calls the [custom-engine.html#custom-engine-api-reference-delete_row `delete_row()`] method to delete the row:

```
int ha_foo::delete_row(const byte *buf)
```

The `*buf` parameter contains the contents of the row to be deleted. For non-indexed storage engines the parameter can be ignored, but transactional storage engines may need to store the deleted data for rollback purposes.

The following example is from the `CSV` storage engine:

```
int ha_tina::delete_row(const byte * buf)
{
    DEBUG_ENTER("ha_tina::delete_row");
    statistic_increment(table->in_use->status_var.ha_delete_count,
                        &LOCK_status);

    if (chain_append())
        DEBUG_RETURN(-1);

    --records;

    DEBUG_RETURN(0);
}
```

The steps of note in the preceding example are the update of the `delete_count` statistic and the record count.

Supporting Non-Sequential Reads

In addition to table scanning, storage engines can implement methods for non-sequential reading. (Note: this is not "can" but rather a "must" because certain operations rely on proper implementation of `position()` and `rnd_pos()` calls. Two examples of such operations are multi-table UPDATE and `SELECT .. table.blob_column ORDER BY something`). The MySQL server uses these methods for certain sort operations.

Implementing the `position()` Method

The [custom-engine.html#custom-engine-api-reference-position `position()`] method is called after every call to `rnd_next()` if the data needs to be reordered:

```
void ha_foo::position(const byte *record)
```

It stores a "position" of a record in `this->ref`. The contents of this "position" is up to you, whatever value you provide will be returned in a later call to retrieve the row. The only rule - "position" or a row must contain enough information to allow you later to retrieve this very row. Most storage engines will store some form of offset or a primary key value.

Implementing the `rnd_pos()` Method

The [custom-engine.html#custom-engine-api-reference-rnd_pos `rnd_pos()`] method behaves in a similar fashion to the `rnd_next()` method but takes an additional parameter:

```
int ha_foo::rnd_pos(byte * buf, byte *pos)
```

The `*pos` parameter contains positioning information previously recorded using the `position()` method.

A storage engine must locate the row specified by the position and return it through `*buf` in the internal MySQL row format.

Supporting Indexing

Once basic read/write operations are implemented in a storage engine, the next stage is to add support for indexing. Without indexing, a storage engine's performance is quite limited.

This section documents the methods that must be implemented to add support for indexing to a storage engine.

Indexing Overview

Adding index support to a storage engine revolves around two tasks: providing information to the optimizer and implementing index-related methods. The information provided to the optimizer helps the optimizer to make better decisions about which index to use or even to skip using an index and instead perform a table scan.

The indexing methods either read rows that match a key, scan a set of rows by index order, or read information directly from the index.

The following example shows the method calls made during an `UPDATE` query that uses an index, such as `UPDATE foo SET ts = now() WHERE id = 1`:

```
ha_foo::index_init
ha_foo::index_read
ha_foo::index_read_idx
ha_foo::rnd_next
ha_foo::update_row
```

In addition to index reading methods, your storage engine must support the creation of new indexes and be able to keep table indexes up to date as rows are added, modified, and removed from tables.

Getting Index Information During `CREATE TABLE` Operations

It is preferable for storage engines that support indexing to read the index information provided during a `CREATE TABLE` operation and store it for future use. The reasoning behind this is that the index information is most readily available during table and index creation and is not as easily retrieved afterward.

The table index information is contained within the `key_info` structure of the `TABLE` argument of the `create()` method.

Within the `key_info` structure there is a `flag` that defines index behavior:

```
#define HA_NOSAME          1 /* Set if not duplicated records */
#define HA_PACK_KEY        2 /* Pack string key to previous key */
#define HA_AUTO_KEY        16
#define HA_BINARY_PACK_KEY 32 /* Packing of all keys to prev key */
#define HA_FULLTEXT        128 /* For full-text search */
#define HA_UNIQUE_CHECK    256 /* Check the key for uniqueness */
#define HA_SPATIAL         1024 /* For spatial search */
#define HA_NULL_ARE_EQUAL  2048 /* NULL in key are cmp as equal */
#define HA_GENERATED_KEY   8192 /* Automatically generated key */
```

In addition to the `flag`, there is an enumerator named `algorithm` that specifies the desired index type:

```
enum ha_key_alg {
    HA_KEY_ALG_UNDEF= 0, /* Not specified (old file) */
    HA_KEY_ALG_BTREE= 1, /* B-tree, default one */
    HA_KEY_ALG_RTREE= 2, /* R-tree, for spatial searches */
    HA_KEY_ALG_HASH= 3, /* HASH keys (HEAP tables) */
    HA_KEY_ALG_FULLTEXT= 4 /* FULLTEXT (MyISAM tables) */
};
```

In addition to the `flag` and `algorithm`, there is an array of `key_part` elements that describe the individual parts of a composite key.

The key parts define the field associated with the key part, whether the key should be packed, and the data type and length of the index part. See `ha_myisam.cc` for an example of how this information is parsed.

As an alternative, a storage engine can instead read index information from the `TABLE` structure of the handler during each operation.

Creating Index Keys

As part of every table-write operation (`INSERT`, `UPDATE`, `DELETE`), the storage engine is required to update its internal index information.

The method used to update indexes will vary from storage engine to storage engine, depending on the method used to store the index.

In general, the storage engine will have to use row information passed in methods such as [custom-engine.html#custom-engine-api-reference-write_row `write_row()`], [custom-engine.html#custom-engine-api-reference-delete_row `delete_row()`], and [custom-engine.html#custom-engine-api-reference-update_row `update_row()`] in combination with index information for the table to determine what index data needs to be modified, and make the needed changes.

The method of associating an index with its row will depend on your storage approach. Current storage engines store the row offset.

Parsing Key Information

Many of the index methods pass a byte array named `*key` that identifies the index entry to be read in a standard format. Your storage engine will need to extract the information stored in the key and translate it into its internal index format to identify the row associated with the index.

The information in the key is obtained by iterating through the key, which is formatted the same as the definition in `table->key_info[index]->key_part[part_num]`.

Along with the key, handler methods pass a `keypart_map` parameter to indicate which parts of the key are present in the `key` parameter. `keypart_map` is a `ulonglong` bitmap with one bit per key part: 1 for `keypart[0]`, 2 for `keypart[1]`, 4 for `keypart[2]`, and so forth. If a bit in `keypart_map` is set, the value for this key part is present in the key buffer. Bits following the bit for the last key part don't matter, so `~0` can be used for all keyparts. Currently, only key prefixes are supported. That is, `assert((keypart_map + 1) & keypart_map == 0)`.

A `keypart_map` is part of the `key_range` structure used by `records_in_range()`, and a `keypart_map` value is passed directly to the `index_read()`, `index_read_idx()`, and `index_read_last()` methods.

Older handlers have a `key_len` parameter instead of `keypart_map`. The `key_len` value is a `uint` that indicates the prefix length when matching by prefix.

Providing Index Information to the Optimizer

In order for indexing to be used effectively, storage engines need to provide the optimizer with information about the table and its indexes. This information is used to choose whether to use an index, and if so, which index to use.

Implementing the `info()` Method

The optimizer requests an update of table information by calling the [custom-engine.html#custom-engine-api-reference-info `handler::info()`] method. The `info()` method does not have a return value, instead it is expected that the storage engine will set a variety of public variables that the server will then read as needed. These values are also used to populate certain `SHOW` outputs such as `SHOW TABLE STATUS` and for queries of the `INFORMATION_SCHEMA`.

All variables are optional but should be filled if possible:

- `records` - The number of rows in the table. If you cannot provide an accurate number quickly you should set the value to be greater than 1 so that the optimizer does not perform optimizations for zero or one row tables.
- `deleted` - Number of deleted rows in table. Used to identify table fragmentation, where applicable.
- `data_file_length` - Size of the data file, in bytes. Helps optimizer calculate the cost of reads.
- `index_file_length` - Size of the index file, in bytes. Helps optimizer calculate the cost of reads.
- `mean_rec_length` - Average length of a single row, in bytes.
- `scan_time` - Cost in I/O seeks to perform a full table scan.
- `delete_length` -
- `check_time` -

When calculating values, speed is more important than accuracy, as there is no sense in taking a long time to give the optimizer clues as to what approach may be the fastest. Estimates within an order of magnitude are usually good enough.

Implementing the `records_in_range` Method

The `[custom-engine.html#custom-engine-api-reference-records_in_range records_in_range()]` method is called by the optimizer to assist in choosing which index on a table to use for a query or join. It is defined as follows:

```
ha_rows ha_foo::records_in_range(uint inx, key_range *min_key, key_range *max_key)
```

The `inx` parameter is the index to be checked. The `*min_key` and `*max_key` parameters are `key_range` structures that indicate the low and high ends of the range. The `key_range` structure looks like this:

```
typedef struct st_key_range
{
    const byte *key;
    uint length;
    key_part_map keypart_map;
    enum ha_rkey_function flag;
} key_range;
```

`key_range` members are used as follows:

- `key` is a pointer to the key buffer.
- `length` is the key length.
- `keypart_map` is a bitmap that indicates which key parts are passed in `key` (as described in [Parsing Key Information](#)).
- `flag` indicates whether to include the key in the range. Its value differs for `min_key` and `max_key`, as described following.

`min_key.flag` can have one of the following values:

- `HA_READ_KEY_EXACT` - Include the key in the range
- `HA_READ_AFTER_KEY` - Don't include key in range

`max_key.flag` can have one of the following values:

- `HA_READ_BEFORE_KEY` - Don't include key in range
- `HA_READ_AFTER_KEY` - Include all 'end_key' values in the range

The following return values are allowed:

- 0 - There are no matching keys in the given range
- *number* > 0 - There are approximately *number* matching rows in the range
- `HA_POS_ERROR` - Something is wrong with the index tree

When calculating values, speed is more important than accuracy.

Preparing for Index Use with `index_init()`

The [custom-engine.html#custom-engine-api-reference-index_init `index_init()`] method is called before an index is used to allow the storage engine to perform any necessary preparation or optimization:

```
int ha_foo::index_init(uint keynr, bool sorted)
```

Most storage engines do not need to make special preparations, in which case a default implementation will be used if the method is not explicitly implemented in the storage engine:

```
int handler::index_init(uint idx) { active_index=idx; return 0; }
```

Cleaning up with `index_end()`

The [custom-engine.html#custom-engine-api-reference-index_end `index_end()`] method is a counterpart to the `index_init()` method. The purpose of the `index_end()` method is to clean up any preparations made by the `index_init()` method.

If a storage engine does not implement `index_init()` it does not need to implement `index_end()`.

Implementing the `index_read()` Method

The [custom-engine.html#custom-engine-api-reference-index_read `index_read()`] method is used to retrieve a row based on a key:

```
int ha_foo::index_read(byte * buf, const byte * key,
                      ulonglong keypart_map,
                      enum ha_rkey_function find_flag)
```

The **buf* parameter is a byte array that the storage engine populates with the row that matches the index key specified in **key*. The *keypart_map* parameter is a bitmap that indicates which parts of the key are present in the *key* parameter. The *find_flag* parameter is an enumerator that dictates the search behavior to be used, as discussed in [Parsing Key Information](#).

The index to be used is previously defined in the [custom-engine.html#custom-engine-index-init `index_init()`] call and is stored in the `active_index` handler variable.

The following values are allowed for *find_flag*:

```
HA_READ_AFTER_KEY
HA_READ_BEFORE_KEY
HA_READ_KEY_EXACT
HA_READ_KEY_OR_NEXT
HA_READ_KEY_OR_PREV
HA_READ_PREFIX
HA_READ_PREFIX_LAST
HA_READ_PREFIX_LAST_OR_PREV
```

Storage engines must convert the **key* parameter to a storage engine-specific format, use it to find the matching row according to the *find_flag*, and then populate **buf* with the matching row in the MySQL internal row format. For more information on the internal row format, see [#Implementing the `rnd_next\(\)` Method](#).

In addition to returning a matching row, the storage engine must also set a cursor to support sequential index reads.

If the `*key` parameter is null, the storage engine should read the first key in the index.

Implementing the `index_read_idx()` Method

The [custom-engine.html#custom-engine-api-reference-index_read_idx `index_read_idx()`] method is identical to [custom-engine.html#custom-engine-index-read `index_read()`] with the exception that `index_read_idx()` accepts an additional `keynr` parameter:

```
int ha_foo::index_read_idx(byte * buf, uint keynr, const byte * key,
                          ulonglong keypart_map,
                          enum ha_rkey_function find_flag)
```

The `keynr` parameter specifies the index to be read, as opposed to `index_read()`, where the index is already set.

As with the `index_read()` method, the storage engine must return the row that matches the key according to the `find_flag` and set a cursor for future reads.

Implementing the `index_read_last()` Method

The [custom-engine.html#custom-engine-api-reference-index_read_last `index_read_last()`] method works like [custom-engine.html#custom-engine-index-read `index_read()`] but finds the last row with the current key value or prefix:

```
int ha_foo::index_read_last(byte * buf, const byte * key,
                           key_part_map keypart_map)
```

`index_read_last()` is used when optimizing away the `ORDER BY` clause for statements such as this:

```
SELECT * FROM t1 WHERE a=1 ORDER BY a DESC,b DESC;
```

Implementing the `index_next()` Method

The [custom-engine.html#custom-engine-api-reference-index_next `index_next()`] method is used for index scanning:

```
int ha_foo::index_next(byte * buf)
```

The `*buf` parameter is populated with the row that corresponds to the next matching key value according to the internal cursor set by the storage engine during operations such as `index_read()` and `index_first()`.

Implementing the `index_prev()` Method

The [custom-engine.html#custom-engine-api-reference-index_prev `index_prev()`] method is used for reverse index scanning:

```
int ha_foo::index_prev(byte * buf)
```

The `*buf` parameter is populated with the row that corresponds to the previous matching key value according to the internal cursor set by the storage engine during operations such as `index_read()` and `index_last()`.

Implementing the `index_first()` Method

The [custom-engine.html#custom-engine-api-reference-index_first `index_first()`] method is used for index scanning:

```
int ha_foo::index_first(byte * buf)
```

The `*buf` parameter is populated with the row that corresponds to the first key value in the index.

Implementing the `index_last()` Method

The `[custom-engine.html#custom-engine-api-reference-index_last index_last()]` method is used for reverse index scanning:

```
int ha_foo::index_last(byte * buf)
```

The `*buf` parameter is populated with the row that corresponds to the last key value in the index.

Supporting Transactions

This section documents the methods that must be implemented to add support for transactions to a storage engine.

Please note that transaction management can be complicated and involve methods such as row versioning and redo logs, which is beyond the scope of this document. Instead coverage is limited to a description of required methods and not their implementation. For examples of implementation, please see ha_innodb.cc.

Transaction Overview

Transactions are not explicitly started on the storage engine level, but are instead implicitly started through calls to either `start_stmt()` or `external_lock()`. If the preceding methods are called and a transaction already exists the transaction is not replaced.

The storage engine stores transaction information in per-connection memory and also registers the transaction in the MySQL server to allow the server to later issue `COMMIT` and `ROLLBACK` operations.

As operations are performed the storage engine will have to implement some form of versioning or logging to permit a rollback of all operations executed within the transaction.

After work is completed, the MySQL server will call either the `commit()` method or the `rollback()` method defined in the storage engine's handler.

Starting a Transaction

A transaction is started by the storage engine in response to a call to either the `start_stmt()` or `external_lock()` methods.

If there is no active transaction, the storage engine must start a new transaction and register the transaction with the MySQL server so that `ROLLBACK` or `COMMIT` can later be called.

Starting a Transaction from a `start_stmt()` Call

The first method call that can start a transaction is the `[custom-engine.html#custom-engine-transactions-starting-start-stmt start_stmt()]` method.

The following example shows how a storage engine could register a transaction:

```
int my_handler::start_stmt(THD *thd, thr_lock_type lock_type)
{
    int error= 0;
    my_txn *txn= (my_txn *) thd->ha_data[my_handler_hton.slot];

    if (txn == NULL)
    {
        thd->ha_data[my_handler_hton.slot]= txn= new my_txn;
    }
    if (txn->stmt == NULL && !(error= txn->tx_begin()))
    {
```

```

    txn->stmt= txn->new_savepoint();
    trans_register_ha(thd, FALSE, &my_handler_hton);
}
return error;
}

```

THD is the current client connection. It holds state relevant data for the current client, such as identity, network connection and other per-connection data.

`thd->ha_data[my_handler_hton.slot]` is a pointer in `thd` to the connection-specific data of this storage engine. In this example we use it to store the transaction context.

An additional example of implementing `start_stmt()` can be found in `ha_innodb.cc`.

Starting a Transaction from a `external_lock()` Method

MySQL calls `[custom-engine.html#custom-engine-api-reference-external_lock handler::external_lock()]` for every table it is going to use at the beginning of every statement. Thus, if a table is touched for the first time, it implicitly starts a transaction.

Note that because of pre-locking, all tables that can be potentially used between the beginning and the end of a statement are locked before the statement execution begins and `handler::external_lock()` is called for all these tables. That is, if an `INSERT` fires a trigger, which calls a stored procedure, that invokes a stored method, and so forth, all tables used in the trigger, stored procedure, method, etc., are locked in the beginning of the `INSERT`. Additionally, if there's a construct like

```

IF
.. use one table
ELSE
.. use another table

```

both tables will be locked.

Also, if a user calls `LOCK TABLES`, MySQL will call `handler::external_lock` only once. In this case, MySQL will call `handler::start_stmt()` at the beginning of the statement.

The following example shows how a storage engine can start a transaction and take locking requests into account:

```

int my_handler::external_lock(THD *thd, int lock_type)
{
    int error= 0;
    my_txn *txn= (my_txn *) thd->ha_data[my_handler_hton.slot];

    if (txn == NULL)
    {
        thd->ha_data[my_handler_hton.slot]= txn= new my_txn;
    }

    if (lock_type != F_UNLCK)
    {
        bool all_tx= 0;
        if (txn->lock_count == 0)
        {
            txn->lock_count= 1;
            txn->tx_isolation= thd->variables.tx_isolation;

            all_tx= test(thd->options & (OPTION_NOT_AUTOCOMMIT | OPTION_BEGIN | OPTION_TABLE_LOCK));
        }

        if (all_tx)
        {
            txn->tx_begin();
            trans_register_ha(thd, TRUE, &my_handler_hton);
        }
        else
    }
}

```

```

    if (txn->stmt == 0)
    {
        txn->stmt= txn->new_savepoint();
        trans_register_ha(thd, FALSE, &my_handler_hton);
    }
    else
    {
        if (txn->stmt != NULL)
        {
            /* Commit the transaction if we're in auto-commit mode */
            my_handler_commit(thd, FALSE);

            delete txn->stmt; // delete savepoint
            txn->stmt= NULL;
        }
    }

    return error;
}

```

Every storage engine must call `trans_register_ha()` every time it starts a transaction. The `trans_register_ha()` method registers a transaction with the MySQL server to allow for future `COMMIT` and `ROLLBACK` calls.

An additional example of implementing `external_lock()` can be found in `ha_innodb.cc`.

Implementing ROLLBACK

Of the two major transactional operations, `ROLLBACK` is the more complicated to implement. All operations that occurred during the transaction must be reversed so that all rows are unchanged from before the transaction began.

To support `ROLLBACK`, create a method that matches this definition:

```
int (*rollback)(THD *thd, bool all);
```

The method name is then listed in the `rollback` (thirteenth) entry of `[custom-engine.html#custom-engine-handlerton the handlerton]`.

The `THD` parameter is used to identify the transaction that needs to be rolled back, while the `bool all` parameter indicates whether the entire transaction should be rolled back or just the last statement.

Details of implementing a `ROLLBACK` operation will vary by storage engine. Examples can be found in `ha_innodb.cc`.

Implementing COMMIT

During a commit operation, all changes made during a transaction are made permanent and a rollback operation is not possible after that. Depending on the transaction isolation used, this may be the first time such changes are visible to other threads.

To support `COMMIT`, create a method that matches this definition:

```
int (*commit)(THD *thd, bool all);
```

The method name is then listed in the `commit` (twelfth) entry of `[custom-engine.html#custom-engine-handlerton the handlerton]`.

The `THD` parameter is used to identify the transaction that needs to be committed, while the `bool all` parameter indicates if this is a full transaction commit or just the end of a statement that is part of the transaction.

Details of implementing a `COMMIT` operation will vary by storage engine. Examples can be found in `ha_innodb.cc`.

If the server is in auto-commit mode, the storage engine should automatically commit all read-only statements such as `SELECT`.

In a storage engine, "auto-committing" works by counting locks. Increment the count for every call to `external_lock()`, decrement when `external_lock()` is called with an argument of `F_UNLCK`. When the count drops to zero, trigger a commit.

Adding Support for Savepoints

First, the implementor should know how many bytes are required to store savepoint information. This should be a fixed size, preferably not large as the MySQL server will allocate space to store the savepoint for all storage engines with each named savepoint.

The implementor should store the data in the space preallocated by `mysqld` - and use the contents from the preallocated space for rollback or release savepoint operations.

When a `COMMIT` or `ROLLBACK` operation occurs (with `bool all` set to `true`), all savepoints are assumed to be released. If the storage engine allocates resources for savepoints, it should free them.

The following handler elements need to be implemented to support savepoints (elements 7,9,10,11):

```
uint savepoint_offset;
int  (*savepoint_set)(THD *thd, void *sv);
int  (*savepoint_rollback)(THD *thd, void *sv);
int  (*savepoint_release)(THD *thd, void *sv);
```

Specifying the Savepoint Offset

The seventh element of the handler is the `savepoint_offset`:

```
uint savepoint_offset;
```

The `savepoint_offset` must be initialized statically to the size of the needed memory to store per-savepoint information.

Implementing the `savepoint_set` Method

The `savepoint_set()` method is called whenever a user issues the `SAVEPOINT` statement:

```
int (*savepoint_set)(THD *thd, void *sv);
```

The `*sv` parameter points to an uninitialized storage area of the size defined by `savepoint_offset`.

When `savepoint_set()` is called, the storage engine needs to store savepoint information into `sv` so that the server can later roll back the transaction to the savepoint or release the savepoint resources.

Implementing the `savepoint_rollback()` Method

The `savepoint_rollback()` method is called whenever a user issues the `ROLLBACK TO SAVEPOINT` statement:

```
int (*savepoint_rollback)(THD *thd, void *sv);
```

The `*sv` parameter points to the storage area that was previously passed to the `savepoint_set()` method.

Implementing the `savepoint_release()` Method

The `savepoint_release()` method is called whenever a user issues the `RELEASE SAVEPOINT` statement:

```
int (*savepoint_release)(THD *thd, void *sv);
```


The `*sv` parameter points to the storage area that was previously passed to the `savepoint_set()` method.

API Reference

bas_ext

Purpose

Defines the file extensions used by the storage engine.

Synopsis

```
virtual const char **      |);  
bas_ext (                  |  
  
                           |;  
                           |;
```

Description

This is the `bas_ext` method. It is called to provide the MySQL server with a list of file extensions used by the storage engine. The list returned is a null-terminated string array.

By providing a list of extensions, storage engines can in many cases omit the [custom-engine.html#custom-engine-api-reference-delete_table `delete_table()`] method as the MySQL server will close all references to the table and delete all files with the specified extension.

Parameters

There are no parameters for this method.

Return Values

- Return value is a null-terminated string array of storage engine extensions. The following is an example from the `CSV` engine:

```
static const char *ha_tina_exts[] =  
{  
    ".CSV",  
    NullS  
};
```

Usage

```
static const char *ha_tina_exts[] =  
{  
    ".CSV",  
    NullS  
};  
  
const char **ha_tina::bas_ext() const  
{  
    return ha_tina_exts;  
}
```

Default Implementation

```
static const char *ha_example_exts[] = {  
    NullS  
};  
  
const char **ha_example::bas_ext() const
```

```
{
    return ha_example_exts;
}
```

close

Purpose

Closes an open table.

Synopsis

```
virtual int close (      |void);
                        |
                        |void;
```

Description

This is the `close` method.

Closes a table. A good time to free any resources that we have allocated.

Called from `sql_base.cc`, `sql_select.cc`, and `table.cc`. In `sql_select.cc` it is only used to close up temporary tables or during the process where a temporary table is converted over to being a [MyISAM](#) table. For `sql_base.cc` look at `close_data_tables()`.

Parameters

- `void`

Return Values

There are no return values.

Usage

Example from the [CSV](#) engine:

```
int ha_example::close(void)
{
    DEBUG_ENTER("ha_example::close");
    DEBUG_RETURN(free_share(share));
}
```

create

Purpose

Creates a new table.

Synopsis

<code>virtual int create (</code>	<code>name,</code>	
	<code>form,</code>	
	<code>info);</code>	
<code>const char *</code>	<code>name ;</code>	
<code>TABLE *</code>	<code>form ;</code>	
<code>HA_CREATE_INFO *</code>	<code>info ;</code>	

Description

This is the `create` method.

`create()` is called to create a table. The variable name will have the name of the table. When `create()` is called you do not need to open the table. Also, the `.frm` file will have already been created so `adjusting create_info` is not recommended.

Called from `handler.cc` by `ha_create_table()`.

Parameters

- `name`
- `form`
- `info`

Return Values

There are no return values.

Usage

Example from the `CSV` storage engine:

```
int ha_tina::create(const char *name, TABLE *table_arg,
                   HA_CREATE_INFO *create_info)
{
    char name_buff[FN_REFLen];
    File create_file;
    DEBUG_ENTER("ha_tina::create");

    if ((create_file= my_create(fn_format(name_buff, name, "", ".CSV",
                                         MY_REPLACE_EXT|MY_UNPACK_FILENAME),0,
                                O_RDWR | O_TRUNC,MYF(MY_WME))) < 0)
        DEBUG_RETURN(-1);

    my_close(create_file,MYF(0));

    DEBUG_RETURN(0);
}
```

delete_row

Purpose

Deletes a row.

Synopsis

```
virtual int delete_row ( |buf);  
  
const byte * |buf ;
```

Description

This is the `delete_row` method.

`buf` will contain a copy of the row to be deleted. The server will call this right after the current row has been called (from either a previous `rnd_next()` or index call). If you keep a pointer to the last row or can access a primary key it will make doing the deletion quite a bit easier. Keep in mind that the server does not guarantee consecutive deletions. `ORDER BY` clauses can be used.

Called in `sql_acl.cc` and `sql_udf.cc` to manage internal table information. Called in `sql_delete.cc`, `sql_insert.cc`, and `sql_select.cc`. In `sql_select` it is used for removing duplicates, while in `insert` it is used for `REPLACE` calls.

Parameters

- `buf`

Return Values

There are no return values.

Usage

This section is yet to be written.

Default Implementation

```
{ return HA_ERR_WRONG_COMMAND; }
```

delete_table

Purpose

Delete all files with extension from [custom-engine.html#custom-engine-api-reference-bas_ext `bas_ext()`].

Synopsis

```
virtual int delete_table | name);  
(  
  
const char * | name ;
```

Description

This is the `delete_table` method.

Used to delete a table. By the time `delete_table()` has been called all opened references to this table will have been closed (and your globally shared references released). The variable name will be the name of the table. You will need to remove any files you have created at this point.

If you do not implement this, the default `delete_table()` is called from `handler.cc`, and it will delete all files with the file extensions returned by `bas_ext()`. We assume that the handler may return more extensions than were actually used for the file.

Called from `handler.cc` by `delete_table` and `ha_create_table()`. Only used during create if the `table_flagHA_DROP_BEFORE_CREATE` was specified for the storage engine.

Parameters

- `name` : Base name of table

Return Values

- 0 if we successfully deleted at least one file from `base_ext` and didn't get any other errors than `ENOENT`
- # : Error

Usage

Most storage engines can omit implementing this method.

external_lock

Purpose

Handles table locking for transactions.

Synopsis

<code>virtual int</code>	<code>thd,</code>	
<code>external_lock (</code>	<code>lock_type);</code>	
THD *	<code>thd;</code>	
int	<code>lock_type;</code>	

Description

This is the `external_lock` method.

The locking methods for mysql section in `lock.cc` has additional comments on this topic that may be useful to read.

This creates a lock on the table. If you are implementing a storage engine that can handle transactions, look at `ha_innodb.cc` to see how you will want to go about doing this. Otherwise you should consider calling `flock()` here.

Called from `lock.cc` by `lock_external()` and `unlock_external()`. Also called from `sql_table.cc` by `copy_data_between_tables()`.

Parameters

- `thd`
- `lock_type`

Return Values

There are no return values.

Default Implementation

```
{ return 0; }
```

extra

Purpose

Passes hints from the server to the storage engine.

Synopsis

<code>virtual int extra (</code>	<code>operation);</code>	
enum ha_extra_function	<code>operation;</code>	

Description

This is the [extra](#) method.

[extra\(\)](#) is called whenever the server wishes to send a hint to the storage engine. The [MyISAM](#) engine implements the most hints. [ha_innodb.cc](#) has the most exhaustive list of these hints.

Parameters

- [operation](#)

Return Values

There are no return values.

Usage

Most storage engines will simply return 0.

```
{ return 0; }
```

Default Implementation

By default your storage engine can opt to implement none of the hints.

```
{ return 0; }
```

index_end

This section is yet to be written.

Purpose

Indicates end of index scan, clean up any resources used.

Synopsis

```
virtual int index_end ( | );  
  
| ;
```

Description

This is the [index_end](#) method. Generally it is used as a counterpart to the [index_init](#) method, cleaning up any resources allocated for index scanning.

Parameters

This method has no parameters.

Return Values

This method has no return values.

Usage

Clean up all resources allocated, return 0.

Default Implementation

```
{ active_index=MAX_KEY; return 0; }
```

index_first

Purpose

Retrieve first row in index and return.

Synopsis

```
virtual int index_first ( buf);  
  
byte * buf ;
```

Description

This is the `index_first` method.

`index_first()` asks for the first key in the index.

Called from `opt_range.cc`, `opt_sum.cc`, `sql_handler.cc`, and `sql_select.cc`.

Parameters

- `buf` - byte array to be populated with row.

Return Values

There are no return values.

Usage

Implementation depends on indexing method used.

Default Implementation

```
{ return HA_ERR_WRONG_COMMAND; }
```

index_init

Purpose

Signals the storage engine that an index scan is about to occur. Storage engine should allocate any resources needed.

Synopsis

```
virtual int index_init ( idx,  
sorted);  
  
uint idx;  
bool sorted;
```

Description

This is the `index_init` method. This method is called before an index scan, allowing the storage engine to allocate resources and make preparations.

Parameters

- `idx`
- `sorted`

Return Values

-

Usage

This method can typically just return 0 if there is no preparation needed.

Default Implementation

```
{ active_index=idx; return 0; }
```

index_last

Purpose

Return the last row in the index.

Synopsis

```
virtual int index_last ( |buf);  
  
byte * |buf ;
```

Description

This is the `index_last` method.

`index_last()` asks for the last key in the index.

Called from `opt_range.cc`, `opt_sum.cc`, `sql_handler.cc`, and `sql_select.cc`.

Parameters

- `buf` - byte array to be populated with matching row.

Return Values

This method has no return values.

Usage

Advance to last row in index and return row in buffer.

Default Implementation

```
{ return HA_ERR_WRONG_COMMAND; }
```

index_next

Purpose

Return next row in index.

Synopsis

```
virtual int index_next ( |buf);  
  
byte * |buf ;
```

Description

This is the `index_next` method.

Used to read forward through the index.

Parameters

- `buf`

Return Values

This method has no return values.

Usage

Advance to next row in index using pointer or cursor, return row in buffer.

Default Implementation

```
{ return HA_ERR_WRONG_COMMAND; }
```

index_prev

Purpose

Advance to previous row in index.

Synopsis

```
virtual int index_prev ( |buf);  
  
byte * |buf ;
```

Description

This is the `index_prev` method.

Used to read backward through the index.

Parameters

- `buf`

Return Values

This method has no return values.

Usage

Move to previous row in index, return in buffer.

Default Implementation

```
{ return HA_ERR_WRONG_COMMAND; }
```

index_read

Purpose

Find a row based on a key and return.

Synopsis

<code>virtual int</code>	<code>index_read (</code>	<code>buf,</code>	
		<code>key,</code>	
		<code>keypart_map,</code>	
		<code>find_flag);</code>	
<code>byte *</code>		<code>buf ;</code>	
<code>const byte *</code>		<code>key ;</code>	
<code>ulonglong</code>		<code>keypart_map ;</code>	
<code>enum ha_rkey_function</code>		<code>find_flag ;</code>	

Description

This is the `index_read` method.

Positions an index cursor to the index specified in the handle. Fetches the row if available. If the key value is null, begin at the first key of the index.

Parameters

- `buf`
- `key`
- `keypart_map`
- `find_flag`

Return Values

This method has no return values.

Usage

This section is still to be written.

Default Implementation

```
{ return HA_ERR_WRONG_COMMAND; }
```

index_read_idx

Purpose

Find a row based on a key and return.

Synopsis

<code>virtual int index_read_idx (</code>	<code>buf,</code>	
	<code>index,</code>	
	<code>key,</code>	
	<code>keypart_map,</code>	
	<code>find_flag);</code>	
<code>byte *</code>	<code>buf ;</code>	
<code>uint</code>	<code>index ;</code>	
<code>const byte *</code>	<code>key ;</code>	
<code>ulonglong</code>	<code>keypart_map ;</code>	
<code>enum ha_rkey_function</code>	<code>find_flag ;</code>	

Description

This is the `index_read_idx` method.

Positions an index cursor to the index specified in key. Fetches the row if any. This is only used to read whole keys.

Parameters

- `buf`
- `index`
- `key`
- `keypart_map`
- `find_flag`

Return Values

This method has no return values.

Usage

Locate the row that matches the key passed and return it in the buffer provided.

index_read_last

Purpose

This section is still to be written.

Synopsis

<code>virtual int index_read_last (</code>	<code>buf,</code>	
	<code>key,</code>	
	<code>keypart_map);</code>	
<code>byte *</code>	<code>buf ;</code>	

const byte *	key ;
ulonglong	keypart_map ;

Description

This is the `index_read_last` method.

Parameters

- `buf`
- `key`
- `keypart_map`

Return Values

-

Usage

This section is still to be written.

Default Implementation

```
{ return (my_errno=HA_ERR_WRONG_COMMAND); }
```

info

Purpose

Prompts storage engine to report statistics.

Synopsis

```
virtual int info (          |flag);
uint          |flag;
```

Description

This is the `info` method.

`info()` is used to return information to the optimizer. Currently, this table handler doesn't implement most of the fields really needed. `SHOW` also makes use of this data. Another note, you will probably want to have the following in your code: `if (records < 2) records = 2`. The reason is that the server will optimize for cases of only a single record. If in a table scan you don't know the number of records it will probably be better to set records to two so you can return as many records as you need. Along with records a few more variables you may wish to set are: records deleted, data_file_length, index_file_length, delete_length, check_time. See public variables in `handler.h` for more information.

Called in: `filesort.cc` `ha_heap.cc` `item_sum.cc` `opt_sum.cc` `sql_delete.cc` `sql_delete.cc` `sql_derived.cc` `sql_select.cc` `sql_select.cc` `sql_select.cc` `sql_select.cc` `sql_select.cc` `sql_show.cc` `sql_show.cc` `sql_show.cc` `sql_show.cc` `sql_table.cc` `sql_union.cc` `sql_update.cc`

Parameters

- `flag`

Return Values

0 for if no error occurred, [HA_ERR_xxx](#) if an error occurred.

Usage

This example is from the [CSV](#) storage engine:

```
int ha_tina::info(uint flag)
{
    DEBUG_ENTER("ha_tina::info");
    /* This is a lie, but you don't want the optimizer to see zero or 1 */
    if (!records_is_known && stats.records < 2)
        stats.records= 2;
    DEBUG_RETURN(0);
}
```

open

Purpose

Opens a table.

Synopsis

<code>virtual int open (</code>	<code>name,</code>	
	<code>mode,</code>	
	<code>test_if_locked);</code>	
<code>const char *</code>	<code>name ;</code>	
<code>int</code>	<code>mode ;</code>	
<code>uint</code>	<code>test_if_locked ;</code>	

Description

This is the [open](#) method.

Used for opening tables. The name will be the name of the file. A table is opened when it needs to be opened. For instance when a request comes in for a select on the table (tables are not open and closed for each request, they are cached).

Called from handler.cc by handler::ha_open(). The server opens all tables by calling ha_open() which then calls the handler specific open().

A handler object is opened as part of its initialization and before being used for normal queries (not before meta-data changes always.) If the object was opened it will also be closed before being deleted.

This is the [open](#) method. [open](#) is called to open a database table.

The first parameter is the name of the table to be opened. The second parameter determines what file to open or what operation to take. The values are defined in [handler.h](#) and are copied here for your convenience:

```
#define HA_OPEN_KEYFILE      1
#define HA_OPEN_RNDFILE     2
#define HA_GET_INDEX        4
#define HA_GET_INFO         8 /* do a ha_info() after open */
#define HA_READ_ONLY       16 /* File opened as readonly */
#define HA_TRY_READ_ONLY   32 /* Try readonly if can't open with read and write */
#define HA_WAIT_IF_LOCKED  64 /* Wait if locked on open */
```

```
#define HA_ABORT_IF_LOCKED 128 /* skip if locked on open.*/
#define HA_BLOCK_LOCK      256 /* unlock when reading some records */
#define HA_OPEN_TEMPORARY  512
```

The final option dictates whether the handler should check for a lock on the table before opening it.

Typically your storage engine will need to implement some form of shared access control to prevent file corruption in a multi-threaded environment. For an example of how to implement file locking, see the [get_share\(\)](#) and [free_share\(\)](#) methods of [sql/examples/ha_tina.cc](#).

Parameters

- `name`
- `mode`
- `test_if_locked`

Return Values

There are no return values.

Usage

This example is from the [CSV](#) storage engine:

```
int ha_tina::open(const char *name, int mode, uint test_if_locked)
{
    DEBUG_ENTER("ha_tina::open");

    if (!(share= get_share(name, table)))
        DEBUG_RETURN(1);
    thr_lock_data_init(&share->lock,&lock,NULL);
    ref_length=sizeof(off_t);

    DEBUG_RETURN(0);
}
```

position

Purpose

Provide the MySQL server with position/offset information for last-read row.

Synopsis

```
virtual void position ( |record);

const byte * |record;
```

Description

This is the [position](#) method.

`position()` is called after each call to `rnd_next()` if the data needs to be ordered. You can do something like the following to store the position: `my_store_ptr(ref, ref_length, current_position);`

The server uses `ref` to store data. `ref_length` in the above case is the size needed to store `current_position`. `ref` is just a byte array that the server will maintain. If you are using offsets to mark rows, then `current_position` should be the offset. If it is a primary key, then it needs to be a primary key.

Called from `filesort.cc`, `sql_select.cc`, `sql_delete.cc` and `sql_update.cc`.

Parameters

- `record`

Return Values

This method has no return values.

Usage

Return offset or retrieval key information for last row.

records_in_range

Purpose

For the given range how many records are estimated to be in this range.

Synopsis

<code>virtual ha_rows</code>	<code>inx,</code>	
<code>records_in_range (</code>		
	<code>min_key,</code>	
	<code>max_key);</code>	
<code>uint</code>	<code>inx;</code>	
<code>key_range *</code>	<code>min_key;</code>	
<code>key_range *</code>	<code>max_key;</code>	

Description

This is the `records_in_range` method.

Given a starting key, and an ending key estimate the number of rows that will exist between the two. `end_key` may be empty which in case determine if `start_key` matches any rows.

Used by optimizer to calculate cost of using a particular index.

Called from `opt_range.cc` by `check_quick_keys()`.

Parameters

- `inx`
- `min_key`
- `max_key`

Return Values

Return the approxamite number of rows.

Usage

Determine an approxamite count of the rows between the key values and return.

Default Implementation

```
{ return (ha_rows) 10; }
```

`rnd_init`

Purpose

Initializes a handler for table scanning.

Synopsis

```
virtual int rnd_init (scan);  
  
bool scan;
```

Description

This is the `rnd_init` method.

`rnd_init()` is called when the system wants the storage engine to do a table scan.

Unlike `index_init()`, `rnd_init()` can be called two times without `rnd_end()` in between (it only makes sense if `scan=1`). then the second call should prepare for the new table scan (e.g if `rnd_init` allocates the cursor, second call should position it to the start of the table, no need to deallocate and allocate it again

Called from `filesort.cc`, `records.cc`, `sql_handler.cc`, `sql_select.cc`, `sql_table.cc`, and `sql_update.cc`.

Parameters

- `scan`

Return Values

There are no return values.

Usage

This example is from the `CSV` storage engine:

```
int ha_tina::rnd_init(bool scan)  
{  
    DEBUG_ENTER("ha_tina::rnd_init");  
  
    current_position= next_position= 0;  
    records= 0;  
    chain_ptr= chain;  
    DEBUG_RETURN(0);  
}
```

`rnd_next`

Purpose

Reads the next row from a table and returns it to the server.

Synopsis

```
virtual int rnd_next ( buf);  
  
byte * buf;
```

Description

This is the `rnd_next` method.

This is called for each row of the table scan. When you run out of records you should return HA_ERR_END_OF_FILE. Fill buff up with the row information. The Field structure for the table is the key to getting data into buf in a manner that will allow the server to understand it.

Called from filesort.cc, records.cc, sql_handler.cc, sql_select.cc, sql_table.cc, and sql_update.cc.

Parameters

- `buf`

Return Values

There are no return values.

Usage

This example is from the [ARCHIVE](#) storage engine:

```
int ha_archive::rnd_next(byte *buf)
{
    int rc;
    DEBUG_ENTER("ha_archive::rnd_next");

    if (share->crashed)
        DEBUG_RETURN(HA_ERR_CRASHED_ON_USAGE);

    if (!scan_rows)
        DEBUG_RETURN(HA_ERR_END_OF_FILE);
    scan_rows--;

    statistic_increment(table->in_use->status_var.ha_read_rnd_next_count,
                        &LOCK_status);
    current_position= gztell(archive);
    rc= get_row(archive, buf);

    if (rc != HA_ERR_END_OF_FILE)
        records++;

    DEBUG_RETURN(rc);
}
```

rnd_pos

Purpose

Return row based on position.

Synopsis

<code>virtual int rnd_pos (</code>	<code>buf,</code>	
	<code>pos);</code>	
<code>byte *</code>	<code>buf ;</code>	
<code>byte *</code>	<code>pos ;</code>	

Description

This is the `rnd_pos` method.

Used for finding row previously marked with position. This is useful for large sorts.

This is like `rnd_next`, but you are given a position to use to determine the row. The position will be of the type that you stored in ref. You can use `ha_get_ptr(pos,ref_length)` to retrieve whatever key

or position you saved when position() was called. Called from filesort.cc records.cc sql_insert.cc sql_select.cc sql_update.cc.

Parameters

- `buf`
- `pos`

Return Values

This method has no return values.

Usage

Locate row based on position value and return in buffer provided.

start_stmt

Purpose

Called at the beginning of a statement for transaction purposes.

Synopsis

<code>virtual int start_stmt (</code>	<code>thd,</code>	
	<code>lock_type);</code>	
THD *	<code>thd ;</code>	
thr_lock_type	<code>lock_type ;</code>	

Description

This is the `start_stmt` method.

When table is locked a statement is started by calling `start_stmt` instead of `external_lock`

Parameters

- `thd`
- `lock_type`

Return Values

This method has no return values.

Usage

Make any preparations needed for a transaction start (if there is no current running transaction).

Default Implementation

```
{return 0;}
```

store_lock

Purpose

Creates and releases table locks.

Synopsis

<code>virtual THR_LOCK_DATA **</code>	<code>thd,</code>	
<code>store_lock (</code>		
	<code>to,</code>	
	<code>lock_type);</code>	
<code>THD *</code>	<code>thd;</code>	
<code>THR_LOCK_DATA **</code>	<code>to;</code>	
<code>enum thr_lock_type</code>	<code>lock_type;</code>	

Description

This is the `store_lock` method.

The idea with `handler::store_lock()` is the following:

The statement decided which locks we should need for the table for updates/deletes/inserts we get WRITE locks, for SELECT... we get read locks.

Before adding the lock into the table lock handler **mysql** calls store lock with the requested locks. Store lock can modify the lock level, e.g. change blocking write lock to non-blocking, ignore the lock (if we don't want to use MySQL table locks at all) or add locks for many tables (like we do when we are using a MERGE handler).

When releasing locks, `store_lock()` are also called. In this case one usually doesn't have to do anything.

If the argument of `store_lock` is `TL_IGNORE`, it means that MySQL requests the handler to store the same lock level as the last time.

Called from `lock.cc` by `get_lock_data()`.

Parameters

- `thd`
- `to`
- `lock_type`

Return Values

There are no return values.

Usage

The following example is from the [ARCHIVE](#) storage engine:

```
/*
  Below is an example of how to setup row level locking.
*/
THR_LOCK_DATA **ha_archive::store_lock(THD *thd,
                                         THR_LOCK_DATA **to,
                                         enum thr_lock_type lock_type)
{
    if (lock_type == TL_WRITE_DELAYED)
        delayed_insert= TRUE;
    else
        delayed_insert= FALSE;
```

```

if (lock_type != TL_IGNORE && lock.type == TL_UNLOCK)
{
    /*
     Here is where we get into the guts of a row level lock.
     If TL_UNLOCK is set
     If we are not doing a LOCK TABLE or DISCARD/IMPORT
     TABLESPACE, then allow multiple writers
    */

    if ((lock_type >= TL_WRITE_CONCURRENT_INSERT &&
        lock_type <= TL_WRITE) && !thd->in_lock_tables
        && !thd->tablespace_op)
        lock_type = TL_WRITE_ALLOW_WRITE;

    /*
     In queries of type INSERT INTO t1 SELECT ... FROM t2 ...
     MySQL would use the lock TL_READ_NO_INSERT on t2, and that
     would conflict with TL_WRITE_ALLOW_WRITE, blocking all inserts
     to t2. Convert the lock to a normal read lock to allow
     concurrent inserts to t2.
    */

    if (lock_type == TL_READ_NO_INSERT && !thd->in_lock_tables)
        lock_type = TL_READ;

    lock.type=lock_type;
}

*to++= &lock;

return to;
}

```

The following is the minimal implementation, for a storage engine that does not need to downgrade locks:

```

THR_LOCK_DATA **ha_tina::store_lock(THD *thd,
                                     THR_LOCK_DATA **to,
                                     enum thr_lock_type lock_type)
{
    /* Note that if the lock type is TL_IGNORE we don't update lock.type,
     preserving the previous lock level */
    if (lock_type != TL_IGNORE && lock.type == TL_UNLOCK)
        lock.type=lock_type;
    /* the heart of the store_lock() method and it's main purpose -
     storing the (possibly changed) lock level into the provided
     memory */
    *to++= &lock;
    return to;
}

```

See also [ha_myisammrg::store_lock\(\)](#) for more complex implementation

update_row

Purpose

Updates the contents of an existing row.

Synopsis

<code>virtual int</code>	<code>update_row (</code>	<code>old_data,</code>	
		<code>new_data);</code>	
<code>const byte *</code>		<code>old_data ;</code>	
<code>byte *</code>		<code>new_data ;</code>	

Description

This is the `update_row` method.

`old_data` will have the previous row record in it, while `new_data` will have the newest data in it.

The server can do updates based on ordering if an ORDER BY clause was used. Consecutive ordering is not guaranteed.

Currently, `new_data` will not have an updated `auto_increment` record, or and updated timestamp field. You can do these for example by doing these: if (table->timestamp_field_type & TIMESTAMP_AUTO_SET_ON_UPDATE) table->timestamp_field->set_time(); if (table->next_number_field && record == table->record[0]) update_auto_increment();

Called from `sql_select.cc`, `sql_acl.cc`, `sql_update.cc`, and `sql_insert.cc`.

Parameters

- `old_data`
- `new_data`

Return Values

There are no return values.

Usage

This section is still to be written.

Default Implementation

```
{ return HA_ERR_WRONG_COMMAND; }
```

write_row

Purpose

Adds a new row to a table.

Synopsis

```
virtual int write_row ( |buf);

byte * |buf ;
```

Description

This is the `write_row` method.

`write_row()` inserts a row. No [custom-engine.html#custom-engine-api-reference-extra `extra()`] hint is given currently if a bulk load is happening. `buf` is a byte array of data with a size of `table->s->reclength`

You can use the field information to extract the data from the native byte array type. Example of this would be: `for (Field **field=table->field ; *field ; field++) { ... }`

BLOBs must be handled specially:

```
for (ptr= table->s->blob_field, end= ptr + table->s->blob_fields ; ptr != end ; ptr++)
```

```
{
    char *data_ptr;
    uint32 size= ((Field_blob*)table->field[*ptr])->get_length();
    ((Field_blob*)table->field[*ptr])->get_ptr(&data_ptr);
    ...
}
```

See `ha_tina.cc` for an example of extracting all of the data as strings.

See the note for `update_row()` on `auto_increments` and `timestamps`. This case also applied to `write_row()`.

Called from `item_sum.cc`, `item_sum.cc`, `sql_acl.cc`, `sql_insert.cc`, `sql_insert.cc`, `sql_select.cc`, `sql_table.cc`, `sql_udf.cc`, and `sql_update.cc`.

Parameters

- `buf` byte array of data

Return Values

There are no return values.

Usage

This section is still to be written.

Default Implementation

```
{ return HA_ERR_WRONG_COMMAND; }
```

FAQ

Q: I've written my own storage engine, compiled it as shared object but when loading it I get an error like "undefined symbol: `_ZTI7handler`." What's wrong?

A: Make sure you compile and link your extension using the same flags as the server uses. The usual reason for this error is that `LD_FLAGS` are missing the `-fno-rtti` option.

Chapter 25. Test Synchronization

Table of Contents

Sleep	321
Wait Condition	322
Dbug Sleep	322
Error Injection	323
User-Level Locks	324
Debug Sync Point	325
Backup Breakpoint	326
Debug Sync Facility	326
Formal Syntax for DEBUG_SYNC Values	328
Debug Sync Activation/Deactivation	328
Debug Sync Implementation	329
A Typical Synchronization Pattern	329
Co-work With the DBUG Facility	330
Debug Sync Further Reading	330

There is a class of problems that require two or more cooperating threads to reproduce them.

A subclass of these problems is known as "race conditions". They require one thread to execute a certain piece of code while another thread executes another certain piece of code.

The vast majority of race conditions cannot be repeated reliably without some sort of synchronization of the involved threads. In most cases it is unlikely that the threads run through these code pieces at the right time. In this context 'synchronization' means to force the threads to meet at the critical code places.

In this chapter I'll describe some synchronization mechanisms:

- Sleep
- Wait Condition
- Dbug Sleep
- Error Injection
- User-Level Locks
- Debug Sync Point
- Backup Breakpoint
- Debug Sync Facility

Sleep

In some cases race conditions can be repeated when all but one thread are blocked (for example waiting for an SQL lock). Then the remaining thread has plenty of time to go through the critical piece of code.

The problem here is to assure that the blocking threads run until they reach their blocking point before the remaining thread reaches the critical code.

One solution is to use the 'sleep' command of 'mysqltest' in front of the SQL statement that drives the remaining thread into the critical code.

Example:

```
--connection conn1
LOCK TABLE t1 WRITE;
--connection conn2
# This will block in wait_for_lock().
send INSERT INTO t1 VALUES (1);
--connection conn1
# Sleep until we can be sure that conn2 reached wait_for_lock().
--sleep 2
# Run through the critical code.
FLUSH TABLE t1;
```

The BIG, BIG problem with 'sleep' is that you need to specify a fixed time. It must be big enough so that the test works as intended even on a very slow machine that is under heavy load. Hence it is much too big for the average machine. A major waste of time.

The bottom line is: AVOID 'SLEEP' WHEREVER POSSIBLE.

Wait Condition

Like 'sleep', this method can also be used, when all but one thread reach a blocked state.

If you are able to detect that the threads are in their blocked state by using SQL statements, then you can use this method. The remaining thread runs the statement(s) until the expected result is returned. Then it continues with the test.

Example:

```
--connection conn1
LOCK TABLE t1 WRITE;
--connection conn2
# Get the id of this thread.
let $conn2_id= `SELECT CONNECTION_ID`;
# This will block in wait_for_lock().
send INSERT INTO t1 VALUES (1);
--connection conn1
# Specify the condition that shows if conn2 reached wait_for_lock().
let $wait_condition= SELECT 1 FROM INFORMATION_SCHEMA.PROCESSLIST
WHERE ID = $conn2_id AND STATE = 'Locked';
# Run the condition in a loop until it becomes true.
--source include/wait_condition.inc
# Run through the critical code.
FLUSH TABLE t1;
```

In conn2 we get the thread ID first. In conn1 we use a SELECT statement that returns '1' when the processlist shows that conn2 reached the 'Locked' state. With this setup we call the wait_condition method. It runs the statement and checks the result. If the condition is not met, it sleeps for 0.1 second and retries.

The maximum waste of time is 0.1 seconds. This is much better than the 'sleep' method, but could still waste a little time.

Another problem is that the condition could be "fuzzy" in some situations. In the example above, the thread state (proc_info) is set to "Locked" right before the locking function is called. In theory it could happen that conn1 continues before conn2 did acquire the lock. The test would then fail to repeat what it was intended to do.

The "Debug Sync Facility" should be able to replace most of the "wait condition" uses.

Dbg Sleep

In cases where the normal server code does not have a block point at the critical place, one can insert an artificial synchronization point.


```

open_tables(...)

DEBUG_EXECUTE_IF("sleep_open_and_lock_after_open", {
    const char *old_proc_info= thd->proc_info;
    thd->proc_info= "DEBUG sleep";
    my_sleep(6000000);
    thd->proc_info= old_proc_info;});

lock_tables(...)

```

In this case, if the 'debug' keyword 'sleep_open_and_lock_after_open' is set, a thread sleeps for 6 seconds after open_tables() and before lock_tables(). Before sleeping, it sets the thread state (proc_info) to 'DEBUG sleep'. The test file that uses this synchronization point looks like so:

```

--connection conn1
let $conn1_id= `SELECT CONNECTION_ID()`;
# System variable 'debug' exists only in debug servers
--error 0, ER_UNKNOWN_SYSTEM_VARIABLE
SET SESSION debug="+d,sleep_open_and_lock_after_open";
send INSERT INTO t1 VALUES (1);
--connection conn2
# Specify the condition that shows if conn1 reached the sync point.
let $wait_condition= SELECT 1 FROM INFORMATION_SCHEMA.PROCESSLIST
    WHERE ID = $conn1_id AND STATE = 'DEBUG sleep';
# Run the condition in a loop until it becomes true.
--source include/wait_condition.inc
# Run through the critical code.
FLUSH TABLE t1;

```

So one can add synchronization points almost everywhere. But only at the cost of the wasted time of a sleep + a wait condition.

This method requires that you modify and recompile the server code. Another problem is that the synchronization point does not exist in non-debug servers. Not even the system variable 'debug' exists in a non-debug server. Each test must be written so that it works on a debug server as well as on a non-debug server. If this is not possible, the test must be moved into a test file that includes 'have_debug.inc'. Setting the possibly not existing variable can be protected by the --error 0, ER_UNKNOWN_SYSTEM_VARIABLE command. It says that the next statement can either succeed (0) or fail (ER_UNKNOWN_SYSTEM_VARIABLE).

Finally the method is bad when the execution should be traced with the DEBUG facility. Setting one (or more) 'debug' keywords disables all other keywords. One would need to add a pretty long list for a meaningful trace.

The bottom line is: Use the "Dbug Sleep" method when there is no other way to repeat a problem. However, the "Debug Sync Facility" should be able to replace all "Dbug Sleep" synchronization points.

Error Injection

Note: The ERROR_INJECT framework has been removed in an early 6.0 version. It is also not available from 5.5 upwards. It might be added back later.

The error injection method is based on the DEBUG framework just like the Dbug Sleep method. In the code you can use the following macros:

```

ERROR_INJECT_ACTION(keyword,action)
ERROR_INJECT_CRASH(keyword)
ERROR_INJECT(keyword)
SET_ERROR_INJECT_VALUE(value)
ERROR_INJECT_VALUE_ACTION(value,action)
ERROR_INJECT_VALUE_CRASH(value)
ERROR_INJECT_VALUE(value)

```

'keyword' is the debug keyword that you set in the test file with:

```
SET SESSION debug='+d,keyword1,keyword2,keyword3';
```

'value' is an unsigned long integer value. It is stored in `THD::error_inject_value` by `SET_ERROR_INJECT_VALUE(value)` and examined by the other `*_VALUE*` macros.

All of the `ERROR_INJECT_*` macros can/must be used in an expression. Their value is 0 (zero) in most cases. Exceptions are mentioned below.

Most of the `ERROR_INJECT_*` macros remove the keyword from the debug keyword list or clear `THD::error_inject_value` respectively before they executes their action. This means each of them will never execute twice within one SQL statement. But if multiple non-VALUE macros are run through in a statement, each can execute once if they use distinct keywords. There is just one `THD::error_inject_value`, not a list. So when any `*_VALUE*` macro clears it, all other `*_VALUE*` macros are disabled. Unless a new value is set by `SET_ERROR_INJECT_VALUE` somewhere. Obvious exceptions of keyword/value removal are `SET_ERROR_INJECT_VALUE` and the CRASH macros.

The `ERROR_INJECT_ACTION` macro is very similar to the `DEBUG_EXECUTE_IF` macro (see the "Dbug Sleep section). But remember the removal of the keyword/value.

The `ERROR_INJECT_VALUE_ACTION` is similar to `ERROR_INJECT_ACTION`. But it is controlled by the thread local value set by `SET_ERROR_INJECT_VALUE`. Also the action must be written as an expression. You can call a function that returns a value valid in the expression in which `ERROR_INJECT_VALUE_ACTION` appears. But if you want to open a block "{...}" you need to make an expression from it: "({...}, 0)". Also, if `ERROR_INJECT_VALUE_ACTION` is executed, it returns the value that the 'action' expression returns, not just 0 (zero) like `ERROR_INJECT_ACTION` does.

`ERROR_INJECT_CRASH` and `ERROR_INJECT_VALUE_CRASH` are pretty self-explanatory.

`ERROR_INJECT` and `ERROR_INJECT_VALUE` are for expression evaluation. They return 1 if the keyword is set or the value matches `THD::error_inject_value` respectively. Otherwise 0.

`SET_ERROR_INJECT_VALUE` copies the argument to `THD::error_inject_value`.

Downsides: The error injection method is NOT enabled in the server by default. You need to `./configure --with-error-inject`

The method is currently not used in the standard test suite anywhere. So you cannot copy and modify an example, but have to learn it the hard way.

When controlling error injection from the test files, explicit debug keywords are required, which has the same downsides as mentioned under Dbug Sleep.

If not using the `ERROR_INJECT` macros in an expression, expect the compiler warning "statement has no effect".

User-Level Locks

User-level locks are controlled with the SQL functions

```
GET_LOCK(str,timeout)
IS_FREE_LOCK(str)
IS_USED_LOCK(str)
RELEASE_LOCK(str)
```

They can be used at places where SQL statements accept SQL functions. Depending on their appearance in the select list, the where clause, the group by clause, etc, of select, update or other statements, these statements can be blocked at different code points. The set of blockable places is limited. Nevertheless, a couple of synchronization problems can be solved with user-level locks.

Example:

```
# Using InnoDB table with innodb_lock_wait_timeout=1 second.
--connection conn1
```

```
# Take an share lock on t1.
LOCK TABLE t1 IN SHARE MODE;
--connection conn2
# Acquire the user level lock "mysqltest1".
SELECT GET_LOCK("mysqltest1", 10);
# INSERT must wait in background for the SQL lock on t1 to go away.
send INSERT INTO t1 VALUES (1);
--connection conn1
# Wait in background until the insert times out and releases the
# user level lock. conn1 will then own the lock.
send SELECT GET_LOCK("mysqltest1", 10);
--connection conn2
# Wait for INSERT to timeout.
--error ER_LOCK_WAIT_TIMEOUT
reap;
# Now let conn1 get the lock and continue.
SELECT RELEASE_LOCK("mysqltest1");
COMMIT;
--connection conn1
reap;
# We do not need the lock any more.
SELECT RELEASE_LOCK("mysqltest1");
# Commit releases the share lock on t1.
COMMIT;
```

A good article about possible uses of user-level locks is from Martin Friebe. MySQL Internals mailing list, 10 Dec 2007: <http://lists.mysql.com/internals/35220>

One limitation of user-level locks is that a thread can have one lock at a time only. This limits the method to relatively simple cases.

Debug Sync Point

Note: Debug Sync Points were based on user-level locks. They were part of the MySQL code until the 6.0.5 and 5.1.46 versions. Debug Sync Points have been removed from the code in favor of the Debug Sync Facility.

Debug Sync Points give user-level locks the ability to synchronize at arbitrary points in code.

```
open_tables(...)

DEBUG_SYNC_POINT("debug_lock.after_open_tables", 10);

lock_tables(...)
```

The synchronization points behave similar to

```
RELEASE_LOCK(<whatever the thread has>);
IS_FREE_LOCK(str) OR (GET_LOCK(str,timeout) AND RELEASE_LOCK(str))
```

This means that the synchronization point releases any lock that the thread may have, waits to acquire the lock if another thread has it, and releases it immediately. If the lock is free (not used by any thread), the synchronization point does nothing but release any user-level lock of the current thread.

So the idea of `DEBUG_SYNC_POINT` is that it does nothing when the user-level lock is not in use by any thread, and does wait for it to become free when it is in use. That way you can block a thread at a synchronization point by acquiring the user-level lock and let it continue by releasing the lock.

This can be used as a "signal". The thread acquires a lock (the "signal" lock) and releases it implicitly when reaching the synchronization point. The other thread, which tried to get the "signal" lock after this thread, gets the lock at the same moment and can continue.

It can be used as a "wait". The other thread has the "synchronization point" lock ("debug_lock.after_open_tables" in this example) and this thread blocks on it in the synchronization point.

Unfortunately I was not able to figure out, how to use it for "signal" _plus_ "wait". While the other thread could have the "synchronization point" lock and this thread have the "signal" lock, and hence reaching the synchronization point would release the "signal" lock and wait on the "synchronization point" lock, the other thread would not be able to wait on the "signal" lock, because it has the "synchronization point" lock. A thread can have one user lock only. When the other thread tries to wait for the "signal" lock, it implicitly releases the "synchronization point" lock. This would be okay if one could be sure that this thread reached the synchronization point before the other thread releases the "synchronization point" lock. Otherwise no wait would happen at the synchronization point. The test would not test what it should test.

A possible workaround might be a third thread, which takes the "synchronization point" lock in the beginning and releases it at the right moment. But this could easily lead to a big number of threads for more complex situations. Tests using this method are likely to become ununderstandable.

It is probably a bug in the implementation that `DEBUG_SYNC_POINT` releases any lock unconditionally. The method is not widely used. I found just one single use in `sql_repl.cc`. I guess lock releasing was added to prevent that a synchronization point could wait on the threads own lock. The behavior could be fixed easily if the method should find more use.

The `DEBUG_SYNC_POINT` method is available in debug servers only. If it is used in the test suite, similar precautions for writing tests have to be taken as mentioned in the "Dbug Sleep" section.

Backup Breakpoint

Note: Backup Breakpoints were based on `DEBUG_SYNC_POINT`. They were part of the MySQL code in some early 6.0 versions. Backup Breakpoints have been removed from the code in favor of the Debug Sync Facility.

```
open_tables(...)  
  
BACKUP_BREAKPOINT("bp_after_open_tables");  
  
lock_tables(...)
```

The `BACKUP_BREAKPOINT` macro consists basically of:

```
DEBUG_EXECUTE_IF("backup_debug", DEBUG_SYNC_POINT((S), 300))
```

Opportunities and downsides of the `DEBUG_SYNC_POINT` method apply here too.

In addition we had the downside that `DEBUG` tracing was hampered as explained in the "Dbug Sleep" section.

Debug Sync Facility

The Debug Sync Facility is available as of MySQL 5.1.41, 5.5.0, and 6.0.6. Please note that - in spite of the "debug" in its name - this facility is completely independent from the `DEBUG` facility (except that it uses `DEBUG` to trace its operation, if `DEBUG` is also configured in the server). With a properly configured server (see [Debug Sync Activation/Deactivation](#)), this facility allows placement of synchronization points in the server code by using the `DEBUG_SYNC` macro:

```
open_tables(...)  
DEBUG_SYNC(thd, "after_open_tables");  
lock_tables(...)
```

When activated, a synchronization point can

- Emit a signal and/or
- Wait for a signal

Nomenclature:

signal	A value of a global variable that persists until overwritten by a new signal. The global variable can also be seen as a "signal post" or "flag mast". Then the signal is what is attached to the "signal post" or "flag mast".
emit a signal	Assign the value (the signal) to the global variable ("set a flag") and broadcast a global condition to wake those waiting for a signal.
wait for a signal	Loop over waiting for the global condition until the global value matches the wait-for signal.

By default, all synchronization points are inactive. They do nothing (except burn a couple of CPU cycles for checking if they are active).

A synchronization point becomes active when an action is requested for it. To do so, assign a value to the `DEBUG_SYNC` system variable:

```
SET DEBUG_SYNC= 'after_open_tables SIGNAL opened WAIT_FOR flushed';
```

This activates the synchronization point named 'after_open_tables'. The activation requests the synchronization point to emit the signal 'opened' and wait for another thread to emit the signal 'flushed' when the thread's execution runs through the synchronization point.

For every synchronization point there can be one action per thread only. Every thread can request multiple actions, but only one per synchronization point. In other words, a thread can activate multiple synchronization points.

Here is an example how to activate and use the synchronization points:

```
--connection conn1
SET DEBUG_SYNC= 'after_open_tables SIGNAL opened WAIT_FOR flushed';
send INSERT INTO t1 VALUES(1);
--connection conn2
SET DEBUG_SYNC= 'now WAIT_FOR opened';
SET DEBUG_SYNC= 'after_abort_locks SIGNAL flushed';
FLUSH TABLE t1;
```

When conn1 runs through the `INSERT` statement, it hits the synchronization point 'after_open_tables'. It notices that it is active and executes its action. It emits the signal 'opened' and waits for another thread to emit the signal 'flushed'.

conn2 waits immediately at the special synchronization point 'now' for another thread to emit the 'opened' signal.

A signal remains in effect until it is overwritten. If conn1 signals 'opened' before conn2 reaches 'now', conn2 will still find the 'opened' signal. It does not wait in this case.

When conn2 reaches 'after_abort_locks', it signals 'flushed', which lets conn1 awake.

Normally the activation of a synchronization point is cleared when it has been executed. Sometimes it is necessary to keep the synchronization point active for another execution. You can add an execute count to the action:

```
SET DEBUG_SYNC= 'name SIGNAL sig EXECUTE 3';
```

This sets the synchronization point's activation counter to 3. Each execution decrements the counter. After the third execution the synchronization point becomes inactive.

One of the primary goals of this facility is to eliminate sleeps from the test suite. In most cases it should be possible to rewrite test cases so that they do not need to sleep. (Note that Debug Sync can synchronize only multiple threads within a single process. It cannot synchronize multiple processes.) However, to support test development, and as a last resort, synchronization point waiting times out. There is a default timeout, but it can be overridden:

```
SET DEBUG_SYNC= 'name WAIT_FOR sig TIMEOUT 10 EXECUTE 2';
```

TIMEOUT 0 is special: If the signal is not present, the wait times out immediately.

If a wait timeout occurs (even on TIMEOUT 0), a warning is generated so that it shows up in the test result.

You can throw an error message and kill the query when a synchronization point is hit a certain number of times:

```
SET DEBUG_SYNC= 'name HIT_LIMIT 3';
```

Or combine it with signal and/or wait:

```
SET DEBUG_SYNC= 'name SIGNAL sig EXECUTE 2 HIT_LIMIT 3';
```

Here the first two hits emit the signal, the third hit returns the error message and kills the query.

For cases where you are not sure that an action is taken and thus cleared in any case, you can forcibly clear (deactivate) a synchronization point:

```
SET DEBUG_SYNC= 'name CLEAR';
```

If you want to clear all actions and clear the global signal, use:

```
SET DEBUG_SYNC= 'RESET';
```

This is the only way to reset the global signal to an empty string.

For testing of the facility itself you can execute a synchronization point just as if it had been hit:

```
SET DEBUG_SYNC= 'name TEST';
```

Formal Syntax for DEBUG_SYNC Values

The string to "assign" to the DEBUG_SYNC variable can contain:

```
{RESET |
  <sync point name> TEST |
  <sync point name> CLEAR |
  <sync point name> {{SIGNAL <signal name> |
    WAIT_FOR <signal name> [TIMEOUT <seconds>]}}
    [EXECUTE <count>] &| HIT_LIMIT <count>}
```

Here '&|' means 'and/or'. This means that one of the sections separated by '&|' must be present or both of them.

Debug Sync Activation/Deactivation

The Debug Sync facility is an optional part of the MySQL server. To cause Debug Sync to be compiled into the server, use the `--enable-debug-sync` option:

```
./configure --enable-debug-sync
```

Debug Sync is also compiled in if you configure with the `--with-debug` option (which implies `--enable-debug-sync`), unless you also use the `--disable-debug-sync` option.

The Debug Sync Facility, when compiled in, is disabled by default. To enable it, start `mysqld` with the `--debug-sync-timeout[=N]` option, where `N` is a timeout value greater than 0. `N` becomes the default timeout for the `WAIT_FOR` action of individual synchronization points. If `N` is 0, Debug Sync stays disabled. If the option is given without a value, the timeout is set to 300 seconds.

The `DEBUG_SYNC` system variable is the user interface to the Debug Sync facility. If Debug Sync is not compiled in, this variable is not available. If compiled in, the global `DEBUG_SYNC` value is read

only and indicates whether the facility is enabled. By default, Debug Sync is disabled and the value of `DEBUG_SYNC` is "OFF". If the server is started with `--debug-sync-timeout=N`, where `N` is a timeout value greater than 0, Debug Sync is enabled and the value of `DEBUG_SYNC` is "ON - current signal" followed by the signal name. Also, `N` becomes the default timeout for individual synchronization points.

The session value can be read by any user and will have the same value as the global variable. The session value can be set by users that have the `SUPER` privilege to control synchronization points.

Setting the `DEBUG_SYNC` system variable requires the 'SUPER' privilege. You cannot read back the string that you assigned to the variable, unless you assign the value that the variable does already have. But that would give a parse error. A syntactically correct string is parsed into a debug synchronization point action and stored apart from the variable value.

The Debug Sync facility is enabled by default in the test suite, but can be disabled with:

```
mysql-test-run.pl ... --debug-sync-timeout=0 ...
```

Likewise, the default wait timeout can be set:

```
mysql-test-run.pl ... --debug-sync-timeout=10 ...
```

For test cases that require the Debug Sync facility, include the following line in the test case file:

```
--source include/have_debug_sync.inc
```

Debug Sync Implementation

Pseudo code for a synchronization point:

```
#define DEBUG_SYNC(thd, sync_point_name)
    if (unlikely(opt_debug_sync_timeout))
        debug_sync(thd, STRING_WITH_LEN(sync_point_name))
```

The synchronization point performs a binary search in a sorted array of actions for this thread.

The `SET DEBUG_SYNC` statement adds a requested action to the array or overwrites an existing action for the same synchronization point. When it adds a new action, the array is sorted again.

A Typical Synchronization Pattern

There are quite a few places in MySQL, where we use a synchronization pattern like this:

```
pthread_mutex_lock(&mutex);
thd->enter_cond(&condition_variable, &mutex, new_message);
#if defined(ENABLE_DEBUG_SYNC)
    if (!thd->killed && !end_of_wait_condition)
        DEBUG_SYNC(thd, "sync_point_name");
#endif
while (!thd->killed && !end_of_wait_condition)
    pthread_cond_wait(&condition_variable, &mutex);
thd->exit_cond(old_message);
```

Here some explanations:

`thd->enter_cond()` is used to register the condition variable and the mutex in `thd->mysys_var`. This is done to allow the thread to be interrupted (killed) from its sleep. Another thread can find the condition variable to signal and mutex to use for synchronization in this thread's `THD::mysys_var`.

`thd->enter_cond()` requires the mutex to be acquired in advance.

`thd->exit_cond()` unregisters the condition variable and mutex and releases the mutex.

If you want to have a Debug Sync point with the wait, please place it behind `enter_cond()`. Only then you can safely decide, if the wait will be taken. Also you will have `THD::proc_info` correct when the

sync point emits a signal. `DEBUG_SYNC` sets its own `proc_info`, but restores the previous one before releasing its internal mutex. As soon as another thread sees the signal, it does also see the `proc_info` from before entering the sync point. In this case it will be "new_message", which is associated with the wait that is to be synchronized.

In the example above, the wait condition is repeated before the sync point. This is done to skip the sync point, if no wait takes place. The sync point is before the loop (not inside the loop) to have it hit once only. It is possible that the condition variable is signaled multiple times without the wait condition to be true.

A bit off-topic: At some places, the loop is taken around the whole synchronization pattern:

```
while (!thd->killed && !end_of_wait_condition)
{
    pthread_mutex_lock(&mutex);
    thd->enter_cond(&condition_variable, &mutex, new_message);
    if (!thd->killed [&& !end_of_wait_condition])
    {
        [DEBUG_SYNC(thd, "sync_point_name");]
        pthread_cond_wait(&condition_variable, &mutex);
    }
    thd->exit_cond(old_message);
}
```

Note that it is important to repeat the test for `thd->killed` after `enter_cond()`. Otherwise the killing thread may kill this thread after it tested `thd->killed` in the loop condition and before it registered the condition variable and mutex in `enter_cond()`. In this case, the killing thread does not know that this thread is going to wait on a condition variable. It would just set `THD::killed`. But if we would not test it again, we would go asleep though we are killed. If the killing thread would kill us when we are after the second test, but still before sleeping, we hold the mutex, which is registered in `mysys_var`. The killing thread would try to acquire the mutex before signaling the condition variable. Since the mutex is only released implicitly in `pthread_cond_wait()`, the signaling happens at the right place. We have a safe synchronization.

Co-work With the DBUG Facility

When running the MySQL test suite with a "debug" server (the DBUG facility is configured in) and the `--debug` command line option, the Debug Sync Facility writes trace messages to the DBUG trace. The following shell commands proved very useful in extracting relevant information:

```
egrep 'query:|debug_sync_exec:' mysql-test/var/log/mysqld.1.trace
```

It shows all executed SQL statements and all actions executed by synchronization points.

Sometimes it is also useful to see, which synchronization points have been run through (hit) with or without executing actions. Then add "**|debug_sync_point:**" to the `egrep` pattern.

Debug Sync Further Reading

For complete syntax tests, functional tests, and examples see the test case `debug_sync.test`.

See also worklog entry [WL#4259 - Debug Sync Facility](#)

Reference manual 5.1

- 2.3.2 Typical configure Options (`--enable-debug-sync`)
- 5.1.2 Command Options (`--debug-sync-timeout`)
- 5.1.4 System Variables (`debug_sync`)

Test framework manual

- 4.14 Thread Synchronization in Test Cases (have_debug_sync.inc)
- 5.3 mysql-test-run.pl (--debug-sync-timeout)

Chapter 26. Injecting Test Faults

Table of Contents

Test Fault Macros	333
Test Fault Macro Usage	333

Status of this section: up to date 2010-09-08

The assessment of the replication code in the presence of faults is extremely important to increase reliability. In particular, one needs to know if servers will either correctly recover or print out appropriate error messages thus avoiding unexpected problems in a production environment. To fulfill this need, we use the macros presented in the following discussion.

Test Fault Macros

- **DEBUG_EXECUTE_IF (keyword, code)** allows to execute a piece of code if the appropriate debug instruction is set. In this case, the debug instruction should be **+d,keyword**.
- **DEBUG_EVALUATE_IF (keyword, val1, val2)** is used in "if" expressions and returns "val2" if the appropriate debug instruction is set. Otherwise, it returns "val1". In this case, the debug instruction should be **+d,keyword**.
- **DEBUG_CRASH_ENTER (function)** is equivalent to **DEBUG_ENTER** which registers the beginning of a function but in addition to it allows for crashing the server while entering the function if the appropriate debug instruction is set. In this case, the debug instruction should be **+d,function_crash_enter**.
- **DEBUG_CRASH_RETURN (value)** is equivalent to **DEBUG_RETURN** which notifies the end of a function but in addition to it allows for crashing the server while returning from the function if the appropriate debug instruction is set. In this case, the debug instruction should be **+d,function_crash_return**. Note that "function" should be the same string used by the **DEBUG_ENTER**.
- **DEBUG_CRASH_VOID_RETURN** is equivalent to **DEBUG_VOID_RETURN** which notifies the end of a function but in addition to it allows for crashing the server while returning from the function if the appropriate debug instruction is set. In this case, the debug instruction should be **+d,function_crash_return**. Note that "function" should be the same string used by the **DEBUG_ENTER**.

Test Fault Macro Usage

Let us assume the following function:

```
void function(void)
{
    DEBUG_CRASH_ENTER("function");
    if (DEBUG_EVALUATE_IF("process_if", 1, 0))
    {
        DEBUG_EXECUTE_IF("process_code", {
            const char *old_proc_info= thd->proc_info;
            thd->proc_info= "DEBUG sleep";
            my_sleep(6000000);
            thd->proc_info= old_proc_info;});
    }
    DEBUG_CRASH_VOID_RETURN;
}
```

To crash the server in the fault points defined above, we need to be **SUPER user** and execute one of the following commands:

- SET SESSION debug="+d,function_crash_enter"; will crash the server while entering the function.
- SET SESSION debug="+d,function_crash_return"; will crash the server while returning from function.
- SET SESSION debug="+d,process_if"; will execute the code inside the "if".
- SET SESSION debug="+d,process_code"; will execute the "sleep" code.

Chapter 27. How to Create Good Test Cases

Table of Contents

Formal Stuff	337
Coding Style	337
SQL Statement Example	337
Please Avoid too Dense Code	337
Header of Scripts	338
Header Example	338
Comments Within Boxes	338
Comments Everywhere	338
Complicated test Architecture or Tricky Code	339
A Subtest Case is Able to Reveal a Bug	339
Subtests	339
Make Test Protocols More Understandable	339
Better Protocol Example	340
Some Rules Outside of the Formal Stuff	340
Error Masking	340
Use of the Option <code>--disable_abort_on_error</code>	340
Perfect Cleanup at the End of a Test	341
Use of OS-Specific Commands	341
"Negative" Tests	341
Tests with Several Variants	342
Variation of the Storage Engine	342
Variation of the Protocol Option	343
Miscellaneous Tips	344
Stability Checks for Random Timing Problems	344
Stability Checks for Hostname-Related Problems	345
Stability Checks for Unexpected Problems	345
Row Order Within Result Sets	345
Test Case Behavior Should Depend on the Result of a <code>SHOW</code>	345
Does a Test Execution Hang?	345
If Your Test is Complicated, Support Debugging	345
You Need to Know Some Syntax, an Option, Etc., but Hate Reading Manuals	346
Notes on MTR2	346
Examples of "Suspicious" Tests and Scripts	346
Probably Mismatch of Focus of Test and Code Sequence	347
Too Greedy Test	347
Risky Handling Around Additional Sessions	348

There are a lot of strict rules and rules of thumb which may increase the quality of tests written for the tool pair `mysqltest/mysqltest-run.pl` (MTR). The following presentation should give you some hints.

- This document should not replace reading our excellent manual about testing <http://dev.mysql.com/doc/mysqltest/en/writing-tests.html>.
- Some of the examples might not work
 - with MTR1 or
 - on some operating system.
- MTR2 means the second version of `mysql-test-run.pl/mysqltest`. MTR2 replaced MTR1 ~ Jan 2009 in MySQL 5.1 and 6.0.
- The term "protocol" means in most cases "output from a test case."

October 2007 - Matthias Leich - Create this document for a MySQL University session
March 2009 - Matthias Leich - Updates and corrections, add "Examples of suspicious scripts"
March 2009 - Patrick Crews - Corrections, add "Notes on MTR"

- [Formal Stuff](#)
 - [Coding Style](#)
 - [SQL Statement Example](#)
 - [Header of Scripts](#)
- [Comments Everywhere](#)
 - [Complicated test Architecture or Tricky Code](#)
 - [A Subtest Case is Able to Reveal a Bug](#)
 - [Subtests](#)
 - [Make Test Protocols More Understandable](#)
 - [Better Protocol Example](#)
- [Some Rules Outside of the Formal Stuff](#)
 - [Error Masking](#)
 - [Use of the Option `--disable_abort_on_error`](#)
 - [Perfect Cleanup at the End of a Test](#)
 - [Use of OS-Specific Commands](#)
- [“Negative” Tests](#)
- [Tests with Several Variants](#)
 - [Variation of the Storage Engine](#)
 - [Variation of the Protocol Option](#)
- [Miscellaneous Tips](#)
 - [Stability Checks for Random Timing Problems](#)
 - [Stability Checks for Hostname-Related Problems](#)
 - [Stability Checks for Unexpected Problems](#)
 - [Row Order Within Result Sets](#)
 - [Test Case Behavior Should Depend on the Result of a `SHOW`](#)
 - [Does a Test Execution Hang?](#)
 - [If Your Test is Complicated, Support Debugging](#)
 - [You Need to Know Some Syntax, an Option, Etc., but Hate Reading Manuals](#)
- [Notes on MTR2](#)
- [Examples of “Suspicious” Tests and Scripts](#)

- [Probably Mismatch of Focus of Test and Code Sequence](#)
- [Too Greedy Test](#)
- [Risky Handling Around Additional Sessions](#)

Formal Stuff

Rules about formatting within scripts and similar stuff.

- [Coding Style](#)
- [SQL Statement Example](#)
- [Header of Scripts](#)

Coding Style

Please have a look at [the C/C++ coding guidelines](#) and apply them when they make sense.

Some more or less strict rules:

- Try to have a good readable text flow. Lines with very different length are a pain.
- Avoid lines longer than ~ 100 characters unless there is no other choice.

Think about a comparison of the new and the old version of a test script within a graphical diff tool.

Having the difference frequent at the end of long lines is very uncomfortable.

- mysqltest only accepts comment lines starting with '#'.
- Use spaces, not tabs.
- Lines must have no trailing spaces.
- Write SQL statements in the style of the MySQL Reference Manual
 - SQL keywords and reserved words: uppercase
 - Identifiers (table names, column names, etc.): lowercase

Ignore this guideline if your intent is to check the processing of mixed lettercases ;-)

Please follow this guideline for new tests. Rewriting existing tests to use better formatting is good, but can be tedious - a rule of thumb: don't rewrite unless you intend to touch the entire subtest (not the entire .test file, but the statements that make up a complete testing 'unit')

- If an SQL statement is long, add line breaks to reformat it and make it easier to read.

SQL Statement Example

```
SELECT f1 AS "my_column", f10 ....
FROM mysqltest1.t5
WHERE (f2 BETWEEN 17 AND 25 OR f2 = 61)
      AND f3 IN (SELECT ....
                  FROM mysqltest1.t4
                  WHERE .....)
```

```
ORDER BY ...
```

Please Avoid too Dense Code

except you intend to check the parser or the performance of the reader of the code ;-).

```

Unfortunate example:
select 1+1,1-1,1+1*2,8/5,8%5,mod(8,5),mod(8,5)|0,-(1+1)*-2;
Improved example:
SELECT 1 + 1, 1 - 1, 1 + 1 * 2, 8 / 5, 8 % 5,
      MOD(8,5), MOD(8,5) | 0, -(1+1) * -2;

```

Header of Scripts

For each test or auxiliary script, create a header that contains the following information:

- Purpose of the test or script
- Corresponding WL task, if there is any
- Creator of the test and date of creation
- Author of last significant change + date of change + what was changed
- Dates should be in ISO format (ISO 8601): YYYY-MM-DD
- In case that the script assigns values to some variables and sources some master test script, please explain the purpose of these variables and why you use these values

Note: The header, like the rest of the test, should not mention confidential information. Remember, our tests are available publicly.

Header Example

```

##### suite/funcs_1/t/processlist_val_no_prot.test #####
#
# Testing of values within INFORMATION_SCHEMA.PROCESSLIST
#
# The prepared statement variant of this test is
# suite/funcs_1/t/processlist_val_ps.test.
#
# There is important documentation within
#   suite/funcs_1/datadict/processlist_val.inc
#
#
# Creation:
# 2007-08-09 mleich Implement this test as part of
#               WL#3982 Test information_schema.processlist
#
#####

```

Comments Within Boxes

Please work 100% perfect when using boxes made of '#'. Such boxes "jump" into the eye of most readers. Bad example similar to what I found in some tests:

```

#####
#
# Some text
# Some text
#
# Author : .....
#
# Some text
#
#####

```

Please note that this is caused by spaces, not tabs.

Comments Everywhere

Write comments, they save the time of others.

- [Complicated test Architecture or Tricky Code](#)
- [A Subtest Case is Able to Reveal a Bug](#)
- [Subtests](#)
- [Make Test Protocols More Understandable](#)
- [Better Protocol Example](#)

Complicated test Architecture or Tricky Code

If your test is very complicated, write a sufficient explanation of the test architecture. This helps avoid having someone else come along and unintentionally change the test in ways that:

- Shift the goal of the check
- Destroy your test architecture
- etc.

Comment your code when you do or check something that someone else may think is not trivial. Please write what you intend to check if it is not obvious.

A Subtest Case is Able to Reveal a Bug

Write a comment if the next subtest revealed a bug in history. Please mention number and title.

Please use exact the formatting "# Bug#nnnnn <Title>"

```
Example: # Bug#3671 Stored procedure crash if function has "set @variable=param"
```

because it does not look nice to see so many different variants like "BUG #nnnnn", "bug #nnnnn - Title" often even within the same test.

Subtests

At least in cases where your file contains many subtests

- Mark these subtests for better readability of the script
- Write also a message into the protocol.

And please explain what each subtest checks, unless it is obvious.

Example:

```
--echo #----- Testcase 3.3.1.52 -----#
#####
# Ensure that a view that is a subset of every column and some rows of a single
# underlying table, contains the correct row-and-column data; such a view has
# a definition that is semantically equivalent to CREATE VIEW <view_name>
# AS SELECT * FROM <table_name> WHERE ...
#####
```

Make Test Protocols More Understandable

Please have the test write comments into the protocol if this makes the surrounding protocol content much more understandable. This is especially true for tests that do the following:

- Suppress the printing of SQL statements and result sets
- Work with more than one connection

Better Protocol Example

Good script with message about switching the connection:

```
SET @aux = 1;
SELECT @aux AS "content of @aux is";
--echo # Establish session con1 (user=root)
connect (con1,localhost,root,,,,);
SELECT @aux AS "content of @aux is";
```

Protocol of good script:

```
SET @aux = 1;
SELECT @aux AS "content of @aux is";
content of @aux is
1
# Establish session con1 (user=root)
SELECT @aux AS "content of @aux is";
content of @aux is
NULL
```

This is what the protocol looks like without the comment. The output becomes less clear because there is no indication that the connection has been changed.

```
SET @aux = 1;
SELECT @aux AS "content of @aux is";
content of @aux is
1
SELECT @aux AS "content of @aux is";
content of @aux is
NULL
```

Some Rules Outside of the Formal Stuff

- Prerequisites checks
- [Error Masking](#)
- [Use of the Option `--disable_abort_on_error`](#)
- [Perfect Cleanup at the End of a Test](#)
- [Use of OS-Specific Commands](#)

Error Masking

Use error names instead of error numbers whenever possible.

Example:

```
--error ER_BAD_DB_ERROR
USE <not existing database>;
```

You can find the error names in the `include/mysqld_error.h` file of a MySQL source distribution, or in the [Server Error Codes and Messages](#) section of the MySQL Reference Manual]. NOTE: Don't do this with error numbers > 2000. Use of the error name does not seem to work well in these cases.

Use of the Option `--disable_abort_on_error`

This option is very useful when starting to write a new test because the test will not abort if your script contains some failing statements or SQL syntax errors. By disabling the abort, you get to see more of the errors per test run and can fix more of them at a time.

But there are only rare situations where the final versions of a test should use `"--disable_abort_on_error"`, either at all or during a sequence of several SQL statements.

Perfect Cleanup at the End of a Test

- Cleaning up becomes much more comfortable and less error prone if you create your "own" database and create all tables there. At the end of the test, you need to drop just this database.
- Do not forget to remove all users you created and all permissions you granted. Otherwise the next test might fail when checking grants in general.
- Close all connections which you have explicitly created and please ensure that the disconnect is finished = the sessions are no more visible within the processlist.
- Remove all auxiliary files created within your test.

MTR2 runs now by default with "check-testcases" enabled. It checks if there are additional objects like user or tables, modified system table etc. "check-testcases" will be soon improved. However, proper cleanup should still be a responsibility of both the test author and reviewer.

Use of OS-Specific Commands

The `exec` and `system` commands enable tests to execute external commands. However, many of such commands are available only on certain platforms. (For example, `rm` is Unix-specific and not available on Windows.) Please avoid these commands if possible. They harm the portability and stability of tests.

Have a look at the `mysqltest` manual and the `t/mysqltest.test` test file. There are now several `mysqltest` built-in commands such as

```
--write_file", "--cat_file", "--remove_file", ...
```

which are reliable on all operating systems. Sometimes OS commands could be also replaced by SQL statement sequences writing to and reading from files and some SQL string functions.

“Negative” Tests

A "negative" test is a test for which you expect to see a failure. If an error does not occur, that itself indicates a problem.

DO NOT FORGET "NEGATIVE" TESTS where we expect to see fine error messages from the server.

This section contains just a few examples of what to test. Please be creative and imagine what could go wrong in rough reality.

Column with numeric data type:

- `NULL`, `0`
- `Minimum - 1`, `Minimum`, `Minimum + 1`
- `Maximum - 1`, `Maximum`, `Maximum + 1`
- Negative values if data type is unsigned
- Garbage like `'abc'`, `'1a'`, `'a1'`

Column with string data type:

- `Null`, `<empty string>`, "exotic" characters like `'ä'`, single quotes, ...
- String longer than column

Limited number of tables, columns, variables, ...

- `Maximum - 1`, `Maximum`, `Maximum + 1`

Assume an SQL operation that affects the filesystem (LOAD DATA, CREATE SCHEMA or TABLE, backup/restore, ...). What will happen if the following conditions occur for a file or directory to be read, written, or implicitly created or deleted?

- Exists/does not exist
- Is writable/not writable(missing permission)
- Is empty/not empty
- Contains the expected content (Example: text)/unexpected content like maybe a JPG)
- Is a regular file/a directory/a symlink pointing to .../a pipe
- Is assigned via an extremely long path
- Becomes victim of file system full

Tests with Several Variants

Please think twice before you create

- engine-specific variants of a test

or

- let runs with ps-protocol/sp-protocol/cursor-protocol/view-protocol

happen. They might be

- of low value and a permanent wasting of resources and/or
- fail (protocol variants only)
- [Variation of the Storage Engine](#)
- [Variation of the Protocol Option](#)

Variation of the Storage Engine

- Checks of the INFORMATION_SCHEMA

The storage engines to be used for information_schema tables are hardcoded. Therefore tests focussed on permissions, optimizer strategies, column data types etc. when selecting on INFORMATION_SCHEMA tables should not run with storage engine variations.

- Please prepare the scripts at least for use with different storage engines.

That means mostly:

DO NOT use hardcoded storage engine assignments within CREATE TABLE statements.

- If you assume that there is no significant impact of storage engines on your testing object, create/run with storage engine variants and check this.

Solution 1:

Do not assign the storage engine within your CREATE TABLE statements at all. The default storage engine MyISAM will be used for your tables. Check your assumption with:

```
./mysql-test-run.pl --mysqld="--default-storage-engine=<engine>" <test_case>
```

Solution 2:

Assign the storage engine to be used via \$variable.

Top level script:

```
let $engine_type= MyISAM;
```

The same script or a sourced script:

```
eval CREATE TABLE .... ENGINE = $engine_type ...
```

Check your assumption by creating and running storage engine variants of the top level test.

```
--source include/have_falcon.inc  
let $engine_type= Falcon;
```

Variation of the Protocol Option

It is usual to run tests with and without the mysql-test-run.pl startup option "--ps-protocol". And there are also attempts of System QA to run the other protocols.

- Effect of the "--ps-protocol" option: mysqltest will run as many SQL statements as possible as prepared statements.
- Effect of the "--sp-protocol"/"--cursor-protocol"/"--view-protocol" options: As far as I know, mysqltest takes many statements and transforms them into a statement sequence checking the corresponding feature (stored procedures, cursors, or views).

We have a lot of tests running very similar and extremely simple SQL just for the creation of a situation to be tested, check of table content, etc. So it can be assumed that the n'th test running again simple statements does not improve the coverage.

Conclusion:

If your test contains

- Simple statements: Prevent the non valuable runs for "protocols".
- Storage engine variations: Prevent the protocol runs for all except one storage engine (my take would be MyISAM).
- "Unique" and complex statements:
 - ps-protocol statements = all SQL
 - sp-protocol statements = DML
 - view/cursor-protocols statements = SELECTs

Do not prevent the protocol runs.

Another problem around runs with such protocols is that there are cases where we get different protocol content.

Example:

Script:

```
SELECT * FROM processlist ...
```

Protocol content if running without any "--*-protocol":

```
ID USER ... COMMAND ... STATE INFO  
1 root ... Query ... preparing SELECT * FROM processlist ...
```

Protocol content if running with "--ps-protocol":

```
ID USER ... COMMAND ... STATE INFO
1 root ... Execute ... preparing SELECT * FROM processlist ...
```

So please check whether every new test gives the same result with every protocol but at least with "--ps-protocol". If not, do something of the following:

- Exclude some protocol variants from execution.
- Disable the use of the protocols for problematic statements.
- Write protocol variant specific tests.

Example solution (code within the top level scripts):

```
# The file with expected results fits only to a run without
if (`SELECT $PS_PROTOCOL + $SP_PROTOCOL
    + $CURSOR_PROTOCOL + $VIEW_PROTOCOL > 0`)
{
    --skip Test requires: ps-protocol/sp-protocol/cursor-protocol/view-protocol disabled
}
--source include/<whatever>.inc
```

or

```
# The file with expected results fits only to a run with "--ps-protocol".
if (`SELECT $SP_PROTOCOL + $CURSOR_PROTOCOL + $VIEW_PROTOCOL > 0
    OR $PS_PROTOCOL = 0`)
{
    --skip Test requires: ps-protocol enabled, other protocols disabled
}
--source include/<whatever>.inc
```

Miscellaneous Tips

- [Stability Checks for Random Timing Problems](#)
- [Stability Checks for Hostname-Related Problems](#)
- [Stability Checks for Unexpected Problems](#)
- [Row Order Within Result Sets](#)
- [Test Case Behavior Should Depend on the Result of a SHOW](#)
- [Does a Test Execution Hang?](#)
- [If Your Test is Complicated, Support Debugging](#)
- [You Need to Know Some Syntax, an Option, Etc., but Hate Reading Manuals](#)

Stability Checks for Random Timing Problems

- **Recommended** Several test runs (<number>), use an memory based (extreme fast) filesystem if available

```
./mysql-test-run.pl --mem .... --repeat=<number> <your test>
```

- Several test runs, use a "slow" disk based filesystem

```
rmdir var # "var" must not be a symlink pointing into a RAM based filesystem
./mysql-test-run.pl .... --repeat=<number> <your test>
```

- **Recommended** Generate parallel load on the CPUs and/or the disk where MTR2 reads and writes data (in most cases directory "var"). Examples:

- all OS: Compile a MySQL version from source
- Unix derivatives (OpenSolaris, Linux, OSX,...): `tar -chvf - <some path> > /dev/null`
- Unix derivatives: root: `dd if=<area of disk containing the "var" directory> of=/dev/null`
- Windows: Defragmentation

```
rmmdir var # "var" must not be a symlink pointing into a RAM based filesystem
./mysql-test-run.pl .... --repeat=<number> <your test>
```

- all OS:

```
./mysql-test-run.pl .... --repeat=<number> --parallel=8 <your test>
```

Stability Checks for Hostname-Related Problems

```
sudo hostname <other hostname>
./mysql-test-run.pl --mem <name of test>
```

Good values for <other hostname> to check are "0", "bbb", "mmm", "zzz".

Stability Checks for Unexpected Problems

Ask a colleague to run your test or try another testing box (preferably one with a different operating system).

Row Order Within Result Sets

Please keep in mind that the row order within a significant number of result sets depends on storage engine properties and in worst cases the current load (NDB!, partitioning?, parallel query?) on the testing box.

Either use

```
--sorted_result
SELECT * FROM t1;
```

or, if possible, "decorate" your SELECT with ORDER BY to make the row order static.

Test Case Behavior Should Depend on the Result of a [SHOW](#)

Please have a look at the `t/mysqltest.test` file where `"query_get_value()"` is checked.

Does a Test Execution Hang?

Symptom: No text flooding through the command window where a test is running.

Run this command in a different window:

```
tail -f <var_dir>/mysqld.1/data/mysql/general_log.CSV
```

If you see some changes, your test is alive.

If Your Test is Complicated, Support Debugging

Top level test script:

```
...
##### Option, for debugging support #####
let $debug= 0;
```

```
...
```

At various places within your scripts:

```
if ($debug)
{
  --echo # var1: $var1 , var2: $var2, .....
  SELECT .....
}
```

or

```
--disable_query_log
if ($debug)
{
  --enable_query_log
}
```

You Need to Know Some Syntax, an Option, Etc., but Hate Reading Manuals

Try this:

```
grep -i '<one keyword you know or guess>' t/* include/* | less
```

Notes on MTR2

MTR2 has brought some new utility that can improve testing, but also some elements that can be pitfalls. This section attempts to describe both.

- Changes
 - \$MYSQLTEST_VARDIR/master-data -> \$MYSQLD_DATADIR
 - This change was introduced to work with parallel. Since MTR2 can run several instances of itself, one path is needed for each instance.
 - You must also use let \$MYSQLD_DATADIR= `SELECT @@datadir`;
- New options
 - max-test-fail: The number of test cases that can fail before the test run aborts. Default=10, set to 0 to continue the run regardless of failure count.

NOTE: --force alone is no longer sufficient to guarantee a full test suite run. If the number of failing tests == max-test-fail, then the test run will abort.
 - parallel: Set n number of workers to run tests in parallel. The current recommendation is 2x the number of processors on your machine.

NOTE: You can no longer be sure of executing test cases in a particular order due to this and other changes brought on by MTR2. *Do not* try to create multi-part tests (test_pt1.test, test_pt2.test, etc).

NOTE: Be cautious when using a high (>2) value for parallel and --mem in combination. It is possible to use vast amounts of resources on your system and cause problematic performance till a crash of the operating system while MTR2 is running.
 - repeat: Run the input test(s) n times in succession. Very good for diagnosing random failures - set a high value for 'n', create a high load on your machine, and see if you can duplicate the failure.

Examples of “Suspicious” Tests and Scripts

- [Probably Mismatch of Focus of Test and Code Sequence](#)

- [Too Greedy Test](#)
- [Risky Handling Around Additional Sessions](#)

Probably Mismatch of Focus of Test and Code Sequence

```
CREATE TABLE t1 (  
  id INT NOT NULL AUTO_INCREMENT,  
  my_column VARCHAR(30),  
  name LONGTEXT,  
  PRIMARY KEY (id));  
  
INSERT INTO t1(my_column,name) VALUES('2','two');  
INSERT INTO t1(my_column,name) VALUES('1','one');  
INSERT INTO t1(my_column,name) VALUES('4','four');  
INSERT INTO t1(my_column,name) VALUES('2','two');  
INSERT INTO t1(my_column,name) VALUES('3','three');
```

This test will

- fail if the feature AUTO_INCREMENT is temporary broken
- will (depending on existence of prerequisite checks) fail or get skipped if the default storage engine does not support AUTO_INCREMENT or LONGTEXT

==> no coverage for other features or properties checked within this test

Some questions with recommended action depending on the answer:

- Is the table t1 just an auxiliary table and not the test object?

Yes: Please ensure that the test does not break or gets skipped if the default storage engine does not support AUTO_INCREMENT or LONGTEXT and you are done.

No: no action
- Do we check AUTO_INCREMENT and the corresponding column is t1.id?

Yes: no action

No: Remove the use of AUTO_INCREMENT
- Do we check LONGTEXT and the corresponding column is t1.name?

Yes: no action

No: Remove the use of LONGTEXT
- Do we check AUTO_INCREMENT and LONGTEXT in combination?

Yes: no action

No: Split the test at least if it should be a test of basic functionality

Too Greedy Test

Let's assume we have to check that every new created table causes a row in information_schema.tables.

```
--replace_column 15 <CREATE_TIME> 16 <UPDATE_TIME> 17 <CHECK_TIME>  
SELECT * FROM information_schema.tables;
```

The SELECT above makes the test extreme "greedy" for changes in behaviour which is maybe good for general bug hunting but not for smart and frequent automatic tests. It is to be expected that such a test requires frequent maintenance like adjustment of expected results.

The problems with the select above: We will (sometimes only maybe) get a result set difference whenever

- we get a new system table

Completeness of result sets without WHERE are outside of the focus.

Refined statement:

```
--replace_column 15 <CREATE_TIME> 16 <UPDATE_TIME> 17 <CHECK_TIME>
SELECT * FROM information_schema.tables
WHERE TABLE_SCHEMA = 'test' AND TABLE_NAME = 't1';
```

- the number of columns within information_schema.tables changes
- the data type of a column within ... changes
- the content within one of the columns within ... changes

Lets assume something intentional like VERSION is now 11 instead of 10.

This means we should try to avoid "SELECT *".

Final statement:

```
Variant 1:
SELECT TABLE_SCHEMA, TABLE_NAME FROM information_schema.tables
WHERE TABLE_SCHEMA = 'test' AND TABLE_NAME = 't1';
```

Are we really consequent? No, but we could

```
Variant 2:
SELECT COUNT(*) FROM information_schema.tables
WHERE TABLE_SCHEMA = 'test' AND TABLE_NAME = 't1';
```

It's on you which variant you choose. I decide depending on the situation.

Risky Handling Around Additional Sessions

Disconnects are performed asynchronous. This is most probably good for the throughput of usual application but has some annoying consequences for the development of tests.

If a

- disconnect <connection> or
- KILL <connection_id> or
- --exec <client program which connects to the server>

occurs within a script for mysqltest than the next statements will be executed even if the disconnect or kill is not complete finished.

I guess this also applies to the end (-> end of statements to be executed by mysqltest) of a test. It triggers some disconnect for any open session but this is also asynchronous.

Not complete finished means that the session is visible within the processlist and the entry about the 'quit' of the session is not already appended to the general log.

So it could happen that a succeeding test suffers by unexpected events like

- observing the session belonging to the previous test within the processlist
- being surprised by an entry (the 'quit' of the disconnected session) within the general log which was not caused by its own activity etc.

According to my experiences the likelihood that some test could harm a succeeding test by unfinished disconnects increases with

- additional parallel load on the testing box
- "disconnect <connection>" is missing within the test
- "shorter" distance between "disconnect <connection>", "KILL <connection_id>" or "--exec ..." and end of test
- the session runs a "heavy" SQL statement just before its disconnect

In my opinion the most robust solution is to add a

```
--source include/wait_until_disconnected.inc
```

just after the disconnect.

In case you don't believe me, run the snippet following later

```
./mtr --skip-ndb --no-check-testcases --repeat=100 <snippet>
```

and you will most probably observe something like

```
TEST                                RESULT    TIME (ms)
-----
...
<snippet>                          [ pass ]    4
<snippet>                          [ fail ]
...
CURRENT_TEST: <snippet>
--- <result>
+++ <reject>
@@ -1,4 +1,4 @@
SELECT COUNT(*) FROM information_schema.processlist WHERE id < CONNECTION_ID();
COUNT(*)
-0
+1
SELECT SLEEP(10);
...
mysqltest: Result content mismatch
```

Snippet for demonstration purposes:

```
SELECT COUNT(*) FROM information_schema.processlist WHERE id < CONNECTION_ID(); # (1)
connect (con1,localhost,root,,);
send SELECT SLEEP(10); # (2)
connection default;
# Wait till the SQL statement of con1 is "in work"
let $wait_condition= SELECT COUNT(*) = 1 FROM information_schema.processlist
WHERE State = 'User sleep' AND Info = 'SELECT SLEEP(10)';
--source include/wait_condition.inc # (3)
SELECT COUNT(*) = 1 FROM information_schema.processlist
WHERE State = 'User sleep' AND Info = 'SELECT SLEEP(10)'; # (4)
disconnect con1; (5)
<end of script>
```

The snippet is intended for demonstration purposes and contains code which is obvious "crap" but needed to enforce the intended effect. Some notes:

- (1) Such a statement at the begin of a test makes this test very sensitive to not finished disconnects.
- (2) The wait routine (3) prevents that our look on the process list (4) happens too early = before the server has started the processing of (2). Otherwise we get random result set differences for (4) under high parallel load.
- How to prevent that this test harms the succeeding test?

```
(5) should be replaced by  
connection con1;  
disconnect con1;  
--source include/wait_until_disconnected.inc  
connection default;
```

- Where is the "crap"?
 - There should be a 'reap' for every preceeding 'send <statement>'. This is not just for academic completeness it also ensures here that (2) is really finished.
 - Do we really need 10 seconds within the SLEEP? According to my experiences sleep times ≤ 2 seconds are critical under high load. What about SLEEP(5)?

Chapter 28. Porting to Other Systems

Table of Contents

Debugging a MySQL Server	352
Compiling MySQL for Debugging	352
Creating Trace Files	353
Debugging <code>mysqld</code> under <code>gdb</code>	354
Using a Stack Trace	355
Using Server Logs to Find Causes of Errors in <code>mysqld</code>	356
Making a Test Case If You Experience Table Corruption	356
Debugging a MySQL Client	357
The DBUG Package	357
Comments about RTS Threads	359
Differences Between Thread Packages	360

This appendix helps you port MySQL to other operating systems. Do check the list of currently supported operating systems first. See [Operating Systems Supported by MySQL Community Server](#). If you have created a new port of MySQL, please let us know so that we can list it here and on our Web site (<http://www.mysql.com/>), recommending it to other users.

Note: If you create a new port of MySQL, you are free to copy and distribute it under the GPL license, but it does not make you a copyright holder of MySQL.

A working POSIX thread library is needed for the server. On Solaris 2.5 we use Sun PThreads (the native thread support in 2.4 and earlier versions is not good enough), on Linux we use LinuxThreads by Xavier Leroy, <Xavier.Leroy@inria.fr>.

The hard part of porting to a new Unix variant without good native thread support is probably to port MIT-pthreads. See [mit-pthreads/README](#) and Programming POSIX Threads (<http://www.humanfactor.com/pthreads/>).

Up to MySQL 4.0.2, the MySQL distribution included a patched version of Chris Provenzano's Pthreads from MIT (see the MIT Pthreads Web page at <http://www.mit.edu/afs/sipb/project/pthreads/> and a programming introduction at http://www.mit.edu:8001/people/proven/IAP_2000/). These can be used for some operating systems that do not have POSIX threads. See [MIT-pthreads Notes](#).

It is also possible to use another user level thread package named FSU Pthreads (see <http://moss.csc.ncsu.edu/~mueller/pthreads/>). This implementation is being used for the SCO port.

See the `thr_lock.c` and `thr_alarm.c` programs in the `mysys` directory for some tests/examples of these problems.

Both the server and the client need a working C++ compiler. We use `gcc` on many platforms. Other compilers that are known to work are SPARCworks, Sun Forte, Irix `cc`, HP-UX `aCC`, IBM AIX `xlC_r`, Intel `ecc/icc` and Compaq `cxx`).

Important: If you are trying to build MySQL 5.1 with `icc` on the IA64 platform, and need support for MySQL Cluster, you should first ensure that you are using `icc` version 9.1.043 or later. (For details, see [\[http://bugs.mysql.com/21875 Bug#21875\]](http://bugs.mysql.com/21875).)

To compile only the client use `./configure --without-server`.

There is currently no support for only compiling the server, nor is it likely to be added unless someone has a good reason for it.

If you want/need to change any `Makefile` or the configure script you also need GNU Automake and Autoconf. See [Installing from the Development Source Tree](#).

All steps needed to remake everything from the most basic files.

```
/bin/rm */.deps/*.P
/bin/rm -f config.cache
aclocal
autoheader
aclocal
automake
autoconf
./configure --with-debug=full --prefix='your installation directory'

# The makefiles generated above need GNU make 3.75 or newer.
# (called gmake below)
gmake clean all install init-db
```

If you run into problems with a new port, you may have to do some debugging of MySQL! See [Debugging a MySQL Server](#).

Note: Before you start debugging `mysqld`, first get the test programs `mysys/thr_alarm` and `mysys/thr_lock` to work. This ensures that your thread installation has even a remote chance to work!

Debugging a MySQL Server

If you are using some functionality that is very new in MySQL, you can try to run `mysqld` with the `--skip-new` (which disables all new, potentially unsafe functionality) or with `--safe-mode` which disables a lot of optimization that may cause problems. See [What to Do If MySQL Keeps Crashing](#).

If `mysqld` doesn't want to start, you should verify that you don't have any `my.cnf` files that interfere with your setup! You can check your `my.cnf` arguments with `mysqld --print-defaults` and avoid using them by starting with `mysqld --no-defaults`

If `mysqld` starts to eat up CPU or memory or if it “hangs,” you can use `mysqladmin processlist status` to find out if someone is executing a query that takes a long time. It may be a good idea to run `mysqladmin -i10 processlist status` in some window if you are experiencing performance problems or problems when new clients can't connect.

The command `mysqladmin debug` dumps some information about locks in use, used memory and query usage to the MySQL log file. This may help solve some problems. This command also provides some useful information even if you haven't compiled MySQL for debugging! dumps some information about locks in use, used memory and query usage to the MySQL log file. This may help solve some problems. This command also provides some useful information even if you haven't compiled MySQL for debugging!

If the problem is that some tables are getting slower and slower you should try to optimize the table with `OPTIMIZE TABLE` or `myisamchk`. See <http://dev.mysql.com/doc/en/database-administration>. You should also check the slow queries with `EXPLAIN`.

You should also read the OS-specific section in this manual for problems that may be unique to your environment. See [Operating System-Specific Notes](#).

Compiling MySQL for Debugging

If you have some very specific problem, you can always try to debug MySQL. To do this you must configure MySQL with the `--with-debug` or the `--with-debug=full` option. You can check whether MySQL was compiled with debugging by doing: `mysqld --help`. If the `--debug` flag is listed with the options then you have debugging enabled. `mysqladmin ver` also lists the `mysqld` version as `mysql ... --debug` in this case.

If you are using `gcc` or `egcs`, the recommended `configure` line is:

```
CC=gcc CFLAGS="-O2" CXX=gcc CXXFLAGS="-O2 -felide-constructors \
-fno-exceptions -fno-rtti" ./configure --prefix=/usr/local/mysql \
```

```
--with-debug --with-extra-charsets=complex
```

This avoids problems with the `libstdc++` library and with C++ exceptions (many compilers have problems with C++ exceptions in threaded code) and compile a MySQL version with support for all character sets.

If you suspect a memory overrun error, you can configure MySQL with `--with-debug=full`, which installs a memory allocation (`SAFEMALLOC`) checker. However, running with `SAFEMALLOC` is quite slow, so if you get performance problems you should start `mysqld` with the `--skip-safemalloc` option. This disables the memory overrun checks for each call to `malloc()` and `free()`. Note: `SAFEMALLOC` was removed in MySQL 5.5.6 because it causes excessive overhead in certain situations.

If `mysqld` stops crashing when you compile it with `--with-debug`, you probably have found a compiler bug or a timing bug within MySQL. In this case, you can try to add `-g` to the `CFLAGS` and `CXXFLAGS` variables above and not use `--with-debug`. If `mysqld` dies, you can at least attach to it with `gdb` or use `gdb` on the core file to find out what happened.

When you configure MySQL for debugging you automatically enable a lot of extra safety check functions that monitor the health of `mysqld`. If they find something unexpected, an entry is written to `stderr`, which `mysqld_safe` directs to the error log! This also means that if you are having some unexpected problems with MySQL and are using a source distribution, the first thing you should do is to configure MySQL for debugging! (The second thing is to send mail to a MySQL mailing list and ask for help. See [MySQL Mailing Lists](#). If you believe that you have found a bug, please use the instructions at [How to Report Bugs or Problems](#).)

In the Windows MySQL distribution, `mysqld.exe` is by default compiled with support for trace files.

Creating Trace Files

If the `mysqld` server doesn't start or if you can cause it to crash quickly, you can try to create a trace file to find the problem.

To do this, you must have a `mysqld` that has been compiled with debugging support. You can check this by executing `mysqld -V`. If the version number ends with `-debug`, it's compiled with support for trace files. (On Windows, the debugging server is named `mysqld-debug` rather than `mysqld` as of MySQL 4.1.)

Start the `mysqld` server with a trace log in `/tmp/mysqld.trace` on Unix or `C:\mysqld.trace` on Windows:

```
shell> mysqld --debug
```

On Windows, you should also use the `--standalone` flag to not start `mysqld` as a service. In a console window, use this command:

```
C:\> mysqld-debug --debug --standalone
```

After this, you can use the `mysql.exe` command-line tool in a second console window to reproduce the problem. You can stop the `mysqld` server with `mysqladmin shutdown`.

Note that the trace file become **very big**! If you want to generate a smaller trace file, you can use debugging options something like this:

```
mysqld --debug=d,info,error,query,general,where:0,/tmp/mysqld.trace
```

This only prints information with the most interesting tags to the trace file.

If you make a bug report about this, please only send the lines from the trace file to the appropriate mailing list where something seems to go wrong!

The trace file is made with the [DEBUG](#) package by Fred Fish. See [The DEBUG Package](#).

Debugging `mysqld` under `gdb`

On most systems you can also start `mysqld mysqld` from `gdb` to get more information if `mysqld` crashes.

With some older `gdb` versions on Linux you must use `run --one-thread` if you want to be able to debug `mysqld` threads. In this case, you can only have one thread active at a time. We recommend you to upgrade to `gdb` 5.1 ASAP as thread debugging works much better with this version!

NPTL threads (the new thread library on Linux) may cause problems while running `mysqld` under `gdb`. Some symptoms are:

- `mysqld` hangs during startup (before it writes `ready for connections`).
- `mysqld` crashes during a `pthread_mutex_lock()` or `pthread_mutex_unlock()` call.

In this case, you should set the following environment variable in the shell before starting `gdb`:

```
LD_ASSUME_KERNEL=2.4.1
export LD_ASSUME_KERNEL
```

When running `mysqld` under `gdb`, you should disable the stack trace with `--skip-stack-trace` to be able to catch segfaults within `gdb`.

In MySQL 4.0.14 and above you should use the `--gdb` option to `mysqld`. This installs an interrupt handler for `SIGINT` (needed to stop `mysqld` with `^C` to set breakpoints) and disable stack tracing and core file handling.

It's very hard to debug MySQL under `gdb` if you do a lot of new connections the whole time as `gdb` doesn't free the memory for old threads. You can avoid this problem by starting `mysqld` with `--thread_cache_size='max_connections+1'`. In most cases just using `--thread_cache_size=5` helps a lot!

If you want to get a core dump on Linux if `mysqld` dies with a `SIGSEGV` signal, you can start `mysqld` with the `--core-file` option. This core file can be used to make a backtrace that may help you find out why `mysqld` died:

```
shell> gdb mysqld core
gdb> backtrace full
gdb> quit
```

See [What to Do If MySQL Keeps Crashing](#).

If you are using `gdb` 4.17.x or above on Linux, you should install a `.gdb` file, with the following information, in your current directory:

```
set print sevenbit off
handle SIGUSR1 nostop noprint
handle SIGUSR2 nostop noprint
handle SIGWAITING nostop noprint
handle SIGLWP nostop noprint
handle SIGPIPE nostop
handle SIGALRM nostop
handle SIGHUP nostop
handle SIGTERM nostop noprint
```

If you have problems debugging threads with `gdb`, you should download `gdb` 5.x and try this instead. The new `gdb` version has very improved thread handling!

Here is an example how to debug `mysqld`:

```
shell> gdb /usr/local/libexec/mysqld
```



```
gdb> run
...
backtrace full # Do this when mysqld crashes
```

Include the above output in a bug report, which you can file using the instructions in [How to Report Bugs or Problems](#).

If `mysqld` hangs you can try to use some system tools like `strace` or `/usr/proc/bin/pstack` to examine where `mysqld` has hung.

```
strace /tmp/log libexec/mysqld
```

If you are using the Perl `DBI` interface, you can turn on debugging information by using the `trace` method or by setting the `DBI_TRACE` environment variable.

Using a Stack Trace

On some operating systems, the error log contains a stack trace if `mysqld` dies unexpectedly. You can use this to find out where (and maybe why) `mysqld` died. See <http://dev.mysql.com/doc/en/error-log>. To get a stack trace, you must not compile `mysqld` with the `-fomit-frame-pointer` option to `gcc`. See [Compiling MySQL for Debugging](#).

If the error file contains something like the following:

```
mysqld got signal 11;
The manual section 'Debugging a MySQL server' tells you how to
use a stack trace and/or the core file to produce a readable
backtrace that may help in finding out why mysqld died
Attempting backtrace. You can use the following information
to find out where mysqld died.  If you see no messages after
this, something went terribly wrong...
stack range sanity check, ok, backtrace follows
0x40077552
0x81281a0
0x8128f47
0x8127be0
0x8127995
0x8104947
0x80ff28f
0x810131b
0x80ee4bc
0x80c3c91
0x80c6b43
0x80c1fd9
0x80c1686
```

you can find where `mysqld` died by doing the following:

1. Copy the preceding numbers to a file, for example `mysqld.stack`.
2. Make a symbol file for the `mysqld` server:

```
nm -n libexec/mysqld > /tmp/mysqld.sym
```

Note that most MySQL binary distributions (except for the "debug" packages, where this information is included inside of the binaries themselves) ship with the above file, named `mysqld.sym.gz`. In this case, you can simply unpack it by doing:

```
gunzip < bin/mysqld.sym.gz > /tmp/mysqld.sym
```

- a. Execute `resolve_stack_dump -s /tmp/mysqld.sym -n mysqld.stack`.

This prints out where `mysqld` and preferable a test case so that we can repeat the problem! See `mysqld` died, you should make a bug report and include the output from the above command with the bug report. Note however that in most cases it does not help us to just have a stack trace to find

the reason for the problem. To be able to locate the bug or provide a workaround, we would in most cases need to know the query that killed `mysqld` and preferably a test case so that we can repeat the problem! See [How to Report Bugs or Problems](#).

Using Server Logs to Find Causes of Errors in `mysqld`

Note that before starting `mysqld` with `--log` you should check all your tables with `myisamchk`. See [Database Administration](#).

If `mysqld` dies or hangs, you should start `mysqld` with `--log`. When `mysqld` dies again, you can examine the end of the log file for the query that killed `mysqld`.

If you are using `--log` without a file name, the log is stored in the database directory as `host_name.log`. In most cases it is the last query in the log file that killed `mysqld`, but if possible you should verify this by restarting `mysqld` and executing the found query from the `mysql` command-line tools. If this works, you should also test all complicated queries that didn't complete.

You can also try the command `EXPLAIN` on all `SELECT` statements that takes a long time to ensure that `mysqld` is using indexes properly. See [Optimizing Queries with EXPLAIN](#).

You can find the queries that take a long time to execute by starting `mysqld` with `--log-slow-queries`. See <http://dev.mysql.com/doc/en/slow-query-log>.

If you find the text `mysqld restarted` in the error log file (normally named `hostname.err`) you probably have found a query that causes `mysqld` to fail. If this happens, you should check all your tables with `myisamchk` (see [Database Administration](#)), and test the queries in the MySQL log files to see whether one fails. If you find such a query, try first upgrading to the newest MySQL version. If this doesn't help and you can't find anything in the `mysql` mail archive, you should report the bug to a MySQL mailing list. The mailing lists are described at <http://lists.mysql.com/>, which also has links to online list archives.

If you have started `mysqld` with `myisam-recover`, MySQL automatically checks and tries to repair `MyISAM` tables if they are marked as 'not closed properly' or 'crashed'. If this happens, MySQL writes an entry in the `hostname.err` file 'Warning: Checking table ...' which is followed by `Warning: Repairing table` if the table needs to be repaired. If you get a lot of these errors, without `mysqld` having died unexpectedly just before, then something is wrong and needs to be investigated further. See [Command Options](#).

It is not a good sign if `mysqld` did die unexpectedly, but in this case, you should not investigate the `Checking table...` messages, but instead try to find out why `mysqld` died.

Making a Test Case If You Experience Table Corruption

If you get corrupted tables or if `mysqld` always fails after some update commands, you can test whether this bug is reproducible by doing the following:

- Take down the MySQL daemon (with `mysqladmin shutdown`).
- Make a backup of the tables (to guard against the very unlikely case that the repair does something bad).
- Check all tables with `myisamchk -s database/*.MYI`. Repair any wrong tables with `myisamchk -r database/table.MYI`.
- Make a second backup of the tables.
- Remove (or move away) any old log files from the MySQL data directory if you need more space.
- Start `mysqld` with `--log-bin`. See <http://dev.mysql.com/doc/en/binary-log>. If you want to find a query that crashes `mysqld`, you should use `--log --log-bin`.

- When you have gotten a crashed table, stop the `mysqld` server.
- Restore the backup.
- Restart the `mysqld` server **without** `--log-bin`
- Re-execute the commands with `mysqlbinlog update-log-file | mysql`. The update log is saved in the MySQL database directory with the name `hostname-bin.#`.
- If the tables are corrupted again or you can get `mysqld` to die with the above command, you have found reproducible bug that should be easy to fix!

You can also use the script `mysql_find_rows` to just execute some of the update statements if you want to narrow down the problem.

Debugging a MySQL Client

To be able to debug a MySQL client with the integrated debug package, you should configure MySQL with `--with-debug` or `--with-debug=full`. See <http://dev.mysql.com/doc/en/configure-options>.

Before running a client, you should set the `MYSQL_DEBUG` environment variable:

```
shell> MYSQL_DEBUG=d:t:O,/tmp/client.trace
shell> export MYSQL_DEBUG
```

This causes clients to generate a trace file in `/tmp/client.trace`.

If you have problems with your own client code, you should attempt to connect to the server and run your query using a client that is known to work. Do this by running `mysql` in debugging mode (assuming that you have compiled MySQL with debugging on):

```
shell> mysql --debug=d:t:O,/tmp/client.trace
```

This provides useful information in case you mail a bug report. See [How to Report Bugs or Problems](#).

If your client crashes at some 'legal' looking code, you should check that your `mysql.h` include file matches your MySQL library file. A very common mistake is to use an old `mysql.h` file from an old MySQL installation with new MySQL library.

The DBUG Package

The MySQL server and most MySQL clients are compiled with the DBUG package originally created by Fred Fish. When you have configured MySQL for debugging, this package makes it possible to get a trace file of what the program is debugging. See [Creating Trace Files](#).

This section summaries the argument values that you can specify in debug options on the command line for MySQL programs that have been built with debugging support. For more information about programming with the DBUG package, see the DBUG manual in the `debug` directory of MySQL source distributions. It's best to use a recent distribution to get the most updated DBUG manual.

You use the debug package by invoking a program with the `--debug="..."` or the `-#...` option.

Most MySQL programs have a default debug string that is used if you don't specify an option to `--debug`. The default trace file is usually `/tmp/program_name.trace` on Unix and `\program_name.trace` on Windows.

The debug control string is a sequence of colon-separated fields as follows:

```
<field_1>:<field_2>:...:<field_N>
```

Each field consists of a mandatory flag character followed by an optional `' , '` and comma-separated list of modifiers:

```
flag[,modifier,modifier,...,modifier]
```

The currently recognized flag characters are:

Flag	Description
<code>d</code>	Enable output from <code>DEBUG_<N></code> macros for the current state. May be followed by a list of keywords which selects output only for the DEBUG macros with that keyword. An empty list of keywords implies output for all macros.
<code>D</code>	Delay after each debugger output line. The argument is the number of tenths of seconds to delay, subject to machine capabilities. For example, <code>-#D,20</code> specifies a delay of two seconds.
<code>f</code>	Limit debugging, tracing, and profiling to the list of named functions. Note that a null list disables all functions. The appropriate <code>d</code> or <code>t</code> flags must still be given; this flag only limits their actions if they are enabled.
<code>F</code>	Identify the source file name for each line of debug or trace output.
<code>i</code>	Identify the process with the PID or thread ID for each line of debug or trace output.
<code>g</code>	Enable profiling. Create a file called <code>dbugmon.out</code> containing information that can be used to profile the program. May be followed by a list of keywords that select profiling only for the functions in that list. A null list implies that all functions are considered.
<code>L</code>	Identify the source file line number for each line of debug or trace output.
<code>n</code>	Print the current function nesting depth for each line of debug or trace output.
<code>N</code>	Number each line of debug output.
<code>o</code>	Redirect the debugger output stream to the specified file. The default output is <code>stderr</code> .
<code>O</code>	Like <code>o</code> , but the file is really flushed between each write. When needed, the file is closed and reopened between each write.
<code>p</code>	Limit debugger actions to specified processes. A process must be identified with the <code>DEBUG_PROCESS</code> macro and match one in the list for debugger actions to occur.
<code>P</code>	Print the current process name for each line of debug or trace output.
<code>r</code>	When pushing a new state, do not inherit the previous state's function nesting level. Useful when the output is to start at the left margin.
<code>S</code>	Do function <code>_sanity(_file_,_line_)</code> at each debugged function until <code>_sanity()</code> returns something that differs from 0. (Mostly used with <code>safemalloc</code> to find memory leaks)

`t`

Enable function call/exit trace lines. May be followed by a list (containing only one modifier) giving a numeric maximum trace level, beyond which no output occurs for either debugging or tracing macros. The default is a compile time option.

Some examples of debug control strings that might appear on a shell command line (the `-#` is typically used to introduce a control string to an application program) are:

```
-#d:t
-#d:f,main,subr1:F:L:t,20
-#d,input,output,files:n
-#d:t:i:O,\\mysqld.trace
```

In MySQL, common tags to print (with the `d` option) are `enter`, `exit`, `error`, `warning`, `info`, and `loop`.

Comments about RTS Threads

I have tried to use the RTS thread packages with MySQL but stumbled on the following problems:

They use old versions of many POSIX calls and it is very tedious to make wrappers for all functions. I am inclined to think that it would be easier to change the thread libraries to the newest POSIX specification.

Some wrappers are currently written. See [mysys/my_pthread.c](#) for more info.

At least the following should be changed:

`pthread_get_specific` should use one argument. `sigwait` should take two arguments. A lot of functions (at least `pthread_cond_wait`, `pthread_cond_timedwait()`) should return the error code on error. Now they return -1 and set `errno`.

Another problem is that user-level threads use the `ALRM` signal and this aborts a lot of functions (`read`, `write`, `open`...). MySQL should do a retry on interrupt on all of these but it is not that easy to verify it.

The biggest unsolved problem is the following:

To get thread-level alarms I changed [mysys/thr_alarm.c](#) to wait between alarms with `pthread_cond_timedwait()`, but this aborts with error `EINTR`. I tried to debug the thread library as to why this happens, but couldn't find any easy solution.

If someone wants to try MySQL with RTS threads I suggest the following:

- Change functions MySQL uses from the thread library to POSIX. This shouldn't take that long.
- Compile all libraries with the `-DHAVE_rts_threads`.
- Compile [thr_alarm](#).
- If there are some small differences in the implementation, they may be fixed by changing [my_pthread.h](#) and [my_pthread.c](#).
- Run [thr_alarm](#). If it runs without any warning, error, or aborted messages, you are on the right track. Here is a successful run on Solaris:

```
Main thread: 1
Thread 0 (5) started
Thread: 5 Waiting
process_alarm
Thread 1 (6) started
Thread: 6 Waiting
process_alarm
```

```
process_alarm
thread_alarm
Thread: 6 Slept for 1 (1) sec
Thread: 6 Waiting
process_alarm
process_alarm
thread_alarm
Thread: 6 Slept for 2 (2) sec
Thread: 6 Simulation of no alarm needed
Thread: 6 Slept for 0 (3) sec
Thread: 6 Waiting
process_alarm
process_alarm
thread_alarm
Thread: 6 Slept for 4 (4) sec
Thread: 6 Waiting
process_alarm
thread_alarm
Thread: 5 Slept for 10 (10) sec
Thread: 5 Waiting
process_alarm
process_alarm
thread_alarm
Thread: 6 Slept for 5 (5) sec
Thread: 6 Waiting
process_alarm
process_alarm

...
thread_alarm
Thread: 5 Slept for 0 (1) sec
end
```

Differences Between Thread Packages

MySQL is very dependent on the thread package used. So when choosing a good platform for MySQL, the thread package is very important.

There are at least three types of thread packages:

- User threads in a single process. Thread switching is managed with alarms and the threads library manages all non-thread-safe functions with locks. Read, write and select operations are usually managed with a thread-specific select that switches to another thread if the running threads have to wait for data. If the user thread packages are integrated in the standard libs (FreeBSD and BSDI threads) the thread package requires less overhead than thread packages that have to map all unsafe calls (MIT-pthreads, FSU Pthreads and RTS threads). In some environments (for example, SCO), all system calls are thread-safe so the mapping can be done very easily (FSU Pthreads on SCO). Downside: All mapped calls take a little time and it's quite tricky to be able to handle all situations. There are usually also some system calls that are not handled by the thread package (like MIT-pthreads and sockets). Thread scheduling isn't always optimal.
- User threads in separate processes. Thread switching is done by the kernel and all data are shared between threads. The thread package manages the standard thread calls to allow sharing data between threads. LinuxThreads is using this method. Downside: Lots of processes. Thread creating is slow. If one thread dies the rest are usually left hanging and you must kill them all before restarting. Thread switching is somewhat expensive.
- Kernel threads. Thread switching is handled by the thread library or the kernel and is very fast. Everything is done in one process, but on some systems, [ps](#) may show the different threads. If one thread aborts, the whole process aborts. Most system calls are thread-safe and should require very little overhead. Solaris, HP-UX, AIX and OSF/1 have kernel threads.

In some systems kernel threads are managed by integrating user level threads in the system libraries. In such cases, the thread switching can only be done by the thread library and the kernel isn't really thread aware.

Chapter 29. Error Messages

Table of Contents

Adding New Error Messages to MySQL	361
Adding an Error Message to Multiple MySQL Versions	362
Adding Error Messages to MySQL 5.0.3 or Higher	363
Adding Error Messages to Old (< MySQL 5.0) Versions	365
Adding Storage Engine Error Messages	366

This chapter describes how error messages are defined and how to add the capability of generating error messages to a table handler.

Adding New Error Messages to MySQL

The procedure for adding error messages depends on which version of MySQL you are using:

- Before MySQL 5.0.3, error messages are stored in `errmsg.txt` files in the language directories under `sql/share`. That is, the files have names like `czech/errmsg.txt`, `danish/errmsg.txt`, and so forth, and each one is language-specific. Each of these language-specific files must contain a line for each error message, so adding a new message involves adding a line to the `errmsg.txt` file for every language. The procedure involves adding the English message to the `english/errmsg.txt` file and running a script that adds the message to the other language-specific files. Translators may translate the message in other `errmsg.txt` files later.
- Beginning with MySQL 5.0.3, error messages are stored in a single `errmsg.txt` file in the `sql/share` directory, and it contains the error messages for all languages. The messages are grouped by error symbol. For each symbol, there must be an English message, and messages can be present for other languages as well. If there is no message for a given language, the English version is used. Messages for a given language are written in the character set used for that language, so the file contains information in several character sets.
- From MySQL 5.5 on, error messages are stored in a single `errmsg-utf8.txt` file in the `sql/share` directory. The file format is similar to `errmsg.txt`, but the contents of the file are written in a single character set, utf8. An added feature is that error messages may contain positional arguments.

The **comp_err** program compiles the text error message file or files into language-specific `errmsg.sys` files that each are located in the appropriate language directory under `sql/share`:

- Before MySQL 5.5, **comp_err** reads `errmsg.txt` and writes `errmsg.sys` files, each in the character set associated with the language for the file. For example, `sql/share/korean/errmsg.sys` is written using `euckr`.
- From MySQL 5.5 on, **comp_err** reads `errmsg-utf8.txt` and writes `errmsg.sys` files using `utf8`.

comp_err also generates a number of header files in the `include` directory. The MySQL build process runs **comp_err** automatically.

Note: You should observe some general considerations regarding error messages that apply no matter your version of MySQL:

- **Always** use parameter constructs such as `"%.64s"` to guard against buffer overflows. The maximum error message length is `MYSQL_ERRMSG_SIZE`.

- **Never** add new parameters (such as %s) to existing error messages. Error messages must always be backward compatible. If a parameter were added, older servers would crash when run with a newer error message file.
- If you need to add new parameters to an existing error message, you can introduce a new error message but continue to use the old error code. For example, in MySQL 5.1, the `handler::print_keydup_error()` method is passed a message like `ER(ER_DUP_ENTRY_WITH_KEY_NAME)`, but always uses the `ER_DUP_ENTRY` error code.

Adding an Error Message to Multiple MySQL Versions

It is critical that error codes are identical in all versions. That is, the value of `ER_SOME_ERROR` must be the same in all versions for which it is defined. The following procedure follows from this requirement. (The discussion uses the name `errmsg.txt`, but when you cross the boundary from before MySQL 5.5 to 5.5 or higher, you should use `errmsg-utf8.txt` instead.)

Let *GA* reference the most recent stable version. For example, if MySQL 5.1 is GA, and MySQL 5.2 is -beta, then *GA* refers to MySQL 5.1.

To add a new error message in multiple versions, first add the specific error message to the *GA* version, at the end of `errmsg.txt`.

Next, add the specific error message to all versions > *GA*, at the **same position** in `errmsg.txt`, **not** at the end of the file. This will ensure that the error code value is the same in *GA* and all later versions. This will have the side effect that any error codes which are not in *GA* will change their values in post-*GA* versions. This is acceptable, in non-stable versions. We consider pre-RC versions non-stable in this context.

Finally, for any version < *GA*, **do not add** the specific error message. Instead, use the `ER_UNKNOWN_ERROR` error code, and print a helpful error text with code similar to this:

```
my_printf_error(ER_UNKNOWN_ERROR,
               "Some error text here, with the '%.64s' parameter value"
               MYF(0), a_parameter);
```

This will provide the user with useful information, while ensuring that all versions have consistent error codes.

If you need to merge error messages up from 5.1 to 5.5, the merge operation will update the 5.5 `errmsg.txt` file with information that then needs to be added to the `errmsg-utf8.txt` file. To handle this, you can use the `errmsg-cnv.sh` script, which converts `errmsg.txt` to `errmsg-utf8.txt`. However, you must be careful not to wipe out the existing 5.5-specific information in `errmsg-utf8.txt`.

Suppose that you added a new error message into 5.1 and have merged this change into 5.5. You can do the following:

1. Make a backup of `errmsg-utf8.txt`:

```
cp errmsg-utf8.txt errmsg-utf8.txt.sav
```

2. Convert 5.1 messages:

```
sh errmsg-cnv.sh
```

Note: This will add 5.1 messages into `errmsg-utf8.txt`, but remove 5.5 messages.

3. Open `errmsg-utf8.txt` and `errmsg-utf8.txt.sav` in a text editor, then cut and paste 5.5-specific messages into `errmsg-utf8.txt`.

4. Make sure everything went fine:


```
bzr diff errmsg-utf8.txt
```

Adding Error Messages to MySQL 5.0.3 or Higher

The file you edit to add a new error message depends on your version of MySQL:

- MySQL 5.0.3 up to 5.5: Edit [errmsg.txt](#)
- MySQL 5.5 and up: Edit [errmsg-utf8.txt](#)

In either case, **comp_err** generates the header files automatically during the MySQL build process.

The [errmsg.txt](#) (or [errmsg-utf8.txt](#)) file begins with some lines that define general characteristics of error messages, followed by sections for particular messages. The following example shows a partial listing of the file. (The [languages](#) line is wrapped here; it must be given all on one line.)

```
languages czech=cze latin2, danish=dan latin1, dutch=nla latin1,
english=eng latin1, estonian=est latin7, french=fre latin1, german=ger
latin1, greek=greek greek, hungarian=hun latin2, italian=ita latin1,
japanese=jpn ujis, japanese-sjis=jps sjis, korean=kor euckr,
norwegian-ny=norwegian-ny latin1, norwegian=nor latin1, polish=pol
latin2, portuguese=por latin1, romanian=rur latin2, russian=rus
koi8r, serbian=serbian cp1250, slovak=slo latin2, spanish=spa latin1,
swedish=swe latin1, ukrainian=ukr koi8u;

default-language eng

start-error-number 1000

ER_HASHCHK
    eng "hashchk"
ER_NISAMCHK
    eng "isamchk"
ER_NO
    cze "NE"
    dan "NEJ"
    nla "NEE"
    eng "NO"
    est "EI"
    ...
```

A line beginning with a '#' character is taken as a comment. Comments and blank lines are ignored.

Indentation is significant. Unless otherwise specified, leading whitespace should not be used.

The grammar of the [errmsg.txt](#) file looks like this:

```
languages langspec [, langspec] ... ;

start-error-number number

default-language langcode

error-message-section
error-message-section
...
```

The [languages](#) line lists the languages for which language-specific [errmsg.sys](#) files should be generated. A language specification [langspec](#) in the [languages](#) line has this syntax:

```
langspec: langname=langcode langcharset
```

langname is the long language name, *langcode* is the short language code, and *langcharset* is the character set to use for error messages in the language. *langcharset* is irrelevant as of MySQL 5.5 because all [errmsg.sys](#) files are written using [utf8](#).

The `default-language` line specifies the short language code for the default language. (If there is no translation into a given language for a given error message, the message from the default language will be used.)

The `start-error-number` line indicates the number to be assigned to the first error message. Messages that follow the first one are numbered consecutively from this value.

Each *error-message-section* begins with a line that lists an error (or warning) symbol, optionally followed by one or two SQLSTATE values. The error symbol must begin with `ER_` for an error or `WARN_` for a warning. Lines following the error symbol line provide language-specific error messages that correspond to the error symbol. Each message line consists of a tab, a short language code, a space, and the text of the error message within double quote (") characters. Presumably, there must be a message in the default language. There may be message translations for other languages. Order of message lines within a section does not matter. If no translation is given for a given language, the default language message will be used. The following example defines several language translations for the `ER_BAD_FIELD_ERROR` symbol:

```
ER_BAD_FIELD_ERROR 42S22 S0022
    dan "Ukendt kolonne '%-.64s' i tabel %s"
    nla "Onbekende kolom '%-.64s' in %s"
    eng "Unknown column '%-.64s' in '%-.64s'"
    est "Tundmatu tulp '%-.64s' '%-.64s'-s"
    fre "Champ '%-.64s' inconnu dans %s"
    ger "Unbekanntes Tabellenfeld '%-.64s' in %-.64s"
```

In the preceding example, two SQLSTATE values are given following the error symbol (`42S22`, `S0022`). Internally (in `sql/sql_state.c`), these are known as `odbc_state` and `jdbc_state`. Currently, only the first appears ever to be used.

In `errmsg-utf8.txt`, the entire file is written in `utf8`. As long as your editor can handle `utf8`, there should be no problem editing the file.

In `errmsg.txt`, message strings for a given language must be written in the character set indicated for that language in the `languages` line. For example, the language information for Japanese in that line is `japanese=jpn ujis`, so messages with a language code of `jpn` must be written in the `ujis` character set. You might need to be careful about the editor you use for editing the `errmsg.txt` file. For example, there is a report that using **Emacs** will mangle the file, whereas **vi** will not.

Within a message string, C-style escape sequences are allowed:

```
\\  \
\"  "
\n  newline
\N  N, where N is an octal number
\X  X, for any other X
```

In MySQL 5.5 and up, error messages can contain positional constructs for arguments. This is convenient when arguments are most naturally specified in different orders in different languages. Positional arguments enable an error-message writer to avoid awkward language or unnecessarily long messages that result from having the arguments in the same order in all languages.

- To specify a positional argument, include a single digit 0 to 9 and a dollar sign in the formatting specifier for each argument in the error message: `"%1$.32s %2$.64s"` includes the first argument, then the second, in that order. `"%2$.64s %1$.32s"` includes the second argument, then the first.
- If any argument is positional, all must be positional. This is illegal: `"%1$.32s %.64s"`
- Duplicates are allowed. This is legal even though it includes the first argument twice: `"%1$.32s %2$.64s %1$.32s"`
- Gaps in positional number are not allowed. This is illegal because `%2` is missing: `"%1$.32s %3$.64s"`

Use the following procedure to add new error messages:

1. To add a new language translation for an existing error message, find the section for the appropriate error symbol. Then add a new message line to the section. For example:

Before:

```
ER_UNKNOWN_COLLATION
    eng "Unknown collation: '%-.64s'"
    ger "Unbekannte Kollation: '%-.64s'"
    por "Collation desconhecida: '%-.64s'"
```

After (with a new Spanish translation):

```
ER_UNKNOWN_COLLATION
    eng "Unknown collation: '%-.64s'"
    ger "Unbekannte Kollation: '%-.64s'"
    por "Collation desconhecida: '%-.64s'"
    spa "Collation desconocida: '%-.64s'"
```

1. To add an entirely new error message, go to the end of the `errmsg.txt` file. Add a new error symbol line, followed by a message line for the default language, and message lines for any translations that you can supply.
2. Make a full build (`configure + make`). A `make all` is insufficient to build the `sql/share/*/errmsg.sys` files.

`comp_err` will generate the `errmsg.sys` files, as well as the header files `mysqld_error.h`, `mysqld_ename.h`, and `sql_state.h` in the `include` directory.

Be aware that if you make a mistake editing a message text file, `comp_err` prints a cryptic error message and gives you no other feedback. For example, it does not print the input line number where it found a problem. It's up to you to figure this out and correct the file. Perhaps that is not a serious difficulty: `errmsg.txt` (or `errmsg-utf8.txt`) tends to grow by gradual accretion, so if an error occurs when `comp_err` processes it, the problem is likely due to whatever change you just made.

Adding Error Messages to Old (< MySQL 5.0) Versions

Note: This section is included for historical purposes. In practice, no new error messages should be added to error message files for versions of MySQL prior to 5.0.3. Doing so and merging the messages upward would cause error numbers to change in more recent versions.

Before MySQL 5.0.3, error messages are stored in `errmsg.txt` files in the language directories under `sql/share`. The files have names like `czech/errmsg.txt`, `danish/errmsg.txt`, and so forth, and each one is language-specific. Each of these language-specific files must contain a line for each error message, so adding a new message involves adding a line to the `errmsg.txt` file for every language. The procedure involves adding the English message to the `english/errmsg.txt` file and running a script that adds the message to the other language-specific files. Translators may translate the message in other `errmsg.txt` files later.

1. Open the file `sql/share/english/errmsg.txt` in an editor.
2. Add new error messages at the end of this file. Each message should be on a separate line, and it must be quoted within double quote (") characters. By convention, every message line except the last should end with a comma (,) following the second double quote.
3. For each new error message, add a `#define` line to the `include/mysqld_error.h` file before the last line (`#define ER_ERROR_MESSAGES`).
4. Adjust the value of `ER_ERROR_MESSAGES` to the new number of error messages.
5. Add the defined error symbols to `include/sql_state.h`. This file contains the SQL states for the error messages. If the new errors don't have SQL states, add a comment instead. Note that this file must be kept sorted according to the value of the error number. That is, although the

`sql_state.h` file might not contain an entry for every symbol in `mysqld_error.h`, those entries that are present in `sql_state.h` must appear in the same order as those for the corresponding entries in `mysqld_error.h`.

6. Go to the `sql` directory in a terminal window and type `.Jadd_errmsg N`. This will copy the last *N* error messages from `share/english.txt` to all the other language files in `share/`.
7. Translate the error message for those languages that you know by editing the files `share/language/errmsg.txt`.
8. Make a full build (**configure + make**). A **make all** is insufficient to build the `sql/share/*/errmsg.sys` files.

Adding Storage Engine Error Messages

To add error messages for table handlers, the following example may be helpful.

Purpose: Implement the `handler::get_error_message` function as `ha_federated::get_error_message` to return the handler-specific error message.

Example:

1. When an error occurs you return an error code. (It should not be in the range of those that `HA_ERR` uses, which currently is 120-159.)
2. When `handler::print_error` is called to convert the handler error code to a MySQL error code, it will enter the default label of the `switch(error)` statement:

```
handler.cc:1721
default:
{
    /* The error was "unknown" to this function.
    Ask handler if it has got a message for this error */
    bool temporary= FALSE;
    String str;
    temporary= get_error_message(error, &str);
    if (!str.is_empty())
    {
        const char* engine= table_type();
        if (temporary)
            my_error(ER_GET_TEMPORARY_ERRMSG, MYF(0), error, str.ptr(), engine);
        else
            my_error(ER_GET_ERRMSG, MYF(0), error, str.ptr(), engine);
    }
    else
        my_error(ER_GET_ERRNO, errflag, error);
    DEBUG_VOID_RETURN;
}
}
```

1. Thus the `handler::get_error_message` is called and you can return the handler-specific error message, which is either a static error message that you retrieve from an error/string array, or a dynamic one that you format when the error occurs.

When you have returned the error message it will be passed to MySQL and formatted as `Got error %d '%-.100s' from %s`. For example:

```
Got error 788 'Could not connect to remote server fed.bb.pl' from FEDERATED
```

The `Got error %d` part will be returned in the user's selected language, but the handler-specific one will use English (unless the handler supports returning the handler error message in the user's selected language).

Appendix A. MySQL Source Code Distribution

Table of Contents

Directory Listing	367
The <code>BUILD</code> Directory	369
The <code>client</code> Directory	369
The <code>config</code> Directory	369
The <code>cmd-line-utils</code> Directory	370
The <code>debug</code> Directory	371
The <code>Docs</code> Directory	371
The <code>extra</code> Directory	371
The <code>heap</code> Directory	372
The <code>include</code> Directory	373
The <code>innobase</code> Directory	373
The <code>libmysql</code> Directory	373
The <code>libmysql_r</code> Directory	374
The <code>libmysqld</code> Directory	374
The <code>man</code> Directory	374
The <code>myisam</code> Directory	374
The <code>myisammrg</code> Directory	376
The <code>mysql-test</code> Directory	377
The <code>mysys</code> Directory	377
The <code>ndb</code> Directory	381
The <code>netware</code> Directory	382
The <code>NEW-RPMS</code> Directory	382
The <code>os2</code> Directory	382
The <code>pstack</code> Directory	382
The <code>regex</code> Directory	383
The <code>SCCS</code> Directory	383
The <code>scripts</code> Directory	383
The <code>server-tools</code> Directory	384
The <code>sql</code> Directory	384
The <code>sql-bench</code> Directory	388
The <code>sql-common</code> Directory	388
The <code>SSL</code> Directory	388
The <code>strings</code> Directory	388
The <code>support-files</code> Directory	390
The <code>tests</code> Directory	390
The <code>tools</code> Directory	391
The <code>VC++Files</code> Directory	391
The <code>vio</code> Directory	392
The <code>zlib</code> Directory	392

This is a description of the files that you get when you download the source code of MySQL. This description begins with a list of the main directories and a short comment about each one. Then, for each directory, in alphabetical order, a longer description is supplied. When a directory contains significant program files, a list of each C program is given along with an explanation of its intended function.

Directory Listing

Directory — Short Comment

- `bdb` — The Berkeley Database table handler

- BitKeeper — BitKeeper administration (not part of the source distribution)
- BUILD — Frequently used build scripts
- client — Client library
- cmd-line-utils — Command-line utilities (libedit and readline)
- config — Some files used during build
- debug — Fred Fish's debug library
- Docs — documentation files
- extra — Some minor standalone utility programs
- heap — The HEAP table handler
- include — Header (*.h) files for most libraries; includes all header files distributed with the MySQL binary distribution
- innobase — The Innobase (InnoDB) table handler
- libmysql — For producing MySQL as a library (e.g. a Windows .DLL)
- libmysql_r — For building a thread-safe libmysql library
- libmysqld — The MySQL Server as an embeddable library
- man — Some user-contributed manual pages
- myisam — The [MyISAM](#) table handler
- myisammrg — The [MyISAM](#) Merge table handler
- mysql-test — A test suite for mysqld
- mysys — MySQL system library (Low level routines for file access etc.)
- ndb — MySQL Cluster
- netware — Files related to the Novell NetWare version of MySQL
- NEW-RPMS — Directory to place RPMs while making a distribution
- os2 — Routines for working with the OS/2 operating system
- pstack — Process stack display (not currently used)
- regex — Henry Spencer's Regular Expression library for support of REGEXP function
- SCCS — Source Code Control System (not part of source distribution)
- scripts — SQL batches, e.g. mysqlbug and mysql_install_db
- server-tools — instance manager
- sql — Programs for handling SQL commands; the "core" of MySQL
- sql-bench — The MySQL benchmarks
- sql-common — Some .c files related to sql directory
- SSL — Secure Sockets Layer; includes an example certification one can use to test an SSL (secure) database connection

- strings — Library for C string routines, e.g. `atof`, `strchr`
- support-files — Files used to build MySQL on different systems
- tests — Tests in Perl and in C
- tools — `mysqlmanager.c` (tool under development, not yet useful)
- VC++Files — Includes this entire directory, repeated for VC++ (Windows) use
- vio — Virtual I/O Library
- zlib — Data compression library, used on Windows

The **BUILD** Directory

Frequently used build scripts.

This directory contains the build switches for compilation on various platforms. There is a subdirectory for each set of options. The main ones are:

- alpha
- ia64
- pentium (with and without debug or bdb, etc.)
- solaris

The **client** Directory

Client library.

The client library includes `mysql.cc` (the source of the `mysql` executable) and other utilities. Most of the utilities are mentioned in the MySQL Reference Manual. Generally these are standalone C programs which one runs in "client mode", that is, they call the server.

The C program files in the directory are:

- `get_password.c` --- ask for a password from the console
- `mysql.cc` --- "The MySQL command tool"
- `mysqladmin.cc` --- maintenance of MySQL databases
- `mysqlcheck.c` --- check all databases, check connect, etc.
- `mysqldump.c` --- dump table's contents as SQL statements, suitable to backup a MySQL database
- `mysqlimport.c` --- import text files in different formats into tables
- `mysqlmanager-pwgen.c` --- pwgen stands for "password generation" (not currently maintained)
- `mysqlmanagerc.c` --- entry point for mysql manager (not currently maintained)
- `mysqlshow.c` --- show databases, tables or columns
- `mysqltest.c` --- test program used by the mysql-test suite, `mysql-test-run`

The **config** Directory

Macros for use during build.

There is a single subdirectory: `\ac-macros`. All the files in it have the extension `.m4`, which is a normal expectation of the GNU autoconf tool.

The `cmd-line-utils` Directory

Command-line utilities (libedit and readline).

There are two subdirectories: `\readline` and `\libedit`. All the files here are "non-MySQL" files, in the sense that MySQL AB didn't produce them, it just uses them. It should be unnecessary to study the programs in these files unless you are writing or debugging a tty-like client for MySQL, such as `mysql.exe`.

The `\readline` subdirectory contains the files of the GNU Readline Library, "a library for reading lines of text with interactive input and history editing". The programs are copyrighted by the Free Software Foundation.

The `\libedit` (library of edit functions) subdirectory has files written by Christos Zoulas. They are distributed and modified under the BSD License. These files are for editing the line contents.

These are the program files in the `\libedit` subdirectory:

- `chared.c` --- character editor
- `common.c` --- common editor functions
- `el.c` --- editline interface functions
- `emacs.c` --- emacs functions
- `fgetln.c` --- get line
- `hist.c` --- history access functions
- `history.c` --- more history access functions
- `key.c` --- procedures for maintaining the extended-key map
- `map.c` --- editor function definitions
- `parse.c` --- parse an editline extended command
- `prompt.c` --- prompt printing functions
- `read.c` --- terminal read functions
- `readline.c` --- read line
- `refresh.c` --- "lower level screen refreshing functions"
- `search.c` --- "history and character search functions"
- `sig.c` --- for signal handling
- `strlcpy.c` --- string copy
- `term.c` --- "editor/termcap-curses interface"
- `tokenizer.c` --- Bourne shell line tokenizer
- `tty.c` --- for a tty interface
- `unvis.c` --- reverse effect of `vis.c`
- `vi.c` --- commands used when in the `vi` (editor) mode

- vis.c --- encode characters

The **debug** Directory

Fred Fish's debug library.

This is not really part of the MySQL package. Rather, it's a set of public-domain routines which are useful for debugging MySQL programs. The MySQL Server and all .c and .cc programs support the use of this package.

How it works: One inserts a function call that begins with `DEBUG_*` in one of the regular MYSQL programs. For example, in `get_password.c`, you will find this line:

```
DEBUG_ENTER("get_tty_password");
```

at the start of a routine, and this line:

```
DEBUG_RETURN(my_strdup(to,MYF(MY_FAE)));
```

at the end of the routine. These lines don't affect production code. Features of the debug library include extensive reporting and profiling (the latter has not been used by the MySQL team).

The C programs in this directory are:

- debug.c --- The main module
- debug_analyze.c --- Reads a file produced by trace functions
- example1.c --- A tiny example
- example2.c --- A tiny example
- example3.c --- A tiny example
- factorial.c --- A tiny example
- main.c --- A tiny example
- my_main.c --- MySQL-specific main.c variant
- sanity.c --- Declaration of a variable

The **Docs** Directory

With the BitKeeper downloads, /Docs is nearly empty. Binary and source distributions include some pre-formatted documentation files, such as the MySQL Reference manual in Info format (for Unix) or CHM format (for Windows).

The **extra** Directory

Some minor standalone utility programs.

These programs are all standalone utilities, that is, they have a `main()` function and their main role is to show information that the MySQL server needs or produces. Most are unimportant. They are as follows:

- charset2html.c --- checks your browser's character set
- comp_err.c --- makes error-message files from a multi-language source
- my_print_defaults.c --- print parameters from my.ini files. Can also be used in scripts to enable processing of my.ini files.

- `mysql_waitpid.c` --- wait for a program to terminate. Useful for shell scripts when one needs to wait until a process terminates.
- `perror.c` --- "print error" --- given error number, display message
- `replace.c` --- replace strings in text files or pipe
- `resolve_stack_dump.c` --- show symbolic information from a MySQL stack dump, normally found in the `mysql.err` file
- `resolveip.c` --- convert an IP address to a hostname, or vice versa

The [heap](#) Directory

The HEAP (MEMORY) table handler.

All the MySQL table handlers (i.e. the handlers that MySQL itself produces) have files with similar names and functions. Thus, this ([heap](#)) directory contains a lot of duplication of the [myisam](#) directory (for the [MyISAM](#) table handler). Such duplicates have been marked with an "*" in the following list. For example, you will find that `\heap\hp_extra.c` has a close equivalent in the [myisam](#) directory (`\myisam\mi_extra.c`) with the same descriptive comment. (Some of the differences arise because [HEAP](#) has different structures. [HEAP](#) does not need to use the sort of B-tree indexing that [ISAM](#) and [MyISAM](#) use; instead there is a hash index. Most importantly, [HEAP](#) is entirely in memory. File-I/O routines lose some of their vitality in such a context.)

- `hp_block.c` --- Read/write a block (i.e. a page)
- `hp_clear.c` --- Remove all records in the table
- `hp_close.c` --- * close database
- `hp_create.c` --- * create a table
- `hp_delete.c` --- * delete a row
- `hp_extra.c` --- * for setting options and buffer sizes when optimizing
- `hp_hash.c` --- Hash functions used for saving keys
- `hp_info.c` --- * Information about database status
- `hp_open.c` --- * open database
- `hp_panic.c` --- * the `hp_panic` routine, for shutdowns and flushes
- `hp_rename.c` --- * rename a table
- `hp_rfirst.c` --- * read first row through a specific key (very short)
- `hp_rkey.c` --- * read record using a key
- `hp_rlast.c` --- * read last row with same key as previously-read row
- `hp_rnext.c` --- * read next row with same key as previously-read row
- `hp_rprev.c` --- * read previous row with same key as previously-read row
- `hp_rrnd.c` --- * read a row based on position
- `hp_rsame.c` --- * find current row using positional read or key-based read
- `hp_scan.c` --- * read all rows sequentially
- `hp_static.c` --- * static variables (very short)

- `hp_test1.c` --- * testing basic functions
- `hp_test2.c` --- * testing database and storing results
- `hp_update.c` --- * update an existing row
- `hp_write.c` --- * insert a new row

There are fewer files in the heap directory than in the `mysam` directory, because fewer are necessary. For example, there is no need for a `\mysam\mi_cache.c` equivalent (to cache reads) or a `\mysam\mi_log.c` equivalent (to log statements).

The `include` Directory

Header (`*.h`) files for most libraries; includes all header files distributed with the MySQL binary distribution.

These files may be included in C program files. Note that each individual directory will also have its own `*.h` files, for including in its own `*.c` programs. The `*.h` files in the `include` directory are ones that might be included from more than one place.

For example, the `mysys` directory contains a C file named `rijndael.c`, but does not include `rijndael.h`. The `include` directory contains `rijndael.h`. Looking further, you'll find that `rijndael.h` is also included in other places: by `my_aes.c` and `my_aes.h`.

The `include` directory contains 55 `*.h` (header) files.

The `innobase` Directory

The Innobase (InnoDB) table handler.

A full description of these files can be found elsewhere in this document.

The `libmysql` Directory

The MySQL Library, Part 1.

The files here are for producing MySQL as a library (e.g. a Windows DLL). The idea is that, instead of producing separate `mysql` (client) and `mysqld` (server) programs, one produces a library. Instead of sending messages, the client part merely calls the server part.

The `libmysql` files are split into three directories: `libmysql` (this one), `libmysql_r` (the next one), and `libmysqld` (the next one after that).

The "library of mysql" has some client-connection modules. For example, as described in an earlier section of this manual, there is a discussion of `libmysql/libmysql.c` which sends packets from the client to the server. Many of the entries in the `libmysql` directory (and in the following `libmysqld` directory) are 'symlinks' on Linux, that is, they are in fact pointers to files in other directories.

The program files on this directory are:

- `conf_to_src.c` --- has to do with charsets
- `dll.c` --- initialization of the dll library
- `errmsg.c` --- English error messages, compare `\mysys\errors.c`
- `get_password.c` --- get password
- `libmysql.c` --- the code that implements the MySQL API, i.e. the functions a client that wants to connect to MySQL will call

- `manager.c` --- initialize/connect/fetch with MySQL manager

The `libmysql_r` Directory

The MySQL Library, Part 2.

There is only one file here, used to build a thread-safe `libmysql` library:

- `makefile.am`

The `libmysqld` Directory

The MySQL library, Part 3.

The Embedded MySQL Server Library. The product of `libmysqld` is not a client/server affair, but a library. There is a wrapper to emulate the client calls. The program files on this directory are:

- `libmysqld.c` --- The called side, compare the `mysqld.exe` source
- `lib_sql.c` --- Emulate the `vio` directory's communication buffer

The `man` Directory

Some user-contributed manual pages

These are user-contributed "man" (manual) pages in a special markup format. The format is described in a document with a heading like "man page for man" or "macros to format man pages" which you can find in a Linux directory or on the Internet.

The `myisam` Directory

The `MyISAM` table handler.

The C files in this subdirectory come in six main groups:

- `ft*.c` files --- `ft` stands for "Full Text", code contributed by Sergei Golubchik
- `mi*.c` files --- `mi` stands for "My Isam", these are the main programs for `Myisam`
- `myisam*.c` files --- for example, "myisamchk" utility routine functions source
- `rt*.c` files --- `rt` stands for "rtree", some code was written by Alexander Barkov
- `sp*.c` files --- `sp` stands for "spatial", some code was written by Ramil Kalimullin
- `sort.c` --- this is a single file that sorts keys for index-create purposes

The "full text" and "rtree" and "spatial" program sets are for special purposes, so this document focuses only on the `mi*.c` "myisam" C programs. They are:

- `mi_cache.c` --- for reading records from a cache
- `mi_changed.c` --- a single routine for setting a "changed" flag (very short)
- `mi_check.c` --- for checking and repairing tables. Used by the `myisamchk` program and by the MySQL server.
- `mi_checksum.c` --- calculates a checksum for a row
- `mi_close.c` --- close database
- `mi_create.c` --- create a table

- `mi_dbug.c` --- support routines for use with "dbug" (see \dbug description)
- `mi_delete.c` --- delete a row
- `mi_delete_all.c` --- delete all rows
- `mi_delete_table.c` --- delete a table (very short)
- `mi_dynrec.c` --- functions to handle space-packed records and blobs
- `mi_extra.c` --- setting options and buffer sizes when optimizing
- `mi_info.c` --- return useful base information for an open table
- `mi_key.c` --- for handling keys
- `mi_keycache.c` --- for handling key caches
- `mi_locking.c` --- lock database
- `mi_log.c` --- save commands in a log file which `myisamlog` program can read. Can be used to exactly replay a set of changes to a table.
- `mi_open.c` --- open database
- `mi_packrec.c` --- read from a data file compressed with `mysampack`
- `mi_page.c` --- read and write pages containing keys
- `mi_panic.c` --- the `mi_panic` routine, probably for sudden shutdowns
- `mi_preload.c` --- preload indexes into key cache
- `mi_range.c` --- approximate count of how many records lie between two keys
- `mi_rename.c` --- rename a table
- `mi_rfirst.c` --- read first row through a specific key (very short)
- `mi_rkey.c` --- read a record using a key
- `mi_rlast.c` --- read last row with same key as previously-read row
- `mi_rnext.c` --- read next row with same key as previously-read row
- `mi_rnext_same.c` --- same as `mi_rnext.c`, but abort if the key changes
- `mi_rprev.c` --- read previous row with same key as previously-read row
- `mi_rrnd.c` --- read a row based on position
- `mi_rsame.c` --- find current row using positional read or key-based read
- `mi_rsamepos.c` --- positional read
- `mi_scan.c` --- read all rows sequentially
- `mi_search.c` --- key-handling functions
- `mi_static.c` --- static variables (very short)
- `mi_statrec.c` --- functions to handle fixed-length records
- `mi_test1.c` --- testing basic functions

- [mi_test2.c](#) --- testing database and storing results
- [mi_test3.c](#) --- testing locking
- [mi_unique.c](#) --- functions to check if a row is unique
- [mi_update.c](#) --- update an existing row
- [mi_write.c](#) --- insert a new row

The [myisammrg](#) Directory

[MyISAM](#) Merge table handler.

As with other table handlers, you'll find that the [*.c](#) files in the [myisammrg](#) directory have counterparts in the [myisam](#) directory. In fact, this general description of a [myisammrg](#) program is almost always true: The [myisammrg](#) function checks an argument, the [myisammrg](#) function formulates an expression for passing to a [myisam](#) function, the [myisammrg](#) calls a [myisam](#) function, the [myisammrg](#) function returns.

These are the 22 files in the [myisammrg](#) directory, with notes about the [myisam](#) functions or programs they're connected with:

- [myrg_close.c](#) --- [mi_close.c](#)
- [myrg_create.c](#) --- [mi_create.c](#)
- [myrg_delete.c](#) --- [mi_delete.c](#) / delete last-read record
- [myrg_extra.c](#) --- [mi_extra.c](#) / "extra functions we want to do ..."
- [myrg_info.c](#) --- [mi_info.c](#) / display information about a mymerge file
- [myrg_locking.c](#) --- [mi_locking.c](#) / lock databases
- [myrg_open.c](#) --- [mi_open.c](#) / open a [MyISAM MERGE](#) table
- [myrg_panic.c](#) --- [mi_panic.c](#) / close in a hurry
- [myrg_queue.c](#) --- read record based on a key
- [myrg_range.c](#) --- [mi_range.c](#) / find records in a range
- [myrg_rfirst.c](#) --- [mi_rfirst.c](#) / read first record according to specific key
- [myrg_rkey.c](#) --- [mi_rkey.c](#) / read record based on a key
- [myrg_rlast.c](#) --- [mi_rlast.c](#) / read last row with same key as previous read
- [myrg_rnext.c](#) --- [mi_rnext.c](#) / read next row with same key as previous read
- [myrg_rnext_same.c](#) --- [mi_rnext_same.c](#) / read next row with same key
- [myrg_rprev.c](#) --- [mi_rprev.c](#) / read previous row with same key
- [myrg_rrnd.c](#) --- [mi_rrnd.c](#) / read record with random access
- [myrg_rsame.c](#) --- [mi_rsame.c](#) / call [mi_rsame](#) function, see [\myisam\mi_rsame.c](#)
- [myrg_static.c](#) --- [mi_static.c](#) / static variable declaration
- [myrg_update.c](#) --- [mi_update.c](#) / call [mi_update](#) function, see [\myisam\mi_update.c](#)
- [myrg_write.c](#) --- [mi_write.c](#) / call [mi_write](#) function, see [\myisam\mi_write.c](#)

The `mysql-test` Directory

A test suite for `mysqld`.

The directory has a [README](#) file which explains how to run the tests, how to make new tests (in files with the filename extension `*.test`), and how to report errors.

There are four subdirectories:

- `\misc` --- contains one minor Perl program
- `\ndb` --- for MySQL Cluster tests
- `\r` --- contains `*.result`, i.e. "what happened" files and `*.required`, i.e. "what should happen" file
- `\std_data` --- contains standard data for input to tests
- `\t` --- contains tests

There are 400 `*.test` files in the `\t` subdirectory. Primarily these are SQL scripts which try out a feature, output a result, and compare the result with what's required. Some samples of what the test files check are: latin1_de comparisons, date additions, the [HAVING](#) clause, outer joins, openssl, load data, logging, truncate, and [UNION](#).

There are other tests in these directories:

- `sql-bench`
- `tests`

The `mysys` Directory

MySQL system library. Low level routines for file access and so on.

There are 125 `*.c` programs in this directory:

- `array.c` --- Dynamic array handling
- `charset.c` --- Using dynamic character sets, set default character set, ...
- `charset-def.c` --- Include character sets in the client using
- `checksum.c` --- Calculate checksum for a memory block, used for `pack_isam`
- `default.c` --- Find defaults from `*.cnf` or `*.ini` files
- `default_modify.c` --- edit option file
- `errors.c` --- English text of global errors
- `hash.c` --- Hash search/compare/free functions "for saving keys"
- `list.c` --- Double-linked lists
- `make-conf.c` --- "Make a charset `.conf` file out of a `ctype-charset.c` file"
- `md5.c` --- MD5 ("Message Digest 5") algorithm from RSA Data Security
- `mf_brkhanth.c` --- Prevent user from doing a Break during critical execution (not used in MySQL; can be used by standalone [MyISAM](#) applications)
- `mf_cache.c` --- "Open a temporary file and cache it with `io_cache`"

- `mf_dirname.c` --- Parse/convert directory names
- `mf_fn_ext.c` --- Get filename extension
- `mf_format.c` --- Format a filename
- `mf_getdate.c` --- Get date, return in yyyy-mm-dd hh:mm:ss format
- `mf_iocache.c` --- Cached read/write of files in fixed-size units
- `mf_iocache2.c` --- Continuation of `mf_iocache.c`
- `mf_keycache.c` --- Key block caching for certain file types
- `mf_keycaches.c` --- Handling of multiple key caches
- `mf_loadpath.c` --- Return full path name (no `..\` stuff)
- `mf_pack.c` --- Packing/unpacking directory names for create purposes
- `mf_path.c` --- Determine where a program can find its files
- `mf_qsort.c` --- Quicksort
- `mf_qsort2.c` --- Quicksort, part 2 (allows the passing of an extra argument to the sort-compare routine)
- `mf_radix.c` --- Radix sort
- `mf_same.c` --- Determine whether filenames are the same
- `mf_sort.c` --- Sort with choice of Quicksort or Radix sort
- `mf_soundex.c` --- Soundex algorithm derived from EDN Nov. 14, 1985 (pg. 36)
- `mf_strip.c` --- Strip trail spaces from a string
- `mf_tempdir.c` --- Initialize/find/free temporary directory
- `mf_tempfile.c` --- Create a temporary file
- `mf_unixpath.c` --- Convert filename to UNIX-style filename
- `mf_util.c` --- Routines, `#ifdef`'d, which may be missing on some machines
- `mf_wcomp.c` --- Comparisons with wildcards
- `mf_wfile.c` --- Finding files with wildcards
- `mulalloc.c` --- Malloc many pointers at the same time
- `my_access.c` --- Check if file or path is accessible
- `my_aes.c` --- AES encryption
- `my_alarm.c` --- Set a variable value when an alarm is received
- `my_alloc.c` --- malloc of results which will be freed simultaneously
- `my_append.c` --- one file to another
- `my_bit.c` --- smallest X where $2^X \geq \text{value}$, maybe useful for divisions
- `my_bitmap.c` --- Handle `uchar` arrays as large bitmaps

- `my_chsize.c` --- Truncate file if shorter, else fill with a filler character
- `my_clock.c` --- Time-of-day ("clock()") function, with OS-dependent `#ifdef`'s
- `my_compress.c` --- Compress packet (see also description of `\zlib` directory)
- `my_copy.c` --- Copy files
- `my_crc32.c` --- Include `\zlib\crc32.c`
- `my_create.c` --- Create file
- `my_delete.c` --- Delete file
- `my_div.c` --- Get file's name
- `my_dup.c` --- Open a duplicated file
- `my_error.c` --- Return formatted error to user
- `my_file.c` --- See how many open files we want
- `my_fopen.c` --- File open
- `my_fstream.c` --- Streaming file read/write
- `my_gethostbyname.c` --- Thread-safe version of standard `net gethostbyname()` func
- `my_gethwaddr.c` --- Get hardware address for an interface
- `my_getopt.c` --- Find out what options are in effect
- `my_getsystemtime.c` --- Time-of-day functions, portably
- `my_getwd.c` --- Get working directory
- `my_handler.c` --- Compare two keys in various possible formats
- `my_init.c` --- Initialize variables and functions in the `mysys` library
- `my_largepage.c` --- Gets the size of large pages from the OS
- `my_lib.c` --- Compare/convert directory names and file names
- `my_lock.c` --- Lock part of a file
- `my_lockmem.c` --- "Allocate a block of locked memory"
- `my_lread.c` --- Read a specified number of bytes from a file into memory
- `my_lwrite.c` --- Write a specified number of bytes from memory into a file
- `my_malloc.c` --- Malloc (memory allocate) and dup functions
- `my_messnc.c` --- Put out a message on `stderr` with "no curses"
- `my_mkdir.c` --- Make directory
- `my_mmap.c` --- Memory mapping
- `my_net.c` --- Thread-safe version of `net inet_ntoa` function
- `my_netware.c` --- Functions used only with the Novell Netware version of MySQL
- `my_once.c` --- Allocation / duplication for "things we don't need to free"

- `my_open.c` --- Open a file
- `my_os2cond.c` --- OS2-specific: "A simple implementation of posix conditions"
- `my_os2dirsrch.c` --- OS2-specific: Emulate a Win32 directory search
- `my_os2dlfcn.c` --- OS2-specific: Emulate UNIX dynamic loading
- `my_os2file64.c` --- OS2-specific: For File64bit setting
- `my_os2mutex.c` --- OS2-specific: For mutex handling
- `my_os2thread.c` --- OS2-specific: For thread handling
- `my_os2tls.c` --- OS2-specific: For thread-local storage
- `my_port.c` --- OS/machine-dependent porting functions, e.g. AIX-specific `my_ulonglong2double()`
- `my_pread.c` --- Read a specified number of bytes from a file
- `my_pthread.c` --- A wrapper for thread-handling functions in different OSs
- `my_quick.c` --- Read/write (labeled a "quicker" interface, perhaps obsolete)
- `my_read.c` --- Read a specified number of bytes from a file, possibly retry
- `my_realloc.c` --- Reallocate memory allocated with `my_alloc.c` (probably)
- `my_redel.c` --- Rename and delete file
- `my_rename.c` --- Rename without delete
- `my_seek.c` --- Seek, i.e. point to a spot within a file
- `my_semaphore.c` --- Semaphore routines, for use on OS that doesn't support them
- `my_sleep.c` --- Wait n microseconds
- `my_static.c` --- Static variables used by the `mysys` library
- `my_symlink.c` --- Read a symbolic link (symlinks are a UNIX thing, I guess)
- `my_symlink2.c` --- Part 2 of `my_symlink.c`
- `my_sync.c` --- Sync data in file to disk
- `my_thr_init.c` --- initialize/allocate "all `mysys` & debug thread variables"
- `my_wincond.c` --- Windows-specific: emulate Posix conditions
- `my_windac.c` --- Windows NT/2000 discretionary access control functions
- `my_winsem.c` --- Windows-specific: emulate Posix threads
- `my_winthread.c` --- Windows-specific: emulate Posix threads
- `my_write.c` --- Write a specified number of bytes to a file
- `ptr_cmp.c` --- Point to an optimal byte-comparison function
- `queues.c` --- Handle priority queues as in Robert Sedgewick's book
- `raid2.c` --- RAID support (the true implementation is in `raid.cc`)
- `rijndael.c` --- "Optimized ANSI C code for the Rijndael cipher (now AES)"

- `safemalloc.c` --- A version of the standard `malloc()` with safety checking
- `sha1.c` --- Implementation of Secure Hashing Algorithm 1
- `string.c` --- Initialize/append/free dynamically-sized strings; see also `sql_string.cc` in the `/sql` directory
- `testhash.c` --- Standalone program: test the hash library routines
- `test_charset.c` --- Standalone program: display character set information
- `test_dir.c` --- Standalone program: placeholder for "test all functions" idea
- `test_fn.c` --- Standalone program: apparently tests a function
- `test_xml.c` --- Standalone program: test XML routines
- `thr_alarm.c` --- Thread alarms and signal handling
- `thr_lock.c` --- "Read and write locks for Posix threads"
- `thr_mutex.c` --- A wrapper for mutex functions
- `thr_rwlock.c` --- Synchronizes the readers' thread locks with the writer's lock
- `tree.c` --- Initialize/search/free binary trees
- `typelib.c` --- Find a string in a set of strings; returns the offset to the string found

You can find documentation for the main functions in these files elsewhere in this document. For example, the main functions in [my_getwd.c](#) are described thus:

```
"int my_getwd _A((string buf, uint size, myf MyFlags));
  int my_setwd _A((const char *dir, myf MyFlags));
  Get and set working directory."
```

The [ndb](#) Directory

The `ndb` (MySQL Cluster) source code.

MySQL's shared-nothing in-memory feature is practically a DBMS by itself. We generally use the term "`ndb`" when referring to the storage engine, and the term "MySQL Cluster" when referring to the combination of the storage engine and the rest of the MySQL facilities.

The sub-directories within `ndb` are:

- `bin` --- Two script files
- `config` --- Files needed for building
- `demos` --- Demonstration scripts
- `docs` --- A doxygen output and a `.txt` file
- `home` --- Some scripts and `.pl` files
- `include` --- The `.h` files
- `lib` --- empty
- `ndbapi-examples` --- Examples for the API
- `src` --- The `.cpp` files
- `test` --- Files for testing

- tools --- Programs for testing select, drop, and so on

The **netware** Directory

Files related to the Novell NetWare version of MySQL.

There are 43 files on this directory. Most have filename extensions of `*.def`, `*.sql`, or `*.c`.

The twenty-eight `*.def` files are all from Novell Inc. They contain import or export symbols. (`.def` is a common filename extension for "definition".)

The three `*.sql` files are short scripts of SQL statements used in testing.

These are the five `*.c` files, all from Novell Inc.:

- `libmysqlmain.c` --- Only one function: `init_available_charsets()`
- `my_manage.c` --- Standalone management utility
- `mysql_install_db.c` --- Compare `\scripts\mysql_install_db.sh`
- `mysql_test_run.c` --- Short test program
- `mysqld_safe.c` --- Compare `\scripts\mysqld_safe.sh`

Perhaps the most important files are:

- `netware/BUILD/*patch` --- NetWare-specific build instructions and switches (compare the files in the BUILD directory)

For instructions about basic installation, see "Deployment Guide For NetWare AMP" at: <http://developer.novell.com/ndk/whitepapers/namp.htm>

The **NEW-RPMS** Directory

Directory to place RPMs while making a distribution.

This directory is not part of the Windows distribution. It is a temporary directory used during RPM builds with Linux distributions. You only see it after you've done a "build".

The **os2** Directory

Routines for working with the OS2 operating system.

The files in this directory are the product of the efforts of three people from outside MySQL: Yuri Dario, Timo Maier, and John M Alfredsson. There are no `.c` program files in this directory.

The contents of `\os2` are:

- A `Readme.Txt` file
- An `\include` subdirectory containing `.h` files which are for OS/2 only
- Files used in the build process (configuration, switches, and one `.obj`)

The README file refers to MySQL version 3.23, which suggests that there have been no updates for MySQL 4.0 for this section.

The **pstack** Directory

Process stack display (not currently used).

This is a set of publicly-available debugging aids which all do pretty well the same thing: display the contents of the stack, along with symbolic information, for a running process. There are versions for various object file formats (such as ELF and IEEE-695). Most of the programs are copyrighted by the Free Software Foundation and are marked as "part of GNU Binutils".

In other words, the pstack files are not really part of the MySQL library. They are merely useful when you re-program some MYSQL code and it crashes.

The [regex](#) Directory

Henry Spencer's Regular Expression library for support of REGEXP function.

This is the copyrighted product of Henry Spencer from the University of Toronto. It's a fairly-well-known implementation of the requirements of POSIX 1003.2 Section 2.8. The library is bundled with Apache and is the default implementation for regular-expression handling in BSD Unix. MySQL's Monty Widenius has made minor changes in three programs (debug.c, engine.c, regexec.c) but this is not a MySQL package. MySQL calls it only in order to support two MySQL functions: REGEXP and RLIKE.

Some of Mr Spencer's documentation for the regex library can be found in the README and WHATSNEW files.

One MySQL program which uses regex is `\cmd-line-utils\libedit\search.c`

This program calls the 'regcomp' function, which is the entry point in `\regex\regex.c`.

The [sccs](#) Directory

Source Code Control System (not part of source distribution).

You will see this directory if and only if you used BitKeeper for downloading the source. The files here are for BitKeeper administration and are not of interest to application programmers.

The [scripts](#) Directory

SQL batches, e.g. `mysqlbug` and `mysql_install_db`.

The `*.sh` filename extension stands for "shell script". Linux programmers use it where Windows programmers would use a `*.bat` (batch filename extension).

Some of the `*.sh` files on this directory are:

- `fill_help_tables.sh` --- Create help-information tables and insert
- `make_binary_distribution.sh` --- Get configure information, make, produce tar
- `msql2mysql.sh` --- Convert (partly) mSQL programs and scripts to MySQL
- `mysqlbug.sh` --- Create a bug report and mail it
- `mysqld_multi.sh` --- Start/stop any number of mysqld instances
- `mysqld_safe-watch.sh` --- Start/restart in safe mode
- `mysqld_safe.sh` --- Start/restart in safe mode
- `mysqldumpslow.sh` --- Parse and summarize the slow query log
- `mysqlhotcopy.sh` --- Hot backup
- `mysql_config.sh` --- Get configuration information that might be needed to compile a client
- `mysql_convert_table_format.sh` --- Conversion, e.g. from [ISAM](#) to [MyISAM](#)

- `mysql_explain_log.sh` --- Put a log (made with `--log`) into a MySQL table
- `mysql_find_rows.sh` --- Search for queries containing `<regex>`
- `mysql_fix_extensions.sh` --- Renames some file extensions, not recommended
- `mysql_fix_privilege_tables.sh` --- Fix `mysql.user` etc. when upgrading. Can be safely run during any upgrade to get the newest MySQL privilege tables
- `mysql_install_db.sh` --- Create privilege tables and func table
- `mysql_secure_installation.sh` --- Disallow remote root login, eliminate test, etc.
- `mysql_setpermission.sh` --- Aid to add users or databases, sets privileges
- `mysql_tableinfo.sh` --- Puts info re MySQL tables into a MySQL table
- `mysql_zap.sh` --- Kill processes that match pattern

The `server-tools` Directory

The instance manager.

Quoting from the README file within this directory: "Instance Manager - manage MySQL instances locally and remotely. File description: `mysqlmanager.cc` - entry point to the manager, main, options. `{h,cc}` - handle startup options. `manager.{h,cc}` - manager process. `mysql_connection.{h,cc}` - handle one connection with mysql client. See also instance manager architecture description in `mysqlmanager.cc`.

The `sql` Directory

Programs for handling SQL commands. The "core" of MySQL.

These are the `.c` and `.cc` files in the `sql` directory:

- `derror.cc` --- read language-dependent message file
- `des_key_file.cc` --- load DES keys from plaintext file
- `discover.cc` --- Functions for discovery of frm file from handler
- `field.cc` --- "implement classes defined in `field.h`" (long); defines all storage methods MySQL uses to store field information into records that are then passed to handlers
- `field_conv.cc` --- functions to copy data between fields
- `filesort.cc` --- sort a result set, using memory or temporary files
- `frm_crypt.cc` --- contains only one short function: `get_crypt_for_frm`
- `gen_lex_hash.cc` --- Knuth's algorithm from Vol 3 Sorting and Searching, Chapter 6.3; used to search for SQL keywords in a query
- `gstream.cc` --- `GTextReadStream`, used to read GIS objects
- `handler.cc` --- handler-calling functions
- `hash_filo.cc` --- static-sized hash tables, used to store info like hostname -> ip tables in a FIFO manner
- `ha_berkeley.cc` --- Handler: BDB
- `ha_blackhole.cc` --- Handler: Black Hole

- `ha_federated.cc` --- Handler: Federated
- `ha_heap.cc` --- Handler: Heap
- `ha_innodb.cc` --- Handler: InnoDB
- `ha_myisam.cc` --- Handler: MyISAM
- `ha_myisammrg.cc` --- Handler: (MyISAM MERGE)
- `ha_ndbcluster.cc` --- Handler: NDB
- `hostname.cc` --- Given IP, return hostname
- `init.cc` --- Init and dummy functions for interface with `unireg`
- `item.cc` --- Item functions
- `item_buff.cc` --- Buffers to save and compare item values
- `item_cmpfunc.cc` --- Definition of all compare functions
- `item_create.cc` --- Create an item. Used by `lex.h`.
- `item_func.cc` --- Numerical functions
- `item_geofunc.cc` --- Geometry functions
- `item_row.cc` --- Row items for comparing rows and for `IN` on rows
- `item_strfunc.cc` --- String functions
- `item_subselect.cc` --- Subqueries
- `item_sum.cc` --- Set functions (`SUM()`, `AVG()`, etc.)
- `item_strfunc.cc` --- String functions
- `item_subselect.cc` --- Item subquery
- `item_timefunc.cc` --- Date/time functions, e.g. week of year
- `item_uniq.cc` --- Empty file, here for compatibility reasons
- `key.cc` --- Functions to create keys from records and compare a key to a key in a record
- `lock.cc` --- Locks
- `log.cc` --- Logs
- `log_event.cc` --- Log event (a binary log consists of a stream of log events)
- `matherr.c` --- Handling overflow, underflow, etc.
- `mf_iocache.cc` --- Caching of (sequential) reads and writes
- `mysqld.cc` --- Source for `mysqld.exe`; includes the `main()` program that starts `mysqld`, handling of signals and connections
- `mf_decimal.cc` --- New decimal and numeric code
- `my_lock.c` --- Lock part of a file (like `/mysys/my_lock.c`, but with timeout handling for threads)
- `net_serv.cc` --- Read/write of packets on a network socket

- `nt_servc.cc` --- Initialize/register/remove an NT service
- `opt_range.cc` --- Range of keys
- `opt_sum.cc` --- Optimize functions in presence of (implied) `GROUP BY`
- `parse_file.cc` --- Text `.frm` files management routines
- `password.c` --- Password checking
- `procedure.cc` --- Procedure interface, as used in `SELECT * FROM Table_name PROCEDURE ANALYSE()`
- `protocol.cc` --- Low level functions for PACKING data that is sent to client; actual sending done with `net_serv.cc`
- `protocol_cursor.cc` --- Low level functions for storing data to be sent to the MySQL client
- `records.cc` --- Functions for easy reading of records, possible through a cache
- `repl_failsafe.cc` --- Replication fail-save (not yet implemented)
- `set_var.cc` --- Set and retrieve MySQL user variables
- `slave.cc` --- Procedures for a slave in a master/slave (replication) relation
- `sp.cc` --- DB storage of stored procedures and functions
- `sp_cache.cc` --- For stored procedures
- `sp_head.cc` --- For stored procedures
- `sp_pcontext.cc` --- For stored procedures
- `sp_rcontext.cc` --- For stored procedures
- `spatial.cc` --- Geometry stuff (lines, points, etc.)
- `sql_acl.cc` --- Functions related to ACL security; checks, stores, retrieves, and deletes MySQL user level privileges
- `sql_analyse.cc` --- Implements the `PROCEDURE ANALYSE()`, which analyzes a query result and returns the 'optimal' data type for each result column
- `sql_base.cc` --- Basic functions needed by many modules, like opening and closing tables with table cache management
- `sql_cache.cc` --- SQL query cache, with long comments about how caching works
- `sql_class.cc` --- SQL class; implements the SQL base classes, of which THD (THREAD object) is the most important
- `sql_client.cc` --- A function called by `my_net_init()` to set some check variables
- `sql_crypt.cc` --- Encode / decode, very short
- `sql_db.cc` --- Create / drop database
- `sql_delete.cc` --- The `DELETE` statement
- `sql_derived.cc` --- Derived tables, with long comments
- `sql_do.cc` --- The `DO` statement

- `sql_error.cc` --- Errors and warnings
- `sql_handler.cc` --- Implements the `HANDLER` interface, which gives direct access to rows in `MyISAM` and `InnoDB`
- `sql_help.cc` --- The `HELP` statement
- `sql_insert.cc` --- The `INSERT` statement
- `sql_lex.cc` --- Does lexical analysis of a query; i.e. breaks a query string into pieces and determines the basic type (number, string, keyword, etc.) of each piece
- `sql_list.cc` --- Only `list_node_end_of_list`, short (the rest of the list class is implemented in `sql_list.h`)
- `sql_load.cc` --- The `LOAD DATA` statement
- `sql_manager.cc` --- Maintenance tasks, e.g. flushing the buffers periodically; used with `BDB` table logs
- `sql_map.cc` --- Memory-mapped files (not yet in use)
- `sql_olap.cc` --- `ROLLUP`
- `sql_parse.cc` --- Parse an SQL statement; do initial checks and then jump to the function that should execute the statement
- `sql_prepare.cc` --- Prepare an SQL statement, or use a prepared statement
- `sql_rename.cc` --- Rename table
- `sql_repl.cc` --- Replication
- `sql_select.cc` --- Select and join optimization
- `sql_show.cc` --- The `SHOW` statement
- `sql_state.c` --- Functions to map `mysqld` `errno` to `sqlstate`
- `sql_string.cc` --- String functions: `alloc`, `realloc`, `copy`, `convert`, etc.
- `sql_table.cc` --- The `DROP TABLE` and `ALTER TABLE` statements
- `sql_test.cc` --- Some debugging information
- `sql_trigger.cc` --- Triggers
- `sql_udf.cc` --- User-defined functions
- `sql_union.cc` --- The `UNION` operator
- `sql_update.cc` --- The `UPDATE` statement
- `sql_view.cc` --- Views
- `stacktrace.c` --- Display stack trace (Linux/Intel only)
- `strfunc.cc` --- String functions
- `table.cc` --- Table metadata retrieval; read the table definition from a `.frm` file and store it in a `TABLE` object
- `thr_malloc.cc` --- Thread-safe interface to `/mysys/my_alloc.c`
- `time.cc` --- Date and time functions

- `udf_example.cc` --- Example file of user-defined functions
- `uniques.cc` --- Function to handle quick removal of duplicates
- `unireg.cc` --- Create a unireg form file (`.frm`) from a `FIELD` and field-info struct

The `sql-bench` Directory

The MySQL Benchmarks.

This directory has the programs and input files which MySQL uses for its comparisons of MySQL, PostgreSQL, mSQL, Solid, etc. Since MySQL publishes the comparative results, it's only right that it should make available all the material necessary to reproduce all the tests.

There are five subdirectories and sub-subdirectories:

- `\Comments` --- Comments about results from tests of Access, Adabas, etc.
- `\Data\ATIS` --- `.txt` files containing input data for the "ATIS" tests
- `\Data\Wisconsin` --- `.txt` files containing input data for the "Wisconsin" tests
- `\Results` --- old test results
- `\Results-win32` --- old test results from Windows 32-bit tests

There are twenty-four `*.sh` (shell script) files, which involve Perl programs.

There are three `*.bat` (batch) files.

There is one README file and one TODO file.

The `sql-common` Directory

Three files: `client.c`, `my_time.c`, `pack.c`. You will find symlinks to these files in other directories.

The `SSL` Directory

Secure Sockets Layer; includes an example certification one can use to test an SSL (secure) database connection.

This isn't a code directory. It contains a short note from Tonu Samuel (the `NOTES` file) and seven `*.pem` files. PEM stands for "Privacy Enhanced Mail" and is an Internet standard for adding security to electronic mail. Finally, there are two short scripts for running clients and servers over SSL connections.

The `strings` Directory

The string library.

Many of the files in this subdirectory are equivalent to well-known functions that appear in most C string libraries. For those, there is documentation available in most compiler handbooks.

On the other hand, some of the files are MySQL additions or improvements. Often the MySQL changes are attempts to optimize the standard libraries. It doesn't seem that anyone tried to optimize for recent Pentium class processors, though.

The `.C` files are:

- `bchange.c` --- short replacement routine written by Monty Widenius in 1987
- `bcmp.c` --- binary compare, rarely used

- `bcopy-duff.c` --- block copy: attempt to copy memory blocks faster than `memcpy`
- `bfill.c` --- byte fill, to fill a buffer with (length) copies of a byte
- `bmove.c` --- block move
- `bmove512.c` --- "should be the fastest way to move a multiple of 512 bytes"
- `bmove_upp.c` --- `bmove.c` variant, starting with last byte
- `bzero.c` --- something like `bfill` with an argument of 0
- `conf_to_src.c` --- reading a configuration file
- `ctype*.c` --- string handling programs for each char type MySQL handles
- `decimal.c` --- for decimal and numeric conversions
- `do_ctype.c` --- display case-conversion and sort-conversion tables
- `dump_map.c` --- standalone file
- `int2str.c` --- integer-to-string
- `is_prefix.c` --- checks whether `string1` starts with `string2`
- `llstr.c` --- convert long long to temporary-buffer string, return pointer
- `longlong2str.c` --- ditto, but to argument-buffer
- `memcmp.c` --- memory compare
- `memcpy.c` --- memory copy
- `memset.c` --- memory set
- `my_strtoll10.c` --- `longlong2str` for radix 10
- `my_vsnprintf.c` --- variant of `printf`
- `r_strinstr.c` --- see if one string is within another
- `str2int.c` --- convert string to integer
- `strappend.c` --- fill up a string to `n` characters
- `strcat.c` --- concatenate strings
- `strcend.c` --- point to where a character `C` occurs within `str`, or `NULL`
- `strchr.c` --- point to first place in string where character occurs
- `strcmp.c` --- compare two strings
- `strcont.c` --- point to where any one of a set of characters appears
- `strend.c` --- point to the `'\0'` byte which terminates `str`
- `strfill.c` --- fill a string with `n` copies of a byte
- `strinstr.c` --- find string within string
- `strlen.c` --- return length of string in bytes
- `strmake.c` --- create new string from old string with fixed length, append end `\0` if needed

- `strmov.c` --- move source to dest and return pointer to end
- `strlen.c` --- return min(length of string, n)
- `strnmov.c` --- move source to dest for source size, or for n bytes
- `strrchr.c` --- find a character within string, searching from end
- `strstr.c` --- find an instance of pattern within source
- `strto.c` --- string to long, to long long, to unsigned long, etc.
- `strtod.c` --- string to double
- `strtol.c` --- string to long
- `strtoll.c` --- string to long long
- `strtoul.c` --- string to unsigned long
- `strtoull.c` --- string to unsigned long long
- `strxmov.c` --- move a series of concatenated source strings to dest
- `strxnmov.c` --- like `strxmov.c` but with a maximum length n
- `str_test.c` --- test of all the string functions encoded in assembler
- `uca-dump.c` --- shows unicode collation algorithm dump
- `udiv.c` --- unsigned long divide, for operating systems that don't support these
- `utr11-dump.c` --- dump east Asian wide text file
- `xml.c` --- read and parse XML strings; used to read character definition information stored in `/sql/share/charsets`

There are also four .ASM files --- `macros.asm`, `ptr_cmp.asm`, `strings.asm`, and `strxmov.asm` --- which can replace some of the C-program functions. But again, they look like optimizations for old members of the Intel processor family.

The [support-files](#) Directory

Files used to build MySQL on different systems.

The files here are for building ("making") MySQL given a package manager, compiler, linker, and other build tools. The support files provide instructions and switches for the build processes. They include example `my.cnf` files one can use as a default setup for MySQL.

The [tests](#) Directory

Tests in Perl and in C.

The files in this directory are test programs that can be used as a base to write a program to simulate problems in MySQL in various scenarios: forks, locks, big records, exporting, truncating, and so on. Some examples are:

- `connect_test.c` --- test that a connect is possible
- `insert_test.c` --- test that an insert is possible
- `list_test.c` --- test that a select is possible

- `select_test.c` --- test that a select is possible
- `showdb_test.c` --- test that a show-databases is possible
- `ssl_test.c` --- test that SSL is possible
- `thread_test.c` --- test that threading is possible

The `tools` Directory

Tools --- well, actually, one tool.

The only file is:

- `mysqlmanager.c` --- A "server management daemon" by Sasha Pachev. This is a tool under development and is not yet useful. Related to fail-safe replication.

The `VC++Files` Directory

Visual C++ Files.

Includes this entire directory, repeated for VC++ (Windows) use.

VC++Files includes a complete environment to compile MySQL with the VC++ compiler. To use it, just copy the files on this directory; the `make_win_src_distribution.sh` script uses these files to create a Windows source installation.

This directory has subdirectories which are copies of the main directories. For example, there is a subdirectory `\VC++Files\heap`, which has the Microsoft developer studio project file to compile `\heap` with VC++. So for a description of the files in `\VC++Files\heap`, see the description of the files in `\heap`. The same applies for almost all of VC++Files's subdirectories (`bdb`, `client`, `isam`, `libmysql`, etc.). The difference is that the `\VC++Files` variants are specifically for compilation with Microsoft Visual C++ in 32-bit Windows environments.

In addition to the "subdirectories which are duplicates of directories", VC++Files contains these subdirectories, which are not duplicates:

- `comp_err` --- (nearly empty)
- `contrib` --- (nearly empty)
- `InstallShield` --- script files
- `isamchk` --- (nearly empty)
- `libmysqltest` --- one small non-MySQL test program: `mytest.c`
- `myisamchk` --- (nearly empty)
- `myisamlog` --- (nearly empty)
- `myisammrg` --- (nearly empty)
- `mysqlbinlog` --- (nearly empty)
- `mysqlmanager` --- MFC foundation class files created by AppWizard
- `mysqlserver` --- (nearly empty)
- `mysqlshutdown` --- one short program, `mysqlshutdown.c`
- `mysqlwatch.c` --- Windows service initialization and monitoring

- my_print_defaults --- (nearly empty)
- pack_isam --- (nearly empty)
- perror --- (nearly empty)
- prepare --- (nearly empty)
- replace --- (nearly empty)
- SCCS --- source code control system
- test1 --- tests connecting via X threads
- thr_insert_test --- (nearly empty)
- thr_test --- one short program used to test for memory-allocation bug
- winmysqladmin --- the winmysqladmin.exe source

The "nearly empty" subdirectories noted above (e.g. comp_err and isamchk) are needed because VC ++ requires one directory per project (i.e. executable). We are trying to keep to the MySQL standard source layout and compile only to different directories.

The **vio** Directory

Virtual I/O Library.

The VIO routines are wrappers for the various network I/O calls that happen with different protocols. The idea is that in the main modules one won't have to write separate bits of code for each protocol. Thus vio's purpose is somewhat like the purpose of Microsoft's winsock library.

The underlying protocols at this moment are: TCP/IP, Named Pipes (for WindowsNT), Shared Memory, and Secure Sockets (SSL).

The C programs are:

- test-ssl.c --- Short standalone test program: SSL
- test-sslclient.c --- Short standalone test program: clients
- test-sslserver.c --- Short standalone test program: server
- vio.c --- Declarations + open/close functions
- viosocket.c --- Send/retrieve functions
- viossl.c --- SSL variations for the above
- viosslfactories.c --- Certification / Verification
- viotest.cc --- Short standalone test program: general
- viotest-ssl.c --- Short standalone test program: SSL
- viotest-sslconnect.cc --- Short standalone test program: SSL connect

The older functions --- raw_net_read, raw_net_write --- are now obsolete.

The **zlib** Directory

Data compression library, used on Windows.

`zlib` is a data compression library used to support the compressed protocol and the COMPRESS/UNCOMPRESS functions under Windows. On Unix, MySQL uses the system `libgz.a` library for this purpose.

`Zlib` --- which presumably stands for "Zip Library" --- is not a MySQL package. It was produced by the GNU Zip (gzip.org) people. `Zlib` is a variation of the famous "Lempel-Ziv" method, which is also used by "Zip". The method for reducing the size of any arbitrary string of bytes is as follows:

- Find a substring which occurs twice in the string.
- Replace the second occurrence of the substring with (a) a pointer to the first occurrence, plus (b) an indication of the length of the first occurrence.

There is a full description of the library's functions in the gzip manual at <http://www.gzip.org/zlib/manual.html>. There is therefore no need to list the modules in this document.

The MySQL program `\mysys\my_compress.c` uses `zlib` for packet compression. The client sends messages to the server which are compressed by `zlib`. See also: `\sql\net_serv.cc`.

Appendix B. InnoDB Source Code Distribution

The InnoDB source files are the best place to look for information about internals of the file structure that MySQLers can optionally use for transaction support. But when you first look at all the subdirectories and file names you'll wonder: Where Do I Start? It can be daunting.

Well, I've been through that phase, so I'll pass on what I had to learn on the first day that I looked at InnoDB source files. I am very sure that this will help you grasp, in overview, the organization of InnoDB modules. I'm also going to add comments about what is going on -- which you should mistrust! These comments are reasonable working hypotheses; nevertheless, they have not been subjected to expert peer review.

Here's how I'm going to organize the discussion. I'll take each of the 31 InnoDB subdirectories that come with the MySQL 5.0 source code in `\innobase` (on my Windows directory). The format of each section will be like this every time:

`\subdirectory-name (LONGER EXPLANATORY NAME)`

File Name	What Name Stands For	Size	Comment Inside File
file-name	my-own-guess	in-bytes	from-the-file-itself

... My-Comments

For example:

```
"\ha (HASHING)
  File Name   What Name Stands For   Size      Comment Inside File
  -----
  ha0ha.c     Hashing/Hashing                8,145     Hash table with external chains

  Comments about hashing will be here.
"
```

The "Comment Inside File" column is a direct copy from the first `/* comment */` line inside the file. All other comments are mine. After I've discussed each directory, I'll finish with some notes about naming conventions and a short list of URLs that you can use for further reference.

Now let's begin.

`\btr (B-TREE)`

File Name	What Name Stands For	Size	Comment Inside File
-----	-----	-----	-----
btr0btr.c	B-tree / B-tree	82,400	B-tree
btr0cur.c	B-tree / Cursor	103,233	index tree cursor
btr0sea.c	B-tree / Search	41,788	index tree adaptive search
btr0pcur.c	B-tree / persistent cursor	16,720	index tree persistent cursor

If you total up the sizes of the C files, you'll see that `\btr` is the second-largest file group in InnoDB. This is understandable because maintaining a B-tree is a relatively complex task. Luckily, there has been a lot of work done to describe efficient management of B-tree and B+-tree structures, much of it open-source or public-domain, since their original invention over thirty years ago.

InnoDB likes to put everything in B-trees. This is what I'd call a "distinguishing characteristic" because in all the major DBMSs (like IBM DB2, Microsoft SQL Server, and Oracle), the main or default or classic structure is the heap-and-index. In InnoDB the main structure is just the index. To put it another way: InnoDB keeps the rows in the leaf node of the index, rather than in a separate file. Compare Oracle's Index Organized Tables, and Microsoft SQL Server's Clustered Indexes.

This, by the way, has some consequences. For example, you may as well have a primary key since otherwise InnoDB will make one anyway. And that primary key should be the shortest of the candidate keys, since InnoDB will use it as a pointer if there are secondary indexes.

Most importantly, it means that rows have no fixed address. Therefore the routines for managing file pages should be good. We'll see about that when we look at the \row (ROW) program group later.

\buf (BUFFERING)

File Name	What Name Stands For	Size	Comment Inside File
-----	-----	-----	-----
buf0buf.c	Buffering / Buffering	65,582	The database buffer buf_pool
buf0flu.c	Buffering / Flush	29,583	... flush algorithm
buf0lru.c	/ least-recently-used	27,515	... replacement algorithm
buf0rea.c	Buffering / read	21,504	... read

There is a separate file group (\mem MEMORY) which handles memory requests in general. A "buffer" usually has a more specific definition, as a memory area which contains copies of pages that ordinarily are in the main data file. The "buffer pool" is the set of all buffers (there are lots of them because InnoDB doesn't depend on the operating system's caching to make things faster).

The pool size is fixed (at the time of this writing) but the rest of the buffering architecture is sophisticated, involving a host of control structures. In general: when InnoDB needs to access a new page it looks first in the buffer pool; InnoDB reads from disk to a new buffer when the page isn't there; InnoDB chucks old buffers (basing its decision on a conventional Least-Recently-Used algorithm) when it has to make space for a new buffer.

There are routines for checking a page's validity, and for read-ahead. An example of "read-ahead" use: if a sequential scan is going on, then a DBMS can read more than one page at a time, which is efficient because reading 32,768 bytes (two pages) takes less than twice as long as reading 16,384 bytes (one page).

\data (DATA)

File Name	What Name Stands For	Size	Comment Inside File
-----	-----	-----	-----
data0data.c	Data / Data	15,344	SQL data field and tuple
data0type.c	Data / Type	7,417	Data types

This is a collection of minor utility routines affecting rows.

\db (DATABASE)

There are no .c files in \db, just one .h file with some definitions for error codes.

\dict (DICTIONARY)

File Name	What Name Stands For	Size	Comment Inside File
-----	-----	-----	-----
dict0dict.c	Dictionary / Dictionary	114,263	Data dictionary system
dict0boot.c	Dictionary / boot	11,704	... booting
dict0crea.c	Dictionary / Create	37,278	... creation
dict0load.c	Dictionary / load	34,049	... load to memory cache
dict0mem.c	Dictionary / memory	7,470	... memory object creation

The data dictionary (known in some circles as the catalog) has the metadata information about objects in the database --- column sizes, table names, and the like.

\dyn (DYNAMICALLY ALLOCATED ARRAY)

File Name	What Name Stands For	Size	Comment Inside File
-----	-----	-----	-----
dyn0dyn.c	Dynamic / Dynamic	994	dynamically allocated array

There is a single function in the dyn0dyn.c program, for adding a block to the dynamically allocated array. InnoDB might use the array for managing concurrency between threads.

At the moment, the \dyn program group is trivial.

\eval (EVALUATING)

File Name	What Name Stands For	Size	Comment Inside File
-----	-----	-----	-----
eval0eval.c	Evaluating/Evaluating	17,061	SQL evaluator
eval0proc.c	Evaluating/Procedures	5,001	Executes SQL procedures

The evaluating step is a late part of the process of interpreting an SQL statement --- parsing has already occurred during \pars (PARSING).

The ability to execute SQL stored procedures is an InnoDB feature, but MySQL handles stored procedures in its own way, so the eval0proc.c program is unimportant.

\fil (FILE)

File Name	What Name Stands For	Size	Comment Inside File
-----	-----	-----	-----
fil0fil.c	File / File	118,312	The low-level file system

The reads and writes to the database files happen here, in coordination with the low-level file i/o routines (see os0file.c in the \os program group).

Briefly: a table's contents are in pages, which are in files, which are in tablespaces. Files do not grow; instead one can add new files to the tablespace. As we saw earlier (discussing the \btr program group) the pages are nodes of B-trees. Since that's the case, new additions can happen at various places in the logical file structure, not necessarily at the end. Reads and writes are asynchronous, and go into buffers, which are set up by routines in the \buf program group.

\fsp (FILE SPACE)

File Name	What Name Stands For	Size	Comment Inside File
-----	-----	-----	-----
fsp0fsp.c	File Space Management	110,495	File space management

I would have thought that the \fil (FILE) and \fsp (FILE SPACE) MANAGEMENT programs would fit together in the same program group; however, I guess the InnoDB folk are splitters rather than lumpers.

It's in fsp0fsp.c that one finds some of the descriptions and comments of extents, segments, and headers. For example, the "descriptor bitmap of the pages in the extent" is in here, and you can find as well how the free-page list is maintained, what's in the bitmaps, and what various header fields' contents are.

\fut (FILE UTILITY)

File Name	What Name Stands For	Size	Comment Inside File
-----	-----	-----	-----
fut0fut.c	File Utility / Utility	293	File-based utilities
fut0lst.c	File Utility / List	14,176	File-based list utilities

Mainly these small programs affect only file-based lists, so maybe saying "File Utility" is too generic. The real work with data files goes on in the \fsp program group.

\ha (HASHING)

File Name	What Name Stands For	Size	Comment Inside File
-----	-----	-----	-----
ha0ha.c	Hashing / Hashing	8,145	Hash table with external chains
hash0hash.c	Hashing / Hashing	3,283	Simple hash table utility

The two C programs in the \ha directory --- ha0ha.c and hash0hash.c --- both refer to a "hash table" but hash0hash.c is specialized, it is mostly about accessing points in the table under mutex control.

When a "database" is so small that InnoDB can load it all into memory at once, it's more efficient to access it via a hash table. After all, no disk i/o can be saved by using an index lookup, if there's no disk.

\ibuf (INSERT BUFFER)

File Name	What Name Stands For	Size	Comment Inside File
-----	-----	-----	-----
ibuf0ibuf.c	Insert Buffer /	91,397	Insert buffer

The words "Insert Buffer" mean not "buffer used for INSERT" but "insertion of a buffer into the buffer pool" (see the `ibuf BUFFER` program group description). The matter is complex due to possibilities for deadlocks, a problem to which the comments in the `ibuf0ibuf.c` program devote considerable attention.

\include (INCLUDE)

All `.h` and `.ic` files are in the `\include` directory. It's habitual to put comments along with the descriptions, so if (for example) you want to see comments about operating system file activity, the place to look is `\include\os0file.h`.

\lock (LOCKING)

File Name	What Name Stands For	Size	Comment Inside File
-----	-----	-----	-----
lock0lock.c	Lock / Lock	139,207	The transaction lock system

If you've used DB2 or SQL Server, you might think that locks have their own in-memory table, that row locks might need occasional escalation to table locks, and that there are three lock types: Shared, Update, Exclusive.

All those things are untrue with **InnoDB**! Locks are kept in the database pages. A bunch of row locks can't be rolled together into a single table lock. And most importantly there's only one lock type. I call this type "Update" because it has the characteristics of DB2 / SQL Server Update locks, that is, it blocks other updates but doesn't block reads. Unfortunately, **InnoDB** comments refer to them as "x-locks" etc.

To sum it up: if your background is Oracle you won't find too much surprising, but if your background is DB2 or SQL Server the locking concepts and terminology will probably confuse you at first.

You can find my online article about the differences between Oracle-style and DB2/SQL-Server-style locks at: <http://dbazine.com/gulutzan6.html>

Now here is a notice from Heikki Tuuri of InnoDB. It concerns lock categories rather than `lock0lock.c`, but I place it in this section because this is the place that people are most likely to look for it.

Errata notice about **InnoDB** row locks:

```
#define LOCK_S 4 /* shared */
#define LOCK_X 5 /* exclusive */
.../* Waiting lock flag */
#define LOCK_WAIT 256
/* this wait bit should be so high that it can be ORed to the lock
mode and type; when this bit is set, it means that the lock has not
yet been granted, it is just waiting for its turn in the wait queue */
...
/* Precise modes */
#define LOCK_ORDINARY 0
/* this flag denotes an ordinary next-key lock in contrast to LOCK_GAP
or LOCK_REC_NOT_GAP */
#define LOCK_GAP 512
/* this gap bit should be so high that it can be ORed to the other
flags; when this bit is set, it means that the lock holds only on the
gap before the record; for instance, an x-lock on the gap does not
give permission to modify the record on which the bit is set; locks of
this type are created when records are removed from the index chain of
records */
#define LOCK_REC_NOT_GAP 1024
/* this bit means that the lock is only on the index record and does
NOT block inserts to the gap before the index record; this is used in
the case when we retrieve a record with a unique key, and is also used
in locking plain SELECTs (not part of UPDATE or DELETE) when the user
has set the READ COMMITTED isolation level */
```

```
#define LOCK_INSERT_INTENTION 2048
/* this bit is set when we place a waiting gap type record lock
request in order to let an insert of an index record to wait until
there are no conflicting locks by other transactions on the gap; note
that this flag remains set when the waiting lock is granted, or if the
lock is inherited to a neighboring record */
```

Errata notice about [InnoDB](#) row locks ends.

\log (LOGGING)

File Name	What Name Stands For	Size	Comment Inside File
log0log.c	Logging / Logging	86,043	Database log
log0recv.c	Logging / Recovery	91,352	Recovery

I've already written about the \log program group, so here's a link to my previous article: "How Logs work with MySQL and InnoDB": <http://www.devarticles.com/c/a/MySQL/How-Logs-Work-On-MySQL-With-InnoDB-Tables>

\mach (MACHINE FORMAT)

File Name	What Name Stands For	Size	Comment Inside File
mach0data.c	Machine/Data	2,335	Utilities for converting

The mach0data.c program has two small routines for reading compressed ulints (unsigned long integers).

\mem (MEMORY)

File Name	What Name Stands For	Size	Comment Inside File
mem0mem.c	Memory / Memory	10,310	The memory management
mem0dbg.c	Memory / Debug	22,054	... the debug code
mem0pool.c	Memory / Pool	16,511	... the lowest level

There is a long comment at the start of the mem0pool.c program, which explains what the memory-consumers are, and how InnoDB tries to satisfy them. The main thing to know is that there are really three pools: the buffer pool (see the \buf program group), the log pool (see the \log program group), and the common pool, which is where everything that's not in the buffer or log pools goes (for example the parsed SQL statements and the data dictionary cache).

\mtr (MINI-TRANSACTION)

File Name	What Name Stands For	Size	Comment Inside File
mtr0mtr.c	Mini-transaction /	12,620	Mini-transaction buffer
mtr0log.c	Mini-transaction / Log	8,090	... log routines

The mini-transaction routines are called from most of the other program groups. I'd describe this as a low-level utility set.

\os (OPERATING SYSTEM)

File Name	What Name Stands For	Size	Comment Inside File
os0file.c	OS / File	104,081	To i/o primitives
os0thread.c	OS / Thread	7,754	To thread control primitives
os0proc.c	OS / Process	16,919	To process control primitives
os0sync.c	OS / Synchronization	14,256	To synchronization primitives

This is a group of utilities that other modules may call whenever they want to use an operating-system resource. For example, in os0file.c there is a public InnoDB function named `os_file_create_simple()`, which simply calls the Windows-API function `CreateFile`. Naturally the call is within an `"#ifdef __WIN__ ... #endif"` block; the effective routines are somewhat different for other operating systems.

\page (PAGE)

File Name	What Name Stands For	Size	Comment Inside File
-----	-----	-----	-----
page0page.c	Page / Page	51,731	Index page routines
page0cur.c	Page / Cursor	38,127	The page cursor

It's in the [page0page.c](#) program that you'll learn as follows: index pages start with a header, entries in the page are in order, at the end of the page is a sparse "page directory" (what I would have called a slot table) which makes binary searches easier.

Incidentally, the program comments refer to "a page size of 8 kB" which seems obsolete. In [univ.i](#) (a file containing universal constants) the page size is now `#defined` as 16KB.

\pars (PARSING)

File Name	What Name Stands For	Size	Comment Inside File
-----	-----	-----	-----
pars0pars.c	Parsing/Parsing	45,376	SQL parser
pars0grm.c	Parsing/Grammar	62,685	A Bison parser
pars0opt.c	Parsing/Optimizer	31,268	Simple SQL Optimizer
pars0sym.c	Parsing/Symbol Table	5,239	SQL parser symbol table
lexyy.c	Parsing/Lexer	62,071	Lexical scanner

The job is to input a string containing an SQL statement and output an in-memory parse tree. The EVALUATING (subdirectory `\eval`) programs will use the tree.

As is common practice, the Bison and Flex tools were used --- [pars0grm.c](#) is what the Bison parser produced from an original file named [pars0grm.y](#) (also supplied), and [lexyy.c](#) is what Flex produced.

Since [InnoDB](#) is a DBMS by itself, it's natural to find SQL parsing in it. But in the MySQL/InnoDB combination, MySQL handles most of the parsing. These files are unimportant.

\que (QUERY GRAPH)

File Name	What Name Stands For	Size	Comment Inside File
-----	-----	-----	-----
que0que.c	Query Graph / Query	30,774	Query graph

The program [que0que.c](#) ostensibly is about the execution of stored procedures which contain commit/rollback statements. I took it that this has little importance for the average MySQL user.

\read (READ)

File Name	What Name Stands For	Size	Comment Inside File
-----	-----	-----	-----
read0read.c	Read / Read	9,935	Cursor read

The [read0read.c](#) program opens a "read view" of a query result, using some functions in the `\trx` program group.

\rem (RECORD MANAGER)

File Name	What Name Stands For	Size	Comment Inside File
-----	-----	-----	-----
rem0rec.c	Record Manager	38,573	Record Manager
rem0cmp.c	Record Manager / Comparison	26,617	Comparison services for records

There's an extensive comment near the start of [rem0rec.c](#) title "Physical Record" and it's recommended reading. At some point you'll ask what are all those bits that surround the data in the rows on a page, and this is where you'll find the answer.

\row (ROW)

File Name	What Name Stands For	Size	Comment Inside File
row0row.c	Row / Row	18,375	General row routines
row0uins.c	Row / Undo Insert	6,799	Fresh insert undo
row0umod.c	Row / Undo Modify	19,712	Undo modify of a row
row0undo.c	Row / Undo	10,512	Row undo
row0vers.c	Row / Version	14,385	Row versions
row0mysql.c	Row / MySQL	112,462	Interface [to MySQL]
row0ins.c	Row / Insert	42,829	Insert into a table
row0sel.c	Row / Select	111,719	Select
row0upd.c	Row / Update	51,824	Update of a row
row0purge.c	Row / Purge	15,609	Purge obsolete records

Rows can be selected, inserted, updated/deleted, or purged (a maintenance activity). These actions cause following actions, for example after insert there can be an index-update test, but it seems to me that sometimes the following action has no MySQL equivalent (yet) and so is inoperative.

Speaking of MySQL, notice that one of the larger programs in the \row program group is the "interface between Innobase row operations and MySQL" (row0mysql.c) --- information interchange happens at this level because rows in InnoDB and in MySQL are analogous, something which can't be said for pages and other levels.

\srv (Server)

File Name	What Name Stands For	Size	Comment Inside File
srv0srv.c	Server / Server	75,633	Server main program
srv0que.c	Server / Query	2,463	Server query execution
srv0start.c	Server / Start	50,154	Starts the server

This is where the server reads the initial configuration files, splits up the threads, and gets going. There is a long comment deep in the program (you might miss it at first glance) titled "IMPLEMENTATION OF THE SERVER MAIN PROGRAM" in which you'll find explanations about thread priority, and about what the responsibilities are for various thread types.

InnoDB has many threads, for example "user threads" (which wait for client requests and reply to them), "parallel communication threads" (which take part of a user thread's job if a query process can be split), "utility threads" (background priority), and a "master thread" (high priority, usually asleep).

\sync (SYNCHRONIZATION)

File Name	What Name Stands For	Size	Comment Inside File
sync0sync.c	Synchronization /	37,940	Mutex, the basic sync primitive
sync0arr.c	... / array	26,455	Wait array used in primitives
sync0rw.c	... / read-write	22,846	read-write lock for thread sync

A mutex (Mutual Exclusion) is an object which only one thread/process can hold at a time. Any modern operating system API has some functions for mutexes; however, as the comments in the sync0sync.c code indicate, it can be faster to write one's own low-level mechanism. In fact the old assembly-language XCHG trick is in sync0sync.c's helper file, \include\sync0sync.ic. This is the only program that contains any assembly code.

The i/o and thread-control primitives are called extensively. The word "synchronization" in this context refers to the mutex-create and mutex-wait functionality.

\thr (Thread Local Storage)

File Name	What Name Stands For	Size	Comment Inside File
thr0loc.c	Thread / Local	5,334	The thread local storage

InnoDB doesn't use the Windows-API thread-local-storage functions, perhaps because they're not portable enough.

\trx (Transaction)

File Name	What Name Stands For	Size	Comment Inside File
-----	-----	-----	-----
trx0trx.c	Transaction /	50,480	The transaction
trx0purge.c	Transaction / Purge	29,133	... Purge old versions
trx0rec.c	Transaction / Record	37,346	... Undo log record
trx0roll.c	/ Rollback	31,448	... Rollback
trx0sys.c	Transaction / System	27,018	... System
trx0rseg.c	/ Rollback segment	6,445	... Rollback segment
trx0undo.c	Transaction / Undo	51,519	... Undo log

InnoDB's transaction management is supposedly "in the style of Oracle" and that's close to true but can mislead you.

- First: InnoDB uses rollback segments like Oracle8i does — but Oracle9i uses a different name.
- Second: InnoDB uses multi-versioning like Oracle does — but I see nothing that looks like an Oracle ITL being stored in the InnoDB data pages.
- Third: InnoDB and Oracle both have short (back-to-statement-start) versioning for the [READ COMMITTED](#) isolation level and long (back-to-transaction-start) versioning for higher levels — but InnoDB and Oracle have different "default" isolation levels.
- Finally: InnoDB's documentation says it has to lock "the gaps before index keys" to prevent phantoms — but any Oracle user will tell you that phantoms are impossible anyway at the [SERIALIZABLE](#) isolation level, so key-locks are unnecessary.

The main idea, though, is that InnoDB has multi-versioning. So does Oracle. This is very different from the way that DB2 and SQL Server do things.

\usr (USER)

File Name	What Name Stands For	Size	Comment Inside File
-----	-----	-----	-----
usr0sess.c	User / Session	1,740	Sessions

One user can have multiple sessions (the session being all the things that happen between a connect and disconnect). This is where InnoDB used to track session IDs, and server/client messaging. It's another of those items which is usually MySQL's job, though. So now usr0sess.c merely closes.

\ut (UTILITIES)

File Name	What Name Stands For	Size	Comment Inside File
-----	-----	-----	-----
ut0ut.c	Utilities / Utilities	9,728	Various utilities
ut0byte.c	Utilities / Debug	793	Byte utilities
ut0rnd.c	Utilities / Random	1,474	Random numbers and hashing
ut0mem.c	Utilities / Memory	10,358	Memory primitives
ut0dbg.c	Utilities / Debug	2,579	Debug utilities

The two functions in ut0byte.c are just for lower/upper case conversion and comparison. The single function in ut0rnd.c is for finding a prime slightly greater than the given argument, which is useful for hash functions, but unrelated to randomness. The functions in ut0mem.c are wrappers for "malloc" and "free" calls — for the real "memory" module see section \mem (MEMORY). Finally, the functions in ut0ut.c are a miscellany that didn't fit better elsewhere: `get_high_bytes`, `clock`, `time`, `difftime`, `get_year_month_day`, and `sprintf` for various diagnostic purposes.

In short: the \ut group is trivial.

This is the end of the section-by-section account of InnoDB subdirectories.

Some Notes About Structures

InnoDB's job, as a storage engine for MySQL, is to provide: commit-rollback, crash recovery, row-level locking, and consistent non-blocking reads. How? With locks, a paged-file structure with buffer pooling, and undo/redo logs,

The locks are kept in bit maps in main memory. Thus InnoDB differs from Oracle in one respect: instead of storing lock information on the page as Oracle does with Interested Transaction Lists, InnoDB keeps it in a separate and more volatile structure. But both Oracle and InnoDB try to achieve a similar goal: "writers don't block readers". So a typical InnoDB row-read involves: (a) if the reading is for writing, then check if the row is locked and if so wait; (b) if according to the information in the row header the row has been changed by some newer transaction, then get the older version from the log. We call the (b) part "versioning" because it means that a reader can get the older version of a row and thus will have a temporally consistent view of all rows.

The InnoDB workspace consists of: tablespace and log files. A tablespace consists of: segments, as many as necessary. A segment is usually a file, but might be a raw disk partition. A segment consists of: extents. An extent consists of: 64 pages. A page's length is always 16KB, for both data and index. A page consists of: a page header, and some rows. The page and row formats are the subjects of later chapters.

InnoDB keeps two logs, the redo log and the undo log.

The redo log is for re-doing data changes that had not been written to disk when a crash occurred. There is one redo log for the entire workspace, it contains multiple files (the number depends on `innodb_log_files_in_group`), it is circular (that is, after writing to the last file InnoDB starts again on the first file). The file header includes the last successful checkpoint. A redo log record's contents are: Page Number (4 bytes = page number within tablespace), Offset of change within page (2 bytes), Log Record Type (insert, update, delete, "fill space with blanks", etc.), and the changes on that page (only redo values, not old values).

The undo log is primarily for removing data changes that had been written to disk when a crash occurred, but should not have been written, because they were for uncommitted transactions. Sometimes InnoDB calls the undo log the "rollback segment". The undo log is inside the tablespace. The "insert" section of the undo log is needed only for transaction rollback and can be discarded at COMMIT time. The "update/delete" section of the undo log is also useful for consistent reads, and can be discarded when InnoDB has ended all transactions that might need the undo log records to reconstruct earlier versions of rows. An undo log record's contents are: Primary Key Value (not a page number or physical address), Old Transaction ID (of the transaction that updated the row), and the changes (only old values).

COMMIT will write the contents of the log buffer to disk, and put undo log records in a history list. ROLLBACK will delete undo log records that are no longer needed. PURGE (an internal operation that occurs outside user control) will no-longer-necessary undo log records and, for data records that have been marked for deletion and are no longer necessary for consistent read, will remove the records. CHECKPOINT causes -- well, see the article "How Logs Work On MySQL With InnoDB Tables".

You might be able to find a slide show, "ACID Transactions in MySQL With InnoDB", via this page: <http://dev.mysql.com/tech-resources/presentations/>. If the links are broken, notify MySQL.

A Note About File Naming

There appears to be a naming convention. The first letters of the file name are the same as the subdirectory name, then there is a '0' separator, then there is an individual name. For the main program in a subdirectory, the individual name may be a repeat of the subdirectory name. For example, there is a file named `ha0ha.c` (the first two letters `ha` mean "it's in subdirectory `..ha`", the next letter `0` means "0 separator", the next two letters mean "this is the main `ha` program"). This naming convention is not strict, though: for example the file `lexyy.c` is in the `\pars` subdirectory.

A Note About Copyrights

Most of the files begin with a copyright notice or a creation date, for example "Created 10/25/1995 Heikki Tuuri". I don't know a great deal about the history of InnoDB, but found it interesting that most creation dates were between 1994 and 1998.

References

-
- Ryan Bannon, Alvin Chin, Faryaz Kassam and Andrew Roszko. "InnoDB Concrete Architecture" http://www.swen.uwaterloo.ca/~mrbannon/cs798/assignment_02/innodb.pdf

A student paper. It's an interesting attempt to figure out InnoDB's architecture using tools, but I didn't end up using it for the specific purposes of this article.

- Peter Gulutzan. "How Logs Work With MySQL And InnoDB" <http://www.devarticles.com/c/a/MySQL/How-Logs-Work-On-MySQL-With-InnoDB-Tables>
- Heikki Tuuri. "InnoDB Engine in MySQL-Max-3.23.54 / MySQL-4.0.9: The Up-to-Date Reference Manual of InnoDB" <http://www.innodb.com/ibman.html>

This is the natural starting point for all InnoDB information. Mr Tuuri also appears frequently on MySQL forums.

Index

C

- clients
 - debugging, 357
- crash, 352

D

- DBI->trace, 355
- DBI_TRACE environment variable, 355
- DEBUG package, 357
- debugging
 - client, 357
 - server, 352

E

- Environment variable
 - DBI_TRACE, 355
 - MYSQL_DEBUG, 357
- error messages, 361
 - defining, 361

F

- filesort optimization, 94

G

- gdb
 - using, 354

M

- MYSQL_DEBUG environment variable, 357

O

- optimizing
 - filesort, 94

R

- RTS-threads, 359

S

- server
 - debugging, 352

T

- thread packages
 - differences between, 360
- threads
 - RTS, 359
- trace DBI method, 355

