



Advanced MySQL Query Tuning

Alexander Rubin
Principal Architect, Percona
April 18, 2015

About Me

My name is Alexander Rubin

- Working with MySQL for over 10 years
 - Started at MySQL AB, then Sun Microsystems,
 - then Oracle (MySQL Consulting)
 - Joined Percona 2 years ago
- Helping customers improve MySQL performance
 - performance tuning
 - full text search
 - high availability

Agenda

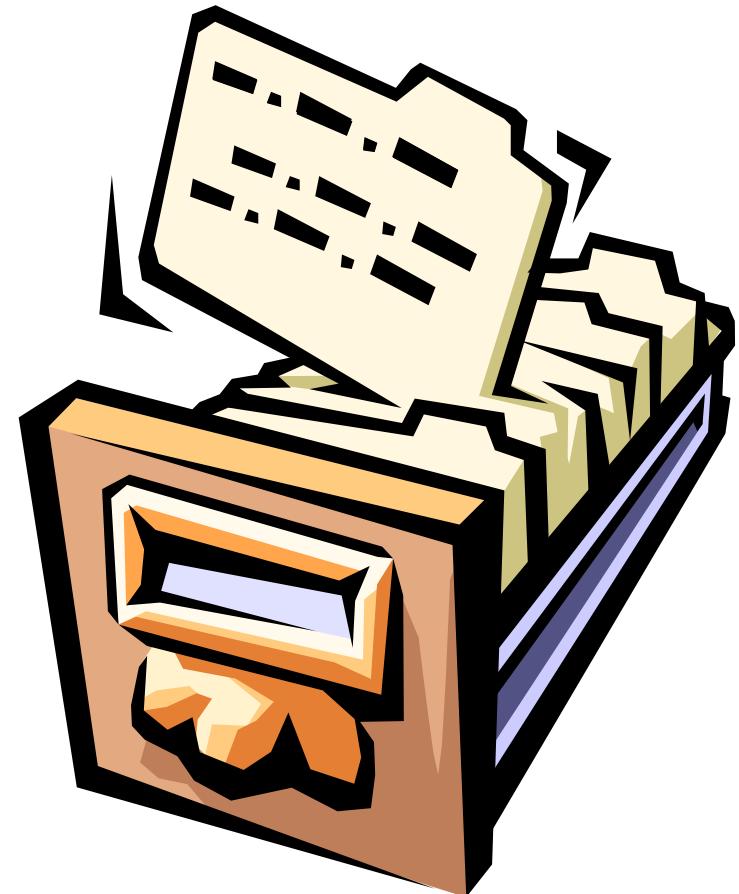
- Queries
 - Temporary Tables and Filesort
 - GROUP BY Optimizations
 - ORDER BY Optimizations
 - DISTINCT Queries
 - “Calculated” fields – MySQL 5.7

New this year

- Removed basic stuff to focus on advanced
- MySQL 5.6 by default
 - Will *not* talk about subqueries issues
- Improvements in MySQL 5.7

Query Tuning: basics

**Indexes and
Explain examples**



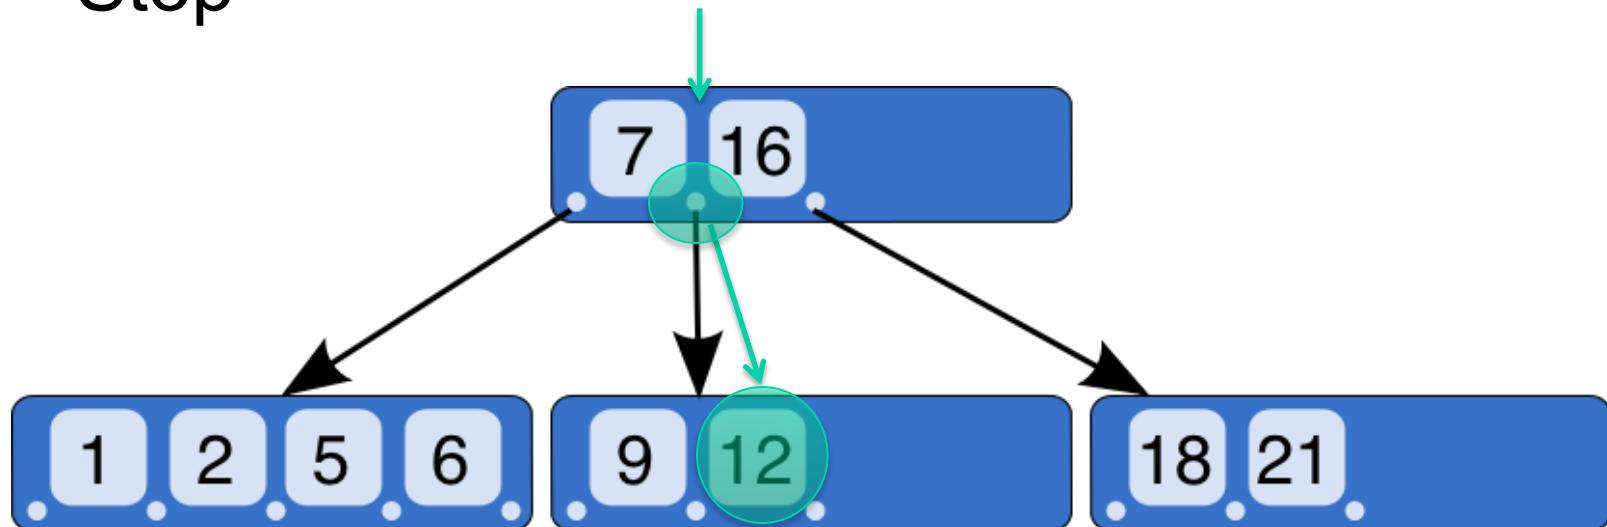
ORACLE

www.percona.com

MySQL Index: B-Tree

Equality search: select * from table where id = 12

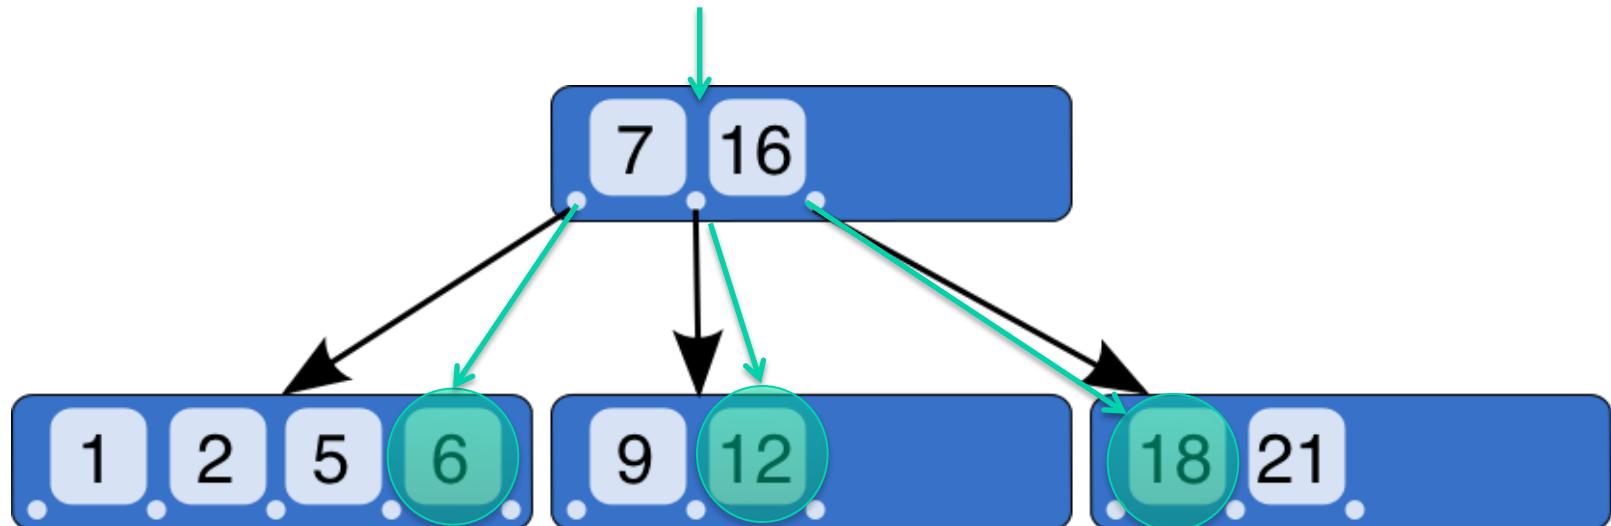
- Scan thru the tree and go directly to 1 leaf
- Stop



MySQL Index: B-Tree

Range: select * from table where id in (6, 12, 18)

- Scan thru the tree and visit many leafs/nodes



Example Table

```
CREATE TABLE City (
    ID int(11) NOT NULL AUTO_INCREMENT,
    Name char(35) NOT NULL DEFAULT '',
    CountryCode char(3) NOT NULL DEFAULT '',
    District char(20) NOT NULL DEFAULT '',
    Population int(11) NOT NULL DEFAULT '0',
    PRIMARY KEY (ID),
    KEY CountryCode (CountryCode)
) Engine=InnoDB;
```

Indexes: Example

- MySQL will use 1 (best) index

```
mysql> explain select * from City where ID = 1;
```

table	type	possible_keys	key	key_len	ref	rows	Extra
City	const	PRIMARY	PRIMARY	4	const	1	

```
mysql> explain select * from City where CountryCode = 'USA';
```

table	type	possible_keys	key	key_len	ref	rows	Extra
City	ref	CountryCode	CountryCode	3	const	274	Using where

Combined Indexes: Example

- Leftmost part of combined index

```
mysql> alter table City add key  
comb(CountryCode, District, Population),  
drop key CountryCode;
```

Combined Indexes: Example

- Leftmost part of combined index

```
mysql> explain select * from City
      where CountryCode = 'USA' \G
*****
   1. row *****
table: City
type: ref
possible_keys: comb
key: comb
key_len: 3 ←
ref: const
rows: 273
```

Uses first field from the comb key

Combined Indexes: Example

- Key_len = total size (in bytes) of index parts used

Index: comb(CountryCode, District, Population)

Explain:

key: comb

key_len: 3

Fields:

CountryCode **char(3)**

District char(20)

Population int(11)

3 -> Char(3) -> First field is used

Combined Indexes: Example

- 2 Leftmost Fields

```
mysql> explain select * from City
where CountryCode = 'USA' and District = 'California'\G
*****
   table: City
   type: ref
possible_keys: comb
      key: comb
key_len: 23
      ref: const,const
     rows: 68
```

Uses 2 first fields from the comb key
CountryCode = 3 chars
District = 20 chars
Total = 23

key_len: 23 ↗

Combined Indexes: Example

- 3 Leftmost Fields

```
mysql> explain select * from City  
where CountryCode = 'USA' and District = 'California'  
and population > 10000\G
```

```
***** 1. row *****
```

```
    table: City  
    type: range  
possible_keys: comb  
    key: comb  
key_len: 27 ←  
    ref: NULL  
rows: 68
```

Uses ***all*** fields from the comb key
CountryCode = 3 chars/bytes
District = 20 chars/bytes
Population = 4 bytes (INT)
Total = 27

Combined Indexes: Example

- Can't use combined index – not a leftmost part

```
mysql> explain select * from City where  
District = 'California' and population > 10000\G  
***** 1. row *****
```

```
table: City  
      type: ALL  
possible_keys: NULL  
      key: NULL ←  
key_len: NULL  
      ref: NULL  
rows: 3868
```

Does not have the ***CountryCode***
in the where clause
= can't use comb index

Covered Index: Example

- Covered index = cover all fields in query

```
select name from City where CountryCode = 'USA'  
and District = 'Alaska' and population > 10000
```

```
mysql> alter table City add key  
cov1(CountryCode, District, population, name);
```



Uses **all** fields in the query in particular order:

1. Where part
2. Group By/Order (not used now)
3. Select part (here: **name**)

Covered Index: Example

- Explain

```
mysql> explain select name from City where CountryCode =  
'USA' and District = 'Alaska' and population > 10000\G  
***** 1. row *****
```

table: City

type: range

possible_keys: cov1

key: cov1

key_len: 27

ref: NULL

rows: 1

Extra: Using where; Using index

Using index = covered index is used

MySQL will only use index
Will not go to the data file



Query tuning example

Calculated Expressions In MySQL Queries

Air Traffic Statistics table

```
CREATE TABLE `ontime` (
  `id` int(11) NOT NULL AUTO_INCREMENT,
  `YearD` year(4) NOT NULL,
  `FlightDate` date DEFAULT NULL,
  `Carrier` char(2) DEFAULT NULL,
  `OriginAirportID` int(11) DEFAULT NULL,
  `OriginCityName` varchar(100) DEFAULT NULL,
  `OriginState` char(2) DEFAULT NULL,
  `DestAirportID` int(11) DEFAULT NULL,
  `DestCityName` varchar(100) DEFAULT NULL,
  `DestState` char(2) DEFAULT NULL,
  `DepDelayMinutes` int(11) DEFAULT NULL,
  `ArrDelayMinutes` int(11) DEFAULT NULL,
  `Cancelled` tinyint(4) DEFAULT NULL,
  `CancellationCode` char(1) DEFAULT NULL,
  `Diverted` tinyint(4) DEFAULT NULL,
  PRIMARY KEY (`id`),
  KEY `FlightDate` (`FlightDate`)
) ENGINE=InnoDB
```

All flights in 2013, group by airline

```
mysql> EXPLAIN SELECT carrier, count(*) FROM ontime_sm
      WHERE year(FlightDate) = 2013 group by carrier\G
***** 1. row *****
    id: 1
  select_type: SIMPLE
        table: ontime_sm
         type: ALL
possible_keys: NULL
           key: NULL
      key_len: NULL
         ref: NULL
       rows: 151253427
     Extra: Using where; Using temporary; Using filesort
```

Year() = “calculated expression”
MySQL can't use index

Results:

```
16 rows in set (1 min 49.48 sec)
```

Rewritten: flights in 2013, group by airline

```
mysql> EXPLAIN SELECT carrier, count(*) FROM ontime_sm  
WHERE FlightDate between '2013-01-01' and '2013-12-31'  
GROUP BY carrier\G
```

```
***** 1. row *****  
      id: 1  
select_type: SIMPLE  
      table: ontime_sm  
      type: range  
possible_keys: FlightDate  
          key: FlightDate  
     key_len: 4  
       ref: NULL  
      rows: 10434762  
    Extra: Using index condition; Using temporary; Using filesort
```

Results:

16 rows in set (11.98 sec)

RANGE condition
MySQL can use index

Almost 10x faster

All flights on Sunday example

```
mysql> EXPLAIN SELECT carrier, count(*) FROM ontime_sm
      WHERE dayofweek(FlightDate) = 7 group by carrier\G
***** 1. row *****
      id: 1
  select_type: SIMPLE
        table: ontime_sm
         type: ALL
possible_keys: NULL
           key: NULL
      key_len: NULL
         ref: NULL
        rows: 151253427
    Extra: Using where; Using temporary; Using filesort
```

Year() = “calculated expression”
MySQL can’t use index

Results:

32 rows in set (1 min 57.93 sec)

Storing additional field

```
CREATE TABLE `ontime_sm` (
  `id` int(11) NOT NULL AUTO_INCREMENT,
  `YearD` year(4) NOT NULL,
  `FlightDate` date DEFAULT NULL,
  `Carrier` char(2) DEFAULT NULL,
  `OriginAirportID` int(11) DEFAULT NULL,
  `OriginCityName` varchar(100) DEFAULT NULL,
  `OriginState` char(2) DEFAULT NULL,
  `DestAirportID` int(11) DEFAULT NULL,
  `DestCityName` varchar(100) DEFAULT NULL,
  `DestState` char(2) DEFAULT NULL, DEFAULT
  `DepDelayMinutes` int(11) NULL,
  `ArrDelayMinutes` int(11) DEFAULT NULL,
  `Cancelled` tinyint(4) DEFAULT NULL,
  ...
  Flight_dayofweek tinyint NOT NULL,
  PRIMARY KEY (`id`),
  KEY `Flight_dayofweek` (`Flight_dayofweek`)
) ENGINE=InnoDB DEFAULT CHARSET=latin1;
```

Materialize the field...



Storing additional field ... and populating it with trigger

```
CREATE DEFINER = CURRENT_USER
TRIGGER ontime_insert
BEFORE INSERT ON ontime_sm_triggers
FOR EACH ROW
SET
NEW.Flight_dayofweek = dayofweek(NEW.FlightDate);
```

May be slower on insert

Virtual Columns in MySQL 5.7.7 (labs release)

```
CREATE TABLE `ontime_sm_virtual` (
  `id` int(11) NOT NULL AUTO_INCREMENT,
  `YearD` year(4) NOT NULL,
  `FlightDate` date DEFAULT NULL,
  `Carrier` char(2) DEFAULT NULL,
  `OriginAirportID` int(11) DEFAULT NULL,
  `OriginCityName` varchar(100) DEFAULT NULL,
  `OriginState` char(2) DEFAULT NULL,
  `DestAirportID` int(11) DEFAULT NULL,
  `DestCityName` varchar(100) DEFAULT NULL,
  `DestState` char(2) DEFAULT NULL,
  `DepDelayMinutes` int(11) DEFAULT NULL,
  `ArrDelayMinutes` int(11) DEFAULT NULL,
  `Cancelled` tinyint(4) DEFAULT NULL,
  ...
  `Flight_dayofweek` tinyint(4)
  GENERATED ALWAYS AS (dayofweek(FlightDate)) VIRTUAL,
  PRIMARY KEY (`id`),
  KEY `Flight_dayofweek` (`Flight_dayofweek`)
) ENGINE=InnoDB;
```

Does not store the column
But INDEX it

<http://mysqlserverteam.com/generated-columns-in-mysql-5-7-5/>
<https://dev.mysql.com/worklog/task/?id=8114>
<http://labs.mysql.com/>

Virtual Columns in MySQL 5.7.7 (labs release)

```
mysql> insert into ontime_sm_triggers (id, YearD, FlightDate, Carrier,  
OriginAirportID, OriginCityName, OriginState, DestAirportID, DestCityName,  
DestState, DepDelayMinutes, ArrDelayMinutes, Cancelled,  
CancellationCode,Diverted, CRSElapsedTime, ActualElapsedTime, AirTime,  
Flights, Distance) select * from ontime_sm;
```

```
Query OK, 999999 rows affected (27.86 sec)  
Records: 999999 Duplicates: 0 Warnings: 0
```

```
mysql> insert into ontime_sm_virtual (id, YearD, FlightDate, Carrier,  
OriginAirportID, OriginCityName, OriginState, DestAirportID, DestCityName,  
DestState, DepDelayMinutes, ArrDelayMinutes, Cancelled,  
CancellationCode,Diverted, CRSElapsedTime, ActualElapsedTime, AirTime,  
Flights, Distance) select * from ontime_sm;
```

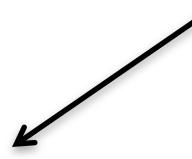
```
Query OK, 999999 rows affected (16.29 sec)  
Records: 999999 Duplicates: 0 Warnings: 0
```

Much faster on load
Compared to triggers

Virtual Columns in MySQL 5.7.7 (labs release)

```
mysql> EXPLAIN SELECT carrier, count(*) FROM ontime_sm_virtual
WHERE Flight_dayofweek = 7 group by carrier\G
***** 1. row *****
      id: 1
select_type: SIMPLE
      table: ontime_sm_virtual
    partitions: NULL
        type: ref
possible_keys: Flight_dayofweek
          key: Flight_dayofweek
      key_len: 2
        ref: const
      rows: 165409
filtered: 100.00
   Extra: Using where; Using temporary; Using filesort
1 row in set, 1 warning (0.00 sec)
```

Using index



Virtual Columns in MySQL 5.7.7

Limitations

```
mysql> alter table ontime_sm_virtual  
    add key comb(Flight_dayofweek, carrier);
```

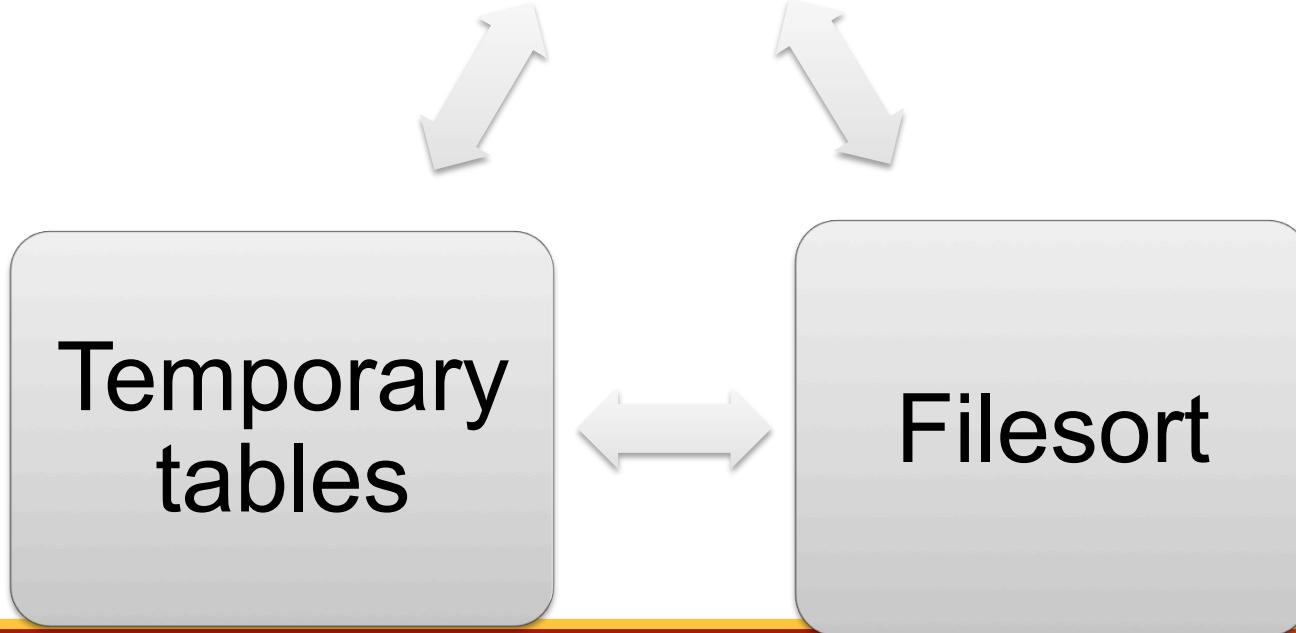
ERROR 3105 (HY000): 'Virtual generated column combines with other columns to be indexed together' is not supported for generated columns.



Can't create a combined index
on virtual and “normal” columns

Complex Slow Queries

... Group By ...
... Order By ...
Select distinct ...



GROUP BY Queries



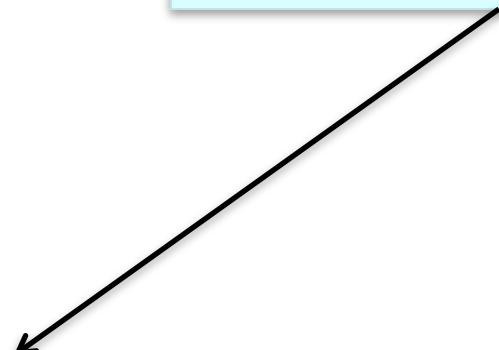
GROUP BY and Temporary Tables

How many cities in each country?

```
mysql> explain select CountryCode, count(*) from City  
group by CountryCode\G
```

```
      id: 1  
select_type: SIMPLE  
      table: City  
        type: ALL  
possible_keys: NULL  
          key: NULL  
key_len: NULL  
        ref: NULL  
      rows: 4079  
Extra: Using temporary; Using filesort
```

Temporary tables are slow!



Temporary Tables: Theory



Temporary Tables, I

Main performance issues

- MySQL can create temporary tables when query uses:
 - GROUP BY
 - Range + ORDER BY
 - Some other expressions

- 2 types of temporary tables
 - MEMORY
 - On-disk

Temporary Tables, II

- First, MySQL tries to create temporary table in memory
- MySQL configuration variables:
 - **tmp_table_size**
 - maximum size for in Memory temporary tables
 - **max_heap_table_size**
 - Sets the maximum size for **MEMORY** tables

Temporary Tables III

MySQL
temp table
 $>$
tmp_table_size

OR

MySQL
temp table
 $>$
max_heap_table_size

=

convert to
MyISAM
temporary
table on
disk

Temporary Tables

- MEMORY engine does not support BLOB/TEXT
- select blob_field from table group by field1
- select concat(...string>512 chars) group by field1
 - Create on-disk temporary table right away

Temporary Tables: Practice



Air Traffic Statistics Table for Testing

6M rows, ~2G in size

```
CREATE TABLE ontime_2012 (
    YearD int(11) DEFAULT NULL,
    MonthD tinyint(4) DEFAULT NULL,
    DayofMonth tinyint(4) DEFAULT NULL,
    DayOfWeek tinyint(4) DEFAULT NULL,
    Carrier char(2) DEFAULT NULL,
    Origin char(5) DEFAULT NULL,
    DepDelayMinutes int(11) DEFAULT NULL,
    ...
) ENGINE=InnoDB DEFAULT CHARSET=latin1
http://www.transtats.bts.gov/DL\_SelectFields.asp?  
Table\_ID=236&DB\_Short\_Name=On-Time
```

GROUP BY Query Example

- Find maximum delay for flights on Sunday
- Group by airline

```
SELECT max(DepDelayMinutes) ,  
carrier, dayofweek  
FROM ontime_2012  
WHERE dayofweek = 7  
GROUP BY Carrier
```

GROUP BY Query Example

```
select max(DepDelayMinutes), carrier, dayofweek  
from ontime_2012  
where dayofweek = 7  
group by Carrier
```

```
    type: ALL  
possible_keys: NULL  
        key: NULL  
key_len: NULL  
        ref: NULL  
rows: 4833086  
Extra: Using where; Using temporary; Using  
filesort
```

Full table scan!
Temporary table!

Adding Index: Fixing Full Table Scan

Better!
!

```
mysql> alter table ontime_2012 add key (dayofweek);  
  
explain select max(DepDelayMinutes), Carrier,  
dayofweek from ontime_2012 where dayofweek =7  
group by Carrier\G
```

```
    type: ref  
possible_keys: DayOfWeek  
      key: DayOfWeek  
key_len: 2  
    ref: const  
   rows: 817258
```

Index is used = better
BUT: Large temporary table!



Extra: Using where; Using temporary; Using filesort

Best
!

GROUP BY: Adding Covered Index

```
mysql> alter table ontime_2012
add key covered(dayofweek, Carrier, DepDelayMinutes);
explain select max(DepDelayMinutes), Carrier, dayofweek from
ontime_2012 where dayofweek =7 group by Carrier\G
...
possible_keys: DayOfWeek,covered
      key: covered
key_len: 2
      ref: const
     rows: 905138
Extra: Using where; Using index
```

No temporary table!
MySQL will only use index



- Called “tight index scan”

When Covered Indexes Aren't Good ...

```
mysql> explain select max(DepDelayMinutes), Carrier,  
dayofweek from ontime_2012  
where dayofweek > 3 group by Carrier, dayofweek\G
```

...

```
type: range ← Range scan  
possible_keys: covered  
key: covered  
key_len: 2  
ref: NULL  
rows: 2441781  
Extra: Using where; Using index; Using temporary;  
Using filesort
```

Converting query to Union

Hack!
!

```
(select max(DepDelayMinutes), Carrier, dayofweek  
from ontime_2012  
where dayofweek = 3  
group by Carrier, dayofweek)  
union  
(select max(DepDelayMinutes), Carrier, dayofweek  
from ontime_2012  
where dayofweek = 4  
group by Carrier, dayofweek)
```

Converting query to Union

Hack!
!

```
***** 1. row *****

```

UNION ALL
Will not create temp table

GROUP BY: Loose index scan

- Loose index scan:
 - considers only a *fraction of the* keys in an index
- Following rules apply:
 - The query is over a single table.
 - The GROUP BY columns should form a leftmost prefix of the index
 - The only aggregate functions = MIN() and MAX(), same column

Loose index scan example

```
mysql> alter table ontime_2012 add key loose_index_scan  
(Carrier, dayofweek, DepDelayMinutes);
```

```
mysql> explain select max(DepDelayMinutes), Carrier,  
    dayofweek from ontime_2012 where dayofweek > 5 group  
by Carrier, dayofweek\G
```

```
...  
    table: ontime_2012  
    type: range  
possible_keys: NULL  
    key: loose_index_scan  
key_len: 5  
    ref: NULL  
    rows: 201  
Extra: Using where; Using index for group-by
```

Loose index scan
Very fast
“Range” works!

Loose index scan vs. tight index scan

Table: `ontime_2012`, 6M rows, data: 2G, index: 210M

```
CREATE TABLE ontime_2012 (
    YearD int(11) DEFAULT NULL,
    MonthD tinyint(4) DEFAULT NULL,
    DayofMonth tinyint(4) DEFAULT NULL,
    DayOfWeek tinyint(4) DEFAULT NULL,
    Carrier char(2) DEFAULT NULL,
    Origin char(5) DEFAULT NULL,
    DepDelayMinutes int(11) DEFAULT NULL,
    ...
    KEY loose_index_scan (Carrier,DayOfWeek,DepDelayMinutes),
    KEY covered (DayOfWeek,Carrier,DepDelayMinutes)
) ENGINE=InnoDB DEFAULT CHARSET=latin1
```

Loose index scan vs. tight index scan

Loose index scan

```
select max(DepDelayMinutes) as ddm, Carrier, dayofweek from ontime_2012  
where dayofweek = 5 group by Carrier, dayofweek
```

```
table: ontime_2012  
type: range  
possible_keys: covered  
key: loose_index_scan  
key_len: 5  
ref: NULL  
rows: 201  
Extra: Using where; Using index for group-by
```

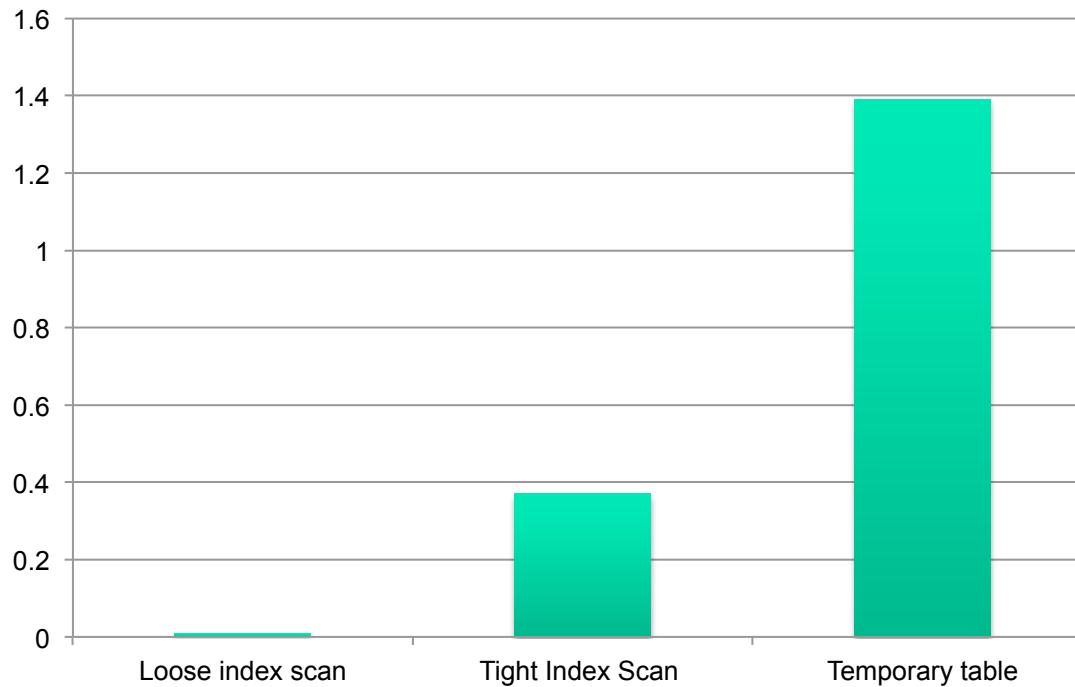
Carrier,
DayOfWeek,
DepDelayMinutes



```
mysql> select ...  
+-----+-----+-----+  
| ddm | Carrier | dayofweek |  
+-----+-----+-----+  
| 1606 | AA      |        7 |  
..  
30 rows in set (0.00 sec)
```

Loose index scan vs. tight index scan

Results



Where loose index scan is not supported

- AVG() + Group By – loose index scan is not supported

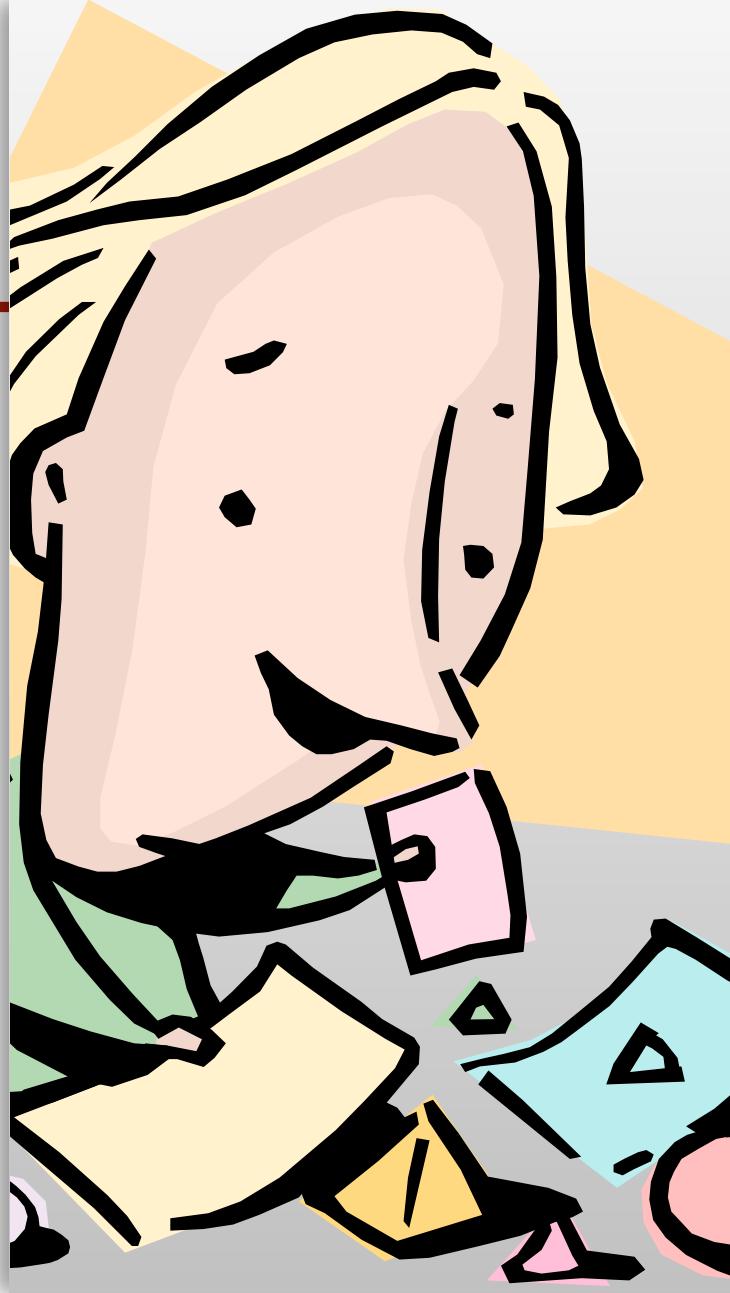
```
mysql> explain select avg(DepDelayMinutes) as ddm, Carrier, dayofweek  
from ontime_2012 where dayofweek >5 group by Carrier, dayofweek \G
```

```
table: ontime_2012  
type: range  
key: covered  
key_len: 2  
ref: NULL  
rows: 2961617  
Extra: Using where; Using index; Using temporary; Using filesort
```

1. No loose index scan
2. Filter by key
3. Group by filesort

```
mysql> select ...  
+-----+-----+-----+  
| ddm   | Carrier | dayofweek |  
+-----+-----+-----+  
| 10.8564 | AA      |       6 |  
...  
30 rows in set (1.39 sec)
```

ORDER BY and filesort



ORDER BY and filesort

Find 10 cities in the US with the largest population

```
mysql> explain select district, name, population from City  
where CountryCode = 'USA' order by population desc limit 10\G
```

```
    table: City  
      type: ALL  
possible_keys: NULL  
          key: NULL  
key_len: NULL  
        ref: NULL  
      rows: 4079
```

Extra: Using where; *Using filesort*

Fixing Filesort: Adding Index

```
mysql> alter table City  
add key my_sort2 (CountryCode, population);
```

```
mysql> explain select district, name, population from City  
where CountryCode = 'USA' order by population desc limit 10\G
```

```
table: City  
type: ref  
key: my_sort2  
key_len: 3  
ref: const  
rows: 207  
Extra: Using where
```

No filesort

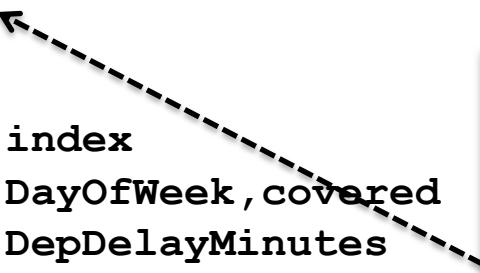


Sorting and Limit

```
mysql> alter table ontime_2012 add key (DepDelayMinutes);
Query OK, 0 rows affected (38.68 sec)
```

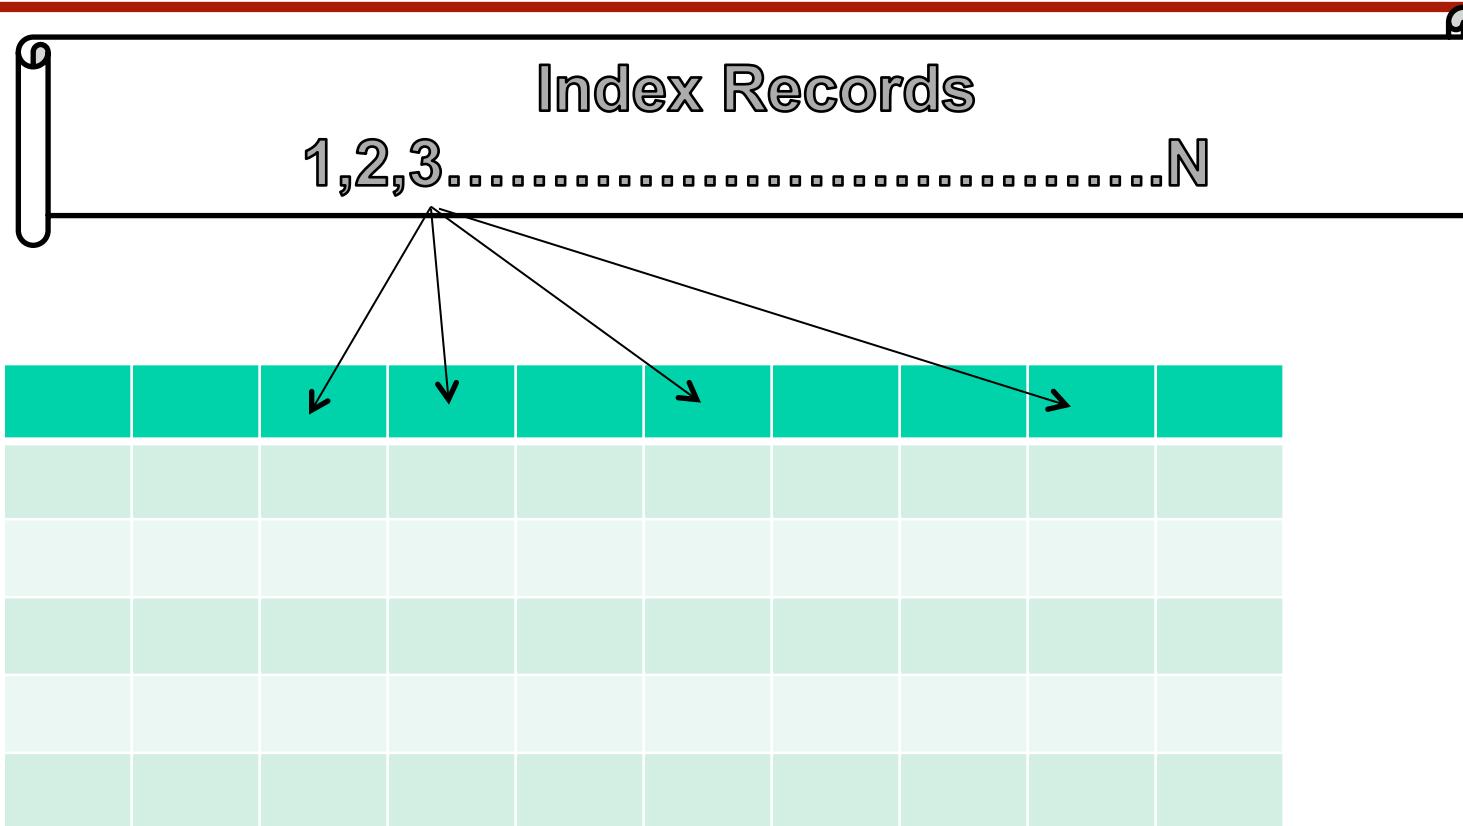
```
mysql> explain select * from ontime_2012
where dayofweek in (6,7) order by DepDelayMinutes desc
limit 10\G
```

```
      type: index
possible_keys: DayOfWeek,covered
      key: DepDelayMinutes
    key_len: 5
      ref: NULL
     rows: 24
   Extra: Using where
```

- 
1. Index is sorted
 2. Scan the **whole table** in the order of the index
 3. Filter results
 4. Stop after finding 10 rows matching the “where” condition

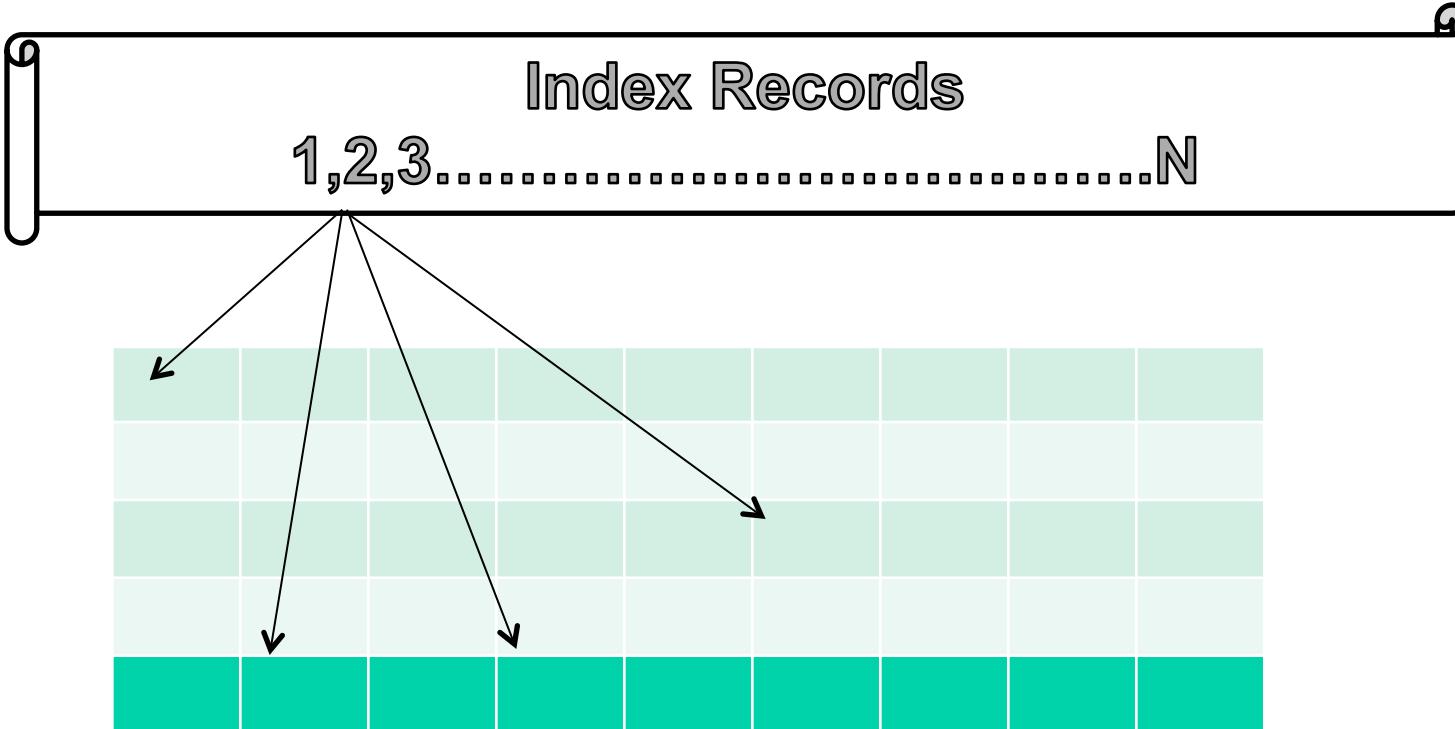
```
10 rows in set (0.00 sec)
```

Sorting and Limit



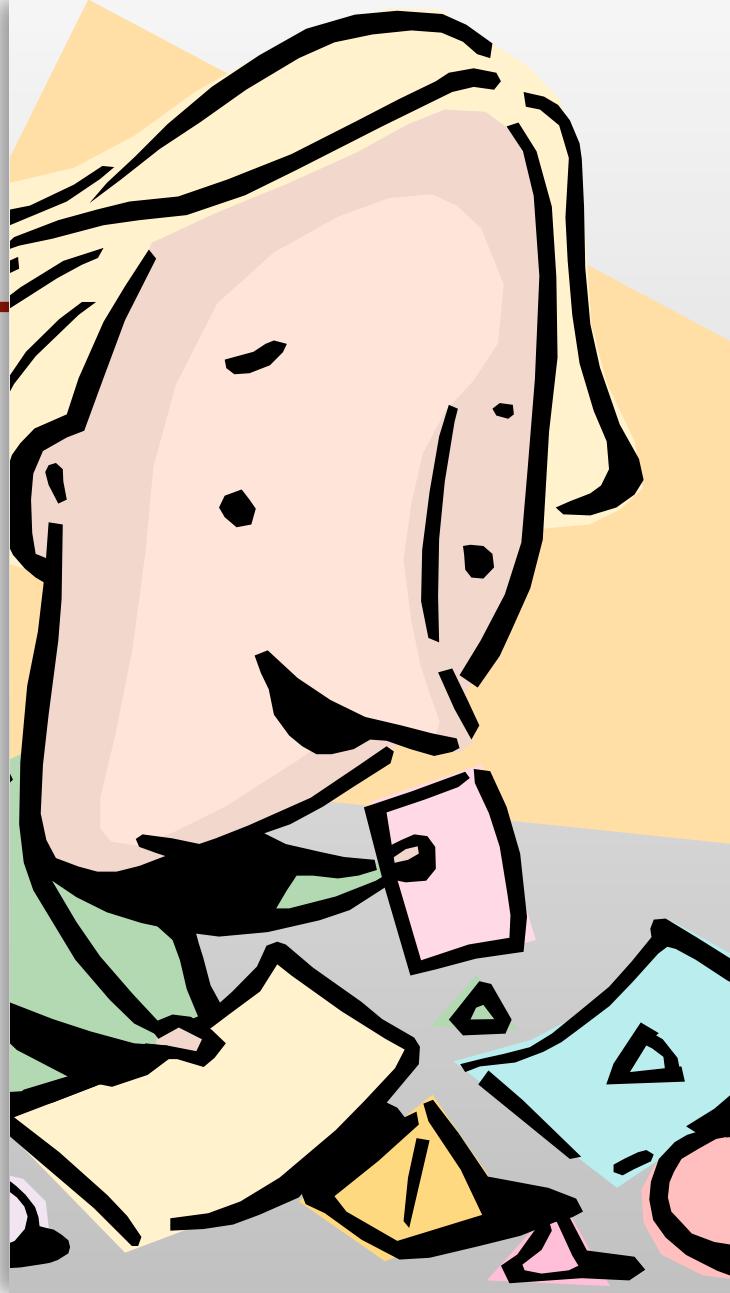
If Index points to the beginning of the table (physically) = fast
As it stops after 10 rows (LIMIT 10)

Sorting and Limit



If Index points to the end of table (physically) or random = slower
Much more rows to scan (and skip)

SELECT DISTINCT



DISTINCT example

```
select t.* , c.name from title t
join movie_companies m on t.id = m.movie_id
join company_name c on m.company_id = c.id
join company_type ct on m.company_type_id = ct.id
where
production_year > 1960 and ct.kind <> 'distributors'
order by production_year desc limit 10\G
...
10 rows in set (0.00 sec)
```

JOINS are for filtering
Result set may contain duplicates
DISTINCT?

DISTINCT example

```
select DISTINCT t.* , c.name from title t
join movie_companies m on t.id = m.movie_id
join company_name c on m.company_id = c.id
join company_type ct on m.company_type_id = ct.id
where
production_year > 1960 and ct.kind <> 'distributors'
order by production_year desc limit 10\G
...
10 rows in set (37.27 sec)
```

DISTINCT example

```
***** 1. row *****
    id: 1
select_type: SIMPLE
    table: t
      type: range
possible_keys: PRIMARY,production_year
      key: production_year
key_len: 5
      ref: NULL
     rows: 1163888
   Extra: Using index condition; Using temporary
***** 2. row *****
    id: 1
select_type: SIMPLE
    table: m
...
...
```

DISTINCT example

```
mysql> show status like 'created%';
+-----+-----+
| Variable_name          | Value |
+-----+-----+
| Created_tmp_disk_tables | 1      |
| Created_tmp_files      | 0      |
| Created_tmp_tables     | 1      |
+-----+-----+
3 rows in set (0.00 sec)
```

DISTINCT example: Hack

```
select distinct * from
(select t.*, c.name from title t
 join movie_companies m on t.id = m.movie_id
 join company_name c on m.company_id = c.id
 join company_type ct on m.company_type_id = ct.id
 where production_year > 1960
 and ct.kind <> 'distributors'
 order by production_year desc limit 1000)
as a limit 10;
...
10 rows in set (0.02 sec)
```

Questions?



Thank you!