



# MySQL 101

*Designing effective schema for InnoDB*

*Yves Trudeau*

*April 2015*



PERCONA  
LIVE



## About myself : Yves Trudeau


- *Principal architect at Percona since 2009*
- *With MySQL then Sun, 2007 to 2009*
- *Focus on MySQL HA and distributed systems*
- *Database and science background*




# Plan

- *What's so special about InnoDB?*
- *Design of a web file sharing application*

# A brief introduction to InnoDB Internals



# Life cycle of an update query





# Where's the data in InnoDB?

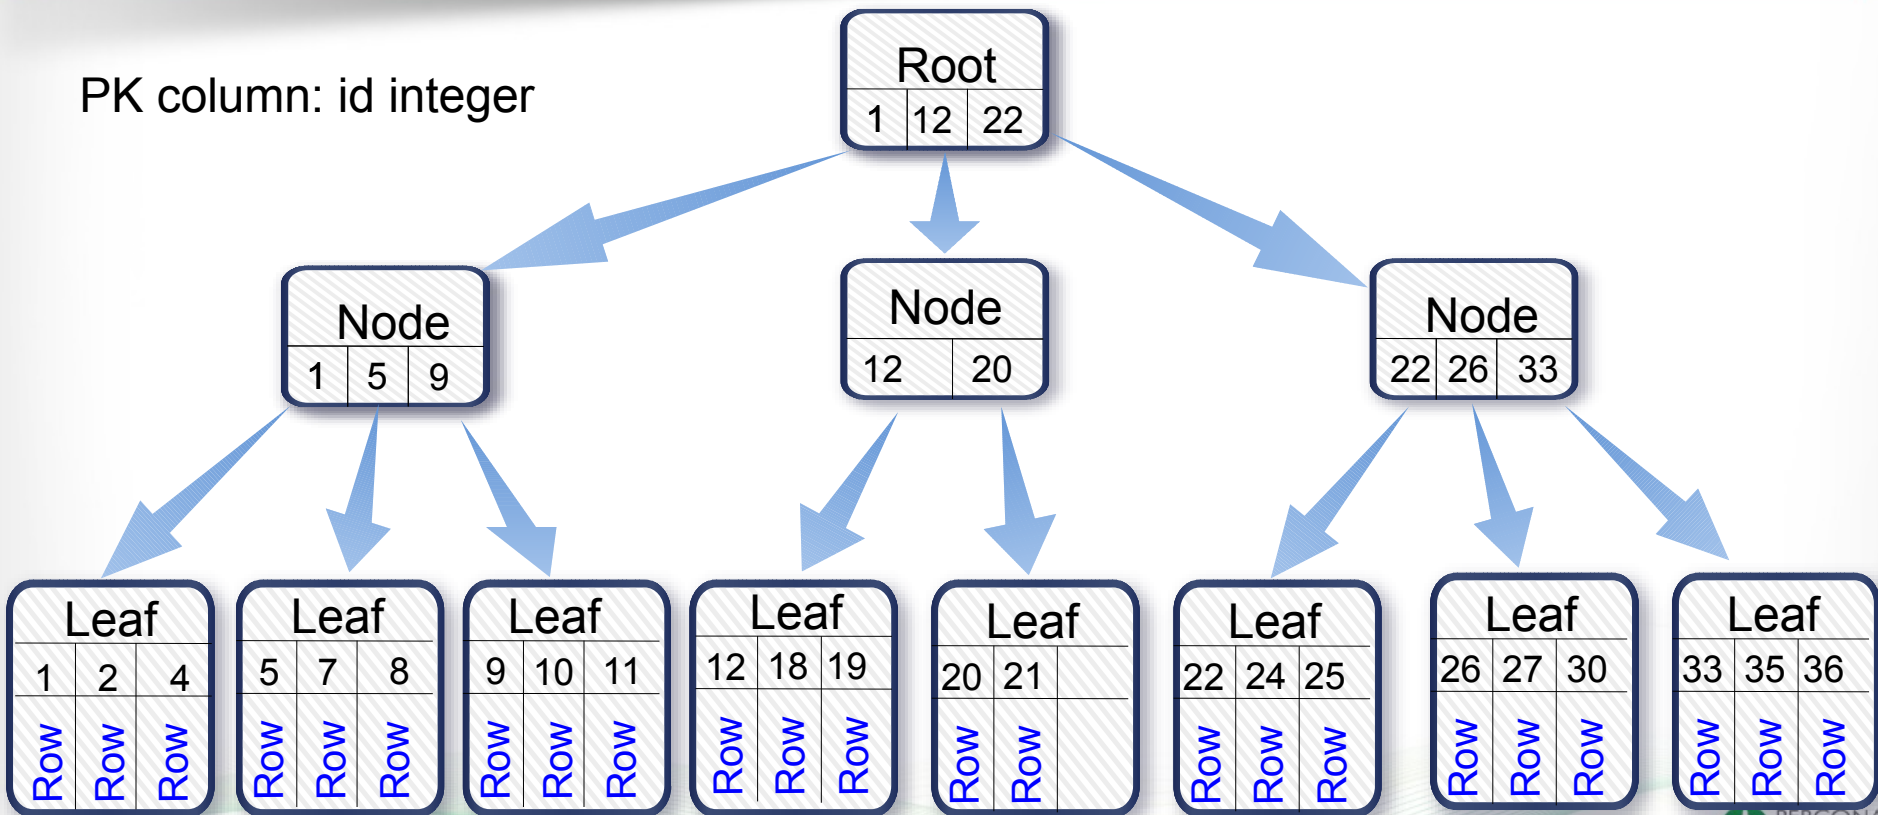
- *The rows are stored as values in the B-tree of the primary key*
- *The secondary keys store as values the primary keys of the matching rows*

*Can't be true*

*I don't have PKs and it works!!*

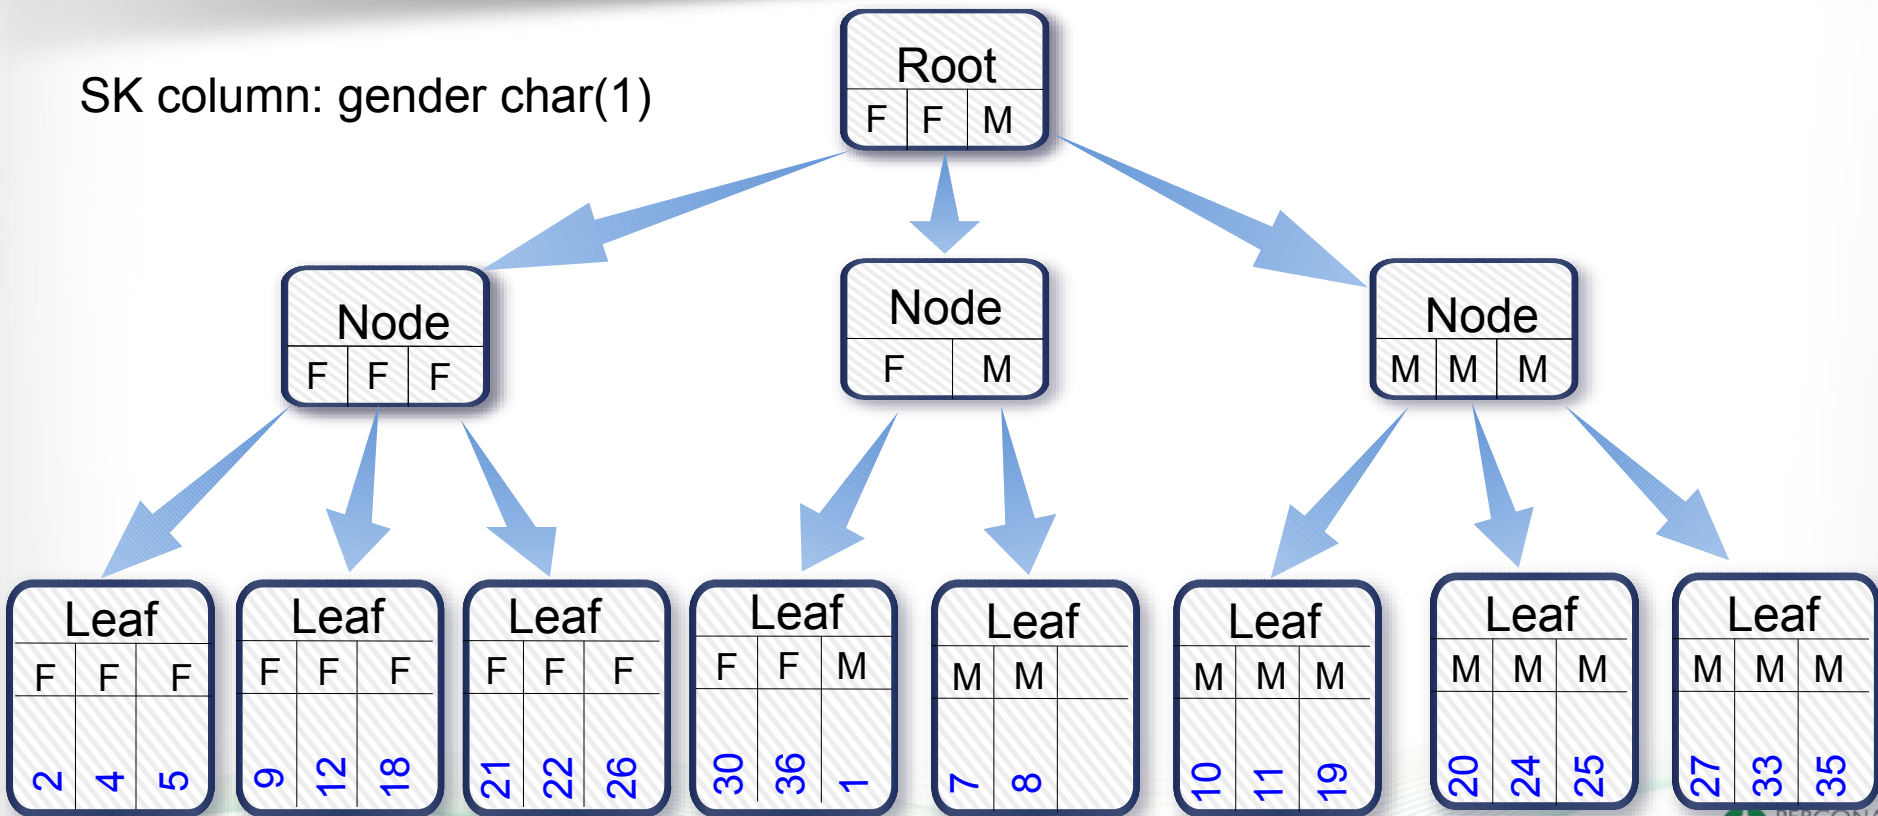
# The primary key B-tree

PK column: id integer



# A secondary key B-tree

SK column: gender char(1)



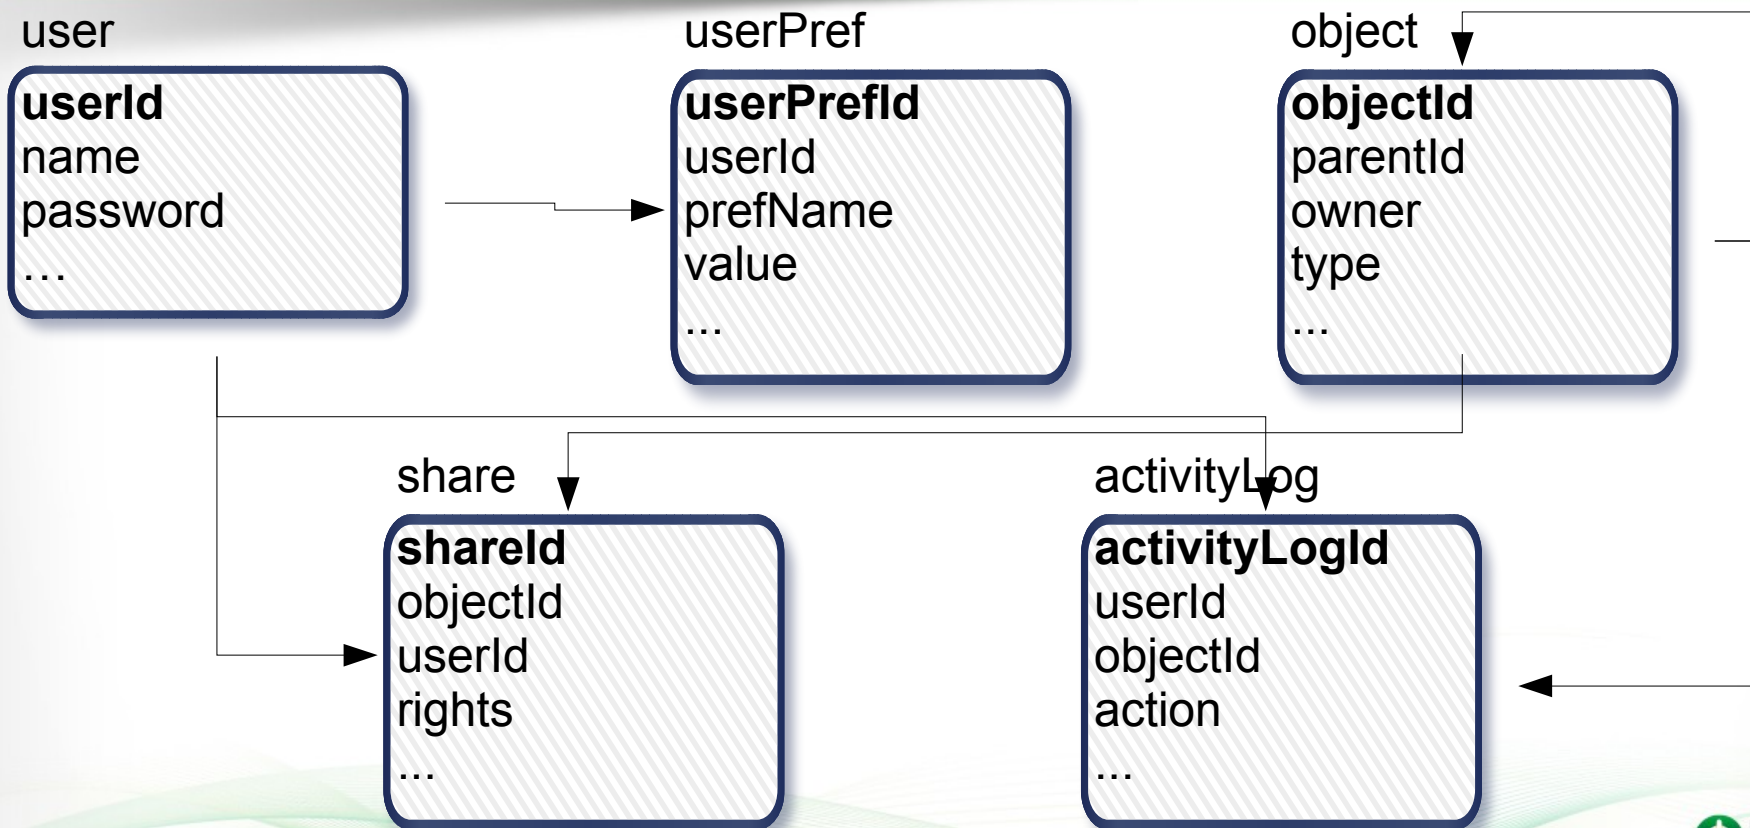


# Enough about InnoDB internals...





# NewBox application, schema v1





# NewBox application, table user

```
CREATE TABLE `user` (  
  `userId` char(36) NOT NULL,  
  `name` varchar(255) DEFAULT NULL,  
  `password` char(32) DEFAULT NULL,  
  `srvSchema` varchar(20) DEFAULT  
NULL,  
  `email` varchar(255) DEFAULT NULL,  
  `updatedAt` datetime DEFAULT NULL,  
  `createdAt` datetime DEFAULT NULL,  
  `lastLogin` datetime DEFAULT NULL,  
  `gender` char(1) DEFAULT NULL,
```

```
PRIMARY KEY (`userId`),  
KEY `idx_name` (`name`),  
KEY `idx_password` (`password`)  
) ENGINE=InnoDB DEFAULT  
CHARSET=utf8
```



# NewBox application, table userPref

```
CREATE TABLE `userPref` (  
  `userPrefId` char(36) NOT NULL,  
  `userId` char(36) DEFAULT NULL,  
  `prefName` varchar(255) DEFAULT  
NULL,  
  `value` varchar(255) DEFAULT NULL,  
  `updateddAt` datetime DEFAULT  
NULL,  
  `createdAt` datetime DEFAULT NULL,
```

```
PRIMARY KEY (`userPrefId`),  
KEY `idx_userId` (`userId`  
) ENGINE=InnoDB DEFAULT  
CHARSET=utf8
```



# NewBox application, table object

```
CREATE TABLE `object` (  
  `objectId` char(36) NOT NULL,  
  `parentId` char(36) DEFAULT NULL,  
  `ownerId` char(36) DEFAULT NULL,  
  `type` varchar(20) DEFAULT NULL,  
  `name` varchar(255) DEFAULT NULL,  
  `urlStore` varchar(255) DEFAULT  
NULL,  
  `version` int(11) DEFAULT NULL,  
  `deleted` int(11) DEFAULT NULL,  
  `updatedAt` datetime DEFAULT NULL,  
  `createdAt` datetime DEFAULT NULL,
```

```
PRIMARY KEY (`objectId`),  
KEY `idx_owner` (`ownerId`),  
KEY `idx_name` (`name`),  
KEY `idx_urlStore` (`urlStore`),  
KEY `idx_deleted` (`deleted`)  
) ENGINE=InnoDB DEFAULT  
CHARSET=utf8
```





# NewBox application, table share

```
CREATE TABLE `share` (  
  `shareId` char(36) NOT NULL,  
  `objectId` char(36) DEFAULT NULL,  
  `userId` char(36) DEFAULT NULL,  
  `ownerId` char(36) DEFAULT NULL,  
  `rights` varchar(20) DEFAULT NULL,  
  `updatedAt` datetime DEFAULT NULL,  
  `createdAt` datetime DEFAULT NULL,
```

```
  PRIMARY KEY (`shareId`),  
  KEY `idx_user` (`userId`)  
) ENGINE=InnoDB DEFAULT  
  CHARSET=latin1
```

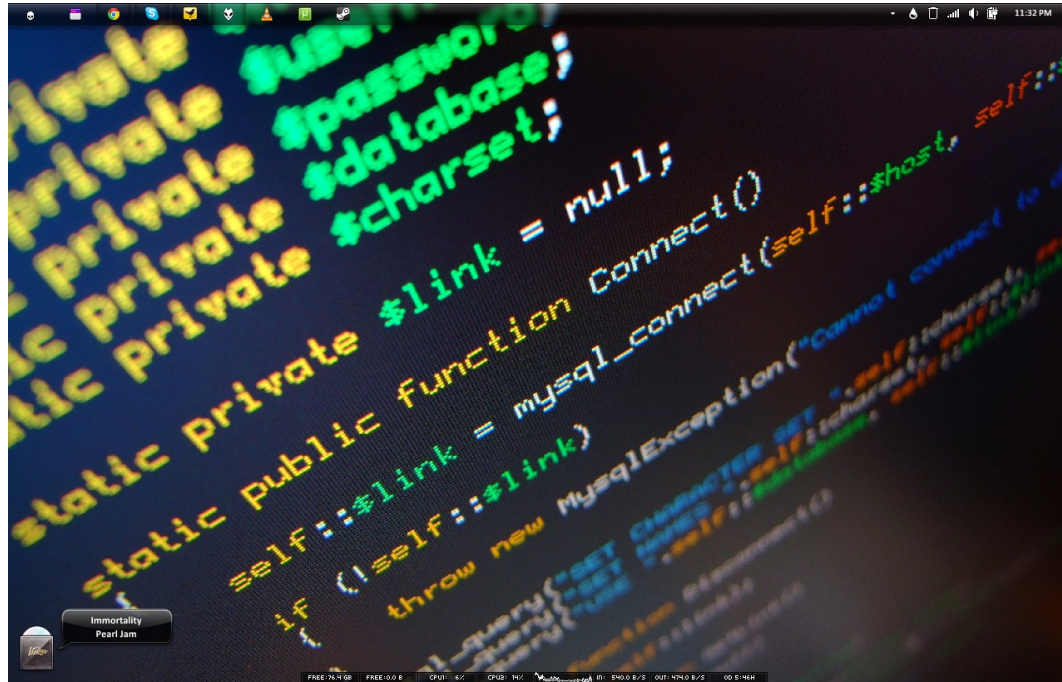


# NewBox application, table activityLog

```
CREATE TABLE `activityLog` (  
  `ActivityId` char(36) NOT NULL,  
  `userId` char(36) DEFAULT NULL,  
  `objectId` char(36) DEFAULT NULL,  
  `action` varchar(255) DEFAULT NULL,  
  `returnCode` int(11) DEFAULT NULL,  
  `error` varchar(255) DEFAULT NULL,  
  `IP` varchar(16) DEFAULT NULL,  
  `createdAt` datetime DEFAULT NULL,
```

```
  PRIMARY KEY (`ActivityId`),  
  KEY `idx_user` (`userId`),  
  KEY `idx_object` (`objectId`),  
  KEY `idx_created` (`createdAt`)  
) ENGINE=InnoDB DEFAULT  
CHARSET=utf8
```

# NewBox application, Coding



A screenshot of a code editor window displaying PHP code for a MySQL connection. The code is color-coded: keywords like 'static', 'public', 'function', 'if', and 'throw' are in blue; variables like '\$link' and '\$database' are in green; and strings and function names like 'mysql\_connect' and 'MySQLException' are in red. The code includes a static private variable for the database link, a public function to connect to the database, and an if-statement to handle connection errors by throwing a MySQLException. The editor's interface includes a top toolbar with icons for file operations, a bottom status bar showing system metrics (free space, CPU usage, memory usage), and a small media player window in the bottom-left corner.

```
private $database;  
private $password;  
private $charset;  
static private $link = null;  
public function Connect()  
{  
    self::$link = mysql_connect(self::$host, self::$  
    if (!self::$link)  
        throw new MySQLException("Cannot connect to  
        query("SET CHARACTER SET " . self::$charset . "  
        query("SET NAMES " . self::$charset . "  
        query("USE " . self::$database . "  
        return self::$link;  
    }  
}
```

# NewBox application, pilot testing



# NewBox application, stage load test







# NewBox application, what's wrong?

- *Dataset is bigger than expected*
- *Database uses more CPU*
- *Database becomes slow when buffer pool is full*
- *Got lockings contention and even deadlocks!!!*
- *Disks are very busy*

# NewBox application, what can we do?

- *More RAM?*
- *Faster drives?*
- *Shard earlier/more?*
- *Maybe my schema isn't that great...*





# NewBox application, review of the data types

## Importance of using the correct types

- *Optimal size = more data in cache*
- *Less reads and writes to disk*
- *Faster comparisons (less CPU)*



# NewBox application, review of the data types

## char with utf8

- *char type uses 3 bytes per char!!!*
- *uuid columns are thus char(108)*
- *keys on uuid columns with uuid pk are 216 bytes per entry*
- *change to varchar or use latin1 for the columns*



# NewBox application, review of the data types

## varchar with utf8

- *Why varchar(255)?*
- *a second length byte after 85*
- *Use proper length or stop at 85*





# NewBox application, review of the data types

## low cardinality columns

- *object.type* → {file, folder, link}
- *userPref.prefName* → {theme, itemPerPage, defaultSort, etc}
- Use *ENUM* or a dictionary table



# NewBox application, review of the data types

## Datetime

- *Arbitrary date and time*
- *8 bytes with 5.5.x, 5 bytes with 5.6.4+*
- *Timestamp ok for [1970,2036]*
- *Use timestamp*



# NewBox application, review of the data types

## Int types

- *Use the correct type → object.deleted tinyint*
- *No negative → unsigned*
- *bigint... is big*
- *using int unsigned for IPs (inet\_aton and inet\_ntoa functions)*



# NewBox application, review of the data types

## Blob/text types

- *Split storage → overlay page*
- *More iops per row*
- *More on disk temp tables for queries (join/sort/group)*
- *Use compression if possible*



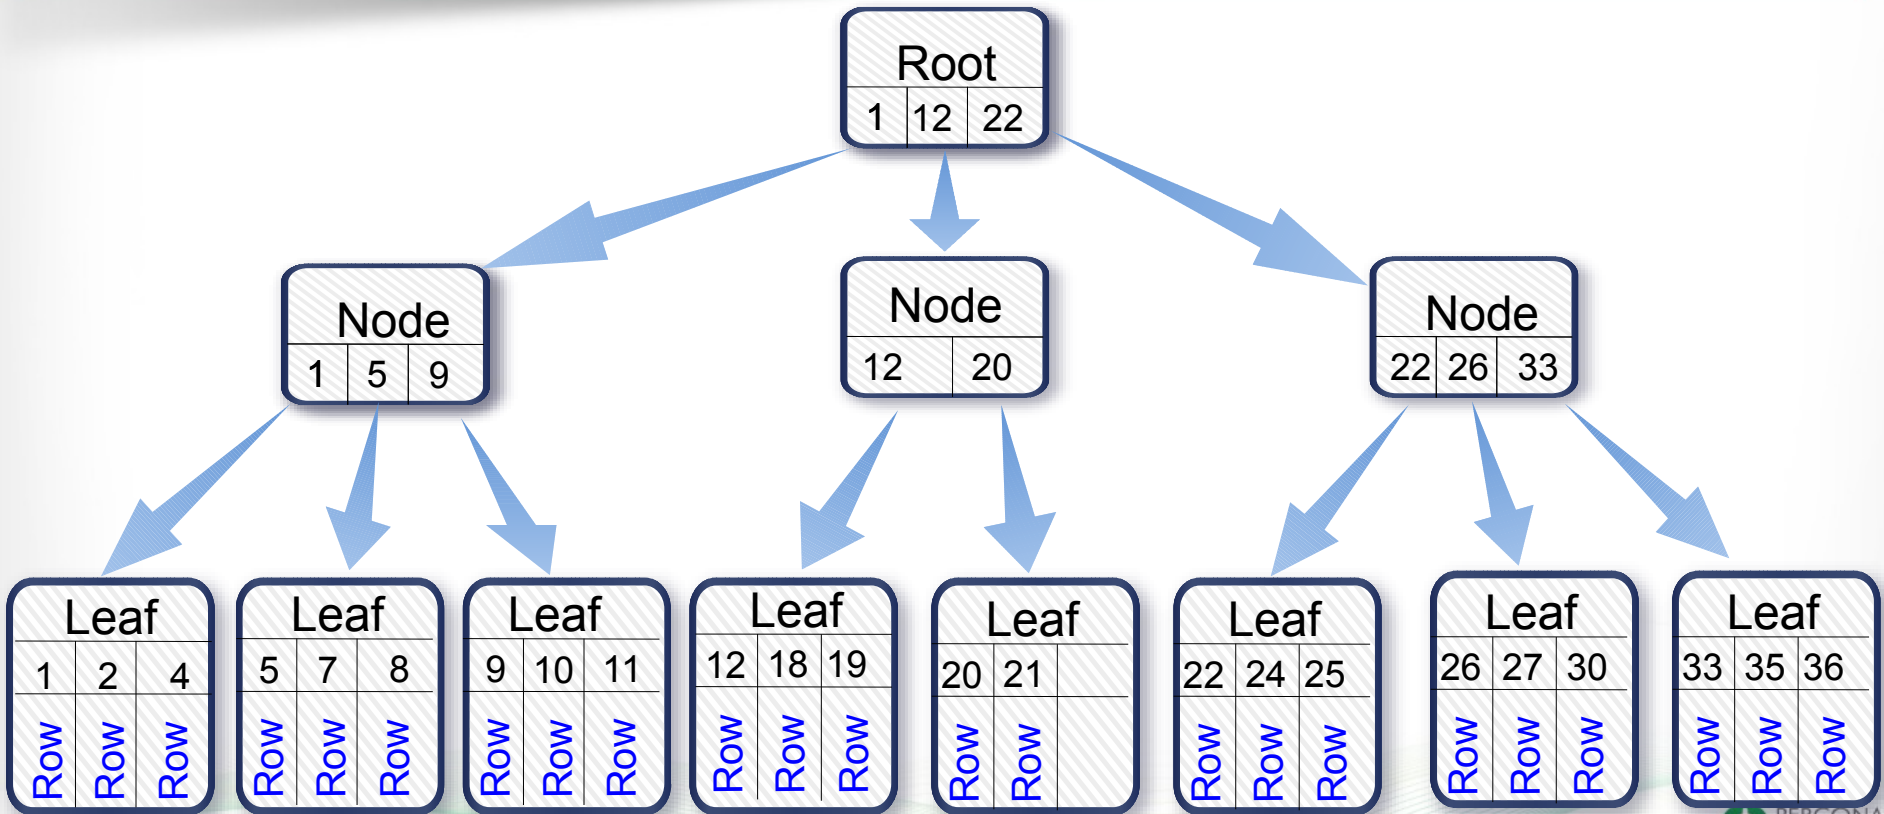
# NewBox application, review of the data types

Is uuid a good thing?

- *Large varchar → slow to compare*
- *hex has low cardinality per byte*
- *inflate the size of the Sks*
- *random insert order*
- **Should use int unsigned auto\_increment**



# NewBox application, review of the PKs



# NewBox application, review of the PKs

userPref, object, share

- Retrieving object rows for a given userId or ownerId*
- SKey on userId = 3 gives us: {8,12,27}*

Leaf		
1	2	4
Row	Row	Row

Leaf		
5	7	8
Row	Row	Row

Leaf		
9	10	11
Row	Row	Row

Leaf		
12	18	19
Row	Row	Row

Leaf		
20	21	
Row	Row	

Leaf		
22	24	25
Row	Row	Row

Leaf		
26	27	30
Row	Row	Row

Leaf		
33	35	36
Row	Row	Row

# NewBox application, review of the PKs

userPref, object, share

- *Reordering the Pks: objectId → UK*
- *PK → (userId, objectId)*

Leaf			Leaf			Leaf			Leaf			Leaf			Leaf			Leaf					
1	1	1	1	2	2	3	3	3	5	6	6	6	6		7	8	8	8	8	8	8	9	9
7	19	20	22	2	11	8	12	27	10	1	4	9	24		35	5	18	25	26	30	36	21	33
Row	Row	Row	Row	Row	Row	Row	Row	Row	Row	Row	Row	Row	Row		Row	Row	Row	Row	Row	Row	Row	Row	Row



# NewBox application, review of the PKs

## activityLog

- *Lots of inserts*
- *Ok as auto\_increment → merges writes*
- *minimize keys on master → use a slave*
- *Good idea to use partitions on ranges of activityLogId*



# NewBox application, be sharding ready

## Sharding

- *The ultimate scaling*
- *Start with 2 schema, NewBox\_common and NewBox\_data\_1*
- *NewBox\_common: { user, userPref }*
- *NewBox\_data\_1: { object, share, activityLog }*





# NewBox application, indexing correctly

## On large varchar

- *Slow to compare and big*
- *object.idx\_name and object.idx\_urlStore*
- *prefix issue with objstore, all start with 'http://'*
- *md5 hash?*
- *Better with a CRC32*



# NewBox application, indexing correctly

## Redundant keys

*PRIMARY KEY (`userId`, `shareId`),  
KEY `idx\_user` (`userId`)*

- *idx\_user is useless, covered by the Primary key*
- *pt-duplicate-key-checker is your friend*



# NewBox application, indexing correctly

## Covering keys

*Select o.\* from object o inner join share s on  
o.objectId = s.objectId where s.userId = 12345;*

- *idx\_userId is there, not bad*
- *For each userId, needs to dive in s PK btree*
- *What about: (userId,objectId)*



# NewBox application, indexing correctly

## Index for sorting

*Select o.\* from object where ownerId = 12345  
order by createdAt;*

- *idx\_owner is there*
- *Still has to sort the rows*
- *What if the key is : (ownerId,createdAt)*



# NewBox application, indexing correctly

## Over indexing...

- *No workload is the same*
- *Write intensive → be greedy on keys*
- *Read intensive → be generous on keys but careful not to harm cache*



# NewBox application, indexing correctly

## Tools

- *explain*
- *pt-query-digest*
- *Percona cloud tool*





# Questions

