

Tuning MySQL

It's About Performance

Bill Karwin

Senior Database Architect @ SchoolMessenger



PERCONA
LIVE

MySQL Tuning

INTRODUCTION



What Are We Talking About?

3

- Configuration variables.
 - We change them to allocate resources for specific features in MySQL.
 - Or enable/disable optional behavior.
 - Make settings persistent by editing /etc/my.cnf



Types of Tuning Changes

4

- Many configuration variables can be GLOBAL or SESSION.
 - Sessions copy global values at connect time.
 - Sessions can change tuning values for the scope of one connection.



Tuning Advantages

- No changes required to schema.
- No changes required to code.
- Some tuning changes possible without restarting mysqld.

Agenda

- InnoDB Buffer Pool
- InnoDB Redo Log
- InnoDB IO Capacity
- InnoDB Other
- Optimizer
- Logging
- Replication
- Connections and Threads
- Tables
- Query Cache
- Operating System
- Tuning Tools
- Monitoring Tools
- Tuning vs. Architecture

MySQL Tuning

INNODB **BUFFER POOL**



InnoDB Buffer Pool

- Largest single use of memory in the server
- In-memory cache of InnoDB data, indexes, undo pages, change buffer – anything that is stored in *pages* in the tablespace.

Buffer Pool Size

9

- Enough to hold most frequently-used pages.
 - `innodb_buffer_pool_size = 10240M`

Buffer Pool Size

- Watch the ratio of page reads to page reads that needed I/O:


```
mysql> SHOW GLOBAL STATUS LIKE  
'Innodb_buffer_pool_read%';
```

variable_name	value	efficiency
Innodb_buffer_pool_read_requests	30887	
Innodb_buffer_pool_reads	411	

98.67%
efficiency

Buffer Pool as a Portion of RAM

11



Buffer Pool Sizing Tips

12

- Assume the buffer pool uses +10% RAM for internal metadata.
- You can reach a point of diminishing returns.
 - A small database doesn't need a large buffer pool.
 - Pages are not stored more than once.
- Don't oversize it and cause swapping!

What's In My Buffer Pool?

```
USE information_schema;
SET @page_size = @@innodb_page_size;
SET @bp_pages = @@innodb_buffer_pool_size/@page_size;

/* MySQL or Percona Server 5.6 – 5.7 */
SELECT P.TABLE_NAME, P.PAGE_TYPE,
CASE WHEN P.INDEX_NAME IS NULL THEN NULL WHEN P.TABLE_NAME LIKE '`SYS_%' THEN P.INDEX_NAME WHEN
P.INDEX_NAME <> 'PRIMARY' THEN 'SECONDARY' ELSE 'PRIMARY' END AS INDEX_TYPE,
COUNT(DISTINCT P.PAGE_NUMBER) AS PAGES,
ROUND(100*COUNT(DISTINCT P.PAGE_NUMBER)/@bp_pages,2) AS PCT_OF_BUFFER_POOL,
CASE WHEN P.TABLE_NAME IS NULL THEN NULL WHEN P.TABLE_NAME LIKE 'SYS\_%' THEN NULL ELSE
ROUND(100*COUNT(DISTINCT P.PAGE_NUMBER)/CASE P.INDEX_NAME WHEN 'PRIMARY' THEN TS.DATA_LENGTH/
@page_size ELSE TS.INDEX_LENGTH/@page_size END, 2) END AS PCT_OF_INDEX
FROM INNODB_BUFFER_PAGE AS P
JOIN INNODB_SYS_TABLES AS T ON PSPACE = TSPACE
JOIN TABLES AS TS ON T.NAME = CONCAT(TS.TABLE_SCHEMA, '/', TS.TABLE_NAME)
WHERE TS.TABLE_SCHEMA <> 'mysql'
GROUP BY TABLE_NAME, PAGE_TYPE, INDEX_TYPE;
```

MySQL or
Percona Server
5.6 – 5.7

What's In My Buffer Pool?

```
USE information_schema;
SET @page_size = @@innodb_page_size;
SET @bp_pages = @@innodb_buffer_pool_size/@page_size;
```

```
/* Percona Server 5.1 - 5.5 */
SELECT P.TABLE_NAME, P.PAGE_TYPE,
CASE WHEN P.INDEX_NAME IS NULL THEN NULL WHEN P.TABLE_NAME LIKE '`SYS_%' THEN P.INDEX_NAME WHEN
P.INDEX_NAME <> 'PRIMARY' THEN 'SECONDARY' ELSE 'PRIMARY' END AS INDEX_TYPE,
COUNT(DISTINCT P.PAGE_NUMBER) AS PAGES,
ROUND(100*COUNT(DISTINCT P.PAGE_NUMBER)/@bp_pages,2) AS PCT_OF_BUFFER_POOL,
CASE WHEN P.TABLE_NAME IS NULL THEN NULL WHEN P.TABLE_NAME LIKE 'SYS\_%' THEN NULL ELSE
ROUND(100*COUNT(DISTINCT P.PAGE_NUMBER)/CASE P.INDEX_NAME WHEN 'PRIMARY' THEN TS.DATA_LENGTH/
@page_size ELSE TS.INDEX_LENGTH/@page_size END, 2) END AS PCT_OF_INDEX
FROM INNODB_BUFFER_PAGE AS P
JOIN INNODB_SYS_TABLES AS T ON P.SPACE = T.SPACE
JOIN TABLES AS TS ON (T.SCHEMA, T.NAME) = (TS.TABLE_SCHEMA, TS.TABLE_NAME)
WHERE TS.TABLE_SCHEMA <> 'mysql'
GROUP BY TABLE_NAME, PAGE_TYPE, INDEX_TYPE;
```

Percona Server
5.1 – 5.5

What's in My Buffer Pool?

15

TABLE_NAME	PAGE_TYPE	INDEX_TYPE	PAGES	PCT_OF_BUFFER_POOL	PCT_OF_INDEX
NULL	FILE_SPACE_HEADER	NULL	1	0.00	NULL
NULL	IBUF_BITMAP	NULL	1	0.00	NULL
NULL	INODE	NULL	1	0.00	NULL
`test`.`foo`	INDEX	PRIMARY	2176	3.32	98.37
`test`.`foo`	INDEX	SECONDARY	2893	4.41	88.47



*how much of
BP is full of
each index*

*how much of
each index is
cached in BP*

Buffer Pool Instances

16

- Many threads can queue up, in contention for exclusive access to the buffer pool.
- Scalability issue more than performance issue.

Buffer Pool Instances

17

- Split the buffer pool into a fixed number of sub-pools, and distributes pages among them.
 - `buffer_pool_instances = <integer>`

How Many Buffer Pool Instances?

18

- Typically set this to the number of CPU cores.
- Default BP instances = 1
 - MySQL 5.6 auto-defaults to 8 when BP > 1GB
- You still specify *total* RAM used in `innodb_buffer_pool_size`,
 - Automatically splits evenly between BP instances.

InnoDB Buffer Pool Save & Restore

19

- After a server restart, an empty buffer pool causes low performance until it “warms up.”

InnoDB Buffer Pool Save & Restore

20

- Dump & restore automatically:
 - SET innodb_buffer_pool_dump_at_shutdown = ON;
 - SET innodb_buffer_pool_load_at_startup = ON;
- Dump & restore manually (e.g. in an EVENT)
 - SET innodb_buffer_pool_dump_now = ON;
 - SET innodb_buffer_pool_load_now = ON;

InnoDB Buffer Pool Save & Restore

21

```
CREATE EVENT mysql.buffer_pool_dump  
ON SCHEDULE EVERY 1 HOUR  
DO  
    SET GLOBAL  
        innodb_buffer_pool_dump_now=ON;
```

good if you anticipate crashes!

MySQL Tuning

INNODB REDO LOG




Simon Law  

InnoDB Log File

23

- The log file records changes to InnoDB pages.
- The file(s) are fixed size, and are overwritten.





InnoDB Log Size

24

- Dirty pages in the BP must be accounted for by log entries.
- Log entries may not be overwritten until the corresponding dirty pages are flushed.
- Thus a larger log file allows more dirty pages.




InnoDB Log Size

25

At 0-75% log file usage, query threads work freely, while page cleaner thread does adaptive flushing continually.

75-88% log file usage, page cleaner runs *async flush*.

88%+ log file usage, *all* query threads block for *sync flush*.



https://blogs.oracle.com/mysqlinnodb/entry/introducing_page_cleaner_thread_in

Log File Size Indicator

26

- Enable InnoDB metrics for the buffer pool:
`SET GLOBAL innodb_monitor_enable='module_buffer';`
- Monitor for the number of sync waits:
`SELECT name, count_reset
FROM INFORMATION_SCHEMA.INNODB_METRICS
WHERE name LIKE 'buffer_flush_sync%';`
- If the counts are regularly greater than zero,
increase `innodb_log_file_size`.

Log Buffer Size


27

- Small (8MB) buffer for writing redo log records to the log file.
- If it's full, a COMMIT has to wait for it to flush.
 - `SHOW GLOBAL STATUS LIKE 'Innodb_log_waits';`
 - If you get frequent waits (> 1/minute), increase `innodb_log_buffer_size`.

InnoDB Log Flush at Transaction Commit

28


- `innodb_flush_log_at_trx_commit = 1`
 - Every transaction COMMIT is fully synchronous



InnoDB Log Flush at Transaction Commit

29


- `innodb_flush_log_at_trx_commit = 2`
 - Each COMMIT flushes to filesystem



InnoDB Log Flush at Transaction Commit

30


- `innodb_flush_log_at_trx_commit = 2`
 - Each COMMIT flushes to filesystem
 - Background thread fsyncs every 1 second



InnoDB Log Flush at Transaction Commit

31


- `innodb_flush_log_at_trx_commit = 0`
 - COMMIT does not flush, only writes to log buffer



InnoDB Log Flush at Transaction Commit

32

- `innodb_flush_log_at_trx_commit = 0`
 - COMMIT does not flush, only writes to log buffer
 - Background thread fsyncs every 1 second



InnoDB Log Flush at Transaction Commit

33

- Tradeoff between durability and performance:
 - Sync on commit (=1) limits commits per second.
 - Flush on commit (=2) risks data loss if OS crashes.
 - No flushing (=0) risks data loss if mysqld aborts.

MySQL Tuning

INNODB IO CAPACITY



Shyaulis Andrjus

InnoDB IO Capacity

35

- Limits the IOPS InnoDB uses while:
 - Flushing dirty pages from the buffer pool to the tablespace.
 - Merging change buffer entries to secondary indexes.

InnoDB IO Capacity Tradeoffs

36

- Raising IO Capacity
 - Causes flushing to become more aggressive.
 - Uses more IO load.
 - **Good when your write load is constantly high.**

InnoDB IO Capacity Tradeoffs

37

- Lowering IO Capacity:
 - Causes flushing to become more gradual.
 - Spreads out the IO load.
 - Allows multiple writes to the same page to be merged into fewer flushes.
 - **Good when your write load has ups and downs.**

InnoDB IO Capacity

38


- `innodb_io_capacity = 200 /* default */`
 - Limit on rate of flushing pages during idle time, or during shutdown.
 - Change buffer merges at a rate of 5-55% of `innodb_io_capacity`.

InnoDB IO Capacity

39

- `innodb_io_capacity_max = 2000 /* default */`
 - Limit on rate of flushing during busy time.

Adaptive Flushing



<http://www.percona.com/blog/2013/10/30/innodb-adaptive-flushing-in-mysql-5-6-checkpoint-age-and-io-capacity/>

InnoDB IO Capacity Caveats

41

- Internals of flushing dirty pages change in each major version of MySQL (5.1/5.5/5.6/5.7/...).
- Different goals? Minimizing IO vs. minimizing log checkpoint age.
- Don't go so high that you cause IO queuing!
http://www.mysqlplus.net/2013/01/07/play-innodb_io_capacity/

InnoDB Buffer Pool LRU Scan Depth

42

- When the BP is full, reading new pages must evict pages currently in the BP.
- Which pages? The least recently used (LRU).

InnoDB Buffer Pool LRU Scan Depth

43

- How many pages may be evicted per second?
`innodb_lru_scan_depth = 1024 /* default */`
- Increase this if you have spare IO capacity.
- This setting is per BP instance
 - Unlike `innodb_io_capacity*`, which are for total capacity of flushing for all BP instances.


<http://mysqlha.blogspot.com/2013/05/configuring-innodb-for-mysql-56.html>

MySQL Tuning

INNODB OTHER CONFIGURATION


InnoDB Flush Method

- `innodb_flush_method=fdatasync /* default */`
- Page flush writes to filesystem, and fsyncs



InnoDB Flush Method

- `innodb_flush_method=O_DIRECT`
- Page flush bypasses the filesystem



InnoDB Flush Method

47

- The best setting depends on your hardware
 - O_DIRECT is often good on caching RAID
 - O_DIRECT not good when IO has latency (e.g. SAN, DRBD, or Amazon EBS when not EBS-optimized)
- To get optimal results, benchmark *your* application workload on *your* hardware



InnoDB Change Buffer

48

- If you INSERT/UPDATE/DELETE in a non-unique index, the *change buffer* helps to delay changes to the index.
 - Makes writes faster
 - The more indexes, the greater the benefit

InnoDB Change Buffer

- The queue of changes can grow up to 25% of the size of the buffer pool.

```
innodb_change_buffer_max_size=25 /* default */
```

- In practice, this fills up if you have a *lot* of writes for a sustained period.

InnoDB Change Buffer

- Background thread merges buffered changes into indexes, at a rate of 5% IO capacity.*
 - Increases gradually up to 55% of IO capacity if the change buffer is more than half of `innodb_change_buffer_max_size`, as a percentage of the BP size.

* It's supposed to merge at a rate of 100% IO capacity when the system is idle, but merging itself counts as "not idle" so I'm not sure it can ever do that.

To shrink the change buffer, try:

- Increase the merge rate by raising `innodb_io_capacity`, or lowering `innodb_change_buffer_max_size`

InnoDB Change Buffer

52

- ...or reduce the growth of the change buffer:
 - innodb_change_buffering = { inserts | updates | deletes | changes | purges | none }
 - Global option only, not per session or per table.
- ...or change write-heavy tables to a different storage engine.

InnoDB Adaptive Hash Index

53

- Cache of frequently-requested index values.
 - Speeds up secondary index searches automatically – nothing to enable.
 - See it working in InnoDB status:
-

INSERT BUFFER AND ADAPTIVE HASH INDEX

60608.42 hash searches/s, 86753.09 non-hash searches/s



~41% of searches use the AHI

InnoDB Adaptive Hash Index


But –

- The mutex for the AHI can become a bottleneck.
 - SEMAPHORES section of InnoDB status reports waits in btr0sea.c
- You can disable it:
 - skip_innodb_adaptive_hash_index
- Percona Server can split it:
 - innodb_adaptive_hash_index_partitions = <N>

InnoDB Doublewrite Buffer

55


- Every dirty page flush writes twice:
 - First, write to the doublewrite buffer on disk



InnoDB Doublewrite Buffer

56

- Every dirty page flush writes twice:
 - Second, write to the respective pages on disk



- The doublewrite buffer adds overhead.
- Alternative:
 - Put datadir on a transactional filesystem.
 - Disable doublewrite buffer:
`skip_innodb_doublewrite`

<http://www.percona.com/blog/2014/05/23/improve-innodb-performance-write-bound-loads/>

InnoDB Read/Write Threads

58

- Background threads to read and write pages.
- See them working in InnoDB Status:

```
-----  
FILE I/O  
-----  
I/O thread 2 state: waiting for completed aio requests (read thread)  
I/O thread 3 state: waiting for completed aio requests (read thread)  
I/O thread 4 state: waiting for completed aio requests (read thread)  
I/O thread 5 state: waiting for completed aio requests (read thread)  
I/O thread 6 state: waiting for completed aio requests (write thread)  
I/O thread 7 state: waiting for completed aio requests (write thread)  
I/O thread 8 state: waiting for completed aio requests (write thread)  
I/O thread 9 state: waiting for completed aio requests (write thread)  
Pending normal aio reads: 0 [0, 0, 0, 0] , aio writes: 0 [0, 0, 0, 0],
```

Watch the number of pending reads and writes.

If these go too high (~64), then increase number of IO threads.

MySQL Tuning

OPTIMIZER CONFIGURATION

Optimizer Switches

60

- Enable/disable optimizer features that aren't doing what you want. Example:
 - `optimizer_switch='index_merge_intersection=off';`

<http://www.percona.com/blog/2012/12/14/the-optimization-that-often-isnt-index-merge-intersection/>

Sort Buffer Size

- In-memory buffer per thread for sorting query results.
 - `SET sort_buffer_size = 256K /* default */`
- If the result is too large, subsets are sorted and merged on disk.
 - `SHOW GLOBAL STATUS LIKE 'Sort_merge_passes';`

variable_name	value
Sort_merge_passes	6060842

“Measure twice, cut once.”




Step 1

- Choose a measurable indicator of performance
 - e.g. sort_buffer_size effectiveness is measured by sort_merge_passes
- Measure the impact of performance before changing the configuration parameter.

Step 2

- Measure the rate of increase:
Sort merge passes



Step 3


- Research the range of reasonable values for the corresponding configuration variable.
- Make a *modest* change.
 - For example, this variable was 256KB by default.
 - Let's raise it to 384KB.

```
mysql> SET GLOBAL sort_buffer_size = 384*1024;
```

Step 4

- Re-measure the rate after the change:

Sort merge passes



How to Decide on a Size?

67

- Must the rate of increase be zero? **No.**
- Using the disk for sort merge passes *occasionally* is normal.
 - Size the sort buffer so the rate of sort merge passes is low – but no need to make it zero.
- This principle applies to other tuning as well.

- Design a graph to show the bell curve of result sizes, and how the first modest increase handles more cases than the subsequent incremental increases.

Join Buffer Size

- Used by Batched Key Access in 5.6

```
SET optimizer_switch =  
'mrr=on,mrr_cost_based=off';
```

```
SET optimizer_switch = 'batched_key_access=on';  
SET join_buffer_size = 256K; /* default */
```

<https://dev.mysql.com/doc/refman/5.6/en/bnl-bka-optimization.html>

Read Rnd Buffer Size

70

- Used by Multi-Range Reads in 5.6:

```
SET optimizer_switch =
'mrr=on,mrr_cost_based=off';
```

```
SET read_rnd_buffer_size = 256K; /* default */
```

<https://dev.mysql.com/doc/refman/5.6/en/mrr-optimization.html>

Index Dives

- Searching for multiple disjoint values in a query
 - WHERE foo IN (1, 2, 3, ... N)
 - WHERE foo=1 OR foo=2 ... OR foo=N
 - MySQL 5.5 searches the index for each value individually, so very long lists result in many index “dives” – and slow queries.
- MySQL 5.6 uses index statistics instead, when the list grows longer than a configurable limit.
 - SET eq_range_index_dive_limit = 200;
 - Default in MySQL 5.6 is 10. In MySQL 5.7, it's 200.

MySQL Tuning

LOGGING CONFIGURATION

Slow Query Log

73

- Write information about every “slow” query.
 - SET GLOBAL slow_query_log = ON;
- Choose FILE or TABLE output.
 - SET GLOBAL log_output = 'FILE';
 - TABLE is much slower.

Slow Query Log Filtering

74

- Limit logging by time.
 - `SET GLOBAL long_query_time = 10;`

Slow Query Log Filtering

75

- (Percona Server) Limit logging by query execution plan.
 - SET GLOBAL log_slow_filter = 'tmp_table_on_disk,filesort_on_disk';

Slow Query Log Filtering

76

- (Percona Server) Limit logging by sampling,
e.g. 1/100 queries or sessions.
 - `SET GLOBAL log_slow_rate_limit = 100;`

MySQL Tuning

REPLICATION CONFIGURATION

Sync Binlog

- Choose how strictly durable the binlog is – trading off performance.
- Sync writes to the binlog every N commits.
 - `SET sync_binlog = 0; /* default */`
- Many people assume this is boolean and set it to 1, which causes the highest overhead.

- `binlog_format` is a tradeoff between deterministic changes vs. efficiency.
 - `STATEMENT` writes less in the log, but requires parsing, optimization and execution on slave.
 - `ROW` writes more log for complex updates, but has less overhead and more reliability.

- A buffer for uncommitted binlog writes.

```
SET GLOBAL binlog_cache_size = 32768; /* default */
```

- A transaction must save buffer to disk when the buffer is full.

Binlog Cache

- Monitor ratio of disk use.

- SHOW GLOBAL VARIABLES LIKE 'Binlog_cache_%use';

variable_name	value
Binlog_cache_disk_use	327
Binlog_cache_use	6060842



Multi-Threaded Slave

82

- Intended to replay binlog events as fast as the master created them.
 - `SET GLOBAL slave_parallel_workers = 10;`
 - Updates to a given schema still happen serially in one thread – use as many workers as schemas.*


https://blogs.oracle.com/MySQL/entry/benchmarking_mysql_replication_with_multi

MySQL Tuning

CONNECTION AND THREAD CONFIGURATION

Max Connections

RRDTOOL / TOBI DETIKER



Max Connections	Current:	120.00
Max Used	Current:	55.00
Aborted Clients	Current:	0.00
Aborted Connects	Current:	0.00
Threads Connected	Current:	26.00
New Connections	Current:	230.56 m
	Average:	5.04 m
	Max:	66.34 m
	Average:	0.00
	Max:	0.00
	Average:	21.47
	Max:	40.33
	Average:	240.14 m
	Max:	958.05 m

Max Connections

- Should you set `max_connections` very high?
 - Have you tested what happens when your server spikes to 10,000+ threads?
 - Resource contention could make performance exponentially worse.
 - It could be better to refuse some connections, so those that are connected can finish.

InnoDB Thread Concurrency


86

- How many threads are allowed to work in the InnoDB engine in parallel?
 - `SET GLOBAL innodb_thread_concurrency = 0;`
/* default, no limit */
 - If too much contention, try 2× CPU cores.

Thread Pool

- (Percona Server) Serve more connections with a limited number of MySQL user threads.
 - `thread_handling=pool-of-threads`
 - `thread_pool_size=36`
 - `thread_concurrency=0`
- Test with the thread pool if you typically see `Threads_running` at 64 or more.

Thread Pool



<http://www.percona.com/blog/2014/01/23/percona-server-improve-scalability-percona-thread-pool/>



Don't Overallocate

- Some buffers are allocated globally:
 - innodb_buffer_pool_size
 - innodb_log_buffer_size
 - max_heap_table_size*
 - query_cache_size
- Some are allocated per SQL thread:
 - sort_buffer_size
 - binlog_cache_size
 - join_buffer_size*
 - read_buffer_size
 - read_rnd_buffer_size
 - thread_stack
 - tmp_table_size*

* may be allocated multiple times

Global vs. Per-Thread Resources

90



× Max_connections
or more practically,
Max_used_connections

MySQL Tuning

TABLES CONFIGURATION



Table Open Cache

92

- Cache for table metadata. Limits the number of tables in use for all threads.
 - `SET GLOBAL table_open_cache = 2000;`
`/* default */`
 - Set to `Threads_running × tables per query.`



Table Open Cache

93

- Watch the number of opened tables per second.
 - SHOW GLOBAL STATUS LIKE 'Opened_tables';
 - If the rate of increase is too sharp, increase the table_open_cache.

Table Open Cache Instances

94

- MySQL 5.6 feature to split the table cache.
 - `table_open_cache_instances = 1; /*default*/`
 - May help reduce contention when you have `Threads_running` more than ~64.
 - Set to 8 or 16, the number of CPU cores.

Table Definition Cache

- Cache for table metadata (contents of .frm).
 - Small overhead, so safe to increase this.
 - `SET GLOBAL table_definition_cache = 400 + (@@table_open_cache/2); /* default */`
 - InnoDB also uses this as a soft limit for the data dictionary cache.

InnoDB Open Tables

96

- Limits the number of open InnoDB tablespaces.
 - `innodb_open_tables = -1 /* default, autosized */`
 - Self-adjusts up to `table_open_cache`.

InnoDB File Per Table

97

- Store each InnoDB data in a separate .ibd file (tablespace). Default in 5.6+.
 - Required for some table options, transportable tablespaces, recovering disk space.
 - No significant performance impact, unless you have tens of thousands of tables.

MySQL Tuning

QUERY CACHE CONFIGURATION

Query Cache Effectiveness: Number

99

- Monitor ratio of QC hits vs. misses:

```
SHOW GLOBAL STATUS LIKE 'QCache%';
```

variable_name	value
Qcache_hits	8675309
Qcache_inserts	606084

15:1 hits – good!

Query Cache Effectiveness: Magnitude

100

- `pt-query-digest` outputs response time histogram per query.

```
# QC Hit      63% yes,  36% no
. .
# Query_time distribution
#   1us
# 10us  #####
# 100us #####
#   1ms
# 10ms
# 100ms #####
#     1s  #####
# 10s+ #####
query cache hits
query cache misses –
4-5 orders of magnitude
slower (for this query)
```

Query Cache Scalability

101

- Every thread that reads or writes the query cache must acquire a mutex.
 - SELECT must check if the query result is cached
 - INSERT/UPDATE/DELETE must evict query results
- This can become a bottleneck when you have many threads running – worse than not having the query cache!



Query Cache Tuning

102

- Better to disable it unless you can confirm it's giving you a lot of benefit.
 - `query_cache_type = 0`
 - `query_cache_size = 0`

*you should set **both** in my.cnf!*

Query Cache Defaults

103

- MySQL 5.5:
 - query_cache_type = 1 /* still enables mutex */
 - query_cache_size = 0 /* allocates no memory */
- MySQL 5.6.8+:
 - query_cache_type = 0 /* off */
 - query_cache_size = 1M /* still allocates memory */

MySQL Tuning

TEMPORARY SPACE CONFIGURATION

Temporary Table Size

- A temporary table created by a query may fit in memory up to tmp_table_size
 - SET tmp_table_size = 16M; /* default */
 - Larger temp tables are written to disk.
- Exception: temp tables with BLOB/TEXT columns always write to disk.

Temporary Table Size

106

- Any in-memory table is limited by:
 - `SET max_heap_table_size = 16M; /* default */`
 - So it's futile to set `tmp_table_size` greater.

Temporary Table Size

- Monitor ratio of temp tables to temp tables on disk:

```
SHOW GLOBAL STATUS LIKE 'Created%tables';
```

variable_name	value
created_tmp_disk_tables	123
created_tmp_tables	6060842



Temporary Table Size

108

- Like other buffers, increase tmp_table_size to handle more cases, until the rate of disk tables drops.

Temporary Tables per Statement

109

- (Percona Server) Verbose slow-query log includes fields for Tmp_tables, Tmp_disk_tables, Tmp_table_sizes.
- Performance_Schema.events_statements_% includes columns created_tmp_tables and created_tmp_disk_tables.
 - But not tmp table sizes (<http://bugs.mysql.com/74484>).

Tmpfs Partition

- Trick MySQL into writing “on disk” temp tables into memory (even with BLOB/TEXT):
 - In /etc/my.cnf:

```
tmpdir=/var/lib/mysqltmp
```
 - In /etc/fstab:

```
tmpfs /var/lib/mysqltmp tmpfs \
noatime,size=1G,mode=700,uid=mysql,gid=mysql 0 0
```

MySQL Tuning


OPERATING SYSTEM CONFIGURATION

- XFS consistently performs better in Percona benchmarks.
 - Especially for multi-threaded IO while using `innodb_flush_method=O_DIRECT`

<http://www.percona.com/blog/2011/12/16/setting-up-xfs-the-simple-edition/>

IO Scheduler

113



<http://www.percona.com/blog/2009/01/30/linux-schedulers-in-tpcc-like-benchmark/>

Mount Options

114

- Reduce unneeded writes to update access times (but less important to tune if you use XFS)
 - noatime
 - nodiratime
- If you use ext4, and battery-backed RAID cache, disable barriers
 - nobarrier

- Reduce the kernel's tendency to swap proactively. Edit `/etc/sysctl.conf`:
 - `vm.swappiness = 1`
 - Some old advice was to set swappiness to 0, but in Linux kernel 2.6.32+, the value 0 can lead to OOM conditions that kill mysqld.

<http://www.percona.com/blog/2014/04/28/oom-relation-vm-swappiness0-new-kernel/>

- Some random stalls have been fixed by:
 - numactl --interleave=all
 - sysctl -q -w vm.drop_caches=3
 - forcing the NUMA node allocation decisions to be made on startup
- These changes were developed at Twitter
 - <http://blog.jcole.us/2012/04/16/a-brief-update-on-numa-and-mysql/>
- These patches are available in Percona Server 5.5 – 5.6
 - http://www.percona.com/doc/percona-server/5.6/performance/innodb_numa_support.html

MySQL Tuning

TUNING TOOLS



- Recommends rough tuning changes based on activity rates in SHOW GLOBAL STATUS.
 - Somewhat outdated – claims to support MySQL 5.4 and 6.0, which never existed.
 - Averages over uptime fail to account for spikes.

<http://mysqltuner.com/>

Percona Configuration Wizard

The screenshot shows the Percona Configuration Wizard interface. At the top, there's a navigation bar with links for About, Products, Software, Resources, News/Events, Community, Contact, Percona Live, and a search icon. Below the navigation is a banner for "Percona Tools for MySQL" with the subtext "Free online productivity tools for MySQL DBAs, SysAdmins and Developers". A "Feedback" button is visible on the left. The main content area displays a configuration file with the following content:

```

# Generated by Percona Configuration Wizard (http://tools.percona.com/) VERSION REL5_20120208
# Configuration name my.cnf generated for Bill.kurwill@percona.com at 2014-11-12 18:16:46
[mysqld]
# CLIENT #
port                = 3306
socket              = /var/lib/mysql/mysql.sock

[mysqld]
# GENERAL #
user                = mysql
default-storage-engine = InnoDB
socket              = /var/lib/mysql/mysql.sock
pid-file            = /var/lib/mysql/mysql.pid

# MyISAM #
key-buffer-size     = 32M
myisam-recover      = FORCE,BACKUP

# SAFETY #
max-allowed-packet = 16M
max-connect-errors  = 1000000
skip-name-resolve
skip-pwd
  
```

Below the configuration file, there are buttons for "Configure another server", "Share this file", "Email me this file", and "Email to a Friend".

<https://tools.percona.com/>

Usually good as a starting point, but not a substitute for monitoring and incremental tuning.

MySQL Tuning

MONITORING TOOLS

- Console tool to measure SHOW GLOBAL STATUS and display incremental changes.
- Nothing to install, part of the Percona Toolkit scripts.

<http://www.percona.com/doc/percona-toolkit/pt-mext.html>

- Real-time monitoring tool for queries, transactions, memory, IO, etc.

<http://www.percona.com/blog/2013/10/14/innotop-real-time-advanced-investigation-tool-mysql/>



Percona Monitoring Plugins

123

- Templates for Cacti and Zabbix to graph GLOBAL STATUS values.

<http://www.percona.com/software/percona-monitoring-plugins>



- New Beta SaaS offering from Percona.
- Graph many performance metrics like PMP.
- Query analysis and trending.

<https://cloud.percona.com/>

MySQL Tuning

TUNING VERSUS ARCHITECTURE



Tuning Advantages

126

- No changes required to schema.
- No changes required to code.
- Some tuning changes possible without restarting mysqld.

Tuning Disadvantages

127

- Most changes global or per session – but cannot be set per user, database, or query.
- Can improve performance – but only so far.



What's the Next Step?

128

- Optimizing by greater order of magnitude:
 - Caching
 - Denormalizing
 - Indexing
 - Sharding or partitioning
 - Query design

- Retain often-used query results.
- Store in memory, for faster retrieval.
- Performance strategy: avoid repeat queries.

Denormalizing

130

- Store data redundantly, for a partially-completed query.
- Avoid expensive expressions, group summaries, or joins.
- Performance strategy: leverage early work.

- Store a pre-sorted copy of some column(s).
- Searching and sorting is now orders of magnitude faster.
- Performance strategy: data structure.

Partitioning

132

- One logical table maps to multiple physical tables of similar type.
- Queries are designed to limit to one partition.
- Performance strategy: divide and conquer.

Sharding

133

- Split a table into subsets of rows.
- Host each subset on a separate instance.
- Performance strategy: horizontal scaling with instances.

Query Design

134

- “There’s more than one way to do it.”
- Experiment with alternative query designs for the same result.
- Performance strategy: take advantage of SQL implementation idiosyncrasies.

MySQL Tuning

CONCLUSIONS

Recap

136

- Measure relevant performance indicators before and after tuning changes.
- Re-test when you upgrade.
- Monitor continually.
- Architecture changes can improve performance beyond tuning.

Thank you!

Questions?

Copyright 2014-2015 Bill Karwin
www.slideshare.net/billkarwin

Released under a Creative Commons 3.0 License:
<http://creativecommons.org/licenses/by-nc-nd/3.0/>

You are free to share - to copy, distribute and
transmit this work, under the following conditions:



Attribution.

You must attribute this work
to Bill Karwin.

Noncommercial.

You may not use this work for
commercial purposes.

No Derivative Works.

You may not alter, transform,
or build upon this work.