

# Why we're excited about MySQL 8

Practical Look for Devs and Ops

---

Peter Zaitsev, CEO

October 2, 2017



# About Percona

Solutions for your success with MySQL, MariaDB and MongoDB

Support, Managed Services, Consulting, Training, Software

Our Software is 100% Free and Open Source

Support Broad Ecosystem – MySQL, MariaDB, Amazon RDS, Google CloudSQL

In Business for 11 years

More than 3000 customers, including top Internet companies and enterprises

# In the Presentation

---

**Practical view  
on MySQL 8**

**Exciting things  
for Devs**

**Exciting things  
for Ops**

# Warning

---

**This assessment is done for Pre-GA MySQL 8, based on documentation and limited testing. We're yet to see how they behave in production**

# Source Notes

---

**Examples liberally borrowed  
from Oracle team  
presentations and Blog Posts**

# MySQL 8 for Ops

---

# Ops care about

---

Stability

High Availability

Performance

Security

Observability

Manageability

# Native Data Dictionary

---

About 10 years overdue

Atomic (Crash Save) DDLs


Much Faster Information Schema

No more MyISAM System Tables!



# Fast Information Schema

```
1 SELECT t.table_schema, t.table_name, c.column_name
2 FROM information_schema.tables t,
3      information_schema.columns c
4 WHERE t.table_schema = c.table_schema
5        AND t.table_name = c.table_name
6        AND t.engine='InnoDB';
```






Scaling from 5.000 to  
1.000.000 tables

# Much Better and Faster UTF8

- utf8mb4 as Default Charset

MySQL 8.0 vs MySQL 5.7 utf8mb4



+300-350% in OLTP RO  
+176-233% in OLTP RW  
+1500-1800% in SELECT DISTINCT\_RANGES

# Security

---

ROLEs

Breakdown of SUPER Privileges

Password History

Faster cached-SHA2 Authentication

--skip-grants blocks remote connections

Redo and Undo Logs are now encrypted if Table Encryption is enabled

# Persistent Auto Increment

---

**Another feature 10 years overdue**

**Do not reset AUTO INCREMENT to the max value in the table on restart**

# Auto-Managed Undo Tablespace

---

**Do not use system table space for undo space any more**

**Automatically reclaim space on disk from large transactions**

# Self Tuning (limited to Innodb)

---

Set  
innodb\_dedicated\_server  
to auto-tune

- innodb\_buffer\_pool\_size
- innodb\_log\_file\_size
- innodb\_flush\_method

# Partial In-Place Update for JSON

---

Can update field in JSON object without full re-write

Great for counters, statuses, timestamps etc

Only update and removal of element is supported

**Only Optimizer and Replication support so far**

# Invisible Indexes

---

- Test impact of dropping indexes before actually dropping them
- Can use use\_invisible\_indexes to use invisible indexes in a session

```
1  CREATE TABLE t1 (  
2      i INT,  
3      j INT,  
4      k INT,  
5      INDEX i_idx (i) INVISIBLE  
6  ) ENGINE = InnoDB;  
7  CREATE INDEX j_idx ON t1 (j) INVISIBLE;  
8  ALTER TABLE t1 ADD INDEX k_idx (k) INVISIBLE;
```



# TmpTable Storage Engine

---

More efficient storage engine for Internal Temporary tables

Efficient storage for VARCHAR and VARBINARY columns

BLOB/TEXT Columns are not supported (yet?)

# Backup Locks

---

**Prevent operation which may result in inconsistent backups**

**LOCK INSTANCE FOR BACKUP**

# Optimizer Histograms

---

- Detailed Statistics on Columns, not just Indexes

```
1  {
2    "buckets": [
3      [
4        1,
5        0.3333333333333333
6      ],
7      [
8        2,
9        0.6666666666666666
10     ],
11     [
12       3,
13       1
14     ]
15   ],
16   "null-values": 0,
17   "last-updated": "2017-03-24 13:32:40.000000",
18   "sampling-rate": 1,
19   "histogram-type": "singleton",
20   "number-of-buckets-specified": 128,
21   "data-type": "int",
22   "collation-id": 8
23 }
```

# Improved Optimizer Cost Model

---

**Keep in account how  
much of data is cached vs  
on disk**

# More on MySQL 8 Optimizer

---

<http://www.unofficialmysqlguide.com/>

# Performance Schema

---

**(Fake) Indexes for Faster Access**

**Error Instrumentation**

**Response Time Histograms (Global and Per Query Digest)**


**Query Examples for Summary by Digest**

# Performance Schema Performance

- Now is Interactively Usable at Scale

## Performance Comparison

Over 30x faster!



# Persistent Global Variables

---

SET PERSIST

innodb buffer pool size =

1024 \* 1024 \* 1024;



# Assumes storage is SSD by Default

---

**Start of the long journey**

# Binary Log On by Default

---

bin\_log is enabled by default

log\_slave\_updates is enabled by default

Expire logs after 30 days by default

# Query Cache Removed

---

It's design caused more problems than it fixed

Use ProxySQL (or other) external query cache instead

# Native Partitioning Only

---

Only “Native” Partitioning supported, not Generic One

Remove partitions from MyISAM partitioned tables or convert them


ALTER TABLE ... REMOVE PARTITIONING

ALTER TABLE ... ENGINE=INNODB

# Resource Groups

- Isolation and Better Performance

MySQL 8.0 Resource Groups - 100% Faster




# Plain Better Performance at Scale

Sysbench: OLTP\_RO Point-Selects

**2.1x Faster than MySQL 5.7**

**2.8x Faster than MySQL 5.6**



# Feature Requests

---

**Better Single Thread Performance**

**Parallel Single Query Processing  
Please**

# MySQL 8 for Devs

---



# InnoDB NO WAIT and SKIP LOCKED

---

## New! Better Handling of Hot Row Contention

```
SELECT * FROM tickets  
WHERE id IN (1,2,3,4)  
AND order_id IS NULL  
FOR UPDATE  
NOWAIT;
```

Error  
immediately if a  
row is already  
locked

```
SELECT * FROM tickets  
WHERE id IN (1,2,3,4)  
AND order_id IS NULL  
FOR UPDATE  
SKIP LOCKED;
```

Non  
deterministically  
skip over locked  
rows

# Descending Indexes

---

**Descending flag in index definition is no more ignored**

**Allows efficient handling of ORDER BY A ASC, B DESC queries**

# Aggregation of Relational into JSON

- **JSON\_ARRAYAGG() and JSON\_OBJECTAGG()**

```
1  mysql> SELECT id, col FROM t1;
2  +-----+-----+
3  | id    | col                                     |
4  +-----+-----+
5  |     1 | {"key1": "value1", "key2": "value2"} |
6  |     2 | {"keyA": "valueA", "keyB": "valueB"} |
7  +-----+-----+
8  2 rows in set (0.00 sec)
9
10 mysql> SELECT JSON_OBJECTAGG(id, col) FROM t1;
11 +-----+
12 | JSON_OBJECTAGG(id, col) |
13 +-----+
14 | {"1": {"key1": "value1", "key2": "value2"}, "2": {"keyA": "valueA", "keyB": "valueB"}} |
15 +-----+
16 1 row in set (0.00 sec)
```

# JSON to Table Conversion (Labs)

```
SET @doc=(SELECT doc->"$.properties.amenities"  
FROM seats WHERE id = 28100);
```

Labs

```
SELECT * FROM json_table(@doc, "$[*]" columns (  
    id for ordinality, amenity_type varchar(100) path "$.type",  
    distance float path '$.distance_in_meters')  
) AS amenities  
WHERE amenity_type IN ('snacks', 'bar')  
ORDER BY distance;
```

```
+-----+-----+-----+  
| id    | amenity_type | distance |  
+-----+-----+-----+  
|      2 | bar          | 100.538 |  
|      3 | snacks       | 136.647 |  
+-----+-----+-----+  
2 rows in set (0.00 sec)
```

# Better JSON Document Data Extraction

---

```
1 mysql> CREATE TABLE t1 (doc JSON);
2 Query OK, 0 rows affected (0.01 sec)
3
4 mysql> INSERT INTO t1 VALUES ('[1, 2, 3, 4, 5]');
5 Query OK, 1 row affected (0.00 sec)
6
7 mysql> SELECT doc->"$[1 to 3]" FROM t1;
8 +-----+
9 | doc->"$[1 to 3]" |
10 +-----+
11 | [2, 3, 4]       |
12 +-----+
13 1 row in set (0.00 sec)
14
15 mysql> SELECT doc->"$[last-2]" FROM t1;
16 +-----+
17 | doc->"$[last-2]" |
18 +-----+
19 | 3               |
20 +-----+
21 1 row in set (0.00 sec)
```

# Common Table Expression

- Recursive and Non-Recursive

```
1 WITH RECURSIVE cte AS
2 (
3   SELECT category_id, name, 0 AS depth FROM category WHERE parent IS NULL
4   UNION ALL
5   SELECT c.category_id, c.name, cte.depth+1 FROM category c JOIN cte ON
6     cte.category_id=c.parent
7 )
8 SELECT * FROM cte ORDER BY depth;
```

category_id	name	depth
1	ELECTRONICS	0
2	TELEVISIONS	1
6	PORTABLE ELECTRONICS	1
7	MP3 PLAYERS	2
9	CD PLAYERS	2
10	2 WAY RADIOS	2
3	TUBE	2
4	LCD	2
5	PLASMA	2
8	FLASH	3

# Window Functions

- Like GROUP BY, But Preserving Rows rather than collapsing them

```
1 mysql> CREATE TABLE t(i INT);
2 mysql> INSERT INTO t VALUES (1),(2),(3),(4);
3
4 mysql> SELECT SUM(i) AS sum FROM t;
5 +-----+
6 | sum |
7 +-----+
8 | 10 |
9 +-----+
10
11 mysql> SELECT i, SUM(i) OVER () AS sum FROM t;
12 +-----+-----+
13 | i | sum |
14 +-----+-----+
15 | 1 | 10 |
16 | 2 | 10 |
17 | 3 | 10 |
18 | 4 | 10 |
19 +-----+-----+
```

# Much Better GIS

---

- “Matching or Exceeding PostgreSQL GIS Feature Set”

## 5.7

- The world is flat
- The world is infinite
- Axes are unitless
- Axes are orthogonal
- Axis order is irrelevant
- Axis direction is irrelevant

## 8.0

- The world can be flat or ellipsoidal
- Geographic coordinate systems wrap around
- Axes have units
- Geographic axes are not orthogonal
- Geographic axis order matters
- Axis direction may be relevant



# MySQL Document Store

---

Full Text Indexing

GeoJSON Support

Anyone Using Document Store ?

# Summary

---

**MySQL 8 looks like release to be excited about**

**Has a lot of new features both for Devs and Ops**

**Thank You!**

---