

17 Things Developers should know about Databases

Peter Zaitsev

CEO, Percona

OpenSource 101 Home

May 12th, 2020



Lets start with some Humor



The Great Split

Developers

Technical
Operations

Operations

Focused on Database (DBA etc)

Generalist (Sysadmin, SRE, etc)

Devs vs Ops

DevOps suppose to have solved it but tension is still common between Devs and Ops

Especially with Databases which are often special snowflake

Especially with larger organizations

Large Organizations

Ops vs Ops have conflict
too

Devs vs Ops Conflict

Devs

- Why is this stupid database always the problem.
- Why can't it just work and work fast

Ops

- Why do not learn schema design
- Why do not you write optimized queries
- Why do not you think about capacity planning

Database Responsibility

Shared Responsibility for
Ultimate Success

Top Recommendations for Developers

Learn Database Basics

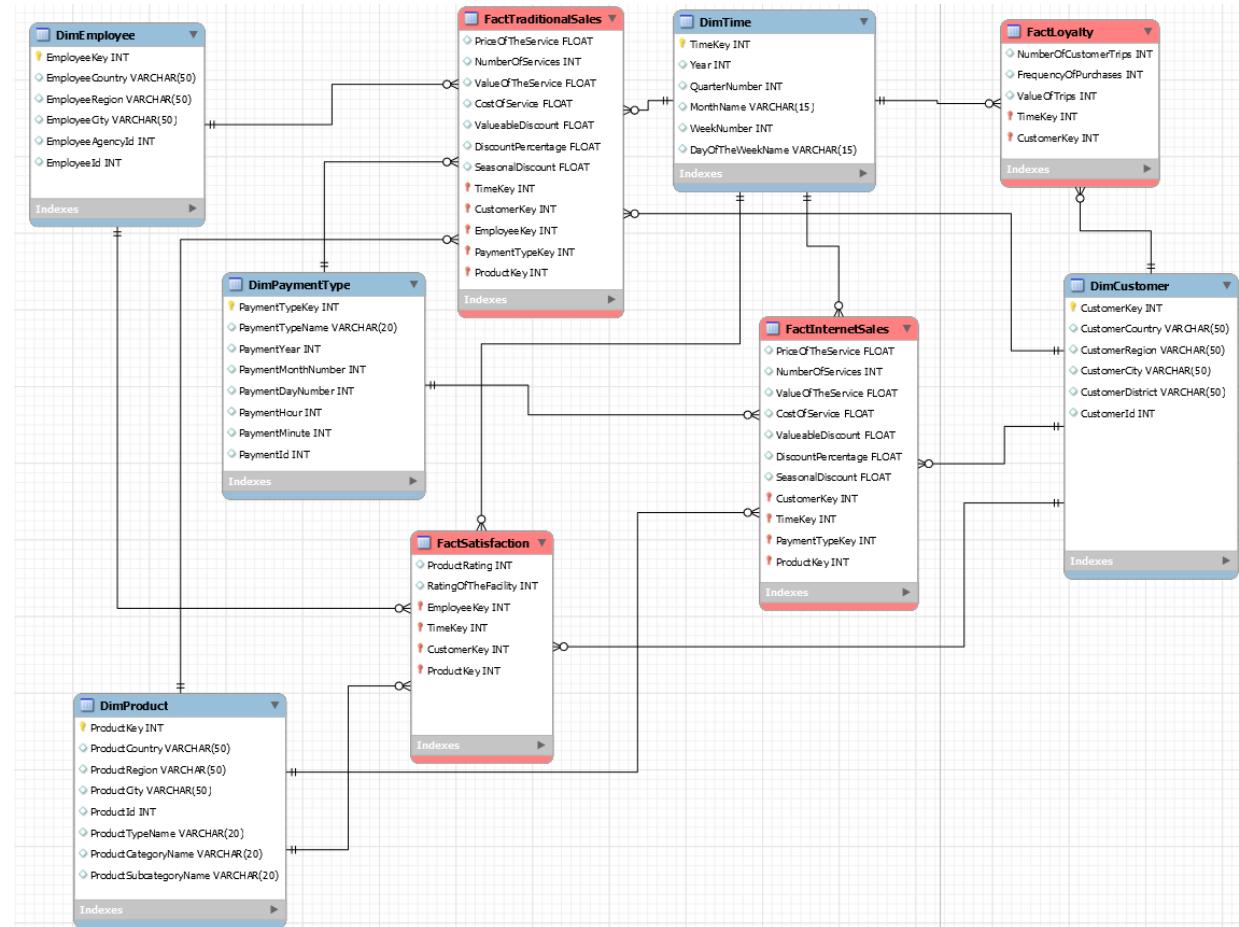
You can't build great database powered applications if you do not understand how databases work

Schema Design

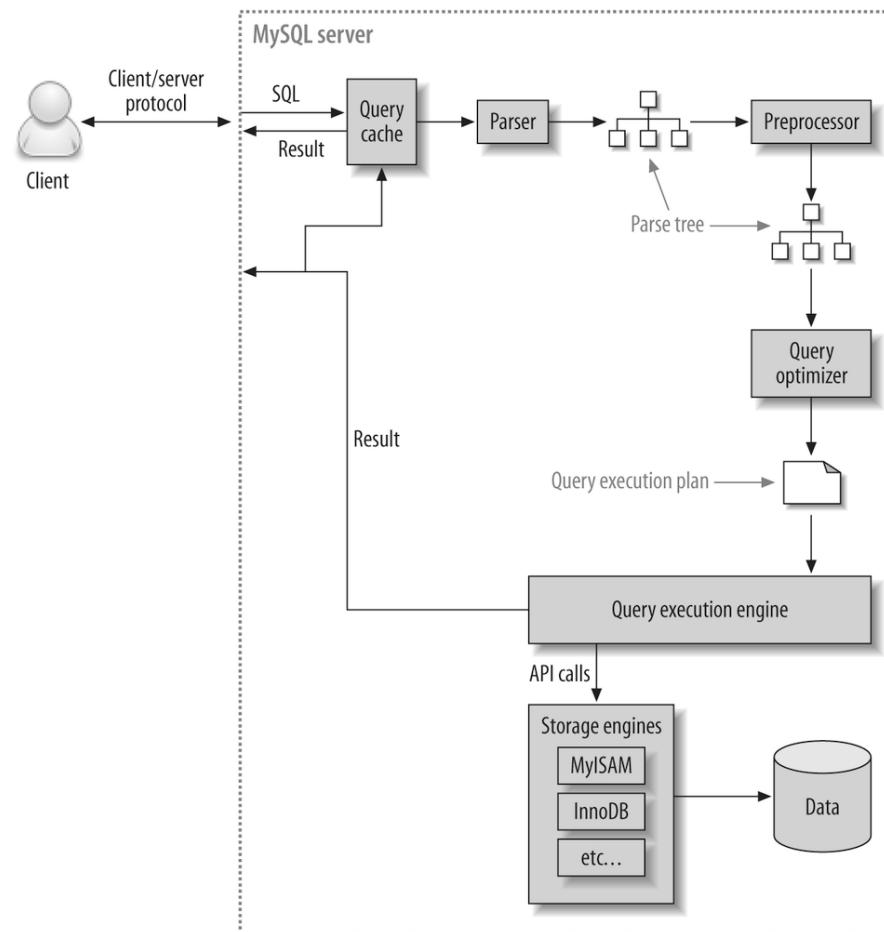
Power of the Database Language

How Database Executes the Query

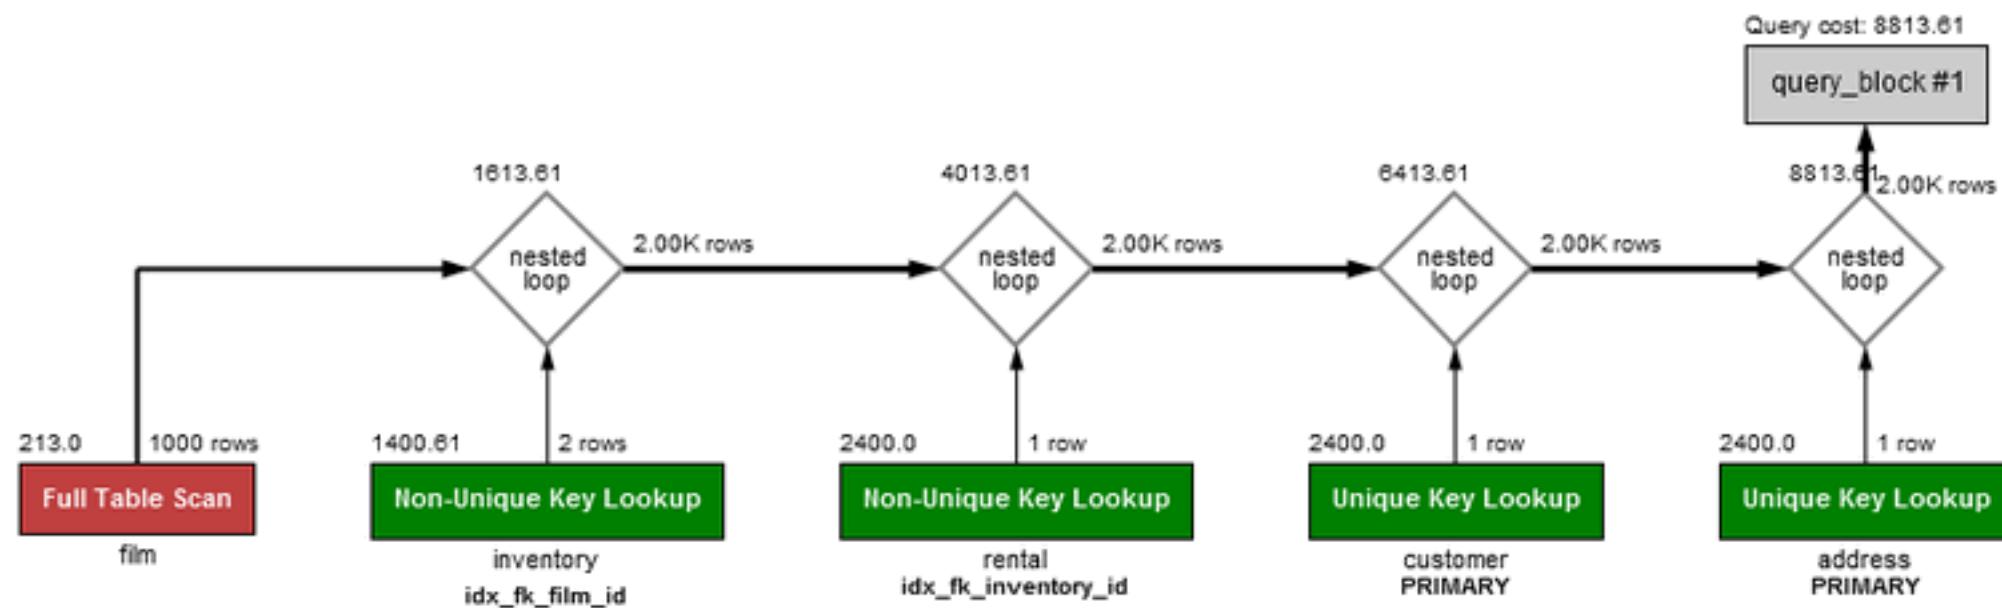
Example of Relational Database Schema



Query Execution Diagram

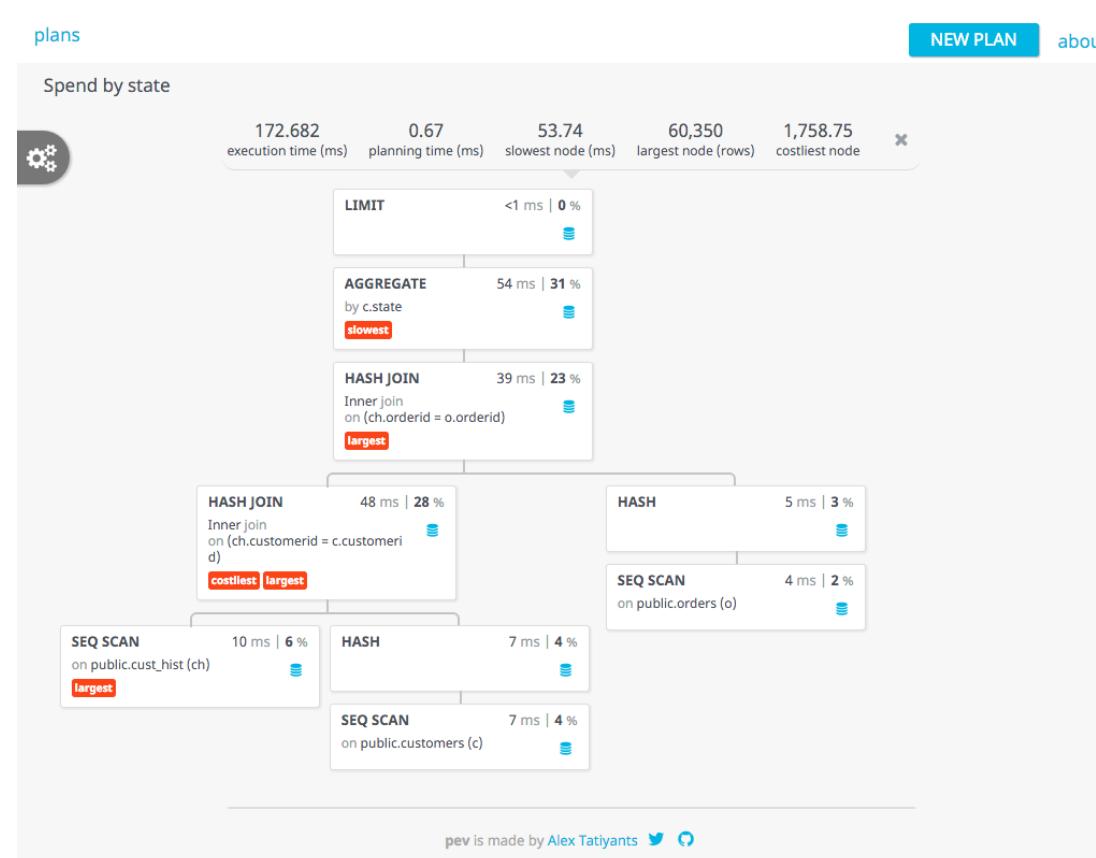


EXPLAIN



<https://dev.mysql.com/doc/refman/8.0/en/execution-plan-information.html>

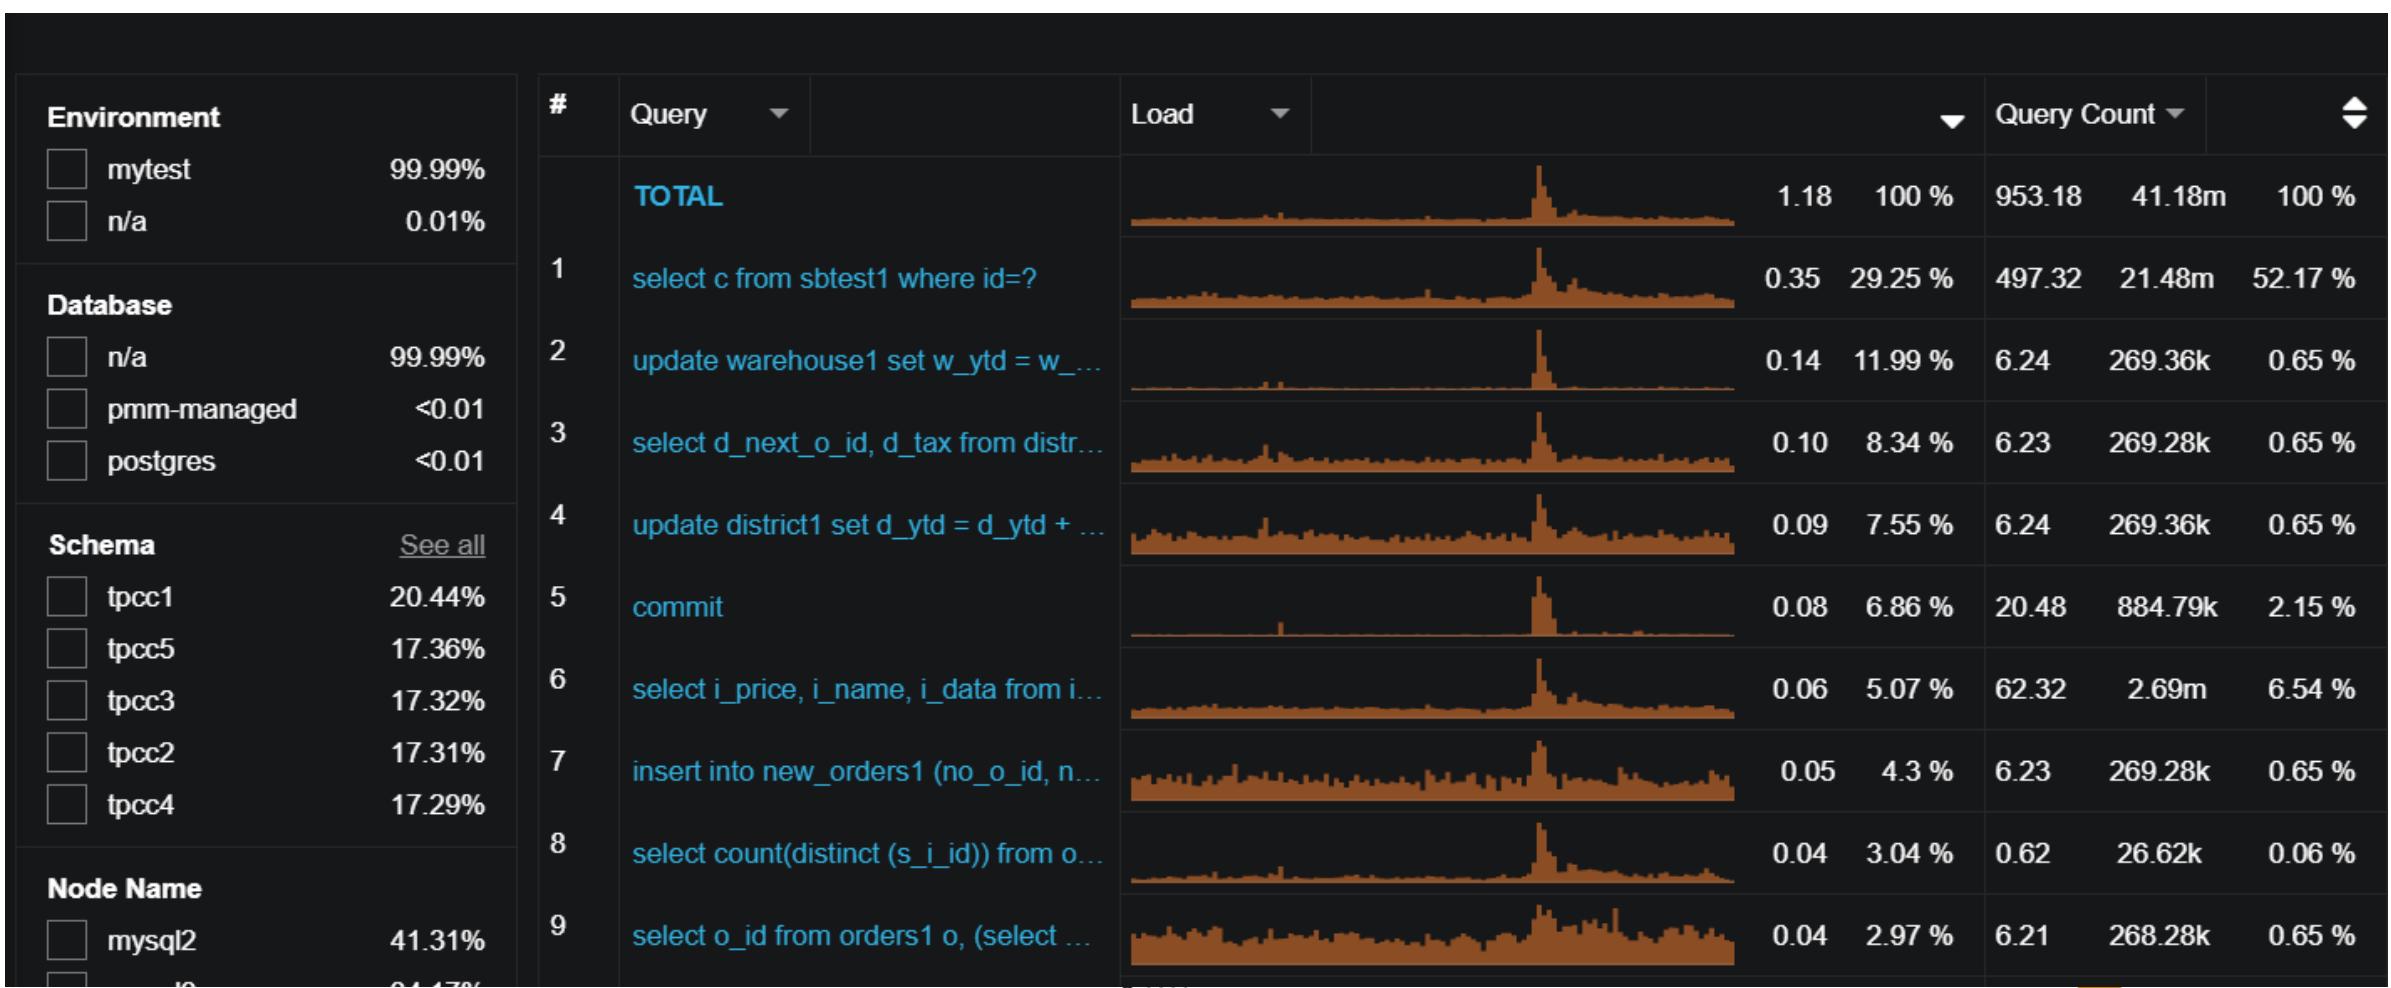
Visualizing PostgreSQL Plan with PEV



<http://tatiyants.com/postgres-query-plan-visualization/>

© 2020 Percona.

Which Queries are Causing the Load



Why Are they Causing this Load

SELECT sbtest			737F39F04B198EF6
Metrics			Query first seen: ⏲ Aug 3, 2017 1:55 PM ... Last seen: ⏲ Today at 9:46 AM
Metrics	Rate/Sec	Sum	Per Query Stats
Query Count	104.05 (per sec)	374.58 k 4.27% of total	
Query Time	19.00 load	18:59:56 29.73% of total	183.66 ms avg
Lock Time	0.11 (avg load)	0:06:42 1.35% of total 0.61% of query time	1.13 ms avg
Innodb IO Read Wait	0.61 (avg load)	0:36:44 9.10% of total 3.38% of query time	6.20 ms avg
Innodb Read Ops	52.35 (per sec)	188.45 k 7.62% of total	0.00 avg
Innodb Read Bytes	857.64 KB (per sec)	3.09 GB 7.62% of total 16.38 KB avg io size	8.22 KB avg
Innodb Distinct Pages	-	-	4.69 avg
Rows Sent	10.41 k (per sec)	37.46 m 30.52% of total	100.00 avg
Bytes Sent	1.30 MB (per sec)	4.67 GB 30.78% of total 124.71 Bytes bytes/row	12.47 KB avg
Rows Examined	1.14 m (per sec)	4.11 b 39.17% of total 109.79 per row sent	10.47 k avg
External Sorts (Filesort)	104.05 (per sec)	374.58 k 49.93% of total 100.00% of queries	-
Full Table Scans	0.01 (per sec)	40.00 0.17% of total 0.01% of queries	-
Queries Requiring Tmp Table In Memory	104.05 (per sec)	374.58 k 95.17% of total 100.00% of queries	-

How to Improve their Performance

✓Example

```
SELECT DISTINCT c
FROM sbtest1
WHERE id
    BETWEEN 5559
        AND 5658
ORDER BY c
```

✓CREATE

```
CREATE TABLE `sbtest1` (
  `id` int(10) unsigned NOT NULL AUTO_INCREMENT,
  `k` int(10) unsigned NOT NULL DEFAULT '0',
  `c` char(120) NOT NULL DEFAULT '',
  `pad` char(60) NOT NULL DEFAULT '',
  PRIMARY KEY (`id`),
  KEY `k_1` (`k`)
) ENGINE=InnoDB AUTO_INCREMENT=100000001 DEFAULT
```

✓JSON

Expand All

```
-{
  "query_block": -{
    "select_id": 1,
    "cost_info": +{...},
    "ordering_operation": -{
      "using_filesort": false,
      "duplicates_removal": -{
        "using_temporary_table": true,
        "using_filesort": true,
        "cost_info": +{...},
        "table": +{...}
      }
    }
  }
}
```

Check out Percona Monitoring and Management

<http://pmmdemo.percona.com>

PMM v 2.6 is fresh off the press

How are Queries Executed ?

Single Threaded

Single Node

Distributed

Indexes

Indexes are
Must

Indexes are
Expensive

Capacity Planning

No Database can handle “unlimited scale”

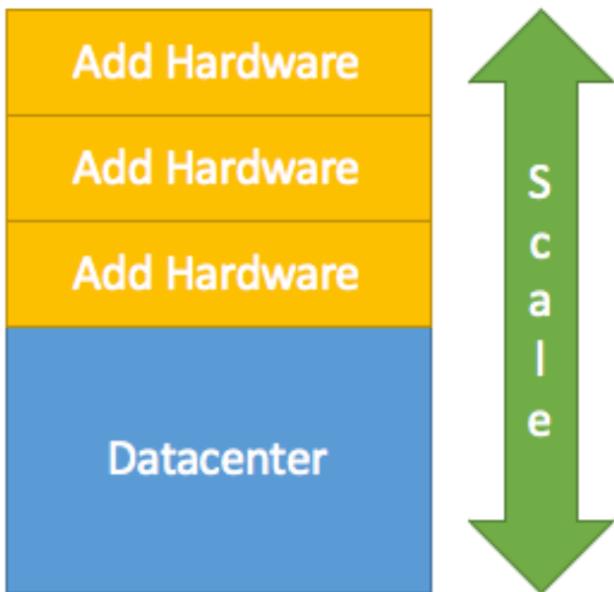
Scalability is very application dependent

Trust Measurements more than Promises

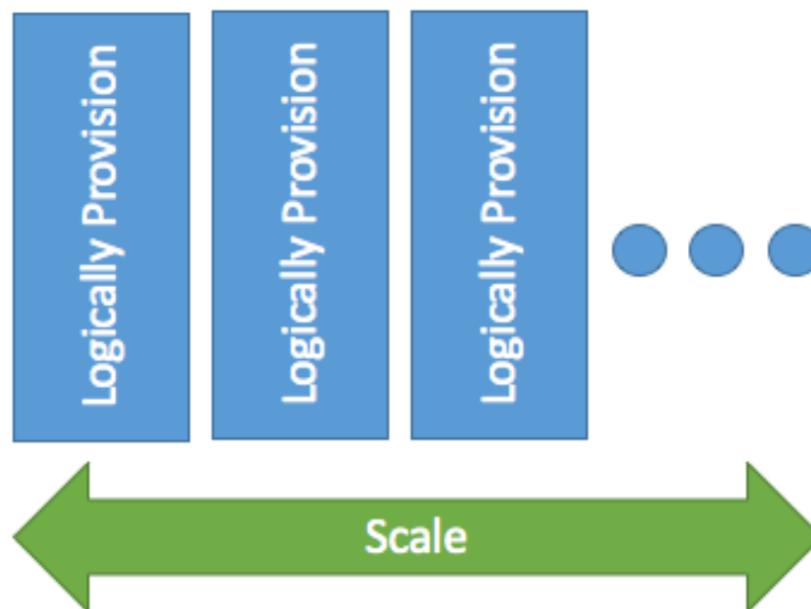
Can be done or can be done Efficiently ?

Vertical and Horizontal Scaling

Vertical Scaling



Horizontal Scaling



Scalable != Efficient

The Systems which promote a scalable can be less efficient

Hadoop, Cassandra, TiDB are great examples

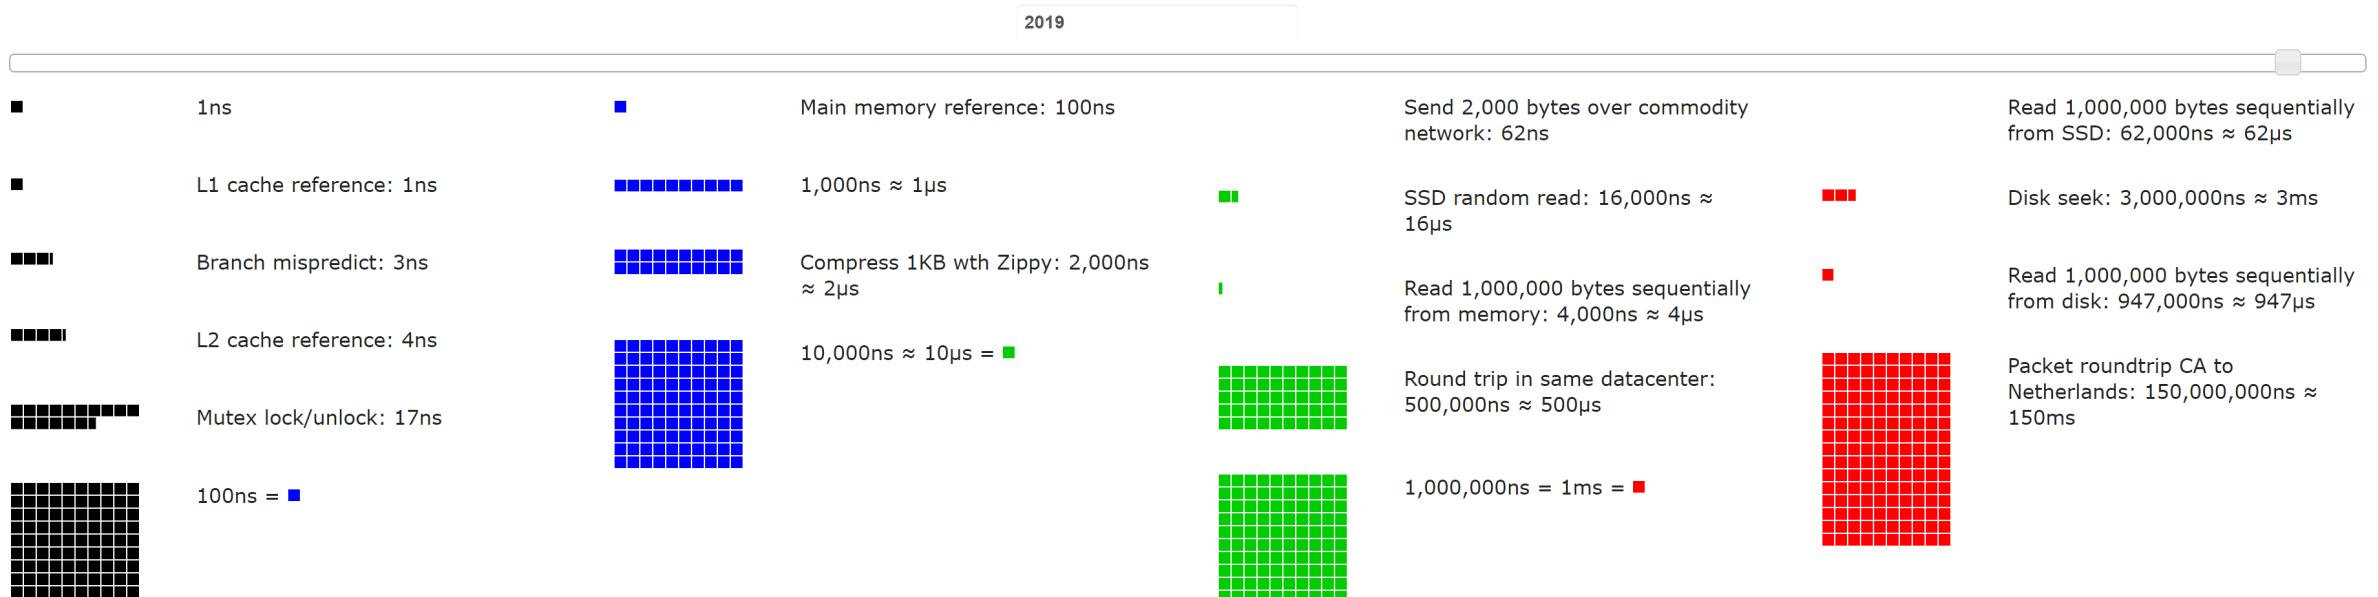
By only the wrong thing you can get in trouble

Throughput != Latency

If I tell you system can do
100.000 queries/sec
would you say it is fast ?

Latency Numbers to Know

Latency Numbers Every Programmer Should Know



Very cool Interactive Diagram

https://colin-scott.github.io/personal_website/research/interactive_latency.html

Speed of Light Limitations

High Availability Design Choices

You want instant durable replication over wide geography or Performance ?

Understanding Difference between High Availability and Disaster Recovery protocols

Network Bandwidth is not the same as Latency

Wide Area Networks

Tend not to be as stable as Local Area Networks (LAN)

Expect increased Jitter

Expect Short term unavailability

Also Understand

Connections to the database are expensive

Especially if doing TLS Handshake

Query Latency Tends to Add Up

Especially on real network and not your laptop

ORM (Object-Relational-Mapping)

Allows Developers to query the database without need to understand SQL

Can create SQL which is very inefficient

Learn SQL Generation “Hints”, Learn JPQL/HQL advanced features

Be ready to manually write SQL if there is no other choice

Do not Leave Transactions Open

Open Connection is rather inexpensive

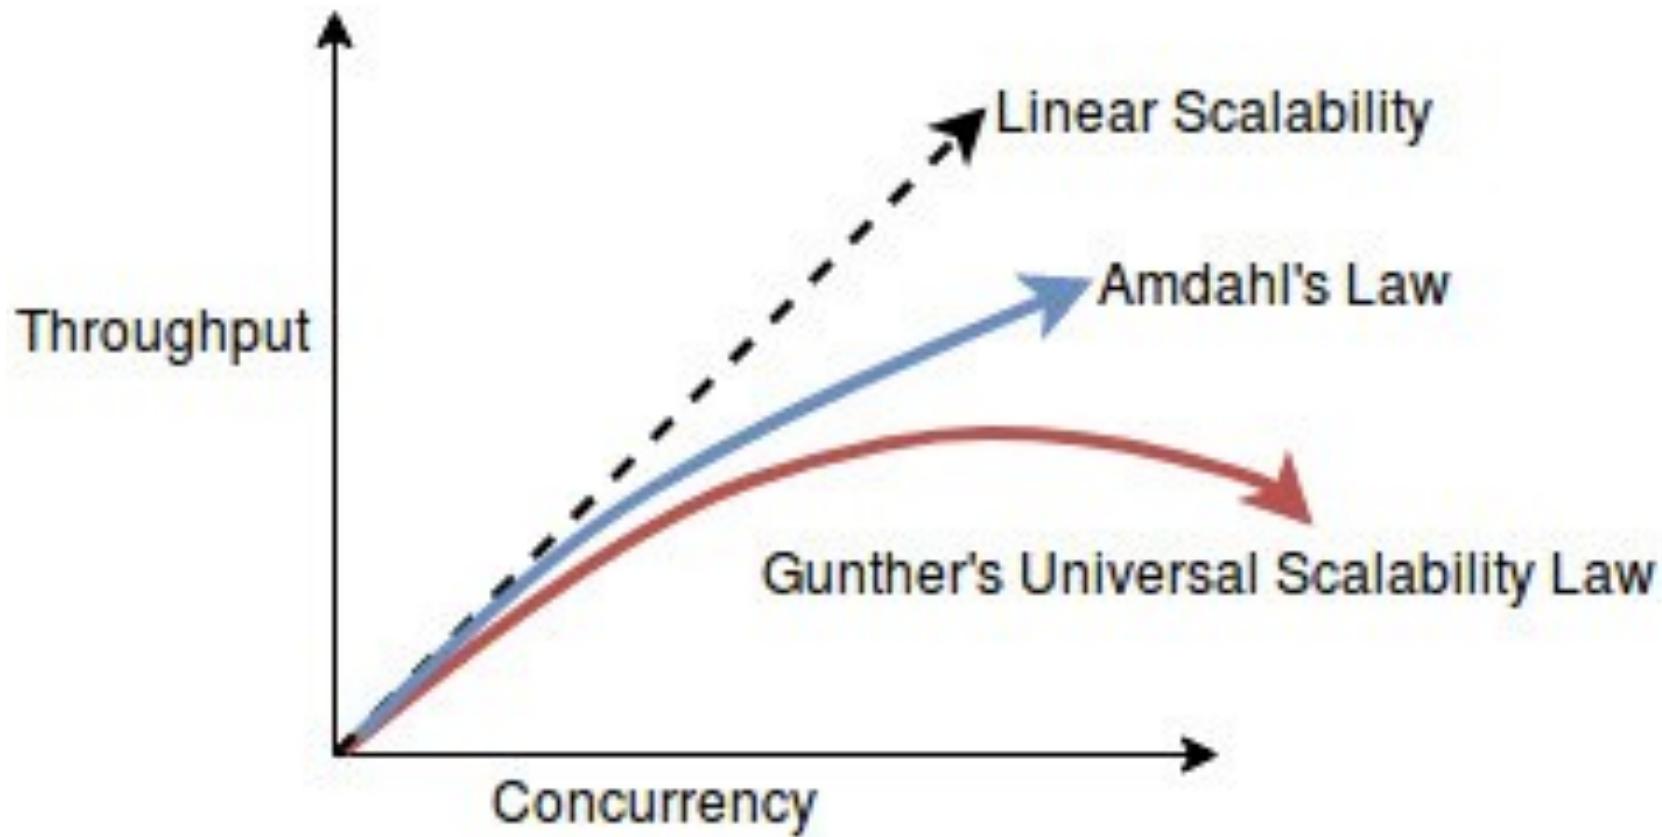
Transaction open for Long Time can get very expensive (even if it performed no writes)

Isolation Mode Matters

SET AUTOCOMMIT=0 - Any SELECT query will Open Transaction

COMMIT/ROLLBACK closes transaction

Understanding Optimal Concurrency



Embrace Queueing

Request Queueing is Normal

With requests coming at “Random Arrivals” some queueing will happen with any system scale

Should not happen to often or for very long

Queueing is “Cheaper” Close to the User

Benefits of Connection Pooling

Avoiding
Connection
Overhead,
especially TLS/SSL

Avoiding using
Excessive Number
of Database
Connections

Multiplexing/Load
Management

Law of Gravity

Shitty Application at
scale will bring down any
Database

Scale Matters

Developing and Testing with Toy Database is risky

Queries Do not slow down linearly

The slowest query may slow down most rapidly

Memory or Disk

Data Accessed in memory is much faster than on disk

It is true even with modern SSDs

SSD accesses data in large blocks, memory does not

Fitting data in Working Set

Newer is not Always Faster

Upgrading to the
new Software/
Hardware is not
always faster

Test it out

Defaults Change are
often to blame

Upgrades are needed but not seamless

Major Database Upgrades often
require application changes

Having Conversation on Application
Lifecycle is a key

Character Sets

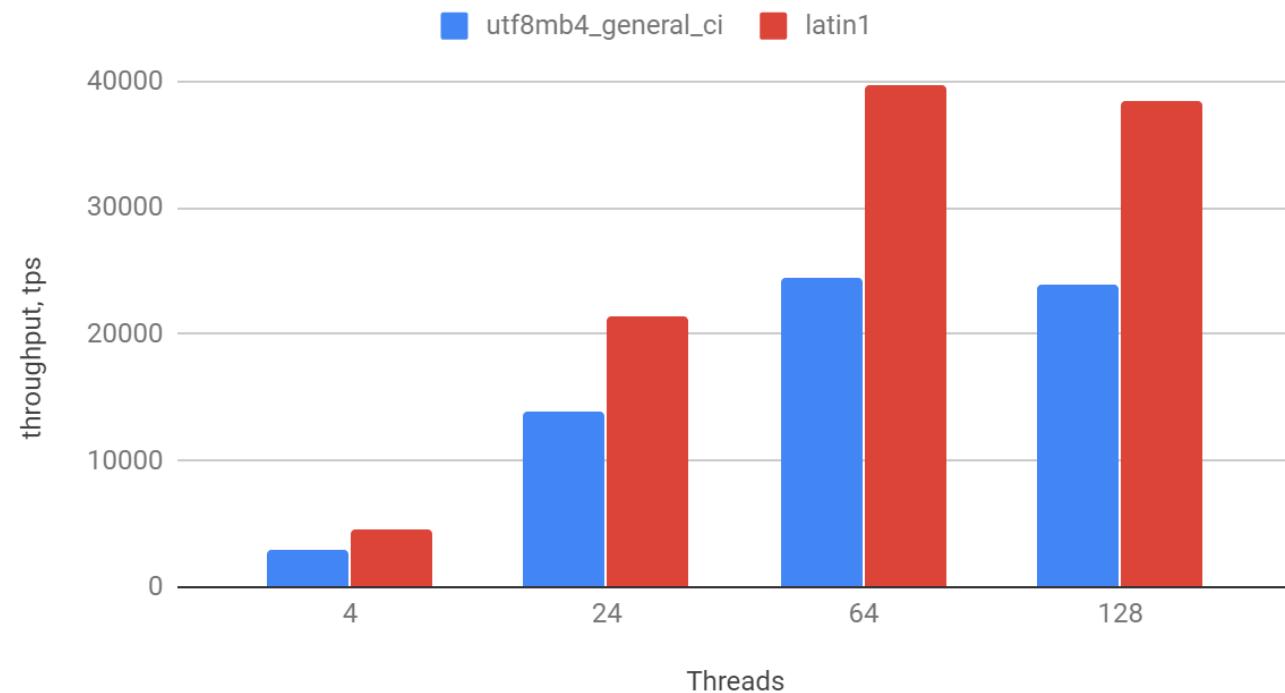
Performance Impact

Pain to Change

Wrong Character Set can cause Data Loss

Character Sets

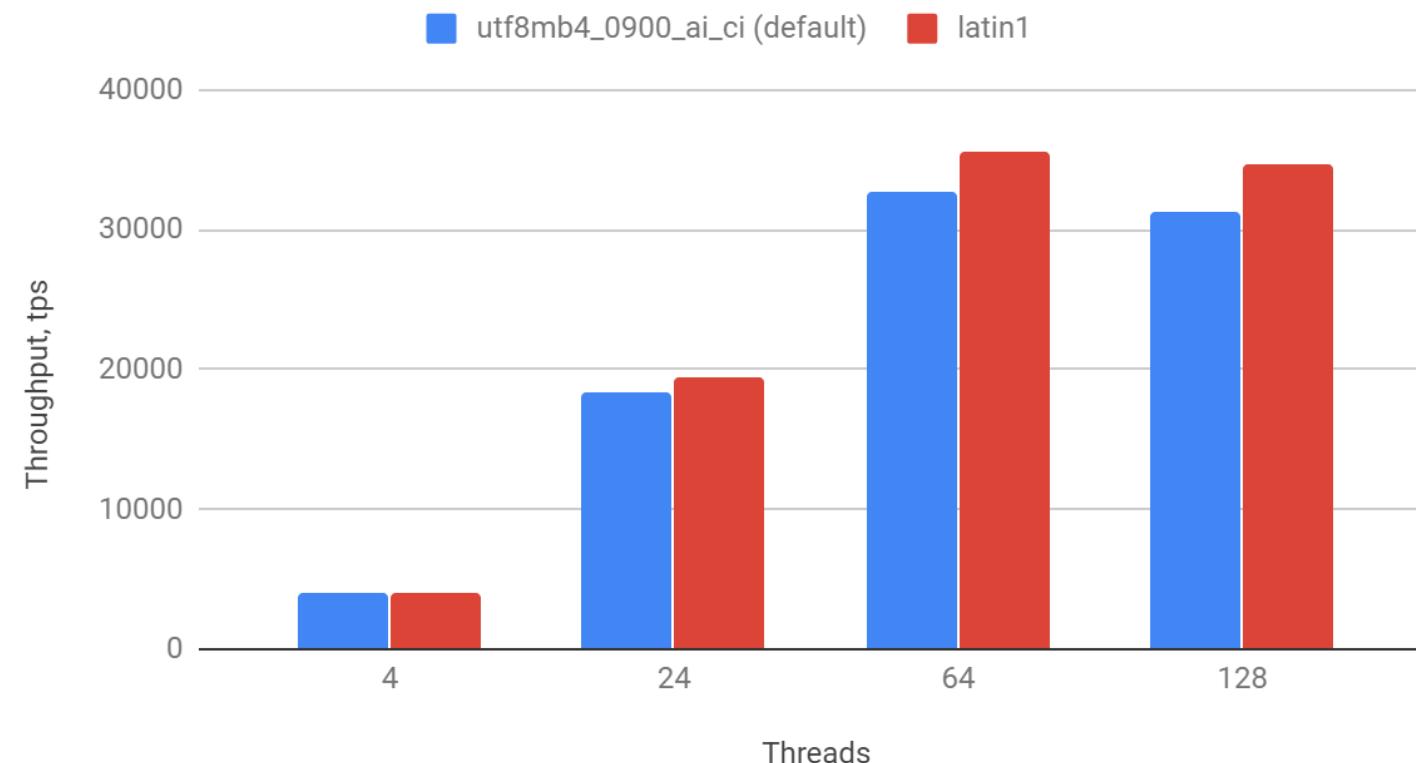
MySQL 5.7 utf8mb4_general_ci (default) and latin1



<https://per.co.na/MySQLCharsetImpact>

Less impact In MySQL 8

MySQL 8.0 utf8mb4_0900_ai_ci and latin1



Operational Overhead

Operations Take Time, Cost Money, Cause Overhead

10TB Database Backup ?

Adding The Index to Large Table ?

Risks of Automation

Automation is
Must

Mistakes can
destroy
database at scale

Security

Database is where the most sensitive
data tends to live

Shared Devs and Ops Responsibility



PERCONA

LIVEONLINE

May 19, 10 AM EDT to May 20, 10 AM EDT

THE FIRST 24-HOUR, AROUND-THE-WORLD OPEN SOURCE DATABASE CONFERENCE





PERCONA | LIVEONLINE

SPEAKER COMPANIES

19 -20 May 2020



ORACLE



Yandex



Pythian



Open Source Database Survey



Help us to understand your Open Source Database Usage Better



This will help us to build better Open Source Database technology, so You can benefit



Results are 100% Anonymous and Open Source (CC BY-SA)



Do it now - it takes just 5-10 minutes

Thank You!

Twitter: @percona @peterzaitsev
