

# Learn how to Analyze & Tune MySQL Queries for Better Performance

Tools and techniques for faster queries

---

Bradley Mickel

MySQL DBA

Thursday June 21st



# Tools

---

Overview of common tools used to identify query optimizations

# EXPLAIN

---

EXPLAIN SELECT \* FROM staff;

EXPLAIN SELECT \* FROM staff\G

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	staff	NULL	ALL	NULL	NULL	NULL	NULL	2	100.00	NULL

```
      id: 1
select_type: SIMPLE
      table: staff
  partitions: NULL
         type: ALL
possible_keys: NULL
          key: NULL
       key_len: NULL
          ref: NULL
         rows: 2
   filtered: 100.00
      Extra: NULL
1 row in set, 1 warning (0.00 sec)
```

# JSON FORMAT EXPLAIN

EXPLAIN FORMAT=JSON SELECT count(1) FROM sakila.rental WHERE last\_update > '2006-02-18 00:00:00'\G

```
EXPLAIN: {
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "1.16"
    },
    "table": {
      "table_name": "rental",
      "access_type": "range",
      "possible_keys": [
        "last_update"
      ],
      "key": "last_update",
      "used_key_parts": [
        "last_update"
      ],
      "key_length": "4",
      "rows_examined_per_scan": 1,
      "rows_produced_per_join": 1,
      "filtered": "100.00",
      "using_index": true,
      "cost_info": {
        "read_cost": "1.06",
        "eval_cost": "0.10",
        "prefix_cost": "1.16",
        "data_read_per_join": "33"
      },
      "used_columns": [
        "last_update"
      ],
      "attached_condition": "(`sakila`.`rental`.`last_update` > '2006-02-18 00:00:00')"
```

# EXPLAIN Fields

Field	JSON FIELD	Description
id	select_id	The sequential number of the SELECT within the query
select_type		The Type of select being performed.
table	table_name	The name, or alias of the table related to this step of the queries
partitions	partitions	The partitions from which records would be matched by the query. NULL for non partitioned tables
type	access_type	The type of access being used to retrieve data or join tables
possible_keys	possible_keys	Which indexes are being reviewed for use in the query
key	key	The index which will be used for the query
key_len	key_length	The length of the key that MySQL has decided to use. This enables you to determine how many parts of a composite index are being used.*
ref	ref	Shows which columns or constants are compared to the index named in the key column to select rows from the table
rows	rows	The number of rows MySQL estimates will be need to be reviewed
filtered	filtered	An estimated percentage of table rows that will be filtered by the condition. 100% means all the rows returned by the condition will be joined with previous tables
extra		Additional information related to the optimizer's chosen path.

# JSON FORMAT EXPLAIN

```
EXPLAIN: {
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "0.62"
    },
    "table": {
      "table_name": "rental",
      "access_type": "range",
      "possible_keys": [
        "last_update"
      ],
      "key": "last_update",
      "used_key_parts": [
        "last_update"
      ],
      "key_length": "4",
```

## ■ Additional Fields:

- *cost\_Info*
  - *query\_cost*
- *used\_key\_parts*

# JSON FORMAT EXPLAIN

```
    "rows_examined_per_scan": 1,  
    "rows_produced_per_join": 1,  
    "filtered": "100.00",  
    "using_index": true,  
    "cost_info": {  
      "read_cost": "0.52",  
      "eval_cost": "0.10",  
      "prefix_cost": "0.62",  
      "data_read_per_join": "32"  
    },  
    "used_columns": [  
      "last_update"  
    ],  
    "attached_condition": "(`sakila`.`rent  
al`.`last_update` > '2006-02-18 00:00:00')"  
  }  
}  
}
```

## ■ Additional Fields:

- *rows\_examined\_per\_scan*
- *rows\_examined\_per\_join*
- *cost\_info*
  - *read\_cost*
  - *eval\_cost*
  - *prefix\_cost*
  - *data\_read\_per\_join*
- *used\_columns*

# Access Types

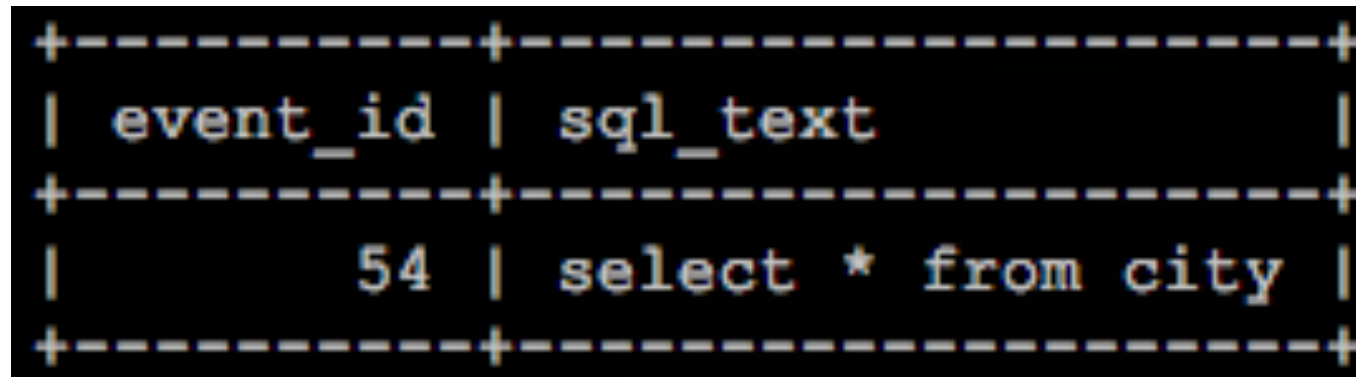
Type	Description
constant	The table has at most one matching row which is read at the start of the query
system	A type of constant table used for systems tables
eq_ref	Used for table joins with a 1:1 relationship. Requires the joining index by a Primary or Unique Key
ref	All rows with matching index values are read from this table for each combination of rows from the previous tables
fulltext	A fulltext Index is being used
ref_or_null	Similar to ref, but will find null values in addition to the selected value
index_merge	Index Merge is being used for the join
unique_subquery	Replaces eq_ref for certain subqueries using IN
index_subquery	Similar to unique_subquery but for when the index does not have a unique constraint
range	Only rows that are in a given range are retrieved, using an index to select the rows. The key column in the output row indicates which index is used. The key_len contains the longest key part that was used. The ref column is NULL for this type.
index	The entire index is scanned
all	A full table scan is required



# Performance Schema Profiling

---

```
SELECT event_id, sql_text FROM  
performance_schema.events_statements_history_long  
WHERE sql_text like '%select * from city%';
```

A terminal window screenshot showing a table with two columns: event\_id and sql\_text. The table has one data row where event\_id is 54 and sql\_text is 'select \* from city'. The table is enclosed in a dashed border with '+' characters at the corners and intersections.

event_id	sql_text
54	select * from city

# Performance Schema Profiling cont.

```
use performance_schema;

SELECT
  event_name AS Stage,
  round(timer_wait/
    pow(10,12),6) AS Duration
FROM
  events_stages_history_long
WHERE nesting_event_id = 54;
```

Stage	Duration
stage/sql/starting	0.000098
stage/sql/checking permissions	0.000001
stage/sql/Opening tables	0.000029
stage/sql/init	0.000007
stage/sql/System lock	0.000005
stage/sql/cptimizing	0.000001
stage/sql/statistics	0.000014
stage/sql/preparing	0.000009
stage/sql/executing	0.000006
stage/sql/Sending data	0.000676
stage/sql/end	0.000002
stage/sql/query end	0.000007
stage/sql/closing tables	0.000010
stage/sql/freeing items	0.000047
stage/sql/cleaning up	0.000003

# Pt-query-digest

---

- Analyses MySQL queries from the slow, general, and binary logs
- Groups queries by fingerprint
- provides average runtime information

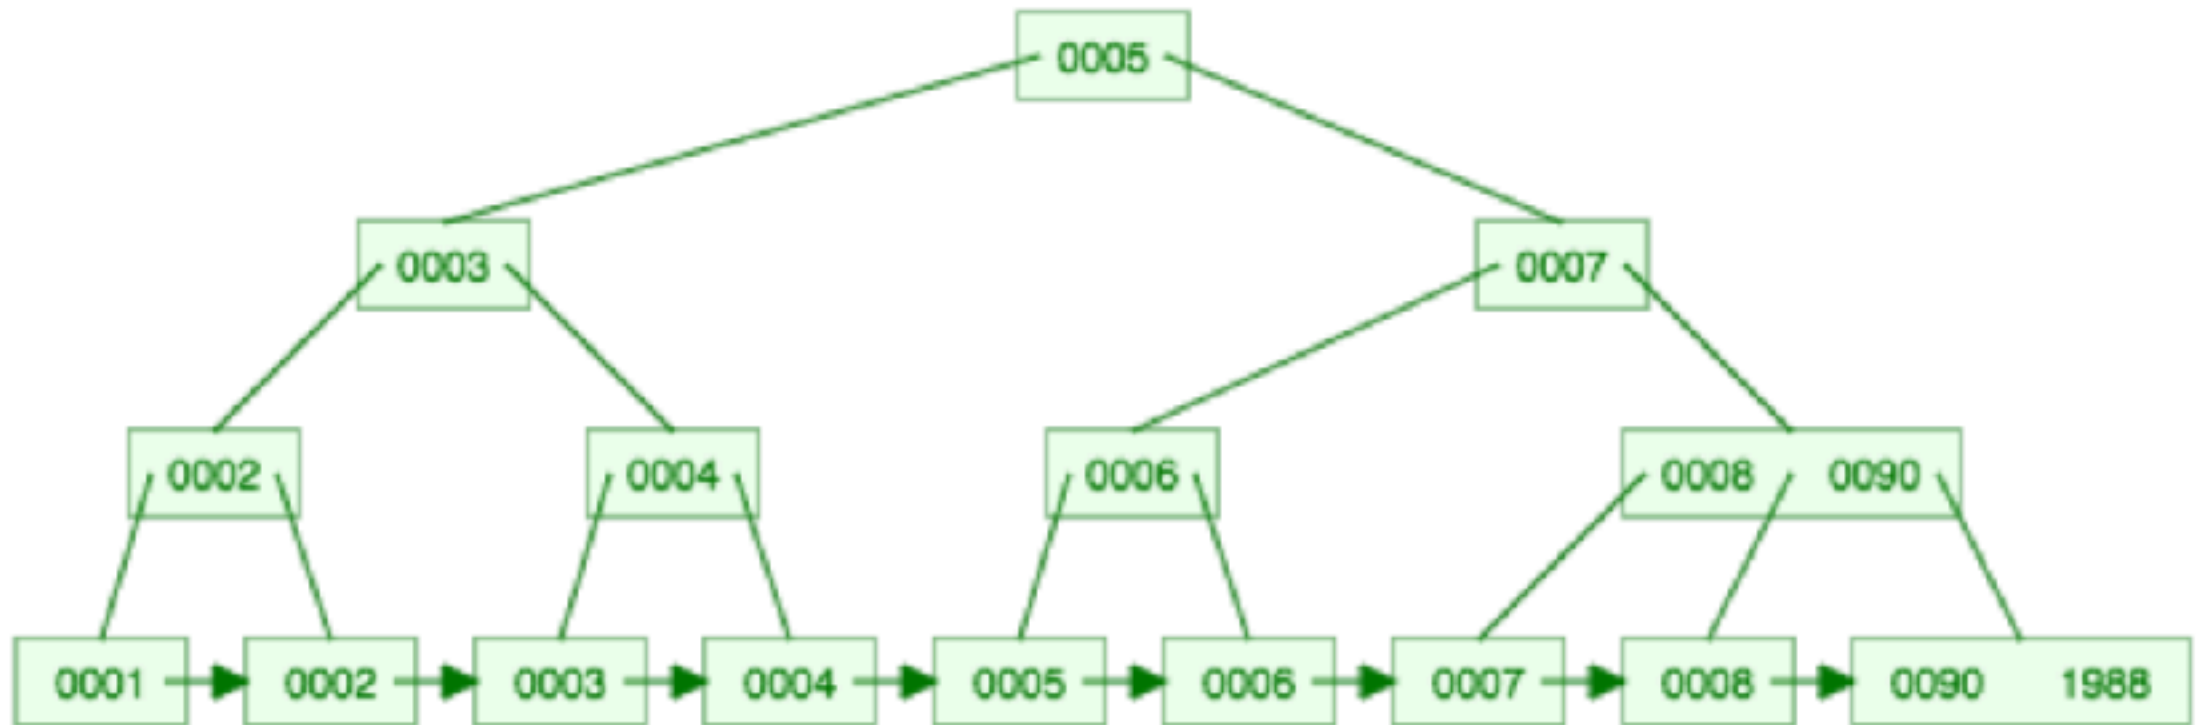
# Indexes

---

## Index concepts and limitations

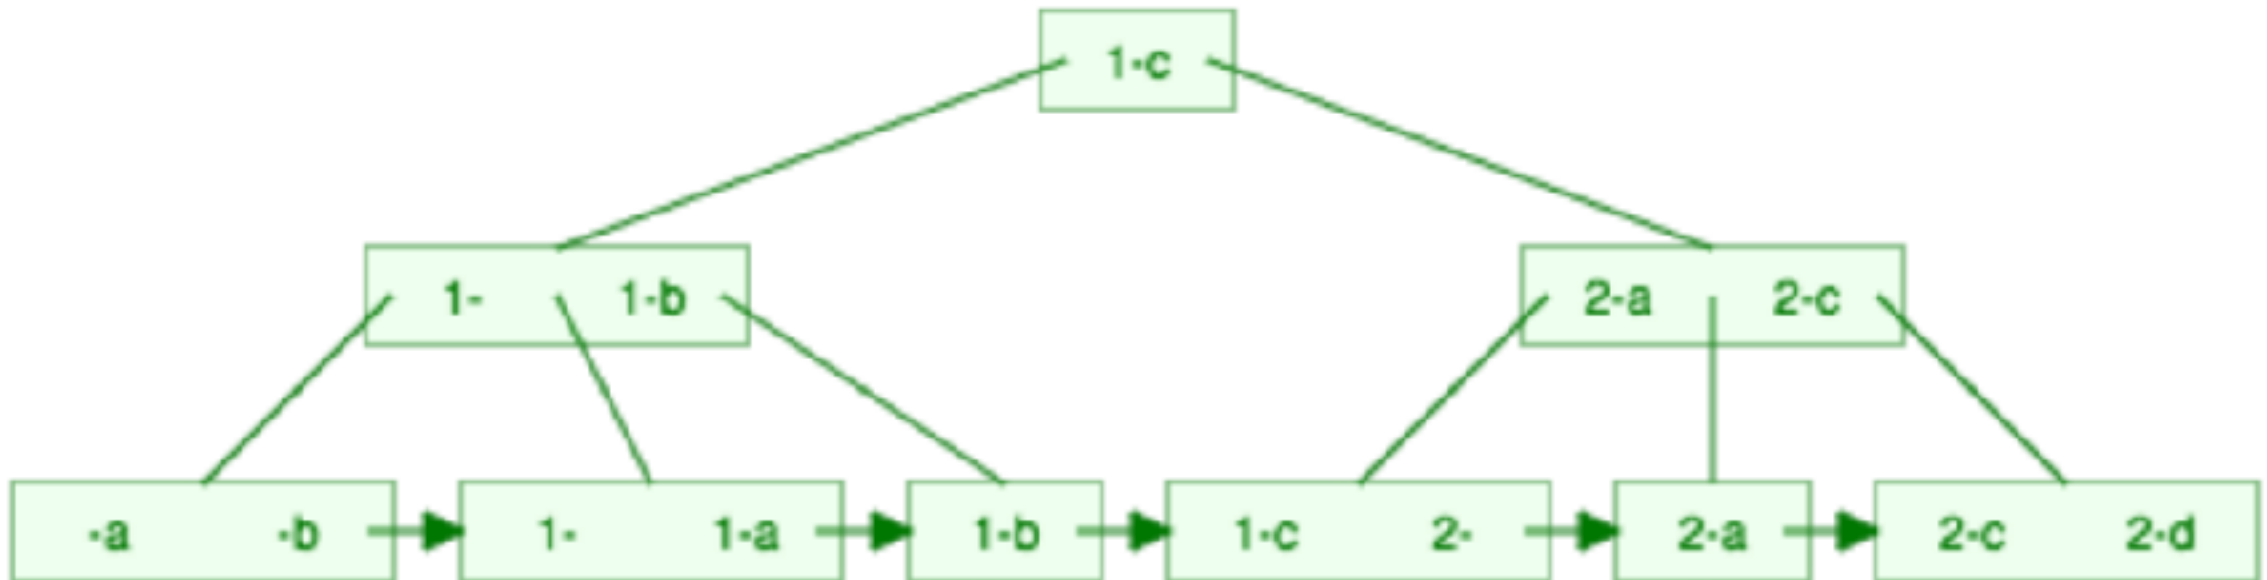
# Balanced Tree Index

---



# Composite Balanced Tree Index

---



# Primary Key

---

## Requirements and constraints

- Field(s) MUST be UNIQUE
- NULL is not valid
- Only 1 per table
- Can be a single or multiple fields

**Controls how data is stored physically**

**Is utilized as part of ALL secondary indexes**

# Index Impact

```
SELECT SQL_NO_CACHE count(1) FROM sakila.rental WHERE last_update>'2006-02-18 00:00:00'\G
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	rental	NULL	ALL	NULL	NULL	NULL	NULL	16005	33.33	Using where

```
# Query_time distribution
```

```
# 1us
```

```
# 10us
```

```
# 100us
```

```
# 1ms #####
```

```
# 10ms #####
```

```
# 100ms
```

```
# 1s
```

```
# 10s+
```



# Index Impact cont

SELECT count(1) FROM sakila.rental WHERE last\_update>'2006-02-18 00:00:00'\G

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	rental	NULL	range	last_update	last_update	4	NULL	1	100.00	Using index condition

```
# Query_time distribution
# 1us
# 10us
# 100us #####
# 1ms #
# 10ms
# 100ms
# 1s
# 10s+
```

# Multiple Tables

SELECT f.rating,sum(p.amount) as income FROM film f JOIN inventory i ON i.film\_id = f.film\_id JOIN rental r ON r.inventory\_id = i.inventory\_id JOIN payment p ON p.rental\_id=r.rental\_id GROUP BY f.rating ORDER BY income DESC;

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	f	NULL	ALL	PRIMARY	NULL	NULL	NULL	1000	100.00	Using temporary; Using filesort
1	SIMPLE	i	NULL	ref	PRIMARY,idx_fk_film_id	idx_fk_film_id	2	sakila.f.film_id	4	100.00	Using index
1	SIMPLE	r	NULL	ref	PRIMARY,idx_fk_inventory_id	idx_fk_inventory_id	3	sakila.i.inventory_id	3	100.00	Using index
1	SIMPLE	p	NULL	ref	fk_payment_rental	fk_payment_rental	5	sakila.r.rental_id	1	100.00	NULL

```
# Query_time distribution
# 1us
# 10us
# 100us
# 1ms
# 10ms #####
# 100ms ##
# 1s
# 10s+
```

# Multiple Tables

SELECT f.rating,sum(p.amount) as income FROM film f JOIN inventory i ON i.film\_id = f.film\_id JOIN rental r ON r.inventory\_id = i.inventory\_id JOIN payment p ON p.rental\_id=r.rental\_id GROUP BY f.rating ORDER BY income DESC;

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	f	NONE	index	PRIMARY, rating	rating	2	NONE	1000	100.00	Using index; Using temporary; Using filesort
1	SIMPLE	i	NONE	ref	PRIMARY, idx_fk_film_id	idx_fk_film_id	2	sakila.f.film_id	4	100.00	Using index
1	SIMPLE	r	NONE	ref	PRIMARY, idx_fk_inventory_id	idx_fk_inventory_id	3	sakila.i.inventory_id	3	100.00	Using index
1	SIMPLE	p	NONE	ref	PRIMARY, idx_fk_payment_rental	fk_payment_rental	5	sakila.r.rental_id	1	100.00	NONE

```
# Query_time distribution
# 1us
# 10us
# 100us
# 1ms
# 10ms #####
# 100ms #
# 1s
# 10s+
```

# Index Limitations

---

## Functions

- Aggregates
  - *SUM, AVG*
  - *Does not include MIN, MAX*
- Multiple Range scans
  - *multiple OR statements*
  - *multiple less than or greater than*
- Character Set Differences
- Data Type Differences

# Multiple Range Scans

SELECT count(1) FROM rental WHERE customer\_id=148 AND return\_date BETWEEN '2005-08-29 00:00:00' AND '2005-09-05 23:59:59' AND last\_update BETWEEN '2006-02-15 00:00:00' AND '2006-02-15 23:59:59';

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	rental	NULL	ref	idx_fk_customer_id,last_update	idx_fk_customer_id	2	const	46	5.55	Using where

Add a composite index on customer\_id,return\_date,last\_update

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	rental	NULL	range	idx_fk_customer_id,last_update,c_r_1	c_r_1	12	NULL	4	50.00	Using where; Using index

# Multiple Range Scans

```
EXPLAIN: [
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "1.81"
    },
    "table": {
      "table_name": "rental",
      "access_type": "range",
      "possible_keys": [
        "idx_fk_customer_id",
        "last_update",
        "c_r",
        "c_l",
        "c_r_l"
      ],
      "key": "c_r_l",
      "used_key_parts": [
        "customer_id",
        "return_date"
      ],
      "key_length": "12",
      "rows_examined_per_scan": 4,
      "rows_produced_per_join": 1,
      "filtered": "50.00",
      "using_index": true,
      "cost_info": {
        "read_cost": "1.62",
        "eval_cost": "0.20",
        "prefix_cost": "1.82",
        "data_read_per_join": "63"
      },
      "used_columns": [
        "customer_id",
        "return_date",
        "last_update"
      ],
      "attached_condition": "(('sakila','rental','customer_id' = 148) and (('sakila','rental','return_date' between '2005-08-29 00:00:00' and '2005-09-05 23:59:59') and (('sakila','rental','last_update' between '2005-02-15 00:00:00' and '2006-02-15 23:59:59')))"
    },
    "1"
  ]
}
```

# Multiple Range Scans

SELECT COUNT(1) FROM rental WHERE customer\_id=148 AND (return\_date BETWEEN '2005-08-29 00:00:00' AND '2005-09-05 23:59:59' OR last\_update BETWEEN '2006-02-15 00:00:00' AND '2006-02-15 23:59:59');

```
EXPLAIN: {
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "5.65"
    },
    "table": {
      "table_name": "rental",
      "access_type": "ref",
      "possible_keys": [
        "idx_fk_customer_id",
        "last_update",
        "c_r_1"
      ],
      "key": "c_r_1",
      "used_key_parts": [
        "customer_id"
      ],
      "key_length": "2",
      "ref": [
        "const"
      ],
      "rows_examined_per_scan": 46,
      "rows_produced_per_join": 9,
      "filtered": "20.99%",
      "using_index": true,
      "cost_info": {
        "read_cost": "1.05",
        "eval_cost": "0.97",
        "prefix_cost": "5.65",
        "data_read_per_join": "308"
      },
      "used_columns": [
        "customer_id",
        "return_date",
        "last_update"
      ],
      "attached_condition": "(((askila,'rental','return_date' between '2005-08-29 00:00:00' and '2005-09-05 23:59:59') or ('askila','rental','last_update' between '2006-02-15 00:00:00' and '2006-02-15 23:59:59')))"
    },
    "join": {
      "join_type": "inner",
      "on_condition": "1"
    }
  }
}
```

# Multiple Range Scans

```
SELECT COUNT(1) FROM rental WHERE customer_id=148 AND return_date  
BETWEEN '2005-08-29 00:00:00' AND '2005-09-05 23:59:59'
```

UNION

```
SELECT COUNT(1) FROM rental WHERE customer_id=148 AND last_update  
BETWEEN '2006-02-15 00:00:00' AND '2006-02-15 23:59:59';
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	PRIMARY	rental	NULL	range	idx_fk_customer_id,c_r_1	c_r_1	8	NULL	4	100.00	Using where; Using index
2	UNION	rental	NULL	ref	idx_fk_customer_id,last_update,c_r_1	c_r_1	2	const	46	50.00	Using where; Using index
NULL	UNION RESULT	<union1,2>	NULL	ALL	NULL	NULL	NULL	NULL	NULL	NULL	Using temporary

Indexes are read left to right so a new index is needed on  
customer\_id,last\_update



# Multiple Range Scans

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	PRIMARY	rental	NULL	range	idx_fk_customer_id,o_r,o_l	c_r	8	NULL	4	100.00	Using where; Using index
2	UNION	rental	NULL	range	idx_fk_customer_id,last_update,o_r,o_l	c_l	5	NULL	46	100.00	Using where; Using index
NULL	UNION RESULT	<union1,2>	NULL	ALL	NULL	NULL	NULL	NULL	NULL	NULL	Using temporary

```
# Query_time distribution
# 1us
# 10us
# 100us #####
# 1ms ####
# 10ms
# 100ms
# 1s
# 10s+
```

OR

UNION

```
# Query_time distribution
# 1us
# 10us
# 100us #####
# 1ms ####
# 10ms
# 100ms
# 1s
# 10s+
```

# References

---

<https://dev.mysql.com/doc/sakila/en/sakila-installation.html>

<https://dev.mysql.com/doc/refman/5.7/en/explain-output.html#explain-join-types>

<https://dev.mysql.com/doc/mysql-perfschema-excerpt/5.7/en/performance-schema-query-profiling.html>

<https://dev.mysql.com/doc/refman/5.7/en/general-thread-states.html>

<https://www.cs.usfca.edu/~galles/visualization/BPlusTree.html>

<https://www.percona.com/software/database-tools/percona-toolkit>

<https://bugs.mysql.com/bug.php?id=83062>