

MySQL Query Optimization

Peter Zaitsev, CEO, Percona

March 8, 2019

SCALE 17x
Pasadena, CA



About Percona

Open Source Database Solutions Company

Support, Managed Services, Consulting, Training, Engineering

Focus on MySQL, MariaDB, MongoDB, PostgreSQL

Support Cloud DBaaS Variants on major clouds

Develop Database Software and Tools

Release Everything as 100% Free and Open Source

Widely Deployed Open Source Software



PERCONA
Server for MySQL

5,000,000+ downloads



PERCONA
Monitoring and Management

175,000+ downloads



PERCONA
XtraBackup

4,500,000+ downloads



PERCONA
Server for MongoDB

450,000+ downloads



PERCONA
Toolkit

2,000,000+ downloads



PERCONA
XtraDB Cluster

1,500,000+ downloads

About the Presentation

Cover the Basics

How MySQL Executes Queries

How To Find Queries to Optimize

How to Optimize Them

The Basics

Grand Goal

**Application which Has a
Great Performance**

Great Performance Defined

Responds With Low Response Time

At All Times

For All Users

Response Time and Database

Database is not always at fault

Database Making your Application Slow

Dev Issues

Ops Issues

Dev Issues:

Many Queries executed serially

Expensive Queries

Poorly Designed Queries

Poorly Optimized Queries

Saturation with Additional Load

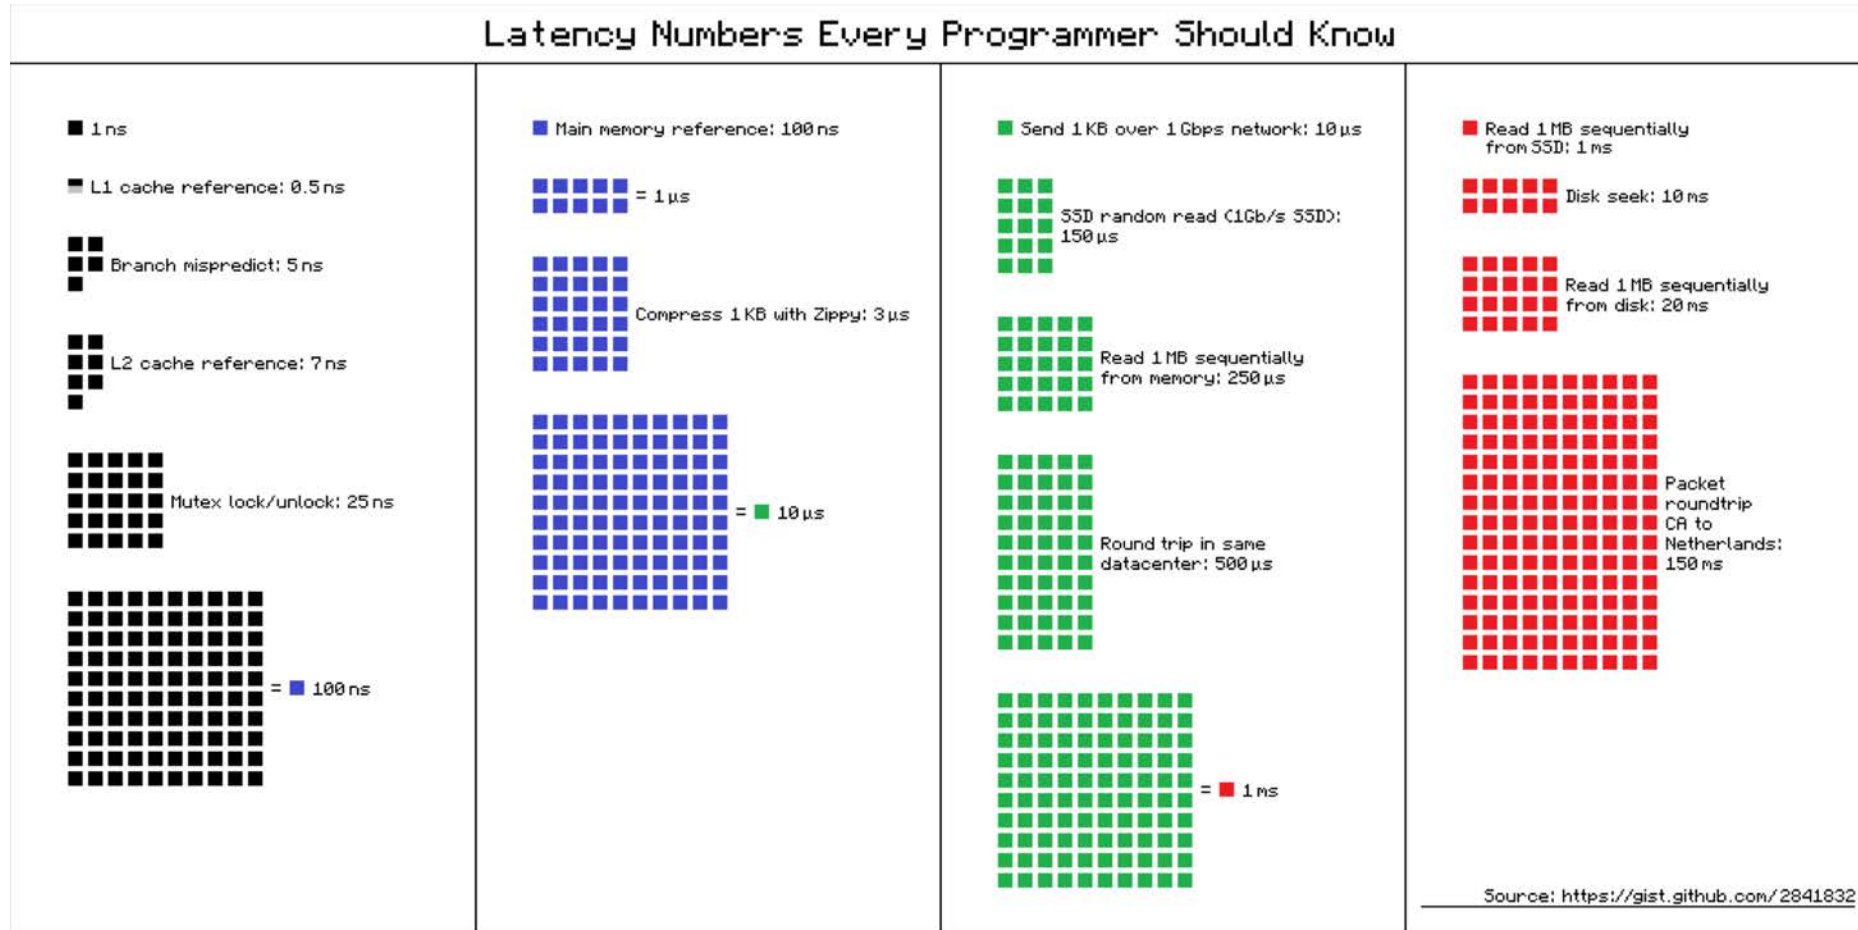
Ops Issues

Problems with System, Storage or Network

Saturation with Additional Load

Capacity Planning

Mind Network Latency



Query Optimization Goals

Specific User Interaction

Application As a Whole

Improving Efficiency

Assuring Scalability

Not Query Optimization Alone

General Architecture

Right Choice of Technology (Not Only MySQL)

Hardware/Instance Properties

OS and MySQL Configuration

Database Schema

How MySQL Executes Queries

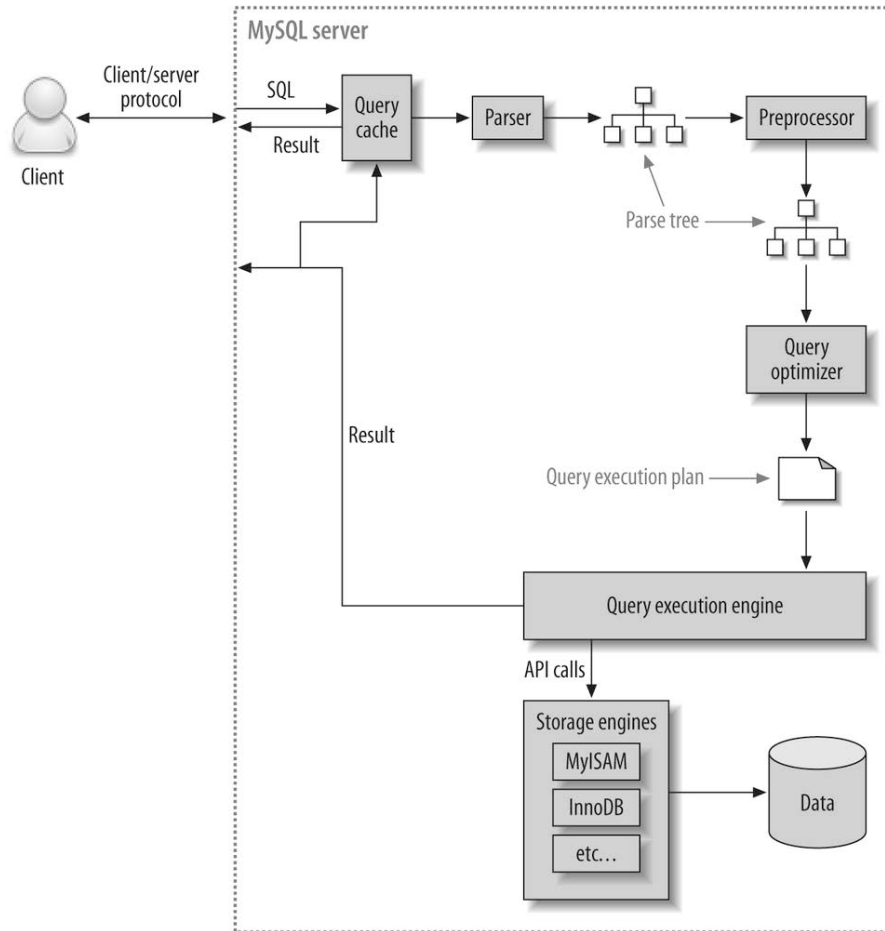
Execution Basics

Single Server

Single Thread (Using Single CPU Core)

No Intermediate Results Caching Between Query Executions

Query Execution Diagram



Added Complexities

UDFs (User Defined Functions)

Stored Programs

VIEWs

Use the LIMIT

Do not just stop fetching rows at the application side

MySQL Client-Server Protocol is NOT cursor based

Join Order

Permanent and “Derived” Tables are going to be “Joined in Order”

MySQL Starts from one table, finding all needed rows in it, and iterating finding matching rows from the next one

Join Order Is Critical For Performance

SELECT STRAIGHT_JOIN to force join order

Indexes

Proper Indexes are must have for Optimal Query Execution

Can improve Query Performance 1000x or more

Expensive to Maintain... so Do not Overdo

Covering Indexes to speed up data reads

Indexes are not Free

Space on Disk

Space in Memory

Extra Optimizer Load to Evaluate Them

Expensive to Maintain with Updates

Columns

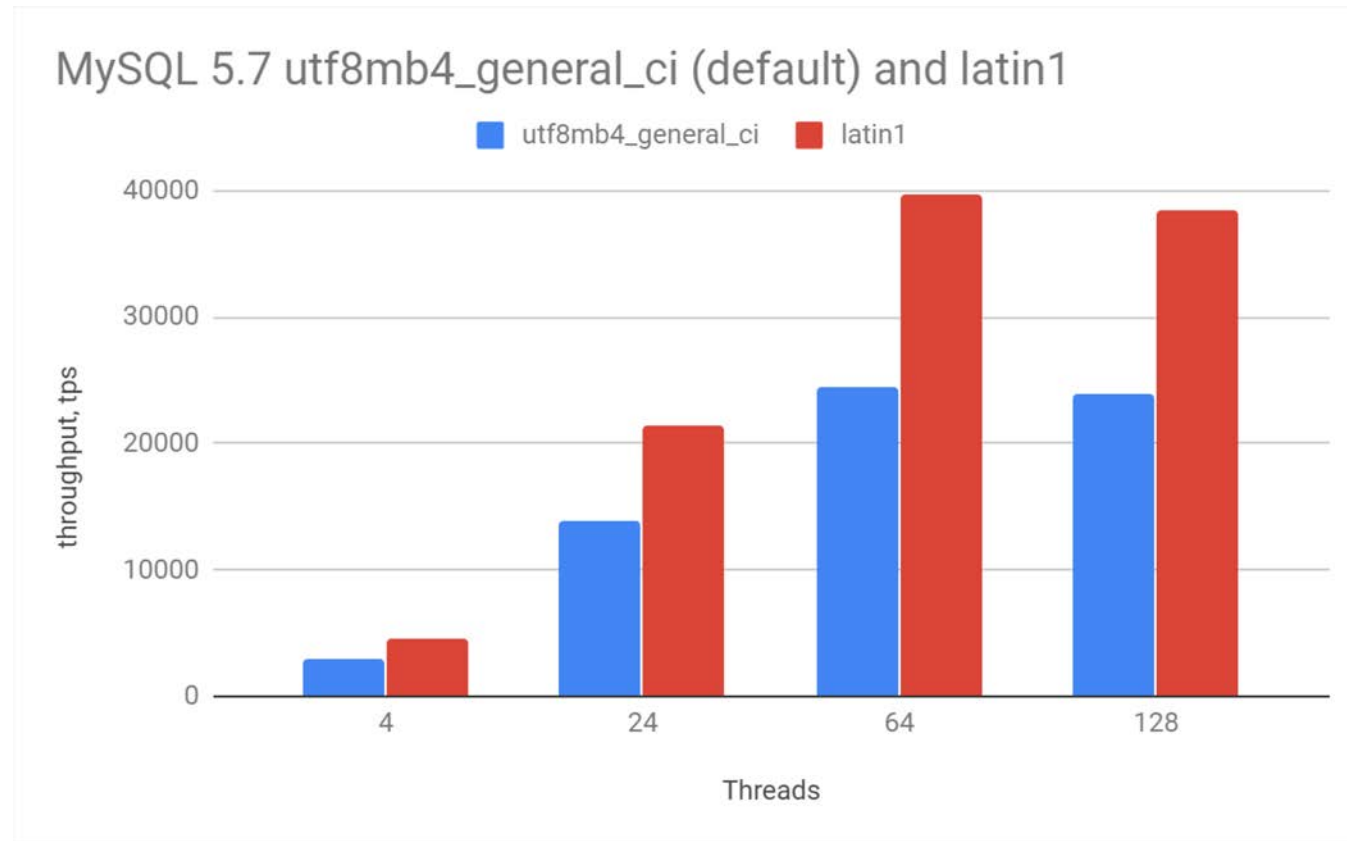
MySQL (Innodb, MyISAM, MyRocks etc) store data row by row

All columns must be read on every row access (excluding Blobs for Innodb)

Number of Total Columns, Their Size Impacts Query Performance a Lot

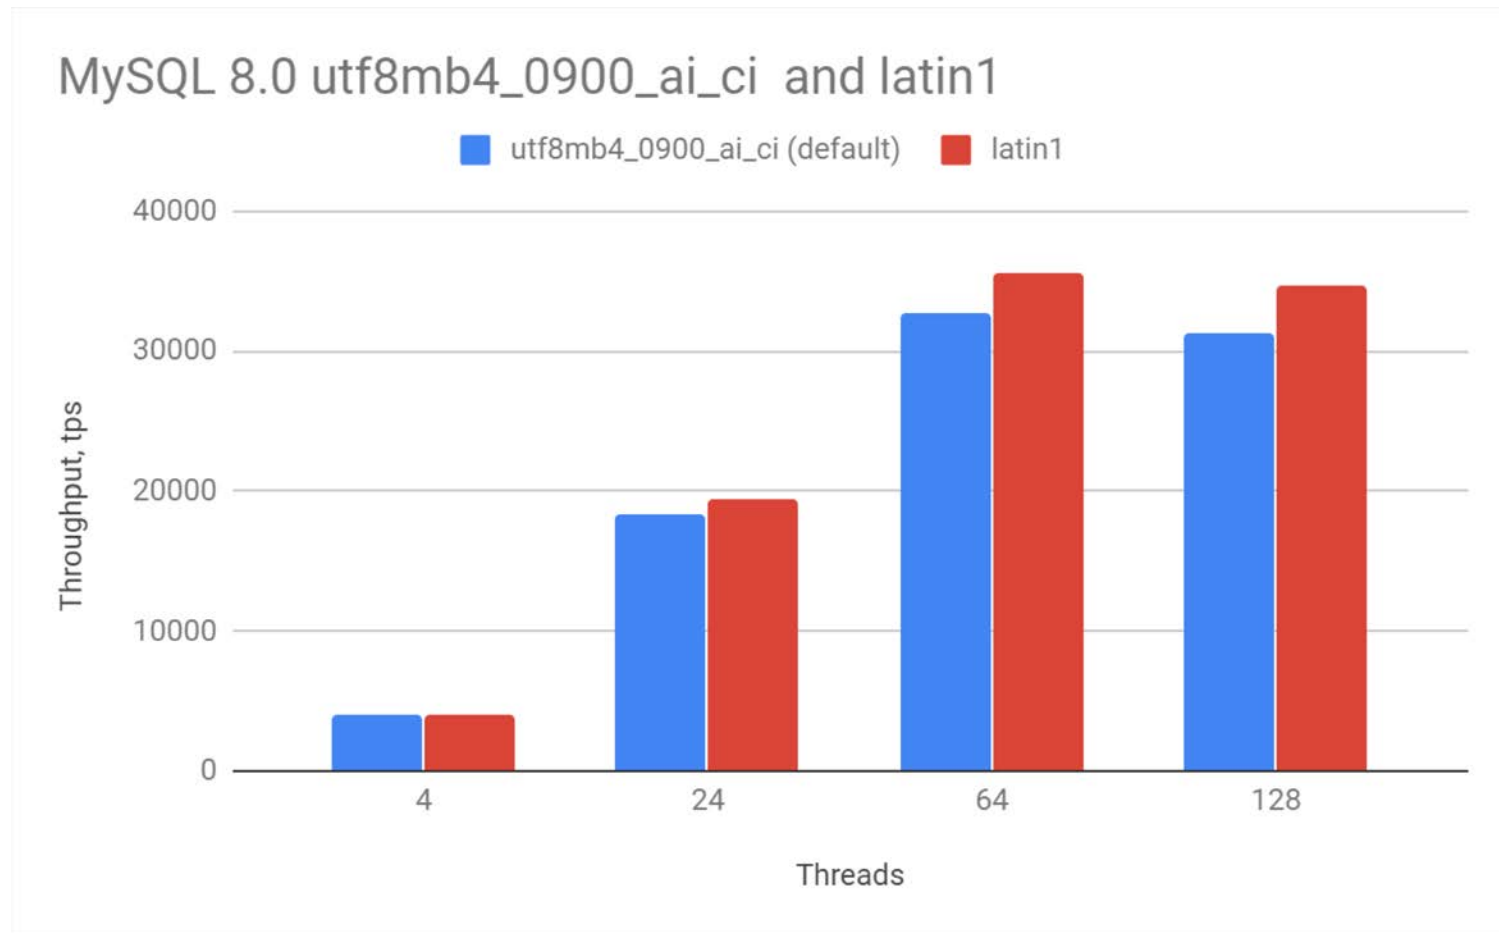
Covering Indexes are great to reduce amount of data query Touches

Character Sets



<https://per.co.na/MySQLCharsetImpact>

Less impact In MySQL 8



Grouping and Sorting

Can use Index, External Sort, Temporary File

Temporary Table can be in memory or on disk

Amount of Data you Sort, Group Matters A Lot

Too many Different Algorithms to Cover in Details

The Mysterious Optimizer

No one knows how MySQL Optimizer Really works

Designed to Choose Best Plan Based on Cost

Cost Model is just a model

Relies on Statistics which can be very wrong

Learn what MySQL Execution Can Do

Not Everything you can imagine can be done by MySQL during execution

Though it also has tricks in its sleeve you may not aware of

Are you Smarter than Optimizer ?

Use Optimizer Hints to Execute Query The way you Like

Often the plan you think is faster is not

<https://dev.mysql.com/doc/refman/8.0/en/optimizer-hints.html>

Learn EXPLAIN

A way to understand how MySQL Expects to Execute Query

Plan May Change based on constants, time server

EXPLAIN SELECT ...

EXPLAIN FORMAT=JSON SELECT ...

<https://dev.mysql.com/doc/refman/8.0/en/using-explain.html>

Explain Example

```
mysql> explain select max(season_nr) from title group by production_year;
```

| id | select_type | table | type | possible_keys | key | key_len | ref | rows | Extra |
|----|-------------|-------|-------|---------------|-----------------|---------|------|------|--------------------------|
| 1 | SIMPLE | title | range | NULL | production_year | 5 | NULL | 201 | Using index for group-by |

1 row in set (0.01 sec)

EXPLAIN EXTENDED

```
mysql> EXPLAIN
      SELECT t1.a, t1.a IN (SELECT t2.a FROM t2) FROM t1\G
```

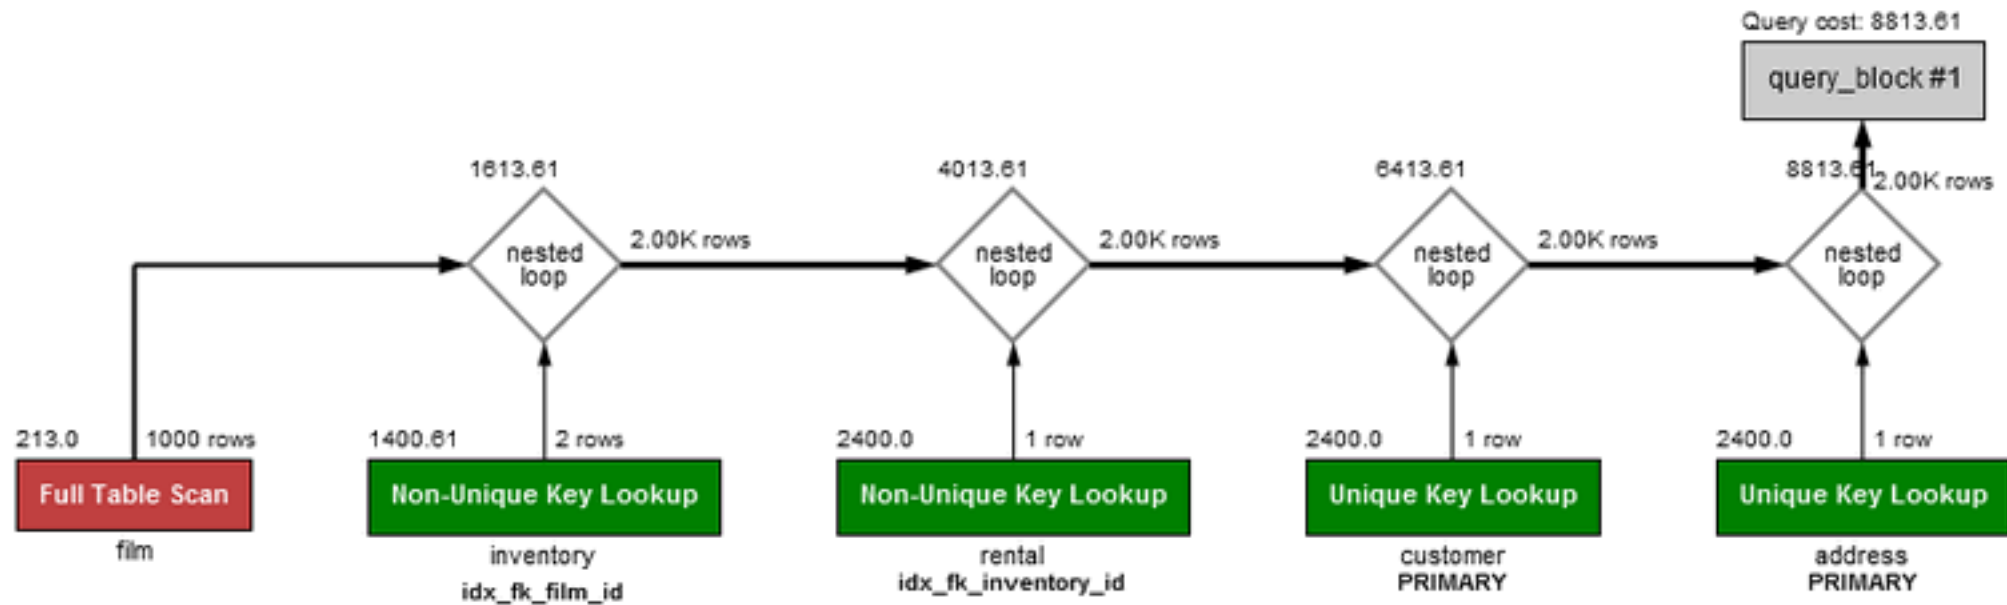
```
mysql> SHOW WARNINGS\G
```

```
***** 1. row *****
Level: Note
Code: 1003
Message: /* select#1 */ select `test`.`t1`.`a` AS `a`,
      <in_optimizer>(`test`.`t1`.`a`,`test`.`t1`.`a` in
      ( <materialize> (/* select#2 */ select `test`.`t2`.`a`
      from `test`.`t2` where 1 having 1 ),
      <primary_index_lookup>(`test`.`t1`.`a` in
      <temporary table> on <auto_key>
      where ((`test`.`t1`.`a` = `materialized-subquery`.`a`)))) AS `t1.a`
      IN (SELECT t2.a FROM t2)` from `test`.`t1`
1 row in set (0.00 sec)
```


EXPLAIN FORMAT=JSON – More Details

```
1  mysql> explain format=json select dept_name from departments where dept_no in (select dept_no from dep
2  **** 1. row ****
3  EXPLAIN: {
4    "query_block": {
5      "select_id": 1,
6      "cost_info": {
7        "query_cost": "16.72"
8      },
9      "nested_loop": [
10       {
11         "table": {
12           "table_name": "departments",
13           "<skipped>"
14         },
15         {
16           "table": {
17             "table_name": "<subquery2>",
18             "access_type": "eq_ref",
19             "key": "<auto_key>",
20             "key_length": "4",
21             "ref": [
22               "employees.departments.dept_no"
23             ],
24             "rows_examined_per_scan": 1,
25             "materialized_from_subquery": {
26               "using_temporary_table": true,
27               "query_block": {
28                 "table": {
```

MySQL WorkBench Visualization



<https://www.mysql.com/products/workbench/>

Optimizer Trace

- **Advanced Optimizer Debugging if you can't figure out why given plan is chosen**

```
1  # Turn tracing on (it's off by default):  
2  SET optimizer_trace="enabled=on";  
3  SELECT ...; # your query here  
4  SELECT * FROM INFORMATION_SCHEMA.OPTIMIZER_TRACE;  
5  # possibly more queries...  
6  # When done with tracing, disable it:  
7  SET optimizer_trace="enabled=off";
```

How to Find Queries to Optimize

Where Should you Optimize your Queries ?

**Development
Environment**

**Production
Environment**

Reality

Both

- You want to ensure unoptimized queries never make it to Production
- But you will have Query Performance Issues in Production anyway

Development

**Can use MySQL Log (Slow Query Log)
or Application Debugging**

**Can help not only to Optimize Slow
Queries but also Eliminate Waste**

Development and Production

Using Query Analyzes Tools

Percona's Open Source Solution

A screenshot of the Percona Monitoring and Management (PMM) web interface. The background is dark with a grid pattern. In the center, there is a large orange circular logo featuring a stylized parrot. Below the logo, the word "PERCONA" is written in large, white, sans-serif capital letters. Underneath "PERCONA", the text "Monitoring and Management" is displayed in a smaller, white, sans-serif font. Below this, a line of text reads "Free and open-source platform for managing and monitoring MySQL®, MariaDB® and MongoDB® performance." At the bottom of the interface, there are two orange buttons with white text: "VIEW THE PMM DEMO" and "DOWNLOAD PMM LATEST". The interface also shows several line graphs with multiple colored lines (yellow, green, red) representing different performance metrics over time. Time stamps like "06:00", "06:05", "06:10", "06:15", "06:35", and "06:40" are visible along the bottom of the graphs. The text "Load Average" is visible on the left side of the interface, and "Memory Dis" is partially visible on the right side.

PERCONA
Monitoring and Management

Free and open-source platform for managing and monitoring MySQL®, MariaDB® and MongoDB® performance.

[VIEW THE PMM DEMO](#) [DOWNLOAD PMM LATEST](#)

See it Live!

<https://pmmdemo.percona.com>

Top Queries



Things to Consider

Outliers may not be causing the most load




Victims and queries causing the problem

Queries not Finished yet

Query Profile

| ▼ Metrics | | | Query first |
|---------------------|-------------------|---|---|
| Metrics | Rate/Sec | | Sum |
| Query Count | 1.56k (per sec) |  | 67.53m 45.63% of total |
| Query Time | 6.01 load |  | 3 days, 0:07:38 27.67% of total |
| Lock Time | 0.70 (avg load) |  | 8:21:52 23.81% of total 10.40% of query time |
| Innodb IO Read Wait | 3.27 (avg load) |  | 1 days, 15:17:49 39.57% of total 56.09% of query time |
| Innodb Read Ops | 320.50(per sec) |  | 13.85m 34.27% of total |
| Innodb Read Bytes | 5.25 MB (per sec) |  | 226.85 GB 34.53% of total 16.38 KB avg io size |

How Efficiently Query Produces Result ?

| | | | |
|---------------|------------------|---|--|
| Rows Sent | <0.01 (per sec) |  | 57.00 <0.01 of total |
| Bytes Sent | 0.09 (per sec) |  | 3.93 KB <0.01 of total 69.00 Bytes bytes/row |
| Rows Examined | 39.58k (per sec) |  | 1.71b 25.89% of total 30.00m per row sent |

Explain and Table Details

▼ VISUAL

```
Filter with WHERE
+- Bookmark lookup
  +- Table
    | table      sbtest1
    | possible_keys PRIMARY
  +- Index range scan
    key          sbtest1->PRIMARY
    possible_keys PRIMARY
    key_len      4
    rows         49315032
```

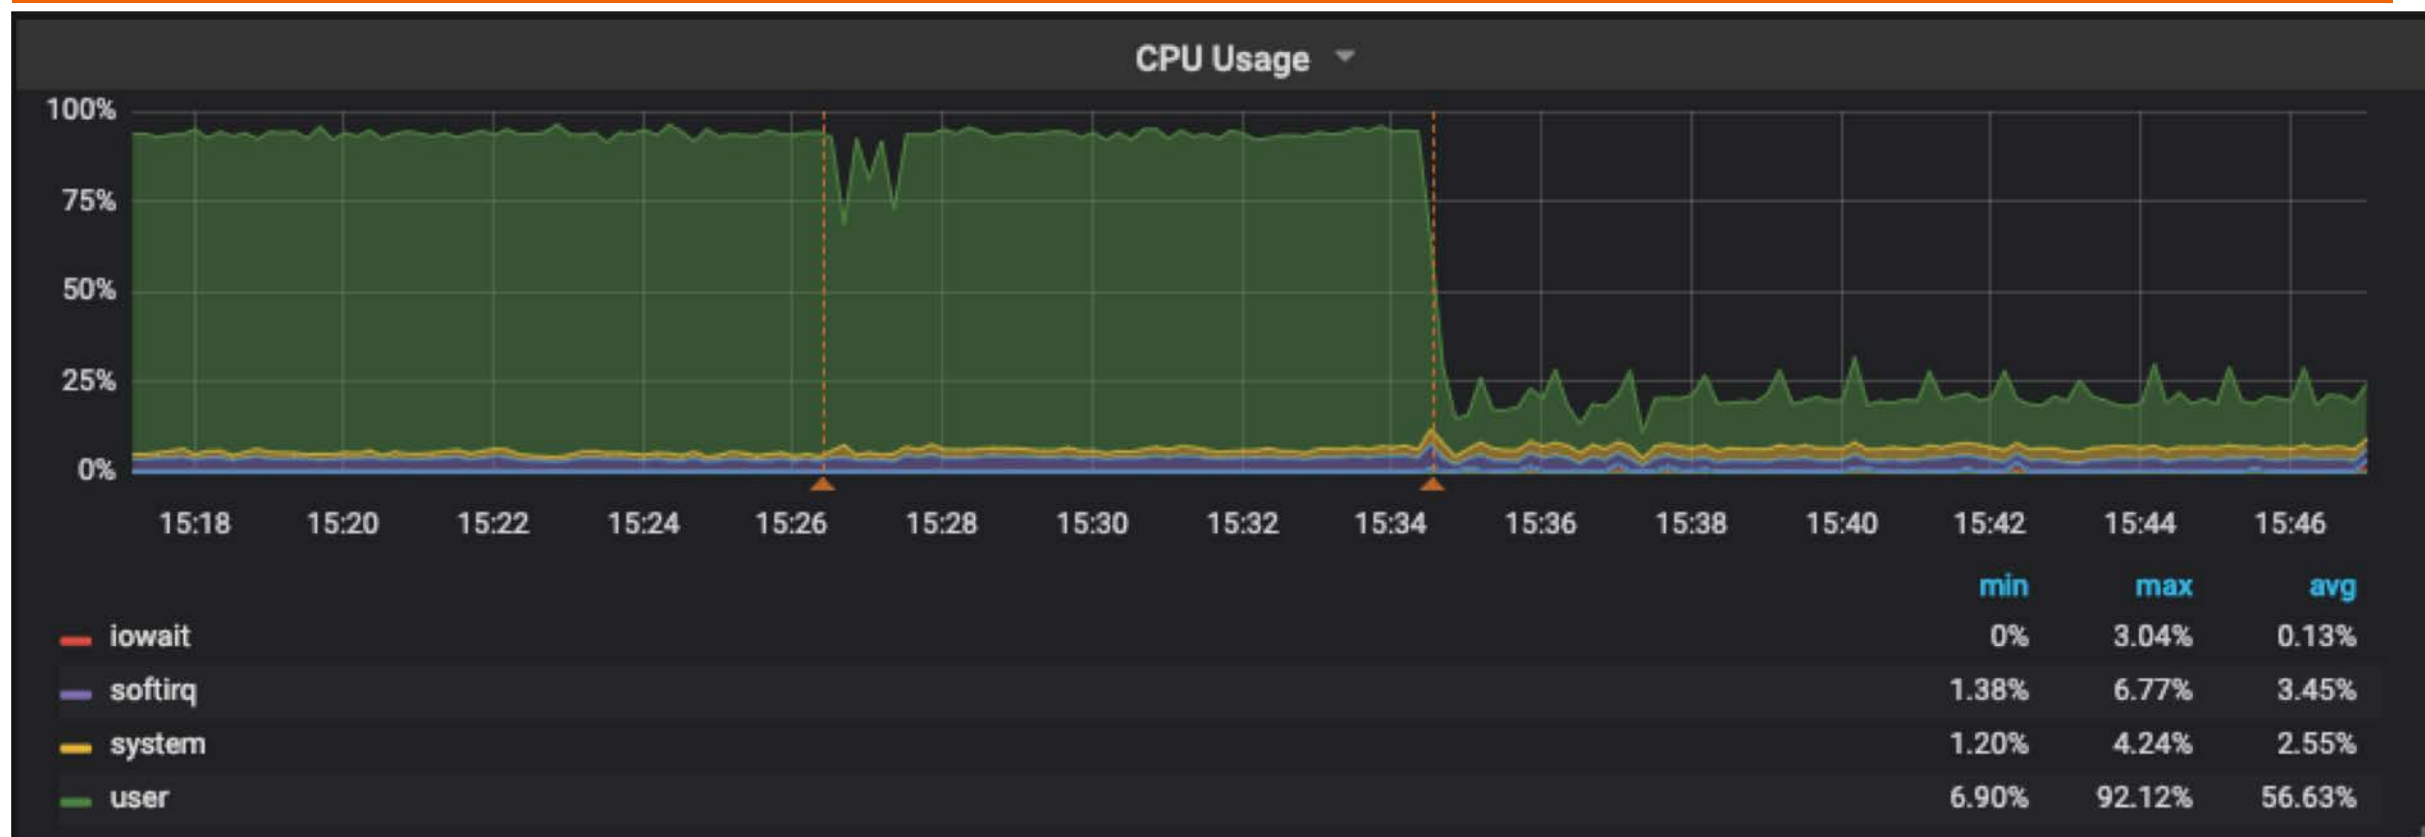
▼ CREATE

```
CREATE TABLE `sbtest1` (
  `id` int(10) unsigned NOT NULL AUTO_INCREMENT,
  `k` int(10) unsigned NOT NULL DEFAULT '0',
  `c` char(120) NOT NULL DEFAULT '',
  `pad` char(60) NOT NULL DEFAULT '',
  PRIMARY KEY (`id`),
  KEY `k_1` (`k`)
) ENGINE=InnoDB AUTO_INCREMENT=100000001 DEFAULT CHARSET=latin1 MAX_ROWS=1000000
```

▼ INDEXES

| KeyName | Type | Unique | Packed | Column | Cardinality |
|---------|-------|--------|--------|--------|-------------|
| PRIMARY | BTREE | Yes | No | id | 98630064 |
| k_1 | BTREE | No | No | k | 31560404 |

CPU Usage Reduction with PMM



<https://per.co.na/PMMCPU>

How to Optimize Them

Query Stats

How Many Rows does it crunch ?

How Many Rows it returns ?

How Much IO is Required ?

Is Temporary Table Required ? Temporary Sort File ?

Run EXPLAIN

Is Plan Reasonable ?

Bad Plan

Missing Indexes

Bad Optimizer Statistics

Bad Query Practices ie “WHERE col+1=10”

Expensive Queries

**Queries
which
Naturally
Require A lot
of Work to
Do**

- `SELECT AVG(value)
FROM ORDERS WHERE
ORDER_DATE < "2018-
01-01"`

Dealing with Expensive Queries

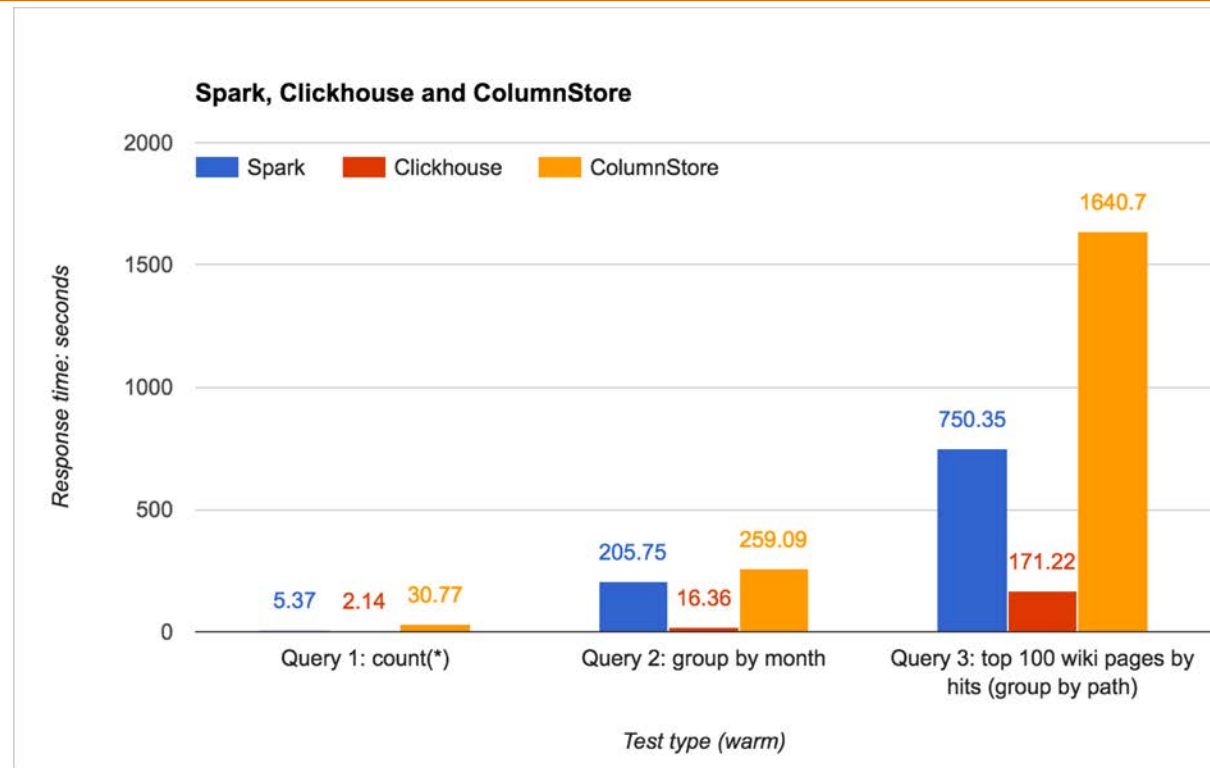
Getting Rid of Them!

Caching Them

Pre-Generating Results

Using Systems which support such queries better

Beyond MySQL for Expensive Queries



<https://per.co.na/jOMbko>

Optimizing Writes

Less Indexes

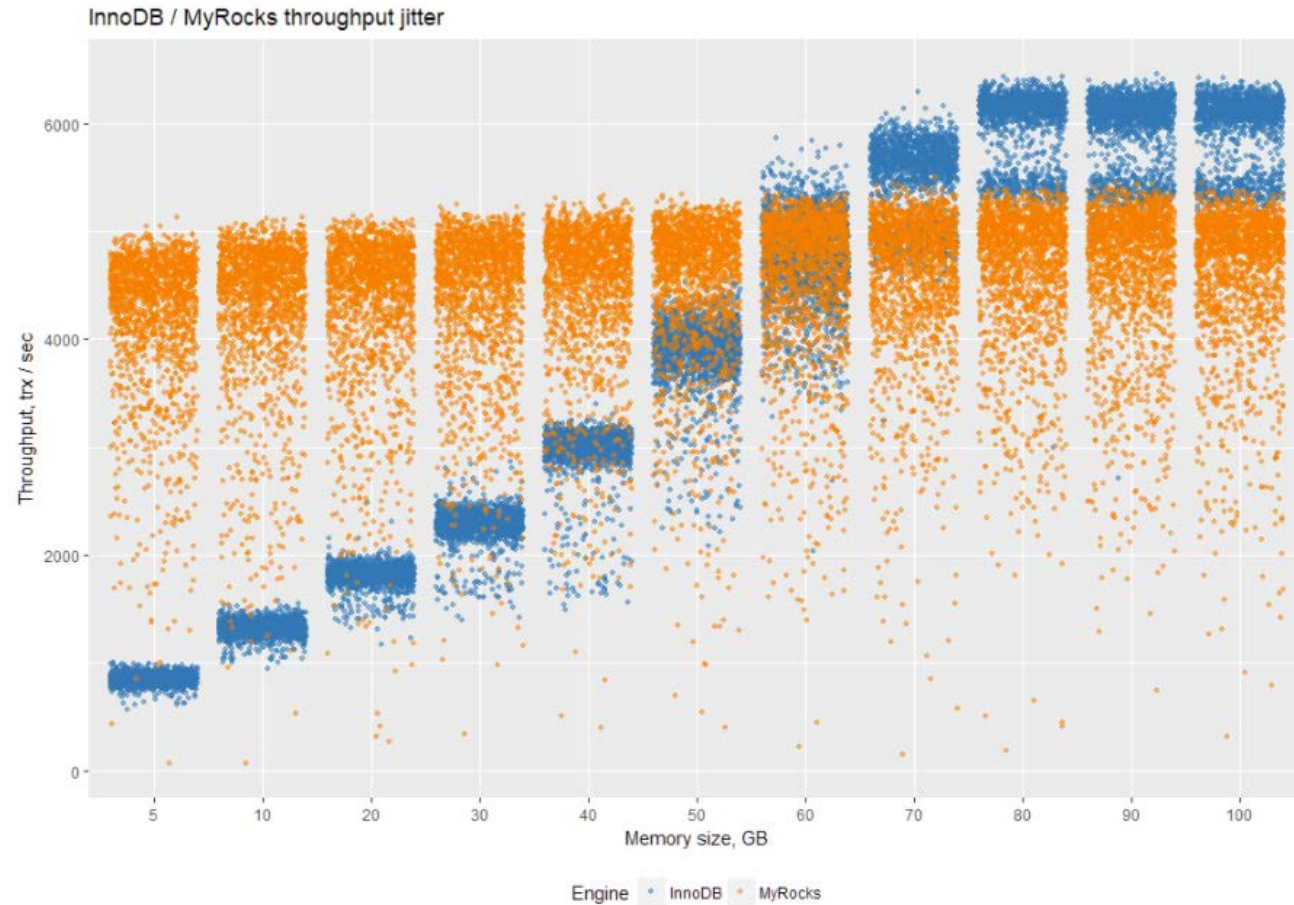
Data Fits in Memory

Batching

Partitioning

Different Storage Engine

MyRocks – Better Performance with Large Data



Share your insight and join the debate!



Percona Open Source Data Management Software Survey

This survey is intended to be of use to the open source community, and we want you all to contribute and share. The more responses we receive, the more useful this will be.

The resulting data will be 100% open and 100% anonymous.

[Take the survey](#)

<https://per.co.na/survey>

Join Us at Percona Live



Percona Live 2019 takes place in Austin, Texas from May 28-30, 2019 at the Hyatt Regency.

Percona Live provides an opportunity to network with peers and technology professionals. Mingle with all types of database community members: DBAs, developers, C-level executives and the latest database technology trend-setters.

SUPER SAVER TICKETS ON SALE!

<https://www.percona.com/live/19/>



Connect. Accelerate. Innovate.

Thank You!
