10x Performance Improvements in 10 steps

A Case Study

Ronald Bradford
http://ronaldbradford.com

FOSDEM - 2010.02

Application

Typical Web 2.0 social media site (Europe based)

- Users Visitors, Free Members, Paying Members
- Friends
- User Content Video, Pictures
- Forums, Chat, Email

Server Environment

- 1 Master Database Server (MySQL 5.0.x)
- 3 Slave Database Servers (MySQL 5.0.x)
- 5 Web Servers (Apache/PHP)
- 1 Static Content Server (Nginx)
- 1 Mail Server

Step I

Monitor, Monitor, Monitor

- What's happened?
- What's happening now?
- What's going to happen?

Past, Present, Future

Monitoring Software

- Installation of Cacti http://www.cacti.net/
- Installation of MySQL Cacti Templates http://code.google.com/p/mysql-cacti-templates/
- (Optional) Installation of MONyog http://www.webyog.com/

Custom Dashboard

- Most important The state of NOW
- Single Page Alerts GREEN YELLOW RED

Database Servers

host - Load Avg (1m*,5m,15m)/ Ping (min/avg*/max/div) / DB Connections (con*,idle*,query*,locked*) / Slave Status (IO,SQL

	Masters	Slaves			
db1	sy:0.87 (22s),0.133, db:[34,25,2,0] (4s),0	, slave1 sy: 1.78 (18s),0.148, db:[11,7,1,0] (1s),0,			
db2	sy:0.00 (20s),0.252, db:[6,1,1,0] (2s),0,	slave2 sy:0.26 (15s),0.090, db:[12,8,1,0] (0s),0,			
dbmoi	n sy:0.33 (11s),0.019, db:[18,16,1,0] (5s),	slave3 sy:1.06 (13s),0.142, db:[14,10,1,0] (6s),0,			

Members Pages

Members Pages from www.example.com - location user (load time, size), non-location user (load time, size) www.example.com m:0.088s, 84K, 2-3-0.0744 0.152s, 76K, 3-3-0.1293 (2s),

Web Servers

```
WWW Server Status (Requests, Requests Per Sec) / index.htm (load time, size) www1 w:, (10s), i:0.089s, 36K, 0-1-0.0666 (10s), www2 w:, (8s), i:0.063s, 36K, 0-1-0.0472 (8s), www3 w:, (6s), i:0.054s, 36K, 0-1-0.0386 (6s), www4 w:, (4s), i:0.049s, 36K, 0-1-0.0357 (4s), www5 w:, (2s), i:0.050s, 36K, 0-1-0.0359 (2s),
```

Dashboard Example

Alerting Software

- Installation of Nagios http://www.nagios.org/
- MONyog also has some DB specific alerts

Application Metrics

Total page generation time

Step 2

Identify problem SQL

Identify SQL Statements

- Slow Query Log
- Processlist
- Binary Log
- Status Statistics

Problems

- Sampling
- Granularity

Solution

tcpdump + mk-query-digest

- Install maatkit http://www.maatkit.org
- Install OS tcpdump (if necessary)
- Get sudo access to tcpdump

http://ronaldbradford.com/blog/take-a-look-at-mk-query-digest-2009-10-08/

#	Rank	Query ID	Response ti	me	Calls	R/Call	Item	
#	====	=======================================	========	=====	======	========	====	
#	1	0xB8CE56EEC1A2FBA0	14.0830	26.8%	78	0.180552	SELECT	c u
#	2	0x195A4D6CB65C4C53	6.7800	12.9%	257	0.026381	SELECT	<u>u</u>
#	3	0xCD107808735A693C	3.7355	7.1%	8	0.466943	SELECT	c u
#	4	0xED55DD72AB650884	3.6225	6.9%	77	0.047046	SELECT	u
#	5	0xE817EFFFF5F6FFFD	3.3616	6.4%	147	0.022868	SELECT	UNION C
#	6	0x15FD03E7DB5F1B75	2.8842	5.5%	2	1.442116	SELECT	c u
#	7	0x83027CD415FADB8B	2.8676	5.5%	70	0.040965	SELECT	c u
#	8	0x1577013C472FD0C6	1.8703	3.6%	61	0.030660	SELECT	С
#	9	0xE565A2ED3959DF4E	1.3962	2.7%	5	0.279241	SELECT	c t u
#	10	0xE15AE2542D98CE76	1.3638	2.6%	6	0.227306	SELECT	С
#	11	0x8A94BB83CB730494	1.2523	2.4%	148	0.008461	SELECT	hv u
#	12	0x959C3B3A967928A6	1.1663	2.2%	5	0.233261	SELECT	c t u
#	13	0xBC6E3F701328E95E	1.1122	2.1%	4	0.278044	SELECT	c t u

```
# Query 2: 4.94 QPS, 0.13x concurrency, ID 0x195A4D6CB65C4C53 at byte 4851683
 This item is included in the report because it matches --limit.
                    min
                              max avg 95% stddev median
           pct total
           3 257
# Count
 Exec time 10 7s
                       35us 492ms 26ms 189ms 78ms
                                                      332us
 Time range 2009-10-16 11:48:55.896978 to 2009-10-16 11:49:47.760802
# bytes
             2 10.75k 41
                               43 42.85 42.48 0.67 42.48
 Errors
                     none
Rows affe 0
 Warning c
 Query time distribution
  1us
 10us #
 1ms ####
 10ms ###
 100ms ########
  1s
 10s+
 Tables
   SHOW TABLE STATUS LIKE 'u'\G
   SHOW CREATE TABLE `u`\G
# EXPLAIN
SELECT ... FROM u ...\G
```

- Wrappers to capture SQL
- Re-run on single/multiple servers
 - e.g. Different slave configurations



- Enable General Query Log in Development/Testing
- Great for testing Batch Jobs

Application Logic

Action 3

- Show total master/slave SQL statements executed
- Show all SQL with execution time (admin user only)

Have abstracted class/method to execute ALL SQL

Step 3

Analyze problem SQL

- Query Execution Plan (QEP)
 - EXPLAIN [EXTENDED] SELECT ...
- Table/Index Structure
 - SHOW CREATE TABLE <tablename>
- Table Statistics
 - SHOW TABLE STATUS <tablename>

Good

```
mysql> EXPLAIN SELECT id FROM example table WHERE id=1\G
*********************** 1. row ********************
           id: 1
  select type: SIMPLE
        table: example table
         type: const
possible keys: PRIMARY
          key: PRIMARY
      key len: 4
         ref: const
         rows: 1
        Extra: Using index
```

Bad

```
mysql> EXPLAIN SELECT * FROM example table \G
********************** 1. row *****************
           id: 1
  select type: SIMPLE
        table: example table
         type: ALL
possible keys: NULL
          key: NULL
      key len: NULL
          ref: NULL
         <u>rows: 59</u>
        Extra:
```

Tip

- SQL Commenting
 - Identify batch statement SQL
 - Identify cached SQL

```
SELECT /* Cache: 10m */ ....

SELECT /* Batch: EOD report */ ...

SELECT /* Func: 123 */ ....
```

Step 4

The Art of Indexes

- Different Types
 - Column
 - Concatenated
 - Covering
 - Partial

http://ronaldbradford.com/blog/understanding-different-mysql-index-implementations-2009-07-22/

- EXPLAIN Output
 - Possible keys
 - Key used
 - Key length
 - Using Index

Tip

- Generally only 1 index used per table
- Make column NOT NULL when possible
- Statistics affects indexes
- Storage engines affect operations

Before (7.88 seconds)

After (0.04 seconds)

```
****** 2. row **
                                   ****** 2. row ***
         id: 2
                                             id: 2
 select type: DEPENDENT SUBQUERY
                                     select type: DEPENDENT SUBQUERY
       table: h p
                                          table: h p
        type: ALL
                                           type: index subquery
possible keys: NULL
                                   possible keys: UId
        key: NULL
                                            key: UId
     key len: NULL
                                        key len: 4
        ref: NULL
                                            ref: func
        rows: 33789
                                           rows: 2
       Extra: Using where
                                          Extra: Using index
```

ALTER TABLE h_p ADD INDEX (UId);

ALTER TABLE f DROP INDEX UID, ADD INDEX (UID, FUID)

Indexes can hurt performance

Step 5

Offloading Master Load

5. Offloading Master Load

- Identify statements for READ ONLY slave(s)
 - e.g. Long running batch statements

Single point v scalable solution

Step 6

Improving SQL

6. Improving SQL

- Poor SQL Examples
 - ORDER BY RAND()
 - SELECT *
 - Lookup joins
 - ORDER BY

The database is best for storing and retrieving data not logic

Step 7

Storage Engines

- MyISAM is default
- Table level locking
 - Concurrent SELECT statements
 - INSERT/UPDATE/DELETE blocked by long running SELECT
 - All SELECT's blocked by INSERT/UPDATE/DELETE
- Supports FULLTEXT

- InnoDB supports transactions
- Row level locking with MVCC
- Does not support FULLTEXT
- Different memory management
- Different system variables

- There are other storage engines
 - Memory
 - Archive
 - Blackhole
 - Third party

Using Multiple Engines

- Different memory management
- Different system variables
- Different monitoring
- Affects backup strategy

- Configure InnoDB correctly
 - innodb buffer pool size
 - innodb log file size
 - innodb_flush_log_at_trx_commit

- Converted the two primary tables
 - Users
 - Content

Locking eliminated

Step 8

Caching

8. Caching

- Memcache is your friend http://memcached.org/
 - Cache query results
 - Cache lookup data (eliminate joins)
 - Cache aggregated per user information
- Caching Page Content
 - Top rated (e.g. for 5 minutes)

8. Caching

- MySQL has a Query Cache
 - Determine the real benefit
 - Turn on or off dynamically
 - SET GLOBAL query cache size = 1024*1024*32;

8. Caching



The best performance improvement for an SQL statement is to eliminate it.

Step 9

Sharding

9. Sharding

- Application level horizontal and vertical partitioning
- Vertical Partitioning
 - Grouping like structures together (e.g. logging, forums)
- Horizontal Partitioning
 - Affecting a smaller set of users (i.e. not 100%)

9. Sharding

Separate Logging

Reduced replication load on primary server

Step 10

Database Management

10. Database Management

Database Maintenance

- Adding indexes (e.g. ALTER)
- OPTIMIZE TABLE
- Archive/purging data (e.g DELETE)

Blocking Operations

10. Database Maintenance

- Automate slave inclusion/exclusion
- Ability to apply DB changes to slaves
- Master still a problem

10. Database Maintenance

- Install Fail-Over Master Server
 - Slave + Master features
 - Master extra configuration
- Scripts to switch slaves
- Scripts to enable/disable Master(s)
- Scripts to change application connection

10. Database Maintenance

Higher Availability &

Testing Disaster Recovery



Front End Improvements

11. Front End Improvements

- Know your total website load time http://getfirebug.com/
 - How much time is actually database related?
- Reduce HTML page size 15% improvement
 - Remove full URL's, inline css styles
- Reduce/combine css & js files
- Identify blocking elements (e.g. js)

11. Front End Improvements

- Split static content to different ServerName
- Spread static content over multiple ServerNames (e.g. 3)
- Sprites Combining lightweight images http://spriteme.org/
- Cookie-less domain name for static content

Conclusion

Before

- Users experienced slow or unreliable load times
- Management could observe, but no quantifiable details
- Concern over load for increased growth
- Release of some new features on hold

Now

- Users experienced consistent load times (~60ms)
 - Quantifiable and visible real-time results
- Far greater load now supported (Clients + DB)
- Better testability and verification for scaling
- New features can be deployed

Consulting Available Now

http://ronaldbradford.com